

Topics on econometrics and causal inference

Xiang Ao

2026-08-10

Table of contents

Preface	1
How to cite	1
 I. Regression Foundations	 3
1. Interpreting interaction in a regression model	5
1.1. Interaction with two binary variables	5
1.2. Interaction with continuous variables	8
2. Interaction term in a non-linear model	13
2.1. Interaction term in a non-linear model	13
2.2. A worked example	15
3. Chow test and more	17
3.1. Chow test	17
3.2. A comparison with two different outcomes	20
3.3. Overlapping samples	25
3.4. IV regression	28
3.5. IV with fixed effects	30
3.6. Chow test with different covariates	40
3.7. Conclusion	44
4. Weights in OLS in causal inference	47
4.1. Unit level weights for OLS	47
4.2. Strata level weights	49
4.3. Recommended estimation approaches	49
4.4. A worked example	50
 II. Marginal Effects & Fixed Effects	 53
5. Marginal effects in models with fixed effects	55
5.1. Marginal effects in a linear model	55
5.2. Marginal effects in a non-linear model	57
6. Comparing Marginal effects with margins command	67
6.1. Comparing Marginal effects	67
7. Fixed or Random Effect, or Both?	77
7.1. Panel data	77

Table of contents

7.2. Time-invariant variables	77
7.3. Between-within model	78
7.4. BW model in R	78
7.5. BW model in Stata	81
7.6. BW model in non-linear models	83
8. Mundlak device in fixed effect models	85
8.1. TWFE	85
8.2. TWM	85
8.3. An important qualification: TWM needs a balanced panel	89
8.4. More than two fixed effects: one set of means, not one per dimension	91
8.5. Applications of TWM to DID in common treatment timing	94
9. Correlated Random Effect	119
9.1. Panel data	119
9.2. Fixed effect	119
9.3. Random effect	119
9.4. Correlated random effect	120
9.5. Example	120
10. More on Correlated Random Effect (CRE)	125
10.1. linear model	125
10.2. nonlinear model	129
10.3. References	133
III. Treatment Effects & Matching	135
11. Treatment effects and matching	137
11.1. Treatment effects in observational studies	137
11.2. Matching	150
11.3. Modern matching in R: balance diagnostics and entropy weighting	155
12. When is OLS coefficient the ATE	161
12.1. OLS cannot be ATE (most of the time)	161
12.2. Regression adjustment	164
13. Matching and Weighting Part 1: matching	167
13.1. Assumptions	167
13.2. Problem	167
13.3. Distance	167
13.4. matching methods	168
13.5. Example	169
13.6. Estimation	185
14. Matching and Weighting - Part 2: weighting	187
14.1. Assumptions	187
14.2. Problem	187
14.3. weighting methods	187

14.4. Example	189
14.5. comparing matching and weighting	194
15. Sensitivity analysis of OLS with sensemakr	195
15.1. OVB	195
15.2. Partial R^2 of treatment on outcome	195
15.3. Robust value	196
15.4. Bounds using observed covariate(s)	196
15.5. Example	196
15.6. sensitivity analysis for IV	200
16. Gender Wage Gap Analysis Using US Current Population Survey Data	205
16.1. Oaxaca-Blinder decomposition	205
16.2. Example	206
16.3. Beyond the mean: where in the distribution is the gap?	232
16.4. What is any of this measuring?	235
IV. Panel Data & DiD	239
17. Causal Forest in panel data	241
17.1. Introduction	241
17.2. Linear effect	241
17.3. Nonlinear effect case 1	247
17.4. Nonlinear effect case 2	253
17.5. Summary	259
18. Synthetic Control in R and Stata	261
18.1. Synthetic control	261
18.2. inference	261
18.3. an example	262
18.4. Generalized Synthetic Control	268
19. Same Data, Different Estimators: A Synthetic Control Comparison	275
19.1. One panel, four estimators	275
19.2. Four estimators, one panel	275
19.3. What changes across estimators	276
19.4. The figure	276
19.5. What the comparison teaches	277
19.6. Setup code	278
20. Bartik Instrument and Rotemberg weights	279
20.1. Bartik Instrument	279
20.2. Bartik Instrument in special cases	280
20.3. Rotemberg weights	280
20.4. Simulations	281
20.5. Why is good to have Rotemberg weights?	287
20.6. What if you are using a reduced form?	287

21. Extended TWFE and other DiD estimators	291
21.1. Staggered DiD without treatment reversal	291
21.2. Callaway and Sant’Anna	297
21.3. Staggered DiD with treatment reversal	309
22. Causal Panel (why DDDiD)	311
22.1. Why DDDiD	311
22.2. Synthetic Control	311
22.3. Synthetic DiD	311
22.4. Example from synthdid	312
22.5. Summary	314
V. Count Data & Specialized Models	315
23. Which count data model to use	317
23.1. A comparison of various count data models with extra zeros	317
23.2. Quick worked example	320
23.3. Mixed count models with glmmTMB	322
24. What model to use for rare events	323
24.1. Introduction	323
24.2. Simulation	324
25. Poisson model with exposure	327
25.1. Poisson with exposure or Poisson rate	327
25.2. Simulation	328
25.3. Fitting the three models	328
25.4. When $\gamma \neq 1$	329
25.5. Stata equivalent	330
25.6. Which model to use?	331
26. IV in Fixed effect Poisson with many fixed effects	333
26.1. Poisson regression with many fixed effects	333
26.2. IV poisson with many fixed effects	333
26.3. Example	335
26.4. Bootstrapping standard errors	339
26.5. How to do it with R	341
26.6. What if the endogenous variable is binary?	341
VI. Causal Inference Methods	345
27. Using machine learning for causal effect in observational study	347
27.1. A simulation for an OLS model	347
28. Mediation analysis in R and Stata	353
28.1. Mediation analysis	353
28.2. Causal Mediation analysis	356

28.3. Binary outcome	359
28.4. Going further: formal causal mediation	360
29. Endogeneity with censored variable	361
29.1. Bunching for the outcome variable	361
29.2. Bunching for the treatment variable	361
29.3. Simulation	363
29.4. Conclusion	366
30. G-computation (g-formula) with time varying covariates	367
30.1. The problem	367
30.2. Assumptions	368
30.3. Identifiability	369
30.4. Estimation	369
30.5. Parametric g-formula	369
30.6. an example	370
30.7. A second example	376
31. Policy learning by policytree	381
31.1. Optimal policy	381
31.2. IPW loss	382
31.3. AIPW loss	382
31.4. Policytree	382
32. Recent causal inference tools	387
32.1. Identification	387
32.2. Nonparametric estimation in a nutshell	387
32.3. Simulation with a known treatment effect	389
32.4. EAGeR trial simulation (AIPW package)	396
33. Intro to proximal causal inference	403
33.1. Unobserved confounder	403
33.2. Proximal causal inference	403
33.3. example	407
34. Simulation on proximal causal inference with panel data	411
34.1. A simulation with panel data	412
35. longitudinal modified treatment policy (LMTP)	417
35.1. Definitions and assumptions	418
35.2. Using lmtip package	424
35.3. another example	426
VII. Advanced Topics	429
36. Doing the same multi-level model with R and stata	431
36.1. Stata and R syntax comparison for the same three level model	431
36.2. A growth model in R and stata	435

36.3. cem in R and stata	440
37. Modeling variation, not just the mean, in Likert-scale data	443
37.1. Modeling the spread, not the mean	443
37.2. The ceiling is exact, not empirical	443
37.3. Dividing by the mean does not fix it	445
37.4. The right idea: spread after removing the ceiling effect	445
37.5. Two models: one with the boundary, one without	446
37.6. Worked example: same dispersion, different means	448
37.7. Takeaways	454
37.8. References	455
38. Conjoint design and analysis using R	457
38.1. Conjoint Analysis	457
38.2. Example	458
38.3. Conjoint design of experiment	459
38.4. Conjoint analysis	462
38.5. Causal effect in conjoint analysis	465
39. Stata 19 CATE command	467
39.1. Stata 19 new CATE command	467
39.2. AIPW	475
40. Spatial econometrics introduction	479
40.1. Why do we need spatial econometrics?	479
40.2. A general linear spatial model	479
40.3. Estimation with S2SLS for a less general model	480
40.4. Estimation with a more general model	480
40.5. an example	481
41. Causal simulation	485
41.1. examples	486
41.2. a more complicated model	490
42. Frengression and Engression	495
42.1. Engression	495
42.2. From Engression to Frengression	497
42.3. The Architecture	497
42.4. Using Rfrengression	497
42.5. Frengression as DGP	499
42.6. Real Data: LaLonde	500
42.7. Comparison: causl vs frengression	501
42.8. Benchmarking: Frengression vs Semiparametric Methods	502
42.9. Summary	512
43. Partial Interference	513
43.1. The Setup	513
43.2. Estimation	514
43.3. Using CausalModel	515

43.4. Monte Carlo Validation	519
43.5. What If You Ignore Interference?	521
43.6. Summary	523
44. Power analysis with list experiment	525
44.1. List experiment	525
44.2. Power analysis	525
45. Causal mediation analysis	529
45.1. Traditional mediation analysis	529
45.2. Causal mediation analysis	529
45.3. Summary	537
46. Uplift Modeling and CATE	539
46.1. Python: causalmml uplift trees	539
46.2. R: causal forest for CATE	541
46.3. Uplift vs CATE: the same thing with different vocabulary	543
47. Using Numpyro	545
47.1. Example 1	545
47.2. Example 2: Numpyro hierarchical model	547
47.3. Comparing the two approaches	549
48. Pre-registration, TOST, and why Bayes doesn't let you skip the hard part	551
48.1. Why this came up	551
48.2. What pre-registration is actually for	551
48.3. A non-significant result is not “no effect”	552
48.4. Equivalence testing (TOST), plainly	552
48.5. The Bayesian alternative	554
48.6. Bayes doesn't let you skip the hard part	556
48.7. So how do you actually justify Δ ?	556
48.8. “Why not just report the CI and let the reader decide?”	557
48.9. What I'd actually do	558

Preface

This is a working notebook rather than a textbook. The chapters are posts I wrote while reading, teaching, and consulting on applied econometrics, often because a colleague or a student asked a question that did not have a tidy answer in the books I had. Each one stands on its own. The order did not follow a plan, and the same topic sometimes appears more than once as my understanding changed.

The common question is the one Angrist and Pischke emphasize: what is being identified, by what variation, under what assumptions. The posts ask that question in different settings: interactions and marginal effects, fixed and correlated random effects, matching and weights, difference-in-differences, synthetic control, count models, TMLE and AIPW, mediation, proximal causal inference, partial interference, conjoint and list experiments, uplift modeling, and policy trees.

Most examples are in R. I use Stata where it is more natural, and sometimes Python or Julia when the package is better there. The two companion books, [Introduction to Causal Econometrics](#) and [Causal Econometrics with Julia](#), organize the same material more systematically. This collection is the messier version: useful when I want to see how a specific method behaves on a specific problem.

How to cite

Ao, Xiang. *Topics on Econometrics and Causal Inference*. https://xiangao.github.io/blog_book/.

```
@book{ao_topics_econometrics,  
  author = {Ao, Xiang},  
  title  = {Topics on Econometrics and Causal Inference},  
  year   = {2026},  
  url    = {https://xiangao.github.io/blog_book/}  
}
```


Part I.

Regression Foundations

1. Interpreting interaction in a regression model

1.1. Interaction with two binary variables

In a regression model with interaction term, people tend to pay attention to only the coefficient of the interaction term.

Let's start with the simplest situation: x_1 and x_2 are binary and coded 0/1.

$$E(y) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 x_2 \quad (1.1)$$

In this case, we have a saturated model; that is, we have three coefficients representing additive effects from the baseline situation (both x_1 and x_2 being 0). There are four different situations, with four combinations of x_1 and x_2 .

A lot of people just pay attention to the interaction term. In the case of studying treatment effects between two groups, say female and male, that makes sense, the interaction term representing the difference between male and female in terms of treatment effect.

In this model:

$$E(y) = \beta_0 + \beta_1 \text{female} + \beta_2 \text{treatment} + \beta_{12} \text{female} * \text{treatment} \quad (1.2)$$

The two dummy-coded binary variables, female and treatment, form four combinations. The following 2x2 table represents the expected means of the four cells(combinations).

	male		female	
control	β_0	(1.2)	$\beta_0 + \beta_1$	(1.2)
treatment	$\beta_0 + \beta_2$	(1.2)	$\beta_0 + \beta_1 + \beta_2 + \beta_{12}$	(1.2)

We can see from this table that, for example,

$$\beta_0 = E(Y|(0,0))$$

1. Interpreting interaction in a regression model

{#eq-interaction-regression-7};

that is, β_0 is the expected mean of the cell (0,0) (male and control).

$$\beta_0 + \beta_1 = E(Y|(1,0))$$

{#eq-interaction-regression-8};

that is, $\beta_0 + \beta_1$ is the expected mean of the cell (1,0) (female and control). And so on.

Now,

$$\beta_{12} = (E(Y|(1,1)) - E(Y|(0,1))) - (E(Y|(1,0)) - E(Y|(0,0))) \quad (1.3)$$

that is, the coefficient on the interaction term is actually the difference in difference. That's why in many situations, people are only interested in the interaction coefficient, since they are only interested in the diff-in-diff estimates. The usual diff-in-diff estimator in the causal inference literature refers to something similar, instead of female vs. male, people are interested in the treatment effect difference in before and after treatment. If we simply replace female/male dummy with before/after dummy, we can use the same logic. In those situations, it's fine to mainly focus on the interaction term coefficient.

In some other situations, the three coefficients are equally important. It depends on your interest. For example, if we are interested in studying differences between union member and non-union member and black vs. non-black, we may not be only interested in the interaction effect. Instead, we might be interested in all four cells, maybe all possible pairwise comparisons. In that case, we should pay attention to all three coefficients. Stata's "margins" command is of great help if we'd like to compare the cell means.

Let's take a look from a sample example in Stata:

```
webuse union3
reg ln_wage i.union##i.black, r
margins union#black
margins union#black, pwcompare
```

```
. webuse union3
(NLS Women 14-24 in 1968)
```

```
. reg ln_wage i.union##i.black, r
```

Linear regression	Number of obs	=	1,244
	F(3, 1240)	=	34.76
	Prob > F	=	0.0000
	R-squared	=	0.0762
	Root MSE	=	.37699

1.1. Interaction with two binary variables

		Robust				
ln_wage	Coefficient	std. err.	t	P> t	[95% conf. interval]	
1.union	.2045053	.0291682	7.01	0.000	.1472808	.2617298
1.black	-.1709034	.0308067	-5.55	0.000	-.2313425	-.1104644
union#black						
1 1	.0386275	.0516609	0.75	0.455	-.062725	.13998
_cons	1.657525	.0138278	119.87	0.000	1.630396	1.684653

. margins union#black

Adjusted predictions Number of obs = 1,244
Model VCE: Robust

Expression: Linear prediction, predict()

		Delta-method				
	Margin	std. err.	t	P> t	[95% conf. interval]	
union#black						
0 0	1.657525	.0138278	119.87	0.000	1.630396	1.684653
0 1	1.486621	.027529	54.00	0.000	1.432613	1.54063
1 0	1.86203	.0256822	72.50	0.000	1.811644	1.912415
1 1	1.729754	.0325611	53.12	0.000	1.665873	1.793635

. margins union#black, pwcompare

Pairwise comparisons of adjusted predictions Number of obs = 1,244
Model VCE: Robust

Expression: Linear prediction, predict()

		Delta-method	Unadjusted	
	Contrast	std. err.	[95% conf. interval]	
union#black				
(0 1) vs (0 0)	-.1709034	.0308067	-.2313425	-.1104644
(1 0) vs (0 0)	.2045053	.0291682	.1472808	.2617298
(1 1) vs (0 0)	.0722294	.0353756	.0028268	.141632
(1 0) vs (0 1)	.3754087	.0376487	.3015466	.4492709
(1 1) vs (0 1)	.2431328	.0426388	.1594807	.326785
(1 1) vs (1 0)	-.1322759	.0414705	-.2136359	-.0509159

What we get by using “margins union#black” is the four cell means of $E(Y)$, in this case, log of wage. Then “margins union#black, pwcompare” tells us all pairwise comparison of these four cell means. Instead of only paying attention to the interaction coefficient, in this case we might be interested in some comparisons of the four different situations of union and black. In this example, all six pairwise comparisons of the cell means have 95% confidence intervals that do not include zero — and yet the interaction term is insignificant ($p = 0.455$). The two facts are compatible, and seeing why is the point. Each pairwise comparison asks whether two cells differ *from each other*. The interaction asks whether the *difference between two of those differences* is non-zero. Every simple contrast can be large and precisely estimated while the difference of differences is small relative to its own standard error, which is exactly what happens here: the union premium is clearly positive for both groups, but the *gap* between the two premiums (0.039) is not distinguishable from zero.

1.2. Interaction with continuous variables

Let’s start with the simplest situation: x_1 and x_2 are continuous.

$$E(y) = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 * x_2 \quad (1.4)$$

In this case, we recommend “centering” x_1 and x_2 if they are continuous; that is, subtracting the mean value from each continuous independent variable when they are involved in the interaction term. Centering is a linear transformation, so it does not change the model’s fitted values, R^2 , or the coefficient, standard error, t-statistic, or p-value of the interaction term β_{12} itself. What it does change is the interpretation of the main effects:

1. To reduce apparent multi-collinearity between the main effects and the interaction term. If the range of x_1 and x_2 include only positive numbers, $x_1 * x_2$ can be highly correlated with x_1 and/or x_2 , which inflates the standard errors of β_1 and β_2 (but not β_{12}) and makes those two coefficients numerically unstable to estimate, even though the model’s predictions are unaffected.

“Centering” reduces the correlation between the interaction term and the main-effect variables. If the original variables are normally distributed, the interaction term after centering is actually uncorrelated with the (centered) main effects. When they are not normally distributed, centering will still reduce the correlation to a large degree. This is an interpretational and numerical-stability convenience for β_1 and β_2 , not a fix that improves estimation of the interaction effect itself.

2. To help with interpretation. In a model with interaction, β_1 represents the effect of x_1 when x_2 is zero. However, in many situations, zero is not within the range of x_2 . After centering, centered x_2 at zero simply means original x_2 at its mean value, so β_1 is now the effect of x_1 evaluated at the mean of x_2 instead of at $x_2 = 0$.

When we have dummy variable interacting with continuous variable, only continuous variable should be centered.

Again, Stata's margins command is helpful.

```
sysuse auto
sum mpg
gen mpg_centered=mpg-r(mean)
sum mpg_centered
reg price i.foreign##c.mpg_centered
margins foreign, at(mpg_centered=(-3 (1) 3))
marginsplot
```

```
. sysuse auto
(1978 automobile data)
```

```
. sum mpg
```

Variable	Obs	Mean	Std. dev.	Min	Max
mpg	74	21.2973	5.785503	12	41

```
. gen mpg_centered=mpg-r(mean)
```

```
. sum mpg_centered
```

Variable	Obs	Mean	Std. dev.	Min	Max
mpg_centered	74	-4.03e-08	5.785503	-9.297297	19.7027

```
. reg price i.foreign##c.mpg_centered
```

Source	SS	df	MS	Number of obs	=	74
Model	183435285	3	61145094.9	F(3, 70)	=	9.48
Residual	451630112	70	6451858.74	Prob > F	=	0.0000
Total	635065396	73	8699525.97	R-squared	=	0.2888
				Adj R-squared	=	0.2584
				Root MSE	=	2540.1

price	Coefficient	Std. err.	t	P> t	[95% conf. interval]
foreign					
Foreign	1666.519	717.217	2.32	0.023	236.0751 3096.963
mpg_centered	-329.2551	74.98545	-4.39	0.000	-478.8088 -179.7013
foreign#					

1. Interpreting interaction in a regression model

c.							
mpg_centered							
Foreign		78.88826	112.4812	0.70	0.485	-145.4485	303.225
_cons		5588.295	369.0945	15.14	0.000	4852.159	6324.431

```
. margins foreign, at(mpg_centered=(-3 (1) 3))
```

Adjusted predictions

Number of obs = 74

Model VCE: OLS

Expression: Linear prediction, predict()

1._at: mpg_centered = -3

2._at: mpg_centered = -2

3._at: mpg_centered = -1

4._at: mpg_centered = 0

5._at: mpg_centered = 1

6._at: mpg_centered = 2

7._at: mpg_centered = 3

		Delta-method					
		Margin	std. err.	t	P> t	[95% conf. interval]	
_at#foreign							
1#Domestic		6576.06	370.446	17.75	0.000	5837.229	7314.891
1#Foreign		8005.915	766.8178	10.44	0.000	6476.545	9535.284
2#Domestic		6246.805	354.4734	17.62	0.000	5539.83	6953.78
2#Foreign		7755.548	709.9327	10.92	0.000	6339.632	9171.464
3#Domestic		5917.55	354.0032	16.72	0.000	5211.513	6623.587
3#Foreign		7505.181	658.8306	11.39	0.000	6191.185	8819.177
4#Domestic		5588.295	369.0945	15.14	0.000	4852.159	6324.431
4#Foreign		7254.814	614.9548	11.80	0.000	6028.325	8481.303
5#Domestic		5259.04	397.981	13.21	0.000	4465.292	6052.788
5#Foreign		7004.447	579.9479	12.08	0.000	5847.778	8161.117
6#Domestic		4929.785	437.9413	11.26	0.000	4056.338	5803.231
6#Foreign		6754.081	555.4891	12.16	0.000	5646.192	7861.969
7#Domestic		4600.53	486.253	9.46	0.000	3630.729	5570.331
7#Foreign		6503.714	543.0057	11.98	0.000	5420.723	7586.704

```
. marginsplot
```

Variables that uniquely identify margins: mpg_centered foreign

.

1.2. Interaction with continuous variables

In this example, the graph shows the predicted price for foreign and domestic cars at different level of mpg.

2. Interaction term in a non-linear model

2.1. Interaction term in a non-linear model

In a non-linear model (for example, logit or poisson model), the interpretation of the coefficient on the interaction term is tricky. [Ai and Norton \(2003\)](#) points out that the interaction term coefficient is not the same as people can interpret as in a linear model; that is, how much effect of x_1 changes with the value of x_2 . They interpret this as a cross partial derivative of the nonlinear response function $E(y)$ with respect to x_1 and x_2 — a quantity that involves the model's other coefficients and the values of all covariates, not just the estimated interaction coefficient β_{12} . As a result, this cross partial derivative can differ in magnitude, sign, and even statistical significance from β_{12} itself.

If we have a linear model with interaction:

$$E(y) = \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 * x_2 \quad (2.1)$$

Then, the marginal effect

$$\frac{\partial^2 E(y)}{\partial x_1 \partial x_2} = \beta_{12} \quad (2.2)$$

That is, β_{12} is the second derivative of $E(y)$ on x_1 and x_2 . The marginal effect of x_1 is $\frac{\partial E(y)}{\partial x_1} = \beta_1 + \beta_{12} x_2$, which is linear in x_2 but does not itself depend on any nonlinear transformation, since the model is linear.

In a non-linear model,

$$F(E(y)) = \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 * x_2 \quad (2.3)$$

$$\frac{\partial^2 F(E(y))}{\partial x_1 \partial x_2} = \beta_{12} \quad (2.4)$$

Here, the partial derivative of $F(E(y))$ on x_1 and x_2 is still β_{12} . However, most people are interested in $\frac{\partial^2 E(y)}{\partial x_1 \partial x_2}$.

$$\frac{\partial^2 E(y)}{\partial x_1 \partial x_2} = \beta_{12} G'(\eta) + (\beta_1 + \beta_{12} x_2)(\beta_2 + \beta_{12} x_1) G''(\eta) \quad (2.5)$$

where $G()$ is the inverse function of $F()$ and η is the linear predictor.

2. Interaction term in a non-linear model

It is true that in a non-linear model with interaction, the marginal effect of x_1 differs with different values of x_2 . However, even if we have a non-linear model without interaction, the marginal effect of x_1 is still different with different values of x_2 . To see this, consider a nonlinear model with no interaction term at all:

$$F(E(y)) = \beta_1 x_1 + \beta_2 x_2 \quad (2.6)$$

$$\frac{\partial^2 E(y)}{\partial x_1 \partial x_2} = (\beta_1 \beta_2) G''(\eta) \quad (2.7)$$

Therefore, when we set up our model,

$$F(E(y)) = \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 * x_2 \quad (2.8)$$

we have in mind that we allow interaction of x_1 and x_2 to interact for the effect on $F(E(y))$; not on $E(y)$.

We agree with Bill Greene (2013), “Interaction Effects in Models for Nonlinear Outcomes” (Lugano lecture notes; the copy formerly at people.stern.nyu.edu/wgreene/Lugano2013/ is no longer online, but the argument appears in Greene’s *Econometric Analysis*, 8th ed., sec. 17.3). In a nonlinear model, the partial effects (as Greene calls it) is nonlinear, regardless of the model. For example, in a logit model, even if you don’t have an interaction term in your model, the effect of x_1 will still be different for every value of x_2 , simply because it’s a nonlinear model.

As Greene put it at the summary section, “Build the model based on appropriate statistical procedures and principles. Statistical testing about the model specification is done at this step Hypothesis tests are about model coefficients and about the structural aspects of the model specifications. Partial effects are neither coefficients nor elements of the specification of the model. They are implications of the specified and estimated model.”

We also agree with [Maarten Buis 2010](#), that we should use multiplicative effect in a non-linear model. That is, in a non-linear model,

$$F(E(y)) = \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 * x_2 \quad (2.9)$$

We should pay more attention to

$$\frac{\partial^2 F(E(y))}{\partial x_1 \partial x_2} = \beta_{12} \quad (2.10)$$

For example, in a logit model,

$$\log(P(y = 1)/(1 - P(y = 1))) = \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 * x_2 \quad (2.11)$$

That is, the log of odds is a linear function of x_1 and x_2 and interaction. The interaction effect has the same interpretation as the linear model, in terms of log of odds.

Or, it becomes multiplicative effect when we talk about odds ratios. Stata’s “margins” command is a great tool to calculate marginal effects in various situations, as shown in [Maarten Buis 2010](#).

2.2. A worked example

Let’s illustrate the Ai-Norton point directly: in a logit model with an interaction term, the cross partial derivative $\frac{\partial^2 E(y)}{\partial x_1 \partial x_2}$ (what we actually care about) is not the same object as the estimated coefficient β_{12} on $x_1 x_2$ in the linear index. We simulate data with a true interaction effect on the log-odds scale, fit a logit model, and then use the `marginalEffects` package for two distinct things: first the marginal effect of x_1 on the probability scale at several values of x_2 , and then the cross partial derivative itself.

```
library(dplyr)
library(marginalEffects)

set.seed(42)
n <- 2000
x1 <- rnorm(n)
x2 <- rnorm(n)
eta <- -0.3 + 0.8 * x1 + 0.5 * x2 + 0.6 * x1 * x2
p <- plogis(eta)
y <- rbinom(n, 1, p)
dat <- data.frame(y, x1, x2)

m <- glm(y ~ x1 * x2, data = dat, family = binomial)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-0.2348658	0.05017155	-4.681255	2.851237e-06
x1	0.7792602	0.05930028	13.140920	1.918795e-39
x2	0.5296149	0.05537644	9.563903	1.133977e-21
x1:x2	0.6624955	0.06403580	10.345705	4.376826e-25

The coefficient on `x1:x2` above is $\hat{\beta}_{12}$, the interaction effect on the log-odds (linear index) scale — constant by construction. On the *probability* scale nothing is constant. Start with the marginal effect of x_1 itself, evaluated at three values of x_2 :

```
slopes(m, variables = "x1", by = "x2", newdata = datagrid(x2 = c(-1, 0, 1)))
```

x2	Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
-1	0.0253	0.0159	1.59	0.112	3.2	-0.00591	0.0565
0	0.1919	0.0147	13.09	<0.001	127.6	0.16314	0.2206
1	0.3538	0.0243	14.59	<0.001	157.7	0.30630	0.4014

2. Interaction term in a non-linear model

Term: x1
Type: response
Comparison: dY/dX

The marginal effect of x_1 on $P(y = 1)$ changes across the three values of x_2 — and notice that at $x_2 = -1$ it is not significant ($p \approx 0.11$) even though $\hat{\beta}_{12}$ is significant at better than 0.001. Already the probability-scale answer and the coefficient tell different stories.

The cross partial derivative is a *further* step: not the effect of x_1 at a given x_2 , but how much that effect changes as x_2 moves. That is a difference of slopes, so it needs a hypothesis test on top of the `by` argument rather than a second call to `slopes()`:

```
slopes(m, variables = "x1", by = "x2",  
       newdata = datagrid(x2 = c(-1, 1)),  
       hypothesis = "b2 - b1 = 0")
```

Hypothesis	Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
b2-b1=0	0.329	0.0299	11	<0.001	90.7	0.27	0.387

Type: response

Here the cross partial comes out large and highly significant, which is worth being honest about: in *this* example it agrees with $\hat{\beta}_{12}$ in significance, differing only in scale and magnitude. Ai and Norton's warning is that the two quantities *need not* agree — not that they never do. The dependable lesson is the one visible in the table above: on the probability scale the answer depends on where you evaluate it, so a single coefficient cannot summarise it.

3. Chow test and more

3.1. Chow test

Comparing coefficients across regressions is common. Chow test is one of them. If you'd like to compare coefficients of regressions for two subsets, that's the original Chow test.

The idea is to interact the subset indicator with all the covariates or only the covariate you are interested (treatment). If you only interact the dummy with the treatment variable, then you are assuming all other covariates have the same effect across the two subsets. This may or may not be reasonable.

This post is inspired by Austin Nichols (<https://www.stata.com/statalist/archive/2009-11/msg01485.html>). The case with overlapping samples is all from his code.

Let's see a simple example:

```
est clear
sysuse nlsw88, clear
reg wage hours if south
est sto south
reg wage hours if !south
est sto nonsouth
suest south nonsouth
est sto suest
gen hours1=hours*(south==1)
gen hours2=hours*(south==0)
reg wage south hours?
est sto chow
test _b[hours1]-_b[hours2]=0
esttab south nonsouth suest chow, nogaps mti
```

(NLSW, 1988 extract)

Source	SS	df	MS	Number of obs	=	938
Model	344.732583	1	344.732583	F(1, 936)	=	12.47
Residual	25866.3404	936	27.634979	Prob > F	=	0.0004
Total	26211.0729	937	27.973397	R-squared	=	0.0132
				Adj R-squared	=	0.0121
				Root MSE	=	5.2569

3. Chow test and more

wage	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
hours	.0623497	.0176532	3.53	0.000	.0277053	.0969941
_cons	4.520583	.6957145	6.50	0.000	3.155242	5.885923

Source	SS	df	MS	Number of obs	=	1,304
				F(1, 1302)	=	55.93
Model	1929.41943	1	1929.41943	Prob > F	=	0.0000
Residual	44919.1023	1,302	34.5000785	R-squared	=	0.0412
				Adj R-squared	=	0.0404
Total	46848.5217	1,303	35.9543528	Root MSE	=	5.8737

wage	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
hours	.1107536	.01481	7.48	0.000	.0816995	.1398076
_cons	4.357811	.564756	7.72	0.000	3.24988	5.465743

Simultaneous results for south, nonsouth

Number of obs = 2,242

		Robust				
	Coefficient	std. err.	z	P> z	[95% conf. interval]	
south_mean						
hours	.0623497	.0174426	3.57	0.000	.0281629	.0965366
_cons	4.520583	.6599379	6.85	0.000	3.227128	5.814037
south_lnvar						
_cons	3.319082	.1435305	23.12	0.000	3.037768	3.600397
nonsouth_m~n						
hours	.1107536	.0134936	8.21	0.000	.0843066	.1372006
_cons	4.357811	.4633326	9.41	0.000	3.449696	5.265927
nonsouth_l~r						
_cons	3.540962	.1006119	35.19	0.000	3.343766	3.738157

(4 missing values generated)

(4 missing values generated)

Source	SS	df	MS	Number of obs	=	2,242
				F(3, 2238)	=	36.91
Model	3502.37892	3	1167.45964	Prob > F	=	0.0000
Residual	70785.4426	2,238	31.6288841	R-squared	=	0.0471
				Adj R-squared	=	0.0459
Total	74287.8215	2,241	33.1494072	Root MSE	=	5.624

wage	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
south	.1627711	.9199871	0.18	0.860	-1.641346	1.966888
hours1	.0623497	.0188858	3.30	0.001	.0253142	.0993852
hours2	.1107536	.0141803	7.81	0.000	.0829456	.1385616
_cons	4.357811	.5407453	8.06	0.000	3.297397	5.418226

(1) hours1 - hours2 = 0

F(1, 2238) = 4.20
 Prob > F = 0.0405

	(1) south	(2) nonsouth	(3) suest	(4) chow
main				
hours	0.0623*** (3.53)	0.111*** (7.48)	0.0623*** (3.57)	
south				0.163 (0.18)
hours1				0.0623*** (3.30)
hours2				0.111*** (7.81)
_cons	4.521*** (6.50)	4.358*** (7.72)	4.521*** (6.85)	4.358*** (8.06)
south_lnvar				
_cons			3.319*** (23.12)	
nonsouth_m~n				
hours			0.111*** (8.21)	
_cons			4.358***	

3. Chow test and more

(9.41)				

nonsouth_1~r				
_cons			3.541***	
			(35.19)	

N	938	1304	2242	2242

t statistics in parentheses				
* p<0.05, ** p<0.01, *** p<0.001				

In the above example, we are interested in the effect of hours on wage for south and non-south subsets. The Chow test is to use the entire sample which has both south and nonsouth data. Then use the interaction of south indicator and hours to find the effect of hours for south and nonsouth. By including these two subsamples in the same regression, we can test the equality of the two coefficients.

The other way to do this is to use Stata's "suest" command. This command basically take the two regressions and the variance covariance structure; then a test of the difference between two coefficients can be done. However, "suest" does not work for some commands. In my opinion, using interaction can be more flexible.

3.2. A comparison with two different outcomes

We can also use the same idea to compare the effect of some treatment on two different outcomes, if we have the same set of covariates. We just need to "stack" the two outcomes and run a pooled regression with some interactions.

Here is an example.

```
est clear
sysuse nlsw88, clear
reg south wage hours
est sto south
reg smsa wage hours
est sto smsa
suest south smsa
est sto suest
preserve
gen Y1=south
gen Y2=smsa
gen id=_n
reshape long Y, i(id) j(subsample)
gen wage1=wage*(subsample==1)
gen wage2=wage*(subsample==2)
```

3.2. A comparison with two different outcomes

```

gen hours1=hours*(subsample==1)
gen hours2=hours*(subsample==2)
* Each individual contributes TWO rows here (one per stacked outcome), so the
* errors are correlated within id. Clustering on the stacking id is what makes
* the cross-equation test comparable to suest -- see the note below.
reg Y wage? hours? subsample, cluster(id)
test _b[wage1]-_b[wage2]=0
est sto stacked
* Note: suest stores coefficients under equation-qualified names (e.g.
* south:hours), while the single-sample models store the same
* coefficient simply as "hours"; esttab matches by exact coefficient
* name, so these land on separate rows instead of being aligned. Use
* esttab's coelabel()/eqlabel() options, or rename coefficients on the
* stored estimates first, to get a directly comparable table.
esttab south smsa suest stacked, nogaps mti

```

(NLSW, 1988 extract)

Source	SS	df	MS	Number of obs	=	2,242
Model	14.513685	2	7.25684251	F(2, 2239)	=	30.60
Residual	531.049205	2,239	.237181423	Prob > F	=	0.0000
				R-squared	=	0.0266
				Adj R-squared	=	0.0257
Total	545.56289	2,241	.24344618	Root MSE	=	.48701

south	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
wage	-.0124053	.0018099	-6.85	0.000	-.0159545	-.008856
hours	.0047722	.0009916	4.81	0.000	.0028277	.0067167
_cons	.3372108	.0387354	8.71	0.000	.2612497	.4131718

Source	SS	df	MS	Number of obs	=	2,242
Model	14.6274602	2	7.31373012	F(2, 2239)	=	36.17
Residual	452.719551	2,239	.202197209	Prob > F	=	0.0000
				R-squared	=	0.0313
				Adj R-squared	=	0.0304
Total	467.347012	2,241	.208543959	Root MSE	=	.44966

smsa	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
wage	.0139103	.0016711	8.32	0.000	.0106332	.0171873
hours	.0003663	.0009155	0.40	0.689	-.0014291	.0021616
_cons	.5820585	.0357648	16.27	0.000	.511923	.6521941

3. Chow test and more

Simultaneous results for south, smsa

Number of obs = 2,242

		Coefficient	Robust std. err.	z	P> z	[95% conf. interval]	
south_mean							
	wage	-.0124053	.0019287	-6.43	0.000	-.0161854	-.0086251
	hours	.0047722	.0009825	4.86	0.000	.0028465	.0066978
	_cons	.3372108	.0379918	8.88	0.000	.2627483	.4116733
south_lnvar							
	_cons	-1.43893	.0096645	-148.89	0.000	-1.457872	-1.419988
smsa_mean							
	wage	.0139103	.0018313	7.60	0.000	.010321	.0174996
	hours	.0003663	.0009099	0.40	0.687	-.0014172	.0021497
	_cons	.5820585	.0361037	16.12	0.000	.5112965	.6528205
smsa_lnvar							
	_cons	-1.598512	.0195103	-81.93	0.000	-1.636751	-1.560272

(j = 1 2)

Data	Wide	->	Long
Number of observations	2,246	->	4,492
Number of variables	23	->	23
j variable (2 values)		->	subsample
xij variables:	Y1 Y2	->	Y

(8 missing values generated)

(8 missing values generated)

3.2. A comparison with two different outcomes

```

Linear regression                Number of obs    =      4,484
                                F(5, 2241)        =      96.34
                                Prob > F          =      0.0000
                                R-squared          =      0.1091
                                Root MSE       =      .46871

```

(Std. err. adjusted for 2,242 clusters in id)

		Robust				
Y	Coefficient	std. err.	t	P> t	[95% conf. interval]	
wage1	-.0124053	.0019298	-6.43	0.000	-.0161896	-.008621
wage2	.0139103	.0018323	7.59	0.000	.010317	.0175035
hours1	.0047722	.000983	4.85	0.000	.0028444	.0066999
hours2	.0003663	.0009104	0.40	0.688	-.0014191	.0021517
subsample	.2448477	.0548754	4.46	0.000	.1372358	.3524597
_cons	.092363	.0872219	1.06	0.290	-.0786812	.2634072

(1) wage1 - wage2 = 0

```

F( 1, 2241) = 72.14
Prob > F = 0.0000

```

	(1) south	(2) smsa	(3) suest	(4) stacked
main				
wage	-0.0124*** (-6.85)	0.0139*** (8.32)	-0.0124*** (-6.43)	
hours	0.00477*** (4.81)	0.000366 (0.40)	0.00477*** (4.86)	
wage1				-0.0124*** (-6.43)
wage2				0.0139*** (7.59)
hours1				0.00477*** (4.85)
hours2				0.000366 (0.40)
subsample				0.245*** (4.46)
_cons	0.337***	0.582***	0.337***	0.0924

3. Chow test and more

	(8.71)	(16.27)	(8.88)	(1.06)

south_lnvar				
_cons			-1.439*** (-148.89)	

smsa_mean				
wage			0.0139*** (7.60)	
hours			0.000366 (0.40)	
_cons			0.582*** (16.12)	

smsa_lnvar				
_cons			-1.599*** (-81.93)	

N	2242	2242	2242	4484

t statistics in parentheses				
* p<0.05, ** p<0.01, *** p<0.001				

In this example, we are interested in comparing the effect of wage on south vs. smsa (not interesting, but just as an example). What I did is to reshape it to long format, stacking south and smsa as “Y”. Then creat interaction of other covariates with subsample indicator. Then run the regression with Y on the interaction terms.

One detail matters a great deal here, and it is easy to skip. Stacking puts **two rows per person** in the data, so the two error terms belonging to the same person are correlated by construction. The cross-equation test is precisely a comparison *across* those two rows, so its standard error has to account for that correlation. Clustering on the stacking id does it, and it is not a cosmetic adjustment:

	test of the wage coefficient across equations
suest	$\chi^2(1) = 72.22$
stacked, no clustering	$F(1, 4478) = 114.12$
stacked, cluster (id)	$F(1, 2241) = 72.14$

Without clustering the statistic is inflated by more than half. With it, the stacked test reproduces **suest** almost exactly – the small remaining gap is just the F -versus- χ^2 finite-sample adjustment. Note that **robust** alone will *not* fix this; the cluster variable has to be the id you stacked on.

3.3. Overlapping samples

What if we'd like to compare coefficients for two overlapping subsamples? As I mentioned, Austin Nichols gave the following example:

```
est clear
sysuse nlsw88, clear
ta south smsa
reg wage hours if south
est sto south
reg wage hours if smsa
est sto smsa
suest south smsa
est sto suest
preserve
expand 2
bys idcode: g n=_n
keep if (n==1&south)|(n==2&smsa)
* hours1 is hours in the south subsample (n==1), hours2 is hours in the
* smsa subsample (n==2); within this restricted sample n==1 <=> south==1
* and n==2 <=> smsa==1, so this is equivalent to but clearer than negating
* the complementary condition.
g hours1=hours*(n==1)
g hours2=hours*(n==2)
reg wage hours? n, cl(idcode)
est sto stacked
restore
esttab south smsa suest stacked, nogaps mti
```

(NLSW, 1988 extract)

Lives in	Lives in SMSA		Total
	Not SMSA	SMSA	
Not south	308	996	1,304
South	357	585	942
Total	665	1,581	2,246

Source	SS	df	MS	Number of obs	=	938
Model	344.732583	1	344.732583	F(1, 936)	=	12.47
Residual	25866.3404	936	27.634979	Prob > F	=	0.0004
				R-squared	=	0.0132
				Adj R-squared	=	0.0121

3. Chow test and more

Total | 26211.0729 937 27.973397 Root MSE = 5.2569

wage	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
hours	.0623497	.0176532	3.53	0.000	.0277053	.0969941
_cons	4.520583	.6957145	6.50	0.000	3.155242	5.885923

Source	SS	df	MS	Number of obs	=	1,578
Model	1594.14881	1	1594.14881	F(1, 1576)	=	46.48
Residual	54048.2539	1,576	34.2945773	Prob > F	=	0.0000
				R-squared	=	0.0286
				Adj R-squared	=	0.0280
Total	55642.4027	1,577	35.2837049	Root MSE	=	5.8562

wage	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
hours	.0953636	.0139872	6.82	0.000	.0679281	.1227991
_cons	4.861519	.5443826	8.93	0.000	3.793729	5.929309

Simultaneous results for south, smsa

Number of obs = 1,934

		Robust				
	Coefficient	std. err.	z	P> z	[95% conf. interval]	
south_mean						
hours	.0623497	.0174432	3.57	0.000	.0281617	.0965378
_cons	4.520583	.6599613	6.85	0.000	3.227082	5.814083
south_lnvar						
_cons	3.319082	.1435356	23.12	0.000	3.037758	3.600407
smsa_mean						
hours	.0953636	.0132806	7.18	0.000	.069334	.1213931
_cons	4.861519	.4842914	10.04	0.000	3.912325	5.810713
smsa_lnvar						
_cons	3.534987	.0910825	38.81	0.000	3.356469	3.713506

3.3. Overlapping samples

(2,246 observations created)

(1,969 observations deleted)

(7 missing values generated)

(7 missing values generated)

Linear regression	Number of obs	=	2,516
	F(3, 1933)	=	40.81
	Prob > F	=	0.0000
	R-squared	=	0.0399
	Root MSE	=	5.6403

(Std. err. adjusted for 1,934 clusters in idcode)

		Robust					
	wage	Coefficient	std. err.	t	P> t	[95% conf. interval]	
hours1		.0623497	.0174536	3.57	0.000	.0281198	.0965796
hours2		.0953636	.0132886	7.18	0.000	.0693022	.121425
n		.3409364	.6124145	0.56	0.578	-.8601261	1.541999
_cons		4.179646	1.177889	3.55	0.000	1.869579	6.489713

	(1)	(2)	(3)	(4)
	south	smsa	suest	stacked
main				
hours	0.0623*** (3.53)	0.0954*** (6.82)	0.0623*** (3.57)	
hours1				0.0623*** (3.57)
hours2				0.0954*** (7.18)
n				0.341 (0.56)
_cons	4.521*** (6.50)	4.862*** (8.93)	4.521*** (6.85)	4.180*** (3.55)
south_lnvar				
_cons			3.319***	

3. Chow test and more

			(23.12)	

smsa_mean				
hours			0.0954***	
			(7.18)	
_cons			4.862***	
			(10.04)	

smsa_lnvar				
_cons			3.535***	
			(38.81)	

N	938	1578	1934	2516

t statistics in parentheses				
* p<0.05, ** p<0.01, *** p<0.001				

3.4. IV regression

What about IV regression?

```
sysuse nlsw88, clear
ivregress 2sls wage (hours=union) if south
ivregress 2sls wage (hours=union) if !south
gen hours1=hours*(south==1)
gen hours2=hours*(south==0)
gen union1=union*(south==1)
gen union2=union*(south==0)

ivregress 2sls wage south (hours1 hours2 = union1 union2)
test _b[hours1]-_b[hours2]=0
```

```
already preserved
r(621);
```

(NLSW, 1988 extract)

Instrumental-variables 2SLS regression	Number of obs	=	798
	Wald chi2(1)	=	2.61
	Prob > chi2	=	0.1060
	Root MSE	=	9.5807

wage	Coefficient	Std. err.	z	P> z	[95% conf. interval]

-----+-----							
hours		.97819	.6050822	1.62	0.106	-.2077492	2.164129
_cons		-30.96026	23.32846	-1.33	0.184	-76.6832	14.76268

Endogenous: hours

Exogenous: union

Instrumental-variables 2SLS regression	Number of obs	=	1,079
	Wald chi2(1)	=	6.03
	Prob > chi2	=	0.0140
	Root MSE	=	7.0372

wage		Coefficient	Std. err.	z	P> z	[95% conf. interval]	
hours		.6255843	.2546731	2.46	0.014	.1264342	1.124735
_cons		-14.9159	9.401508	-1.59	0.113	-33.34252	3.510714

Endogenous: hours

Exogenous: union

(4 missing values generated)

(4 missing values generated)

(368 missing values generated)

(368 missing values generated)

Instrumental-variables 2SLS regression	Number of obs	=	1,877
	Wald chi2(3)	=	21.75
	Prob > chi2	=	0.0001
	Root MSE	=	8.2154

wage		Coefficient	Std. err.	z	P> z	[95% conf. interval]	
hours1		.97819	.5188511	1.89	0.059	-.0387394	1.995119
hours2		.6255843	.297311	2.10	0.035	.0428654	1.208303
south		-16.04435	22.81705	-0.70	0.482	-60.76495	28.67625
cons		-14.9159	10.97553	-1.36	0.174	-36.42754	6.595735

Endogenous: hours1 hours2

Exogenous: south union1 union2

3. Chow test and more

```
( 1)  hours1 - hours2 = 0

      chi2( 1) =      0.35
      Prob > chi2 =      0.5554
```

We can see same Chow kind of test works, with IV regression, if we have the right interaction terms.

3.5. IV with fixed effects

However, when doing with a fixed effect IV, I seem to have difficulties. In this example, I use “`reghdfe`” to do an IV regression with fixed effect. We can also use “`xtivreg2`”, but “`ivreghdfe`” is supposed to be faster.

Warning

The examples below cluster standard errors on `race`, which in `nls88` has only three categories (White, Black, Other). Cluster-robust standard errors rely on asymptotics in the number of clusters $G \rightarrow \infty$; with $G = 3$ they are severely downward-biased, inflating t-statistics and false-positive rates. These chunks use `race` purely to illustrate the fixed-effect/IV mechanics — in an actual analysis, cluster on a variable with many more categories (e.g., `industry` or `occupation`), or on the true level of treatment assignment/sampling design.

```
sysuse nls88, clear
gen hours1=hours*(south==1)
gen hours2=hours*(south==0)
gen union1=union*(south==1)
gen union2=union*(south==0)
ivreghdfe wage (hours=union) if south, a(race) cluster(race)
ivreghdfe wage (hours=union) if !south, a(race) cluster(race)
ivreghdfe wage south (hours1 hours2 = union1 union2) , a(race) cluster(race)
```

```
already preserved
r(621);
```

```
(NLSW, 1988 extract)
```

```
(4 missing values generated)
```

```
(4 missing values generated)
```

```
(368 missing values generated)
```

```
(368 missing values generated)
```


(MWFE estimator converged in 1 iterations)

IV (2SLS) estimation

Estimates efficient for homoskedasticity only

Statistics robust to heteroskedasticity and clustering on race

Number of clusters (race) =	3	Number of obs =	798
		F(1, 2) =	69.28
		Prob > F =	0.0141
Total (centered) SS =	12186.8806	Centered R2 =	-6.4527
Total (uncentered) SS =	12186.8806	Uncentered R2 =	-6.4527
Residual SS =	90825.44631	Root MSE =	10.68

		Robust				
wage	Coefficient	std. err.	t	P> t	[95% conf. interval]	
hours	1.107099	.133012	8.32	0.014	.5347947	1.679404

Underidentification test (Kleibergen-Paap rk LM statistic): 1.756
Chi-sq(1) P-val = 0.1852

Weak identification test (Cragg-Donald Wald F statistic): 3.233
(Kleibergen-Paap rk Wald F statistic): 13.977
Stock-Yogo weak ID test critical values: 10% maximal IV size 16.38
15% maximal IV size 8.96
20% maximal IV size 6.66
25% maximal IV size 5.53

Source: Stock-Yogo (2005). Reproduced by permission.

NB: Critical values are for Cragg-Donald F statistic and i.i.d. errors.

Hansen J statistic (overidentification test of all instruments): 0.000
(equation exactly identified)

Instrumented: hours

Excluded instruments: union

Partialled-out: _cons

nb: total SS, model F and R2s are after partialling-out;
any small-sample adjustments include partialled-out
variables in regressor count K

Absorbed degrees of freedom:

Absorbed FE | Categories - Redundant = Num. Coefs |

3. Chow test and more

```

-----+-----|
      race |      3      3      0  *|
-----+-----

```

* = FE nested within cluster; treated as redundant for DoF computation

(MWFE estimator converged in 1 iterations)

IV (2SLS) estimation

Estimates efficient for homoskedasticity only

Statistics robust to heteroskedasticity and clustering on race

```

Number of clusters (race) =      3      Number of obs =      1079
                                     F(  1,    2) =      1.41
                                     Prob > F      =      0.3564
Total (centered) SS      = 19086.22115   Centered R2      = -2.3302
Total (uncentered) SS   = 19086.22115   Uncentered R2    = -2.3302
Residual SS              = 63560.97528   Root MSE        =      7.682

```

```

-----+-----
      |      Robust
      | Coefficient std. err.      t    P>|t|      [95% conf. interval]
-----+-----
hours |   .7023322   .5905772    1.19  0.356   -1.838717    3.243381
-----+-----

```

```

Underidentification test (Kleibergen-Paap rk LM statistic):      0.789
                                     Chi-sq(1) P-val =      0.3745
-----+-----

```

```

Weak identification test (Cragg-Donald Wald F statistic):      5.501
      (Kleibergen-Paap rk Wald F statistic):      3.772
Stock-Yogo weak ID test critical values: 10% maximal IV size    16.38
                                     15% maximal IV size      8.96
                                     20% maximal IV size      6.66
                                     25% maximal IV size      5.53

```

Source: Stock-Yogo (2005). Reproduced by permission.

NB: Critical values are for Cragg-Donald F statistic and i.i.d. errors.

```

-----+-----
Hansen J statistic (overidentification test of all instruments):      0.000
                                     (equation exactly identified)
-----+-----

```

```

Instrumented:      hours
Excluded instruments: union
Partialled-out:    _cons
                  nb: total SS, model F and R2s are after partialling-out;
                      any small-sample adjustments include partialled-out
                      variables in regressor count K
-----+-----

```

Absorbed degrees of freedom:

-----+-----				
Absorbed FE		Categories	- Redundant	= Num. Coefs
-----+-----				
race		3	3	0 *
-----+-----				

* = FE nested within cluster; treated as redundant for DoF computation

(MWFE estimator converged in 1 iterations)

IV (2SLS) estimation

Estimates efficient for homoskedasticity only

Statistics robust to heteroskedasticity and clustering on race

Number of clusters (race) =	3	Number of obs =	1877
		F(3, 2) =	1036.24
		Prob > F =	0.0010
Total (centered) SS =	32142.319	Centered R2 =	-3.9041
Total (uncentered) SS =	32142.319	Uncentered R2 =	-3.9041
Residual SS =	157627.5718	Root MSE =	9.174

-----+-----							
		Robust					
wage		Coefficient	std. err.	t	P> t	[95% conf. interval]	
-----+-----							
hours1		1.142036	.0650507	17.56	0.003	.8621454	1.421927
hours2		.6880691	.5265185	1.31	0.321	-1.577357	2.953495
south		-19.74773	22.10811	-0.89	0.466	-114.8712	75.37577
-----+-----							

Underidentification test (Kleibergen-Paap rk LM statistic): 1.793
Chi-sq(1) P-val = 0.1806

Weak identification test (Cragg-Donald Wald F statistic): 3.584

(Kleibergen-Paap rk Wald F statistic): 8.484

Stock-Yogo weak ID test critical values: 10% maximal IV size 7.03

15% maximal IV size 4.58

20% maximal IV size 3.95

25% maximal IV size 3.63

Source: Stock-Yogo (2005). Reproduced by permission.

NB: Critical values are for Cragg-Donald F statistic and i.i.d. errors.

Warning: estimated covariance matrix of moment conditions not of full rank.
overidentification statistic not reported, and standard errors and
model tests should be interpreted with caution.

Possible causes:

3. Chow test and more

number of clusters insufficient to calculate robust covariance matrix
singleton dummy variable (dummy with one 1 and N-1 0s or vice versa)
partial option may address problem.

```
-----
Instrumented:      hours1 hours2
Included instruments: south
Excluded instruments: union1 union2
Partialled-out:    _cons
                  nb: total SS, model F and R2s are after partialling-out;
                     any small-sample adjustments include partialled-out
                     variables in regressor count K
-----
```

Absorbed degrees of freedom:

```
-----+
Absorbed FE | Categories - Redundant = Num. Coefs |
-----+-----+-----+-----+
      race |           3           3           0   *|
-----+-----+-----+-----+
```

* = FE nested within cluster; treated as redundant for DoF computation

We can see I failed to replicate the first two regressions in the third regression. Why? Because we'll need the fixed effect to be interacted with the subsample indicator to make it right.

Here is another try:

```
sysuse nlsw88, clear
gen hours1=hours*(south==1)
gen hours2=hours*(south==0)
gen union1=union*(south==1)
gen union2=union*(south==0)

gen race1=race*(south==1)
gen race2=race*(south==0)

ivreghdfe wage (hours=union) if south, a(race) cluster(race)
ivreghdfe wage (hours=union) if !south, a(race) cluster(race)
ivregress 2sls wage south (hours1 hours2 = union1 union2) i.race1 i.race2, cluster(race)
```

```
already preserved
r(621);
```

(NLSW, 1988 extract)

(4 missing values generated)

(4 missing values generated)

(368 missing values generated)

(368 missing values generated)

(MWFE estimator converged in 1 iterations)

IV (2SLS) estimation

Estimates efficient for homoskedasticity only

Statistics robust to heteroskedasticity and clustering on race

Number of clusters (race) =	3	Number of obs =	798
		F(1, 2) =	69.28
		Prob > F =	0.0141
Total (centered) SS =	12186.8806	Centered R2 =	-6.4527
Total (uncentered) SS =	12186.8806	Uncentered R2 =	-6.4527
Residual SS =	90825.44631	Root MSE =	10.68

		Robust				
wage	Coefficient	std. err.	t	P> t	[95% conf. interval]	
hours	1.107099	.133012	8.32	0.014	.5347947	1.679404

Underidentification test (Kleibergen-Paap rk LM statistic): 1.756
Chi-sq(1) P-val = 0.1852

Weak identification test (Cragg-Donald Wald F statistic): 3.233
(Kleibergen-Paap rk Wald F statistic): 13.977
Stock-Yogo weak ID test critical values: 10% maximal IV size 16.38
15% maximal IV size 8.96
20% maximal IV size 6.66
25% maximal IV size 5.53

Source: Stock-Yogo (2005). Reproduced by permission.

NB: Critical values are for Cragg-Donald F statistic and i.i.d. errors.

Hansen J statistic (overidentification test of all instruments): 0.000
(equation exactly identified)

Instrumented: hours
Excluded instruments: union
Partialled-out: _cons

3. Chow test and more

nb: total SS, model F and R2s are after partialling-out;
any small-sample adjustments include partialled-out
variables in regressor count K

Absorbed degrees of freedom:

-----+				
Absorbed FE		Categories	- Redundant	= Num. Coefs
-----+				
race		3	3	0 *
-----+				

* = FE nested within cluster; treated as redundant for DoF computation

(MWFE estimator converged in 1 iterations)

IV (2SLS) estimation

Estimates efficient for homoskedasticity only

Statistics robust to heteroskedasticity and clustering on race

Number of clusters (race) =	3	Number of obs =	1079
		F(1, 2) =	1.41
		Prob > F =	0.3564
Total (centered) SS	= 19086.22115	Centered R2	= -2.3302
Total (uncentered) SS	= 19086.22115	Uncentered R2	= -2.3302
Residual SS	= 63560.97528	Root MSE	= 7.682

	-----+					
		Robust				
wage		Coefficient	std. err.	t	P> t	[95% conf. interval]
-----+						
hours		.7023322	.5905772	1.19	0.356	-1.838717 3.243381
-----+						

Underidentification test (Kleibergen-Paap rk LM statistic): 0.789
Chi-sq(1) P-val = 0.3745

Weak identification test (Cragg-Donald Wald F statistic): 5.501
(Kleibergen-Paap rk Wald F statistic): 3.772
Stock-Yogo weak ID test critical values: 10% maximal IV size 16.38
15% maximal IV size 8.96
20% maximal IV size 6.66
25% maximal IV size 5.53

Source: Stock-Yogo (2005). Reproduced by permission.

NB: Critical values are for Cragg-Donald F statistic and i.i.d. errors.

Hansen J statistic (overidentification test of all instruments): 0.000
(equation exactly identified)

```

-----
Instrumented:      hours
Excluded instruments: union
Partialled-out:    _cons
                   nb: total SS, model F and R2s are after partialling-out;
                   any small-sample adjustments include partialled-out
                   variables in regressor count K
-----

```

Absorbed degrees of freedom:

```

-----+
Absorbed FE | Categories - Redundant = Num. Coefs |
-----+-----+
      race |           3           3           0   *|
-----+-----+

```

* = FE nested within cluster; treated as redundant for DoF computation

note: 3.race1 omitted because of collinearity.

note: 3.race2 omitted because of collinearity.

```

Instrumental-variables 2SLS regression      Number of obs   =      1,877
                                           Wald chi2(7)      =    101696.08
                                           Prob > chi2       =      0.0000
                                           Root MSE         =      9.0693

```

(Std. err. adjusted for 3 clusters in race)

```

-----
           |               Robust
           | Coefficient std. err.      z    P>|z|    [95% conf. interval]
-----+-----
hours1 |    1.107099   .1085357   10.20  0.000    .8943732    1.319825
hours2 |    .7023322   .4819806    1.46  0.145   - .2423324    1.646997
south  |   -21.52299   22.35225   -0.96  0.336   -65.33259   22.28662
-----+-----
race1   |
  1     |    3.054296   .1451752   21.04  0.000    2.769758    3.338834
  2     |    1.987268   .176538    11.26  0.000    1.64126    2.333277
  3     |           0 (omitted)
-----+-----
race2   |
  1     |   -.4816127   .434723   -1.11  0.268   -1.333654    .3704287
  2     |   -2.065527   .7694573   -2.68  0.007   -3.573635   -.5574181
  3     |           0 (omitted)
-----+-----
_cons  |   -17.01633   18.01689   -0.94  0.345   -52.32879   18.29614
-----

```

Endogenous: hours1 hours2

Exogenous: south 1.race1 2.race1 1.race2 2.race2 union1 union2

3. Chow test and more

This works. Basically we use dummies which are interactions of subsample indicator and the fixed effect dummies. This would not work if we have a lot of fixed effect units.

One caution about how those dummies were built. `race1 = race*(south==1)` multiplies a *categorical code* by an indicator, so non-south observations are assigned `race1 = 0` and then `i.race1` treats 0 as its own category. That is safe here only because `race` has no zero level (it is coded 1/2/3). If the absorbed variable did have a 0 category, those observations and all the non-south ones would be silently collapsed into the same dummy. For a general recipe, build the interactions from the category dummies directly (`i.race#i.south`) rather than from the code.

But we can trick “`reghdfe`” to use a two way fixed effect option:

This way we can do a test to see whether hours effect differs between these two samples.

```
sysuse nlsw88, clear
gen hours1=hours*(south==1)
gen hours2=hours*(south==0)
gen union1=union*(south==1)
gen union2=union*(south==0)
gen race1=race*(south==1)
gen race2=race*(south==0)

ivreghdfe wage south (hours1 hours2 = union1 union2) , a(race1 race2) cluster(race)
test _b[hours1]-_b[hours2]=0
```

```
already preserved
r(621);
```

```
(NLSW, 1988 extract)
```

```
(4 missing values generated)
```

```
(4 missing values generated)
```

```
(368 missing values generated)
```

```
(368 missing values generated)
```

```
(MWFE estimator converged in 2 iterations)
```

```
IV (2SLS) estimation
```

```
-----
```

```
Estimates efficient for homoskedasticity only
```

```
Statistics robust to heteroskedasticity and clustering on race
```



```

Number of clusters (race) =          3
Number of obs =          1877
F( 2, 2) = 13010.62
Prob > F = 0.0001
Centered R2 = -3.9367
Uncentered R2 = -3.9367
Root MSE = 9.089

Total (centered) SS = 31273.10174
Total (uncentered) SS = 31273.10174
Residual SS = 154386.4216

```

```

-----+-----
            |               Robust
            | Coefficient  std. err.      t    P>|t|    [95% conf. interval]
-----+-----
hours1 |    1.107099   .1331773    8.31  0.014   .5340838   1.680115
hours2 |    .7023322   .5914076    1.19  0.357  -1.84229   3.246954
south |              0 (omitted)
-----+-----

```

```

Underidentification test (Kleibergen-Paap rk LM statistic):          0.000
Chi-sq(1) P-val =          1.0000
-----+-----

```

```

Weak identification test (Cragg-Donald Wald F statistic):          3.793
(Kleibergen-Paap rk Wald F statistic):          0.000
Stock-Yogo weak ID test critical values: 10% maximal IV size          7.03
                                           15% maximal IV size          4.58
                                           20% maximal IV size          3.95
                                           25% maximal IV size          3.63

```

Source: Stock-Yogo (2005). Reproduced by permission.

NB: Critical values are for Cragg-Donald F statistic and i.i.d. errors.

```

-----+-----
Warning: estimated covariance matrix of moment conditions not of full rank.
overidentification statistic not reported, and standard errors and
model tests should be interpreted with caution.

```

Possible causes:

```

    number of clusters insufficient to calculate robust covariance matrix
    singleton dummy variable (dummy with one 1 and N-1 0s or vice versa)

```

partial option may address problem.

```

-----+-----
Collinearities detected among instruments: 1 instrument(s) dropped

```

Instrumented: hours1 hours2

Included instruments: south

Excluded instruments: union1 union2

Partialled-out: _cons

```

    nb: total SS, model F and R2s are after partialling-out;
    any small-sample adjustments include partialled-out
    variables in regressor count K
-----+-----

```

Absorbed degrees of freedom:

```

-----+-----

```

3. Chow test and more

Absorbed FE	Categories	- Redundant	= Num. Coefs
race1	4	0	4
race2	4	2	2

(1) hours1 - hours2 = 0

F(1, 2) = 0.31
 Prob > F = 0.6325

3.6. Chow test with different covariates

What if we want to compare coefficients across equations with different covariates? We can still do Chow test.

Say you have

$$Y_1 = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon_1 \quad (3.1)$$

$$Y_2 = \gamma_0 + \gamma_1 X_1 + \epsilon_2 \quad (3.2)$$

The way you can think of the second equation is that you still have X_2 but it is just a constant, which means it will go into the constant term.

$$Y_2 = \gamma_0 + \gamma_1 X_1 + \gamma_2 C + \epsilon_2 \quad (3.3)$$

So we can just replace X_2 with a constant in the sample for Y_2 and still do Chow test.

```
sysuse nlsw88, clear
reg south wage hours tenure
est sto south
reg smsa wage hours
est sto smsa
suest south smsa
est sto suest
preserve
gen Y1=south
gen Y2=smsa
gen id=_n
reshape long Y, i(id) j(subsample)
gen wage1=wage*(subsample==1)
gen wage2=wage*(subsample==2)
gen hours1=hours*(subsample==1)
gen hours2=hours*(subsample==2)
replace tenure=1 if subsample==2
gen tenure1= tenure*(subsample==1)
```

3.6. Chow test with different covariates

```
gen tenure2= tenure*(subsample==2)
reg Y wage? hours? tenure? subsample, cluster(id)
est sto chow
test _b[hours1]-_b[hours2]=0
esttab suest chow, nogaps mti
```

already preserved
r(621);

(NLSW, 1988 extract)

Source	SS	df	MS	Number of obs	=	2,227
				F(3, 2223)	=	20.30
Model	14.4507388	3	4.81691294	Prob > F	=	0.0000
Residual	527.507052	2,223	.23729512	R-squared	=	0.0267
				Adj R-squared	=	0.0254
Total	541.957791	2,226	.243467112	Root MSE	=	.48713

south	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
wage	-.0119584	.0018359	-6.51	0.000	-.0155586	-.0083583
hours	.004828	.0010072	4.79	0.000	.0028528	.0068031
tenure	-.0024032	.0019221	-1.25	0.211	-.0061726	.0013661
_cons	.346361	.0392121	8.83	0.000	.2694648	.4232573

Source	SS	df	MS	Number of obs	=	2,242
				F(2, 2239)	=	36.17
Model	14.6274602	2	7.31373012	Prob > F	=	0.0000
Residual	452.719551	2,239	.202197209	R-squared	=	0.0313
				Adj R-squared	=	0.0304
Total	467.347012	2,241	.208543959	Root MSE	=	.44966

smsa	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
wage	.0139103	.0016711	8.32	0.000	.0106332	.0171873
hours	.0003663	.0009155	0.40	0.689	-.0014291	.0021616
_cons	.5820585	.0357648	16.27	0.000	.511923	.6521941

Simultaneous results for south, smsa

Number of obs = 2,242

3. Chow test and more

		Coefficient	Robust std. err.	z	P> z	[95% conf. interval]	
south_mean							
	wage	-.0119584	.001953	-6.12	0.000	-.0157862	-.0081307
	hours	.004828	.0009971	4.84	0.000	.0028737	.0067822
	tenure	-.0024032	.0018928	-1.27	0.204	-.0061131	.0013066
	_cons	.346361	.038473	9.00	0.000	.2709554	.4217667
south_lnvar							
	_cons	-1.438451	.0096359	-149.28	0.000	-1.457337	-1.419565
smsa_mean							
	wage	.0139103	.0018313	7.60	0.000	.010321	.0174996
	hours	.0003663	.0009099	0.40	0.687	-.0014172	.0021497
	_cons	.5820585	.0361037	16.12	0.000	.5112965	.6528205
smsa_lnvar							
	_cons	-1.598512	.0195103	-81.93	0.000	-1.636751	-1.560272

(j = 1 2)

Data	Wide	->	Long
Number of observations	2,246	->	4,492
Number of variables	23	->	23
j variable (2 values)		->	subsample
xij variables:	Y1 Y2	->	Y

(8 missing values generated)

(8 missing values generated)

(2,222 real changes made)

(15 missing values generated)

3.6. Chow test with different covariates

(15 missing values generated)

note: subsample omitted because of collinearity.

Linear regression	Number of obs	=	4,469
	F(6, 2241)	=	81.10
	Prob > F	=	0.0000
	R-squared	=	0.1091
	Root MSE	=	.4687

(Std. err. adjusted for 2,242 clusters in id)

		Robust				
Y	Coefficient	std. err.	t	P> t	[95% conf. interval]	
wage1	-.0119584	.0019543	-6.12	0.000	-.0157909	-.008126
wage2	.0139103	.0018325	7.59	0.000	.0103166	.0175039
hours1	.004828	.0009978	4.84	0.000	.0028713	.0067846
hours2	.0003663	.0009105	0.40	0.688	-.0014193	.0021519
tenure1	-.0024032	.0018941	-1.27	0.205	-.0061176	.0013111
tenure2	.2356975	.0550137	4.28	0.000	.1278144	.3435806
subsample	0	(omitted)				
_cons	.346361	.0384988	9.00	0.000	.270864	.4218581

(1) hours1 - hours2 = 0

F(1, 2241)	=	10.03
Prob > F	=	0.0016

	(1)	(2)
	suest	chow
main		
wage	-0.0120*** (-6.12)	
hours	0.00483*** (4.84)	
tenure	-0.00240 (-1.27)	
wage1		-0.0120*** (-6.12)
wage2		0.0139*** (7.59)

3. Chow test and more

```

hours1                0.00483***
                      (4.84)
hours2                0.000366
                      (0.40)
tenure1              -0.00240
                      (-1.27)
tenure2               0.236***
                      (4.28)
subsample              0
                      (.)
_cons                 0.346***    0.346***
                      (9.00)      (9.00)
-----
south_lnvar
_cons                -1.438***
                      (-149.28)
-----
smsa_mean
wage                 0.0139***
                      (7.60)
hours                0.000366
                      (0.40)
_cons                0.582***
                      (16.12)
-----
smsa_lnvar
_cons                -1.599***
                      (-81.93)
-----
N                     2242        4469
-----
t statistics in parentheses
* p<0.05, ** p<0.01, *** p<0.001

```

The stacking method reproduces `suest`'s point estimates: the separate `reg south wage hours tenure` gives `wage = -.0119584`, `hours = .004828`, `tenure = -.0024032`, and the stacked model returns the same three numbers. Read them off the two tables separately rather than trying to compare columns, though – as noted earlier, `suest` stores its coefficients under equation-qualified names (`south:wage`) while the stacked model stores plain `wage`, so `esttab` puts them on different rows and the columns do not line up. For standard errors you can use “robust” or “cluster” option in “reg” command.

3.7. Conclusion

For testing cross equations hypotheses, we can use “suest” or Chow type “stacking” method. Sometimes people use “sureg”. I prefer not use “sureg”. It is a GLS estimator for all the equations

together. It relies on assumptions of the error term of the whole system. It assumes homoscedasticity for example. If it's true, then GLS is more efficient; if not, then it is *inefficient* and its reported standard errors are invalid. (It is not biased: a GLS estimator built on the wrong error covariance still satisfies the same moment conditions, exactly as weighted least squares with the wrong weights remains consistent.) If the equations have the same covariates, then it returns the same coefficient estimates as the single equation estimates. If different covariates, then different estimates as the single equation estimates.

The nice thing about Chow test is that it is very flexible, and it does not rely on Stata's internal functions. You can do it with R or other programs. And it works for complicated models too, usually, as long as the single equation works. "suest" needs the results stored before hand, which may not be available even within stata.

4. Weights in OLS in causal inference

4.1. Unit level weights for OLS

Chad Hazlett and Tanvi Shinkre’s paper “Demystifying and avoiding the OLS “weighting problem”: Unmodeled heterogeneity and straightforward solutions” has demystified OLS weights for me.

When we have unconfoundedness, that is, we assume there is no unobserved confounders, we usually run a linear regression of Y on D and X , where D is the treatment variable and X is the covariates. We know in fact the unconfoundedness actually is only true for each level of X . If X is low dimensional and discrete, then it’s possible to stratify and basically calculate the ATE with g-formula. However, usually we just do a linear regression

$$Y = \alpha + \beta D + \gamma X + \epsilon \quad (4.1)$$

and claim we have controlled for X , especially when X is high-dimensional or continuous. At least we claim it’s a linear approximation of the truth. But, how bad is it?

Suppose we are interested in the ATE of D on Y . Say we have discrete variables X . Under unconfoundedness, the ATE is

$$\begin{aligned} \tau_{ATE} &= \sum E[Y(1) - Y(0)|X = x]P(X = x) \\ &= \sum (E[Y|D = 1, X = x] - E[Y|D = 0, X = x])P(X = x) \\ &= \sum DIM_x P(X = x) \end{aligned} \quad (4.2)$$

This is basically the g-formula, or regression adjustment formula. We compute the DIM (difference in means) between treated and control for each strata $X = x$, then the weighted average is the ATE, with $\hat{P}(X = x)$ as weights.

If we have high dimensional X , we usually do

$$Y = \beta_0 + \beta_1 X + \delta_{reg} D + \epsilon \quad (4.3)$$

We would hope δ_{reg} is close to δ_{ATE} , but we are not sure.

OLS coefficients are:

$$\hat{\beta} = (Z'Z)^{-1}Z'y \quad (4.4)$$

Here I use Z to represent all regressors. Each coefficient is a weighted sum of the outcome,

$$\hat{\tau}_{reg} = \sum \omega_i Y_i \quad (4.5)$$

4. Weights in OLS in causal inference

By FWL theorem,

$$\hat{\tau}_{reg} = \frac{\sum Y_i(D_i - \hat{d}(X_i))}{\sum (D_i - \hat{d}(X_i))^2} \quad (4.6)$$

Here $\hat{d}(X_i) = \hat{p}(D_i = 1|X_i)$ which is the predicted value of linear regression of D on X . The numerator is the covariance between Y and residualized D , the denominator is the variance of the residualized D . Note that the original FWL would have residualized Y here, but it turns out it works the same with the original Y .

Therefore the weights are

$$\omega_i = \frac{D_i - \hat{d}(X_i)}{\sum (D_i - \hat{d}(X_i))^2} \quad (4.7)$$

Note for the denominator, it's a constant. It does *not* scale the weights to sum to one — they sum to **zero**, because $D_i - \hat{d}(X_i)$ is an OLS residual from a regression that includes an intercept. What the denominator normalizes is the contrast: since $\sum D_i(D_i - \hat{d}(X_i)) = \sum (D_i - \hat{d}(X_i))^2$ (the residual is orthogonal to the fitted values), we get $\sum_i \omega_i D_i = 1$, i.e. the weights on the treated sum to +1 and those on the control sum to -1. That is what makes $\sum_i \omega_i Y_i$ a weighted *difference* of means. The numerator for treated units are $1 - \hat{d}(X_i)$, for control units it's $-\hat{d}(X_i)$ (negative). It is similar to IPW weights, which are $1/\hat{d}(X_i)$ for treated and $1/(1 - \hat{d}(X_i))$ for control. These two sets of weights go the same direction. For the treated, the higher probability of being treated, the lower weight; for the control, the higher probability of being treated, the larger the magnitude of the (negative) weight.

If these weights are the same as the IPW weights, then we'll get $\tau_{reg} = \tau_{ATE}$. However, they are usually different.

Another way to write this as weighted difference in means:

$$\hat{\tau}_{wdim} = \sum \omega_i Y_i = \sum_{i:D=1} \tilde{\omega}_i Y_i + \sum_{i:D=0} \tilde{\omega}_i Y_i \quad (4.8)$$

where, with $c = \sum (D_i - \hat{d}(X_i))^2$, the per-arm weights are $\tilde{\omega}_i = (1 - \hat{d}(X_i))/c$ for treated units ($D_i = 1$) and $\tilde{\omega}_i = -\hat{d}(X_i)/c$ for control units ($D_i = 0$)

Because $\hat{d}(X_i)$ is from a linear probability model, it's possible that it's not between 0 and 1. In that case, we can have negative weights. The negative weights could mean the observations do not have a match in the other group. For example, $\hat{d}(X_i) > 1$ will cause the weight $(1 - \hat{d}(X_i))/c$ to flip sign and become negative for a treated unit. That means this unit is more like treated units than any control units. In other words, there is no match in the control group. The mirror image happens for a control unit with $\hat{d}(X_i) < 0$: its weight $-\hat{d}(X_i)/c$ flips to *positive*, so it enters the contrast on the same side as the treated. Either way the unit contributes with the wrong sign.

The worst thing is that we don't know which observations have negative weights, unless we look further and actually run a regression of D on X . If we do that, we are doing similar things as in IPW or matching. If we simply do a linear regression, it's like a black box. We have no idea how

bad some negative weights are. We are relying on extrapolation with regression model like this. Even if there is no matched units, we are extrapolating based on linearity.

Therefore, $\hat{\tau}_{reg}$ is another weighted average of the outcome, which is different from the τ_{ATE} , potentially. In addition, for τ_{ATE} , we can calculate the difference in means for that strata, given we have observations from both treated and controls. In the case of linear regression, we might have units that should not be included in the calculation and we don't know about it.

4.2. Strata level weights

Angrist (1998, *Econometrica*) has the strata-wise weights for OLS:

$$\hat{\tau}_{reg} = \frac{\sum_x \hat{\tau}_x \hat{d}(x)(1 - \hat{d}(x)) \hat{P}(X = x)}{\sum_x \hat{d}(x)(1 - \hat{d}(x)) \hat{P}(X = x)} \quad (4.9)$$

this shows that different from ATE, these weights are $\hat{d}(x)(1 - \hat{d}(x)) \hat{P}(X = x)$. The stratas are weighted not only by $P(X_i = x)$, but also by the variance of treatment status. When $\hat{d}(X_i)$ is close to .5, the weight is the largest.

If there is no treatment effect heterogeneity, then it does not matter how you weight. When there is treatment effect heterogeneity, this estimator could be very different from ATE.

Hazlett and Shinkre show that Angrist's estimator is a special case of a general case. It is only valid when the probability of treatment is linear in X. When you have discrete X and have a saturated model in X, then it's linear in X. Hazlett and Shinkre showed a more general formula for the weights, but here we assume Angrist's formula is correct, basically assuming linearity in modeling probability of treatment.

4.3. Recommended estimation approaches

Under no effect heterogeneity and linearity of Y on D and X, namely, $E[Y|X, D] = \beta_0 + \beta_1 X + \beta_2 D$, there is no problem with a regression of Y on D and X.

If this is not true, then we'll need to be more flexible on the specification to make OLS estimators close to ATE.

Hazlett and Shinkre suggested relaxing the assumption to separate linearity. Basically allowing different intercepts for $Y(0)$ and $Y(1)$.

If that is true, then the following estimators are unbiased: interaction, imputation, g-computation, and stratification.

Interaction estimator is the Lin (2013, *Annals of Applied Statistics*) estimator, basically running a regression of Y on D, $\tilde{X}$, and $D * \tilde{X}$. $\tilde{X}$ is the demeaned X. The coefficient of D is the ATE.

Imputation is to run separate regressions for treatment and control group, then use the results to predict outcomes for the other group. The ATE is the difference in the predicted outcomes. It turns out the imputation estimator is the same as the interaction estimator.

4. Weights in OLS in causal inference

Stratification and g-computation work when X is discrete. Basically get the difference in means for each strata and weight by the probability of strata.

The above mentioned estimators all depend on the linearity assumption. If that is not the case, then we can use matching, such as “mean balancing” matching.

Nonparametric estimation is another option if the linearity assumption fails.

4.4. A worked example

Let’s make the regression weights ω_i concrete and compare them to IPW weights $1/\hat{d}(X_i)$ (treated) and $1/(1 - \hat{d}(X_i))$ (control).

```
library(dplyr)

set.seed(1)
n <- 500
X <- rnorm(n)
d_prob <- plogis(0.5 * X)
D <- rbinom(n, 1, d_prob)
Y <- 1 + 2 * D + 1.5 * X + rnorm(n)
dat <- data.frame(Y, D, X)

# Regression-implied weights via the FWL residualization of D on X.
d_hat <- fitted(lm(D ~ X, data = dat))
resid_D <- dat$D - d_hat
omega_reg <- resid_D / sum(resid_D^2)

# IPW weights (Horvitz-Thompson), using the same linear-probability d_hat
# purely for comparability with the regression weights above.
ipw <- ifelse(dat$D == 1, 1 / d_hat, 1 / (1 - d_hat))
# rescale by a single constant so the treated weights average 1; this only
# fixes the axis scale of the plot below and does not change any comparison
ipw <- ipw / sum(ipw[dat$D == 1]) * sum(dat$D == 1)

tau_reg <- sum(omega_reg * dat$Y)
tau_lm <- coef(lm(Y ~ D + X, data = dat))["D"]

data.frame(tau_reg_from_weights = tau_reg, tau_from_lm_coef = tau_lm)
```


```
tau_reg_from_weights tau_from_lm_coef
D                2.084644          2.084644
```

The two are numerically identical (up to floating point), confirming $\hat{\tau}_{reg} = \sum \omega_i Y_i$. Plotting `omega_reg` against `ipw` for treated and control units separately shows the two weighting schemes move in the same direction but are not proportional — confirming that $\hat{\tau}_{reg} \neq \hat{\tau}_{ATE}$ in general, as discussed above.

```

par(mfrow = c(1, 2))
plot(ipw[dat$D == 1], omega_reg[dat$D == 1],
     xlab = "IPW weight", ylab = "Regression (FWL) weight",
     main = "Treated units")
plot(ipw[dat$D == 0], omega_reg[dat$D == 0],
     xlab = "IPW weight", ylab = "Regression (FWL) weight",
     main = "Control units")

```



```

par(mfrow = c(1, 1))

```

Note also what this example does *not* show. Here $\hat{d}(X_i)$ stays inside $[0.06, 0.99]$, so every weight keeps its expected sign and the negative-weight problem discussed above never materializes. That is a property of this particular design, not a reassurance. Strengthen the linear index and the linear probability model starts predicting outside the unit interval:

```

set.seed(2)
X2 <- rnorm(n)
D2 <- rbinom(n, 1, plogis(3 * X2)) # much stronger selection on X
dat2 <- data.frame(D = D2, X = X2)
d_hat2 <- fitted(lm(D ~ X, data = dat2))
omega2 <- (dat2$D - d_hat2) / sum((dat2$D - d_hat2)^2)

data.frame(
  dhat_below_0 = sum(d_hat2 < 0),
  dhat_above_1 = sum(d_hat2 > 1),
  treated_wrong_sign = sum(dat2$D == 1 & omega2 < 0),

```

4. Weights in OLS in causal inference

```
control_wrong_sign = sum(dat2$D == 0 & omega2 > 0)
)
```

	dhat_below_0	dhat_above_1	treated_wrong_sign	control_wrong_sign
1	28	42	42	28

Those units are the ones with no counterpart in the other arm, and a plain regression of Y on D and X gives you no indication that they are there.

Part II.

Marginal Effects & Fixed Effects

5. Marginal effects in models with fixed effects

5.1. Marginal effects in a linear model

Stata's margins command has been a powerful tool for many economists. It can calculate predicted means as well as predicted marginal effects. However, we do need to be careful when we use it when fixed effects are included. In a linear model, everything works out fine. However, in a non-linear model, you may not want to use margins, since it's not calculating what you have in mind.

In a linear model with fixed effects, we can do it either by “demeaning” every variable, or include dummy variables. They return the same results. Fortunately, marginal effects can be calculated the same way in both models.

For example:

```
clear
sysuse auto
xtset rep78
xtreg price c.mpg##c.trunk, fe
margins , dydx(mpg)
reg price c.mpg##c.trunk i.rep78
margins , dydx(mpg)
```

```
. clear
```

```
. sysuse auto
(1978 automobile data)
```

```
. xtset rep78
```

```
Panel variable: rep78 (unbalanced)
```

```
. xtreg price c.mpg##c.trunk, fe
```

```
Fixed-effects (within) regression
Group variable: rep78
```

```
Number of obs      =           69
Number of groups   =            5
```

```
R-squared:
```

```
    Within = 0.2570
    Between = 0.0653
    Overall = 0.2237
```

```
Obs per group:
```

```
    min =            2
    avg  =          13.8
    max  =           30
```

5. Marginal effects in models with fixed effects

```
corr(u_i, Xb) = -0.4133
F(3, 61) = 7.03
Prob > F = 0.0004
```

price	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
mpg	-98.12003	226.8708	-0.43	0.667	-551.7763	355.5362
trunk	295.0544	343.3934	0.86	0.394	-391.6032	981.712
c.mpg#						
c.trunk	-12.23318	15.94713	-0.77	0.446	-44.12143	19.65506
_cons	7574.85	5321.325	1.42	0.160	-3065.797	18215.5
sigma_u	992.2156					
sigma_e	2631.2869					
rho	.12449059	(fraction of variance due to u_i)				

```
F test that all u_i=0: F(4, 61) = 0.86
Prob > F = 0.4948
```

```
. margins , dydx(mpg)
```

```
Average marginal effects
Model VCE: Conventional
Number of obs = 69
```

```
Expression: Linear prediction, predict()
dy/dx wrt: mpg
```

		Delta-method				
	dy/dx	std. err.	z	P> z	[95% conf. interval]	
mpg	-268.4981	74.12513	-3.62	0.000	-413.7807	-123.2156

```
. reg price c.mpg##c.trunk i.rep78
```

Source	SS	df	MS	Number of obs	=	69
Model	154453046	7	22064720.8	F(7, 61)	=	3.19
Residual	422343913	61	6923670.71	Prob > F	=	0.0061
				R-squared	=	0.2678
				Adj R-squared	=	0.1838
Total	576796959	68	8482308.22	Root MSE	=	2631.3

price	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
-------	-------------	-----------	---	------	----------------------	--

5.2. Marginal effects in a non-linear model

mpg		-98.12003	226.8708	-0.43	0.667	-551.7763	355.5362
trunk		295.0544	343.3934	0.86	0.394	-391.6032	981.712
c.mpg#							
c.trunk		-12.23318	15.94713	-0.77	0.446	-44.12143	19.65506
rep78							
2		438.0002	2161.922	0.20	0.840	-3885.031	4761.031
3		987.1363	2022.606	0.49	0.627	-3057.315	5031.587
4		1240.944	2046.417	0.61	0.547	-2851.12	5333.008
5		2605.83	2161.837	1.21	0.233	-1717.031	6928.691
_cons		6355.731	5209.899	1.22	0.227	-4062.105	16773.57

```
. margins , dydx(mpg)
```

Average marginal effects
Model VCE: OLS

Number of obs = 69

Expression: Linear prediction, predict()
dy/dx wrt: mpg

		Delta-method				
		dy/dx	std. err.	t	P> t	[95% conf. interval]
mpg		-268.4981	74.12513	-3.62	0.001	-416.7205 -120.2758

.

All is fine.

5.2. Marginal effects in a non-linear model

In a nonlinear model, we need to be more careful:

```
clear
set seed 42
set obs 250
gen id = ceil(_n/5)
bys id: gen t = _n
gen mpg = rnormal(20,4)
gen trunk = rnormal(15,3)
gen alpha = rnormal(0,1)
```

5. Marginal effects in models with fixed effects

```
bys id: replace alpha = alpha[1]
gen lambda = exp(0.5 + 0.05*mpg - 0.03*trunk + alpha)
gen accidents = rpoisson(lambda)
xtset id
xtpoisson accidents mpg trunk, fe
margins , dydx(mpg)
margins , dydx(mpg) predict(nu0)
poisson accidents mpg trunk i.id
margins , dydx(mpg)
```

```
. clear

. set seed 42

. set obs 250
Number of observations (_N) was 0, now 250.

. gen id = ceil(_n/5)

. bys id: gen t = _n

. gen mpg = rnormal(20,4)

. gen trunk = rnormal(15,3)

. gen alpha = rnormal(0,1)

. bys id: replace alpha = alpha[1]
(200 real changes made)

. gen lambda = exp(0.5 + 0.05*mpg - 0.03*trunk + alpha)

. gen accidents = rpoisson(lambda)

. xtset id

Panel variable: id (balanced)

. xtpoisson accidents mpg trunk, fe
note: 1 group (5 obs) dropped because of all zero outcomes

Iteration 0: Log likelihood = -358.34297
Iteration 1: Log likelihood = -340.16596
Iteration 2: Log likelihood = -340.16413
Iteration 3: Log likelihood = -340.16413
```

5.2. Marginal effects in a non-linear model

Conditional fixed-effects Poisson regression
Group variable: id

Number of obs = 245
Number of groups = 49

Obs per group:

min = 5
avg = 5.0
max = 5

Log likelihood = -340.16413

Wald chi2(2) = 35.90
Prob > chi2 = 0.0000

accidents	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
mpg	.0487946	.0090524	5.39	0.000	.0310521	.0665371
trunk	-.0316066	.0108931	-2.90	0.004	-.0529566	-.0102566

. margins , dydx(mpg)

Average marginal effects
Model VCE: OIM

Number of obs = 245

Expression: Linear prediction, predict()
dy/dx wrt: mpg

		Delta-method				[95% conf. interval]	
	dy/dx	std. err.	z	P> z			
mpg	.0487946	.0090524	5.39	0.000	.0310521	.0665371	

. margins , dydx(mpg) predict(nu0)

Average marginal effects
Model VCE: OIM

Number of obs = 245

Expression: Predicted number of events (assuming u_i=0), predict(nu0)
dy/dx wrt: mpg

		Delta-method				[95% conf. interval]	
	dy/dx	std. err.	z	P> z			
mpg	.0850287	.0340547	2.50	0.013	.0182826	.1517747	

5. Marginal effects in models with fixed effects

```
. poisson accidents mpg trunk i.id
```

```
Iteration 0: Log likelihood = -485.81671
Iteration 1: Log likelihood = -453.45562
Iteration 2: Log likelihood = -453.02481
Iteration 3: Log likelihood = -452.93985
Iteration 4: Log likelihood = -452.92127
Iteration 5: Log likelihood = -452.91701
Iteration 6: Log likelihood = -452.916
Iteration 7: Log likelihood = -452.91579
Iteration 8: Log likelihood = -452.91575
```

Poisson regression

Number of obs = 250

LR chi2(51) = 934.93

Prob > chi2 = 0.0000

Pseudo R2 = 0.5079

Log likelihood = -452.91575

accidents	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
mpg	.0487946	.0090525	5.39	0.000	.031052	.0665371
trunk	-.0316067	.0108931	-2.90	0.004	-.0529568	-.0102565
i.id						
2	.3464221	.5592835	0.62	0.536	-.7497533	1.442598
3	1.123341	.5131078	2.19	0.029	.1176679	2.129014
4	.1098041	.6057603	0.18	0.856	-1.077464	1.297073
5	.4069152	.5702496	0.71	0.475	-.7107536	1.524584
6	1.791221	.4776303	3.75	0.000	.8550829	2.727359
7	-14.00991	475.5637	-0.03	0.976	-946.0976	918.0778
8	.9632705	.5216941	1.85	0.065	-.0592311	1.985772
9	.1039008	.6061094	0.17	0.864	-1.084052	1.291853
10	.7748186	.5273812	1.47	0.142	-.2588296	1.808467
11	.7331684	.5395719	1.36	0.174	-.3243732	1.79071
12	1.114943	.5007908	2.23	0.026	.133411	2.096475
13	-.2054342	.6343396	-0.32	0.746	-1.448717	1.037848
14	1.248116	.4945786	2.52	0.012	.2787602	2.217472
15	1.395876	.4914484	2.84	0.005	.4326552	2.359098
16	1.689054	.4845687	3.49	0.000	.7393167	2.638791
17	.9941942	.5171224	1.92	0.055	-.0193471	2.007735
18	1.742704	.4844911	3.60	0.000	.793119	2.692289
19	.9912213	.5168563	1.92	0.055	-.0217985	2.004241
20	-1.134862	.8375742	-1.35	0.175	-2.776477	.5067532
21	1.484731	.4879013	3.04	0.002	.5284622	2.441
22	2.442088	.4643803	5.26	0.000	1.531919	3.352256
23	.9732572	.5265941	1.85	0.065	-.0588483	2.005363
24	.280425	.5859507	0.48	0.632	-.8680172	1.428867
25	2.472329	.4641555	5.33	0.000	1.562601	3.382057

5.2. Marginal effects in a non-linear model

26		.5602592	.5485592	1.02	0.307	-.5148971	1.635416
27		1.537242	.4861325	3.16	0.002	.5844401	2.490044
28		1.34391	.5026286	2.67	0.008	.3587763	2.329044
29		2.370956	.4660218	5.09	0.000	1.45757	3.284342
30		2.453329	.4642702	5.28	0.000	1.543377	3.363282
31		.9383342	.5171502	1.81	0.070	-.0752616	1.95193
32		.1850002	.5864331	0.32	0.752	-.9643876	1.334388
33		-.2686233	.6711829	-0.40	0.689	-1.584118	1.046871
34		1.547864	.485423	3.19	0.001	.5964525	2.499276
35		1.557431	.4878123	3.19	0.001	.6013366	2.513526
36		1.295147	.4965837	2.61	0.009	.3218611	2.268433
37		2.668556	.4616456	5.78	0.000	1.763748	3.573365
38		.9287953	.5326269	1.74	0.081	-.1151343	1.972725
39		-.9247308	.8366882	-1.11	0.269	-2.56461	.715148
40		1.358213	.4979867	2.73	0.006	.3821769	2.334249
41		-.0472585	.6070096	-0.08	0.938	-1.236975	1.142458
42		1.394504	.4958685	2.81	0.005	.4226201	2.366389
43		-.5414444	.7303533	-0.74	0.458	-1.972911	.8900219
44		.9881123	.513744	1.92	0.054	-.0188073	1.995032
45		1.171458	.5165503	2.27	0.023	.1590378	2.183878
46		.817007	.5269245	1.55	0.121	-.2157461	1.84976
47		2.800173	.4596719	6.09	0.000	1.899233	3.701113
48		.365911	.5704157	0.64	0.521	-.7520831	1.483905
49		2.445633	.4649624	5.26	0.000	1.534324	3.356943
50		.7177863	.5279692	1.36	0.174	-.3170143	1.752587
_cons		-.4227135	.5046249	-0.84	0.402	-1.41176	.566333

```
. margins , dydx(mpg)
```

Average marginal effects
Model VCE: OIM

Number of obs = 250

Expression: Predicted number of events, predict()
dy/dx wrt: mpg

		Delta-method				
		dy/dx	std. err.	z	P> z	[95% conf. interval]
mpg		.2238673	.0420551	5.32	0.000	.1414408 .3062937

.

(We use a simulated count outcome `accidents` and a proper panel with 50 units of 5 periods each,

5. Marginal effects in models with fixed effects

rather than `auto`'s continuous `price` and its 5-category `rep78` — a true panel needs many units, and a fixed-effect count model needs an actual count dependent variable.)

In this example, “`xtpoisson, fe`” and “`poisson ... i.id`” return the same coefficients (both give `mpg` = .0487946 with a standard error of .0090524). Fixed effect Poisson model (sometimes called conditional fixed effect Poisson) is the same models as a Poisson model with dummies, just like a linear model (OLS with dummies is the same as fixed effect OLS). Poisson model and OLS are unique in this sense that there is no “incidental parameter” problem.

We see in this example, `margins` commands do not return the same marginal effects, even though the models are the same. The reason behind this is that in a conditional fixed effect Poisson, the fixed effects are not estimated (they are not in the final likelihood function that gets estimated). Therefore, we'll have to make a decision what values to use as the values of the fixed effects. “`margins, predict(nu0)`” simply set all fixed effects to zero. On the other hand, `margins` after Poisson model with dummies does not do that. The fixed effect in that case gets estimated. Therefore the marginal effects in that case make more sense.

So our advice for a conditional Poisson model is that we should not use `margins` to calculate marginal effects afterwards; instead, we should simply stick with the original coefficient estimates.

The same logic applies to the conditional logit model. Fixed effects are not estimated in that model; simply setting them to zero does not make too much sense. In addition, conditional logit model is not the same model as a logit model with dummies, since there is the “incidental parameter” problem. Again, we should just focus on the coefficient estimates as the effect on the logged odds.

In other words, for fixed effect (conditional) logit model, the situation is worse: you cannot do logit with dummies, unless you have a deep panel. That is, when you have, say, more than 20 observations per group, the “incidental parameter” bias becomes negligible. If you stay with conditional logit model, the fixed effects are not estimated. Unfortunately the predicted probability depends on the fixed effects. Stata's `margins` command after `clogit` (or `xtlogit, fe`) comes with a few options, but none is reasonable for the fixed effects. For example, the `pu0` option is to assume all fixed effects being 0.

In a fixed effect logit model,

$$\log(P(y = 1)/(1 - P(y = 1))) = \alpha_i + \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 * x_2 \quad (5.1)$$

Here α_i is fixed effect for each firm. Therefore,

$$P(y = 1) = F(\alpha_i + \beta_1 x_1 + \beta_2 x_2 + \beta_{12} x_1 * x_2) \quad (5.2)$$

F is the logistic CDF (logit); for probit it would be the normal CDF. Therefore, without estimating α_i , there is no way to predict P in a reasonable way (assuming $\alpha = 0$ is not reasonable to me).

However, if we stick with logged odds ($LO = \log(P(y = 1)/(1 - P(y = 1)))$), then LO is a linear function of α_i and other covariates. In that case, the marginal effects of x_1 or x_2 on Y has nothing to do with α_i .

Therefore, we can use `margins` command to calculate effects on the logged odds, which will be “`predict(xb)`” option. Note that `xb` here is $X_{it}\beta$ only — the individual fixed effect α_i is not estimated in conditional logit and so is not part of the prediction. The logged-odds *effect* of a covariate

$(\partial(X\beta)/\partial x)$ therefore matches the original coefficient exactly and does not depend on α_i , but $\mathbf{x}\beta$ itself is a log-odds relative to each individual's (unidentified) baseline, not the individual's absolute log-odds. This lets you make linear extrapolations of relative log-odds across values of the covariates, but not of absolute predicted probabilities, since those would require α_i .

```
clear
webuse union
clogit union c.age##i.south not_smsa grade, group(idcode)
margins, at( age=(15 20 25 30 35 40) south=(0 1)) predict(xb)
marginsplot
```

```
. clear

. webuse union
(NLS Women 14-24 in 1968)

. clogit union c.age##i.south not_smsa grade, group(idcode)
note: multiple positive outcomes within groups encountered.
note: 2,744 groups (14,165 obs) omitted because of all positive or
      all negative outcomes.
```

```
Iteration 0:  Log likelihood = -4518.8815
Iteration 1:  Log likelihood = -4512.8224
Iteration 2:  Log likelihood = -4512.8192
Iteration 3:  Log likelihood = -4512.8192
```

```
Conditional (fixed-effects) logistic regression      Number of obs = 12,035
                                                    LR chi2(5)      = 74.73
                                                    Prob > chi2     = 0.0000
Log likelihood = -4512.8192                        Pseudo R2      = 0.0082
```

union	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
age	.0096842	.0050265	1.93	0.054	-.0001676	.019536
1.south	-1.382178	.276966	-4.99	0.000	-1.925022	-.8393346
south#c.age						
1	.0208997	.0081247	2.57	0.010	.0049756	.0368238
not_smsa	.0195233	.1131292	0.17	0.863	-.2022058	.2412523
grade	.0822276	.0419062	1.96	0.050	.000093	.1643622

```
. margins, at( age=(15 20 25 30 35 40) south=(0 1)) predict(xb)
```

5. Marginal effects in models with fixed effects

Predictive margins
Model VCE: OIM

Number of obs = 12,035

Expression: Linear prediction, predict(xb)

```
1._at: age   = 15
        south = 0
2._at: age   = 15
        south = 1
3._at: age   = 20
        south = 0
4._at: age   = 20
        south = 1
5._at: age   = 25
        south = 0
6._at: age   = 25
        south = 1
7._at: age   = 30
        south = 0
8._at: age   = 30
        south = 1
9._at: age   = 35
        south = 0
10._at: age  = 35
        south = 1
11._at: age  = 40
        south = 0
12._at: age  = 40
        south = 1
```

		Delta-method					
		Margin	std. err.	z	P> z	[95% conf. interval]	
_at							
1		1.202147	.5190753	2.32	0.021	.184778	2.219516
2		.133464	.5599015	0.24	0.812	-.9639228	1.230851
3		1.250568	.5153257	2.43	0.015	.240548	2.260588
4		.2863834	.5465398	0.52	0.600	-.7848148	1.357582
5		1.298989	.5127819	2.53	0.011	.2939548	2.304023
6		.4393029	.5349589	0.82	0.412	-.6091973	1.487803
7		1.34741	.5114619	2.63	0.008	.3449629	2.349857
8		.5922223	.5252767	1.13	0.260	-.4373011	1.621746
9		1.395831	.5113752	2.73	0.006	.3935538	2.398108
10		.7451418	.5175997	1.44	0.150	-.2693351	1.759619
11		1.444252	.5125224	2.82	0.005	.4397264	2.448777
12		.8980612	.5120182	1.75	0.079	-.1054761	1.901598

```
. marginsplot
```

```
Variables that uniquely identify margins: age south
```

In this example, we have a fixed effect logit on union status, with age and south interaction, age as continuous variable. Suppose we'd like to see the predicted logged odds of union status for different age and south/north, then we can still use margins to predict logged odds. But we cannot use margins to predict probability, since the fixed effects are not estimated.

Read the vertical axis of that plot as a *relative* scale. Because α_i is not estimated, $\mathbf{x}\mathbf{b}$ is each individual's log-odds measured from an unidentified baseline, so the level of any single point is arbitrary; what is interpretable is how the curves move across age and how far apart the south and non-south curves sit. Differences within the plot are the estimable quantities, not heights.

6. Comparing Marginal effects with margins command

6.1. Comparing Marginal effects

Stata's margins command has been a powerful tool for many economists. It can calculate predicted means as well as predicted marginal effects. Sometimes we'd like to compare those marginal effects. People use margins and marginsplot to generate marginal effects; then draw conclusions on whether there is a difference between marginal effects, based on whether the confidence intervals overlap or not. However, that can actually be wrong. In this post, I'd like to introduce a way to compare effects.

For example:

```
webuse nhanes2f, clear
logit diabetes i.female#c.age, nolog
margins female, at(age=(20 30 40 50 60 70))
marginsplot
graph export "marginsplot1.svg", as(svg) replace
margins r.female, at(age=(20 30 40 50 60 70))
marginsplot
graph export "marginsplot2.svg", as(svg) replace
```

Logistic regression

Number of obs = 10,335

LR chi2(3) = 353.06

Prob > chi2 = 0.0000

Log likelihood = -1822.5393

Pseudo R2 = 0.0883

diabetes	Coefficient	Std. err.	z	P> z	[95% conf. interval]		

female							
Female	1.404111	.4858776	2.89	0.004	.4518088	2.356414	
age	.0714166	.0063231	11.29	0.000	.0590235	.0838097	

female#c.age							
Female	-.0206572	.0078369	-2.64	0.008	-.0360172	-.0052972	

_cons	-7.041892	.3967299	-17.75	0.000	-7.819468	-6.264315	

6. Comparing Marginal effects with margins command

Adjusted predictions

Number of obs = 10,335

Model VCE: OIM

Expression: Pr(diabetes), predict()

1._at: age = 20
2._at: age = 30
3._at: age = 40
4._at: age = 50
5._at: age = 60
6._at: age = 70

		Delta-method				
		Margin	std. err.	z	P> z	[95% conf. interval]

_at#female						
1#Male		.0036348	.0009895	3.67	0.000	.0016953 .0055743
1#Female		.0097316	.0018436	5.28	0.000	.0061184 .0133449
2#Male		.007396	.0015622	4.73	0.000	.0043341 .0104579
2#Female		.0160637	.0023434	6.85	0.000	.0114706 .0206568
3#Male		.0149906	.0022831	6.57	0.000	.0105159 .0194653
3#Female		.026406	.0027757	9.51	0.000	.0209657 .0318462
4#Male		.0301469	.0029996	10.05	0.000	.0242677 .0360261
4#Female		.043115	.0030923	13.94	0.000	.0370543 .0491758
5#Male		.0596983	.0040248	14.83	0.000	.0518099 .0675867
5#Female		.069641	.0040284	17.29	0.000	.0617454 .0775366
6#Male		.1147889	.0089467	12.83	0.000	.0972536 .1323242
6#Female		.1106004	.0078699	14.05	0.000	.0951757 .126025

Variables that uniquely identify margins: age female

file marginsplot1.svg saved as SVG format

Contrasts of adjusted predictions

Number of obs = 10,335

Model VCE: OIM

Expression: Pr(diabetes), predict()

1._at: age = 20
2._at: age = 30
3._at: age = 40
4._at: age = 50
5._at: age = 60
6._at: age = 70



	df	chi2	P>chi2
female@_at			
(Female vs Male) 1	1	8.49	0.0036
(Female vs Male) 2	1	9.47	0.0021
(Female vs Male) 3	1	10.09	0.0015
(Female vs Male) 4	1	9.06	0.0026
(Female vs Male) 5	1	3.05	0.0808
(Female vs Male) 6	1	0.12	0.7252
Joint	4	40.48	0.0000

	Contrast	std. err.	[95% conf. interval]	
female@_at				
(Female vs Male) 1	.0060969	.0020923	.0019959	.0101978
(Female vs Male) 2	.0086677	.0028164	.0031476	.0141878
(Female vs Male) 3	.0114154	.003594	.0043712	.0184595
(Female vs Male) 4	.0129682	.0043081	.0045244	.0214119
(Female vs Male) 5	.0099427	.0056945	-.0012183	.0211037
(Female vs Male) 6	-.0041885	.0119155	-.0275425	.0191654

Variables that uniquely identify margins: age

file marginsplot2.svg saved as SVG format

6. Comparing Marginal effects with margins command



Here we run a logit model of diabetes and have female indicator and continuous age variable and their interaction as predictors.

See my other post about interpreting marginal effects in non-linear model. In general we recommend against going after marginal effects in a non-linear model, instead we should focus on original

coefficients. However, in case you still want to do that (many people do), here is an example. Suppose you are interested in comparing the effect of female on probability, at different ages, ranging from 20 to 70. The first margins command above tells you the predicted probability for each gender, the second tells the difference between the two genders.

People tend to draw conclusions based on the confidence intervals overlapping or not. That is, if the confidence intervals of two effects overlap, then there is no statistically significant difference between the two effects. This actually can be wrong. The reason is that the confidence intervals are based on individual variable's distribution, that is, how big the variance of that estimate. The difference between two effects, however, depends on the joint distribution of the two effects. That is, the variances of the two effects, and the covariance between the two effects. We should expect two effects are correlated, since they are derived from estimates from the same model, and often the same data. Therefore, two effects with overlapping confidence intervals do not necessarily mean the difference between the two is not statistically different from zero.

To see this, suppose we are interested in $x - y$, the difference between two effects x and y ,

$$Var(x - y) = Var(x) + Var(y) - 2Cov(x, y) \quad (6.1)$$

The variance of the difference is the sum of the two variances, minus *twice* the covariance of x and y . When we draw confidence intervals of marginal effects, x and y , they are determined by $Var(x)$ and $Var(y)$, respectively. But the confidence interval of $x - y$ should be based on $Var(x - y)$, which is affected by $Cov(x, y)$. When x and y are highly correlated, the variance of the difference can become small, and therefore the estimated $x - y$ can be statistically significant, even if the two confidence intervals overlap (meaning the variances of x and y can be relatively large).

For example, we see in the second margins output, the effect of female on probability is .006 at point 1 (age 20), .0087 at point 2 (age 30). If you'd like to compare these two effects, you see their confidence intervals overlap. It's easy to draw a conclusion that these two effects do not differ statistically because their confidence intervals overlap. However, if we do a statistical test of the difference, we see the difference is significant.

```
webuse nhanes2f, clear
logit diabetes i.female##c.age, nolog
margins r.female, at(age=(20 30 40 50 60 70)) post
lincom _b[r1vs0.female@2._at] - _b[r1vs0.female@1._at]
```

Logistic regression	Number of obs = 10,335
	LR chi2(3) = 353.06
	Prob > chi2 = 0.0000
Log likelihood = -1822.5393	Pseudo R2 = 0.0883

diabetes	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
female						
Female	1.404111	.4858776	2.89	0.004	.4518088	2.356414
age	.0714166	.0063231	11.29	0.000	.0590235	.0838097

6. Comparing Marginal effects with margins command

```
female#c.age |
      Female |  -.0206572   .0078369   -2.64   0.008   -.0360172   -.0052972
      |
      _cons |  -7.041892   .3967299  -17.75   0.000   -7.819468   -6.264315
-----
```

Contrasts of adjusted predictions
Model VCE: OIM

Number of obs = 10,335

Expression: Pr(diabetes), predict()

```
1._at: age = 20
2._at: age = 30
3._at: age = 40
4._at: age = 50
5._at: age = 60
6._at: age = 70
```

```
-----
              |          df          chi2      P>chi2
-----+-----
      female@_at |
(Female vs Male) 1 |          1          8.49      0.0036
(Female vs Male) 2 |          1          9.47      0.0021
(Female vs Male) 3 |          1         10.09      0.0015
(Female vs Male) 4 |          1          9.06      0.0026
(Female vs Male) 5 |          1          3.05      0.0808
(Female vs Male) 6 |          1          0.12      0.7252
      Joint      |          4         40.48      0.0000
-----
```

```
-----
              |          Delta-method
              | Contrast  std. err.    [95% conf. interval]
-----+-----
      female@_at |
(Female vs Male) 1 |   .0060969   .0020923     .0019959     .0101978
(Female vs Male) 2 |   .0086677   .0028164     .0031476     .0141878
(Female vs Male) 3 |   .0114154   .003594     .0043712     .0184595
(Female vs Male) 4 |   .0129682   .0043081     .0045244     .0214119
(Female vs Male) 5 |   .0099427   .0056945     -.0012183     .0211037
(Female vs Male) 6 |  -.0041885   .0119155     -.0275425     .0191654
-----
```

```
( 1)  - r1vs0.female@1bn._at + r1vs0.female@2._at = 0
-----
```

	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
(1)	.0025709	.0007946	3.24	0.001	.0010134	.0041283

We can do this with native Stata syntax: `margins ...`, `post` posts the estimated contrasts as if they were regular coefficients (named `r1vs0.female@1._at`, `r1vs0.female@2._at`, etc. — check `matrix list e(b)` to see the exact names), so `lincom` can test any linear combination of them natively, without depending on a user-written command.

Fortunately, `margins` has an option to show this difference, by `contrast(atcontrast(r))`. This is to say comparing all other points to the reference category.

```
webuse nhanes2f, clear
logit diabetes i.female#c.age, nolog
margins r.female, at(age=(20 30 40 50 60 70)) contrast(atcontrast(r))
marginsplot
graph export "marginsplot3.svg", as(svg) replace
```

Logistic regression

Number of obs = 10,335

LR chi2(3) = 353.06

Prob > chi2 = 0.0000

Log likelihood = -1822.5393

Pseudo R2 = 0.0883

diabetes	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
female						
Female	1.404111	.4858776	2.89	0.004	.4518088	2.356414
age	.0714166	.0063231	11.29	0.000	.0590235	.0838097
female#c.age						
Female	-.0206572	.0078369	-2.64	0.008	-.0360172	-.0052972
_cons	-7.041892	.3967299	-17.75	0.000	-7.819468	-6.264315

Contrasts of adjusted predictions

Number of obs = 10,335

Model VCE: OIM

Expression: `Pr(diabetes), predict()`

1._at: age = 20

2._at: age = 30

3._at: age = 40

4._at: age = 50

5._at: age = 60

6. Comparing Marginal effects with margins command


6._at: age = 70

	df	chi2	P>chi2
-----+-----			
_at#female			
(2 vs 1) (Female vs Male)	1	10.47	0.0012
(3 vs 1) (Female vs Male)	1	9.28	0.0023
(4 vs 1) (Female vs Male)	1	5.48	0.0192
(5 vs 1) (Female vs Male)	1	0.50	0.4807
(6 vs 1) (Female vs Male)	1	0.65	0.4203
Joint	4	20.04	0.0005
-----+-----			

	Contrast	std. err.	Delta-method [95% conf. interval]	
-----+-----				
_at#female				
(2 vs 1) (Female vs Male)	.0025709	.0007946	.0010134	.0041283
(3 vs 1) (Female vs Male)	.0053185	.001746	.0018964	.0087407
(4 vs 1) (Female vs Male)	.0068713	.0029351	.0011187	.0126239
(5 vs 1) (Female vs Male)	.0038458	.0054532	-.0068422	.0145339
(6 vs 1) (Female vs Male)	-.0102854	.0127614	-.0352972	.0147264
-----+-----				

Variables that uniquely identify margins: age

file marginsplot3.svg saved as SVG format



From this we see the difference of effects between point 2 and 1 is .0026, and significant — the same conclusion the `lincom` test gave above.

7. Fixed or Random Effect, or Both?

7.1. Panel data

When we have a panel data (repeated observations over time, or observations clustered at higher level), we usually think of two choices: random effect or fixed effect? Economists usually prefers fixed effect models, since it wipes out all within unit heterogeneity. Economists do not like random effect models since it has a big assumption: the random effects need to be uncorrelated to other covariates in the model. To see this, suppose we have

$$y_{it} = \beta_0 + \beta_1 x_{it} + c_i + \epsilon_{it} \quad (7.1)$$

Suppose we have individuals $i = 1, \dots, n$ measured at time $t = 1, \dots, T$. Here c_i is the unobserved time-invariant individual effects. The difference between fixed and random effects is in how they handle c_i .

Fixed effect models for a linear model can be implemented by one of these two methods: with dummies of individuals, or run an OLS with de-meaned y and x . These two methods are equivalent. In a non-linear model, things are more difficult, except Poisson model, other non-linear model with dummies suffer “incidental parameter” problem. The gold-standard is to do a conditional likelihood (conditional logit for example), which “absorbs” the fixed effects in the likelihood function, therefore it’s not necessary to estimate them. Unfortunately most non-linear models do not have such nice conditional likelihood. In that case we can only hope the bias would be small (it does get smaller when you have deeper panel, that is , number of observations per individual).

Random effect models treat c_i as part of the error term. In that case, it comes the biggest drawback: the covariates have to be uncorrelated with the error term to have a consistent estimator. Therefore in the above equation, x has to be uncorrelated with c_i , which economists in general do not think it’s realistic.

7.2. Time-invariant variables

Sometimes people are interested in the effect of time-invariant variables, thus the model

$$y_{it} = \beta_0 + \beta_1 x_{it} + c_i + \gamma z_i + \epsilon_{it} \quad (7.2)$$

Fixed effect models cannot handle this, because γ is not identified because z_i is perfectly collinear with c_i . Random effect can still be estimated, treating z_i simply as another covariate.

7. Fixed or Random Effect, or Both?

7.3. Between-within model

Usually we were told to do a Hausman test to see whether we should use fixed effect or random effect model. The basic idea is the random effect is more efficient if the assumptions are satisfied. If not, then fixed effect model is still consistent. The Hausman test is to compare the difference between the two. If the difference is small then stick with random effect. If it's big, then fixed effect should be preferred since it's consistent.

However, there is a between-within model (BW) that can incorporate both. [Neuhaus and Kalbfleisch \(1998\)](#) introduced the BW estimator,

$$y_{it} = \beta_0 + \beta_1(x_{it} - \bar{x}_i) + \beta_2\bar{x}_i + c_i + \gamma z_i + \epsilon_{it} \quad (7.3)$$

It can be shown that β_1 is the same as the one in the fixed effect model. It is the effect of within individual deviation of x on within individual deviation of y . β_2 is the effect of mean of x on mean of y , that is, the “between” effect. γ is the effect of time-invariant variable on the mean of y .

The other specification of BW estimator is

$$y_{it} = \beta_0 + \beta_1 x_{it} + \beta_2 \bar{x}_i + c_i + \gamma z_i + \epsilon_{it} \quad (7.4)$$

This is just some transformation of the original specification, it's the same model. β_1 is exactly the same as before, and β_2 becomes the *between* effect minus the *within* effect (substituting $\beta_1(x_{it} - \bar{x}_i) + \beta_2\bar{x}_i = \beta_1 x_{it} + (\beta_2 - \beta_1)\bar{x}_i$ shows this directly). This is called “contextual model”, β_2 is the “contextual” effect. See [Neuhaus and Kalbfleisch \(1998\)](#). In this specification, β_2 is actually similar to a Hausman test. It shows the difference between “between” and “within”.

One advantage of BW model is that it can incorporate fixed effect models along with a random effect estimation, thus including time-invariant covariates becomes possible. A second advantage is that it can do more complicated models, such as cross-level interactions, random slopes, or other multi-level models.

The actual implementation of the simplest form of BW is easy: simply use random effect models on the above two equations.

7.4. BW model in R

R has a package [panelr](#) that implements various kinds of BW models. Let's see an example.

```
library(panelr)
data("WageData")
wages <- panel_data(WageData, id = id, wave = t)
model1 <- wbm(lwage ~ wks + union + ms + occ | blk + fem, data = wages)
summary(model1)
```


MODEL INFO:

Entities: 595

Time periods: 1-7

Dependent variable: lwage

Model type: Linear mixed effects

Specification: within-between

MODEL FIT:

AIC = 2036.78, BIC = 2119.13

Pseudo-R² (fixed effects) = 0.27Pseudo-R² (total) = 0.69

Entity ICC = 0.57

WITHIN EFFECTS:

	Est.	S.E.	t val.	p
wks	0.00	0.00	1.06	0.29
union	0.06	0.03	2.53	0.01
ms	-0.08	0.03	-2.57	0.01
occ	-0.08	0.02	-3.32	0.00

BETWEEN EFFECTS:

	Est.	S.E.	t val.	p
(Intercept)	6.30	0.20	30.85	0.00
imean(wks)	0.01	0.00	2.25	0.02
imean(union)	0.15	0.03	4.67	0.00
imean(ms)	0.17	0.05	3.07	0.00
imean(occ)	-0.41	0.03	-13.31	0.00
blk	-0.15	0.05	-2.81	0.00
fem	-0.32	0.06	-4.96	0.00

p values calculated using df = 4153

RANDOM EFFECTS:

Group	Parameter	Std. Dev.
id	(Intercept)	0.2992
Residual		0.2589

Let's compare this with another popular package "lfe".

7. Fixed or Random Effect, or Both?

```
library(lfe)
model2 <- felm(lwage ~ wks + union + ms + occ | id, data = wages)
summary(model2)
```

Call:

```
felm(formula = lwage ~ wks + union + ms + occ | id, data = wages)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.89500	-0.16174	0.00652	0.17060	1.94521

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
wks	0.001083	0.001019	1.063	0.287816
union	0.064320	0.025378	2.534	0.011305 *
ms	-0.082905	0.032226	-2.573	0.010132 *
occ	-0.077507	0.023359	-3.318	0.000916 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2589 on 3566 degrees of freedom

Multiple R-squared(full model): 0.7304 Adjusted R-squared: 0.6852

Multiple R-squared(proj model): 0.006509 Adjusted R-squared: -0.1601

F-statistic(full model):16.16 on 598 and 3566 DF, p-value: < 2.2e-16

F-statistic(proj model): 5.841 on 4 and 3566 DF, p-value: 0.0001106

We can see these two give the same fixed effect estimation — identically so, to machine precision, because `WageData` is a balanced panel (7 waves for every one of the 595 individuals). “panelr” in addition estimates the effect of “blk” and “fem” which are time-invariant. But “lfe” has an advantage, it allows you to estimate fixed effect with clustered standard errors, which I wish “panelr” can do too.

```
model3 <- felm(lwage ~ wks + union + ms + occ | id | 0 | id, data = wages)
summary(model3)
```

Call:

```
felm(formula = lwage ~ wks + union + ms + occ | id | 0 | id, data = wages)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.89500	-0.16174	0.00652	0.17060	1.94521

Coefficients:

	Estimate	Cluster s.e.	t value	Pr(> t)
--	----------	--------------	---------	----------

```

wks      0.001083      0.001331      0.814      0.4160
union    0.064320      0.040936      1.571      0.1167
ms       -0.082905      0.047399     -1.749      0.0808 .
occ      -0.077507      0.031320     -2.475      0.0136 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2589 on 3566 degrees of freedom
Multiple R-squared(full model): 0.7304   Adjusted R-squared: 0.6852
Multiple R-squared(proj model): 0.006509   Adjusted R-squared: -0.1601
F-statistic(full model, *iid*):16.16 on 598 and 3566 DF, p-value: < 2.2e-16
F-statistic(proj model): 3.456 on 4 and 594 DF, p-value: 0.008358

```

7.5. BW model in Stata

In stata, there is no package to do BW estimator. But we can do it with “xtreg”.

```

webuse nlswork
xtset idcode
xtreg ln_w age, fe cluster(idcode)

```

(National Longitudinal Survey of Young Women, 14-24 years old in 1968)

Panel variable: idcode (unbalanced)

```

Fixed-effects (within) regression      Number of obs   =    28,510
Group variable: idcode                 Number of groups =     4,710

R-squared:                             Obs per group:
    Within = 0.1026                      min =          1
    Between = 0.0877                     avg =         6.1
    Overall = 0.0774                     max =        15

                                F(1, 4709)      =    884.05
corr(u_i, Xb) = 0.0314              Prob > F       =    0.0000

```

(Std. err. adjusted for 4,710 clusters in idcode)

		Robust				
ln_wage	Coefficient	std. err.	t	P> t	[95% conf. interval]	
age	.0181349	.0006099	29.73	0.000	.0169392	.0193306
_cons	1.148214	.0177153	64.81	0.000	1.113483	1.182944

7. Fixed or Random Effect, or Both?

```
sigma_u | .40635023
sigma_e | .30349389
rho     | .64192015   (fraction of variance due to u_i)
```

We then generate the mean of age and run a BW estimation.

```
webuse nlswork
xtset idcode
bysort idcode: center age, prefix(d) mean(m)
xtreg ln_w dage mage i.race, re cluster(idcode)
```

(National Longitudinal Survey of Young Women, 14-24 years old in 1968)

Panel variable: idcode (unbalanced)

(generated variables: dage mage)

Random-effects GLS regression	Number of obs	=	28,510
Group variable: idcode	Number of groups	=	4,710
R-squared:	Obs per group:		
Within = 0.1026	min =		1
Between = 0.1040	avg =		6.1
Overall = 0.0950	max =		15
	Wald chi2(4)	=	1335.89
corr(u_i, X) = 0 (assumed)	Prob > chi2	=	0.0000

(Std. err. adjusted for 4,710 clusters in idcode)

ln_wage	Coefficient	Robust std. err.	z	P> z	[95% conf. interval]	
dage	.0181349	.00061	29.73	0.000	.0169394	.0193304
mage	.022558	.0011405	19.78	0.000	.0203226	.0247933
race						
Black	-.1190246	.0127418	-9.34	0.000	-.1439982	-.0940511
Other	.0974996	.0617364	1.58	0.114	-.0235016	.2185008
_cons	1.037566	.0323185	32.10	0.000	.9742233	1.10091
sigma_u	.36581005					
sigma_e	.30349389					

```
rho | .59230575 (fraction of variance due to u_i)
```

In this BW model, we use the centered (within-group-demeaned) `dage` together with the group mean `mage`, which gives the standard within-between (Mundlak) decomposition: the coefficient on `dage`, .0181, is exactly the fixed-effect (“within”) coefficient on age, and the coefficient on `mage`, .0226, is directly the between effect. And we have the effect of time-invariant covariate race estimated. The advantage of using `xtreg` is that we have clustered standard errors implemented.

Note: if we instead ran `xtreg ln_w age mage i.race, re cluster(idcode)` — using the raw (uncentered) `age` together with `mage` — we would get the *contextual model* instead: the coefficient on `age` stays .0181, but the coefficient on `mage` becomes .0044, the “contextual effect” (the *additional* between-effect on top of the within effect). The two models are algebraically related: contextual-model `mage` = Mundlak-model `mage` – Mundlak-model `dage` (here, .0226 – .0181 = .0044).

7.6. BW model in non-linear models

Paul Allison in [his blog](#) mentioned using BW model for a binary outcome. I have not dig into the literature to see how large the bias can be using the BW, comparing to, say a conditional logit model. But if OLS is a good linear approximation of a logit model, BW model could be a good approximation with a binary outcome with panel data.

8. Mundlak device in fixed effect models

8.1. TWFE

Two-way fixed effect model has been widely used, especially in difference-in-difference situation, when we have a panel data with pre and post data. The basic setup is:

$$y_{it} = X_{it}\beta + c_i + f_t + u_{it} \quad (8.1)$$

with X has k variables, t has T periods, and i has N units.

For a linear model, we know we can put in k unit dummies and T time dummies to estimate the model. When we have a lot of dummies, this becomes inefficient, sometimes infeasible. When we have nonlinear models, such as Poisson, or logit, then this becomes even more difficult.

8.2. TWM

The recent Wooldridge paper (https://www.researchgate.net/publication/353938385_Two-Way_Fixed_Effects_the_Two-Way_Mundlak_Regression_and_Difference-in-Differences_Estimators) showed that two-way Mundlak regression (TWM) can be useful, not only in the common event time situation, but also in staggered TWFE situation. But in this post I'll only talk about common event time situation (which is the case that all treated units get treated at the same period, and stay treated afterwards). The staggered situation (treated units get treated at different periods) will be a different post.

The idea of Mundlak (1978) is the following.

In the one-way fixed effect case,

$$y_{it} = x_{it}\beta + c_i + u_{it} \quad (8.2)$$

Mundlak models c_i as a function of the mean of x_{it} for each unit:

$$c_i = \bar{x}_i\theta + \eta_i \quad (8.3)$$

Then

$$y_{it} = x_{it}\beta + \bar{x}_i\theta + \eta_i + u_{it} \quad (8.4)$$

8. Mundlak device in fixed effect models

This $\eta_i + u_{it}$ becomes the composite error term. This becomes a model including $\bar{x}_i$'s, which are additional k variables, instead of including a set of unit dummies. This model can be estimated by either Pooled OLS (POLS), or random effect. The estimated coefficients β are identical. The standard errors are the same too, after clustering on units. In other words, if we include means of x 's as controls, then POLS and random effect models are identical.

The by-product of this equation is that if $\theta = 0$, then c_i is uncorrelated with x_{it} ; therefore a random effect is good. Otherwise we should do fixed effect. This is the Mundlak form of the Hausman test: under the random-effects assumptions, testing $\theta = 0$ is asymptotically equivalent to the classical Hausman test, and it has the practical advantage of working directly with cluster-robust standard errors, where the classical form can break down.

Here is an example:

```
clear
webuse nlswork
xtset idcode year
drop if union==.
drop if age==.
bysort idcode: egen mean_age=mean(age)
bysort idcode: egen mean_union=mean(union)
xtreg ln_wage age union, fe cluster(idcode)
reg ln_wage age union mean_age mean_union, cluster(idcode)
xtreg ln_wage age union mean_age mean_union, re cluster(idcode)
```

(National Longitudinal Survey of Young Women, 14-24 years old in 1968)

Panel variable: idcode (unbalanced)
Time variable: year, 68 to 88, but with gaps
Delta: 1 unit

(9,296 observations deleted)

(9 observations deleted)

Fixed-effects (within) regression	Number of obs	=	19,229
Group variable: idcode	Number of groups	=	4,150
R-squared:	Obs per group:		
Within = 0.0963	min =		1
Between = 0.0433	avg =		4.6
Overall = 0.0562	max =		12
	F(2, 4149)	=	331.52

corr(u_i, Xb) = 0.0127 Prob > F = 0.0000

(Std. err. adjusted for 4,150 clusters in idcode)

		Robust				
ln_wage	Coefficient	std. err.	t	P> t	[95% conf. interval]	
age	.0153507	.0006912	22.21	0.000	.0139956	.0167058
union	.1055274	.0098582	10.70	0.000	.0862001	.1248546
_cons	1.248435	.0215661	57.89	0.000	1.206154	1.290716
sigma_u	.42353003					
sigma_e	.26213464					
rho	.72302816	(fraction of variance due to u_i)				

Linear regression

Number of obs = 19,229
 F(4, 4149) = 234.77
 Prob > F = 0.0000
 R-squared = 0.0769
 Root MSE = .44953

(Std. err. adjusted for 4,150 clusters in idcode)

		Robust				
ln_wage	Coefficient	std. err.	t	P> t	[95% conf. interval]	
age	.0153507	.0006912	22.21	0.000	.0139956	.0167059
union	.1055274	.0098587	10.70	0.000	.0861991	.1248556
mean_age	-.0064208	.001548	-4.15	0.000	-.0094556	-.0033859
mean_union	.1900785	.0223172	8.52	0.000	.1463249	.2338321
_cons	1.405264	.042734	32.88	0.000	1.321482	1.489046

Random-effects GLS regression

Group variable: idcode

Number of obs = 19,229
 Number of groups = 4,150

R-squared:

Within = 0.0963
 Between = 0.0750
 Overall = 0.0768

Obs per group:

min = 1
 avg = 4.6
 max = 12

corr(u_i, X) = 0 (assumed)

Wald chi2(4) = 1023.99
 Prob > chi2 = 0.0000

(Std. err. adjusted for 4,150 clusters in idcode)

8. Mundlak device in fixed effect models

		Robust				
ln_wage	Coefficient	std. err.	z	P> z	[95% conf. interval]	
age	.0153507	.0006912	22.21	0.000	.0139959	.0167055
union	.1055274	.0098587	10.70	0.000	.0862047	.12485
mean_age	-.0061464	.0014254	-4.31	0.000	-.0089401	-.0033527
mean_union	.212183	.0210098	10.10	0.000	.1710046	.2533614
_cons	1.362101	.0383784	35.49	0.000	1.286881	1.437322
sigma_u	.38505558					
sigma_e	.26213464					
rho	.68331727	(fraction of variance due to u_i)				

Clearly we see in the these three models, the coefficients and standard errors on age and union are the same.

What Wooldridge (2021) shows is that this (Mundlak) works for two-way fixed effect too.

That is, for the TWFE,

$$y_{it} = X_{it}\beta + c_i + f_t + u_{it} \quad (8.5)$$

it's equivalent to estimate

$$y_{it} = X_{it}\beta + X_{i.}\lambda + X_{.t}\theta + u_{it} \quad (8.6)$$

where $X_{i.}$ is the unit specific averages over time, and $X_{.t}$ is the time specific averages over units. This is basically some kind of control function approach. That is, from FWL theorem, these two models are equivalent. See Wooldridge 2021 for proof.

Wooldridge calls this POLS of regressing y_{it} on regressors X_{it} and two sets of additional controls $X_{i.}$ and $X_{.t}$ Two-way Mundlak (TWM) model.

One advantage of TWM is that it has potentially huge cut on the number of regressors, if you use the dummy variable approach for the TWFE.

The other advantage is that you can actually include time-invariant or unit-invariant variables in the regression and the β 's remain unchanged. We all know that in TWFE model you cannot include those time-invariant or unit-invariant variables in the model.

We can see in the following example of POLS (which is TWM), we add in **race**, which is **time-invariant** — constant within a person, but varying across people. (In **nlswork** the within-person standard deviation of **race** is exactly 0, while within any given year it is about 0.50, so it is emphatically not unit-invariant.) It can be estimated in TWM, but would not be estimable in TWFE. Notice that coefficients on age and union remain unchanged.

The reason is *not* that **race** is orthogonal to the other regressors — it is not orthogonal to the unit means **mean_age** and **mean_union**. It is that the variation identifying β is the within-unit variation, and within-unit deviations are orthogonal to **any** unit-level variable by construction, since $\sum_t (x_{it} - \bar{x}_i) = 0$ for every i . Adding a unit-level regressor therefore cannot move β , whatever that

8.3. An important qualification: TWM needs a balanced panel

regressor happens to be. A genuinely unit-invariant variable — one constant across units within a period, such as a time dummy — leaves β unchanged for the mirror-image reason.

```
reg ln_wage age union mean_age mean_union i.race, cluster(idcode)
```

Linear regression	Number of obs	=	19,229
	F(6, 4149)	=	189.66
	Prob > F	=	0.0000
	R-squared	=	0.1069
	Root MSE	=	.4422

(Std. err. adjusted for 4,150 clusters in idcode)

		Robust				
ln_wage	Coefficient	std. err.	t	P> t	[95% conf. interval]	
age	.0153507	.0006913	22.21	0.000	.0139955	.0167059
union	.1055274	.0098592	10.70	0.000	.0861981	.1248566
mean_age	-.0070345	.0015342	-4.59	0.000	-.0100424	-.0040266
mean_union	.2227511	.0220026	10.12	0.000	.1796142	.265888
race						
Black	-.1790995	.0143624	-12.47	0.000	-.2072575	-.1509415
Other	.0882461	.0715315	1.23	0.217	-.0519939	.2284861
_cons	1.466452	.0426087	34.42	0.000	1.382916	1.549988

8.3. An important qualification: TWM needs a balanced panel

The two-way equivalence above holds for a **balanced** panel. It is worth stating this explicitly, because the one-way result does *not* need balance and the contrast is easy to miss.

The one-way version — regressing y_{it} on X_{it} and $X_{i\cdot}$ — is exact whatever the panel looks like, for the reason given below: $\sum_t (x_{it} - \bar{x}_i) = 0$ holds within each unit by construction, no matter how many periods that unit contributes.

The two-way version needs something stronger. Subtracting both margins and adding back the grand mean, $x_{it} - \bar{x}_i - \bar{x}_t + \bar{x}$, is the projection onto *both* sets of dummies only when the unit and time factors are **orthogonal** — that is, when every unit appears in every period. In a balanced panel that holds automatically. In an unbalanced panel it fails, and TWM is then simply a different estimator from TWFE.

The difference is large enough to matter in applied work. Using `nlswork`, which is unbalanced (persons contribute between 1 and 12 observations), the coefficient on **age** differs by more than a factor of two:

8. Mundlak device in fixed effect models

```
* preserve/restore so this illustration leaves the data as later chunks expect
preserve
webuse nlswork, clear
keep if !missing(ln_wage, age, union)
xtset idcode year

bysort idcode: egen m_age    = mean(age)
bysort idcode: egen m_union = mean(union)
egen t_age    = mean(age),    by(year)
egen t_union  = mean(union), by(year)

* unit means of the year dummies -- the terms the unbalanced case needs
levelsof year, local(yrs)
foreach y of local yrs {
    bysort idcode: egen my_`y' = mean(year==`y')
}

quietly xtreg ln_wage age union i.year, fe
scalar b_twfe = _b[age]
quietly regress ln_wage age union m_age m_union t_age t_union
scalar b_twm = _b[age]
quietly regress ln_wage age union m_age m_union i.year my_*
scalar b_fix = _b[age]

di "TWFE                                age = " %9.7f b_twfe
di "TWM (unit + time means)           age = " %9.7f b_twm
di "TWM + unit means of dummies       = " %9.7f b_fix

restore
```

(National Longitudinal Survey of Young Women, 14-24 years old in 1968)

(9,305 observations deleted)

Panel variable: idcode (unbalanced)
Time variable: year, 70 to 88, but with gaps
Delta: 1 unit

70 71 72 73 77 78 80 82 83 85 87 88

8.4. More than two fixed effects: one set of means, not one per dimension

```
TWFE                                age = 0.0276645

TWM (unit + time means)            age = 0.0117968

TWM + unit means of dummies        = 0.0276646
```

What restores the equivalence is including the **unit means of the time dummies**. They have a clear meaning: the unit mean of the dummy for period t is s_{it}/T_i , where s_{it} indicates whether unit i is observed in period t . So adding them amounts to conditioning on *which* periods each unit is observed in — the panel’s selection pattern. This is the conditioning set in Wooldridge’s work on correlated random effects with unbalanced panels.

Two practical notes. In a balanced panel these extra terms have zero variance — every unit is observed in every period, so each mean is $1/T$ for everybody — which is exactly why they never come up in the balanced case. And the number of observations per unit, T_i , is *not* a substitute for them: it does not restore exactness, and it is collinear with them once they are included.

Everything in this section is about the *point estimate*. Standard errors from TWM are not the TWFE standard errors, and should not be read as such.

It is worth pausing on the fact that this chapter now demonstrates both results on the *same* data. The `xtreg` example near the top shows one-way Mundlak matching one-way FE on `nlswork`, and the chunk just above shows two-way Mundlak failing on `nlswork`. Same dataset, opposite verdicts, because the two claims do not have the same requirements. The one-way result needs nothing: $\sum_t (x_{it} - \bar{x}_i) = 0$ holds within each unit whatever periods that unit contributes, so T_i may vary freely. The two-way result needs the unit and time factors to be orthogonal *to each other*, which is a statement about the relationship between two dimensions, and that is what imbalance destroys. Nothing one-way can be affected, because there is no second dimension for the first to fail to be orthogonal to.

Which regressors need means, then? Those that vary within unit. This is why `i.race` above takes no mean: `race` is constant within a person, so its within-person mean *is* `race` — in `nlswork` the two agree to the last digit — and adding it would be perfectly collinear. A year dummy is the opposite case: it does vary within a person (in `nlswork` the within-person standard deviation of the 1971 dummy is about 0.71), so its unit mean is not the dummy but the fraction of that person’s observations falling in that year. That fraction is new information, and in an unbalanced panel it is exactly the information the equivalence needs.

8.4. More than two fixed effects: one set of means, not one per dimension

The natural guess, having seen TWM, is that three fixed effects call for three sets of means — $X_{i\cdot}$, $X_{\cdot t}$, and the averages over the third factor. That guess is wrong in general, and the reason is worth

8. Mundlak device in fixed effect models

spelling out because it also explains the balanced-panel condition above.

The Mundlak device deals with **one** unobserved effect at a time. It replaces c_i by its projection on the within- i means of the other regressors. So the number of sets of means you need is the number of dimensions you are *eliminating*, not the number of dimensions you have. With three fixed effects the robust recipe is:

- pick the dimension you cannot estimate — typically the one with many levels and few observations each — and replace it with means;
- **keep the other dimensions as dummies**, since they have few levels and many observations each;
- take the means over the chosen dimension of *everything else in the model*, the retained dummies included.

That last clause is what does the work, and it is the same clause that fixes the unbalanced two-way case above. The test for whether a given regressor needs a mean is the one stated there: does it vary within the dimension you are eliminating? Time and group dummies do, so their means carry information; a regressor already constant within unit does not, so its mean is itself. Call this construction **B**, and call the three-sets-of-means guess **A**.

```
preserve
capture program drop twm3
program define twm3
    args label
    quietly {
        capture drop x_i x_t x_g mt_* mg_* x y ft hg
        sort t
        by t: gen double ft = rnormal() if _n==1
        by t: replace ft = ft[1]
        sort g
        by g: gen double hg = rnormal() if _n==1
        by g: replace hg = hg[1]
        gen double x = 0.6*ci + 0.5*ft + 0.4*hg + rnormal()
        gen double y = 1*x + ci + ft + hg + rnormal()
        bysort i: egen double x_i = mean(x)
        bysort t: egen double x_t = mean(x)
        bysort g: egen double x_g = mean(x)
        levelsof t, local(ts)
        foreach k of local ts {
            bysort i: egen double mt_`k' = mean(t==`k')
        }
        levelsof g, local(g)
        foreach k of local g {
            bysort i: egen double mg_`k' = mean(g==`k')
        }
        areg y x i.t i.g, absorb(i)
        scalar b_fe = _b[x]
        regress y x x_i x_t x_g
        scalar b_A = _b[x]
```

8.4. More than two fixed effects: one set of means, not one per dimension

```

    regress y x x_i i.t i.g mt_* mg_*
    scalar b_B = _b[x]
}
di "`label'"
di "    three-way FE                        = " %9.6f b_fe
di "    A: one set of means per dimension  = " %9.6f b_A   "    gap " %9.6f (b_A-b_fe)
di "    B: means over i, dummies included = " %9.6f b_B   "    gap " %9.6f (b_B-b_fe)
end

clear
set seed 20260731
set obs 300
gen i = _n
gen double ci = rnormal()
expand 10
bysort i: gen t = _n

* third factor assigned at random: the three factors are not orthogonal
gen g = 1 + int(runiform()*5)
twm3 "g assigned at random (NOT orthogonal):"

* Latin square: every unit sees every g equally, and so does every period
replace g = mod(i+t,5)+1
twm3 "g from a Latin square (orthogonal):"

capture program drop twm3
restore

```

Number of observations (_N) was 0, now 300.

(2,700 observations created)

```

g assigned at random (NOT orthogonal):
    three-way FE                        = 0.994424
    A: one set of means per dimension  = 1.007648   gap 0.013224
    B: means over i, dummies included  = 0.994424   gap 0.000000

```

(2,435 real changes made)

```

g from a Latin square (orthogonal):
    three-way FE                        = 0.996471
    A: one set of means per dimension  = 0.996471   gap 0.000000
    B: means over i, dummies included  = 0.996471   gap 0.000000

```

So **A is not incorrect but conditional**: it is exact when the factors are mutually orthogonal, and inexact otherwise. A full factorial design gives orthogonality, but so does a Latin square, which fills only a fraction of the possible cells — what matters is the balance property, that each level of one factor appears equally often with each level of the others, not the number of cells. Deleting a single row from either design is enough to lose exactness, and the size of the gap tracks how far from orthogonal the design is.

This is the general statement of the balanced-panel condition from the previous section. A balanced two-way panel *is* an orthogonal design: every unit appears in every period. Unbalance it, and the unit and time factors stop being orthogonal, which is why the two-way equivalence fails there.

The practical reading is that orthogonality is a property of **designed** data — factorial experiments, Latin squares, crop trials, where somebody chose the assignment. It requires each level of one factor to appear equally often with each level of the others, which is not something that happens by accident in data you did not design. A third factor tends to make matters worse rather than better, because it is frequently determined in part by the first two. In an audit panel I looked at recently, for instance, the third fixed effect was the running count of audits within a factory, which is mechanically a function of factory and calendar time; that design cannot be orthogonal even in principle. For observational work, B is the construction to use.

8.5. Applications of TWM to DID in common treatment timing

There is an extension of Mundlak device that if x_{it} can be seen as an interaction between a time-invariant variable and an unit-invariant variable,

$$x_{itj} = z_{ij} * m_{tj} \quad (8.7)$$

Then

$$\bar{x}_{i \cdot j} = z_{ij} \bar{m}_j \quad (8.8)$$

$$\bar{x}_{\cdot t j} = \bar{z}_j m_{tj} \quad (8.9)$$

In other words, the two averages of x are constant times z_{ij} and constant times m_{tj} respectively. By Mundlak, we can just include z_{ij} and m_{tj} as control in addition to x_{it} to get the same estimates as TWFE. This makes things much easier, as we'll see later.

Wooldridge 2021 goes on showing that TWM works in common treatment time, and also in staggered timing setting, as long as we allow flexible interactions of treatment cohorts and treatment periods. It's truly remarkable that he showed that TWM can be so powerful! In this post I am going to just use simulation to show TWM in common treatment timing; staggered timing setting will be another post. All the codes are based on Wooldridge's do files he shared in his dropbox.

In a common timing setting, treatment comes in at one time, all the treated units are treated at the same time. There is the parallel trend assumption that assumes that the treatment group and control group have the same trend before treatment. Also there is no anticipation assumption.

8.5. Applications of TWM to DID in common treatment timing

$$y_{it} = w_{it}\beta + c_i + f_t + u_{it} \quad (8.10)$$

where w_{it} is the treatment indicator, with T periods and treatment happens at $t = q$.

$$w_{it} = d_i \cdot p_t \quad (8.11)$$

where d_i is the treatment group dummy, and p_t is the post treatment dummy.

$$\bar{w}_{i\cdot} = d_i \cdot \bar{p} \quad (8.12)$$

$$\bar{w}_{\cdot t} = \bar{d} \cdot p_t \quad (8.13)$$

Therefore, we just need to include d_i and p_t as controls in the regression to be the same estimator as TWFE, by the spirit of Mundlak.

That is, instead of doing TWFE, we can do TWM by regressing y on w , d_i and p_t .

Moreover, in a TWM, we can include any time-invariant or unit-invariant controls, or both and have no effect on β .

We can also allow the effect to change across time post treatment by doing

$$y_{it} = \alpha + \beta_q(w_{it} \cdot f_q) + \dots + \beta_T(w_{it} \cdot f_T) + \eta d_i + \theta_q f_q + \dots + \theta_T f_T + e_{it} \quad (8.14)$$

In the TWM, we allow effects change from $t = q$ to $t = T$, by interacting w_{it} and f_t , which is an dummy for time t . Very flexible, yet easy to do, given everything is linear.

Even more, if we have another variable, x_i , say gender, which is time-invariant, we can allow the effect to be different across gender, across time!

Now I use Wooldridge's code to run on a simulated data set.

```
clear
set obs 100
set seed 1234567
gen id = _n
gen w1 = rnormal()
gen w2 = rnormal()
expand 4
bysort id: gen year=2011+_n
gen post=(year>2013)
gen d=(id>50)
gen w=d*post
gen x1 = rnormal() + w1 + 2*w2
gen u = rnormal()

gen f2014=(year==2014)
gen f2015=(year==2015)
```

8. Mundlak device in fixed effect models

```
gen y1 = 1 + 2*x1 + 3*w1 + u
* let's allow each year's treatment effect being different. ATT for post period would be rough
gen logy = y1 + 2.5*d*f2014 + 3.5*d*f2015

xtset id year
tab year w
```

Number of observations (_N) was 0, now 100.

(300 observations created)

Panel variable: id (strongly balanced)
Time variable: year, 2012 to 2015
Delta: 1 unit

		w		
year		0	1	Total
2012		100	0	100
2013		100	0	100
2014		50	50	100
2015		50	50	100
Total		300	100	400

In this data set, we have 100 units, four years (2012 to 2015). Half units are treated, treatment happens on 2014. Treatment dummy d , w is $d * post$. Treatment effect is 2.5 in year 2014, and 3.5 in year 2015. If we estimate ATT for the entire post treatment period, we should expect to see ATT about 3, given the balanced panel.

All the following code are from Wooldridge. I just used it on the simulated data.

The difficulty is that we have to be careful and mindful about what we want to estimate, then come up with the correct code for all those interaction terms.

```
* CACHE WARNING: this chunk depends on the data built by the earlier
* `collectcode: true` chunks. If those are served from the knitr cache while
* this one re-executes, their code never enters Statamarkdown's accumulated
* buffer, every command below fails on missing data, and the chunk silently
* renders with NO output. If that happens, delete the whole
* `mundlak-device_cache/` directory (not just this chunk's entry) and
* `_freeze/mundlak-device/`, then re-render so all chunks run in order.
```

```
* Basic DID regression. Only need the post-treatment dummy but including
* all time dummies is harmless (and necessary for heterogeneous trends).
* Standard errors change a little because of different degrees-of-freedom.
```

```
reg logy c.d#c.post d post, vce(cluster id)
reg logy w d post, vce(cluster id)
reg logy w d f2014 f2015, vce(cluster id)
reg logy w d i.year, vce(cluster id)
```

```
* Using fixed effects is equivalent:
```

```
xtreg logy w i.year, fe vce(cluster id)
xtreg logy w post, fe vce(cluster id)
```

```
* So is RE with d included:
```

```
xtreg logy w d post, re vce(cluster id)
```

```
* Allow a separate effect in each of the treated time periods:
```

```
* It is the ATT for each treated period:
```

```
xtreg logy c.w#c.f2014 c.w#c.f2015 i.year, fe vce(cluster id)
xtreg logy c.d#c.f2014 c.d#c.f2015 i.year, fe vce(cluster id)
```

```
* POLS version:
```

```
reg logy c.w#c.f2014 c.w#c.f2015 d f2014 f2015, vce(cluster id)
lincom (c.w#c.f2014 + c.w#c.f2015)/2
```

```
* RE still gives same estimates:
```

```
xtreg logy c.w#c.f2014 c.w#c.f2015 d f2014 f2015, re vce(cluster id)
```

```
* Putting in time-constant controls in the levels and even interacted with d
* does not change the estimates. It does boost the R-squared and
* slightly changes the standard errors:
```

8. Mundlak device in fixed effect models

```
reg logy c.w#c.f2014 c.w#c.f2015 d f2014 f2015 x1 c.d#c.x1, vce(cluster id)
reg logy c.w#c.f2014 c.w#c.f2015 d i.year x1 c.d#c.x1, vce(cluster id)

* Now add covariate interacted with everything:

sum x1 if d
gen x1_dm_1 = x1 - r(mean)

xtreg logy c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1_dm_1 c.w#c.f2015#c.x1_dm_1 i.year c.f2014#c.x1 c.f2015#c.x1, vce(cluster id)

reg logy c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1_dm_1 c.w#c.f2015#c.x1_dm_1 i.year c.f2014#c.x1 c.f2015#c.x1, vce(cluster id)

reg logy c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1_dm_1 c.w#c.f2015#c.x1_dm_1 i.year i.year#c.x1, vce(cluster id)

* Now use margins to account for the sampling error in the mean of x1. It is
* important to have w defined as the time-varying treatment variable.
* In practice, it seems to have little effect:

reg logy c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1 c.w#c.f2015#c.x1 i.year c.f2014#c.x1 c.f2015#c.x1, vce(cluster id)
margins, dydx(w) at(f2014 = 1 f2015 = 0) subpop(if d == 1) vce(uncond)
margins, dydx(w) at(f2014 = 0 f2015 = 1) subpop(if d == 1) vce(uncond)

* Apply Callaway & Sant'Anna (2021):

gen first_treat = 0
replace first_treat = 2014 if d
csdid logy x1, ivar(id) time(year) gvar(first_treat)

* Show imputation is equivalent to POLS/ETWFE:

reg logy i.year c.f2014#c.x1 c.f2015#c.x1 d x1 c.d#c.x1 if w == 0
predict tetilda, resid
sum tetilda if (d & f2014)
sum tetilda if (d & f2015)

* Regression produces same ATT estimates, but use the full pooled regression for
* valid standard errors:

reg tetilda c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1_dm_1 c.w#c.f2015#c.x1_dm_1 if w == 1, nocoll

* Now test parallel trends. First, unconditionally.

xtddidreg (logy) (w), group(id) time(year)
* estat trendplots
estat ptrends

* With two control periods, same as the following:
```

8.5. Applications of TWM to DID in common treatment timing

```

* Note: since the pre-period only spans 2012-2013, D.logy is non-missing
* only for year==2013, so this is a single cross-section of first
* differences (2013 vs 2012), not a multi-period trend test - it's a
* valid two-period parallel-trends check, but should not be read as
* testing trends across more than two pre-periods.

sort id year
reg D.logy d if year <= 2013, vce(robust)

* Test/correct for common trends in a general equation:

gen t = year - 2011

xtreg logy c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1_dm_1 c.w#c.f2015#c.x1_dm_1 i.year i.year#c.x1
reg logy c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1_dm_1 c.w#c.f2015#c.x1_dm_1 i.year i.year#c.x1

* Allow the trend to also vary with x1:

reg logy c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1_dm_1 c.w#c.f2015#c.x1_dm_1 i.year i.year#c.x1
test c.d#c.t c.d#c.t#c.x1

* Imputation. Note that the estimates on c.d#c.t and c.d#c.t#c.x1
* are the same as full regression with all data, as is the joint
* F statistic:

drop tetilda

reg logy i.year i.year#c.x1 d x1 c.d#c.x1 c.d#c.t c.d#c.t#c.x1 if w == 0, vce(cluster id)
test c.d#c.t c.d#c.t#c.x1

predict tetilda, resid
sum tetilda if (d & f2014)
sum tetilda if (d & f2015)

* Again, can use regression to obtain the estimates but not the standard errors:

reg tetilda c.w#c.f2014 c.w#c.f2015 c.w#c.f2014#c.x1_dm_1 c.w#c.f2015#c.x1_dm_1 if w == 1, noc

```

Linear regression	Number of obs	=	400
	F(3, 99)	=	34.04
	Prob > F	=	0.0000
	R-squared	=	0.0348
	Root MSE	=	7.0626

(Std. err. adjusted for 100 clusters in id)

8. Mundlak device in fixed effect models

			Robust				
	logy	Coefficient	std. err.	t	P> t	[95% conf. interval]	
c.d#c.post		3.206857	.4314949	7.43	0.000	2.350677	4.063036
d		-.0636461	1.356731	-0.05	0.963	-2.755695	2.628403
post		-.1255431	.3021942	-0.42	0.679	-.725162	.4740758
_cons		.11305	.9272947	0.12	0.903	-1.726904	1.953004

Linear regression

Number of obs	=	400
F(3, 99)	=	34.04
Prob > F	=	0.0000
R-squared	=	0.0348
Root MSE	=	7.0626

(Std. err. adjusted for 100 clusters in id)

			Robust				
	logy	Coefficient	std. err.	t	P> t	[95% conf. interval]	
w		3.206857	.4314949	7.43	0.000	2.350677	4.063036
d		-.0636461	1.356731	-0.05	0.963	-2.755695	2.628403
post		-.1255431	.3021942	-0.42	0.679	-.725162	.4740758
_cons		.11305	.9272947	0.12	0.903	-1.726904	1.953004

Linear regression

Number of obs	=	400
F(4, 99)	=	27.19
Prob > F	=	0.0000
R-squared	=	0.0381
Root MSE	=	7.0597

(Std. err. adjusted for 100 clusters in id)

			Robust				
	logy	Coefficient	std. err.	t	P> t	[95% conf. interval]	
w		3.206857	.4320407	7.42	0.000	2.349594	4.064119
d		-.0636461	1.358447	-0.05	0.963	-2.7591	2.631808
f2014		-.7011363	.3419429	-2.05	0.043	-1.379625	-.0226474
f2015		.4500501	.3583234	1.26	0.212	-.2609413	1.161041
_cons		.11305	.9284678	0.12	0.903	-1.729232	1.955331

8.5. Applications of TWM to DID in common treatment timing

```
Linear regression                                Number of obs    =      400
                                                F(5, 99)         =      21.95
                                                Prob > F          =      0.0000
                                                R-squared         =      0.0383
                                                Root MSE         =      7.0677
```

(Std. err. adjusted for 100 clusters in id)

		Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
logy							
w		3.206857	.4325887	7.41	0.000	2.348507	4.065207
d		-.0636461	1.36017	-0.05	0.963	-2.762519	2.635227
year							
2013		.3260046	.3325494	0.98	0.329	-.3338456	.9858548
2014		-.538134	.3934578	-1.37	0.175	-1.31884	.2425716
2015		.6130524	.3830314	1.60	0.113	-.146965	1.37307
_cons		-.0499523	.9214237	-0.05	0.957	-1.878257	1.778352

```
Fixed-effects (within) regression                Number of obs    =      400
Group variable: id                             Number of groups =      100
```

```
R-squared:                                     Obs per group:
  Within = 0.2539                               min =           4
  Between = 0.0129                             avg  =          4.0
  Overall = 0.0383                               max  =           4
```

```
corr(u_i, Xb) = -0.0027                      F(4, 99)         =      27.29
                                                Prob > F          =      0.0000
```

(Std. err. adjusted for 100 clusters in id)

		Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
logy							
w		3.206857	.4320407	7.42	0.000	2.349594	4.064119
year							
2013		.3260046	.3321282	0.98	0.329	-.3330098	.985019
2014		-.538134	.3929594	-1.37	0.174	-1.317851	.2415828
2015		.6130524	.3825463	1.60	0.112	-.1460024	1.372107
_cons		-.0817754	.2003881	-0.41	0.684	-.4793889	.3158381

8. Mundlak device in fixed effect models

```

sigma_u | 6.7557992
sigma_e | 2.3305153
rho | .89365429 (fraction of variance due to u_i)
-----

Fixed-effects (within) regression      Number of obs   =      400
Group variable: id                    Number of groups =      100

R-squared:                            Obs per group:
    Within = 0.2207                      min =          4
    Between = 0.0129                     avg =         4.0
    Overall = 0.0348                     max =          4

                                         F(2, 99)        =      50.25
corr(u_i, Xb) = -0.0028                 Prob > F         =      0.0000

```

(Std. err. adjusted for 100 clusters in id)

```

-----
              |               Robust
              | Coefficient std. err.      t    P>|t|    [95% conf. interval]
-----+-----
      w |      3.206857   .4309511    7.44  0.000    2.351756    4.061957
    post |     -.1255431   .3018134   -0.42  0.678   -.7244064    .4733201
   _cons |      .0812269   .1077378    0.75  0.453   -.1325482    .295002
-----+-----
sigma_u | 6.7557992
sigma_e | 2.3738231
rho | .89010353 (fraction of variance due to u_i)
-----

```

```

Random-effects GLS regression      Number of obs   =      400
Group variable: id                 Number of groups =      100

R-squared:                          Obs per group:
    Within = 0.2207                      min =          4
    Between = 0.0129                     avg =         4.0
    Overall = 0.0348                     max =          4

                                         Wald chi2(3)     =     102.12
corr(u_i, X) = 0 (assumed)         Prob > chi2       =      0.0000

```

(Std. err. adjusted for 100 clusters in id)

```

-----
              |               Robust
              | Coefficient std. err.      z    P>|z|    [95% conf. interval]
-----+-----

```


8.5. Applications of TWM to DID in common treatment timing

w		3.206857	.4314949	7.43	0.000	2.361142	4.052571
d		-.0636461	1.356731	-0.05	0.963	-2.72279	2.595498
post		-.1255431	.3021942	-0.42	0.678	-.7178329	.4667467
_cons		.11305	.9272947	0.12	0.903	-1.704414	1.930514

sigma_u		6.685563					
sigma_e		2.3738231					
rho		.88804221	(fraction of variance due to u_i)				

Fixed-effects (within) regression	Number of obs	=	400
Group variable: id	Number of groups	=	100

R-squared:	Obs per group:
Within = 0.2575	min = 4
Between = 0.0129	avg = 4.0
Overall = 0.0387	max = 4

	F(5, 99)	=	21.98
corr(u_i, Xb) = -0.0027	Prob > F	=	0.0000

(Std. err. adjusted for 100 clusters in id)

		Robust					
	logy	Coefficient	std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014		2.813593	.5368268	5.24	0.000	1.748412	3.878774
c.w#c.f2015		3.600121	.5765523	6.24	0.000	2.456116	4.744126
year							
2013		.3260046	.3325494	0.98	0.329	-.3338456	.9858548
2014		-.341502	.4488412	-0.76	0.449	-1.2321	.5490964
2015		.4164204	.4300378	0.97	0.335	-.4368679	1.269709
_cons		-.0817754	.2006423	-0.41	0.684	-.4798932	.3163424
sigma_u		6.7557992					
sigma_e		2.3288408					
rho		.89379082	(fraction of variance due to u_i)				

Fixed-effects (within) regression	Number of obs	=	400
Group variable: id	Number of groups	=	100

R-squared:	Obs per group:
------------	----------------

8. Mundlak device in fixed effect models

```

Within   = 0.2575           min =           4
Between  = 0.0129           avg  =          4.0
Overall  = 0.0387           max  =           4

```

```

corr(u_i, Xb) = -0.0027      F(5, 99)      =      21.98
                               Prob > F      =      0.0000

```

(Std. err. adjusted for 100 clusters in id)

	logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.d#c.f2014		2.813593	.5368268	5.24	0.000	1.748412	3.878774
c.d#c.f2015		3.600121	.5765523	6.24	0.000	2.456116	4.744126
year							
2013		.3260046	.3325494	0.98	0.329	-.3338456	.9858548
2014		-.341502	.4488412	-0.76	0.449	-1.2321	.5490964
2015		.4164204	.4300378	0.97	0.335	-.4368679	1.269709
_cons		-.0817754	.2006423	-0.41	0.684	-.4798932	.3163424
sigma_u		6.7557992					
sigma_e		2.3288408					
rho		.89379082	(fraction of variance due to u_i)				

```

Linear regression      Number of obs   =      400
                        F(5, 99)       =      21.97
                        Prob > F       =      0.0000
                        R-squared      =      0.0384
                        Root MSE     =      7.0672

```

(Std. err. adjusted for 100 clusters in id)

	logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014		2.813593	.5368268	5.24	0.000	1.748412	3.878774
c.w#c.f2015		3.600121	.5765523	6.24	0.000	2.456116	4.744126
d		-.0636461	1.36017	-0.05	0.963	-2.762519	2.635227
f2014		-.5045043	.3923088	-1.29	0.201	-1.28293	.2739216
f2015		.2534181	.4206018	0.60	0.548	-.5811472	1.087983
_cons		.11305	.9296453	0.12	0.903	-1.731568	1.957668

8.5. Applications of TWM to DID in common treatment timing

```
( 1) .5*c.w#c.f2014 + .5*c.w#c.f2015 = 0
```

logy	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
(1)	3.206857	.4325887	7.41	0.000	2.348507	4.065207

Random-effects GLS regression Number of obs = 400
Group variable: id Number of groups = 100

R-squared: Obs per group:

Within	= 0.2550	min	=	4
Between	= 0.0129	avg	=	4.0
Overall	= 0.0384	max	=	4

corr(u_i, X) = 0 (assumed) Wald chi2(5) = 109.84
 Prob > chi2 = 0.0000

(Std. err. adjusted for 100 clusters in id)

logy	Coefficient	Robust std. err.	z	P> z	[95% conf. interval]	
c.w#c.f2014	2.813593	.5368268	5.24	0.000	1.761432	3.865754
c.w#c.f2015	3.600121	.5765523	6.24	0.000	2.470099	4.730143
d	-.0636461	1.36017	-0.05	0.963	-2.729531	2.602239
f2014	-.5045043	.3923088	-1.29	0.198	-1.273415	.2644069
f2015	.2534181	.4206018	0.60	0.547	-.5709463	1.077782
_cons	.11305	.9296453	0.12	0.903	-1.709021	1.935121
sigma_u	6.6895238					
sigma_e	2.3287614					
rho	.89191109	(fraction of variance due to u_i)				

Linear regression Number of obs = 400
 F(7, 99) = 182.67
 Prob > F = 0.0000
 R-squared = 0.8438
 Root MSE = 2.8557

8. Mundlak device in fixed effect models

(Std. err. adjusted for 100 clusters in id)						
logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014	2.503103	.3066694	8.16	0.000	1.894605	3.111602
c.w#c.f2015	3.327274	.2902385	11.46	0.000	2.751378	3.90317
d	.881691	.5285779	1.67	0.098	-.1671223	1.930504
f2014	.0840872	.2215839	0.38	0.705	-.3555834	.5237578
f2015	.2391039	.2082552	1.15	0.254	-.1741196	.6523275
x1	2.506694	.1327049	18.89	0.000	2.243379	2.77001
c.d#c.x1	.0743163	.1915921	0.39	0.699	-.305844	.4544765
_cons	.2148708	.3676352	0.58	0.560	-.5145972	.9443389

Linear regression	Number of obs	=	400
	F(8, 99)	=	166.57
	Prob > F	=	0.0000
	R-squared	=	0.8438
	Root MSE	=	2.8593

(Std. err. adjusted for 100 clusters in id)						
logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014	2.503013	.3070762	8.15	0.000	1.893707	3.112319
c.w#c.f2015	3.327292	.2906341	11.45	0.000	2.750611	3.903973
d	.8816195	.5292507	1.67	0.099	-.1685288	1.931768
year						
2013	-.0324553	.1709047	-0.19	0.850	-.3715673	.3066568
2014	.067934	.2343565	0.29	0.773	-.3970801	.5329482
2015	.2228745	.2273177	0.98	0.329	-.2281732	.6739222
x1	2.507011	.1335155	18.78	0.000	2.242087	2.771935
c.d#c.x1	.0738548	.1927512	0.38	0.702	-.3086053	.4563149
_cons	.2311113	.37362	0.62	0.538	-.5102317	.9724544

8.5. Applications of TWM to DID in common treatment timing

Variable	Obs	Mean	Std. dev.	Min	Max	
x1	200	-.4048388	2.493319	-6.95061	4.588676	
Fixed-effects (within) regression						
Group variable: id				Number of obs	= 400	
				Number of groups	= 100	
R-squared:				Obs per group:		
Within = 0.3620				min	= 4	
Between = 0.4276				avg	= 4.0	
Overall = 0.2782				max	= 4	
				F(9, 99)	= 23.70	
corr(u_i, Xb) = 0.3305				Prob > F	= 0.0000	
(Std. err. adjusted for 100 clusters in id)						
logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014	2.938834	.5031086	5.84	0.000	1.940557	3.93711
c.w#c.f2015	3.742608	.534969	7.00	0.000	2.681113	4.804102
c.w#c.f2014#c.x1_dm_1	-.2243583	.1904085	-1.18	0.242	-.60217	.1534535
c.w#c.f2015#c.x1_dm_1	.0296616	.2431933	0.12	0.903	-.4528867	.5122099
year						
2013	.3260046	.3342504	0.98	0.332	-.3372208	.98923
2014	-.16839	.3880607	-0.43	0.665	-.9383866	.6016067
2015	.4360252	.3954162	1.10	0.273	-.3485664	1.220617
c.f2014#c.x1	.6285213	.1342955	4.68	0.000	.3620499	.8949926
c.f2015#c.x1	.5615944	.1919492	2.93	0.004	.1807255	.9424633
_cons	-.0817754	.1900787	-0.43	0.668	-.4589327	.2953819
sigma_u	6.1628203					
sigma_e	2.1734897					
rho	.88937776	(fraction of variance due to u i)				

8. Mundlak device in fixed effect models

```

-----
Linear regression                                Number of obs   =       400
                                                F(12, 99)      =      112.09
                                                Prob > F       =       0.0000
                                                R-squared     =       0.8445
                                                Root MSE     =       2.8679

```

(Std. err. adjusted for 100 clusters in id)

	logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014		2.496168	.3173129	7.87	0.000	1.86655	3.125785
c.w#c.f2015		3.349012	.3132217	10.69	0.000	2.727513	3.970512
c.w#c.f2014#							
c.x1_dm_1		-.2701343	.1759396	-1.54	0.128	-.6192367	.0789681
c.w#c.f2015#							
c.x1_dm_1		-.0683301	.1502763	-0.45	0.650	-.3665108	.2298506
year							
2013		-.0188111	.1738311	-0.11	0.914	-.3637298	.3261075
2014		.0787158	.2377676	0.33	0.741	-.3930668	.5504984
2015		.2325892	.2297992	1.01	0.314	-.2233823	.6885607
c.f2014#c.x1		.0404883	.1103324	0.37	0.714	-.1784352	.2594118
c.f2015#c.x1		.0778513	.1061898	0.73	0.465	-.1328523	.2885549
d		.904209	.5311615	1.70	0.092	-.1497307	1.958149
x1		2.476382	.1544607	16.03	0.000	2.169899	2.782866
c.d#c.x1		.1570952	.2113815	0.74	0.459	-.2623316	.576522
_cons		.2230452	.3747574	0.60	0.553	-.5205549	.9666452

```

-----
Linear regression                                Number of obs   =       400
                                                F(13, 99)      =      103.22
                                                Prob > F       =       0.0000
                                                R-squared     =       0.8445
                                                Root MSE     =       2.8715

```

8.5. Applications of TWM to DID in common treatment timing

(Std. err. adjusted for 100 clusters in id)

logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014	2.500641	.3196658	7.82	0.000	1.866354	3.134927
c.w#c.f2015	3.353485	.3142809	10.67	0.000	2.729884	3.977087
c.w#c.f2014# c.x1_dm_1	-.2674635	.177262	-1.51	0.135	-.6191897	.0842627
c.w#c.f2015# c.x1_dm_1	-.0656593	.1506332	-0.44	0.664	-.3645482	.2332296
year						
2013	-.0261778	.1756818	-0.15	0.882	-.3747686	.3224129
2014	.07109	.2439251	0.29	0.771	-.4129104	.5550904
2015	.2249634	.2325834	0.97	0.336	-.2365326	.6864593
year#c.x1						
2013	-.0298874	.0910807	-0.33	0.743	-.2106112	.1508364
2014	.0227905	.1348423	0.17	0.866	-.2447658	.2903469
2015	.0601535	.1180653	0.51	0.612	-.1741137	.2944207
d	.8986549	.5321658	1.69	0.094	-.1572776	1.954587
x1	2.49408	.1726029	14.45	0.000	2.151599	2.836562
c.d#c.x1	.1544244	.2128387	0.73	0.470	-.2678938	.5767426
_cons	.230671	.3794753	0.61	0.545	-.5222904	.9836324

Linear regression

Number of obs = 400
F(12, 99) = 112.09
Prob > F = 0.0000
R-squared = 0.8445
Root MSE = 2.8679

(Std. err. adjusted for 100 clusters in id)

logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014	2.386807	.3096756	7.71	0.000	1.772343	3.001271
c.w#c.f2015	3.32135	.2964989	11.20	0.000	2.733032	3.909668

8. Mundlak device in fixed effect models

c.w#c.f2014#							
c.x1		-.2701343	.1759396	-1.54	0.128	-.6192367	.0789681
c.w#c.f2015#							
c.x1		-.0683301	.1502763	-0.45	0.650	-.3665109	.2298506
year							
2013		-.0188111	.1738311	-0.11	0.914	-.3637298	.3261075
2014		.0787158	.2377676	0.33	0.741	-.3930668	.5504984
2015		.2325892	.2297992	1.01	0.314	-.2233823	.6885607
c.f2014#c.x1		.0404883	.1103324	0.37	0.714	-.1784352	.2594118
c.f2015#c.x1		.0778513	.1061898	0.73	0.465	-.1328523	.2885549
d		.904209	.5311615	1.70	0.092	-.1497307	1.958149
x1		2.476382	.1544607	16.03	0.000	2.169899	2.782866
c.d#c.x1		.1570952	.2113815	0.74	0.459	-.2623316	.576522
_cons		.2230452	.3747574	0.60	0.553	-.5205549	.9666452

Average marginal effects

Number of obs = 400
Subpop. no. obs = 200

Expression: Linear prediction, predict()
dy/dx wrt: w
At: f2014 = 1
f2015 = 0

(Std. err. adjusted for 100 clusters in id)

		Unconditional					
		dy/dx	std. err.	t	P> t	[95% conf. interval]	
w		2.496168	.3059971	8.16	0.000	1.889003	3.103333

Average marginal effects

Number of obs = 400
Subpop. no. obs = 200

Expression: Linear prediction, predict()
dy/dx wrt: w
At: f2014 = 0

8.5. Applications of TWM to DID in common treatment timing

f2015 = 1

(Std. err. adjusted for 100 clusters in id)

		Unconditional				
		dy/dx	std. err.	t	P> t	[95% conf. interval]
w		3.349012	.3110518	10.77	0.000	2.731818 3.966207

(200 real changes made)

...

Difference-in-difference with Multiple Time Periods

Number of obs = 400

Outcome model : least squares

Treatment model: inverse probability

	Coefficient	Std. err.	z	P> z	[95% conf. interval]
g2014					
t_2012_2013	-1.401593	.6428278	-2.18	0.029	-2.661512 -.1416737
t_2013_2014	3.388112	.5850216	5.79	0.000	2.241491 4.534734
t_2013_2015	3.915051	.6484609	6.04	0.000	2.644091 5.186011

Control: Never Treated

See Callaway and Sant'Anna (2021) for details

Source	SS	df	MS	Number of obs	=	300
				F(8, 291)	=	210.00
Model	12568.7919	8	1571.09898	Prob > F	=	0.0000
Residual	2177.12159	291	7.48151749	R-squared	=	0.8524
				Adj R-squared	=	0.8483
Total	14745.9135	299	49.3174363	Root MSE	=	2.7352

logy	Coefficient	Std. err.	t	P> t	[95% conf. interval]
year					
2013	-.0188111	.3880383	-0.05	0.961	-.7825286 .7449063
2014	.0787158	.513587	0.15	0.878	-.9321002 1.089532
2015	.2325892	.5121166	0.45	0.650	-.7753329 1.240511
c.f2014#c.x1	.0404883	.1798899	0.23	0.822	-.3135619 .3945385

8. Mundlak device in fixed effect models

c.f2015#c.x1		.0778513	.1879703	0.41	0.679	-.2921023	.4478049
d		.904209	.3894073	2.32	0.021	.1377972	1.670621
x1		2.476382	.1100309	22.51	0.000	2.259825	2.69294
c.d#c.x1		.1570952	.1556659	1.01	0.314	-.1492786	.463469
_cons		.2230452	.3355663	0.66	0.507	-.4373996	.8834899

Variable	Obs	Mean	Std. dev.	Min	Max
tetilda	50	2.525512	3.370158	-5.463166	10.58631

Variable	Obs	Mean	Std. dev.	Min	Max
tetilda	50	3.34147	3.105354	-3.463709	10.24051

Source	SS	df	MS	Number of obs	=	100
Model	900.277122	4	225.06928	F(4, 96)	=	21.48
Residual	1005.96245	96	10.4787755	Prob > F	=	0.0000
Total	1906.23957	100	19.0623957	R-squared	=	0.4723
				Adj R-squared	=	0.4503
				Root MSE	=	3.2371

tetilda	Coefficient	Std. err.	t	P> t	[95% conf. interval]
c.w#c.f2014	2.496168	.4582501	5.45	0.000	1.586549 3.405787
c.w#c.f2015	3.349012	.4582306	7.31	0.000	2.439432 4.258593
c.w#c.f2014#c.x1_dm_1	-.2701343	.1881329	-1.44	0.154	-.6435751 .1033065
c.w#c.f2015#c.x1_dm_1	-.0683301	.1811302	-0.38	0.707	-.4278707 .2912104

Treatment and time information

Time variable: year

Control: w = 0

Treatment: w = 1

	Control	Treatment
--	---------	-----------

8.5. Applications of TWM to DID in common treatment timing

Group			
id		50	50
-----+-----			
Time			
Minimum		2012	2014
Maximum		2012	2014

Difference-in-differences regression
Data type: Longitudinal

Number of obs = 400

(Std. err. adjusted for 100 clusters in id)

	logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
ATET							
	w						
(1 vs 0)		3.206857	.4320407	7.42	0.000	2.349594	4.064119

Note: ATET estimate adjusted for panel effects and time effects.

Parallel-trends test (pretreatment time period)
H0: Linear trends are parallel

F(1, 99) = 4.64
Prob > F = 0.0337

Linear regression

Number of obs = 100
F(1, 98) = 4.66
Prob > F = 0.0332
R-squared = 0.0454
Root MSE = 3.2451

-----+-----						
	D.logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]
-----+-----						
	d	-1.401579	.6490187	-2.16	0.033	-2.689536 -.1136226
	_cons	1.026794	.4697368	2.19	0.031	.0946167 1.958972

Fixed-effects (within) regression

Number of obs = 400

8. Mundlak device in fixed effect models

```

Group variable: id                                Number of groups =      100

R-squared:                                       Obs per group:
    Within = 0.8680                               min =      4
    Between = 0.8466                             avg =     4.0
    Overall = 0.8412                             max =      4

corr(u_i, Xb) = 0.4488                          F(12, 99) =     157.25
                                                Prob > F   =      0.0000

```

(Std. err. adjusted for 100 clusters in id)

	logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014		2.456402	.5067512	4.85	0.000	1.450898	3.461907
c.w#c.f2015		3.208291	.7927064	4.05	0.000	1.63539	4.781193
c.w#c.f2014#							
c.x1_dm_1		-.0503407	.1059229	-0.48	0.636	-.2605148	.1598333
c.w#c.f2015#							
c.x1_dm_1		.0924389	.0842389	1.10	0.275	-.0747095	.2595872
year							
2013		.0205346	.1993658	0.10	0.918	-.3750504	.4161195
2014		-.009636	.2174468	-0.04	0.965	-.4410976	.4218256
2015		.2623215	.1940125	1.35	0.179	-.1226414	.6472845
year#c.x1							
2012		1.933064	.0764551	25.28	0.000	1.781361	2.084768
2013		2.003714	.0662191	30.26	0.000	1.872321	2.135107
2014		2.01687	.0849868	23.73	0.000	1.848238	2.185502
2015		1.991922	.072964	27.30	0.000	1.847146	2.136699
c.d#c.t		.0668116	.2957253	0.23	0.822	-.5199715	.6535946
_cons		.4575496	.1764897	2.59	0.011	.1073557	.8077434
sigma_u		3.0627341					
sigma_e		.99359271					
rho		.90477742	(fraction of variance due to u_i)				

```

Linear regression                                Number of obs   =      400
                                                F(14, 99)      =     101.72

```

8.5. Applications of TWM to DID in common treatment timing

```

Prob > F      =      0.0000
R-squared     =      0.8446
Root MSE     =      2.8741

```

(Std. err. adjusted for 100 clusters in id)

logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014	1.833542	.6136741	2.99	0.004	.615879	3.051204
c.w#c.f2015	2.243344	.9499793	2.36	0.020	.358379	4.128309
c.w#c.f2014# c.x1_dm_1	-.2653456	.177891	-1.49	0.139	-.6183199	.0876287
c.w#c.f2015# c.x1_dm_1	-.0635414	.150658	-0.42	0.674	-.3624795	.2353968
year						
2013	-.2472397	.2421693	-1.02	0.310	-.7277561	.2332767
2014	-.0388673	.2538858	-0.15	0.879	-.5426319	.4648973
2015	.1150061	.2472402	0.47	0.643	-.375572	.6055842
year#c.x1						
2013	-.0241251	.0913453	-0.26	0.792	-.2053741	.1571239
2014	.0216128	.1358993	0.16	0.874	-.248041	.2912665
2015	.0589757	.1180401	0.50	0.618	-.1752415	.293193
d	.23577	.7209278	0.33	0.744	-1.194707	1.666247
x1	2.495258	.1736213	14.37	0.000	2.150756	2.83976
c.d#c.x1	.1523065	.2136548	0.71	0.478	-.2716309	.5762438
c.d#c.t	.4430422	.3468509	1.28	0.204	-.2451853	1.13127
_cons	.3406282	.3823786	0.89	0.375	-.418094	1.09935

Linear regression

```

Number of obs   =      400
F(15, 99)       =      95.45
Prob > F        =      0.0000
R-squared       =      0.8450
Root MSE       =      2.8744

```

(Std. err. adjusted for 100 clusters in id)

8. Mundlak device in fixed effect models

logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
c.w#c.f2014	1.917612	.6348799	3.02	0.003	.6578725	3.177352
c.w#c.f2015	2.378839	.9651184	2.46	0.015	.4638351	4.293844
c.w#c.f2014# c.x1_dm_1	-.7263424	.2927708	-2.48	0.015	-1.307263	-.1454217
c.w#c.f2015# c.x1_dm_1	-.8415862	.4868867	-1.73	0.087	-1.807675	.1245027
year						
2013	-.2611685	.2476914	-1.05	0.294	-.7526419	.2303049
2014	-.0671335	.2642967	-0.25	0.800	-.5915555	.4572885
2015	.0867399	.2545516	0.34	0.734	-.4183457	.5918255
year#c.x1						
2013	-.1857772	.1396877	-1.33	0.187	-.4629479	.0913934
2014	-.0735746	.1623912	-0.45	0.651	-.395794	.2486448
2015	-.0362116	.1317686	-0.27	0.784	-.2976691	.2252458
d	.1075441	.732768	0.15	0.884	-1.346427	1.561515
x1	2.590445	.1878499	13.79	0.000	2.21771	2.96318
c.d#c.x1	-.3378406	.3669519	-0.92	0.359	-1.065953	.3902716
c.d#c.t	.5199705	.3447009	1.51	0.135	-.1639908	1.203932
c.d#c.t#c.x1	.317048	.1852193	1.71	0.090	-.0504673	.6845633
_cons	.3688944	.3868765	0.95	0.343	-.3987525	1.136541

```
( 1) c.d#c.t = 0
( 2) c.d#c.t#c.x1 = 0
```

```
F( 2, 99) = 2.47
Prob > F = 0.0898
```

Linear regression	Number of obs	=	300
	F(11, 99)	=	82.82
	Prob > F	=	0.0000
	R-squared	=	0.8531

8.5. Applications of TWM to DID in common treatment timing

Root MSE = 2.7429

(Std. err. adjusted for 100 clusters in id)

	logy	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]	
year							
2013		-.2611685	.2475879	-1.05	0.294	-.7524366	.2300996
2014		-.0671335	.2641863	-0.25	0.800	-.5913364	.4570694
2015		.0867399	.2544452	0.34	0.734	-.4181347	.5916145
year#c.x1							
2013		-.1857773	.1396293	-1.33	0.186	-.4628321	.0912775
2014		-.0735746	.1623234	-0.45	0.651	-.3956594	.2485101
2015		-.0362116	.1317135	-0.27	0.784	-.2975599	.2251366
d		.107544	.7324618	0.15	0.884	-1.345819	1.560907
x1		2.590445	.1877714	13.80	0.000	2.217866	2.963025
c.d#c.x1		-.3378407	.3667986	-0.92	0.359	-1.065649	.3899672
c.d#c.t		.5199705	.3445569	1.51	0.134	-.1637051	1.203646
c.d#c.t#c.x1		.3170481	.1851419	1.71	0.090	-.0503137	.6844098
_cons		.3688944	.3867149	0.95	0.342	-.3984318	1.136221

(1) c.d#c.t = 0

(2) c.d#c.t#c.x1 = 0

F(2, 99) = 2.47
Prob > F = 0.0896

Variable	Obs	Mean	Std. dev.	Min	Max
tetilda	50	1.996512	3.755621	-5.357233	9.968782
Variable	Obs	Mean	Std. dev.	Min	Max
tetilda	50	2.285944	3.772198	-4.526336	10.79229

Source	SS	df	MS	Number of obs	=	100
Model	842.991546	4	210.747887	F(4, 96)	=	20.11
Residual	1005.96248	96	10.4787758	Prob > F	=	0.0000
				R-squared	=	0.4559

8. Mundlak device in fixed effect models

-----+-----					Adj R-squared	=	0.4333
Total		1848.95402	100	18.4895402	Root MSE	=	3.2371
-----+-----							
tetilda		Coefficient	Std. err.	t	P> t	[95% conf. interval]	
-----+-----							
c.w#c.f2014		1.917612	.4582501	4.18	0.000	1.007993	2.827231
c.w#c.f2015		2.378839	.4582306	5.19	0.000	1.469259	3.28842
c.w#c.f2014#							
c.x1_dm_1		-.7263426	.1881329	-3.86	0.000	-1.099783	-.3529017
c.w#c.f2015#							
c.x1_dm_1		-.8415865	.1811302	-4.65	0.000	-1.201127	-.4820459
-----+-----							

9. Correlated Random Effect

9.1. Panel data

For panel data, the usual set up is:

$$y_{it} = \alpha + X_{it}\beta + v_i + \epsilon_{it} \quad (9.1)$$

9.2. Fixed effect

A fixed effect model can be done with OLS on

$$(y_{it} - \bar{y}_i) = (X_{it} - \bar{X}_i)\beta + (\epsilon_{it} - \bar{\epsilon}_i) \quad (9.2)$$

This is basically OLS on de-meaned Y and X . Note that the intercept α and the individual effect v_i both drop out after de-meaning by unit.

9.3. Random effect

The random effect looks at this at different angle: it treated $v_i + \epsilon_{it}$ as the error term. There are two components of the error term. Suppose they are estimated as $\hat{\sigma}_e^2$ (idiosyncratic component), and $\hat{\sigma}_u^2$ (individual component). Then we can do GLS transformation:

$$z_{it}^* = z_{it} - \hat{\theta}_i \bar{z}_i$$

$$\text{and } \hat{\theta}_i = 1 - \sqrt{\frac{\hat{\sigma}_e^2}{T_i \hat{\sigma}_u^2 + \hat{\sigma}_e^2}}$$

where T_i is the number of observations for individual i .

Given estimates of $\hat{\sigma}_e^2$ and $\hat{\sigma}_u^2$, we can run OLS on transformed variables (including y and all X 's). We can iterate the process.

9.4. Correlated random effect

Correlated random effect (CRE) can be done by running a random effect model of y_{it} on X_{it} , $\bar{X}_i$, w_i (time-invariant covariates) and a constant.

The shortcoming of RE model is that it has to assume v_i and X_{it} are uncorrelated to have a consistent estimator. If that is not true (most economists don't think it is), then we have an inconsistent estimator. Meanwhile, FE estimator is consistent, because v_i is not in the error term, it gets cancelled out.

The disadvantage of FE model is that it cannot include any time-invariant covariate, such as race, gender, etc.

The benefit of CRE is that it can include z_i 's which are time-invariant, while remain consistent. In fact, Mundlak and Wooldridge pointed out that CRE estimates on X_{it} are the same as FE estimates.

There is also a Mundlak test to choose between RE or CRE.

9.5. Example

Stata 18 implemented CRE in the `xtreg` command, with “cre” option. The following example is from Stata's website, using the `nlswork` dataset.

```
webuse nlswork

xtreg ln_wage tenure age i.race, cre vce(cluster idcode)
```

(National Longitudinal Survey of Young Women, 14-24 years old in 1968)

note: 2.race omitted from xt_means because of collinearity.

note: 3.race omitted from xt_means because of collinearity.

Correlated random-effects regression	Number of obs	=	28,101
Group variable: idcode	Number of groups	=	4,699

R-squared:	Obs per group:
Within = 0.1296	min = 1
Between = 0.2346	avg = 6.0
Overall = 0.1890	max = 15

	Wald chi2(4)	=	1685.18
corr(xit_vars*b, xt_means*) = 0.5474	Prob > chi2	=	0.0000

(Std. err. adjusted for 4,699 clusters in idcode)

	Robust
--	--------

ln_wage	Coefficient	std. err.	z	P> z	[95% conf. interval]	

xit_vars						
tenure	.0211313	.0012113	17.44	0.000	.0187572	.0235055
age	.0121949	.0007414	16.45	0.000	.0107417	.013648
race						
Black	-.1312068	.0117856	-11.13	0.000	-.1543061	-.1081075
Other	.1059379	.0593177	1.79	0.074	-.0103225	.2221984
_cons	1.2159	.0306965	39.61	0.000	1.155736	1.276064

xt_means						
tenure	.0376991	.002281	16.53	0.000	.0332283	.0421698
age	-.0011984	.0013313	-0.90	0.368	-.0038077	.0014109
race						
Black	0	(omitted)				
Other	0	(omitted)				

sigma_u	.33334407					
sigma_e	.29808194					
rho	.55567161	(fraction of variance due to u_i)				

Mundlak test (xt means = 0): chi2(2) = 331.5144					Prob > chi2 = 0.0000	

To compare with a fixed effect model:

```
webuse nlswork
```

```
xtreg ln_wage tenure age i.race, fe vce(cluster idcode)
```

(National Longitudinal Survey of Young Women, 14-24 years old in 1968)

note: 2.race omitted because of collinearity.

note: 3.race omitted because of collinearity.

Fixed-effects (within) regression

Group variable: idcode

Number of obs = 28,101

Number of groups = 4,699

R-squared:

 Within = 0.1296

 Between = 0.1916

 Overall = 0.1456

Obs per group:

 min = 1

 avg = 6.0

 max = 15

corr(u_i, Xb) = 0.1302

F(2, 4698) = 766.79

Prob > F = 0.0000

9. Correlated Random Effect

(Std. err. adjusted for 4,699 clusters in idcode)							
ln_wage	Coefficient	Robust std. err.	t	P> t	[95% conf. interval]		
tenure	.0211313	.0012112	17.45	0.000	.0187568	.0235059	
age	.0121949	.0007414	16.45	0.000	.0107414	.0136483	
race							
Black	0	(omitted)					
Other	0	(omitted)					
_cons	1.256467	.0194187	64.70	0.000	1.218397	1.294537	
sigma_u	.39034493						
sigma_e	.29808194						
rho	.63165531	(fraction of variance due to u_i)					

We see the coefficient estimates are the same for “tenure” and “age”, but CRE model allows you to estimate the effect of “race”.

We can also manually do it by using a RE model on X , $\bar{X}$ and z :

```
webuse nlswork
* Restrict the mean calculation to the estimation sample (listwise-deleted
* on the model's variables), matching what xtreg's native cre option does
* internally. Both tenure and age have missing values in nlswork, so
* computing age_mean/tenure_mean over all rows for an idcode (rather than
* only rows with complete data on ln_wage, tenure, age, and race) would
* make the manual means diverge from cre's.
egen age_mean = mean(age) if !missing(ln_wage, tenure, age, race), by(idcode)
egen tenure_mean = mean(tenure) if !missing(ln_wage, tenure, age, race), by(idcode)
xtreg ln_wage tenure tenure_mean age age_mean i.race, vce(cluster idcode)
```

(National Longitudinal Survey of Young Women, 14-24 years old in 1968)

(433 missing values generated)

(433 missing values generated)

Random-effects GLS regression	Number of obs	=	28,101
Group variable: idcode	Number of groups	=	4,699

R-squared:	Obs per group:
------------	----------------

Within	= 0.1296	min	= 1
Between	= 0.2346	avg	= 6.0
Overall	= 0.1890	max	= 15

	Wald chi2(6)	= 2688.54
corr(u_i, X) = 0 (assumed)	Prob > chi2	= 0.0000

(Std. err. adjusted for 4,699 clusters in idcode)

ln_wage	Coefficient	Robust std. err.	z	P> z	[95% conf. interval]	
tenure	.0211313	.0012113	17.44	0.000	.0187572	.0235055
tenure_mean	.0376991	.002281	16.53	0.000	.0332283	.0421698
age	.0121949	.0007414	16.45	0.000	.0107417	.013648
age_mean	-.0011984	.0013313	-0.90	0.368	-.0038077	.0014109
race						
Black	-.1312068	.0117856	-11.13	0.000	-.1543061	-.1081075
Other	.1059379	.0593177	1.79	0.074	-.0103225	.2221984
_cons	1.2159	.0306965	39.61	0.000	1.155736	1.276064
sigma_u	.33334407					
sigma_e	.29808194					
rho	.55567161	(fraction of variance due to u_i)				

This is what stata's "cre" option is doing behind the scene.

10. More on Correlated Random Effect (CRE)

Jeff Wooldridge suggested using Correlated Random Effect (CRE) for binary outcomes in this twitter post: <https://x.com/jmwooldridge/status/1986100627454206220>

One word on the name before we start, since it invites a misreading. “Correlated random effects” describes an *assumption*, not an estimator: the unobserved effect v_i is modelled as a function of the group means and is allowed to correlate with the regressors, rather than being treated as a free parameter to estimate. Whether you then estimate by pooled OLS, pooled probit, or an actual random-effects routine is a separate choice. Most of what follows uses **pooled** estimation with group means added, in which no random effect is estimated at all.

Here I summarize what I learned so far:

10.1. linear model

For panel data, the usual set up is:

$$y_{it} = \alpha + X_{it}\beta + v_i + \epsilon_{it} \quad (10.1)$$

Wooldridge (2021) paper on DiD shows that these models are the same: 1. Two way fixed effect OLS. Dummy variable approach is the same too. 2. Mundlak’s CRE approach. This includes a pooled OLS with Mundlak device, and CRE with random effects.

Two qualifications on that second item, both of which matter in practice.

What “the Mundlak device” has to include. For a *one-way* model it is the unit means of the time-varying regressors, and nothing else. That much holds whatever the panel looks like, because $\sum_t (x_{it} - \bar{x}_i) = 0$ within each unit by construction. A *two-way* model needs the time dimension handled as well, and the two usual ways of doing it — adding the time means of the regressors, or keeping a full set of time dummies — reproduce the two-way fixed effect estimate **only when the unit and time factors are orthogonal**, which for a two-way panel means every unit appearing in every period. The `castle` data used below is balanced, so both routes are exact there; the example takes the second one (unit means plus `factor(year)`), which is why it reproduces the *two-way* estimate and not the *one-way* one. Unit means alone would give the one-way estimate.

That orthogonality condition is binding, not decorative. On an unbalanced panel it fails, and the discrepancy can be substantial: in the `nlswork` example in [the Mundlak chapter](#) the coefficient on `age` differs by more than a factor of two. What restores the equivalence there is adding the **unit means of the time dummies**. In a balanced panel those means are the constant $1/T$ and have no variance, which is precisely why the shortcut above reads as though it were unconditional. With **three or more** fixed effects the shortcut is weaker still — one set of means per dimension is exact

10. More on Correlated Random Effect (CRE)

only under a fully orthogonal design, and the construction that survives without one is to pick the dimension being eliminated and project it onto the within-unit means of everything else retained in the model, dummies included. See [More than two fixed effects](#).

The idea of Chamberlain's device is that since we don't observe v_i , we project v_i onto the history of X_i : $v_i = \psi + \bar{X}_i\xi + a_i$, and replace v_i with that projection. Mundlak's device is a special case, where we put equal weight on all X_{it} , so the projection of v_i onto X_i reduces to the mean $\bar{X}_i = \frac{1}{T} \sum_t X_{it}$. Replace v_i with that projection, and what is left (a_i) is by construction uncorrelated with X_{it} .

Here is an example:

10.1.1. example 1

```
library(dplyr)
library(fixest)
library(bacondecomp)
data("castle")

# Fixed-effects model with individual and time fixed effects
fe_model_twoway <- feols(l_homicide ~ poverty + l_police | state + year, data = castle)
summary(fe_model_twoway)
```

OLS estimation, Dep. Var.: l_homicide

Observations: 550

Fixed-effects: state: 50, year: 11

Standard-errors: IID

	Estimate	Std. Error	t value	Pr(> t)
poverty	-0.027068	0.013306	-2.034206	0.042471 *
l_police	0.066033	0.104588	0.631364	0.528098

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

RMSE: 0.176686 Adj. R2: 0.898978

Within R2: 0.009286

```
# random effect
library(lme4)
re_model <- lmer(l_homicide ~ poverty + l_police + factor(year) + (1 | state), data = castle)
summary(re_model)
```

Linear mixed model fit by REML ['lmerMod']

Formula: l_homicide ~ poverty + l_police + factor(year) + (1 | state)

Data: castle

REML criterion at convergence: 3.9

Scaled residuals:

Min	1Q	Median	3Q	Max
-4.5357	-0.4657	0.0151	0.4601	3.6973

Random effects:

Groups	Name	Variance	Std.Dev.
state	(Intercept)	0.30041	0.5481
Residual		0.03551	0.1884

Number of obs: 550, groups: state, 50

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	0.4395941	0.6019389	0.730
poverty	-0.0007991	0.0120087	-0.067
l_police	0.1665913	0.1015581	1.640
factor(year)2001	0.0225820	0.0379597	0.595
factor(year)2002	0.0005681	0.0389059	0.015
factor(year)2003	0.0475588	0.0395876	1.201
factor(year)2004	0.0419220	0.0411546	1.019
factor(year)2005	0.0612232	0.0435803	1.405
factor(year)2006	0.0806576	0.0418469	1.927
factor(year)2007	0.0790850	0.0409066	1.933
factor(year)2008	0.0401268	0.0414964	0.967
factor(year)2009	-0.0424481	0.0488959	-0.868
factor(year)2010	-0.0959167	0.0470940	-2.037

```
# Mundlak
castle2 <- castle |>
  group_by(state) |>
  mutate(poverty_mean = mean(poverty, na.rm = TRUE), l_police_mean = mean(l_police, na.rm = TRUE))
cre_model <- feols(l_homicide ~ poverty + poverty_mean + l_police + l_police_mean + factor(year))
summary(cre_model)
```

OLS estimation, Dep. Var.: l_homicide

Observations: 550

Standard-errors: IID

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-7.800873	0.531186	-14.685774	< 2.2e-16 ***
poverty	-0.027068	0.030004	-0.902125	3.6740e-01
poverty_mean	0.125283	0.030674	4.084272	5.0983e-05 ***
l_police	0.066033	0.235836	0.279996	7.7959e-01
l_police_mean	1.318037	0.253323	5.202983	2.7987e-07 ***
factor(year)2001	0.033071	0.085349	0.387481	6.9855e-01
factor(year)2002	0.022870	0.087972	0.259968	7.9499e-01
factor(year)2003	0.074569	0.089852	0.829907	4.0696e-01
factor(year)2004	0.079080	0.094152	0.839920	4.0133e-01
factor(year)2005	0.109800	0.100737	1.089973	2.7622e-01

10. More on Correlated Random Effect (CRE)

```
factor(year)2006  0.118822  0.095991  1.237857 2.1631e-01
factor(year)2007  0.115412  0.093475  1.234679 2.1749e-01
factor(year)2008  0.077190  0.095054  0.812069 4.1711e-01
factor(year)2009  0.027528  0.114987  0.239400 8.1089e-01
factor(year)2010 -0.034286  0.110153 -0.311255 7.5573e-01
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
RMSE: 0.417154  Adj. R2: 0.486344
```

```
# Mundlak
cre_model2 <- lmer(l_homicide ~ poverty + poverty_mean + l_police + l_police_mean + factor(year)
summary(cre_model2)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: l_homicide ~ poverty + poverty_mean + l_police + l_police_mean +
  factor(year) + (1 | state)
Data: castle2
```

REML criterion at convergence: -30.1

Scaled residuals:

Min	1Q	Median	3Q	Max
-4.5991	-0.4336	0.0219	0.4307	3.6228

Random effects:

Groups	Name	Variance	Std.Dev.
state	(Intercept)	0.14872	0.3856
	Residual	0.03518	0.1876

Number of obs: 550, groups: state, 50

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	-7.80087	1.60978	-4.846
poverty	-0.02707	0.01331	-2.034
poverty_mean	0.12528	0.02360	5.309
l_police	0.06603	0.10459	0.631
l_police_mean	1.31804	0.30140	4.373
factor(year)2001	0.03307	0.03785	0.874
factor(year)2002	0.02287	0.03901	0.586
factor(year)2003	0.07457	0.03985	1.871
factor(year)2004	0.07908	0.04175	1.894
factor(year)2005	0.10980	0.04467	2.458
factor(year)2006	0.11882	0.04257	2.791
factor(year)2007	0.11541	0.04145	2.784
factor(year)2008	0.07719	0.04215	1.831
factor(year)2009	0.02753	0.05099	0.540
factor(year)2010	-0.03429	0.04885	-0.702

In this example, fixed effect model, CRE1 (pooled OLS with the Mundlak device) and CRE2 (random effects with the Mundlak device) all give the same coefficient on poverty.

As we know, RE model assumes v_i is uncorrelated with X_{it} , while CRE and FE allow for correlation between v_i and X_{it} .

10.2. nonlinear model

What I did not realize before is that this is only valid for the linear case. In the linear model the algebra above is an identity: adding the group means reparameterizes the regression so that the coefficient on x_{it} is the within estimator, and no distributional assumption is involved. In a nonlinear model — logit, probit, Poisson — there is no such identity. Substituting a projection for v_i becomes a genuine modelling assumption about $D(v_i | X_i)$, and the likelihood you end up maximizing is not the one you started from.

10.2.1. Binary outcome

For binary outcomes, the model is:

$$\Pr(y_{it} = 1 | X_{it}, v_i) = \Phi(\alpha + X_{it}\beta + v_i) \quad (10.2)$$

I used to think it's the best to use conditional logit. However, it needs “conditional independence” (serial independence could be a better name) to be consistent. That is, we have to assume the series of y_{it} is conditionally independent of each other given X_{it} and v_i . This is a strong assumption. If we think we have serial correlation in the error term, then this would not hold. The other disadvantage of conditional logit (also called fixed effect logit) is that v_i 's are not estimated, it is just a nuisance parameter. Any partial effects cannot be calculated, because any partial effects are functions of v_i 's.

Instead, Wooldridge suggests using CRE probit — that is, the Mundlak device inside a probit. Basically adding Mundlak device (group means of X 's) to the model, and then use pooled probit. Note that we do not use RE probit or logit, since that would need conditional independence too. Lin and Wooldridge (2019, sec. 5.2) put this plainly: “None of random effects probit, RE logit, RE Tobit, RE Poisson, and so on has robustness properties in the presence of serial correlation of the innovations.” Pooled methods with cluster-robust inference are robust to serial dependence of arbitrary form, which is the reason to prefer them here.

For nonlinear panel data with unobserved heterogeneity (meaning there is v_i that we don't observe), we have a few assumptions.

First, strict exogeneity,

$$D(y_{it}|X_{i1}, \dots, X_{iT}, v_i) = D(y_{it}|X_{it}, v_i) \quad (10.3)$$

This means that the distribution of y_{it} only depends on X_{it} , given v_i .

Second, conditional independence,

$$D(y_{i1}, y_{i2}, \dots, y_{iT}|X_i, v_i) = \prod_{t=1}^T D(y_{it}|X_{it}, v_i) \quad (10.4)$$

10. More on Correlated Random Effect (CRE)

This means that the distribution of joint distribution of y_{it} 's can be modeled by each y_{it} independently, given X_i and v_i . We need this in most nonlinear cases, but not in linear case.

Then we need to specify $D(v_i|X_i)$. The random effect assumption is saying

$$D(v_i|X_i) = D(v_i)$$

{#eq-more-cre-5}, that is, v_i is independent of X_i . This is a strong assumption. We should try to avoid this assumption.

Now, most nonlinear models make the second assumption, in the panel setting, when we have v_i . That includes RE model, conditional logit. However, that does not include pooled probit/logit, or CRE with pooled logit/probit. That is why Wooldridge recommend using CRE with pooled probit.

10.2.2. Example 2

```
# generate a binary outcome, high_homocide

castle2 <- castle2 |>
  mutate(high_homocide = ifelse(l_homicide > mean(l_homicide, na.rm = TRUE), 1, 0))

# CRE pooled probit
cre_probit <- glm(high_homocide ~ poverty + poverty_mean + l_police + l_police_mean + factor(year),
  family = binomial)
# Compute Average Marginal Effects (AMEs)
library(marginaleffects)
ame_cre_probit <- avg_slopes(cre_probit)
ame_cre_probit
```

	Term	Contrast	Estimate	Std. Error	z	Pr(> z)	S	2.5 %
	l_police	dY/dX	0.01924	0.2622	0.0734	0.9415	0.1	-0.494686
	l_police_mean	dY/dX	-0.00304	0.2824	-0.0108	0.9914	0.0	-0.556572
	poverty	dY/dX	-0.06477	0.0334	-1.9417	0.0522	4.3	-0.130144
	poverty_mean	dY/dX	0.06620	0.0341	1.9424	0.0521	4.3	-0.000597
	year	2001 - 2000	0.06596	0.0990	0.6659	0.5054	1.0	-0.128159
	year	2002 - 2000	0.09470	0.1019	0.9293	0.3527	1.5	-0.105016
	year	2003 - 2000	0.22958	0.1020	2.2518	0.0243	5.4	0.029753
	year	2004 - 2000	0.15492	0.1082	1.4315	0.1523	2.7	-0.057193
	year	2005 - 2000	0.22388	0.1133	1.9766	0.0481	4.4	0.001880
	year	2006 - 2000	0.26184	0.1068	2.4526	0.0142	6.1	0.052595
	year	2007 - 2000	0.24899	0.1048	2.3756	0.0175	5.8	0.043564
	year	2008 - 2000	-0.03718	0.1099	-0.3383	0.7352	0.4	-0.252646
	year	2009 - 2000	-0.02471	0.1345	-0.1837	0.8542	0.2	-0.288280
	year	2010 - 2000	-0.21753	0.1186	-1.8339	0.0667	3.9	-0.450000
	97.5 %							
			0.533162					

0.550499
 0.000608
 0.132987
 0.260070
 0.294406
 0.429403
 0.367039
 0.445887
 0.471080
 0.454421
 0.178276
 0.238865
 0.014950

Type: response

```
# linear model
cre_lm <- lm(high_homicide ~ poverty + poverty_mean + l_police + l_police_mean + factor(year),
ame_lm <- avg_slopes(cre_lm)
ame_lm
```

	Term	Contrast	Estimate	Std. Error	z	Pr(> z)	S	2.5 %
	l_police	dY/dX	0.01234	0.2666	0.0463	0.9631	0.1	-0.51010
	l_police_mean	dY/dX	0.00427	0.2865	0.0149	0.9881	0.0	-0.55724
	poverty	dY/dX	-0.06406	0.0339	-1.8891	0.0589	4.1	-0.13053
	poverty_mean	dY/dX	0.06578	0.0347	1.8975	0.0578	4.1	-0.00217
	year	2001 - 2000	0.06384	0.0965	0.6618	0.5081	1.0	-0.12522
	year	2002 - 2000	0.09082	0.0994	0.9134	0.3610	1.5	-0.10406
	year	2003 - 2000	0.22452	0.1016	2.2108	0.0270	5.2	0.02548
	year	2004 - 2000	0.14789	0.1064	1.3898	0.1646	2.6	-0.06068
	year	2005 - 2000	0.21659	0.1139	1.9023	0.0571	4.1	-0.00656
	year	2006 - 2000	0.25687	0.1085	2.3677	0.0179	5.8	0.04423
	year	2007 - 2000	0.24413	0.1056	2.3108	0.0208	5.6	0.03706
	year	2008 - 2000	-0.04737	0.1074	-0.4410	0.6592	0.6	-0.25794
	year	2009 - 2000	-0.03438	0.1300	-0.2645	0.7914	0.3	-0.28910
	year	2010 - 2000	-0.20933	0.1245	-1.6814	0.0927	3.4	-0.45335
	97.5 %							
			0.5348					
			0.5658					
			0.0024					
			0.1337					
			0.2529					
			0.2857					
			0.4236					
			0.3565					
			0.4397					

10. More on Correlated Random Effect (CRE)

0.4695
0.4512
0.1632
0.2203
0.0347

Type: response

```
# CRE RE probit
cre_re_probit <- glmer(high_homicide ~ poverty + poverty_mean + l_police + l_police_mean + fac
ame_cre_re_probit <- avg_slopes(cre_re_probit)
ame_cre_re_probit
```

	Term	Contrast	Estimate	Std. Error	z	Pr(> z)	S	2.5 %
	l_police	dY/dX	0.01927	0.2624	0.0735	0.9414	0.1	-0.494931
	l_police_mean	dY/dX	-0.00307	0.2827	-0.0109	0.9913	0.0	-0.557188
	poverty	dY/dX	-0.06477	0.0334	-1.9417	0.0522	4.3	-0.130143
	poverty_mean	dY/dX	0.06619	0.0341	1.9424	0.0521	4.3	-0.000599
	year	2001 - 2000	0.06595	0.0990	0.6659	0.5055	1.0	-0.128160
	year	2002 - 2000	0.09469	0.1019	0.9293	0.3527	1.5	-0.105017
	year	2003 - 2000	0.22958	0.1020	2.2518	0.0243	5.4	0.029753
	year	2004 - 2000	0.15492	0.1082	1.4315	0.1523	2.7	-0.057195
	year	2005 - 2000	0.22388	0.1133	1.9765	0.0481	4.4	0.001878
	year	2006 - 2000	0.26184	0.1068	2.4526	0.0142	6.1	0.052595
	year	2007 - 2000	0.24899	0.1048	2.3756	0.0175	5.8	0.043563
	year	2008 - 2000	-0.03719	0.1099	-0.3383	0.7352	0.4	-0.252648
	year	2009 - 2000	-0.02471	0.1345	-0.1837	0.8542	0.2	-0.288283
	year	2010 - 2000	-0.21752	0.1186	-1.8339	0.0667	3.9	-0.449999

97.5 %
0.533476
0.551049
0.000609
0.132988
0.260069
0.294405
0.429402
0.367037
0.445886
0.471079
0.454420
0.178275
0.238863
0.014953

Type: response

Two caveats about that third model, and the second one matters more than the first.

First, on what `avg_slopes` is averaging. For a `merMod` object `marginalEffects` predicts with `re.form = NULL` by default, which *includes* the estimated random effects rather than setting them to zero; the package warns that its standard errors nevertheless reflect only the uncertainty in the fixed-effect parameters. If you want the AME for a “typical” ($v_i = 0$) unit you have to ask for it with `re.form = NA`, and neither of those is the *population-averaged* AME, which requires integrating the marginal effect over the estimated distribution of v_i . In a nonlinear model like probit those three quantities genuinely differ, because $\Phi(\cdot)$ is nonlinear in v_i .

Second, and more to the point here: on this data set the random-effects probit is a **boundary (singular) fit**. `lme4` reports an estimated state-level standard deviation of about 5×10^{-9} — numerically zero — and `isSingular()` returns `TRUE`. With no variance left in the random effect, this model *is* the pooled probit: its coefficients agree with the pooled fit to about 10^{-4} (poverty -0.17848 either way). So the third specification is not really a third specification on this data, and the distinction between conditional and population-averaged effects has nothing to apply to. It is kept here because the comparison is the one you would want to run in general — but do not read the agreement between models two and three as evidence about random effects; check `VarCorr()` before drawing that conclusion.

Ultimately we are interested in the partial effect, therefore we can compare across models. In this case, we should prefer average marginal effect by CRE pooled probit.

10.3. References

- Chamberlain, G. (1982). Multivariate regression models for panel data. *Journal of Econometrics* 18, 5–46.
- Lin, W., and J. M. Wooldridge (2019). Testing and correcting for endogeneity in nonlinear unobserved effects models. Ch. 2 in *Panel Data Econometrics: Theory*, ed. M. Tsionas, Elsevier, 21–43.
- Mundlak, Y. (1978). On the pooling of time series and cross section data. *Econometrica* 46, 69–85.
- Papke, L. E., and J. M. Wooldridge (2008). Panel data methods for fractional response variables with an application to test pass rates. *Journal of Econometrics* 145, 121–133.
- Wooldridge, J. M. (2010). *Econometric Analysis of Cross Section and Panel Data*, 2nd ed. MIT Press. Ch. 15 covers the nonlinear panel and CRE material.
- Wooldridge, J. M. (2019). Correlated random effects models with unbalanced panels. *Journal of Econometrics* 211, 137–150.

Part III.

Treatment Effects & Matching

11. Treatment effects and matching

11.1. Treatment effects in observational studies

Despite the popularity of randomized experiments in economics nowadays, most situations we have observational data in economic studies. One reason is experiments are expensive; the other reason is that sometimes it is simply not feasible to have experiments. If we have observational data, and we'd like to draw causal conclusions, then we have a few different situations. The worse situation is that we have an endogenous treatment. Or to say, we have unobserved confounders. In that case, we need instrumental variables. However instruments are hard to find, and even harder to justify. If we assume we don't have unobserved confounders. Or to say, we have conditional independence of the treatment variable. That is, conditional on other variables in the model, there is no unobserved confounders. If that assumption holds, then we can make causal inference.

Stata has a set of `eteffects` and `teffects` commands are designed for treatment effects. `eteffects` is for treatment effects when you have endogenous treatment. In that case, you'll need instruments to model treatment at first stage. Then `eteffects` use control function approach for the second stage of modeling treatment effects on outcome. In this blog, we are trying to understand how `teffects` works. That is, when we don't need an instruments, or say we assume no unobserved confounders.

11.1.1. Regression adjustment (RA)

The RA method is to allow all covariates' effects differ between treatment and control. That is, it is the same model as an outcome model with treatment interacting with all other covariates.

11.1.1.1. Example:

```
clear
webuse bweightex
teffects ra (bweight prenatal1 mmarried mage fbaby) (mbsmoke)
reg bweight i.mbsmoke##c.(prenatal1 mmarried mage fbaby)
margins r.mbsmoke
```

```
. clear

. webuse bweightex
(Hypothetical birthweight data)
```

11. Treatment effects and matching

```
. teffects ra (bweight prenatal1 mmarried mage fbaby) (mbsmoke)
```

```
Iteration 0: EE criterion = 1.223e-24
```

```
Iteration 1: EE criterion = 1.792e-25
```

```
Treatment-effects estimation      Number of obs      =          60
Estimator      : regression adjustment
Outcome model   : linear
Treatment model: none
```

		Robust		z	P> z	[95% conf. interval]	
bweight		Coefficient	std. err.				
ATE							
mbsmoke							
(Smoker							
vs							
Nonsmoker)		-389.3099	37.00882	-10.52	0.000	-461.8458	-316.7739
POmean							
mbsmoke							
Nonsmoker		3613.769	29.97438	120.56	0.000	3555.021	3672.518

```
. reg bweight i.mbsmoke##c.(prenatal1 mmarried mage fbaby)
```

Source	SS	df	MS	Number of obs	=	60
				F(9, 50)	=	18.11
Model	1376306.34	9	152922.927	Prob > F	=	0.0000
Residual	422241.592	50	8444.83184	R-squared	=	0.7652
				Adj R-squared	=	0.7230
Total	1798547.93	59	30483.8633	Root MSE	=	91.896

bweight	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
mbsmoke						
Smoker	-32.87702	193.2827	-0.17	0.866	-421.0967	355.3426
prenatal1	-23.8431	38.31014	-0.62	0.537	-100.7913	53.10507
mmarried	-40.39753	50.21942	-0.80	0.425	-141.2662	60.47114
mage	27.9537	7.671423	3.64	0.001	12.54519	43.36221
fbaby	-2.030497	40.76029	-0.05	0.960	-83.89995	79.83896
mbsmoke#						
c.prenatal1						
Smoker	8.436558	56.38855	0.15	0.882	-104.8232	121.6963
mbsmoke#						

c.mmarrried							
Smoker		33.92362	61.34015	0.55	0.583	-89.2817	157.1289
mbsmoke#							
c.mage							
Smoker		-15.98608	9.006494	-1.77	0.082	-34.07616	2.103995
mbsmoke#							
c.fbaby							
Smoker		14.86782	57.58771	0.26	0.797	-100.8005	130.5362
_cons		2976.807	151.0117	19.71	0.000	2673.491	3280.123

```
. margins r.mbsmoke
```

Contrasts of predictive margins
Model VCE: OLS

Number of obs = 60

Expression: Linear prediction, predict()

	df	F	P>F
mbsmoke	1	116.06	0.0000
Denominator	50		

	Contrast	Delta-method std. err.	[95% conf. interval]
mbsmoke			
(Smoker vs Nonsmoker)	-389.3099	36.13675	-461.8927 -316.7271

We can see the teffects return the same ATE(Average Treatment Effect) as the margins command after a regression with treatment interacting with all other covariates.

To estimate ATET (or ATT, Average Treatment effected on the Treated),

```
clear
webuse bweightex
teffects ra (bweight prenatal1 mmarrried mage fbaby) (mbsmoke), atet
reg bweight i.mbsmoke#c.(prenatal1 mmarrried mage fbaby)
margins r.mbsmoke, subpop(mbsmoke)
```

11. Treatment effects and matching

```
. clear

. webuse bweightex
(Hypothetical birthweight data)

. teffects ra (bweight prenatal1 mmarried mage fbaby) (mb smoke), atet
```

Iteration 0: EE criterion = 1.159e-24

Iteration 1: EE criterion = 5.133e-26

```
Treatment-effects estimation          Number of obs      =          60
Estimator      : regression adjustment
Outcome model  : linear
Treatment model: none
```

		Robust				
bweight		Coefficient	std. err.	z	P> z	[95% conf. interval]
-----+-----						
ATET						
mb smoke						
(Smoker						
vs						
Nonsmoker)		-437.4721	53.43513	-8.19	0.000	-542.203 -332.7412
-----+-----						
POmean						
mb smoke						
Nonsmoker		3693.639	53.80537	68.65	0.000	3588.182 3799.095
-----+-----						

```
. reg bweight i.mb smoke##c.(prenatal1 mmarried mage fbaby)
```

Source	SS	df	MS	Number of obs	=	60
-----+-----				F(9, 50)	=	18.11
Model	1376306.34	9	152922.927	Prob > F	=	0.0000
Residual	422241.592	50	8444.83184	R-squared	=	0.7652
-----+-----				Adj R-squared	=	0.7230
Total	1798547.93	59	30483.8633	Root MSE	=	91.896

bweight	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
-----+-----						
mbsmoke						
Smoker	-32.87702	193.2827	-0.17	0.866	-421.0967	355.3426
prenatal1	-23.8431	38.31014	-0.62	0.537	-100.7913	53.10507
mmarried	-40.39753	50.21942	-0.80	0.425	-141.2662	60.47114
mage	27.9537	7.671423	3.64	0.001	12.54519	43.36221
fbaby	-2.030497	40.76029	-0.05	0.960	-83.89995	79.83896

mbsmoke#							
c.prenatal1							
Smoker		8.436558	56.38855	0.15	0.882	-104.8232	121.6963
mbsmoke#							
c.mmarrried							
Smoker		33.92362	61.34015	0.55	0.583	-89.2817	157.1289
mbsmoke#							
c.mage							
Smoker		-15.98608	9.006494	-1.77	0.082	-34.07616	2.103995
mbsmoke#							
c.fbaby							
Smoker		14.86782	57.58771	0.26	0.797	-100.8005	130.5362
_cons		2976.807	151.0117	19.71	0.000	2673.491	3280.123

```
. margins r.mbsmoke, subpop(mbsmoke)
```

Contrasts of predictive margins

Number of obs = 60

Model VCE: OLS

Subpop. no. obs = 30

Expression: Linear prediction, predict()

		df	F	P>F
mbsmoke		1	73.30	0.0000
Denominator		50		

		Delta-method		
		Contrast	std. err.	[95% conf. interval]
mbsmoke				
(Smoker vs Nonsmoker)		-437.4721	51.09606	-540.1015 -334.8426

It is the comparison of treatment's effect on the potential outcomes, for the treated. That's why in margins, we have subpop(mbsmoke) option.

11.1.2. Inverse Probability Weighting (IPW)

IPW estimators have two steps. The first step is to estimate the treatment model, that is, treatment as a function of some covariates. Usually a logit model is used. Then the probability of treatment is estimated. In the second stage, the inverse probability is used as weights to compute the outcome difference between treatment versus control units.

These steps produce consistent estimates of the effect parameters because the treatment is assumed to be independent of the potential outcomes after conditioning on the covariates.

We can manually conduct the two steps, but the nice thing about using Stata's `teffects` is that it takes account of the noise of estimating probability in the first step when calculating standard errors in the second step.

11.1.2.1. example

```
clear
webuse cattaneo2
teffects ipw (bweight) (mbsmoke mmarried c.mage##c.mage fbaby medu, probit)
probit mbsmoke mmarried c.mage##c.mage fbaby medu
predict ps
replace ps = 1/ps if mbsmoke==1
replace ps = 1/(1-ps) if mbsmoke==0
reg bweight mbsmoke [pweight=ps]
```

```
. clear

. webuse cattaneo2
(Excerpt from Cattaneo (2010) Journal of Econometrics 155: 138-154)

. teffects ipw (bweight) (mbsmoke mmarried c.mage##c.mage fbaby medu, probit)
```

```
Iteration 0: EE criterion = 4.622e-21
Iteration 1: EE criterion = 8.453e-26
```

```
Treatment-effects estimation      Number of obs      =      4,642
Estimator      : inverse-probability weights
Outcome model  : weighted mean
Treatment model: probit
```

```
-----
            |               Robust
      bweight | Coefficient  std. err.      z    P>|z|      [95% conf. interval]
-----+-----
ATE          |
      mbsmoke |
```


(Smoker						
vs						
Nonsmoker)		-230.6886	25.81524	-8.94	0.000	-281.2856 -180.0917

P0mean						
mbsmoke						
Nonsmoker		3403.463	9.571369	355.59	0.000	3384.703 3422.222

```
. probit mbsmoke mmarried c.mage##c.mage fbaby medu
```

```
Iteration 0: Log likelihood = -2230.7484
Iteration 1: Log likelihood = -2042.6734
Iteration 2: Log likelihood = -2040.5088
Iteration 3: Log likelihood = -2040.5061
Iteration 4: Log likelihood = -2040.5061
```

Probit regression

Number of obs = 4,642

LR chi2(5) = 380.48

Prob > chi2 = 0.0000

Pseudo R2 = 0.0853

Log likelihood = -2040.5061

mbsmoke		Coefficient	Std. err.	z	P> z	[95% conf. interval]

mmarried		-.6484821	.0526991	-12.31	0.000	-.7517705 -.5451938
mage		.1744327	.0352437	4.95	0.000	.1053562 .2435092

c.mage#						
c.mage		-.0032559	.0006462	-5.04	0.000	-.0045224 -.0019894

fbaby		-.2175962	.0491066	-4.43	0.000	-.3138433 -.121349
medu		-.0863631	.0098692	-8.75	0.000	-.1057064 -.0670198
_cons		-1.558255	.4511589	-3.45	0.001	-2.44251 -.674

```
. predict ps
```

```
(option pr assumed; Pr(mbsmoke))
```

```
. replace ps = 1/ps if mbsmoke==1
```

```
(864 real changes made)
```

```
. replace ps = 1/(1-ps) if mbsmoke==0
```

```
(3,778 real changes made)
```

```
. reg bweight mbsmoke [pweight=ps]
```

```
(sum of wgt is 9,193.96990537643)
```

11. Treatment effects and matching

Linear regression	Number of obs	=	4,642
	F(1, 4640)	=	79.08
	Prob > F	=	0.0000
	R-squared	=	0.0389
	Root MSE	=	573.36

		Robust				
bweight		Coefficient	std. err.	t	P> t	[95% conf. interval]

mbsmoke		-230.6886	25.94182	-8.89	0.000	-281.5469 -179.8303
_cons		3403.463	9.616992	353.90	0.000	3384.609 3422.317

In the above example, we run `teffects ipw` and then manually replicate the two steps. The only thing different is the standard error for the treatment effect.

11.1.3. Doubly robust estimators

RA estimator estimates the outcome directly, IPW estimates the treatment assignment. There is also a group of estimators called doubly-robust estimators. They estimate both stages, and only require one of these stages to be correctly specified.

Stata implemented the augmented-IPW (AIPW) combination proposed by Robins and Rotnitzky (1995) and the IPW-regression-adjustment (IPWRA) combination proposed by Wooldridge (2010).

The AIPW estimator augments the IPW estimator with a correction term built from the outcome model. Double robustness is best stated as a **consistency** property: AIPW is consistent if *either* the treatment model *or* the outcome model is correctly specified (not necessarily both). It is *not* that the correction term literally “goes to zero” when the treatment model is right and the outcome model is wrong — with a misspecified outcome model the correction term does not generally vanish. Rather, the two pieces are constructed so that the estimator’s bias cancels in the limit as long as one of the two models is right.

The IPWRA estimator uses IPW probability weights when performing RA. The weights do not affect the accuracy of the RA estimator if the treatment model is wrong and the outcome model is correct. The weights correct the RA estimator if the treatment model is correct and the outcome model is wrong.

Here we run the two commands, then manually replicate both as a two-step (“plug-in”) procedure. `teffects` estimates the treatment and outcome models jointly by GMM/ML, using the joint estimation to get standard errors that account for first-stage estimation noise; the manual two-step version below gets essentially the same point estimate but not the correct standard errors (for the same reason noted above for IPW).

```
clear
webuse cattaneo2
teffects ipwra (bweight mmarried mage prenatal1 fbaby) (mbsmoke mmarried mage prenatal1 fbaby)
teffects aipw (bweight mmarried mage prenatal1 fbaby) (mbsmoke mmarried mage prenatal1 fbaby)
```

```
. clear

. webuse cattaneo2
(Excerpt from Cattaneo (2010) Journal of Econometrics 155: 138-154)

. teffects ipwra (bweight mmarried mage prenatal1 fbaby) (mbsmoke mmarried mag
> e prenatal1 fbaby)
```

```
Iteration 0: EE criterion = 9.066e-22
Iteration 1: EE criterion = 2.902e-26
```

```
Treatment-effects estimation          Number of obs    =      4,642
Estimator      : IPW regression adjustment
Outcome model  : linear
Treatment model: logit
```

			Robust			
	bweight	Coefficient	std. err.	z	P> z	[95% conf. interval]
ATE						
	mbsmoke					
	(Smoker					
	vs					
	Nonsmoker)	-238.7679	24.38353	-9.79	0.000	-286.5587 -190.977
POmean						
	mbsmoke					
	Nonsmoker	3402.851	9.538741	356.74	0.000	3384.155 3421.546

```
. teffects aipw (bweight mmarried mage prenatal1 fbaby) (mbsmoke mmarried mage
> prenatal1 fbaby)
```

```
Iteration 0: EE criterion = 2.115e-22
Iteration 1: EE criterion = 1.343e-26
```

```
Treatment-effects estimation          Number of obs    =      4,642
Estimator      : augmented IPW
Outcome model  : linear by ML
Treatment model: logit
```

			Robust			
	bweight	Coefficient	std. err.	z	P> z	[95% conf. interval]
ATE						
	mbsmoke					
	(Smoker					

11. Treatment effects and matching

vs							
Nonsmoker)		-239.0294	24.2524	-9.86	0.000	-286.5632	-191.4955

P0mean							
mbsmoke							
Nonsmoker		3402.839	9.538926	356.73	0.000	3384.143	3421.535

Manual IPWRA: first estimate the propensity score, then run *weighted* regression adjustment within each treatment arm, using the inverse probability of the *observed* treatment status as weights:

```
clear
webuse cattaneo2
logit mbsmoke mmarried mage prenatal1 fbaby
predict ps

gen wgt1 = mbsmoke/ps
gen wgt0 = (1-mbsmoke)/(1-ps)
reg bweight mmarried mage prenatal1 fbaby if mbsmoke==1 [pweight=wgt1]
predict mu1_ipwra
reg bweight mmarried mage prenatal1 fbaby if mbsmoke==0 [pweight=wgt0]
predict mu0_ipwra
gen tau_ipwra_i = mu1_ipwra - mu0_ipwra
sum tau_ipwra_i
```

```
. clear
```

```
. webuse cattaneo2
```

(Excerpt from Cattaneo (2010) Journal of Econometrics 155: 138-154)

```
. logit mbsmoke mmarried mage prenatal1 fbaby
```

```
Iteration 0:  Log likelihood = -2230.7484
Iteration 1:  Log likelihood = -2089.7838
Iteration 2:  Log likelihood = -2082.3958
Iteration 3:  Log likelihood = -2082.3879
Iteration 4:  Log likelihood = -2082.3879
```

Logistic regression

Number of obs = 4,642

LR chi2(4) = 296.72

Prob > chi2 = 0.0000

Log likelihood = -2082.3879

Pseudo R2 = 0.0665

```
-----
mbsmoke | Coefficient Std. err.      z    P>|z|    [95% conf. interval]
```

mmarried		-1.10666	.091279	-12.12	0.000	-1.285564	-.9277569
mage		-.0187894	.0082217	-2.29	0.022	-.0349036	-.0026751
prenatal1		-.2721103	.0943067	-2.89	0.004	-.4569481	-.0872725
fbaby		-.5428134	.0864561	-6.28	0.000	-.7122642	-.3733625
_cons		.1297324	.2095812	0.62	0.536	-.2810391	.540504

```
. predict ps
(option pr assumed; Pr(mbsmoke))
```

```
.
. gen wgt1 = mbsmoke/ps
```

```
. gen wgt0 = (1-mbsmoke)/(1-ps)
```

```
. reg bweight mmarried mage prenatal1 fbaby if mbsmoke==1 [pweight=wgt1]
(sum of wgt is 4,620.9067401886)
```

Linear regression	Number of obs	=	864
	F(4, 859)	=	3.12
	Prob > F	=	0.0145
	R-squared	=	0.0142
	Root MSE	=	563.01

		Robust				
bweight	Coefficient	std. err.	t	P> t	[95% conf. interval]	
mmarried	130.5978	39.90558	3.27	0.001	52.27399	208.9217
mage	-6.495147	4.95637	-1.31	0.190	-16.22316	3.232866
prenatal1	23.29089	40.04843	0.58	0.561	-55.31335	101.8951
fbaby	45.19338	48.48447	0.93	0.352	-49.96852	140.3553
_cons	3206.397	126.7646	25.29	0.000	2957.592	3455.202

```
. predict mul_ipwra
(option xb assumed; fitted values)
```

```
. reg bweight mmarried mage prenatal1 fbaby if mbsmoke==0 [pweight=wgt0]
(sum of wgt is 4,643.79376959801)
```

Linear regression	Number of obs	=	3,778
	F(4, 3773)	=	26.89
	Prob > F	=	0.0000
	R-squared	=	0.0336
	Root MSE	=	567.52

11. Treatment effects and matching

		Robust				
bweight	Coefficient	std. err.	t	P> t	[95% conf. interval]	
mmarried	160.9954	26.87191	5.99	0.000	108.3105	213.6803
mage	3.061463	2.181356	1.40	0.161	-1.215289	7.338214
prenatal1	62.4539	28.28456	2.21	0.027	6.99938	117.9084
fbaby	-67.38469	20.29791	-3.32	0.001	-107.1806	-27.58876
_cons	3188.522	56.19956	56.74	0.000	3078.338	3298.707

```
. predict mu0_ipwra
(option xb assumed; fitted values)

. gen tau_ipwra_i = mu1_ipwra - mu0_ipwra

. sum tau_ipwra_i
```

Variable	Obs	Mean	Std. dev.	Min	Max
tau_ipwra_i	4,642	-238.7679	97.94189	-481.7334	6.216553

The mean of tau_ipwra_i matches teffects ipwra's ATE.

Manual AIPW: estimate outcome models $\hat{\mu}_1(X)$, $\hat{\mu}_0(X)$ (unweighted, one per arm) and the propensity score $\hat{p}(X)$, then combine via the augmented-IPW formula

$$\hat{\tau}_i^{aipw} = D_i \frac{Y_i - \hat{\mu}_1(X_i)}{\hat{p}(X_i)} + \hat{\mu}_1(X_i) - (1 - D_i) \frac{Y_i - \hat{\mu}_0(X_i)}{1 - \hat{p}(X_i)} - \hat{\mu}_0(X_i) \quad (11.1)$$

```
clear
webuse cattaneo2
logit mbsmoke mmarried mage prenatal1 fbaby
predict ps

reg bweight mmarried mage prenatal1 fbaby if mbsmoke==1
predict mu1_aipw
reg bweight mmarried mage prenatal1 fbaby if mbsmoke==0
predict mu0_aipw
gen aipw_i = mbsmoke*(bweight-mu1_aipw)/ps + mu1_aipw - (1-mbsmoke)*(bweight-mu0_aipw)/(1-ps)
sum aipw_i
```

```
. clear

. webuse cattaneo2
(Excerpt from Cattaneo (2010) Journal of Econometrics 155: 138-154)
```

```
. logit mbsmoke mmarried mage prenatal1 fbaby
```

```
Iteration 0: Log likelihood = -2230.7484
Iteration 1: Log likelihood = -2089.7838
Iteration 2: Log likelihood = -2082.3958
Iteration 3: Log likelihood = -2082.3879
Iteration 4: Log likelihood = -2082.3879
```

Logistic regression

```
Number of obs = 4,642
LR chi2(4)      = 296.72
Prob > chi2     = 0.0000
Pseudo R2      = 0.0665
```

Log likelihood = -2082.3879

mbsmoke	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
mmarried	-1.10666	.091279	-12.12	0.000	-1.285564	-.9277569
mage	-.0187894	.0082217	-2.29	0.022	-.0349036	-.0026751
prenatal1	-.2721103	.0943067	-2.89	0.004	-.4569481	-.0872725
fbaby	-.5428134	.0864561	-6.28	0.000	-.7122642	-.3733625
_cons	.1297324	.2095812	0.62	0.536	-.2810391	.540504

```
. predict ps
(option pr assumed; Pr(mbsmoke))
```

```
.
. reg bweight mmarried mage prenatal1 fbaby if mbsmoke==1
```

Source	SS	df	MS	Number of obs	=	864
Model	4586519.21	4	1146629.8	F(4, 859)	=	3.69
Residual	266914157	859	310726.609	Prob > F	=	0.0055
Total	271500676	863	314601.015	R-squared	=	0.0169
				Adj R-squared	=	0.0123
				Root MSE	=	557.43

bweight	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
mmarried	133.6617	41.33023	3.23	0.001	52.54164	214.7818
mage	-7.370881	3.983409	-1.85	0.065	-15.18924	.4474735
prenatal1	25.11133	43.64639	0.58	0.565	-60.55473	110.7774
fbaby	41.43991	41.82146	0.99	0.322	-40.6443	123.5241
_cons	3227.169	102.3392	31.53	0.000	3026.305	3428.033

```
. predict mul_aipw
(option xb assumed; fitted values)
```

11. Treatment effects and matching

```
. reg bweight mmarried mage prenatal1 fbaby if mbsmoke==0
```

Source	SS	df	MS	Number of obs	=	3,778
Model	37916015	4	9479003.74	F(4, 3773)	=	30.00
Residual	1.1922e+09	3,773	315979.752	Prob > F	=	0.0000
Total	1.2301e+09	3,777	325683.776	R-squared	=	0.0308
				Adj R-squared	=	0.0298
				Root MSE	=	562.12

bweight	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
mmarried	160.9513	24.55483	6.55	0.000	112.8092	209.0933
mage	2.546828	1.935504	1.32	0.188	-1.247908	6.341564
prenatal1	64.40859	25.83668	2.49	0.013	13.75338	115.0638
fbaby	-71.3286	19.4895	-3.66	0.000	-109.5396	-33.11763
_cons	3202.746	51.20076	62.55	0.000	3102.362	3303.13

```
. predict mu0_aipw
(option xb assumed; fitted values)
```

```
. gen aipw_i = mbsmoke*(bweight-mu1_aipw)/ps + mu1_aipw - (1-mbsmoke)*(bweight-
> mu0_aipw)/(1-ps) - mu0_aipw
```

```
. sum aipw_i
```

Variable	Obs	Mean	Std. dev.	Min	Max
aipw_i	4,642	-239.0294	1618.64	-29389.38	16313.36

The mean of `aipw_i` is close to `teffects aipw`'s ATE (small differences are expected: `teffects` estimates the treatment and outcome models jointly by ML, while this is a two-step plug-in version).

11.2. Matching

First let's emphasize that matching usually does not deal with endogenous treatment problem. Some people have the wrong impression that matching resolves the endogeneity problem, but in fact it only helps for selection on observables. If you have selection on unobservables, or unobserved confounders, matching does not help.

Matching aims to balance the distribution of covariates in the treatment and control groups. That's all it does, to make comparisons between apples, not apples to oranges. In that sense, regression does that too. We use regression to adjust for differences between treatment and control. However,

regression sometimes rely on extrapolation too much, when data do not overlap between treatment and control. It's not that matching can do magic on non-overlapping data, but it can make it clear that how bad the non-overlapping problem is. Simply running regression blindly will not have the researchers realize the non-overlapping problem. Combining matching with regression is probably a better idea. That is, running regression on matched sample is recommended.

Here we introduced a few popular matching methods implemented in Stata, namely `nnmatch`, `psmatch`, and `cem`.

11.2.1. Propensity score matching

The ideal situation of matching is exact matching; that is, treatment and control units can match to each other with exactly the same value of all covariates. This is often not possible, as the number of covariates increase and if they are not discrete. In high dimension (many covariates to match), it is very hard or even impossible to find exact matches. Therefore, how to reduce from high dimension to low dimension is key. Rosenbaum and Rubin 1983 proved that propensity score provides the one-dimension representation of high-dimension of covariates. Therefore, we only need to find matches based on propensity scores.

In Stata, teffects `psmatch` can do estimation after matching on propensity scores. Stata's `psmatch2` command has been popular for propensity score matching too. The nice thing of these commands is that it does two steps in one command: first it estimate the logit or probit model for propensity score, then match the treatment and control groups, then estimate the outcome equation on matched sample. The standard errors are correct based on all these steps. If we do this manually, standard errors will not be correct.

11.2.1.1. example

```
clear
webuse cattaneo2
teffects psmatch (bweight) (mbsmoke mmarried c.mage##c.mage fbaby medu), atet
psmatch2 mbsmoke mmarried c.mage##c.mage fbaby medu, logit ties
reg bweight mbsmoke [aweight=_weight]
```

```
. clear
```

```
. webuse cattaneo2
```

(Excerpt from Cattaneo (2010) Journal of Econometrics 155: 138-154)

```
. teffects psmatch (bweight) (mbsmoke mmarried c.mage##c.mage fbaby medu), atet
```

Treatment-effects estimation	Number of obs	=	4,642
Estimator : propensity-score matching	Matches: requested	=	1
Outcome model : matching	min	=	1
Treatment model: logit	max	=	74

11. Treatment effects and matching

		AI robust				
bweight	Coefficient	std. err.	z	P> z	[95% conf. interval]	
ATET						
mbsmoke						
(Smoker						
vs						
Nonsmoker)	-236.7848	26.57789	-8.91	0.000	-288.8765	-184.693

```
. psmatch2 mbsmoke mmarried c.mage##c.mage fbaby medu, logit ties
```

```
Logistic regression                                Number of obs = 4,642
                                                    LR chi2(5)      = 375.00
                                                    Prob > chi2     = 0.0000
Log likelihood = -2043.2504                        Pseudo R2      = 0.0841
```

mbsmoke	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
mmarried	-1.145706	.0918962	-12.47	0.000	-1.32582	-.965593
mage	.321518	.0638472	5.04	0.000	.1963798	.4466563
c.mage#						
c.mage	-.0060368	.0011849	-5.09	0.000	-.0083592	-.0037144
fbaby	-.3864258	.0880445	-4.39	0.000	-.5589898	-.2138618
medu	-.1420833	.0173215	-8.20	0.000	-.1760328	-.1081338
_cons	-2.950915	.8102504	-3.64	0.000	-4.538976	-1.362853

```
. reg bweight mbsmoke [aweight=_weight]
(sum of wgt is 1,728)
```

Source	SS	df	MS	Number of obs	=	3,671
Model	51455506	1	51455506	F(1, 3669)	=	151.87
Residual	1.2431e+09	3,669	338808.237	Prob > F	=	0.0000
				R-squared	=	0.0397
				Adj R-squared	=	0.0395
Total	1.2945e+09	3,670	352736.492	Root MSE	=	582.07

bweight	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
mbsmoke	-236.7848	19.21387	-12.32	0.000	-274.4557	-199.1138
_cons	3374.444	13.58626	248.37	0.000	3347.807	3401.082

In the above example, suppose we are interesting in mbsmoke's effect on bweight. If we match smoking mothers with non-smoking mothers, by age, first baby, and education, the the teffects psmatch gives us the treatment effect on the treated. We can do this in two steps. First, by psmatch2 we generate __weight, which is a number for how many times a control unit is used to match to a treatment unit. Notice that to match teffects result, we use logit, and ties option. By default, teffects psmatch includes all ties (control units that have the same propensity scores that are close enough), but psmatch2 by default only include one. Then in the second step, we run a regression with __weight as the weight. Notice the standard errors will differ. We should use the standard errors reported in teffects.

11.2.2. Nearest neighbor matching

Stata has the nnmatch option for teffects. nnmatch by default uses Mahalanobis distance, it can also specify to have exact matching. The advantage is this is nonparametric.

```
clear
webuse cattaneo2
teffects nnmatch (bweight mmarried mage fbaby medu) (mbsmoke) , atet
teffects psmatch (bweight) (mbsmoke mmarried mage fbaby medu), atet
```

```
. clear
```

```
. webuse cattaneo2
```

```
(Excerpt from Cattaneo (2010) Journal of Econometrics 155: 138-154)
```

```
. teffects nnmatch (bweight mmarried mage fbaby medu) (mbsmoke) , atet
```

```
Treatment-effects estimation      Number of obs      =      4,642
Estimator      : nearest-neighbor matching      Matches: requested =      1
Outcome model  : matching                      min =      1
Distance metric: Mahalanobis                  max =      74
```

```
-----+-----
          |               AI robust
          | Coefficient  std. err.      z    P>|z|      [95% conf. interval]
-----+-----
ATET      |
      mbsmoke |
      (Smoker |
      vs      |
      Nonsmoker) | -239.2433   25.68524   -9.31   0.000   -289.5854   -188.9011
-----+-----
```

11. Treatment effects and matching

```
. teffects psmatch (bweight) (mbsmoke mmarrried mage fbaby medu), atet
```

```
Treatment-effects estimation      Number of obs      =      4,642
Estimator      : propensity-score matching      Matches: requested =      1
Outcome model  : matching                      min =      1
Treatment model: logit                      max =      74
```

```
-----
      |               AI robust
      | Coefficient  std. err.      z    P>|z|      [95% conf. interval]
-----+-----
ATET  |
      | mbsmoke      |
      | (Smoker      |
      | vs          |
      | Nonsmoker)  |      -245.711    26.38675    -9.31    0.000    -297.4281    -193.9939
-----
```

```
.
```

Notice the specification is different in these two models.

11.2.3. Coarsened Exact Matching (CEM)

CEM is not part of teffects. But Stata does have cem package implemented.

```
clear
webuse cattaneo2
cem mmarrried mage fbaby medu, treatment(mbsmoke)
reg bweight mbsmoke [aweight=cem_weights]
* Note: cem creates the variable "cem_weights" (plural). We spell it out
* in full here rather than relying on Stata's automatic variable-name
* abbreviation matching, to keep the reference unambiguous.
```

```
. clear
```

```
. webuse cattaneo2
```

(Excerpt from Cattaneo (2010) Journal of Econometrics 155: 138-154)

```
. cem mmarrried mage fbaby medu, treatment(mbsmoke)
```

Matching Summary:

Number of strata: 274

Number of matched strata: 135

11.3. Modern matching in R: balance diagnostics and entropy weighting

```

      0      1
All   3778   864
Matched 3421   827
Unmatched 357   37

```

Multivariate L1 distance: .25424705

Univariate imbalance:

	L1	mean	min	25%	50%	75%	max
mmarried	6.4e-15	8.4e-15	0	0	0	0	0
mage	.04955	-.00887	1	0	1	0	-2
fbaby	6.1e-15	6.1e-15	0	0	0	0	0
medu	.03809	-.03316	0	0	0	0	0

```

. reg bweight mbsmoke [aweight=cem_weights]
(sum of wgt is 4,248.000000000005)

```

Source	SS	df	MS	Number of obs	=	4,248
Model	40566772.9	1	40566772.9	F(1, 4246)	=	121.12
Residual	1.4221e+09	4,246	334928.257	Prob > F	=	0.0000
				R-squared	=	0.0277
				Adj R-squared	=	0.0275
Total	1.4627e+09	4,247	344401.26	Root MSE	=	578.73

bweight	Coefficient	Std. err.	t	P> t	[95% conf. interval]
mbsmoke	-246.8017	22.42533	-11.01	0.000	-290.7671 -202.8364
_cons	3383.394	9.894625	341.94	0.000	3363.996 3402.793

```

. * Note: cem creates the variable "cem_weights" (plural). We spell it out
. * in full here rather than relying on Stata's automatic variable-name
. * abbreviation matching, to keep the reference unambiguous.
.

```

In the above example, we use CEM on the same covariates, then apply weights in the final regression. The disadvantage is that we have not accounted for the noise in calculating those weights.

11.3. Modern matching in R: balance diagnostics and entropy weighting

The 2019 vintage of this chapter used Stata's `teffects` throughout. R's `MatchIt` and `WeightIt` packages have since become the standard for matching workflows, and `cobalt` provides publication-quality balance tables and plots that are now expected in applied papers.

11. Treatment effects and matching

11.3.1. Balance checking with cobalt

After any matching or weighting step, always inspect covariate balance. `cobalt`'s `bal.tab()` and `love.plot()` are the go-to tools:

```
library(MatchIt)
library(cobalt)

# Propensity score matching on the lalonde data
data("lalonde", package = "MatchIt")

m_out <- matchit(treat ~ age + educ + race + married + nodegree +
                 re74 + re75,
                 data = lalonde, method = "nearest", distance = "glm",
                 ratio = 1)

# Balance table: standardized mean differences before and after matching
bal.tab(m_out, stats = c("mean.diffs", "variance.ratios"), thresholds = c(m = 0.1))
```

Balance Measures

	Type	Diff.Adj	M.Threshold	V.Ratio.Adj
distance	Distance	0.9739		0.7566
age	Contin.	0.0718	Balanced, <0.1	0.4568
educ	Contin.	-0.1290	Not Balanced, >0.1	0.5721
race_black	Binary	0.3730	Not Balanced, >0.1	.
race_hispan	Binary	-0.1568	Not Balanced, >0.1	.
race_white	Binary	-0.2162	Not Balanced, >0.1	.
married	Binary	-0.0216	Balanced, <0.1	.
nodegree	Binary	0.0703	Balanced, <0.1	.
re74	Contin.	-0.0505	Balanced, <0.1	1.3289
re75	Contin.	-0.0257	Balanced, <0.1	1.4956

Balance tally for mean differences

	count
Balanced, <0.1	5
Not Balanced, >0.1	4


Variable with the greatest mean difference

Variable	Diff.Adj	M.Threshold
race_black	0.373	Not Balanced, >0.1

Sample sizes

	Control	Treated
All	429	185
Matched	185	185
Unmatched	244	0

```
# Love plot: visualize balance improvement
love.plot(m_out, stats = "mean.diffs", threshold = 0.1,
          abs = TRUE, var.order = "unadjusted")
```



The love plot shows each covariate’s standardized mean difference (SMD) before matching (open circles) and after (filled circles). The vertical line at SMD = 0.1 is the conventional threshold for “balanced.” If post-matching points cross the line, the matching is not improving balance and a different specification is needed.

That is exactly what happens here, which is why it is worth running the diagnostic rather than assuming. 1:1 nearest-neighbour matching on the propensity score leaves four of the nine covariates *above* the 0.1 threshold, **race** worst of all (an adjusted SMD of 0.37 for the black indicator), while discarding 244 of the 429 controls. Matching on a propensity score balances the score, not necessarily the covariates that went into it.

11.3.2. Entropy balancing

Propensity score matching discards unmatched units and can reduce effective sample size substantially. **Entropy balancing** (Hainmueller 2012) keeps all units and reweights the control group so that weighted covariate moments exactly match the treatment group — without discarding anyone.

```
library(WeightIt)
library(cobalt)

# Entropy balancing: exact balance on means, variances, and covariances
w_out <- weightit(treat ~ age + educ + race + married + nodegree +
                  re74 + re75,
```

11. Treatment effects and matching

```
data = lalonde, method = "ebal", estimand = "ATT",
moments = 1) # moments=1: balance means only
              # moments=2: balance means + variances

summary(w_out)
```

Summary of weights

- Weight ranges:

	Min	Max
treated 1.	1.	1.
control 0.008	0.008	4.062

- Units with the 5 most extreme weights by group:

	NSW5	NSW4	NSW3	NSW2	NSW1
treated	1	1	1	1	1
	PSID423	PSID196	PSID412	PSID118	PSID226
control	3.073	3.235	3.45	3.897	4.062

- Weight statistics:

	Coef of Var	MAD	Entropy	# Zeros
treated	0.	0.	0.	0
control	1.834	1.287	1.101	0

- Effective Sample Sizes:

	Control	Treated
Unweighted	429.	185
Weighted	98.46	185

```
# Check balance
bal.tab(w_out, stats = "mean.diffs", thresholds = c(m = 0.1))
```

Balance Measures

	Type	Diff.Adj	M.Threshold
age	Contin.	-0	Balanced, <0.1
educ	Contin.	-0	Balanced, <0.1
race_black	Binary	0	Balanced, <0.1
race_hispan	Binary	0	Balanced, <0.1
race_white	Binary	-0	Balanced, <0.1
married	Binary	-0	Balanced, <0.1
nodegree	Binary	0	Balanced, <0.1
re74	Contin.	-0	Balanced, <0.1

11.3. Modern matching in R: balance diagnostics and entropy weighting

```
re75          Contin.          -0 Balanced, <0.1
```

Balance tally for mean differences

	count
Balanced, <0.1	9
Not Balanced, >0.1	0


Variable with the greatest mean difference

Variable	Diff.Adj	M.Threshold
re74	-0	Balanced, <0.1

Effective sample sizes

	Control	Treated
Unadjusted	429.	185
Adjusted	98.46	185

```
love.plot(w_out, threshold = 0.1, abs = TRUE)
```



```
# Weighted outcome regression
library(lmtest); library(sandwich)
m_eb <- lm(re78 ~ treat + age + educ + race + married + nodegree +
           re74 + re75,
           data = lalonde, weights = w_out$weights)
coeftest(m_eb, vcov = vcovHC(m_eb, type = "HC3"))["treat", ]
```

Estimate	Std. Error	t value	Pr(> t)
1273.2618139	820.1848528	1.5524083	0.1210883

11. Treatment effects and matching

Entropy balancing delivers on its promise here: the largest absolute standardized mean difference after weighting is **exactly zero**, since the weights are chosen to satisfy the mean-balance constraints as equalities. All 429 controls are retained.

Method	Pros	Cons
PSM (nearest neighbor)	Intuitive; respects overlap	Discards units; balance not guaranteed
CEM	Exact balance in bins	Requires discretizing continuous vars
Entropy balancing	Exact mean balance; keeps all units	Weights can be extreme; no variance balance by default
IPW (logistic)	Simple; relies only on the treatment model	Not doubly robust on its own; extreme weights if overlap is poor
AIPW / IPWRA	Doubly robust: consistent if <i>either</i> the treatment or outcome model is correct	Needs both models specified; more moving parts

For most applications, entropy balancing or IPW with trimmed weights (via `WeightIt`'s `trim()`) is preferred over PSM: better effective sample size, guaranteed balance, and easy integration with doubly-robust estimators.

See also the [weighting chapter](#) for IPW estimators and the [OLS-ATE chapter](#) for regression-based approaches.

Systematic treatment: [R](#) · [Julia](#).

12. When is OLS coefficient the ATE

12.1. OLS cannot be ATE (most of the time)

Tymon Słoczyński has this [cool paper](#) which helps me understand OLS better.

A common setup in empirical studies:

$$y = \alpha + \tau d + X\beta \quad (12.1)$$

d is the treatment dummy, X is the covariates. This model assumes homogeneous treatment effect. That is, τ is the ATE if treatment effect is homogeneous. But, if not, then τ is not ATE anymore, although most people treat it as ATE. Obviously we are assuming the usual unconfoundedness; that is, there are no other variables other than X are driving both d and y .

In this paper, Słoczyński showed that τ is actually a convex combination of ATT and ATU (average treatment effect for the untreated), if treatment has heterogeneous effects. Therefore it can be close to ATE, or it can be quite different, depending on the weight on ATT and ATU. More surprisingly, he showed that the weights are inversely proportional to the proportion of treated vs untreated. The more number of treated units, the less weight is on ATT!

By definition,

$$\tau_{ATE} = \rho\tau_{ATT} + (1 - \rho)\tau_{ATU} \quad (12.2)$$

The main result from the paper is:

$$\tau = (1 - \rho)\tau_{ATT} + \rho\tau_{ATU} \quad (12.3)$$

where $\rho = P(d = 1)$, the proportion of treated; $\tau_{ATT} = E(y(1) - y(0)|d = 1)$, $\tau_{ATU} = E(y(1) - y(0)|d = 0)$, and $\tau_{ATE} = E(y(1) - y(0))$.

From these two equations, we see that τ and τ_{ATE} are quite different, unless the weight on ATT happens to equal ρ . A caution on reading this too simply: the clean form $\tau = (1 - \rho)\tau_{ATT} + \rho\tau_{ATU}$ with $\rho = P(d = 1)$ is the **no-covariate special case** (or, with covariates, holds in terms of *covariate-adjusted* treatment-probability weights, not the raw marginal share). In Słoczyński's general result the weights depend on the conditional treatment probabilities (the propensity score), so a 50/50 *marginal* treatment share alone does **not** guarantee that OLS recovers the ATE — what matters is the covariate-adjusted weighting. With that caveat, the intuition stands: the more imbalanced the (conditional) treatment probabilities, the further OLS can be from the ATE.

12. When is OLS coefficient the ATE

These results are based on the fact ATT, ATU and ATE can be calculated if we have the true propensity score, then we can estimate the heterogeneous effect by including the propensity score as an additional covariate. If we don't, then we can include the estimated propensity score as an additional covariate.

Słoczyński has a package in R and stata "hettreatreg".

```
library(hettreatreg)
data("nswcps")

summary(lm(re78 ~ treated + age + age2 + educ + black + hispanic + married + nodegree + re74 + re75, data = nswcps))
```

Call:

```
lm(formula = re78 ~ treated + age + age2 + educ + black + hispanic + married + nodegree + re74 + re75, data = nswcps)
```

Residuals:

Min	1Q	Median	3Q	Max
-25130	-3601	1274	3668	55040

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	7634.34415	736.67074	10.363	< 2e-16 ***
treated	793.58704	548.25433	1.447	0.14778
age	-233.67749	41.18067	-5.674	1.42e-08 ***
age2	1.81437	0.56099	3.234	0.00122 **
educ	166.84923	28.65984	5.822	5.94e-09 ***
black	-790.60856	213.24523	-3.708	0.00021 ***
hispanic	-175.97512	218.99126	-0.804	0.42166
married	224.26599	149.84542	1.497	0.13450
nodegree	311.84453	178.51743	1.747	0.08068 .
re74	0.29534	0.01222	24.175	< 2e-16 ***
re75	0.47064	0.01216	38.700	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 7002 on 16166 degrees of freedom

Multiple R-squared: 0.4762, Adjusted R-squared: 0.4758

F-statistic: 1469 on 10 and 16166 DF, p-value: < 2.2e-16

```
outcome <- nswcps$re78
treated <- nswcps$treated
our_vars <- c("age", "age2", "educ", "black", "hispanic", "married", "nodegree", "re74", "re75")
covariates <- subset(nswcps, select = our_vars)

hettreatreg(outcome, treated, covariates, verbose = TRUE)
```

"OLS" is the estimated regression coefficient on treated.

OLS = 793.6

P(d=1) = 0.011

P(d=0) = 0.989

w1 = 0.983

w0 = 0.017

delta = -0.971

ATE = -6751

ATT = 928.4

ATU = -6840

OLS = w1*ATT + w0*ATU = 793.6

Słoczyński's paper itself is an algebraic decomposition of the OLS coefficient τ in terms of ATT/ATU weights; it does not propose or endorse the propensity-score-as-covariate estimator below as a general-purpose way to recover ATT/ATU/ATE. What follows is our own illustrative exercise, done in three steps, to see how such an estimator would behave in practice: First, estimate the propensity score equation, basically

$$d = X\gamma \tag{12.4}$$

Second, include the predicted propensity score as a covariate and estimate the model with treatment and control group separately. Note that regressing Y linearly on the scalar propensity score $p(X)$ assumes $E[Y(d) \mid p(X)]$ is linear in $p(X)$ — a strong and generally untested assumption, unlike matching/weighting directly on the propensity score, or regression adjustment on the covariates X themselves, which are the standard approaches. Treat the numbers below as an exploratory illustration rather than a recommended estimator.

Third, predict the counterfactuals for the full sample. Then we get the ATE, ATT and ATU.

One more reason to treat this as illustration rather than method: the propensity score here comes from a linear probability model, and on these data it runs from -0.034 to 0.154 — some fitted “probabilities” are negative. That is harmless when the score is used merely as a regressor, as it is here, but it would be fatal if the same score were used as an inverse-probability weight.

Let's see with the same data:

```
m_propensity <- lm(treated ~ age + age2 + educ + black + hispanic + married + nodegree + re74
ps <- predict(m_propensity)
df_combined <- data.frame(nswcps, ps)
df.1 <- df_combined[which(treated == 1),]
df.0 <- df_combined[which(treated == 0),]
```

12. When is OLS coefficient the ATE

```
m_ot <- lm(re78 ~ ps, data = df.1)
ot <- suppressWarnings(predict(m_ot, newdata = df_combined)) # add newdata option to predict f

m_oc <- lm(re78 ~ ps, data = df.0)
oc <- suppressWarnings(predict(m_oc, newdata = df_combined)) # add newdata option to predict f

te <- ot - oc
ate <- mean(te)
df_combined$te <- te
att <- as.numeric(mean(df_combined[which(treated == 1), 'te']))
atu <- as.numeric(mean(df_combined[which(treated == 0), 'te']))

print(paste("ate=", signif(ate, 4)))
```

```
[1] "ate= -6751"
```

```
print(paste("att=", signif(att, 4)))
```

```
[1] "att= 928.4"
```

```
print(paste("atu=", signif(atu, 4)))
```

```
[1] "atu= -6840"
```

We see in this case OLS coefficient is hugely different from ATE. This is because $P(d = 1)$ is only .01, about 1 percent; we have much larger control group than treatment group. In this case, the OLS coefficient is far different from ATE; in fact, it's pretty close to ATT.

Now, what's the intuition that OLS coefficient is just the opposite as the ATE in terms of weighting ATT and ATU? The reason is that OLS is to predict actual outcome, not counterfactual outcome. For calculating ATE, we need counterfactual outcome predicted. Now suppose we have a lot of control observations, we'd much better predicting $(y(0)|d == 0)$ rather than $(y(1)|d == 1)$. That is, we are better getting the slope of $(y(0)|d == 0, X = x)$; therefore predicting $(y(0)|d == 1, X = x)$. Therefore although treatment group has less observations, but because the $(y(0)|d == 1)$ is better predicted, the ATT is better estimated, thus more weight. If the opposite happens, then the ATU should be more heavily weighted. Thus the counter-intuitive weighting.

12.2. Regression adjustment

Słoczyński's method is to include propensity score as a covariate. This propensity score serves as a proxy as all covariates. As Rosenbaum and Rubin (1983) showed, the propensity score is a balancing score (treatment is independent of the covariates given the propensity score). Another way to allow treatment effect to differ across all covariates is to allow interaction of treatment dummy with all covariates, then I am doing "Regression Adjustment" (see Stata's `teffects ra`,

used in the [treatment effects and matching chapter](#); note this is not `treatreg`, which was the old endogenous-treatment command, superseded by `etregress`). We can calculate ATE, ATT, ATU after the linear model.

```
library(tidyverse)
lm1 <- lm(re78 ~ treated*(age + age2 + educ + black + hispanic + married + nodegree + re74 +
y1 <- predict(lm1,newdata=nswcps %>% mutate(treated=1))
y0 <- predict(lm1,newdata=nswcps %>% mutate(treated=0))
ate=mean(y1-y0) #ate
print(paste("ate=", signif(ate, 4)))
```

```
[1] "ate= -4930"
```

```
y1.att <- predict(lm1,newdata=nswcps %>% filter(treated==1))
y0.att <- predict(lm1,newdata=nswcps %>% filter(treated==1) %>% mutate(treated=0))
att=mean(y1.att)-mean(y0.att) #att
print(paste('att=', signif(att, 4)))
```

```
[1] "att= 796"
```

```
y1.atu <- predict(lm1,newdata=nswcps %>% filter(treated==0) %>% mutate(treated=1))
y0.atu <- predict(lm1,newdata=nswcps %>% filter(treated==0))
atu=mean(y1.atu)-mean(y0.atu) #atu
print(paste('atu=', signif(atu, 4)))
```

```
[1] "atu= -4996"
```

We see this is somewhat different from Słoczyński's method. But the point is the original OLS coefficient 794 is nowhere close to ATE. In this case, it's very close to ATT, as the treatment group is very small.

Systematic treatment: [R](#) · [Julia](#).

13. Matching and Weighting Part 1: matching

This is based on R's MatchIt and WeightIt packages and their documentations. All credits go to Noah Greifer and coauthors. I write to summarize what I have learned.

13.1. Assumptions

Matching or weighting is usually used in observational studies. We have the following assumptions before using matching or weighting to make a causal claim:

1. SUTVA
2. ignorability (unconfoundedness, or no unobserved confounders)
3. overlap (common support, or positivity)

With these assumptions, we can try to identify the treatment effect in an observational study.

13.2. Problem

The second assumption is saying controlling for X , the treatment assignment D is independent of the potential outcomes $Y(0)$ and $Y(1)$, i.e., we have a pseudo-randomized experiment.

In reality, it's rare we have a balanced distribution of X for treated and control. That's when we need to use matching or weighting to balance the distribution of X between treatment and control groups.

13.3. Distance

The idea of matching is to find a set of control units that are similar(close) to the treatment units in terms of X . The first question is: how to define "close"? There are several distance measures.

13.3.1. Propensity score

The most common distance measure is the propensity score (PS), which is the probability of being treated given X . The benefit is that when there is high dimension of X , we can reduce the dimension to a single number, the PS.

There are many ways to estimate the propensity score, such as logit, or more flexible parametric "gam", or other machine learning methods, such as "gbm", "lasso", "rpart", "randomforest", "cbps", etc. Here CBPS has seen more popularity. The propensity scores are estimated using the

13. Matching and Weighting Part 1: matching

covariate balancing propensity score (CBPS) algorithm, which is a form of logistic regression where balance constraints are incorporated into a generalized method of moments estimation of the model coefficients.

13.3.2. Distance from covariates

Or we can compute distance from covariates directly, not using propensity score. The distance can be computed using Mahalanobis distance, or Euclidean distance, or other distance measures.

13.4. matching methods

With a distance, we then decide which method to use. The first few matching methods all need to specify a distance measure. They are matching based on distance. Then there are stratum matching methods, which do not need a distance measure. The last two methods are subsetting methods.

13.4.1. Nearest neighbor matching

NNM is widely used; it is also known as greedy matching. It matches each treated unit to the closest control unit in terms of distance. The closest control unit is defined as the one with the smallest distance to the treated unit. If there are multiple control units with the same distance and we want a 1:1 match, one is randomly selected.

13.4.2. Optimal pair matching

NNM is not optimal in the sense that it does not necessarily minimize the total distance between matched pairs. Optimal pair matching is a method that minimizes the total distance (or some other overall criterion) between matched pairs. It is more computationally intensive than NNM, but it can lead to better balance.

13.4.3. Optimal full matching

OFM assigns each unit (treated or control) to a subclass. For each subclass then to calculate the sum of within-subclass distances, and then to minimize the total distance across all subclasses. It is more computationally intensive than NNM and optimal pair matching. Weights are used based on subclass membership. Because all units are included, it can be used to estimate ATE.

13.4.4. Generalized Full matching

GFM is a variant of OFM. It uses a faster algorithm for large datasets.

13.4.5. Genetic matching

Genetic matching (`method = "genetic"`, Diamond and Sekhon 2013) runs nearest neighbor matching on a *scaled* generalized Mahalanobis distance, where the scaling weights on the covariates are not fixed in advance but chosen by a genetic algorithm that searches for the weighting giving the best covariate balance. That search is the point of the method — and also why it is far slower than the alternatives above.

13.4.6. Exact matching

It's a stratum matching. First subclasses are defined, by combinations of covariate values. Then units are assigned to subclasses. When a subclass has only either control or treated units, it is discarded.

Exact matching is nonparametric, but it works only when covariates are discrete and it happens that there are enough treated and control units in each subclass.

13.4.7. Coarsened exact matching

CEM is a more popular stratum matching. First bins are created after coarsening the covariates. then use exact matching on the covariate bins. CEM is nonparametric. User needs to decide how many bins to create for each covariate.

13.4.8. Subclassification

Subclassification is another kind of stratum matching. It uses the propensity score to define bins, usually quantiles of PS in the treated group, or control group or overall, depending on target estimand.

13.4.9. Cardinality and profile matching

Cardinality matching involves finding the largest sample that satisfies user-supplied balance constraints and constraints on the ratio of matched treated to matched control units. Profile matching involves identifying a target distribution (e.g., the full sample for the ATE or the treated units for the ATT) and finding the largest subset of the treated and control groups that satisfy user-supplied balance constraints with respect to that target.

13.5. Example

Here is an example with Lalonde data, we are interested in the effect of treatment on “re78” (1978 real earnings).

13. Matching and Weighting Part 1: matching

```
library("MatchIt")
data("lalonge")

head(lalonge)
```

	treat	age	educ	race	married	nodegree	re74	re75	re78
NSW1	1	37	11	black	1	1	0	0	9930.0460
NSW2	1	22	9	hispan	0	1	0	0	3595.8940
NSW3	1	30	12	black	0	0	0	0	24909.4500
NSW4	1	27	11	black	0	1	0	0	7506.1460
NSW5	1	33	8	black	0	1	0	0	289.7899
NSW6	1	22	9	black	0	1	0	0	4056.4940

```
# No matching; constructing a pre-match matchit object
m.out0 <- matchit(treat ~ age + educ + race + married +
  nodegree + re74 + re75,
  data = lalonge,
  method = NULL,
  distance = "glm")

summary(m.out0)
```

Call:

```
matchit(formula = treat ~ age + educ + race + married + nodegree +
  re74 + re75, data = lalonge, method = NULL, distance = "glm")
```

Summary of Balance for All Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.5774	0.1822	1.7941	0.9211	0.3774
age	25.8162	28.0303	-0.3094	0.4400	0.0813
educ	10.3459	10.2354	0.0550	0.4959	0.0347
raceblack	0.8432	0.2028	1.7615	.	0.6404
racehispan	0.0595	0.1422	-0.3498	.	0.0827
racewhite	0.0973	0.6550	-1.8819	.	0.5577
married	0.1892	0.5128	-0.8263	.	0.3236
nodegree	0.7081	0.5967	0.2450	.	0.1114
re74	2095.5737	5619.2365	-0.7211	0.5181	0.2248
re75	1532.0553	2466.4844	-0.2903	0.9563	0.1342

	eCDF Max
distance	0.6444
age	0.1577
educ	0.1114
raceblack	0.6404
racehispan	0.0827
racewhite	0.5577
married	0.3236

```

nodegree    0.1114
re74        0.4470
re75        0.2876

```

Sample Sizes:

	Control	Treated
All	429	185
Matched	429	185
Unmatched	0	0
Discarded	0	0

Or use “cobalt” to see the balance:

```

library("cobalt")
bal.tab(treat ~ age + educ + race + married + nodegree + re74 + re75,
        data = lalonde, estimand = "ATT", thresholds = c(m = .05))

```

Balance Measures

	Type	Diff.Un	M.Threshold.Un
age	Contin.	-0.3094	Not Balanced, >0.05
educ	Contin.	0.0550	Not Balanced, >0.05
race_black	Binary	0.6404	Not Balanced, >0.05
race_hispan	Binary	-0.0827	Not Balanced, >0.05
race_white	Binary	-0.5577	Not Balanced, >0.05
married	Binary	-0.3236	Not Balanced, >0.05
nodegree	Binary	0.1114	Not Balanced, >0.05
re74	Contin.	-0.7211	Not Balanced, >0.05
re75	Contin.	-0.2903	Not Balanced, >0.05

Balance tally for mean differences

	count
Balanced, <0.05	0
Not Balanced, >0.05	9

Variable with the greatest mean difference

Variable	Diff.Un	M.Threshold.Un
re74	-0.7211	Not Balanced, >0.05

Sample sizes

	Control	Treated
All	429	185

First do a PS match.

13. Matching and Weighting Part 1: matching

```
# Matching on logit of a PS estimated with logistic
# regression:
m.out1 <- matchit(treat ~ age + educ + race + married +
                  nodegree + re74 + re75,
                  data = lalonde,
                  distance = "glm",
                  link = "linear.logit")
summary(m.out1)
```

Call:

```
matchit(formula = treat ~ age + educ + race + married + nodegree +
        re74 + re75, data = lalonde, distance = "glm", link = "linear.logit")
```

Summary of Balance for All Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.2657	-2.1500	2.1426	0.5398	0.3774
age	25.8162	28.0303	-0.3094	0.4400	0.0813
educ	10.3459	10.2354	0.0550	0.4959	0.0347
raceblack	0.8432	0.2028	1.7615	.	0.6404
racehispan	0.0595	0.1422	-0.3498	.	0.0827
racewhite	0.0973	0.6550	-1.8819	.	0.5577
married	0.1892	0.5128	-0.8263	.	0.3236
nodegree	0.7081	0.5967	0.2450	.	0.1114
re74	2095.5737	5619.2365	-0.7211	0.5181	0.2248
re75	1532.0553	2466.4844	-0.2903	0.9563	0.1342

eCDF Max

distance	0.6444
age	0.1577
educ	0.1114
raceblack	0.6404
racehispan	0.0827
racewhite	0.5577
married	0.3236
nodegree	0.1114
re74	0.4470
re75	0.2876

Summary of Balance for Matched Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.2657	-0.7706	0.9192	0.7712	0.1321
age	25.8162	25.3027	0.0718	0.4568	0.0847
educ	10.3459	10.6054	-0.1290	0.5721	0.0239
raceblack	0.8432	0.4703	1.0259	.	0.3730
racehispan	0.0595	0.2162	-0.6629	.	0.1568
racewhite	0.0973	0.3135	-0.7296	.	0.2162
married	0.1892	0.2108	-0.0552	.	0.0216


nodegree	0.7081	0.6378	0.1546	.	0.0703
re74	2095.5737	2342.1076	-0.0505	1.3289	0.0469
re75	1532.0553	1614.7451	-0.0257	1.4956	0.0452
	eCDF	Max	Std.	Pair	Dist.
distance	0.4216		0.9193		
age	0.2541		1.3938		
educ	0.0757		1.2474		
raceblack	0.3730		1.0259		
racehispan	0.1568		1.0743		
racewhite	0.2162		0.8390		
married	0.0216		0.8281		
nodegree	0.0703		1.0106		
re74	0.2757		0.7965		
re75	0.2054		0.7381		

Sample Sizes:

	Control	Treated
All	429	185
Matched	185	185
Unmatched	244	0
Discarded	0	0


```
plot(m.out1, type = "jitter", interactive = FALSE)
```

Distribution of Propensity Scores



```
plot(summary(m.out1))
```

13. Matching and Weighting Part 1: matching



We see there are still some imbalance on variables such as “raceblack”. We can use more flexible models for the PS.

```
# GAM logistic PS with smoothing splines (s()):
m.out2 <- matchit(treat ~ s(age) + s(educ) +
  race + married +
  nodegree + re74 + re75,
  data = lalonde,
  distance = "gam")
summary(m.out2)
```

Call:

```
matchit(formula = treat ~ s(age) + s(educ) + race + married +
  nodegree + re74 + re75, data = lalonde, distance = "gam")
```

Summary of Balance for All Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.6566	0.1481	1.9123	1.6161	0.4155
raceblack	0.8432	0.2028	1.7615	.	0.6404
racehispan	0.0595	0.1422	-0.3498	.	0.0827
racewhite	0.0973	0.6550	-1.8819	.	0.5577
married	0.1892	0.5128	-0.8263	.	0.3236
nodegree	0.7081	0.5967	0.2450	.	0.1114
re74	2095.5737	5619.2365	-0.7211	0.5181	0.2248
re75	1532.0553	2466.4844	-0.2903	0.9563	0.1342
age	25.8162	28.0303	-0.3094	0.4400	0.0813
educ	10.3459	10.2354	0.0550	0.4959	0.0347

	eCDF Max
distance	0.7172
raceblack	0.6404
racehispan	0.0827
racewhite	0.5577
married	0.3236
nodegree	0.1114
re74	0.4470
re75	0.2876
age	0.1577
educ	0.1114

Summary of Balance for Matched Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.6566	0.3055	1.3204	1.2324	0.1853
raceblack	0.8432	0.4432	1.1002	.	0.4000
racehispan	0.0595	0.1730	-0.4800	.	0.1135
racewhite	0.0973	0.3838	-0.9667	.	0.2865
married	0.1892	0.3568	-0.4278	.	0.1676
nodegree	0.7081	0.6541	0.1189	.	0.0541
re74	2095.5737	3521.7556	-0.2919	1.0512	0.1346
re75	1532.0553	2056.2000	-0.1628	1.2120	0.0917
age	25.8162	25.7081	0.0151	0.7958	0.0270
educ	10.3459	10.1568	0.0941	0.6186	0.0270


	eCDF Max	Std. Pair Dist.
distance	0.5189	1.3204
raceblack	0.4000	1.1002
racehispan	0.1135	0.8914
racewhite	0.2865	1.1491
married	0.1676	1.0627
nodegree	0.0541	0.9749
re74	0.4270	0.8894
re75	0.2378	0.8453
age	0.0865	1.1619
educ	0.0919	1.2447

Sample Sizes:


	Control	Treated
All	429	185
Matched	185	185
Unmatched	244	0
Discarded	0	0

```
plot(m.out2, type = "jitter", interactive = FALSE)
```

Distribution of Propensity Scores



```
plot(summary(m.out2))
```



Or use CBPS.

```
# CBPS for ATC matching w/replacement, using the just-
# identified version of CBPS (setting method = "exact"):
m.out3 <- matchit(treat ~ age + educ + race + married +
  nodegree + re74 + re75,
```

```

data = lalonde,
distance = "cbps",
estimand = "ATC",
distance.options = list(method = "exact"),
replace = TRUE)
summary(m.out3)

```

Call:

```

matchit(formula = treat ~ age + educ + race + married + nodegree +
  re74 + re75, data = lalonde, distance = "cbps", distance.options = list(method = "exact"),
  estimand = "ATC", replace = TRUE)

```

Summary of Balance for All Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.2801	0.7758	-1.5464	0.9710	0.3525
age	25.8162	28.0303	-0.2053	0.4400	0.0813
educ	10.3459	10.2354	0.0387	0.4959	0.0347
raceblack	0.8432	0.2028	1.5928	.	0.6404
racehispan	0.0595	0.1422	-0.2369	.	0.0827
racewhite	0.0973	0.6550	-1.1732	.	0.5577
married	0.1892	0.5128	-0.6475	.	0.3236
nodegree	0.7081	0.5967	0.2270	.	0.1114
re74	2095.5737	5619.2365	-0.5190	0.5181	0.2248
re75	1532.0553	2466.4844	-0.2838	0.9563	0.1342

	eCDF Max
distance	0.5967
age	0.1577
educ	0.1114
raceblack	0.6404
racehispan	0.0827
racewhite	0.5577
married	0.3236
nodegree	0.1114
re74	0.4470
re75	0.2876

Summary of Balance for Matched Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.7720	0.7758	-0.0117	1.1568	0.0499
age	27.3520	28.0303	-0.0629	0.4383	0.1091
educ	11.0117	10.2354	0.2719	0.7323	0.0566
raceblack	0.3100	0.2028	0.2667	.	0.1072
racehispan	0.1096	0.1422	-0.0934	.	0.0326
racewhite	0.5804	0.6550	-0.1569	.	0.0746
married	0.5967	0.5128	0.1679	.	0.0839
nodegree	0.4499	0.5967	-0.2994	.	0.1469

13. Matching and Weighting Part 1: matching


re74	5696.9438	5619.2365	0.0114	0.8562	0.0994
re75	1935.6090	2466.4844	-0.1613	1.5194	0.1071
	eCDF	Max	Std.	Pair	Dist.
distance	0.3543			0.0194	
age	0.2191			0.8676	
educ	0.2727			1.0719	
raceblack	0.1072			0.3826	
racehispan	0.0326			0.6407	
racewhite	0.0746			0.5492	
married	0.0839			0.5596	
nodegree	0.1469			1.0692	
re74	0.2541			0.7513	
re75	0.2727			0.9781	

Sample Sizes:


	Control	Treated
All	429	185.
Matched (ESS)	429	6.67
Matched	429	93.
Unmatched	0	92.
Discarded	0	0.

```
plot(m.out3, type = "jitter", interactive = FALSE)
```

Distribution of Propensity Scores



```
plot(summary(m.out3))
```



```
# Mahalanobis distance matching - no PS estimated
m.out4 <- matchit(treat ~ age + educ + race + married +
  nodegree + re74 + re75,
  data = lalonde,
  distance = "mahalanobis")
summary(m.out4)
```

Call:

```
matchit(formula = treat ~ age + educ + race + married + nodegree +
  re74 + re75, data = lalonde, distance = "mahalanobis")
```

Summary of Balance for All Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
age	25.8162	28.0303	-0.3094	0.4400	0.0813
educ	10.3459	10.2354	0.0550	0.4959	0.0347
raceblack	0.8432	0.2028	1.7615	.	0.6404
racehispan	0.0595	0.1422	-0.3498	.	0.0827
racewhite	0.0973	0.6550	-1.8819	.	0.5577
married	0.1892	0.5128	-0.8263	.	0.3236
nodegree	0.7081	0.5967	0.2450	.	0.1114
re74	2095.5737	5619.2365	-0.7211	0.5181	0.2248
re75	1532.0553	2466.4844	-0.2903	0.9563	0.1342

	eCDF Max
age	0.1577
educ	0.1114
raceblack	0.6404

13. Matching and Weighting Part 1: matching

```
racehispan  0.0827
racewhite   0.5577
married     0.3236
nodegree    0.1114
re74        0.4470
re75        0.2876
```

Summary of Balance for Matched Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
age	25.8162	24.9081	0.1269	0.5327	0.0695
educ	10.3459	10.4324	-0.0430	0.8345	0.0108
raceblack	0.8432	0.4649	1.0407	.	0.3784
racehispan	0.0595	0.0595	0.0000	.	0.0000
racewhite	0.0973	0.4757	-1.2767	.	0.3784
married	0.1892	0.2486	-0.1518	.	0.0595
nodegree	0.7081	0.6595	0.1070	.	0.0486
re74	2095.5737	3305.4420	-0.2476	0.8705	0.1012
re75	1532.0553	1957.5097	-0.1322	1.0071	0.0654

	eCDF Max	Std. Pair Dist.
age	0.2378	0.6195
educ	0.0486	0.2903
raceblack	0.3784	1.0407
racehispan	0.0000	0.0000
racewhite	0.3784	1.2767
married	0.0595	0.1794
nodegree	0.0486	0.1070
re74	0.3351	0.4546
re75	0.2162	0.4113



Sample Sizes:

	Control	Treated
All	429	185
Matched	185	185
Unmatched	244	0
Discarded	0	0


```
m.out4$distance #NULL
```

NULL


```
plot(m.out4, interactive = FALSE)
```



13. Matching and Weighting Part 1: matching



```
plot(summary(m.out4))
```



```
# Full matching on a probit PS
m.out5 <- matchit(treat ~ age + educ + race + married +
  nodegree + re74 + re75,
  data = lalonde,
  method = "full",
  distance = "glm",
```



```
link = "probit")

summary(m.out5)
```

Call:

```
matchit(formula = treat ~ age + educ + race + married + nodegree +
        re74 + re75, data = lalonde, method = "full", distance = "glm",
        link = "probit")
```

Summary of Balance for All Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.5773	0.1817	1.8276	0.8777	0.3774
age	25.8162	28.0303	-0.3094	0.4400	0.0813
educ	10.3459	10.2354	0.0550	0.4959	0.0347
raceblack	0.8432	0.2028	1.7615	.	0.6404
racehispan	0.0595	0.1422	-0.3498	.	0.0827
racewhite	0.0973	0.6550	-1.8819	.	0.5577
married	0.1892	0.5128	-0.8263	.	0.3236
nodegree	0.7081	0.5967	0.2450	.	0.1114
re74	2095.5737	5619.2365	-0.7211	0.5181	0.2248
re75	1532.0553	2466.4844	-0.2903	0.9563	0.1342

	eCDF Max
distance	0.6413
age	0.1577
educ	0.1114
raceblack	0.6404
racehispan	0.0827
racewhite	0.5577
married	0.3236
nodegree	0.1114
re74	0.4470
re75	0.2876

Summary of Balance for Matched Data:

	Means Treated	Means Control	Std. Mean Diff.	Var. Ratio	eCDF Mean
distance	0.5773	0.5764	0.0045	0.9949	0.0043
age	25.8162	25.5347	0.0393	0.4790	0.0787
educ	10.3459	10.5381	-0.0956	0.6192	0.0253
raceblack	0.8432	0.8389	0.0119	.	0.0043
racehispan	0.0595	0.0492	0.0435	.	0.0103
racewhite	0.0973	0.1119	-0.0493	.	0.0146
married	0.1892	0.1633	0.0660	.	0.0259
nodegree	0.7081	0.6577	0.1110	.	0.0504
re74	2095.5737	2100.2150	-0.0009	1.3467	0.0314
re75	1532.0553	1561.4420	-0.0091	1.5906	0.0536

	eCDF Max	Std. Pair Dist.
distance	0.6413	0.0043
age	0.1577	0.0787
educ	0.1114	0.0253
raceblack	0.6404	0.0043
racehispan	0.0827	0.0103
racewhite	0.5577	0.0146
married	0.3236	0.0259
nodegree	0.1114	0.0504
re74	0.4470	0.0314
re75	0.2876	0.0536

13. Matching and Weighting Part 1: matching


distance	0.0486	0.0198
age	0.2742	1.2843
educ	0.0730	1.2179
raceblack	0.0043	0.0162
racehispan	0.0103	0.4412
racewhite	0.0146	0.3454
married	0.0259	0.4473
nodegree	0.0504	0.9872
re74	0.1881	0.8387
re75	0.1984	0.8240

Sample Sizes:


	Control	Treated
All	429.	185
Matched (ESS)	50.76	185
Matched	429.	185
Unmatched	0.	0
Discarded	0.	0

```
plot(m.out5, type = "jitter", interactive = FALSE)
```

Distribution of Propensity Scores



```
plot(summary(m.out5))
```



13.6. Estimation

After we get a matched data that we think is good enough, we can estimate the treatment effect.

```
m.data <- match_data(m.out5)
```

```
head(m.data)
```

	treat	age	educ	race	married	nodegree	re74	re75	re78	distance
NSW1	1	37	11	black	1	1	0	0	9930.0460	0.6356769
NSW2	1	22	9	hispan	0	1	0	0	3595.8940	0.2298151
NSW3	1	30	12	black	0	0	0	0	24909.4500	0.6813558
NSW4	1	27	11	black	0	1	0	0	7506.1460	0.7690590
NSW5	1	33	8	black	0	1	0	0	289.7899	0.6954138
NSW6	1	22	9	black	0	1	0	0	4056.4940	0.6943658

	weights	subclass
NSW1	1	1
NSW2	1	55
NSW3	1	63
NSW4	1	70
NSW5	1	79
NSW6	1	86

What “match_data” does is to grab the matched data with a few additional variables such as “weights”, “subclass”, and “distance”. Then we can use it in a simple linear model with weights.

13. Matching and Weighting Part 1: matching

```
library("marginaleffects")

fit <- lm(re78 ~ treat * (age + educ + race + married +
                        nodegree + re74 + re75),
        data = m.data,
        weights = weights)

avg_comparisons(fit,
                variables = "treat",
                vcov = ~subclass,
                newdata = subset(m.data, treat == 1))
```

Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
1977	704	2.81	0.00501	7.6	596	3357

Term: treat
Type: response
Comparison: 1 - 0

Note we allow interaction of treatment and covariates, but we do not need to demean the covariates here, since we are relying on “marginaleffects” and its “avg_comparisons” function to compute the average treatment effect. The “newdata” argument is used to specify the effect, in this case ATT. Note the weights are generated this way: since it’s ATT, all treated units are weighted by 1, and control units are weighted by $p/(1 - p)$, where p is proportion of treated units in the subclass. Anyway, for full matching, we get the comparison of potential outcomes for the treated, with weights, and clustered standard errors on subclasses, because that the level the comparison is conducted.

Systematic treatment: [R](#) · [Julia](#).

14. Matching and Weighting - Part 2: weighting

This is based on R's MatchIt and WeightIt packages and their documentations. All credits go to Noah Greifer and coauthors. I write to summarize what I have learned.

14.1. Assumptions

Matching or weighting is usually used in observational studies. We have the following assumptions before using matching or weighting to make a causal claim:

1. SUTVA
2. ignorability (unconfoundedness, or no unobserved confounders)
3. overlap (common support, or positivity)

With these assumptions, we can try to identify the treatment effect in an observational study.

14.2. Problem

The second assumption is saying controlling for X , the treatment assignment D is independent of the potential outcomes $Y(0)$ and $Y(1)$, i.e., we have a pseudo-randomized experiment.

In reality, it's rare we have a balanced distribution of X for treated and control. That's when we need to use matching or weighting to balance the distribution of X between treatment and control groups.

14.3. weighting methods

The idea of weighting is to use weights to create a pseudo population which has a balanced distribution of X for treated and control groups. The most popular weighting method is inverse probability of treatment weighting (IPTW). IPTW uses the propensity score to create weights.

There are a few other weighting methods that are gaining popularity.

14.3.1. Inverse probability of treatment weighting (IPTW)

IPTW is the most common weighting method. It uses the inverse of the propensity score as weights. For the **ATE**, weight each treated unit by $1/p$ and each control unit by $1/(1 - p)$, where p is the propensity score. For the **ATT** — which is what the examples below request via `estimand = "ATT"` — the treated units all get weight 1 and the controls get $p/(1 - p)$; you can see this in the weight summaries, where the treated range is exactly 1 to 1. Either way the point is the same: create a pseudo population in which the distribution of X is balanced between the groups.

14.3.2. Covariate balancing propensity score (CBPS)

The same idea, but with a different way to estimate the propensity score. CBPS is a form of logistic regression where balance constraints are incorporated to a generalized method of moments estimation of the model coefficients.

14.3.3. Entropy balancing

Entropy balancing involves the specification of an optimization problem, the solution to which is then used to compute the weights. The constraints of the primal optimization problem correspond to covariate balance on the means (for binary and multi-category treatments) or treatment-covariate covariances (for continuous treatments), positivity of the weights, and that the weights sum to a certain value. Entropy balancing is doubly robust (for the ATT) in the sense that it is consistent either when the true propensity score model is a logistic regression of the treatment on the covariates or when the true outcome model for the control units is a linear regression of the outcome on the covariates, and it attains a semi-parametric efficiency bound when both are true. Entropy balancing will always yield exact mean balance on the included terms (not necessarily on the distribution of the covariates).

14.3.4. Energy balancing

Energy balancing is a method of estimating weights using optimization without a propensity score.

The primary benefit of energy balancing is that all features of the covariate distribution are balanced, not just means, as with other optimization-based methods like entropy balancing.

14.3.5. SuperLearner

SuperLearner works by fitting several machine learning models to the treatment and covariates and then taking a weighted combination of the generated predicted values to use as the propensity scores, which are then used to construct weights.

14.4. Example

Here is an example with Lalonde data, we are interested in the effect of treatment on “re78” (real earnings 1978).

14.4.1. IPW

```
library("WeightIt")
library(cobalt)
data(lalonde)
w.out1 <- weightit(treat ~ age + educ + race + married + nodegree + re74 + re75,
                  data = lalonde, estimand = "ATT", method = "glm")
w.out1 #print the output
```

A weightit object

- method: "glm" (propensity score weighting with GLM)
- number of obs.: 614
- sampling weights: none
- treatment: 2-category
- estimand: ATT (focal: 1)
- covariates: age, educ, race, married, nodegree, re74, re75

```
summary(w.out1)
```

Summary of weights

- Weight ranges:

	Min	Max
treated 1.		1.
control 0.009	-----	3.743

- Units with the 5 most extreme weights by group:

	5	4	3	2	1
treated	1	1	1	1	1
	597	573	381	411	303
control	3.03	3.059	3.24	3.523	3.743

- Weight statistics:

	Coef of Var	MAD	Entropy	# Zeros
treated	0.	0.	0.	0
control	1.818	1.289	1.098	0

14. Matching and Weighting - Part 2: weighting

- Effective Sample Sizes:

	Control	Treated
Unweighted	429.	185
Weighted	99.82	185

```
library(cobalt)
bal.tab(w.out1, stats = c("m", "v"), thresholds = c(m = .05))
```

Balance Measures

	Type	Diff.Adj	M.Threshold	V.Ratio.Adj
prop.score	Distance	-0.0205	Balanced, <0.05	1.0324
age	Contin.	0.1188	Not Balanced, >0.05	0.4578
educ	Contin.	-0.0284	Balanced, <0.05	0.6636
race_black	Binary	-0.0022	Balanced, <0.05	.
race_hispan	Binary	0.0002	Balanced, <0.05	.
race_white	Binary	0.0021	Balanced, <0.05	.
married	Binary	0.0186	Balanced, <0.05	.
nodegree	Binary	0.0184	Balanced, <0.05	.
re74	Contin.	-0.0021	Balanced, <0.05	1.3206
re75	Contin.	0.0110	Balanced, <0.05	1.3938

Balance tally for mean differences

	count
Balanced, <0.05	9
Not Balanced, >0.05	1

Variable with the greatest mean difference

Variable	Diff.Adj	M.Threshold
age	0.1188	Not Balanced, >0.05

Effective sample sizes

	Control	Treated
Unadjusted	429.	185
Adjusted	99.82	185

14.4.2. entropy balancing

```
w.out2 <- weightit(treat ~ age + educ + race + married + nodegree + re74 + re75,
                  data = lalonde, estimand = "ATT", method = "ebal")
summary(w.out2)
```

Summary of weights

- Weight ranges:

	Min	Max
treated	1.	1.
control	0.008	4.062

- Units with the 5 most extreme weights by group:

	5	4	3	2	1
treated	1	1	1	1	1
	608	381	597	303	411
control	3.073	3.235	3.45	3.897	4.062

- Weight statistics:

	Coef of Var	MAD	Entropy	# Zeros
treated	0.	0.	0.	0
control	1.834	1.287	1.101	0

- Effective Sample Sizes:

	Control	Treated
Unweighted	429.	185
Weighted	98.46	185

```
bal.tab(w.out2, stats = c("m", "v"), thresholds = c(m = .05))
```

Balance Measures

	Type	Diff.Adj	M.Threshold	V.Ratio.Adj
age	Contin.	-0	Balanced, <0.05	0.4096
educ	Contin.	-0	Balanced, <0.05	0.6635
race_black	Binary	0	Balanced, <0.05	.
race_hispan	Binary	0	Balanced, <0.05	.
race_white	Binary	-0	Balanced, <0.05	.
married	Binary	-0	Balanced, <0.05	.
nodegree	Binary	0	Balanced, <0.05	.
re74	Contin.	-0	Balanced, <0.05	1.3265
re75	Contin.	-0	Balanced, <0.05	1.3351

Balance tally for mean differences

	count
Balanced, <0.05	9
Not Balanced, >0.05	0

Variable with the greatest mean difference

Variable	Diff.Adj	M.Threshold
re74	-0	Balanced, <0.05

14. Matching and Weighting - Part 2: weighting

	Control	Treated
Unadjusted	429.	185
Adjusted	98.46	185

Suppose we are satisfied with the balance, we can use the weights to estimate the treatment effect. We can use the “marginaleffects” package to do that.

```
# Fit outcome model
fit <- lm_weightit(re78 ~ treat * (age + educ + race + married +
                                nodegree + re74 + re75),
                  data = lalonde, weightit = w.out2)
# G-computation for the treatment effect
library("marginaleffects")
avg_comparisons(fit, variables = "treat",
                 newdata = subset(lalonde, treat == 1))
```

Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
1273	770	1.65	0.0983	3.3	-236	2783

Term: treat
Type: probs
Comparison: 1 - 0

14.4.3. energy balancing

```
w.out3 <- tryCatch(
  weightit(treat ~ age + educ + race + married +
            nodegree + re74 + re75, data = lalonde,
            method = "energy", estimand = "ATE"),
  error = function(e) { message("Energy balancing failed: ", e$message); NULL }
)
if (!is.null(w.out3)) summary(w.out3)
```

Summary of weights

- Weight ranges:

	Min	Max
treated	0 -----	14.041
control	0 -----	11.065

- Units with the 5 most extreme weights by group:

	10	179	125	137	124
treated	5.984	6.23	6.758	13.202	14.041
	537	553	573	420	608
control	5.499	6.404	6.599	8.274	11.065

- Weight statistics:

	Coef of Var	MAD	Entropy	# Zeros
treated	1.821	1.054	0.92	0
control	1.264	0.866	0.579	0

- Effective Sample Sizes:

	Control	Treated
Unweighted	429.	185.
Weighted	165.42	43.04

```
if (!is.null(w.out3)) {
  fit <- lm_weightit(re78 ~ treat * (age + educ + race + married +
                                   nodegree + re74 + re75),
                    data = lalonde, weightit = w.out3)
  # G-computation for the treatment effect
  avg_comparisons(fit, variables = "treat")
}
```

Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
-73	1085	-0.0672	0.946	0.1	-2200	2054

Term: treat

Type: probs

Comparison: 1 - 0

Two things are worth noticing in that output rather than passing over it. First, the estimand changed: `w.out3` requests the **ATE**, while the entropy-balancing fit above targeted the **ATT**, so the two numbers (-73 against 1273) are not competing estimates of the same quantity. Second, look at the weight summary. Entropy balancing gives every treated unit a weight of exactly 1, as the ATT requires. Energy balancing spreads the treated weights from 0 to 14.0 — one treated unit is effectively discarded and another counts as fourteen — and the resulting standard error (1085) is much larger than entropy balancing's (770). Balancing the whole covariate *distribution* rather than just its means is a genuine advantage, but it is bought with weight variability, and that shows up directly in the precision of the estimate.

14.5. comparing matching and weighting

Matching and weighting are popular in different fields. We used to say matching involves discarding data therefore changing the estimand, but now with full matching and other methods, matching can be done without losing data. We used to believe weighting can be unstable, as weights can be highly variable, but with CBPS, entropy balancing and energy balancing, weighting can perform much better.

So, in practice, we can try different methods, either matching or weighting, and see which one gives us better balance. Then use “marginaleffects” to get the treatment effect.

Systematic treatment: [R](#) · [Julia](#).

15. Sensitivity analysis of OLS with sensemakr

15.1. OVB

Suppose we have a true model,

$$Y = \delta D + X\beta + \gamma Z + \epsilon \quad (15.1)$$

However, we only observe X , not Z . We estimate the model

$$Y = \delta D + X\beta + u \quad (15.2)$$

In this case, we have omitted variable bias (OVB).

Cinelli and Hazlett (2020) shows that

$$|\widehat{\text{bias}}| = \sqrt{\frac{R_{Y \sim Z|D,X}^2 R_{D \sim Z|X}^2}{1 - R_{D \sim Z|X}^2}} \left(\frac{\hat{\sigma}_{y.dx}}{\hat{\sigma}_{d.x}} \right) \quad (15.3)$$

We can see that when we have high R^2 between Y and Z , bias is high. When we have high R^2 between D and Z , bias is high. In other words, the omitted variable needs to be somewhat highly correlated with both Y and D to have a high bias.

Cinelli and Hazlett paper has many stats, but we focus on three things we can report:

15.2. Partial R^2 of treatment on outcome

The original purpose of this stat is to see how much variation treatment explains outcome after partialling out the covariates.

Also, this can be used for an extreme situation. Suppose the unobserved confounder Z explains all residual variance of the outcome, that is, $R_{Y \sim Z|D,X}^2 = 1$. What kind of situation we would have 0 effect?

It is shown if $R_{Y \sim Z|D,X}^2 = 1$, then $R_{D \sim Z|X}^2 \geq R_{Y \sim D|X}^2$ to make the effect go to 0. Roughly speaking, the partial R^2 between the omitted variable and the treatment needs to exceed the partial R^2 between the treatment and the outcome to make the treatment effect go to 0.

15.3. Robust value

The robustness value $RV_{q,\alpha}$ quantifies the minimal strength of association that the confounder needs to have, both with the treatment and with the outcome, so that a confidence interval of level $1 - \alpha$ includes a change of $q\%$ of the current estimated value. q is often 1. That is, for the effect to go to 0, that robust value is the minimum value of the correlation between the omitted variable and the treatment and the correlation between the omitted variable and the outcome. Note this is about both the correlation between Z and D and the correlation between Z and Y . They need to be both high to make the effect go to 0.

15.4. Bounds using observed covariate(s)

For the first two values, they are absolute values. You don't know whether those are large or small, without context. They cannot be compared across studies.

Suppose there is a covariate x that you know from your domain knowledge that it is an important factor in explaining treatment or/and outcome. Then you can compare how bad the confounder needs to be compared to this covariate. This will give us some sense how much a confounder needs to be compared to a strong covariate, in driving the effect to 0.

$$k_D := \frac{R_{D \sim Z | X_{-j}}^2}{R_{D \sim X_j | X_{-j}}^2} \quad (15.4)$$

where X_{-j} represents all covariates other than the j th we pick as the benchmark covariate.

$$k_Y := \frac{R_{Y \sim Z | D, X_{-j}}^2}{R_{Y \sim X_j | D, X_{-j}}^2} \quad (15.5)$$

15.5. Example

Here is an example from sensemakr documentation.

```
# loads package
library(sensemakr)

# loads dataset
data("darfur")

# runs regression model
model <- lm(peacefactor ~ directlyharmed + age + farmer_dar + herder_dar +
            pastvoted + hhsizedarfur + female + village, data = darfur)

# runs sensemakr for sensitivity analysis
sensitivity <- sensemakr(model = model,
```

```

        treatment = "directlyharmed",
        benchmark_covariates = "female",
        kd = 1:3)

# short description of results
sensitivity

```

Sensitivity Analysis to Unobserved Confounding

Model Formula: peacefactor ~ directlyharmed + age + farmer_dar + herder_dar +
pastvoted + hhsize_darfur + female + village

Null hypothesis: $q = 1$ and reduce = TRUE

Unadjusted Estimates of 'directlyharmed':

Coef. estimate: 0.09732
Standard Error: 0.02326
t-value: 4.18445

Sensitivity Statistics:

Partial R2 of treatment with outcome: 0.02187
Robustness Value, $q = 1$: 0.13878
Robustness Value, $q = 1$ alpha = 0.05 : 0.07626

For more information, check summary.

```

# long description of results
summary(sensitivity)

```

Sensitivity Analysis to Unobserved Confounding

Model Formula: peacefactor ~ directlyharmed + age + farmer_dar + herder_dar +
pastvoted + hhsize_darfur + female + village

Null hypothesis: $q = 1$ and reduce = TRUE

-- This means we are considering biases that reduce the absolute value of the current estimate
-- The null hypothesis deemed problematic is $H_0:\tau = 0$

Unadjusted Estimates of 'directlyharmed':

Coef. estimate: 0.0973
Standard Error: 0.0233
t-value ($H_0:\tau = 0$): 4.1844

Sensitivity Statistics:

Partial R2 of treatment with outcome: 0.0219
Robustness Value, $q = 1$: 0.1388

15. Sensitivity analysis of OLS with sensemakr

Robustness Value, q = 1, alpha = 0.05: 0.0763

Verbal interpretation of sensitivity statistics:

-- Partial R2 of the treatment with the outcome: an extreme confounder (orthogonal to the covariates)

-- Robustness Value, q = 1: unobserved confounders (orthogonal to the covariates) that explain

-- Robustness Value, q = 1, alpha = 0.05: unobserved confounders (orthogonal to the covariates)

Bounds on omitted variable bias:

--The table below shows the maximum strength of unobserved confounders with association with the

Bound	Label	R2dz.x	R2yz.dx	Treatment	Adjusted Estimate	Adjusted Se
1x	female	0.0092	0.1246	directlyharmed	0.0752	0.0219
2x	female	0.0183	0.2493	directlyharmed	0.0529	0.0204
3x	female	0.0275	0.3741	directlyharmed	0.0304	0.0187
Adjusted T	Adjusted	Lower CI	Adjusted	Upper CI		
3.4389		0.0323		0.1182		
2.6002		0.0130		0.0929		
1.6281		-0.0063		0.0670		

```
# minimal reporting
ovb_minimal_reporting(sensitivity, format = "html")
```

Outcome: peacefactor

Treatment

Est.

S.E.

t-value

$R^2_{Y \sim D | \mathbf{X}}$

$RV_{q=1}$

$RV_{q=1, \alpha=0.05}$

directlyharmed

0.097

0.023

4.184

2.2%

13.9%

7.6%

Note: $df = 783$; Bound (1x female): $R_{Y \sim Z|X,D}^2 = 12.5\%$, $R_{D \sim Z|X}^2 = 0.9\%$

Let's see how to interpret this.

$RV_{q=1} = 13.9\%$. This means that unobserved confounders that explain 13.9% of the residual variance both of the treatment and of the outcome are sufficiently strong to explain away all the observed effect. On the other hand, unobserved confounders that do not explain at least 13.9% of the residual variance both of the treatment and of the outcome are not sufficiently strong to do so.


The robustness value for testing the null hypothesis that the coefficient of “directlyharmed” is zero ($RV_{q=1, \alpha=0.05}$) is 7.6%. This means that unobserved confounders that explain 7.6% of the residual variance both of the treatment and of the outcome are sufficiently strong to bring the lower bound of the confidence interval to zero (at the chosen significance level of 5%).

The partial R^2 of `directlyharmed` with `peacefactor` is 2.2%. This means that, in an extreme scenario in which we assume that unobserved confounders explain all of the left-out variance of the outcome, these unobserved confounders would need to explain at least 2.2% of the residual variance of the treatment to fully explain away the observed effect.

The above numbers are for the absolute strength the confounder needs to be to explain away the treatment effect. The authors believe in this context, “female” is an important factor. Suppose the confounder is as important as “female”, what the partial R^2 would be? It is shown that $R_{Y \sim Z|X,D}^2 = 12.5\%$ and $R_{D \sim Z|X}^2 = 0.9\%$, so a confounder as strong as “female” is not strong enough to explain away the treatment effect.

It is worth being careful about *why*, because comparing each number to $RV_{q=1} = 13.9\%$ one at a time invites the wrong reading. The RV is defined for a confounder of *equal* strength on both sides, so it is not a threshold that each partial R^2 must separately clear. Here the outcome-side association (12.5%) is in fact close to the RV , which on its own might look like a near miss — but the treatment-side association is only 0.9%, more than an order of magnitude short. It is the treatment side that does the work: “female” predicts the outcome nearly well enough to matter, and barely predicts the treatment at all. The contour plot below is the honest check, since it evaluates the pair jointly rather than either coordinate alone.

```
# plot bias contour of point estimate
plot(sensitivity)
```



This plot shows that even 3 times as strong as “female”, the confounder is not strong enough to explain away the treatment effect.

15.6. sensitivity analysis for IV

Even with an IV model, we need sensitivity analysis. The worry would be there could be an unobserved confounder that affects both the instrument and the outcome.

Cinelli and Hazlett have a package “iv.sensemakr” for this. The usage is similar to “sensemakr”.

```
# loads package
library(iv.sensemakr)

# loads dataset
data("card")

# prepares data
y <- card$lwage # outcome
d <- card$educ  # treatment
z <- card$nearc4 # instrument
x <- model.matrix( ~ exper + expersq + black + south + smsa + reg661 + reg662 +
                    reg663 + reg664 + reg665 + reg666 + reg667 + reg668 + smsa66,
                    data = card) # covariates

# fits IV model
card.fit <- iv_fit(y,d,z,x)

# see results
card.fit
```

Instrumental Variable Estimation
(Anderson-Rubin Approach)

=====

IV Estimates:

Coef. Estimate: 0.132

t-value: 2.33

p-value: 0.02

Conf. Interval: [0.025, 0.285]

Note: $H_0 = 0$, $\alpha = 0.05$, $df = 2994$.

=====

See summary for first stage and reduced form.

```
# runs sensitivity analysis
card.sens <- sensemakr(card.fit, benchmark_covariates = c("black", "smsa"))

# see results
card.sens
```

Sensitivity Analysis for Instrumental Variables
(Anderson-Rubin Approach)

=====

IV Estimates:

Coef. Estimate: 0.132

t-value: 2.33

p-value: 0.02

Conf. Interval: [0.025, 0.285]

Sensitivity Statistics:

Extreme Robustness Value: 0.000523

Robustness Value: 0.00667

Bounds on Omitted Variable Bias:

Bound	Label	R2zw.x	R2y0w.zx	Lower CI	Upper CI	Crit. Thr.
1x	black	0.00221	0.0750	-0.0212	0.402	2.59
1x	smsa	0.00639	0.0202	-0.0192	0.396	2.57

Note: $H_0 = 0$, $q \geq 1$, $\alpha = 0.05$, $df = 2994$.

=====

See summary for first stage and reduced form.

```
summary(card.sens)
```

Sensitivity Analysis for Instrumental Variables

15. Sensitivity analysis of OLS with sensemakr

(Anderson-Rubin Approach)

IV Estimates:

Coef. Estimate: 0.132
t-value: 2.33
p-value: 0.02
Conf. Interval: [0.025, 0.285]

Sensitivity Statistics:

Extreme Robustness Value: 0.000523
Robustness Value: 0.00667

Bounds on Omitted Variable Bias:

Bound Label	R2zw.x	r2y0w.zx	Lower CI	Upper CI	Crit. Thr.
1x black	0.00221	0.0750	-0.0212	0.402	2.59
1x smsa	0.00639	0.0202	-0.0192	0.396	2.57

Note: $H_0 = 0$, $q \geq 1$, $\alpha = 0.05$, $df = 2994$.

FS Estimates:

Coef. Estimate: 0.32
Standard Error: 0.0879
t-value: 3.64
p-value: 0.000276
Conf. Interval: [0.148, 0.492]

Sensitivity Statistics:

Extreme Robustness Value: 0.00313
Robustness Value: 0.0302

Bounds on Omitted Variable Bias:

Bound Label	R2zw.x	R2dw.zx	Lower CI	Upper CI	Crit. Thr.
1x black	0.00221	0.03342	0.109	0.531	2.40
1x smsa	0.00639	0.00498	0.120	0.520	2.27

Note: $H_0 = 0$, $q = 1$, $\alpha = 0.05$, $df = 2994$.

RF Estimates:

Coef. Estimate: 0.0421
Standard Error: 0.0181
t-value: 2.33
p-value: 0.02
Conf. Interval: [0.007, 0.078]

Sensitivity Statistics:

Extreme Robustness Value: 0.000523
Robustness Value: 0.00667


Bounds on Omitted Variable Bias:

Bound	Label	R2zw.x	R2yw.zx	Lower CI	Upper CI	Crit. Thr.
1x	black	0.00221	0.0657	-0.00418	0.0883	2.56
1x	smsa	0.00639	0.0197	-0.00429	0.0884	2.56

Note: $H_0 = 0$, $q = 1$, $\alpha = 0.05$, $df = 2994$.

The basic idea is to run an IV regression, then apply `sensmakr` to that model. In this case, the authors use “black” and “smsa” as benchmark covariates. The authors believe these are important factors. A quick look at the result we can see that a confounder as strong as “black” or “smsa” is not strong enough to explain away the treatment effect. We can also plot a contour plot to see how strong the confounder needs to be to explain away the treatment effect.

```
plot(card.sens, parm = "rf", sensitivity.of = "t-value", lim = 0.08, kz = 1, ky = 1)
```



Unfortunately the package currently only allows for one instrument. Should not be too hard to modify it to be more general.

Systematic treatment: [R](#) · [Julia](#).

16. Gender Wage Gap Analysis Using US Current Population Survey Data

Traditionally Oaxaca-Blinder decomposition is used to decompose the wage gap into explained and unexplained parts. The explained part is due to differences in observable characteristics, such as education, experience, occupation, etc. The unexplained part is due to discrimination or other unobserved factors.

It has a problem that when decomposing, using female as reference group or male as reference group will give different results.

Słoczyński (2015) and Słoczyński (2013) have shown that a regression adjustment type of analysis is equivalent to Oaxaca-Blinder decomposition in a nonstandard way.

16.1. Oaxaca-Blinder decomposition

Assume the regression model is different for male and female:

$$y_i = X_i\beta_a + \epsilon_a \quad (16.1)$$

where a is gender (female is 1), X 's are person characteristics.

$$E[y_i|a_i = 1] - E[y_i|a_i = 0] = E[X_i|a_i = 1](\beta_1 - \beta_0) + (E[X_i|a_i = 1] - E[X_i|a_i = 0])\beta_0 \quad (16.2)$$

or

$$E[y_i|a_i = 1] - E[y_i|a_i = 0] = E[X_i|a_i = 0](\beta_1 - \beta_0) + (E[X_i|a_i = 1] - E[X_i|a_i = 0])\beta_1 \quad (16.3)$$

These two equations are different forms of the Oaxaca-Blinder decomposition. The first element is the “unexplained component”, which is the difference in the coefficients, and the second element is the “explained component”, which is the difference in the means of the covariates.

These two equations can give different results for the two components, one has a baseline of male, and other one female.

We can further shows that

$$\begin{aligned}
 E[y_i|a_i = 1] - E[y_i|a_i = 0] &= E[X_i|a_i = 1](\beta_1 - \beta_0) + (E[X_i|a_i = 1] - E[X_i|a_i = 0])\beta_0 \\
 &= E[y_i(1) - y_i(0)|a_i = 1] + (E[y_i(0)|a_i = 1] - E[y_i(0)|a_i = 0]) \quad (16.4) \\
 &= \tau_{PATT} + (E[y_i(0)|a_i = 1] - E[y_i(0)|a_i = 0])
 \end{aligned}$$

That is, the observed difference is the sum of τ_{PATT} and selection bias. It is also known in the OB decomposition literature as the unexplained component and the explained component. The former as “discrimination”.

We can simply use Regression Adjustment for these quantities. The easiest is an OLS of y_i on a_i , the demeaned X_i , and their interaction; the coefficient on a_i is then the effect evaluated at the mean of the covariates used for the demeaning. Demean at the **overall** mean and that coefficient is τ_{PAGE} (the ATE), which is what the example below does; demean at the mean among **women** instead and it is τ_{PATT} . The choice of centring is what selects the estimand, so it is worth being explicit about which one you have.

He also advocates the “population average gender effect” (PAGE):

$$\tau_{PAGE} = E[\tau(X_i)] \quad (16.5)$$

which is basically ATE, or conditional average effect (CATE), then average over the X ’s.

In this version, the two forms of OB decomposition can be τ_{PAGM} and τ_{PAGW} , for men and women.

The nice thing about this estimator τ_{PAGE} is that first, it is consistent with the treatment effect (causal inference) literature, and it is invariant to the choice of reference group.

The second advantage is that we also do not need to assume homogeneity of treatment effect, which is a strong assumption in the Oaxaca-Blinder decomposition.

The third advantage is that we can use nonparametric methods to estimate the treatment effect. Traditional OB decomposition is based on linear regression, which assumes a linear relationship between the covariates and the outcome.

16.2. Example

I am using the data from <https://github.com/d2cml-ai/CausalAI-Course/tree/main/data>.

“The data set we consider is from the 2015 March Supplement of the U.S. Current Population Survey. We select white non-hispanic individuals, aged 25 to 64 years, and working more than 35 hours per week for at least 50 weeks of the year. We exclude self-employed workers; individuals living in group quarters; individuals in the military, agricultural or private household sectors; individuals with inconsistent reports on earnings and employment status; individuals with allocated or missing information in any of the variables used in the analysis; and individuals with hourly wage below 3.”

I am using this data to study gender wage gap (GWP).

We can model GWP using the potential outcome framework, with some care about what plays the role of “treatment.” Gender itself is an immutable characteristic, not something that can be manipulated or randomized — the standard “no causation without manipulation” principle (Holland, 1986) says we should not treat gender itself as a causal treatment. What we really have in mind, and what the “treatment” indicator below is best understood as a proxy for, is something like *being perceived/treated as female by employers* (e.g., employer discrimination), which is at least conceptually manipulable even though we cannot experimentally assign it here. With that caveat in mind, we let $D = 1$ if female, $D = 0$ if male, and outcome Y is logged wage.

The ATE would be

$$\delta_{ATE} = E[Y(1) - Y(0)] \quad (16.6)$$

where $Y(1)$ is the potential outcome if treated, and $Y(0)$ is the potential outcome if not treated.

16.2.1. Regression adjustment

Since this is to estimate ATE, we can use regression adjustment to estimate the ATE. First let’s use de-meaned X ’s. Then the coefficient on gender is the ATE.

We use These variables as X ’s:

“shs”, “hsg”, “scl”, “clg”, “ad”, “ne”, “mw”, “so”, “we”, “exp1”, “occ”, “occ2”, “ind”, and “ind2”.

They are:

“Less then High School”, “High School Graduate”, “Some College”, “College Graduate”, “Advanced Degree”, “Northeast”, “Midwest”, “South”, “West”, “Experience”, “occupation”, “occupation 2”, “industry”, and “industry 2”.

I decide to use “occ2” and “ind2” as they are broader categories of occupation and industry. We don’t have enough data for finer categories “occ” and “ind”.

“exp1” seems to be the only variable that is continuous, let’s demean it first. Let’s also input “occ2” and “ind2” as factors.

Let’s include demeaned “exp1”, “sohs”, “hsg”, “socl”, “clg”, “mw”, “sout”, and “we” first. I leave out “occ2” and “ind2” since there are too many levels to demean.

```
library(tidyverse)
library(broom)

data <- read_csv("wage2015_subsample_inference.csv") |>
  rename(socl = scl, sohs = shs, sout = so) |>
  mutate(exp_dm = exp1 - mean(exp1, na.rm = TRUE), female=sex, occ2 = factor(occ2), ind2 = fac
  mutate(sohs_dm = sohs - mean(sohs, na.rm = TRUE),
         hsg_dm = hsg - mean(hsg, na.rm = TRUE),
         socl_dm = socl - mean(socl, na.rm = TRUE),
         clg_dm = clg - mean(clg, na.rm = TRUE),
```

16. Gender Wage Gap Analysis Using US Current Population Survey Data

```
mw_dm = mw - mean(mw, na.rm = TRUE),
sout_dm = sout - mean(sout, na.rm = TRUE),
we_dm = we - mean(we, na.rm = TRUE))
```

```
glimpse(data)
```

Rows: 5,150

Columns: 30

```
$ rownames <dbl> 10, 12, 15, 18, 19, 30, 43, 44, 47, 71, 73, 77, 84, 89, 96, 1~
$ wage <dbl> 9.615385, 48.076923, 11.057692, 13.942308, 28.846154, 11.7307~
$ lwage <dbl> 2.263364, 3.872802, 2.403126, 2.634928, 3.361977, 2.462215, 2~
$ sex <dbl> 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1~
$ sohs <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ hsg <dbl> 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 0~
$ socl <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1~
$ clg <dbl> 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0~
$ ad <dbl> 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ mw <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ sout <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ we <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
$ ne <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1~
$ exp1 <dbl> 7.0, 31.0, 18.0, 25.0, 22.0, 1.0, 42.0, 37.0, 31.0, 4.0, 7.0, ~
$ exp2 <dbl> 0.4900, 9.6100, 3.2400, 6.2500, 4.8400, 0.0100, 17.6400, 13.6~
$ exp3 <dbl> 0.343000, 29.791000, 5.832000, 15.625000, 10.648000, 0.001000~
$ exp4 <dbl> 0.24010000, 92.35210000, 10.49760000, 39.06250000, 23.4256000~
$ occ <dbl> 3600, 3050, 6260, 420, 2015, 1650, 5120, 5240, 4040, 3255, 40~
$ occ2 <fct> 11, 10, 19, 1, 6, 5, 17, 17, 13, 10, 13, 14, 11, 11, 1, 19, 1~
$ ind <dbl> 8370, 5070, 770, 6990, 9470, 7460, 7280, 5680, 8590, 8190, 82~
$ ind2 <fct> 18, 9, 4, 12, 22, 14, 14, 9, 19, 18, 18, 18, 18, 18, 17, 4, 4~
$ exp_dm <dbl> -6.760583, 17.239417, 4.239417, 11.239417, 8.239417, -12.7605~
$ female <dbl> 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1~
$ sohs_dm <dbl> -0.02330097, -0.02330097, -0.02330097, -0.02330097, -0.023300~
$ hsg_dm <dbl> -0.2438835, -0.2438835, 0.7561165, -0.2438835, -0.2438835, -0~
$ socl_dm <dbl> -0.2780583, -0.2780583, -0.2780583, -0.2780583, -0.2780583, --
$ clg_dm <dbl> 0.6823301, 0.6823301, -0.3176699, -0.3176699, 0.6823301, 0.68~
$ mw_dm <dbl> -0.2596117, -0.2596117, -0.2596117, -0.2596117, -0.2596117, --
$ sout_dm <dbl> -0.2965049, -0.2965049, -0.2965049, -0.2965049, -0.2965049, --
$ we_dm <dbl> -0.2161165, -0.2161165, -0.2161165, -0.2161165, -0.2161165, --
```

```
# construct matrices for estimation from the data
```

```
# 1. basic model
```

```
reg1 <- lm(lwage ~ female, data=data)
tidy(reg1)
```

```
# A tibble: 2 x 5
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	2.99	0.0107	280.	0
2	female	-0.0383	0.0160	-2.40	0.0165

```
reg2 <- lm(lwage ~ female*(exp_dm + sohs_dm + hsg_dm + socl_dm + clg_dm + mw_dm + sout_dm + we_dm)
tidy(reg2)
```

```
# A tibble: 18 x 5
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	3.02	0.00974	310.	0
2	female	-0.115	0.0147	-7.80	7.39e-15
3	exp_dm	0.00801	0.000957	8.37	7.31e-17
4	sohs_dm	-0.811	0.0620	-13.1	1.63e-38
5	hsg_dm	-0.706	0.0350	-20.1	6.09e-87
6	socl_dm	-0.553	0.0351	-15.7	1.45e-54
7	clg_dm	-0.256	0.0345	-7.41	1.42e-13
8	mw_dm	0.0337	0.0280	1.20	2.29e-1
9	sout_dm	0.0148	0.0271	0.548	5.84e-1
10	we_dm	0.0555	0.0290	1.91	5.59e-2
11	female:exp_dm	0.00212	0.00140	1.51	1.30e-1
12	female:sohs_dm	-0.0180	0.117	-0.154	8.78e-1
13	female:hsg_dm	0.0398	0.0507	0.786	4.32e-1
14	female:socl_dm	0.0466	0.0482	0.968	3.33e-1
15	female:clg_dm	0.0992	0.0468	2.12	3.41e-2
16	female:mw_dm	-0.124	0.0417	-2.98	2.93e-3
17	female:sout_dm	-0.0530	0.0403	-1.31	1.89e-1
18	female:we_dm	-0.0655	0.0435	-1.51	1.32e-1

We see the raw gap is 3.8%. Adjusting for some covariates we have 11% gap.

That direction deserves comment, because it is the opposite of what most readers expect. Adding controls usually shrinks a gap; here it nearly triples it. The reason is composition:

```
library(tidyverse) # re-attached: the chunk that loads it is cached

data |>
  group_by(female) |>
  summarise(n = n(), mean_lwage = mean(lwage), exp1 = mean(exp1),
            `<HS` = mean(sohs), HS = mean(hsg), `some coll` = mean(socl),
            college = mean(clg), advanced = mean(ad)) |>
  mutate(across(where(is.numeric), ~round(.x, 3)))
```

```
# A tibble: 2 x 9
```

female	n	mean_lwage	exp1	`<HS`	HS	`some coll`	college	advanced
--------	---	------------	------	-------	----	-------------	---------	----------

16. Gender Wage Gap Analysis Using US Current Population Survey Data

	<dbl>	<dbl>		<dbl>	<dbl>	<dbl>	<dbl>		<dbl>	<dbl>	<dbl>
1	0	2861		2.99	13.8	0.032	0.294		0.273	0.294	0.107
2	1	2289		2.95	13.7	0.013	0.181		0.284	0.347	0.175

Women in this extract are **more** educated than men — 17.5% hold advanced degrees against 10.7%, and 1.3% are below high school against 3.2% — while potential experience is essentially identical (13.7 against 13.8 years). Education is therefore working in women’s favour in the raw comparison. Holding it fixed removes that advantage and exposes a larger gap within education cells. Nothing subtle is happening; the raw and adjusted numbers simply describe different populations of comparison.

Keep this table in mind, because it comes back twice below: once when we ask what the adjusted number is an estimate *of*, and once when we ask whether conditioning on education might be doing harm rather than good.

Or, we can use “marginaleffects” instead of demeaning.

```
reg3 <- lm(lwage ~ female*(exp1 + sohs + hsg+ soc1 + clg + mw + sout + we ), data=data)

library(marginaleffects)

avg_comparisons(reg3,
                 variables = "female")
```

```
Estimate Std. Error    z Pr(>|z|)    S  2.5 %  97.5 %
-0.115    0.0147 -7.8   <0.001 47.2 -0.143 -0.0858
```

```
Term: female
Type: response
Comparison: 1 - 0
```

We get the same result as the demeaned regression.

If we want τ_{ATT} , then we can use the “newdata” argument to specify the treatment group,

```
avg_comparisons(reg3,
                 variables = "female",
                 newdata = subset(data, female == 1))
```

```
Estimate Std. Error    z Pr(>|z|)    S  2.5 %  97.5 %
-0.113    0.0148 -7.65  <0.001 45.5 -0.142 -0.0844
```

```
Term: female
Type: response
Comparison: 1 - 0
```

16.2.2. Occupation and industry: a bad control problem

Before adding occupation and industry, it is worth asking what kind of variables they are — and the honest answer implicates the covariates we have already been using.

Sex is assigned at birth. Essentially nothing in the covariate vector precedes it:

Covariate	Downstream of sex?
Education (<code>sohs...ad</code>)	Yes — schooling decisions respond to gender
Experience (<code>exp1</code>)	Partly — potential experience is age minus schooling, so it inherits education’s endogeneity; actual experience is fully downstream (interruptions, part-time spells)
Region (<code>mw</code> , <code>sout</code> , <code>we</code>)	Yes — current region reflects migration, and tied-mover patterns are gendered
Occupation, industry (<code>occ2</code> , <code>ind2</code>)	Yes, unambiguously

Genuinely pre-treatment variables would be birth cohort, region of birth, and parental characteristics — and even the last is shaky, since parental investment is known to respond to a child’s sex. So the valid adjustment set for sex is close to **empty**.

Two consequences follow, and the second is the one that matters.

First, there is no confounding to adjust away. The usual reason to control for X is to make $Y(d) \perp D \mid X$ hold. But sex is not confounded by education — education is *caused* by sex. Sex is about as close to randomised-by-nature as a treatment gets, so $Y(d) \perp D$ holds without conditioning on anything. Adjustment cannot improve identification here; it can only change the estimand.

Second, and consequently, every model in this chapter sits somewhere on a **ladder of controlled direct effects**, not on a scale from worse to better estimates of one quantity:

Model	Held fixed	Question answered	Estimate
<code>reg1</code>	nothing	Total effect of the gender signal	−3.8%
<code>reg3</code>	education, experience, region	Gap within education/experience cells	−11.3%
<code>reg4</code>	+ occupation, industry	Gap within occupation cells too	−7.5%

Read that way, movement between rungs is not error to be minimised — it is information about where in the causal chain the gap accumulates. But each rung names a different intervention and needs its own justification, and the argument below applies at *every* rung, not only the last. Occupation and industry are simply the clearest case, so we use them to make the point.

They are *bad controls* in the sense of Cinelli, Forney and Pearl (2022), for two separate reasons at the same time.

16. Gender Wage Gap Analysis Using US Current Population Survey Data

```
library(ggdag)
library(ggplot2)

g <- dagify(
  Wage ~ Female + Occupation + U,
  Occupation ~ Female + U,
  exposure = "Female",
  outcome = "Wage",
  coords = list(x = c(Female = 0, Occupation = 1.1, U = 1.1, Wage = 2.2),
                y = c(Female = 0, Occupation = -1, U = 1, Wage = 0))
)

# node_size has to stay small: ggdag stops each edge at the node boundary, so
# large nodes swallow the arrowheads -- and the arrowheads are the whole point.
ggdag(g, text = FALSE, node_size = 9) +
  geom_dag_text_repel(aes(label = name), colour = "black", size = 3.2,
                      box.padding = 0.9, seed = 42) +
  theme_dag() +
  expand_limits(x = c(-0.45, 2.65), y = c(-1.5, 1.5))
```




Figure 16.1.: Occupation is both a mediator of the gender–wage relationship and a collider on a path through unobserved wage determinants U .

Here U stands for unobserved wage determinants — drive, negotiation propensity, family constraints, expected career continuity — and **Occupation** covers occupation and industry together.

Occupation is a mediator. Part of the total effect of gender on wages runs **Female** → **Occupation** → **Wage**. Conditioning on occupation blocks that path, so the coefficient stops estimating a total

effect. On its own this would be tolerable, as long as we said so: we would be reporting a within-occupation gap, which is a perfectly coherent descriptive quantity.

Occupation is also a collider. It is a common effect of gender and of U . Conditioning on a common effect *induces* an association between its causes, so conditioning on occupation opens the non-causal path `Female -> Occupation <- U -> Wage`, which was closed before we touched anything (Elwert and Winship, 2014).

We do not have to take that on faith. `dagitty` will enumerate the paths and say which are open, with and without conditioning on occupation:

```
paths_table <- function(Z = NULL) {
  p <- dagitty::paths(g, from = "Female", to = "Wage", Z = Z)
  data.frame(path = p$paths, open = p$open)
}

paths_table()                # nothing conditioned on
```

	path	open
1	Female -> Occupation -> Wage	TRUE
2	Female -> Occupation <- U -> Wage	FALSE
3	Female -> Wage	TRUE

```
paths_table(Z = "Occupation") # conditioning on occupation
```

	path	open
1	Female -> Occupation -> Wage	FALSE
2	Female -> Occupation <- U -> Wage	TRUE
3	Female -> Wage	TRUE

The two tables are the whole argument. Conditioning on occupation flips `Female -> Occupation -> Wage` from open to closed — that is the mediation we knowingly gave up — and simultaneously flips `Female -> Occupation <- U -> Wage` from closed to open. That second flip is the damage. Blocking a mediator removes something we wanted; conditioning on a collider manufactures something we did not. After adjustment, gender and U are correlated *within* occupation cells even if they are independent in the population, and because U also drives wages, that induced correlation leaks into the gender coefficient.

Stated formally: conditioning on a mediator identifies a *controlled direct effect* only under **no unmeasured mediator–outcome confounding** (VanderWeele and Robinson, 2014). U is exactly such a confounder, and in a wage setting it is not a remote worry — it is more or less why this literature exists.

One more thing worth flagging before we run the model: this is not a problem a better estimator solves. The AIPW and SuperLearner fits later in this chapter relax functional form, and do nothing at all about a structurally uninterpretable estimand.

16.2.3. Selection is the real threat, and it is not confounding

It is worth separating two problems that both get called “bias,” because they have different fixes and only one of them is live here.

Under the exogeneity argument above, the raw difference in means *is* the total causal effect of sex on wages. That is a genuine causal quantity — the effect of being conceived female rather than male on wages measured decades later — identified without any adjustment. It is enormously composite, bundling biology, socialisation, schooling, sorting, and employer behaviour into one number, which is why it answers no policy question. But it is not confounded.

What does break it is the **sample**. This extract keeps only full-year, full-time, non-self-employed workers, and employment is a common effect of sex and unobserved wage determinants:

```
library(ggdag)
library(ggplot2)
library(dagitty)

g3 <- dagitty('dag {
  SexBirth -> Educ
  SexBirth -> Wage
  SexBirth -> Employed
  U -> Employed
  U -> Wage
  Educ -> Wage
}')

coordinates(g3) <- list(
  x = c(SexBirth = 0, Employed = 1.5, U = 1.5, Educ = 1.5, Wage = 3.2),
  y = c(SexBirth = 0, Employed = 1.2, U = 2.5, Educ = -1.3, Wage = 0))

# the open square marks the node the sample is conditioned on
ggdag(tidy_dagitty(g3), text = FALSE, node_size = 9) +
  geom_point(data = data.frame(x = 1.5, y = 1.2), aes(x = x, y = y),
    inherit.aes = FALSE, shape = 0, size = 10, stroke = 1.2,
    colour = "grey15") +
  geom_dag_text_repel(aes(label = name), colour = "black", size = 3.2,
    box.padding = 1.3, seed = 5) +
  theme_dag() +
  expand_limits(x = c(-0.7, 4.0), y = c(-1.9, 3.0))
```





Figure 16.2.: The sample restriction as a conditioning statement. The **boxed** node is the variable the analysis conditions on — not by adding a covariate, but by the sample definition. **Employed** is a common effect of sex and of unobserved wage determinants U , so boxing it opens $\text{SexBirth} \rightarrow \text{Employed} \leftarrow U \rightarrow \text{Wage}$.

The box is worth dwelling on, because it is drawn by the sampling frame rather than by any modelling choice. Nobody typed `+ employed` into a regression; the extract simply contains only full-year, full-time workers. A sample definition **is** a conditioning statement, and most applied DAGs omit the selection node entirely — which makes the graph look identified when it is not.

`dagitty` will confirm that there is no way out by adjustment. Note that a selection variable has to be declared as such: marking a node `[selected]` tells `dagitty` the sample is restricted on it, which is a different claim from conditioning on an ordinary covariate.

```

g3_sel <- dagitty('dag {
  SexBirth -> Educ
  SexBirth -> Wage
  SexBirth -> Employed
  U -> Employed
  U -> Wage
  Educ -> Wage
  U [latent]
  Employed [selected]
}
```

16. Gender Wage Gap Analysis Using US Current Population Survey Data

```
SexBirth [exposure]
Wage [outcome]
}')

# how many covariate sets identify the total effect under this selection?
length(adjustmentSets(g3_sel, effect = "total"))
```

```
[1] 0
```

```
isAdjustmentSet(g3_sel, c())      # the raw comparison
```

```
[1] FALSE
```

```
isAdjustmentSet(g3_sel, "Educ")   # adding education does not help either
```

```
[1] FALSE
```

```
# and the path that does the damage
p3 <- paths(g3, from = "SexBirth", to = "Wage", Z = "Employed")
data.frame(path = p3$paths, open = p3$open)
```

```

                                path open
1      SexBirth -> Educ -> Wage TRUE
2 SexBirth -> Employed <- U -> Wage TRUE
3      SexBirth -> Wage TRUE
```

Zero valid adjustment sets, and the raw comparison is no longer among them. Conditional on being in the sample, sex is no longer independent of potential wages: sex is exogenous, but *selected* sex is not. Compare this with the occupation problem earlier in the chapter. There the collider opened because we *chose* to add a control, so declining to add it was an available fix. Here the box is imposed before any modelling decision, so “do not condition on it” is not an option — the remedy has to attack the box itself rather than the covariate list, which is why it takes data on non-workers (Lee, 2009, sharp bounds; Neal, 2004; Olivetti and Petrongolo, 2008, imputation for non-workers) rather than a better regression.

The direction is unusually easy to sign here. Women are less likely to hold full-year full-time jobs, so clearing that bar takes more of whatever else helps — meaning employed women are *positively* selected on wage-raising unobservables relative to employed men. If so, the observed gap **understates** the true one. Olivetti and Petrongolo (2008) find exactly this signature across countries: measured gender wage gaps are smaller where female employment rates are lower, and imputing wages for non-workers widens them most in the low-participation countries.

16.2.3.1. How much does this actually cost? A simulation

Everything above is graph-theoretic: the paths are open or closed, the adjustment sets exist or do not. It says nothing about magnitude. Because the non-workers are missing from this extract by construction, the honest way to put numbers on it is to simulate a population where the truth is known, apply the selection, and see what survives.

The design mirrors the graph: sex is randomly assigned, U is an unobserved wage determinant, women are less likely to be employed full-year full-time, and higher U makes employment more likely. The true effect is set at -0.10 log points.

```
set.seed(42)
N <- 300000
tau <- -0.10 # TRUE effect of sex on log wage

sex <- rbinom(N, 1, 0.5) # birth sex, randomly assigned
U <- rnorm(N) # unobserved: health, drive, reservation wage
wage <- 2.5 + tau * sex + 0.5 * U + rnorm(N, 0, 0.4)

pi_sel <- plogis(0.8 - 0.9 * sex + 0.7 * U) # women less likely; high U more likely
emp <- rbinom(N, 1, pi_sel)

d <- data.frame(wage, sex, U, pi_sel)[emp == 1, ]

b <- function(fit) coef(fit)[["sex"]]

# 1. the truth, from the full population
full <- b(lm(wage ~ sex))

# 2. the selected sample, as an analyst would see it
naive <- b(lm(wage ~ sex, data = d))

# 3. reweight so the sex ratio is exactly 50/50
w_bal <- ifelse(d$sex == 1, 0.5 / mean(d$sex), 0.5 / (1 - mean(d$sex)))
rewt <- b(lm(wage ~ sex, data = d, weights = w_bal))

# 4. randomly drop men until the counts match
n_w <- sum(d$sex == 1)
men <- d[d$sex == 0, ]
kept <- men[sample(nrow(men), n_w), ]
subs <- b(lm(wage ~ sex, data = rbind(kept, d[d$sex == 1, ])))

# 5. weight by the inverse probability of selection (needs U -- infeasible in practice)
ipsw <- b(lm(wage ~ sex, data = d, weights = 1 / d$pi_sel))

data.frame(
  approach = c("full population regression", "selected sample, naive",
    "sex ratio reweighted to 50/50", "men randomly subsampled to equal n",
```

16. Gender Wage Gap Analysis Using US Current Population Survey Data

```
"inverse-probability-of-selection weighted"),
estimate = round(c(full, naive, rewt, subs, ipsw), 4)
)
```

	approach	estimate
1	full population regression	-0.0960
2	selected sample, naive	-0.0389
3	sex ratio reweighted to 50/50	-0.0389
4	men randomly subsampled to equal n	-0.0393
5	inverse-probability-of-selection weighted	-0.1014

```
c(employment_rate_men = round(mean(emp[sex == 0]), 3),
employment_rate_women = round(mean(emp[sex == 1]), 3))
```

employment_rate_men	employment_rate_women
0.671	0.479

```
c(mean_U_employed_men = round(mean(d$U[d$sex == 0]), 3),
mean_U_employed_women = round(mean(d$U[d$sex == 1]), 3))
```

mean_U_employed_men	mean_U_employed_women
0.205	0.332

Three things to take from that table.

The full-population regression recovers the true -0.10 up to sampling noise, as it must — sex is randomly assigned, so nothing confounds it. The naive selected-sample estimate is **wrong by a factor of about two and a half**, and in the direction the graph predicted: attenuated toward zero, because employed women are positively selected on U (mean $+0.33$ against $+0.21$ for employed men). The signed prediction of the previous section is not merely plausible; it is what happens.

Balancing does nothing. Reweighting the sex ratio to exactly 50/50 leaves the estimate unchanged to four decimal places, and randomly discarding men until the counts match leaves it unchanged as well. The reason is worth stating as a general rule: **no random sampling operation can remove a bias**. Bias is a property of a distribution, and random sampling preserves distributions — a random subsample of men has the same mean U as all men, $+0.21$ either way. The problem is not that women are outnumbered; it is that the employed women are a *differently selected* subset of women than the employed men are of men. Reweighting on sex changes the marginal distribution of sex, while the bias lives in the conditional distribution of U given sex.

Only the last row recovers the truth, and it does so by weighting on the probability of selection — which requires U . That is exactly what is unavailable.

16.2.3.2. Lee bounds: what you can get without U

Lee (2009) avoids needing U by trimming the over-selected group on the *outcome* rather than on the unobservable, in both extreme directions. The trim fraction is fixed by the selection rates, $q = (p_{\text{men}} - p_{\text{women}})/p_{\text{men}}$: drop the lowest-paid q of men to make them look as good as possible, then the highest-paid q to make them look as bad as possible. The two results bracket the effect.

The important question is not whether the bounds are valid — they are — but how wide they are. That depends entirely on how unbalanced the selection is:

```
lee_run <- function(b_sex, N = 300000, tau = -0.10) {
  sex <- rbinom(N, 1, 0.5)
  U <- rnorm(N)
  wage <- 2.5 + tau * sex + 0.5 * U + rnorm(N, 0, 0.4)
  emp <- rbinom(N, 1, plogis(0.8 + b_sex * sex + 0.7 * U))
  d <- data.frame(wage, sex)[emp == 1, ]

  p_m <- mean(emp[sex == 0]); p_w <- mean(emp[sex == 1])
  q <- (p_m - p_w) / p_m

  men <- d$wage[d$sex == 0]
  wbar <- mean(d$wage[d$sex == 1])
  lo <- wbar - mean(men[men >= quantile(men, q, names = FALSE)])
  hi <- wbar - mean(men[men <= quantile(men, 1 - q, names = FALSE)])

  data.frame(p_men = round(p_m, 3), p_women = round(p_w, 3), q = round(q, 3),
             naive = round(coef(lm(wage ~ sex, data = d))["sex"], 4),
             lower = round(lo, 3), upper = round(hi, 3),
             width = round(hi - lo, 3),
             signed = ifelse(hi < 0, "yes", "no"),
             covers_truth = (lo <= tau && tau <= hi))
}

set.seed(7)
do.call(rbind, lapply(c(-0.9, -0.5, -0.25, -0.10, -0.03), lee_run))
```

	p_men	p_women	q	naive	lower	upper	width	signed	covers_truth
1	0.672	0.477	0.290	-0.0365	-0.337	0.265	0.602	no	TRUE
2	0.674	0.568	0.157	-0.0669	-0.245	0.112	0.357	no	TRUE
3	0.674	0.621	0.079	-0.0857	-0.186	0.015	0.201	no	TRUE
4	0.672	0.652	0.029	-0.0960	-0.139	-0.053	0.086	yes	TRUE
5	0.673	0.667	0.008	-0.0959	-0.110	-0.082	0.028	yes	TRUE

The bounds cover the true -0.10 in every row, so the method does what it promises. But the width scales almost proportionally with the trim fraction, and the effect can only be *signed* once the employment rates differ by a couple of percentage points. At a 19-point gap the interval runs from -0.34 to $+0.26$ and contains zero: valid, and useless.

Now read the `naive` column alongside `width`. The naive estimate degrades in lockstep with the bounds widening — reading down the table, -0.037 , -0.067 , -0.086 , -0.096 , -0.096 against a truth of -0.10 . So Lee bounds are tight exactly when they are not needed and wide exactly when they are. That inverts into what they are actually for here: the width is a **measure of how much the selection problem is costing**, and it is information one otherwise does not have at all. Their value in this setting is more diagnostic than inferential. No method manufactures information; this one stops a badly selected point estimate from passing as a precise one.

Three caveats before anyone reaches for it. The load-bearing assumption is **monotonicity** — being female never *increases* anyone’s employment probability, which is an individual-level claim, stronger than the average pattern. The target quantity changes: the bounds are for the effect among the **always-employed**, whoever they happen to be, not the population. And the procedure needs both selection rates, so it cannot be run on `wage2015_subsample_inference.csv` at all — that file contains only employed people. One would have to go back to the unfiltered CPS for men’s and women’s full-year full-time employment rates.

On real US data those rates differ by something on the order of ten to fifteen percentage points for ages 25–64, putting q near 0.2 — the second row of the table, where the interval contains zero. So the likely honest conclusion from doing this properly is that the selected-sample gender gap is consistent with anything from a large penalty to a modest premium, and that no covariate adjustment narrows it. That is a stronger and more useful statement than the -3.8% point estimate, even though it names no number.

A note on notation, since U now appears in two graphs doing two different jobs. In the occupation graph, U has to drive *occupational sorting* — drive, negotiation propensity, expected career continuity. In the selection graph it has to drive *labour supply* — health, reservation wage (non-labour income, spousal earnings), latent productivity. The two sets overlap but are not the same, and the two collider problems are separate.

So the honest summary is uncomfortable: -3.8% is the best-identified number in this chapter and the least informative about anything one would want to act on, while -7.5% is the worst-identified and the closest to the question people actually ask. Identification improves as conditioning is stripped away; interpretability improves as it is added. No estimator resolves that trade-off — only a different design does, which is what audit and correspondence studies are for.

16.2.4. Two treatments, two adjustment sets

There is a natural objection to everything above, and it is a good one. *Suppose we set aside the fact that sex shapes schooling and experience, and treat those as simply given — this woman’s actual history. Then asking what she would have been paid as a man, holding that history fixed, is exactly the question of interest, and controlling for education is obviously the right thing to do.*

That is right, and it exposes something the discussion so far has glossed. **“Bad control” is not a property of a variable.** It is a property of the triple (*variable, treatment, timing*). Re-date the treatment and the classification flips:

Treatment	Education is	Controlling is
Sex assigned at conception	post-treatment mediator	bad control

Treatment	Education is	Controlling is
Perceived sex at the hiring decision	pre-treatment covariate	appropriate

The second row is Greiner and Rubin’s (2011) framing — treatment as *perception of an immutable characteristic at a specified decision point* — and it is precisely what correspondence studies estimate. It is also the legally operative question for disparate treatment. So it is the better estimand, not a compromise.

Encoding it, however, produces a result sharper than the verbal argument suggests. Let **SexBirth** be sex at birth, **Perceived** the perceived sex facing the employer, and **A** unobserved ability affecting both schooling and wages:

```
library(ggdag)
library(ggplot2)
library(dagitty)

g2 <- dagitty('dag {
  SexBirth -> Perceived
  SexBirth -> Educ
  A -> Educ
  A -> Wage
  Educ -> Wage
  Perceived -> Wage
  A [latent]
  Perceived [exposure]
  Wage [outcome]
}')

coordinates(g2) <- list(
  x = c(SexBirth = 0, Perceived = 1.3, Educ = 1.3, A = 1.0, Wage = 3.1),
  y = c(SexBirth = 0, Perceived = 1.1, Educ = -1.1, A = -2.3, Wage = 0))

ggdag(g2, text = FALSE, node_size = 9) +
  geom_dag_text_repel(aes(label = name), colour = "black", size = 3.1,
    box.padding = 1.4, seed = 11) +
  theme_dag() +
  expand_limits(x = c(-0.7, 3.8), y = c(-3.0, 1.8))
```




Figure 16.3.: Treatment re-dated to the hiring decision. Education now *precedes* the treatment — but it is still a collider between sex at birth and unobserved ability A .

Now ask `dagitty` which covariates would license the comparison:

```

c("{}" = isAdjustmentSet(g2, c()),
  "{Educ}" = isAdjustmentSet(g2, "Educ"))

```

```

{} {Educ}
FALSE FALSE

```

```

adjustmentSets(g2, effect = "total")

```

```

{ SexBirth }

```

```

paths_table2 <- function(Z = NULL) {
  p <- paths(g2, from = "Perceived", to = "Wage", Z = Z)
  data.frame(path = p$paths, open = p$open)
}

```

```

paths_table2()          # nothing conditioned on

```



```

                                path  open
1                               Perceived -> Wage  TRUE
2      Perceived <- SexBirth -> Educ -> Wage  TRUE
3 Perceived <- SexBirth -> Educ <- A -> Wage FALSE

```

```
paths_table2(Z = "Educ") # conditioning on education
```

```

                                path  open
1                               Perceived -> Wage  TRUE
2      Perceived <- SexBirth -> Educ -> Wage FALSE
3 Perceived <- SexBirth -> Educ <- A -> Wage  TRUE

```

Three things follow, and the third is the point.

The objection is half-vindicated. Adjusting for nothing is *invalid* under this estimand. Sex at birth confounds perceived sex and wages through the education channel, so the raw difference is contaminated for this question — part of it is human capital rather than employer conduct. Controlling for education does remove a genuine bias. Everything said earlier about bad controls was conditional on the *other* estimand.

But it introduces a different one. Conditioning on education closes the human-capital backdoor and simultaneously opens `Perceived <- SexBirth -> Educ <- A -> Wage`. Education is still a collider between sex at birth and ability, and re-dating the treatment does nothing about that. One bias is traded for another.

The only valid adjustment set is {SexBirth}, and it has no overlap. Sex at birth determines perceived sex, so conditioning on it leaves no variation in the treatment. Positivity fails outright. The effect of perceived sex holding human capital fixed is therefore **not identified from observational data under any conditioning strategy** — not because this extract is small or its covariates coarse, but because of the structure of the graph.

That last point also says exactly what a design must do: sever `SexBirth -> Perceived`. Randomise the signal independently of biological sex, and the backdoor disappears, so *no adjustment is needed at all*. This is why an audit study requires no controls, and why it answers a question a regression provably cannot. Goldin and Rouse (2000) is the cleanest case — the same musician is evaluated with and without the screen, so ability is held fixed exactly rather than statistically.

One usable consequence in the meantime. The two biases run in opposite directions, which brackets the answer informally. With no conditioning, the education backdoor is open and — since women here are *more* educated — their human-capital advantage offsets part of the treatment effect, biasing toward zero. With education conditioned on, the collider opens and equally-educated women are pushed low on *A*, inflating the residual gap. So the quantity of interest plausibly sits *between* -3.8% and -11.3% , and neither endpoint is it. This is not a formal bound — the sign of each bias rests on assumptions rather than proof — but it is more informative than either number alone.

Finally, a normative point that no amount of graph-reading settles. Holding education and experience fixed accepts whatever produced those differences as the baseline. If they were themselves partly produced by discrimination, then conditioning on them conditions on discrimination's own output. That is the right choice for *is this employer treating equivalent applicants differently* and the

16. Gender Wage Gap Analysis Using US Current Population Survey Data

wrong choice for *is the labour market fair to women* — roughly the disparate-treatment versus disparate-impact distinction. The data does not decide it; the question does.

There is a further objection which holds that the resolution just reached — re-date the treatment to perception, then randomise it — does not work either, and that the normative point in the paragraph above is not a footnote to the analysis but the whole of it. Because it needs the empirical material in place first, it is taken up in the closing section, Section 16.4.

With all that on the table, here is the model.

```
reg4 <- lm(lwage ~ female*(exp1 + sohs + hsg+ soc1 + clg + mw + sout + we + occ2 + ind2), data=
avg_comparisons(reg4,
  variables = "female",
  newdata = subset(data, female == 1))
```

Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
-0.075	0.0162	-4.62	<0.001	18.0	-0.107	-0.0432

Term: female

Type: response

Comparison: 1 - 0

So in the model with full interaction, we have a gap of 7.5%, against 11.3% without occupation and industry.

The tempting reading of that pair — “occupational and industrial sorting accounts for roughly four points of the gap” — is exactly the reading the graph above forbids. Three things are mixed into that four-point move, and nothing in the output separates them:

1. genuine mediation, the part of the gap that really does operate through sorting;
2. collider bias, from the gender- U association induced by conditioning;
3. ordinary confounding of the occupation-wage relationship.

The likely *direction* of (2) is worth spelling out. If women who enter male-dominated, high-paying occupations are positively selected on U — they had more to overcome to get there — then within those occupations they are better on unobservables than the men they are being compared with, which pushes the within-occupation gap **toward zero**. Some of the attenuation from 11.3% to 7.5% may therefore be bias rather than mediation.

The defensible way to report this: neither number is “the” gender gap. -11.3% is a controlled direct effect holding education, experience and region fixed; -7.5% additionally holds occupation and industry fixed. The total effect — the only rung with no mediators held fixed — is the raw -3.8% . What we *can* do for the last step is ask how strong U would have to be to account for the whole 4-point difference between the bottom two rungs.

This also gives a precise form to the objection most readers will have already framed for themselves: *maybe men and women differ in some unobserved, productivity-relevant way, and that is what we are picking up*. Note first that such a trait would be a **mediator**, not a confounder — nothing causes sex, so anything that differs by sex is downstream of it, and the total effect correctly includes

it. The objection is therefore not “the comparison is biased” but “I want the direct effect holding that trait fixed,” which is a different and unidentified estimand. Note second that controlling for observables does not approximate it: education is a *collider* between sex and any such trait, so conditioning on education induces exactly the sex–ability association one hoped to remove. Given the composition table above — women here are more educated — that induced association would push the estimated gap further from zero, which is the direction the numbers in fact move.

What the sensitivity analysis below can do is convert that unfalsifiable worry into a magnitude.

16.2.4.1. How strong would U have to be?

The omitted-variable-bias machinery in `sensemakr` (Cinelli and Hazlett, 2020) needs a single treatment coefficient, so for this exercise we drop the interactions and use additive specifications. That costs us the heterogeneity, but the attenuation we are interrogating survives the simplification almost exactly.

```
library(sensemakr)

Xb <- "exp1 + sohs + hsg + socl + clg + mw + sout + we"

fit_total <- lm(as.formula(paste("lwage ~ female +", Xb)), data = data)
fit_adj   <- lm(as.formula(paste("lwage ~ female +", Xb, "+ occ2 + ind2")), data = data)

c(total      = coef(fit_total)["female"],
   adjusted = coef(fit_adj)["female"],
   attenuation = coef(fit_adj)["female"] - coef(fit_total)["female"])
```

total.female	adjusted.female	attenuation.female
-0.11375231	-0.07285748	0.04089483

```
sens <- sensmakr(model = fit_adj, treatment = "female",
                 benchmark_covariates = c("clg", "exp1"),
                 kd = c(1, 2, 3))
summary(sens)
```

Sensitivity Analysis to Unobserved Confounding

Model Formula: `lwage ~ female + exp1 + sohs + hsg + socl + clg + mw + sout + we + occ2 + ind2`

Null hypothesis: `q = 1` and `reduce = TRUE`

-- This means we are considering biases that reduce the absolute value of the current estimate

-- The null hypothesis deemed problematic is $H_0:\tau = 0$

Unadjusted Estimates of 'female':

Coef. estimate: -0.0729

16. Gender Wage Gap Analysis Using US Current Population Survey Data

Standard Error: 0.015
t-value (H0:tau = 0): -4.8485

Sensitivity Statistics:

Partial R2 of treatment with outcome: 0.0046
Robustness Value, q = 1: 0.0656
Robustness Value, q = 1, alpha = 0.05: 0.0396

Verbal interpretation of sensitivity statistics:

-- Partial R2 of the treatment with the outcome: an extreme confounder (orthogonal to the covariates)
-- Robustness Value, q = 1: unobserved confounders (orthogonal to the covariates) that explain
-- Robustness Value, q = 1, alpha = 0.05: unobserved confounders (orthogonal to the covariates)

Bounds on omitted variable bias:

--The table below shows the maximum strength of unobserved confounders with association with treatment

Bound	Label	R2dz.x	R2yz.dx	Treatment	Adjusted Estimate	Adjusted Se	Adjusted T
1x	clg	0.0001	0.0124	female	-0.0718	0.0149	-4.8081
2x	clg	0.0002	0.0247	female	-0.0708	0.0148	-4.7680
3x	clg	0.0002	0.0371	female	-0.0697	0.0147	-4.7275
1x	exp1	0.0016	0.0338	female	-0.0649	0.0148	-4.3914
2x	exp1	0.0032	0.0676	female	-0.0570	0.0145	-3.9199
3x	exp1	0.0048	0.1014	female	-0.0490	0.0143	-3.4322
Adjusted Lower CI		Adjusted Upper CI					

```

      r2dz.x = grid$r2dz, r2yz.dx = grid$r2yz,
      reduce = FALSE)

grid |>
  group_by(r2dz) |>
  slice_min(abs(adj - target), n = 1) |>
  ungroup() |>
  transmute(`partial R2 with gender` = r2dz,
            `partial R2 with wage`   = round(r2yz, 3),
            `restored estimate`      = round(adj, 4))

```

```

# A tibble: 4 x 3
  `partial R2 with gender` `partial R2 with wage` `restored estimate`
      <dbl>                <dbl>                <dbl>
1      0.0125              0.115              -0.114
2      0.02               0.071              -0.114
3      0.033              0.042              -0.114
4      0.06               0.023              -0.114

```

```

# sensemakr is attached again here on purpose: the chunk above is cached, so on a
# cache hit its library() call never runs and plot() would dispatch to
# plot.default instead of the sensemakr method.
library(sensemakr)

plot(sens)

```




Figure 16.4.: Adjusted estimate of the gender gap as a function of a hypothetical confounder’s partial R^2 with gender (horizontal) and with log wage (vertical), benchmarked against college and experience.

The answer is uncomfortable. A confounder explaining about **3% of the residual variance of gender and about 4% of the residual variance of wages** is enough to restore the full total effect — that is, enough for occupation to be contributing no genuine mediation at all. For scale, `sensemkr` reports a robustness value of 6.6%: a confounder explaining more than that share of both residual variances would drive the adjusted estimate to zero. So a U roughly *half* as strong as the one needed to erase the within-occupation gap entirely is already strong enough to explain the whole $11.3\% \rightarrow 7.5\%$ attenuation.

The benchmark rows give another angle. Experience explains 3.4% of the residual variance of wages but only 0.16% of the residual variance of gender, so along the outcome dimension the required U is only a little stronger than experience. Along the gender dimension the required multiple looks large, but that is mostly because gender is close to orthogonal to experience in this sample; in absolute terms we are still only asking for a few percent of residual variance.

None of this shows that occupational sorting explains nothing. It shows that this design cannot distinguish “sorting explains four points” from “conditioning on a collider costs four points,” and that the confounder needed to tip the balance is unremarkable in size. Which is why the total effect, not the occupation-adjusted number, is the estimate to lead with.

Now let's try nonparametric estimation.

16.2.5. Nonparametric estimation

If we'd like to go for nonparametric, relaxing linearity assumption, then we can use “npcausal”. Note this may take a while, depending on which machine learning packages you decide to include. We can also use other nonparametric estimators, such as TMLE or doubleML.

```
library(npcausal)

#SL.library <- c("SL.earth","SL.glmnet","SL.glm.interaction", "SL.mean","SL.ranger", "SL.xgboost")

SL.library <- c("SL.earth","SL.glmnet","SL.mean","SL.ranger", "SL.xgboost")

# construct W from model matrix of reg3
reg3 <- lm(lwage ~ exp_dm + sohs + hsg+ socl + clg + mw + sout + we + occ2 + ind2, data=data)
W <- as.data.frame(model.matrix(reg3)[, -1]) # remove intercept

Y <- data$lwage
A <- data$female
set.seed(123)
aipw_ate <- ate(y=Y, a=A, x=W, nsplits=5, sl.lib=SL.library)
```

```
|
|
| 0%

|
|====
| 5%
|
|=====
| 10%
|
|=====
| 15%
|
|=====
| 20%
|
|=====
| 25%
|
|=====
| 30%

|
|=====
| 35%
|
```

16. Gender Wage Gap Analysis Using US Current Population Survey Data

```

=====| 40%
|
=====| 45%
|
=====| 50%
|
=====| 55%
|
=====| 60%
|
=====| 65%
|
=====| 70%
|
=====| 75%
|
=====| 80%
|
=====| 85%
|
=====| 90%
|
=====| 95%
|
=====| 100%

```

	parameter	est	se	ci.ll	ci.ul	pval
1	E{Y(0)}	2.99695506	0.01093465	2.97552315	3.01838697	0
2	E{Y(1)}	2.93906408	0.01327342	2.91304817	2.96507999	0
3	E{Y(1)-Y(0)}	-0.05789098	0.01597522	-0.08920241	-0.02657955	0

```
aipw_att <- att(y=Y, a=A, x=W, nsplits=5, sl.lib=SL.library)
```

```

|
|
| 0%

|
|=====| 10%
|
|=====| 20%

|
|=====| 30%
|

```



```

|=====| 40%
|
|=====| 50%
|
|=====| 60%
|
|=====| 70%
|
|=====| 80%
|
|=====| 90%
|
|=====| 100%

```

	parameter	est	se	ci.ll	ci.ul	pval
1	E(Y A=1)	2.94948490	0.01160804	2.92673314	2.97223666	0
2	E{Y(0) A=1}	3.01689509	0.01221863	2.99294658	3.04084360	0
3	E{Y-Y(0) A=1}	-0.06741019	0.01409116	-0.09502886	-0.03979152	0

The two fits above print only SuperLearner's progress bars, so the estimates themselves have to be asked for:

```
aipw_ate$res
```

	parameter	est	se	ci.ll	ci.ul	pval
1	E{Y(0)}	2.99695506	0.01093465	2.97552315	3.01838697	0
2	E{Y(1)}	2.93906408	0.01327342	2.91304817	2.96507999	0
3	E{Y(1)-Y(0)}	-0.05789098	0.01597522	-0.08920241	-0.02657955	0

```
aipw_att$res
```

	parameter	est	se	ci.ll	ci.ul	pval
1	E(Y A=1)	2.94948490	0.01160804	2.92673314	2.97223666	0
2	E{Y(0) A=1}	3.01689509	0.01221863	2.99294658	3.04084360	0
3	E{Y-Y(0) A=1}	-0.06741019	0.01409116	-0.09502886	-0.03979152	0

Two estimates come out of this: an AIPW **ATE** of -5.8% (95% CI -8.9% to -2.7%) and an AIPW **ATT** of -6.7% (95% CI -9.5% to -3.9%).

The comparison worth making is like with like. The -7.5% from the fully interacted linear model above was an ATT (it was computed on `newdata = subset(data, female == 1)`), so it should be set against the nonparametric **ATT** of -6.7% , not against the ATE. Those two agree closely, which is reassuring: dropping the linearity assumption altogether moves the estimate by less than a percentage point, and both use the same covariate set including occupation and industry.

That last clause is also the caveat. Because both specifications condition on occupation and industry, both inherit the bad-control problem set out above: the agreement between -7.5% and -6.7% tells

us the two estimators recover the same quantity, not that the quantity is a causal effect. Relaxing linearity addresses functional form; it leaves the collider untouched.

The ATE and ATT point estimates do differ a little, in the direction of a larger gap among women than in the sample as a whole — which would indicate effect heterogeneity across the covariates. Their confidence intervals overlap heavily, though, so this data cannot really distinguish them.

16.3. Beyond the mean: where in the distribution is the gap?

Everything so far has been a mean estimand — ATE, ATT, PAGE, PAGM, PAGW. But the best-known stylised fact about the gender wage gap is that it is *not uniform* across the distribution: a **glass ceiling** if the gap widens at the top, a **sticky floor** if it widens at the bottom. A chapter that reports only means cannot see either.

There are two distinct questions here, and they are routinely conflated.

16.3.1. The direct comparison

The distributional analogue of the raw gap needs no special machinery at all: compare the two groups' wage quantiles and difference them. Under the exogeneity argument above this is the total effect at each quantile — the well-identified rung of the ladder.

```
library(tidyverse)

taus <- c(.10, .25, .50, .75, .90)

qgap <- function(dat) {
  sapply(taus, function(t)
    quantile(dat$lwage[dat$female == 1], t, names = FALSE) -
    quantile(dat$lwage[dat$female == 0], t, names = FALSE))
}

set.seed(1)
bs <- replicate(2000, qgap(data[sample(nrow(data), replace = TRUE), ]))

tibble(tau = taus,
       gap = qgap(data),
       lo = apply(bs, 1, quantile, .025),
       hi = apply(bs, 1, quantile, .975),
       pct = paste0(round(100 * (exp(qgap(data)) - 1), 1), "%")) |>
  mutate(across(gap:hi, ~round(.x, 4)))
```

A tibble: 5 x 5

	tau	gap	lo	hi	pct
	<dbl>	<dbl>	<dbl>	<dbl>	<chr>
1	0.1	0	-0.0392	0.0282	0%

2	0.25	-0.0521	-0.0909	0	-5.1%
3	0.5	-0.0202	-0.0513	0	-2%
4	0.75	-0.069	-0.105	-0.0177	-6.7%
5	0.9	-0.0692	-0.143	-0.0126	-6.7%

The gap is indistinguishable from zero at the 10th percentile and around -6.7% at the 75th and 90th — a glass-ceiling pattern, though a noisy one. Two features of this table are artefacts worth naming. The exact zeros are heaping: log wages take repeated values, so the empirical quantile lands on the same number for both groups. That heaping is also why the sequence is not monotone (the gap at $\tau = 0.25$ exceeds the gap at $\tau = 0.50$). The direct comparison is exact but coarse in this data.

16.3.2. RIF regression, and what it actually estimates

Firpo, Fortin and Lemieux (2009) provide the machinery for putting covariates into a distributional statistic. For any functional ν of the outcome distribution, the influence function measures how ν responds to adding mass at y ; recentring it gives a transformed outcome whose *mean is the statistic itself*. For the τ -th quantile,

$$\text{RIF}(y; q_\tau) = q_\tau + \frac{\tau - \mathbb{1}\{y \leq q_\tau\}}{f_Y(q_\tau)}. \quad (16.7)$$

Compute the sample quantile, estimate the density there, build the RIF as a new outcome, and regress it on covariates. Because $E[\text{RIF}] = q_\tau$, the coefficients speak to the **unconditional** quantile. This is the key contrast with `quantreg::rq`, which gives *conditional* quantile effects — the effect at the τ -th quantile of the conditional distribution, among people with the same covariates. For “is there a glass ceiling in the wage distribution,” the unconditional version is the one wanted, and fitting `rq` instead answers a different question.

```
library(tidyverse)
library(marginaleffects)

rif_quantile <- function(y, tau) {
  q <- quantile(y, tau, names = FALSE)
  dn <- density(y)
  q + (tau - as.numeric(y <= q)) / approx(dn$x, dn$y, xout = q)$y
}

Xb <- "exp1 + sohs + hsg + soc1 + clg + mw + sout + we"

map_dfr(taus, function(t) {
  dd <- transform(data, rif = rif_quantile(lwage, t))
  m0 <- lm(rif ~ female, data = dd)
  m1 <- lm(as.formula(paste("rif ~ female*(", Xb, ")")), data = dd)
  tibble(tau = t,
         rif_no_covariates = round(coef(m0)["female"], 4),
```

```

rif_with_covariates =
  round(avg_comparisons(m1, variables = "female",
                        vcov = "HC3")$estimate, 4))
})

```

```

# A tibble: 5 x 3
  tau rif_no_covariates rif_with_covariates
<dbl>          <dbl>          <dbl>
1  0.1          -0.012          -0.0822
2  0.25         -0.0425          -0.120
3  0.5          -0.0409          -0.124
4  0.75         -0.0557          -0.134
5  0.9          -0.0775          -0.159

```

Now the important caveat, and it is an estimand distinction rather than a technical one. **The RIF-regression coefficient on female is not the gap between men’s and women’s wage quantiles.** It is the effect on the overall τ -quantile of marginally raising the *share* female, holding conditional distributions fixed — a composition effect. That is why the numbers here disagree with the direct comparison at every τ : at the median, -0.020 directly against -0.041 by RIF. Both are legitimate; they answer different questions. Reporting a RIF coefficient as “the gap at the 90th percentile” is a common slip.

With that said, the RIF panel is the more readable of the two here, precisely because the density smoothing interpolates through the heaping that makes the direct comparison lumpy. Both columns rise monotonically in τ , and the covariate column runs from -8.2% at the 10th percentile to -15.9% at the 90th. The mean estimates from earlier in the chapter — -3.8% raw, -11.4% adjusted — are averages over that gradient.

Three further caveats before anyone leans on these numbers:

- **RIF is a first-order approximation**, valid for small covariate shifts. Flipping a binary regressor for the entire sample is not small, so the approximation degrades. Exact alternatives are Firpo (2007) propensity-weighted quantile treatment effects, or Chernozhukov, Fernández-Val and Melly (2013) counterfactual distributions.
- **The standard errors here are too small.** The density estimate contributes uncertainty that `lm` knows nothing about; bootstrap the whole procedure, as done for the direct comparison above.
- **The covariate column is a distributional rung of the same CDE ladder**, and inherits every bad-control problem from the earlier section. It is not a better estimate of the same thing.

Finally, a quantile treatment effect compares quantiles of two marginal distributions. It is not the effect *for the person at* quantile τ unless one is willing to assume rank invariance — that the same individuals occupy the same ranks in both counterfactual worlds.

16.4. What is any of this measuring?

Two directed acyclic graphs and a **dagitty** verification have now been used to sort covariates into good and bad controls. It is worth closing with an objection to that whole procedure, because it comes from a literature economists rarely read and because it bears directly on the re-dating manoeuvre of Section 16.2.4. The principals are **Lily Hu** (Philosophy, Yale) and **Issa Kohler-Hausmann** (Law, Yale); the two papers to read are Hu and Kohler-Hausmann (2020) and Hu and Kohler-Hausmann (2024).

16.4.1. Causal versus constitutive

Their objection turns on a distinction the causal toolkit has no notation for.

	Causal	Constitutive
Relation	X produces Y	X is part of what it is to be Y
Timing	X precedes Y	simultaneous
Contingency	can have X without Y	cannot have Y without X
Intervention	wiggle X , hold others fixed	wiggling X changes <i>what the thing is</i>

Non-social cases make it clean. Being unmarried does not *cause* someone to be a bachelor — it is part of what being a bachelor consists in. A chess position does not *cause* checkmate; it constitutes it. Ask “what would happen to this checkmate if the king were not trapped?” and there is no answer, not because the quantity is hard to estimate but because there would be no checkmate left to ask about.

16.4.2. Why gender’s features are partly constitutive

Consider the observation that women do more caregiving. On the causal reading there is a category, woman, fixed independently, which then produces caregiving downstream. On the constitutive reading, part of what it **is** to occupy the social position “woman” is to be subject to caregiving expectations — the norm is not produced by the category, it is one of the threads the category is woven from.

Press on what makes someone socially a woman in the sense labour markets respond to. Not chromosomes; markets do not observe those. It is occupying a position defined by a cluster of expectations, treatments and readings. Strip the cluster away and no social category remains, only biology. On this view the cluster is not the category’s output; it is largely the category itself.

The cleanest comparison is a social role. *What is the causal effect of being a manager on having authority?* That is not a causal question — having authority is part of what being a manager consists in. One cannot hold authority fixed and vary managerial status, because nothing would be left for “manager” to refer to.

“Partly” is doing real work, though. Hu and Kohler-Hausmann claim *many*, not all, of the effects attributed to sex are constitutive, so the category is a mixture: caregiving norms and being read as a plausible authority figure are plausibly constitutive, while one firm’s decision to deny one promotion is causal (contingent, could go otherwise, category unchanged) and average upper-body strength is causal-biological. The difficulty is that **no method sorts them**. A usable test is to ask whether removing the feature would leave the same category or a different thing entirely — but applying it is a substantive sociological claim, and drawing a DAG settles it silently.

16.4.3. Why this defeats modularity

Modularity is the formal assumption that the dependencies along one pathway are invariant to interventions on other pathways. It is exactly what licensed every arrow-cutting exercise above.

If **female** → **educ** is causal, “flip gender, hold education fixed” is coherent: imagine changing the assignment while the education-producing mechanism keeps running. If the relation is partly constitutive, the instruction does not describe anything. An intervention that changed a woman’s gender while leaving the associated expectations intact is not an intervention on gender-as-a-social-category — it holds fixed part of the very antecedent it purports to vary. The counterfactual has no determinate content.

This is a different *kind* of problem from everything else in this chapter, and the difference is worth stating flatly. The mediator, collider, selection and positivity problems are all **epistemic**: each says the quantity exists and we cannot reach it, and each has a design remedy — randomise, bound, run a sensitivity analysis. The constitutive objection is **semantic**: the description “same woman, but a man, everything else fixed” fails to pick out a possible state of affairs, so there is no quantity that identification is failing to reach. That is why no design answers it, and why the authors call the isolation project theoretically misguided rather than underpowered.

In place of causal diagrams they propose *constitutive* diagrams, which display how features mutually define one another. Note what that implies for this chapter: the tool itself — the DAG, with its cuttable arrows — encodes the assumption at issue. Drawing one is already taking a side.

16.4.4. The perception move does not escape it

Section 16.2.4 resolved the bad-control problem by re-dating the treatment to perceived sex at the hiring decision, then observed that identification requires randomising the signal. Hu and Kohler-Hausmann (2024) attack precisely that estimand, by asking: a perception of sex or race is a perception *of what, exactly?*

Their answer is that the perception of the protected trait and the perception of other decision-relevant features are **co-constituted**. Their example: present two arrestees with “identical” prior-arrest records. They have not thereby been made equivalent, because the same record carries different meaning depending on whom it describes — an arrestee racialised as Black with *N* priors may be read as having an *unexpectedly low* number relative to imputed risk, while a white arrestee with the same *N* reads as ordinary. Making candidates similar in one respect necessarily makes them different in others.

Transposed to a résumé audit: “identical file, different perceived sex” again holds fixed something — the file’s *significance* — that co-varies with the thing being varied. The files are identical in ink and

non-identical in meaning, and meaning is what the employer responds to. **Which similarity is held fixed is itself the contested question**, not a technical parameter of the design.

Their conclusion is narrower than “audit studies are worthless,” and more damaging for being narrower: they reject the claim that such studies measure perception effects free of confounding, and the claim that a perceived-sex effect simply *is* disparate treatment. The second equation needs an unstated **normative** premise about what treatment ought to look like. The charge is that researchers do normative theorising behind technical machinery, and the demand is that both the sociological theory and the normative one be defended out loud instead of being settled by a covariate list.

16.4.5. Three replies

The argument is live, and the replies are substantive.

- **The signal is modular even if the category is not.** One really can change the name on a résumé holding the rest of the text fixed — a physical operation needing no metaphysics. That the reader’s interpretation is holistic is arguably part of the effect measured, not a defect in the estimand. This is the strongest reply and it protects correspondence studies specifically.
- **Comparison is not optional for law.** Discrimination doctrine is comparative by construction; if no comparison is well-posed the legal concept cannot operate. The authors would do the normative work openly rather than abandon comparison, but this leaves the empirical programme without operational guidance.
- **Constitution is scale-dependent.** Over a lifetime, gender and human capital are mutually constituting and the objection bites hard. At the moment of a single hiring decision with the file fixed, the manipulation is genuinely modular and it bites least. Which suggests the defensible middle: the critique is devastating for lifetime-scale “effect of gender” claims and comparatively weak against tightly scoped decision-stage experiments.

16.4.6. What survives, in this chapter’s own numbers

None of this licenses the conclusion that the gender wage gap cannot be studied, and its authors do not draw it. What is ill-defined is one specific summary statistic, not the subject.

Safe, and unaffected by the objection:

- **“Women in this sample earn 3.8% less than men.”** A description of an observed population — and the statistic UK and EU pay-gap reporting mandate, unadjusted, precisely because adjusted numbers launder the structural component.
- **“Within education, experience and region cells the gap is 11.3%; adding occupation and industry it is 7.5%.”** Within-cell contrasts, correctly labelled as such.
- **“The gap is near zero at the 10th percentile and about 7% at the 90th.”** A distributional description.
- **“A confounder explaining roughly 3% of the residual variance of gender and 4% of wages would account for the entire 11.3% to 7.5% attenuation.”** A sensitivity statement about a stated model.

Ill-defined: “**the effect of being a woman, net of everything legitimate**” — a single number separating discrimination from productivity across a lifetime.

Two independent routes condemn that one number and nothing else on the list. The identification route, worked through earlier in this chapter, shows it is unreachable from these data for three separate reasons: mediators everywhere, a collider at every rung, and a positivity failure at the only valid adjustment set. The constitutive route holds it was never a coherent target. Since both converge, dropping it costs nothing that was ever in hand.

16.4.7. The operational rule

Report what was **manipulated** or **observed**, in the **population it was observed in**. Then make the normative claim as a stated argument, rather than as something a coefficient is supposed to deliver.

Which is what the ladder in this chapter does. -3.8% is the observed total gap; -11.3% and -7.5% are within-cell contrasts; none of them is discrimination, and now the chapter says so.

16.4.8. References for this section

- Hu, L. and I. Kohler-Hausmann (2020). What’s sex got to do with fair machine learning? *ACM FAccT*. arXiv:2006.01770.
- Hu, L. and I. Kohler-Hausmann (2024). What is perceived when race is perceived and why it matters for causal inference and discrimination studies. *Law & Society Review* (open access).
- Hu, L. (2023). What is “race” in algorithmic discrimination on the basis of race? *Journal of Moral Philosophy*.
- Hu, L. (2025). Sex discrimination, normativity, and begging the causal question. *Political Philosophy*.
- Kohler-Hausmann, I. (2019). Eddie Murphy and the dangers of counterfactual causal thinking about detecting racial discrimination. *Northwestern University Law Review*.

Part IV.

Panel Data & DiD

17. Causal Forest in panel data

17.1. Introduction

In this simulation exercise, we use Causal Forest (Now is implemented in Generalized Random Forest) (<https://github.com/swager/grf>) to calculate conditional average treatment effect (or heterogeneous treatment effect). We assume three different data generating processes. The first one is a linear interaction between a variable of interest and the treatment dummy. The second one assumes a nonlinear function (a step function) of a variable of interest, say X , and the treatment dummy W . The third one is also nonlinear, assuming we have variable X , but the real DGP is the interaction of log of X and W .

17.2. Linear effect

Suppose we have a panel of firms over time, treatment assignment is random. We split the data into training and test. Then we generate 10 random variables, the first two are highly correlated. Then we collapse $V2$ into firm means, and use that as the unit effect. That is, we have a DGP as:

$$y_{it} = \alpha_i + V1_{it} + W_{it} + V1_{it} * W_{it} + \epsilon_{it} \quad (17.1)$$

Here α_i is not observed, and it is correlated with $V1$, so leaving it out is in principle a source of bias.

It is worth being precise about how much bias, though, because the results below are less dramatic than that setup suggests. The unit effect is the *firm mean* of $V2$ over all $t = 10$ periods, and averaging attenuates its correlation with contemporaneous $V1$ by a factor of t : with $\text{Cor}(V1_{it}, V2_{it}) = 0.9025$ we get $\text{Cov}(V1_{it}, \bar{V}2_i) = 0.9025/10$, and $\text{Cor}(V1_{it}, \alpha_i) \approx 0.29$. So the omitted-variable bias on $V1$ is only about 0.05–0.09, not a qualitative failure.

That is what the three models show: FE, RE and OLS all recover $V1 \approx 1$, $W \approx 1$ and $V1:W \approx 1.05$, essentially the truth. Note in particular that the *treatment* coefficients are unbiased in all three, because W is randomized and so independent of α_i — omitting a unit effect cannot bias the effect of a randomly assigned treatment. If you want a panel example where FE visibly beats OLS, the unit effect has to be correlated with $V1$ *contemporaneously* (or t has to be small), and the treatment has to be non-random.

```
# This is a simulation exercise to see how causal forest performs
# First we run a linear model, with panel setting
# Then two nonlinear situations, with panel setting.
```

17. Causal Forest in panel data

```
library(lfe, quietly = TRUE)
library(lme4)
require(snowfall)
library(MASS)
library(grf)
library(tidyverse)
set.seed(666)
rm(list = ls())
# t is no. of periods. p is no. of variables, n is no. of firms.
t <- 10
p <- 10
n=400
# generate p random variables
data1 = as.data.frame(matrix(rnorm(n*t*p), n*t, p))
names(data1) <- paste0("X", 1:p)

firm=seq(1,n)
time=seq(1,t)
data <- expand.grid(firm=firm,time=time)

# W is the treatment assignment
data$W <- rbinom(n*t, 1, 0.5)
data$train <- rbinom(n*t, 1, 0.5)

# generate two correlated variables
covarMat = matrix( c(1, .95^2, .95^2, 1 ) , nrow=2 , ncol=2 )
data.2 = as.data.frame(mvrnorm(n=n*t , mu=rep(0,2), Sigma=covarMat ))
names(data.2) <- c("V1", "V2")

data <- bind_cols(data,data.2) %>%
  tibble()

data <- bind_cols(data,data1)

# generate unit effect by firm means of V2, which is correlated with V1
data <- data %>%
  group_by(firm) %>%
  mutate(unit.effect=mean(V2)) %>%
  ungroup() %>%
  arrange(firm, time)

y<-c()
unit<-c()
X<-c()
k <- 0
```

```

for(i in 1:n){
  for(j in 1:t){
    k<-k+1
    unit[k]<-i
    y[k]<-1+ data$V1[k] + data$W[k] + data$V1[k]*data$W[k] + data$unit.effect[k]+rnorm(1,mu)
  }
}
data$unit=unit
data$y=y

data.train <- data %>%
  filter(train==1)
data.test <- data %>%
  filter(train==0)

# run fixed effect, random effect, and ols to see whether they are
# biased.
# in general we should expect fe performs better, since we are
# omitting unit.effect in the other two models
fe <- felm(y ~ V1 * W | unit, data=data.train)
re <- lmer(y~V1*W+(1 | unit), data=data.train)
ols <- lm(y ~ V1 * W , data=data.train)
summary(fe)

```

Call:

```
felm(formula = y ~ V1 * W | unit, data = data.train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-3.0440	-0.6252	-0.0083	0.6221	3.6174

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
V1	0.98354	0.03748	26.24	<2e-16 ***
W	1.02186	0.05093	20.06	<2e-16 ***
V1:W	1.04302	0.05237	19.92	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.018 on 1612 degrees of freedom

Multiple R-squared(full model): 0.7952 Adjusted R-squared: 0.7443

Multiple R-squared(proj model): 0.7269 Adjusted R-squared: 0.659

F-statistic(full model):15.61 on 401 and 1612 DF, p-value: < 2.2e-16

F-statistic(proj model): 1430 on 3 and 1612 DF, p-value: < 2.2e-16

17. Causal Forest in panel data

```
summary(re)
```

Linear mixed model fit by REML ['lmerMod']

Formula: $y \sim V1 * W + (1 | \text{unit})$

Data: data.train

REML criterion at convergence: 5938.9

Scaled residuals:

	Min	1Q	Median	3Q	Max
	-3.4332	-0.6615	-0.0049	0.6640	3.8329

Random effects:

Groups	Name	Variance	Std.Dev.
unit	(Intercept)	0.07166	0.2677
Residual		1.04602	1.0228

Number of obs: 2014, groups: unit, 399

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	1.00010	0.03539	28.26
V1	1.02961	0.03486	29.53
W	1.00333	0.04684	21.42
V1:W	1.05963	0.04814	22.01

Correlation of Fixed Effects:

	(Intr)	V1	W
V1	0.013		
W	-0.656	-0.008	
V1:W	-0.007	-0.727	0.025

```
summary(ols)
```

Call:

```
lm(formula = y ~ V1 * W, data = data.train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-3.8384	-0.7101	0.0081	0.7334	4.0053

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.00363	0.03319	30.24	<2e-16 ***
V1	1.04612	0.03516	29.75	<2e-16 ***
W	0.99657	0.04712	21.15	<2e-16 ***

```
V1:W      1.06458    0.04840    22.00    <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 1.057 on 2010 degrees of freedom
Multiple R-squared:  0.7249,    Adjusted R-squared:  0.7245
F-statistic: 1765 on 3 and 2010 DF,  p-value: < 2.2e-16
```

Now we run Causal Forest on this data. Two clarifications about what is and is not at stake here. First, treatment W is **randomized** (`rbinom(., 0.5)`, independent of the firm effect), so the omitted unit effect α_i does *not* confound the treatment effect — it only enters the nuisance functions (and hence precision), it does not create treatment-effect bias. The panel issue here is therefore mainly *dependence and nuisance fit*, not confounding. Second, the firm structure should be handled in two ways: we include the firm dummies in X so the forest can absorb α_i in the nuisance models, **and** we pass `clusters = firm` so that `grf`'s honest sample splitting and its variance estimates keep each firm's repeated observations together rather than treating them as independent rows. Including the dummies alone does neither of those.

```
# ensure consistent factor levels across train/test
all_firms <- levels(factor(data$firm))
data.train$firm <- factor(data.train$firm, levels = all_firms)
data.test$firm <- factor(data.test$firm, levels = all_firms)

# a is a data set of firm dummies
# Caution: this one-hot-encodes all ~400 firms into individual dummy columns
# of X. Random forests split poorly on such a high-cardinality sparse block
# of binary columns (each dummy is only informative for a tiny fraction of
# observations). A fixed-effects forest (demean y, W, and the continuous
# covariates within firm first, then run causal_forest on the demeaned
# variables) or a method that handles unit effects natively would typically
# perform better than including firm as raw dummies.
a <- as.data.frame(model.matrix(y~firm, data.train))
X <- as.matrix(bind_cols(data.train[,c(5,7:16)],a))
Y <- data.train$y
W <- data.train$W
# causal forest is run on y against V1, to X10 and firm dummies.
# clusters = firm makes honest splitting and the variance estimates respect
# the panel structure (a firm's repeated rows are kept together, not treated
# as independent observations).
tau.forest <- causal_forest(X, Y, W, clusters = as.integer(data.train$firm))

average_treatment_effect(tau.forest, target.sample = "all")
```


```
estimate    std.err
0.93560912 0.08227266
```

17. Causal Forest in panel data

```
average_treatment_effect(tau.forest, target.sample = "treated")
```

```
      estimate      std.err  
0.93168066 0.08158925
```


```
a <- as.data.frame(model.matrix(y~firm, data.test))  
X.test <- as.matrix(bind_cols(data.test[,c(5,7:16)],a))  
Y.test <- data.test$y  
W.test <- data.test$W  
  
tau.hat = predict(tau.forest, X.test, estimate.variance = TRUE)  
sigma.hat = sqrt(tau.hat$variance.estimates)  
  
data.test$pred <- tau.hat$predictions  
ggplot(data.test, aes(x=V1, y=pred)) + geom_point() + geom_abline(intercept = 1, slope = 1)
```




This graph is the predicted value of CATE on the test data, vs. the value of V1.

Now we graph the predicted CATE against the true CATE. The other nice feature of Causal Forest is that it returns variance for CATE for each observation. We also graph the error bar here.

```
graph.data <- as.data.frame(tau.hat )  
graph.data <- graph.data %>%  
  mutate(cate=1+data.test$V1) %>%  
  mutate(lower=predictions - sigma.hat , upper=predictions + sigma.hat )  
  
ggplot(graph.data , aes(x=cate, y=predictions)) + geom_point() + geom_abline(slope = 1)
```

```
ggplot() + geom_errorbar(graph.data , mapping=aes(x=cate, ymin=lower,ymax=upper), alpha = 0.2,
```



17.3. Nonlinear effect case 1

In the first situation, we assume DGP:

17. Causal Forest in panel data

$$y_{it} = 1 + \alpha_i + \max(V1_{it}, 0) * W_{it} + \epsilon_{it} \quad (17.2)$$

```
# if we have a nonlinear effect
# suppose we have the same X, but y is generated nonlinearly.
# in this example, y is a nonlinear function of V1 and W
set.seed(666)
rm(list = ls())
# t is no. of periods. p is no. of variables, n is no. of firms.
t <- 10
p <- 10
n=400
# generate p random variables
data1 = as.data.frame(matrix(rnorm(n*t*p), n*t, p))
names(data1) <- paste0("X", 1:p)

firm=seq(1,n)
time=seq(1,t)
data <- expand.grid(firm=firm,time=time)

# W is the treatment assignment
data$W <- rbinom(n*t, 1, 0.5)
data$train <- rbinom(n*t, 1, 0.5)

# generate two correlated variables
covarMat = matrix( c(1, .95^2, .95^2, 1 ) , nrow=2 , ncol=2 )
data.2 = as.data.frame(mvrnorm(n=n*t , mu=rep(0,2), Sigma=covarMat ))
names(data.2) <- c("V1", "V2")

data <- bind_cols(data,data.2) %>%
  tibble()

data <- bind_cols(data,data1)

# generate unit effect by firm means of V2, which is correlated with V1
data <- data %>%
  group_by(firm) %>%
  mutate(unit.effect=mean(V2)) %>%
  ungroup() %>%
  arrange(firm, time)

y<-c()
unit<-c()
X<-c()
k <- 0
for(i in 1:n){
```

```

    for(j in 1:t){
      k<-k+1
      unit[k]<-i
      y[k]<-1+ pmax(data$V1[k],0)*data$W[k] + data$unit.effect[k]+rnorm(1,mean=0,sd=1)
    }
  }
data$y=y
data$unit <- unit

data.train <- data %>%
  filter(train==1)
data.test <- data %>%
  filter(train==0)

# run fixed effect, random effect, and ols to see whether they are
# biased.
# they all perform poorly
fe <- felm(y ~ V1 * W | unit, data=data.train)
re <- lmer(y~V1*W+(1 | unit), data=data.train)
ols <- lm(y ~ V1 * W , data=data.train)
summary(fe)

```

Call:

```
felm(formula = y ~ V1 * W | unit, data = data.train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.9461	-0.6225	-0.0178	0.6245	3.3537

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
V1	-0.02076	0.03815	-0.544	0.586
W	0.42527	0.05184	8.203	4.73e-16 ***
V1:W	0.53367	0.05331	10.012	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.036 on 1612 degrees of freedom

Multiple R-squared(full model): 0.3744 Adjusted R-squared: 0.2187

Multiple R-squared(proj model): 0.1393 Adjusted R-squared: -0.07481

F-statistic(full model):2.405 on 401 and 1612 DF, p-value: < 2.2e-16

F-statistic(proj model): 86.96 on 3 and 1612 DF, p-value: < 2.2e-16

```
summary(re)
```

Linear mixed model fit by REML ['lmerMod']

17. Causal Forest in panel data

Formula: $y \sim V1 * W + (1 \mid \text{unit})$

Data: data.train

REML criterion at convergence: 6017.2

Scaled residuals:

Min	1Q	Median	3Q	Max
-3.5905	-0.6596	-0.0031	0.6717	3.4143

Random effects:

Groups	Name	Variance	Std.Dev.
unit	(Intercept)	0.07977	0.2824
Residual		1.08359	1.0410

Number of obs: 2014, groups: unit, 399

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	0.99976	0.03624	27.587
V1	0.02797	0.03553	0.787
W	0.40652	0.04774	8.515
V1:W	0.54484	0.04907	11.104

Correlation of Fixed Effects:

	(Intr)	V1	W
V1	0.013		
W	-0.653	-0.008	
V1:W	-0.007	-0.727	0.025

```
summary(ols)
```

Call:

```
lm(formula = y ~ V1 * W, data = data.train)
```

Residuals:

Min	1Q	Median	3Q	Max
-4.0965	-0.7415	0.0095	0.7512	3.8413

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.00363	0.03386	29.643	<2e-16 ***
V1	0.04612	0.03587	1.286	0.199
W	0.39928	0.04807	8.307	<2e-16 ***
V1:W	0.54938	0.04937	11.127	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.078 on 2010 degrees of freedom
 Multiple R-squared: 0.1556, Adjusted R-squared: 0.1543
 F-statistic: 123.5 on 3 and 2010 DF, p-value: < 2.2e-16

These linear models miss the *shape* of the effect, but it overstates things to say they return nothing. Each recovers the best linear approximation to the true CATE $\max(V1, 0)$, which is $0.3989 + 0.5 V1$: the fitted W coefficients are 0.40–0.43 and the $V1:W$ interactions are 0.53–0.55. In particular the W coefficient is essentially the true ATE, $E[\max(V1, 0)] = \phi(0) = 0.3989$. What a linear specification cannot do is see the kink at $V1 = 0$ — and that is what the forest is for.

Now we use Causal Forest. The first graph is the predicted CATE against $V1$; the second graph is the predicted CATE against true CATE. The third one adds in error bars.

```
all_firms <- levels(factor(data$firm))
data.train$firm <- factor(data.train$firm, levels = all_firms)
data.test$firm <- factor(data.test$firm, levels = all_firms)

# a is a data set of firm dummies
a <- as.data.frame(model.matrix(y~firm,data.train))
X <- as.matrix(bind_cols(data.train[,c(5,7:16)],a))
Y <- data.train$y
W <- data.train$W

#tau.forest = causal_forest(X, Y, W, num.trees = 4000)
a <- as.data.frame(model.matrix(y~firm,data.test))
X.test <- as.matrix(bind_cols(data.test[,c(5,7:16)],a))
Y.test <- data.test$y
W.test <- data.test$W

tau.forest3 = causal_forest(X, Y, W, clusters = as.integer(data.train$firm))

average_treatment_effect(tau.forest3, target.sample = "all")
```

```
      estimate      std.err
0.37320336 0.05326215
```

```
average_treatment_effect(tau.forest3, target.sample = "treated")
```


```
      estimate      std.err
0.37223789 0.05285912
```

```
#tau.forest3 = causal_forest(X, Y, W, num.trees = 4000)
tau.hat3 = predict(tau.forest3, X.test, estimate.variance = TRUE)
sigma.hat3 = sqrt(tau.hat3$variance.estimates)

data.test$pred3 <- tau.hat3$predictions
```


17. Causal Forest in panel data

```
seg1 <- data.frame(x1 = -3, x2 = 0, y1 = 0, y2 = 0)
seg2 <- data.frame(x1 = 0, x2 = 2, y1 = 0, y2 = 2)
ggplot(data.test, aes(x=V1, y=pred3)) + geom_point() +
  geom_segment(aes(x = x1, y = y1, xend = x2, yend = y2, colour = "segment"), data = seg1)+
  geom_segment(aes(x = x1, y = y1, xend = x2, yend = y2, colour = "segment"), data = seg2)
```




```
graph.data3 <- as.data.frame(tau.hat3)
graph.data3 <- graph.data3 %>%
  mutate(cate=pmax(data.test$V1,0)) %>%
  mutate(lower=predictions - sigma.hat3, upper=predictions + sigma.hat3)

ggplot(graph.data3, aes(x=cate, y=predictions)) + geom_point() + geom_abline(slope = 1)
```



```
ggplot() + geom_errorbar(graph.data3, mapping=aes(x=cate, ymin=lower,ymax=upper), alpha = 0.2,
```



17.4. Nonlinear effect case 2

In the second situation, we assume DGP:

$$y_{it} = 1 + \alpha_i + \log(V1_{it}^2) * W_{it} + \min(X3_{it}, 0) + \epsilon_{it} \quad (17.3)$$

But we only observe $V1$, not the log of the squared $V1$. The extra $\min(X3_{it}, 0)$ term is a nonlinearity in the *outcome* rather than in the treatment effect, so it leaves the CATE equal to $\log(V1^2)$; it is there to give the forest's nuisance models something nonlinear to fit.

```
# if we have a nonlinear effect
# suppose we have the same X, but y is generated nonlinearly.
# another example: we observe V1, but the true treatment effect is log(V1^2)
set.seed(666)
rm(list = ls())
# t is no. of periods. p is no. of variables, n is no. of firms.
t <- 10
p <- 10
n=400
# generate p random variables
data1 = as.data.frame(matrix(rnorm(n*t*p), n*t, p))
names(data1) <- paste0("X", 1:p)

firm=seq(1,n)
time=seq(1,t)
data <- expand.grid(firm=firm,time=time)

# W is the treatment assignment
data$W <- rbinom(n*t, 1, 0.5)
data$train <- rbinom(n*t, 1, 0.5)

# generate two correlated variables
covarMat = matrix( c(1, .95^2, .95^2, 1 ) , nrow=2 , ncol=2 )
data.2 = as.data.frame(mvrnorm(n=n*t , mu=rep(0,2), Sigma=covarMat ))
names(data.2) <- c("V1", "V2")

data <- bind_cols(data,data.2) %>%
  tibble()

data <- bind_cols(data,data1)

# generate unit effect by firm means of V2, which is correlated with V1
data <- data %>%
  group_by(firm) %>%
  mutate(unit.effect=mean(V2)) %>%
  ungroup() %>%
  arrange(firm, time)

y<-c()
```



```

unit<-c()
X<-c()
k <- 0
for(i in 1:n){
  for(j in 1:t){
    k<-k+1
    unit[k]<-i
    y[k]<-1+ log(data$V1[k]^2)*data$W[k] + pmin(data$X3[k], 0) + data$unit.effect[k]+rnorm
  }
}
data$y=y
data$unit <- unit

data.train <- data %>%
  filter(train==1)
data.test <- data %>%
  filter(train==0)

# run fixed effect, random effect, and ols to see whether they are
# biased.
# they all perform poorly
fe <- felm(y ~ V1 * W | unit, data=data.train)
re <- lmer(y~V1*W+(1 | unit), data=data.train)
ols <- lm(y ~ V1 * W , data=data.train)
summary(fe)

```

Call:

```
felm(formula = y ~ V1 * W | unit, data = data.train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-10.0383	-0.9478	0.0907	1.0865	5.3884

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
V1	-0.01979	0.07088	-0.279	0.780
W	-1.17812	0.09632	-12.231	<2e-16 ***
V1:W	0.01450	0.09904	0.146	0.884

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.925 on 1612 degrees of freedom

Multiple R-squared(full model): 0.2876 Adjusted R-squared: 0.1104

Multiple R-squared(proj model): 0.08495 Adjusted R-squared: -0.1427

F-statistic(full model):1.623 on 401 and 1612 DF, p-value: 6.032e-11

F-statistic(proj model): 49.88 on 3 and 1612 DF, p-value: < 2.2e-16

17. Causal Forest in panel data

```
summary(re)
```

Linear mixed model fit by REML ['lmerMod']

Formula: $y \sim V1 * W + (1 | \text{unit})$

Data: data.train

REML criterion at convergence: 8421.8

Scaled residuals:

Min	1Q	Median	3Q	Max
-6.1184	-0.4977	0.0647	0.6098	2.8959

Random effects:

Groups	Name	Variance	Std.Dev.
unit	(Intercept)	0.09312	0.3052
Residual		3.72467	1.9299

Number of obs: 2014, groups: unit, 399

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	0.57512	0.06301	9.128
V1	0.06152	0.06490	0.948
W	-1.18310	0.08706	-13.590
V1:W	-0.01892	0.08945	-0.212

Correlation of Fixed Effects:

	(Intr)	V1	W
V1	0.015		
W	-0.685	-0.010	
V1:W	-0.009	-0.727	0.028

```
summary(ols)
```

Call:

```
lm(formula =  $y \sim V1 * W$ , data = data.train)
```

Residuals:

Min	1Q	Median	3Q	Max
-12.1215	-0.9805	0.1390	1.1999	5.7125

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.57618	0.06136	9.390	<2e-16 ***
V1	0.07086	0.06502	1.090	0.276
W	-1.18331	0.08712	-13.582	<2e-16 ***

```
V1:W      -0.02005    0.08949  -0.224    0.823
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 1.954 on 2010 degrees of freedom
Multiple R-squared:  0.08508,    Adjusted R-squared:  0.08372
F-statistic: 62.31 on 3 and 2010 DF,  p-value: < 2.2e-16
```

Here the linear models fail more completely, and instructively so. Because $\log(V1^2)$ is *even* in $V1$, its linear interaction with $V1$ is zero by construction: the fitted $V1:W$ terms come out at 0.015, -0.019 and -0.020 , i.e. no detectable heterogeneity at all, even though the true CATE varies enormously with $V1$. The W coefficients (≈ -1.18) still approximate the true ATE, $E[\log(V1^2)] = \psi(1/2) + \log 2 = -1.2704$. So a linear model gets the average right and the heterogeneity completely wrong.

Now we run Causal Forest. Again, first graph is predicted CATE against $V1$, the second one is the predicted CATE against true CATE, the third one adds error bars.

```
all_firms <- levels(factor(data$firm))
data.train$firm <- factor(data.train$firm, levels = all_firms)
data.test$firm <- factor(data.test$firm, levels = all_firms)

# a is a data set of firm dummies
a <- as.data.frame(model.matrix(y~firm,data.train))
X <- as.matrix(bind_cols(data.train[,c(5,7:16)],a))
Y <- data.train$y
W <- data.train$W

#tau.forest = causal_forest(X, Y, W, num.trees = 4000)
a <- as.data.frame(model.matrix(y~firm,data.test))
X.test <- as.matrix(bind_cols(data.test[,c(5,7:16)],a))
Y.test <- data.test$y
W.test <- data.test$W

tau.forest4 = causal_forest(X, Y, W, clusters = as.integer(data.train$firm))

average_treatment_effect(tau.forest4, target.sample = "all")
```

```
      estimate      std.err
-1.20491164  0.08132368
```

```
average_treatment_effect(tau.forest4, target.sample = "treated")
```


```
      estimate      std.err
-1.20241393  0.07981497
```

17. Causal Forest in panel data

```
#tau.forest3 = causal_forest(X, Y, W, num.trees = 4000)
tau.hat4 = predict(tau.forest4, X.test, estimate.variance = TRUE)
sigma.hat4 = sqrt(tau.hat4$variance.estimates)


data.test$pred4 <- tau.hat4$predictions

fun.1 <- function(x) log(x^2)
ggplot(data.test, aes(x=V1, y=pred4)) + geom_point() +
stat_function(fun = fun.1, colour="red") + xlim(-3,3)
```




```
graph.data4 <- as.data.frame(tau.hat4)
graph.data4 <- graph.data4 %>%
  mutate(cate=log(data.test$V1^2)) %>%
  mutate(lower=predictions - sigma.hat4, upper=predictions + sigma.hat4)

ggplot(graph.data4, aes(x=cate, y=predictions)) + geom_point() + geom_abline(slope = 1)
```



```
ggplot() + geom_errorbar(graph.data4, mapping=aes(x=cate, ymin=lower,ymax=upper), alpha = 0.2)
```



17.5. Summary

We see in these examples, Causal Forest performs fairly well, while other linear models will not work in the nonlinear DGP situations.

17. Causal Forest in panel data

Systematic treatment: *R* · *Julia*.

18. Synthetic Control in R and Stata

18.1. Synthetic control

Abadie et al. have a few papers on synthetic control — Abadie, Diamond and Hainmueller (2010, *JASA*) and (2015, *AJPS*), and Abadie’s (2021, *JEL*) survey. The key idea is for a case study situation. A single unit, say a state/firm/country is exposed to a shock, and we are interested in the effect of that shock. For example, Card and Krueger (1994) studied the effect of a minimum wage increase in New Jersey. They used diff-in-diff, with Pennsylvania as the control state. The point of synthetic control is that maybe parallel trend assumption does not hold for the control unit. Is Pennsylvania a good control for New Jersey? Does it follow the same trend as New Jersey? Maybe not. Synthetic control is to generate a more comparable control with some combination of multiple control units, which can have a much closer parallel trend as the treated unit, than using any of the control units.

Now, how to construct this synthetic control unit? Basically it is a weighted average of multiple control units. There are two sets of weights to be determined. One set for the control units, one set to determine the importance of predictors. Abadie et al (2010) propose to choose a set of w_i ’s so that the resulting synthetic control best resembles the pre-intervention values for the treated unit of predictors of the outcome variable, subject to the constraints that the weights are non-negative and sum to 1. Another set of weights v_h for predictor importance is determined such that it minimizes the mean squared prediction error (MSPE) of this synthetic control with respect to the outcome in the pre-treatment period. The intuition of these weights are to pick w_i ’s to make the weighted sum of each predictor as close to the treated unit’s value of predictor as close as possible, given a set of v_h ’s. Then v_h ’s are determined by how important each predictor is to predict the outcome. We don’t use any information post-treatment to determine the weights. But we can use pre-treatment outcome to see the relative importance of each predictor in terms of predictor outcome. We can do this by out-of-sample validation. Basically divide the pre-treatment period into training and validation periods (that’s one reason that the pre-treatment period cannot be too short). Iterate the process of choosing w_i and v_h to achieve a minimization of MSPE in the validation period.

18.2. inference

Since synthetic control is for case study situation, we have essentially one treated unit and one synthetic control. The usual inference method would not work. Abadie et al suggested a permutation based approach. The basic idea is to permute the label of treatment; in other words, suppose one control unit is now a treated unit. Then follow the same procedure to get the synthetic control for that “treated unit”; look at the effect. Now if the treatment is truly for the treated unit only, then there would not be any effect for all the control unit. Therefore we can have a distribution of treatment effects after permutation of treatment label.

18.3. an example

The canonical synthetic control is implemented in both R and Stata (the `Synth` package, [CRAN](#); see Abadie, Diamond and Hainmueller's [JSS paper](#) for the software description). We use a package called “tidysynth” here to study one of the original examples, which evaluates the impact of Proposition 99 on cigarette consumption in California.

```
library(tidysynth)
library(dplyr)
library(tidyverse)
library(ggplot2)
library(tidyr)
data("smoking")
smoking %>% dplyr::glimpse()
```

Rows: 1,209

Columns: 7

```
$ state      <chr> "Rhode Island", "Tennessee", "Indiana", "Nevada", "Louisiana~
$ year       <dbl> 1970, 1970, 1970, 1970, 1970, 1970, 1970, 1970, 1970, 1970, ~
$ cigsale    <dbl> 123.9, 99.8, 134.6, 189.5, 115.9, 108.4, 265.7, 93.8, 100.3, ~
$ lnincome   <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ~
$ beer       <dbl> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ~
$ age15to24  <dbl> 0.1831579, 0.1780438, 0.1765159, 0.1615542, 0.1851852, 0.175~
$ retprice   <dbl> 39.3, 39.9, 30.6, 38.9, 34.3, 38.4, 31.4, 37.3, 36.7, 28.8, ~
```

```
smoking_out <-
  smoking %>%
  # initial the synthetic control object
  synthetic_control(outcome = cigsale, # outcome
                    unit = state, # unit index in the panel data
                    time = year, # time index in the panel data
                    i_unit = "California", # unit where the intervention occurred
                    i_time = 1988, # time period when the intervention occurred
                    generate_placebos=T # generate placebo synthetic controls (for inference)
  ) %>%

  # Generate the aggregate predictors used to fit the weights
  # average log income, retail price of cigarettes, and proportion of the
  # population between 15 and 24 years of age from 1980 - 1988
  generate_predictor(time_window = 1980:1988,
                    ln_income = mean(lnincome, na.rm = T),
                    ret_price = mean(retprice, na.rm = T),
                    youth = mean(age15to24, na.rm = T)) %>%

  # average beer consumption in the donor pool from 1984 - 1988
  generate_predictor(time_window = 1984:1988,
                    beer_sales = mean(beer, na.rm = T)) %>%

  # Lagged cigarette sales
```




```

generate_predictor(time_window = 1975,
                  cigsale_1975 = cigsale) %>%
generate_predictor(time_window = 1980,
                  cigsale_1980 = cigsale) %>%
generate_predictor(time_window = 1988,
                  cigsale_1988 = cigsale) %>%
  # Generate the fitted weights for the synthetic control
generate_weights(optimization_window = 1970:1988, # time to use in the optimization task
                margin_ipop = .02, sigf_ipop = 7, bound_ipop = 6 # optimizer options
) %>%
  # Generate the synthetic control
generate_control()

smoking_out %>% plot_trends()

```




Dashed line denotes the time of the intervention.

```


smoking_out %>% plot_differences()

```

18. Synthetic Control in R and Stata



```
smoking_out %>% plot_weights()
```




```
smoking_out %>% grab_balance_table()
```

```
# A tibble: 7 x 4
  variable      California synthetic_California donor_sample
  <chr>          <dbl>          <dbl>          <dbl>
1 cigsale_1975  0.00000000    0.00000000    0.00000000
2 cigsale_1980  0.00000000    0.00000000    0.00000000
3 ret_price     0.00000000    0.00000000    0.00000000
4 cigsale_1988  0.00000000    0.00000000    0.00000000
5 beer_sales    0.00000000    0.00000000    0.00000000
6 youth         0.00000000    0.00000000    0.00000000
7 ln_income     0.00000000    0.00000000    0.00000000
```

1	ln_income	10.1	9.85	9.83
2	ret_price	89.4	89.4	87.3
3	youth	0.174	0.174	0.173
4	beer_sales	24.3	24.2	23.7
5	cigsale_1975	127.	127.	137.
6	cigsale_1980	120.	120.	138.
7	cigsale_1988	90.1	91.4	114.


```
smoking_out %>% plot_placebos()
```




prune cases with a pre-period RMSPE exceeding two times the treated unit's pre-period RMSPE.

```
smoking_out %>% plot_placebos(prune = FALSE)
```

18. Synthetic Control in R and Stata



```
smoking_out %>% plot_mspe_ratio()
```



```
smoking_out %>% grab_significance()
```

```
# A tibble: 39 x 8
```

```
  unit_name      type pre_mspe post_mspe mspe_ratio rank fishers_exact_pvalue
  <chr>         <chr>   <dbl>   <dbl>     <dbl> <int>         <dbl>
```

```

1 California    Trea~    3.17    392.    124.    1    0.0256
2 Georgia      Donor    3.79    179.    47.2    2    0.0513
3 Indiana      Donor   25.2    770.    30.6    3    0.0769
4 West Virginia Donor    9.52    284.    29.8    4    0.103
5 Wisconsin    Donor   11.1    268.    24.1    5    0.128
6 Missouri     Donor    3.03    67.8    22.4    6    0.154
7 Texas        Donor   14.4    277.    19.3    7    0.179
8 South Carolina Donor   12.6    234.    18.6    8    0.205
9 Virginia     Donor    9.81    96.4    9.83    9    0.231
10 Nebraska     Donor    6.30    52.9    8.40   10    0.256
# i 29 more rows
# i 1 more variable: z_score <dbl>

```

```
smoking_out %>% grab_synthetic_control()
```

```

# A tibble: 31 x 3
  time_unit real_y synth_y
  <dbl>    <dbl>    <dbl>
1   1970    123    117.
2   1971    121    119.
3   1972   124.   124.
4   1973   124.   125.
5   1974   127.   127.
6   1975   127.   127.
7   1976   128.   128.
8   1977   126.   126.
9   1978   126.   125.
10  1979   122.   123.
# i 21 more rows

```

```
smoking_out %>% grab_synthetic_control(placebo = T)
```

```

# A tibble: 1,209 x 5
  .id      .placebo time_unit real_y synth_y
  <chr>      <dbl>    <dbl>    <dbl>    <dbl>
1 California    0    1970    123    117.
2 California    0    1971    121    119.
3 California    0    1972   124.   124.
4 California    0    1973   124.   125.
5 California    0    1974   127.   127.
6 California    0    1975   127.   127.
7 California    0    1976   128.   128.
8 California    0    1977   126.   126.
9 California    0    1978   126.   125.
10 California    0    1979   122.   123.
# i 1,199 more rows

```

```
smoking_out %>%
  tidyr::unnest(cols = c(.outcome))
```

```
# A tibble: 1,482 x 50
  .id      .placebo .type   time_unit California Alabama Arkansas Colorado
  <chr>      <dbl> <chr>      <dbl>      <dbl>    <dbl>    <dbl>    <dbl>
1 California      0 treated    1970      123      NA      NA      NA
2 California      0 treated    1971      121      NA      NA      NA
3 California      0 treated    1972      124      NA      NA      NA
4 California      0 treated    1973      124      NA      NA      NA
5 California      0 treated    1974      127      NA      NA      NA
6 California      0 treated    1975      127      NA      NA      NA
7 California      0 treated    1976      128      NA      NA      NA
8 California      0 treated    1977      126      NA      NA      NA
9 California      0 treated    1978      126      NA      NA      NA
10 California      0 treated    1979      122      NA      NA      NA
# i 1,472 more rows
# i 42 more variables: Connecticut <dbl>, Delaware <dbl>, Georgia <dbl>,
#   Idaho <dbl>, Illinois <dbl>, Indiana <dbl>, Iowa <dbl>, Kansas <dbl>,
#   Kentucky <dbl>, Louisiana <dbl>, Maine <dbl>, Minnesota <dbl>,
#   Mississippi <dbl>, Missouri <dbl>, Montana <dbl>, Nebraska <dbl>,
#   Nevada <dbl>, `New Hampshire` <dbl>, `New Mexico` <dbl>,
#   `North Carolina` <dbl>, `North Dakota` <dbl>, Ohio <dbl>, ...
```


18.4. Generalized Synthetic Control

The idea of generalized SC is to combine both an interactive fixed effect model and synthetic control. It is basically a panel data synthetic control. It naturally takes multiple treated units and uses bootstrap to get standard errors. It generates ATT with standard errors.


Examples are here: <https://yiqingxu.org/packages/gsynth/>

```
## for processing C++ code
require(Rcpp)
## for plotting
require(ggplot2)
require(GGally)
## for parallel computing
require(foreach)
require(future)
require(doParallel)
require(abind)
require(lfe)
library(gsynth)
## ## Syntax has been updated since v.1.2.0.
## ## Comments and suggestions -> yiqingxu@stanford.edu.
```

```
data(gsynth)
library(panelview)
## ## See bit.ly/panelview4r for more info.
## ## Report bugs -> yiqingxu@stanford.edu.
panelview(Y ~ D, data = simdata, index = c("id", "time"), pre.post = TRUE)
```



```
panelview(Y ~ D, data = simdata, index = c("id", "time"), type = "outcome")
```



18. Synthetic Control in R and Stata

```
system.time(
  out <- gsynth(Y ~ D + X1 + X2, data = simdata, index = c("id","time"), force = "two-way", (
)
```

```
      user  system elapsed
1.679    0.000    1.682
```

```
print(out)
```

Call:

```
gsynth(formula = Y ~ D + X1 + X2, data = simdata, index = c("id",
  "time"), force = "two-way", r = c(0, 5), CV = TRUE, se = TRUE,
  nboots = 100, inference = "parametric", parallel = FALSE,
  vartype = "parametric")
```

Average Treatment Effect on the Treated:

	ATT.avg	S.E.	CI.lower	CI.upper	p.value
[1,]	5.084	0.3326	4.432	5.736	0

~ by Period (including Pre-treatment Periods):

	ATT	S.E.	CI.lower	CI.upper	p.value	count
-19	0.05868	0.6716	-1.25757	1.3749	9.304e-01	5
-18	-0.11741	0.4647	-1.02814	0.7933	8.005e-01	5
-17	-0.89029	0.5334	-1.93573	0.1551	9.510e-02	5
-16	1.29462	0.5791	0.15967	2.4296	2.537e-02	5
-15	-1.20350	0.5627	-2.30636	-0.1006	3.245e-02	5
-14	2.21992	0.7481	0.75369	3.6861	3.003e-03	5
-13	-0.18469	0.7557	-1.66574	1.2964	8.069e-01	5
-12	0.87413	0.4833	-0.07307	1.8213	7.049e-02	5
-11	-0.04824	0.6238	-1.27081	1.1743	9.384e-01	5
-10	0.84503	0.6792	-0.48623	2.1763	2.135e-01	5
-9	0.49754	0.6345	-0.74604	1.7411	4.329e-01	5
-8	-0.38001	0.5696	-1.49641	0.7364	5.047e-01	5
-7	-1.03493	0.6606	-2.32974	0.2599	1.172e-01	5
-6	-0.81532	0.6910	-2.16971	0.5391	2.381e-01	5
-5	-1.13329	0.6543	-2.41574	0.1492	8.327e-02	5
-4	-0.17733	0.5964	-1.34631	0.9917	7.662e-01	5
-3	0.17671	0.8271	-1.44445	1.7979	8.308e-01	5
-2	-0.14376	0.5958	-1.31145	1.0239	8.093e-01	5
-1	-1.19697	0.4708	-2.11966	-0.2743	1.100e-02	5
0	1.35910	0.7762	-0.16217	2.8804	7.994e-02	5
1	0.34558	0.5608	-0.75364	1.4448	5.378e-01	5
2	0.15709	0.8485	-1.50592	1.8201	8.531e-01	5
3	3.64127	0.5570	2.54966	4.7329	6.241e-11	5
4	2.74975	0.7010	1.37591	4.1236	8.750e-05	5
5	4.45569	0.8610	2.76823	6.1431	2.276e-07	5

6	5.67231	0.8205	4.06411	7.2805	4.745e-12	5
7	8.17039	0.6550	6.88661	9.4542	0.000e+00	5
8	6.57085	0.8954	4.81599	8.3257	2.154e-13	5
9	9.11624	0.5527	8.03299	10.1995	0.000e+00	5
10	9.96205	0.5275	8.92814	10.9960	0.000e+00	5

Coefficients for the Covariates:

	Coef	S.E.	CI.lower	CI.upper	p.value
X1	1.454	0.03993	1.376	1.533	0
X2	3.507	0.04104	3.427	3.588	0

out\$est.att

	ATT	S.E.	CI.lower	CI.upper	p.value	count
-19	0.05868370	0.6715682	-1.25756576	1.3749332	9.303670e-01	5
-18	-0.11741144	0.4646683	-1.02814458	0.7933217	8.005171e-01	5
-17	-0.89029173	0.5333967	-1.93573004	0.1551466	9.509782e-02	5
-16	1.29461792	0.5790639	0.15967354	2.4295623	2.537089e-02	5
-15	-1.20349647	0.5626948	-2.30635803	-0.1006349	3.245118e-02	5
-14	2.21991721	0.7480902	0.75368734	3.6861471	3.002851e-03	5
-13	-0.18468804	0.7556547	-1.66574411	1.2963680	8.069149e-01	5
-12	0.87413414	0.4832772	-0.07307177	1.8213401	7.048776e-02	5
-11	-0.04824208	0.6237714	-1.27081152	1.1743274	9.383536e-01	5
-10	0.84503247	0.6792288	-0.48623146	2.1762964	2.134606e-01	5
-9	0.49754000	0.6344913	-0.74604016	1.7411202	4.329488e-01	5
-8	-0.38001372	0.5696025	-1.49641402	0.7363866	5.046725e-01	5
-7	-1.03492639	0.6606315	-2.32974035	0.2598876	1.172149e-01	5
-6	-0.81531903	0.6910261	-2.16970525	0.5390672	2.380530e-01	5
-5	-1.13328625	0.6543233	-2.41573628	0.1491638	8.327401e-02	5
-4	-0.17733004	0.5964300	-1.34631145	0.9916514	7.662229e-01	5
-3	0.17670638	0.8271342	-1.44444692	1.7978597	8.308302e-01	5
-2	-0.14376360	0.5957716	-1.31145447	1.0239273	8.093175e-01	5
-1	-1.19696676	0.4707704	-2.11965977	-0.2742737	1.100405e-02	5
0	1.35910371	0.7761745	-0.16217040	2.8803778	7.994100e-02	5
1	0.34558474	0.5608399	-0.75364126	1.4448107	5.377682e-01	5
2	0.15709192	0.8484914	-1.50592063	1.8201045	8.531172e-01	5
3	3.64127124	0.5569529	2.54966368	4.7328788	6.241185e-11	5
4	2.74975180	0.7009533	1.37590853	4.1235951	8.749879e-05	5
5	4.45569195	0.8609635	2.76823443	6.1431495	2.276196e-07	5
6	5.67231032	0.8205269	4.06410723	7.2805134	4.744871e-12	5
7	8.17038502	0.6550005	6.88660758	9.4541625	0.000000e+00	5
8	6.57084917	0.8953552	4.81598521	8.3257131	2.153833e-13	5
9	9.11624492	0.5526911	8.03299024	10.1994996	0.000000e+00	5
10	9.96204576	0.5275143	8.92813677	10.9959548	0.000000e+00	5

18. Synthetic Control in R and Stata

```
out$est.avg
```

	ATT.avg	S.E.	CI.lower	CI.upper	p.value
[1,]	5.084123	0.3325502	4.432336	5.735909	0

```
out$est.beta
```


	Coef	S.E.	CI.lower	CI.upper	p.value
X1	1.454447	0.03992589	1.376194	1.532701	0
X2	3.507197	0.04104218	3.426756	3.587638	0

```
system.time(  
out <- gsynth(Y ~ D + X1 + X2, data = simdata, index = c("id","time"),  
  force = "two-way", CV = TRUE, r = c(0, 5), se = TRUE,  
  inference = "parametric", nboots = 100, parallel = TRUE, cores = 4)  
)
```

	user	system	elapsed
	7.655	0.344	17.941

```
# cumuEff was removed in newer gsynth versions (now a wrapper for fecht)  
# cumu1 <- cumuEff(out, cumu = TRUE, id = NULL, period = c(0,5))  
# cumu1$est.catt
```

```
plot(out)
```



Systematic treatment: *Julia*.

19. Same Data, Different Estimators: A Synthetic Control Comparison

19.1. One panel, four estimators

Here I use the Proposition 99 data to compare four panel estimators on the same outcome matrix: DiD, synthetic DiD, synthetic control, and TASC.

This is the standard California smoking example. There are 39 states, annual cigarette sales per capita from 1970 to 2000, and California is treated after the 1988 anti-smoking law. The useful thing about this example is that the estimators do not give exactly the same answer.

19.2. Four estimators, one panel

```
using CSV, DataFrames, Statistics
using SynthDiD
using TASC

df = CSV.read("california_prop99.csv", DataFrame)
# ... pivot to N×T matrix Y, N0 = 38 donors, T0 = 19 pre-treatment years
# (full setup code below)

tau_did = did_estimate(Y, N0, T0)
tau_sdid = synthdid_estimate(Y, N0, T0)
tau_sc = sc_estimate(Y, N0, T0)

# TASC: treated unit must be row 1
Y_tasc = vcat(Y[(N0+1):end, :], Y[1:N0, :])
model = fit_tasc(Y_tasc; d=2, T0=T0, n_em=200, tol=1e-3)
pred = predict_counterfactual(model, Y_tasc)
```

Estimator	ATT (packs per capita)
Difference-in-differences	−27.35
Synthetic control	−19.62
Synthetic DiD	−15.60
TASC	−17.11

(These are the values computed live in the companion Julia book’s [same analysis](#); the code in this chapter is `eval: false`, so that chapter is the authoritative source for the numbers. The DiD and SC figures also match the published `synthdid` benchmarks for this panel.)

The smallest and largest estimates differ by nearly twelve packs per capita. That is not a rounding issue. It comes from the way each estimator constructs the counterfactual.

19.3. What changes across estimators

DiD (-27.35) assigns equal weight to every control state and every pre-treatment year. California is compared to the average of all 38 donor states. In this data, that average sits well below California before treatment. So the large DiD estimate is not automatically evidence of a larger causal effect. It is partly a warning that equal weights do not fit the pre-period well.

Synthetic control (-19.62) replaces equal weights with donor weights chosen to match California’s 1970-1988 cigarette sales. Once the pre-period fit is better, the post-treatment gap shrinks. The move from about -27 to about -20 is mostly the price of using a better control group.

Synthetic DiD (-15.60) adds time weights. The estimator can put more weight on the pre-treatment years that are more useful for the post-1988 counterfactual. Here that moves the estimate a bit closer to zero.

TASC (-17.11) does something different. Instead of choosing only weights, it fits a state-space model to the pre-treatment panel:

$$x_t = A x_{t-1} + q_t, \quad q_t \sim \mathcal{N}(0, Q) \quad (19.1)$$

$$y_t = H x_t + r_t, \quad r_t \sim \mathcal{N}(0, R) \quad (19.2)$$

The matrix A describes how the latent factors move over time. After treatment, the filter continues using the donor panel, and the treated unit’s row of H maps the latent state to California’s counterfactual. The TASC estimate (-17.11) lands between SC and SDiD in this example.

19.4. The figure

```
# Counterfactual path for each method. omega_* / intercept_* are the donor
# weights and (for DiD/SDiD) intercepts extracted from the fitted objects
# above (e.g. tau_sdid.weights.omega for SynthDiD.jl) - this snippet is
# schematic and omits that extraction step for brevity; it is not meant to
# run as-is (hence eval: false).
cfact_did = omega_did' * Y[1:NO, :] .+ intercept_did # parallel trends
cfact_sdid = omega_sdid' * Y[1:NO, :] .+ intercept_sdid
cfact_sc = omega_sc' * Y[1:NO, :] # no intercept
cfact_tasc = vec(pred.target)
tasc_se = sqrt.(max.(vec(pred.variance), 0.0))
```




Figure 19.1.: Four counterfactual paths for California. Dashed lines (DiD, SDiD) use pre-period intercepts to enforce parallel trends; solid lines (SC, TASC) fit levels directly. The shaded band is the TASC 95% posterior interval $\hat{Y}_{1t}(0) \pm 1.96 \hat{\sigma}_t$.

Two things are visible in the figure that the table cannot show.

First, the pre-treatment paths differ. The DiD control average sits well below California before treatment. SC and TASC track California more closely. This is why the DiD post-treatment gap is larger.

Second, the TASC band is narrow before treatment and wider after 1989. That is what I would expect. Before treatment, the smoother can use information from both directions in time. After treatment, the filter has to move forward without California's observed outcomes, so the uncertainty accumulates through $P_{t+1} = AP_tA^\top + Q$.

19.5. What the comparison teaches

Comparison	What it isolates
DiD vs SC	Unit weights — equal vs optimised donor match
SC vs SDiD	Time weights — equal vs recency-weighted pre-period
SC vs TASC	Model — weight fitting vs latent state-space
SC/SDiD vs TASC	Uncertainty — point estimate vs posterior interval

Which estimator I would use depends on the setting. DiD is transparent, but the pre-period fit here is poor. SC improves the donor match. SDiD also weights time periods. TASC is useful when the

19. Same Data, Different Estimators: A Synthetic Control Comparison

outcome series has persistent dynamics and we want uncertainty for the counterfactual path, not just one average number.

For this panel, the three methods that try to match the pre-period give estimates between about -15.6 and -19.6. DiD is the outlier. I would read this as a problem with equal weights, not as evidence that the effect of Proposition 99 depends on the estimator.

19.6. Setup code

```
df      = CSV.read("california_prop99.csv", DataFrame)
states  = sort(unique(df.State))
years   = sort(unique(df.Year))
N, T    = length(states), length(years)

Y_wide  = zeros(N, T)
treated_vec = zeros(Bool, N)
for row in eachrow(df)
    i = findfirst==(row.State), states)
    t = findfirst==(row.Year), years)
    Y_wide[i, t] = row.PacksPerCapita
    treated_vec[i] |= (row.treated == 1)
end

ctrl_idx = findall(!treated_vec)
trt_idx  = findall(treated_vec)
Y = Y_wide[vcat(ctrl_idx, trt_idx), :] # donors first
N0 = length(ctrl_idx)                 # 38
T0 = findfirst==(1989), years) - 1   # 19
```

The SynthDiD.jl package provides `did_estimate`, `synthdid_estimate`, and `sc_estimate`. The TASC.jl package provides `fit_tasc` and `predict_counterfactual`. Both packages are available on GitHub.

Systematic treatment: [Julia](#).

20. Bartik Instrument and Rotemberg weights

20.1. Bartik Instrument

Bartik instrument is a method to estimate the causal effect of a treatment on an outcome. The method is developed by Bartik (1991). The method is also known as the shift share IV method.

We follow Goldsmith-Pinkham, Sorkin, and Swift (2021) (GPSS) to describe the method. Let's say we are interested in estimating the elasticity of employment on wage growth.

$$y_{lt} = D_{lt}\rho + \beta_0 x_{lt} + \epsilon_{lt} \quad (20.1)$$

y_l is the wage growth in location l , x_l is the employment growth rate. β_0 is the inverse elasticity of labor supply.

The usual challenge is that x_l is endogenous. There could be factors that are driving both wage growth and employment growth at location l . Note that I am omitting the time subscript for simplicity.

There are two accounting identities:

$$x_{lt} = Z_{lt}G_{lt} = \sum_{k=1}^K z_{lkt}g_{lkt} \quad (20.2)$$

and

$$g_{lkt} = g_{kt} + \tilde{g}_{lkt} \quad (20.3)$$

Suppose in location l there are K industries. Then the employment growth can be decomposed to the sum of industry growth g_{lkt} , with z 's as the industry shares.

The second identity is to say the this industry growth rate at this location can be decomposed to a national industry rate and a location-industry rate. The idea of Bartik instrument is to use the product of initial industry location share and industry growth rate as the instrument.

$$B_{lt} = Z_{l0}G_t = \sum_{k=1}^K z_{lk0}g_{kt} \quad (20.4)$$

The intuition is that the initial shares are exogenous, as it's a level variable. And the industry growth rate is exogenous, since it's national, not location related.

20.2. Bartik Instrument in special cases

20.2.1. Two industries and one period

GPSS goes through the simplest case: two industries and one period.

$$B_l = z_{l1}g_1 + z_{l2}g_2 \quad (20.5)$$

Since there are only two industries, $z_{l1} = 1 - z_{l2}$. Then we can rewrite the Bartik instrument as

$$B_l = g_2 + (g_1 - g_2)z_{l1} \quad (20.6)$$

From this simplest case, we see that the instrument has only z_{l1} , that is related to location. If the industry share can be considered exogenous, then the instrument is valid. We have to argue that the industry share is exogenous. In this case, the instrument is a linear function of the industry share. Therefore, the Bartik instrument is the same as using the industry share as an instrument.

GPSS shows that just-identified two-stage least squares using a Bartik instrument can in fact be written as GMM with an overidentified set of instruments, where the weight matrix is the outer product of national growth rates. That is, Bartik estimator is equivalent to a GMM estimator, with each location share as one instrument, and the weight matrix is the outer product of national growth rates.

With K industries and T time periods, it's basically using $K \times T$ instruments. Still, Bartik instrument is the same as GMM with all these instruments.

20.3. Rotemberg weights

GPSS decompose the Bartik estimator:

$$\beta_{Bartik} = \sum_k \hat{\alpha}_k \hat{\beta}_k \quad (20.7)$$

where

$$\hat{\beta}_k = (Z'_k X^\perp)^{-1} Z'_k Y^\perp \quad (20.8)$$

and

$$\hat{\alpha}_k = \frac{g_k Z'_k X^\perp}{\sum_{k'} g_{k'} Z'_{k'} X^\perp} \quad (20.9)$$

Here $X^\perp$ and $Y^\perp$ are the residualized X and Y , from the regression of X and Y on all other variables (controls).

Rotemberg weights provide a measure of how important each share is. Read the direction carefully: if industry k 's own just-identified estimate $\hat{\beta}_k$ is biased by 1 percent, the overall Bartik estimate inherits α_k percent of that bias. So a large α_k means the overall answer is *sensitive* to that industry, not that the industry is itself badly biased. This is what makes the weights useful: they tell you for which instruments the exogeneity assumption has to be defended hardest. (The formal statement is GPSS Corollary D.1, given below.)

20.4. Simulations

Simulations codes are adopted from Kohei Kawaguchi's lecture notes. I modified slightly to fit different scenarios.

Here we simulate a case with 4 industries and 4 periods. We have 1000 locations, 4 periods, and 4 industries. The true β (effect of x on y) is 1.

```
library(ggplot2)
library(tidyverse)
library(modelsummary)
library(kableExtra)
library(estimatr)
set.seed(6)
L <- 1000
T <- 4
K <- 4
beta <- 1
sigma <- c(1, 1, 1)

z_lk0 <- expand_grid(
  l = 1:L,
  k = 1:K
) |>
mutate(z_lk0 = runif(length(l)))

g_kt <- expand_grid(
  t = 1:T,
  k = 1:K
) |>
mutate(g_kt = rnorm(n(), 0, .1))

df <- expand_grid(
  l = 1:L,
  t = 1:T,
  k = 1:K
) |>
left_join(z_lk0, by = c("l", "k")) |>
left_join(g_kt, by = c("k", "t")) |>
mutate(z_lkt = z_lk0 + sigma[1] * runif(length(l)),
  g_lkt = g_kt + sigma[2] * rnorm(length(l), 0, .1)
)

df <- df %>%
  group_by(t, l) |>
  mutate(
    z_lk0 = z_lk0 / sum(z_lk0),
    z_lkt = (z_lkt / sum(z_lkt)),
```

20. Bartik Instrument and Rotemberg weights

```

    x_lt = sum(z_lkt * g_lkt)
  )

df

# A tibble: 16,000 x 8
# Groups:   t, l [4,000]
   l     t     k z_lk0    g_kt z_lkt    g_lkt    x_lt
<int> <int> <int> <dbl>    <dbl> <dbl>    <dbl>    <dbl>
1     1     1     1 0.277  0.130  0.252  0.115  0.0400
2     1     1     2 0.428  0.146  0.302  0.291  0.0400
3     1     1     3 0.121 -0.165  0.321 -0.231  0.0400
4     1     1     4 0.174 -0.110  0.126 -0.0185  0.0400
5     1     2     1 0.277 -0.0289  0.160 -0.00619  0.0402
6     1     2     2 0.428 -0.00130  0.356  0.0560  0.0402
7     1     2     3 0.121  0.0433  0.272 -0.0590  0.0402
8     1     2     4 0.174  0.0840  0.212  0.176  0.0402
9     1     3     1 0.277 -0.101  0.237 -0.0361 -0.0738
10    1     3     2 0.428 -0.152  0.367 -0.126 -0.0738
# i 15,990 more rows

```

This assumes we observe the industry shares at time 0 (z_{lk0}), and the industry growth rates (g_{kt}). We also observe the industry shares at time t (z_{lkt}). The employment growth rate at location l and time t (g_{lkt}) is a function of g_{kt} plus some noise. We can calculate the x_{lt} , which is the weighted sum of industry growth rates.

Next we generate the outcome at l, t level. Outcome (wage growth) is a linear function of x (employment growth) at l, t level.

```

df_lt <- df |>
  mutate(g_tilde_lkt = g_lkt - g_kt) |>
  group_by(l, t) |>
  summarise(across(c(x_lt, g_tilde_lkt), mean), .groups = "drop") |>
  ungroup() |>
  mutate(y_lt = beta * x_lt + sigma[3] * g_tilde_lkt)

```

If we run OLS on raw data, we'll get biased results.

```

result_ols <-
  df_lt |>
  lm(formula = y_lt ~ x_lt)

modelsummary(result_ols)

```

Now we generate the Bartik instrument. We calculate the Bartik instrument as the weighted sum of z_{lk0} and g_{kt} .

	(1)
(Intercept)	0.000
	(0.001)
x_lt	1.336
	(0.008)
Num.Obs.	4000
R2	0.884
R2 Adj.	0.884
AIC	−14 059.9
BIC	−14 041.0
Log.Lik.	7032.945
F	30 603.441
RMSE	0.04

```
b_lt <- df |>
  group_by(l, t) |>
  summarise(b_lt = sum(z_lk0 * g_kt)) |>
  ungroup()

df_lt <- df_lt |>
  left_join(b_lt, by = c("l", "t"))

result_first_stage <- lm(formula = x_lt ~ b_lt, data = df_lt)

modelsummary(result_first_stage)
```

```
result_tsls <- df_lt |>
  estimatr::iv_robust(formula = y_lt ~ x_lt | b_lt)

modelsummary::modelsummary(result_tsls)
```

20.4.1. Calculate Rotemberg weights

Let's do IV using Bartik instruments. No controls here. We use matrix operation to recover the IV estimates, which is the same as the `tsls` function in R.

```
B = as.matrix(cbind(1, df_lt$b_lt))
Y = as.matrix(df_lt$y_lt)
X = as.matrix(cbind(1, df_lt$x_lt))
```

	(1)
(Intercept)	−0.001 (0.001)
b_lt	0.919 (0.013)
Num.Obs.	4000
R2	0.573
R2 Adj.	0.573
AIC	−11 638.5
BIC	−11 619.6
Log.Lik.	5822.262
F	5361.487
RMSE	0.06

	(1)
(Intercept)	0.000 (0.001)
x_lt	0.990 (0.012)
Num.Obs.	4000
R2	0.825
R2 Adj.	0.825
AIC	−12 398.4
BIC	−12 379.5
RMSE	0.05

```

iv = round(solve(t(B)%*%X)%*%t(B)%*%Y,3)

## Label and organize results into a data frame
beta.hat = as.data.frame(cbind(c("Intercept","x"),iv))
names(beta.hat) = c("Coeff. ","Est")
beta.hat

```

```

      Coeff.  Est
1 Intercept    0
2          x 0.99

```

Now we calculate the Rotemberg weights.

```

# This code is the R version of Paul Goldsmith-Pinkham's R code for Rotemberg weights on
# github; which is actually a call from a C++ code by jjchern.
# bw.R calls ComputeAlphaBeta.cpp
# dimensions:  G is KTx1, Z is LTxKT, Y is LTx1, X is LTx1
# so B=ZG is LTx1
# Here I learned the dimensions from the ADH example in Paul's github.
# in ADH example, czone(L) is 722, year(T) is 2, ind(K) is 390.
# so G is 780x1, Z is 1444x780, Y is 1444x1, X is 1444x1
y_lt <- df_lt |>
  select(l,t,y_lt)

Y <- as.matrix(y_lt[,3])

x_lt <- df_lt |>
  select(l,t,x_lt)

X <- as.matrix(x_lt[,3])
#X <- as.matrix(cbind(1,x_lt[,3]))

# global is kt level data set
# # g_data <- df |>
# #   group_by(k,t) |>
# #   summarise(g_kt=sum(g_lkt*z_lkt))
# #
# G <- as.matrix(g_data[,3])

G <- as.matrix(g_kt[,3])

# generate wide Z matrix, LTxKT from df
# z_data <- df |>
#   select(l,t,k,z_lkt) |>
#   pivot_wider(names_from = k, values_from = z_lkt) |>
#   select(-l, -t)

```

20. Bartik Instrument and Rotemberg weights

```
# this makes a wide Z data, which is LTxKT, KT variables are generated by the spread function.
# basically expand the shares to KT variables, filled with 0 if there is no data.
# Use z_lk0 (the INITIAL shares), not z_lkt: the Bartik instrument is B = Z_0 G,
# so the Rotemberg decomposition has to be taken with respect to the same Z that
# builds the instrument. Using z_lkt here would decompose a different estimator.
# z_data <- df |>
#   select(l,t,k,z_lkt) |>
#   mutate(k = str_glue("t{t}_sh_ind_{k}") ) |>
#   spread(k, z_lkt, fill = 0)

z_data <- df |>
  select(l,t,k,z_lk0) |>
  mutate(k = str_glue("t{t}_sh_ind_{k}") ) |>
  pivot_wider(names_from = k, values_from = z_lk0, values_fill = 0)

Z <- as.matrix(z_data[, -c(1,2)])

# Residualize Y and X on the controls before forming the weights. There are
# no controls here other than the constant, so the residualization is just
# demeaning; this is what makes the Rotemberg decomposition exact (GPSS).
Yp <- Y - mean(Y)
Xp <- X - mean(X)
beta = round((t(Z) %*% Yp)/(t(Z) %*% Xp), digits=3)
alpha=(diag(as.numeric(G)) %*% t(Z) %*% Xp) / as.numeric((t(G) %*% t(Z) %*% Xp))

bw <- tibble(g_kt, alpha=as.numeric(alpha), beta=as.numeric(beta)) |>
  mutate(product=alpha*beta)

bw
```

```
# A tibble: 16 x 6
      t     k   g_kt   alpha beta  product
  <int> <int>   <dbl>   <dbl> <dbl>   <dbl>
1     1     1  0.130   0.0103 0.741  0.00760
2     1     2  0.146   0.0188 1.19   0.0223
3     1     3 -0.165   0.0103 0.771  0.00794
4     1     4 -0.110   0.00760 1.04   0.00787
5     2     1 -0.0289 -0.00789 0.886 -0.00699
6     2     2 -0.00130 -0.000387 0.941 -0.000364
7     2     3  0.0433  0.0133 0.898  0.0119
8     2     4  0.0840  0.0280 0.913  0.0256
9     3     1 -0.101   0.138  0.987  0.136
10    3     2 -0.152   0.222  0.986  0.219
11    3     3 -0.0569  0.0758 0.984  0.0746
12    3     4 -0.0994  0.136  0.993  0.135
13    4     1  0.0266  0.0268 1.01   0.0272
```


14	4	2	0.139	0.155	0.984	0.152
15	4	3	0.0263	0.0267	1.01	0.0270
16	4	4	0.127	0.140	1.02	0.143

```
sum(bw$product)
```

```
[1] 0.9894259
```

The sum of the product of α and β should be the same as the Bartik estimator.

Note there is no restriction for α , other than the sum is 1. Therefore it can be negative or positive, and individual weights can exceed 1, as long as they sum to 1. In GPSS the Rotemberg weights routinely take negative values and values greater than 1; this is a feature used to flag problematic instruments.

20.5. Why is good to have Rotemberg weights?

GPSS Corollary D.1. shows the interpretation of the Rotemberg weights.

The percentage of bias can be written as:

$$\frac{E(\tilde{\beta})}{\beta_0} = \sum_k \alpha_k \frac{E(\tilde{\beta}_k)}{\beta_0} \quad (20.10)$$

This is to say the α_k is the share of the bias of the estimate of β that comes from the bias of the estimate of β_k . That is, how sensitive it is to the bias from industry k .

In my simulations, all the industries are random, therefore the weights are random too. But in empirical studies, they are not random, most likely dominated by some industries. This is why it is important to have Rotemberg weights. It tells us which industries are important in the estimation of the effect of interest.

20.6. What if you are using a reduced form?

If you use the reduced form, something like:

$$y_{lt} = D_{lt}\rho + \beta_0 B_{lt} + \epsilon_{lt} \quad (20.11)$$

where B_{lt} is the Bartik instrument, constructed in such as a way:

$$B_{lt} = Z_{l0}G_t = \sum_{k=1}^K z_{lk0}g_{kt} \quad (20.12)$$

Then I think the Rotemberg weights should be calculated as:

	(1)
(Intercept)	−0.001
	(0.002)
b_lt	0.909
	(0.023)
Num.Obs.	4000
R2	0.278
R2 Adj.	0.278
AIC	−6730.0
BIC	−6711.1
Log.Lik.	3367.980
F	1538.779
RMSE	0.10

$$\beta_{Bartik} = \sum_k \hat{\alpha}_k \hat{\beta}_k \quad (20.13)$$

where

$$\hat{\beta}_k = (Z'_k B^\perp)^{-1} Z'_k Y^\perp \quad (20.14)$$

and

$$\hat{\alpha}_k = \frac{g_k Z'_k B^\perp}{\sum_{k'} g_{k'} Z'_{k'} B^\perp} \quad (20.15)$$

where $B^\perp$ and $Y^\perp$ are residualized B and Y .

```
result_reduced <- df_lt |>
  lm(formula = y_lt ~ b_lt)

modelsummary::modelsummary(result_reduced)
```

Now we calculate the Rotemberg weights for the reduced form, for our simulated data.

```
B <- as.matrix(b_lt[,3])

# Residualize Y and B on the controls (here only a constant demean), as in
# the GPSS reduced-form Rotemberg formula above, so the decomposition is exact.
Yp <- Y - mean(Y)
Bp <- B - mean(B)
beta = (t(Z) %*% Yp)/(t(Z) %*% Bp)
alpha=(diag(as.numeric(G)) %*% t(Z) %*% Bp) / as.numeric((t(G) %*% t(Z) %*% Bp))

bw <- tibble(g_kt, alpha=as.numeric(alpha), beta=as.numeric(beta)) |>
```

```
mutate(product = alpha * beta)
```

```
bw
```

```
# A tibble: 16 x 6
      t     k   g_kt   alpha beta product
  <int> <int>   <dbl>   <dbl> <dbl>   <dbl>
1     1     1  0.130   0.0218 0.320  0.00698
2     1     2  0.146   0.0285 0.720  0.0205
3     1     3 -0.165   0.0322 0.226  0.00730
4     1     4 -0.110   0.0141 0.514  0.00723
5     2     1 -0.0289 -0.00628 1.02 -0.00642
6     2     2 -0.00130 -0.000344 0.973 -0.000335
7     2     3  0.0433   0.0132 0.829  0.0110
8     2     4  0.0840   0.0304 0.773  0.0235
9     3     1 -0.101   0.128   0.978  0.125
10    3     2 -0.152   0.211   0.954  0.201
11    3     3 -0.0569   0.0687 0.998  0.0685
12    3     4 -0.0994   0.127   0.982  0.124
13    4     1  0.0266   0.0240 1.04   0.0250
14    4     2  0.139   0.152   0.917  0.140
15    4     3  0.0263   0.0238 1.04   0.0248
16    4     4  0.127   0.133   0.989  0.131
```

```
sum(bw$product)
```

```
[1] 0.9090661
```

Systematic treatment: [R](#) · [Julia](#).

21. Extended TWFE and other DiD estimators

21.1. Staggered DiD without treatment reversal

This blog is adopted from the tutorial of paneltools package. https://yiqingxu.org/tutorials/panel.html#Installing_Packages

I added Wooldridge's ETWFE, which is my favorite estimator with staggered DiD. It is easy to understand and implement. It also works with nonlinear model such as logit and poisson.


```
#devtools:: install_github("xuyiqing/paneltools")
library(paneltools)
data(paneltools)
library(tidyverse)
hh2019 <- as_tibble(hh2019)
hh2019
```

```
# A tibble: 22,971 x 4
   bfs  year nat_rate_ord indirect
  <dbl> <dbl>         <dbl>     <dbl>
1     1  1991             0         0
2     1  1992             0         0
3     1  1993             0         0
4     1  1994             3.45        0
5     1  1995             0         0
6     1  1996             0         0
7     1  1997             0         0
8     1  1998             0         0
9     1  1999             0         0
10    1  2000             0         0
# i 22,961 more rows
```

This data is from Hainmueller and Hopkins (2019), in which the authors investigate the effects of indirect democracy (versus direct democracy) on naturalization rates in Switzerland using municipality-year level panel data from 1991 to 2009. The study shows that a switch from direct to indirect democracy increased naturalization rates by an average of 1.22 percentage points (Model 1 of Table 1 in the paper).

21. Extended TWFE and other DiD estimators

```
#devtools:: install_github("xuyiqing/panelView")
library(panelView)
panelview(nat_rate_ord ~ indirect, data = hh2019, index = c("bfs","year"),
  xlab = "Year", ylab = "Unit", display.all = T,
  gridOff = TRUE, by.timing = TRUE)
```



21.1.1. TWFE

First we do a TWFE model on it.

```
library(fixest)
model_twfe <- feols(nat_rate_ord~indirect|bfs+year,
  data=hh2019, cluster = "bfs")
summary( model_twfe)
```

```
OLS estimation, Dep. Var.: nat_rate_ord
Observations: 22,971
Fixed-effects: bfs: 1,209, year: 19
Standard-errors: Clustered (bfs)

      Estimate Std. Error t value Pr(>|t|)
indirect 1.33932 0.186525 7.18039 1.2117e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
RMSE: 4.09541      Adj. R2: 0.152719
              Within R2: 0.005173
```

TWFE has issues when there is staggered adoption and heterogeneous treatment effects.

Goodman-Bacon (2021) demonstrates that the two-way fixed-effects (TWFE) estimator in a staggered adoption setting can be represented as a weighted average of all possible 2x2 difference-in-differences (DID) estimates between different cohorts. However, when treatment effects change over time heterogeneously across cohorts, the “forbidden” comparisons that use post-treatment data from early adopters as controls for late adopters may introduce bias in the TWFE estimator. We employ the procedure outlined in Goodman-Bacon (2021) and decompose the TWFE estimate using the command `bacon` in the `bacondecomp` package.

“Negative weights” can be the other issue. Basically it is a weighted average of comparison between different cohorts. The weights can be negative in a TWFE model.

```
library(bacondecomp)
df_bacon <- bacon(nat_rate_ord~indirect,
                  data = hh2019,
                  id_var = "bfs",
                  time_var = "year")
```


	type	weight	avg_est
1	Earlier vs Later Treated	0.17605	1.97771
2	Later vs Always Treated	0.31446	0.75233
3	Later vs Earlier Treated	0.05170	0.32310
4	Treated vs Untreated	0.45779	1.61180

```
coef_bacon <- sum(df_bacon$estimate * df_bacon$weight)
coef_bacon
```

```
[1] 1.339325
```

```
ggplot(df_bacon) +
  aes(x = weight, y = estimate, shape = factor(type), color = factor(type)) +
  labs(x = "Weight", y = "Estimate", shape = "Type", color = 'Type') +
  geom_point()
```

21. Extended TWFE and other DiD estimators



The problem shown in Bacon decomposition here is the comparison between “Later” and “Earlier Treated”. It’s a “forbidden” comparison.

TWFE with event history type of graph.

```
df.use <- hh2019 |>
  na.omit()

df <- as.data.frame(
  hh2019 |>
  group_by(bfs) |>
  mutate(treatment_mean = mean(indirect, na.rm = TRUE))
)
df.use <- df[which(df$treatment_mean < 1),]

df.twfe <- as_tibble(get_cohort(df.use, D="indirect", index=c("bfs", "year"), start0 = TRUE))


# time_to_treatment is NA when the unit is always treated, or never treated. It needs to be set to 0
df.twfe$treat <- as.numeric(df.twfe$treatment_mean > 0) # drop always treated units
df.twfe[which(is.na(df.twfe$Time_to_Treatment)), 'Time_to_Treatment'] <- 0 # can be an arbitrary value
twfe.est <- feols(nat_rate_ord ~ i(Time_to_Treatment, treat, ref = -1) | bfs + year,
  data = df.twfe, cluster = "bfs")
summary(twfe.est)
```

```
OLS estimation, Dep. Var.: nat_rate_ord
Observations: 17,594
Fixed-effects: bfs: 926, year: 19
Standard-errors: Clustered (bfs)
```


	Estimate	Std. Error	t value	Pr(> t)
Time_to_Treatment::-18:treat	-0.320651	0.632702	-0.506796	6.1242e-01
Time_to_Treatment::-17:treat	-0.118540	0.391148	-0.303058	7.6191e-01
Time_to_Treatment::-16:treat	-0.395632	0.402921	-0.981910	3.2640e-01
Time_to_Treatment::-15:treat	-0.410976	0.343899	-1.195048	2.3237e-01
Time_to_Treatment::-14:treat	0.244512	0.369127	0.662406	5.0788e-01
Time_to_Treatment::-13:treat	-0.044310	0.354017	-0.125164	9.0042e-01
Time_to_Treatment::-12:treat	-0.545891	0.315172	-1.732042	8.3600e-02
Time_to_Treatment::-11:treat	-0.293138	0.346835	-0.845180	3.9823e-01
Time_to_Treatment::-10:treat	-0.423981	0.319375	-1.327534	1.8466e-01
Time_to_Treatment::-9:treat	-0.467114	0.301017	-1.551784	1.2106e-01
Time_to_Treatment::-8:treat	-0.115633	0.333659	-0.346561	7.2900e-01
Time_to_Treatment::-7:treat	-0.724290	0.305727	-2.369075	1.8037e-02 *
Time_to_Treatment::-6:treat	-0.572950	0.317153	-1.806540	7.1159e-02
Time_to_Treatment::-5:treat	-0.554993	0.309246	-1.794667	7.3033e-02
Time_to_Treatment::-4:treat	0.066209	0.345103	0.191853	8.4790e-01
Time_to_Treatment::-3:treat	-0.509530	0.316296	-1.610929	1.0754e-01
Time_to_Treatment::-2:treat	-0.195091	0.330220	-0.590793	5.5480e-01
Time_to_Treatment::0:treat	0.522715	0.359646	1.453415	1.4645e-01
Time_to_Treatment::1:treat	1.668834	0.367919	4.535879	6.4907e-06 ***
Time_to_Treatment::2:treat	1.724230	0.471363	3.657968	2.6858e-04 ***
Time_to_Treatment::3:treat	1.178336	0.400559	2.941732	3.3452e-03 **
Time_to_Treatment::4:treat	0.756323	0.478697	1.579962	1.1446e-01
Time_to_Treatment::5:treat	2.438229	0.705451	3.456269	5.7263e-04 ***
Time_to_Treatment::6:treat	2.848590	0.805315	3.537238	4.2442e-04 ***
Time_to_Treatment::7:treat	2.409359	0.783716	3.074275	2.1721e-03 **
Time_to_Treatment::8:treat	2.270428	0.797869	2.845617	4.5305e-03 **
Time_to_Treatment::9:treat	2.392678	0.863222	2.771799	5.6867e-03 **
Time_to_Treatment::10:treat	2.002836	0.953654	2.100170	3.5984e-02 *
Time_to_Treatment::11:treat	0.842590	0.560248	1.503961	1.3293e-01
Time_to_Treatment::12:treat	3.831918	2.300422	1.665745	9.6103e-02
Time_to_Treatment::13:treat	4.436437	2.322318	1.910349	5.6397e-02
Time_to_Treatment::14:treat	4.756977	3.485523	1.364781	1.7265e-01
Time_to_Treatment::15:treat	1.541967	1.415284	1.089511	2.7621e-01
Time_to_Treatment::16:treat	7.168943	4.844016	1.479959	1.3922e-01
Time_to_Treatment::17:treat	6.490854	2.726526	2.380632	1.7485e-02 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
RMSE: 4.42406 Adj. R2: 0.146601
Within R2: 0.012807

```
twfe.output <- as.matrix(twfe.est$coefstable)
twfe.output <- as.data.frame(twfe.output)
twfe.output$Time <- c(c(-18:-2),c(0:17))+1
p.twfe <- esplot(twfe.output,Period = 'Time',Estimate = 'Estimate',
                 SE = 'Std. Error', xlim = c(-12,10))
p.twfe
```



21.1.2. Sun and Abraham

Sun and Abraham (2021) estimator is a weighted average of the average treatment effect (ATT) estimates for each cohort, obtained from a TWFE regression that includes cohort dummies fully interacted with indicators of relative time to the treatment's onset. The iw estimator is robust to heterogeneous treatment effects (HTE) and can be implemented using the command `sunab` in the package `fixest`.

```
df.sa <- df.twfe
df.sa[which(is.na(df.sa$FirstTreat)), "FirstTreat"] <- 1000 # replace NA with an arbitrary number
model.sa.1 <- feols(nat_rate_ord ~ unab(FirstTreat, year) | bfs + year,
                    data = df.sa, cluster = "bfs")
summary(model.sa.1, agg = "ATT")
```

OLS estimation, Dep. Var.: nat_rate_ord

Observations: 17,594

Fixed-effects: bfs: 926, year: 19

Standard-errors: Clustered (bfs)

	Estimate	Std. Error	t value	Pr(> t)
ATT	1.3309	0.287971	4.62163	4.3474e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


RMSE: 4.34795 Adj. R2: 0.163888

Within R2: 0.046484

```

sa.output <- as.matrix(model.sa.1$coeftable)
sa.output <- as.data.frame(sa.output)
sa.output$Time <- c(c(-18:-2),c(0:17))+1
p.sa <- esplot(sa.output,Period = 'Time',Estimate = 'Estimate',
               SE = 'Std. Error', xlim = c(-12,10))
p.sa

```



21.2. Callaway and Sant'Anna

Callaway and Sant'Anna (2021) build the estimator from group-time average treatment effects $ATT(g, t)$ — the effect on cohort g at time t — each identified by comparing the treated cohort to a never-treated or not-yet-treated comparison group. The $ATT(g, t)$ are estimated by outcome regression, inverse-probability weighting, or a doubly-robust combination of the two, rather than from a single fully-interacted TWFE regression. These group-time effects are then aggregated with weights into summary parameters (overall, event-study, or by cohort). The estimator is robust to heterogeneous treatment effects (HTE) and can be implemented with the `att_gt` function in the `did` package.

```

library(did)
df.cs <- df.twfe
df.cs[which(is.na(df.cs$FirstTreat)), "FirstTreat"] <- 0 # replace NA with 0
cs.est.1 <- att_gt(yname = "nat_rate_ord",
                  gname = "FirstTreat",
                  idname = "bfs",
                  tname = "year",
                  xformula = ~1,

```

21. Extended TWFE and other DiD estimators

```
control_group = "nevertreated",
allow_unbalanced_panel = TRUE,
data = df.cs,
est_method = "reg")
cs.est.att.1 <- aggte(cs.est.1, type = "simple", na.rm=T, bstrap = F)
print(cs.est.att.1)
```

Call:

```
aggte(MP = cs.est.1, type = "simple", na.rm = T, bstrap = F)
```

Reference: Callaway, Brantly and Pedro H.C. Sant'Anna. "Difference-in-Differences with Multiplicity", 2021. <<https://doi.org/10.1016/j.jeconom.2020.12.001>>, <<https://arxiv.org/abs/1803.09015>>

ATT	Std. Error	[95% Conf. Int.]
1.3309	0.3019	0.7392 1.9226 *

Signif. codes: `*' confidence band does not cover 0

Control Group: Never Treated, Anticipation Periods: 0

Estimation Method: Outcome Regression

```
cs.att.1 <- aggte(cs.est.1, type = "dynamic",
                  bstrap=FALSE, cband=FALSE, na.rm=T)
print(cs.att.1)
```

Call:

```
aggte(MP = cs.est.1, type = "dynamic", na.rm = T, bstrap = FALSE,
      cband = FALSE)
```

Reference: Callaway, Brantly and Pedro H.C. Sant'Anna. "Difference-in-Differences with Multiplicity", 2021. <<https://doi.org/10.1016/j.jeconom.2020.12.001>>, <<https://arxiv.org/abs/1803.09015>>

Overall summary of ATT's based on event-study/dynamic aggregation:

ATT	Std. Error	[95% Conf. Int.]
1.2205	0.6314	-0.0171 2.458

Dynamic Effects:

Event time	Estimate	Std. Error	[95% Pointwise Conf. Band]
-17	0.1350	0.6467	-1.1324 1.4025
-16	-0.0839	0.3583	-0.7861 0.6183

-15	0.1467	0.3964	-0.6303	0.9236
-14	0.7509	0.3219	0.1200	1.3818 *
-13	-0.2318	0.3392	-0.8967	0.4331
-12	-0.5851	0.2814	-1.1367	-0.0335 *
-11	0.2656	0.2705	-0.2647	0.7958
-10	-0.2386	0.2805	-0.7883	0.3111
-9	0.0576	0.2702	-0.4721	0.5872
-8	0.5244	0.3165	-0.0959	1.1448
-7	-0.6282	0.3132	-1.2420	-0.0144 *
-6	0.1081	0.2565	-0.3946	0.6108
-5	0.1623	0.2890	-0.4041	0.7287
-4	0.6485	0.3380	-0.0139	1.3109
-3	-0.6161	0.3674	-1.3362	0.1040
-2	0.1505	0.3355	-0.5071	0.8081
-1	0.0279	0.3425	-0.6435	0.6992
0	0.7141	0.3754	-0.0217	1.4500
1	1.5846	0.3681	0.8632	2.3060 *
2	1.6337	0.5118	0.6306	2.6369 *
3	0.9090	0.4396	0.0473	1.7707 *
4	0.4091	0.5725	-0.7130	1.5311
5	2.1106	0.7976	0.5473	3.6739 *
6	2.4842	0.9078	0.7051	4.2634 *
7	2.3802	0.8538	0.7068	4.0535 *
8	2.5602	0.8723	0.8506	4.2699 *
9	2.0100	0.9769	0.0953	3.9246 *
10	1.6391	1.0257	-0.3712	3.6494
11	0.0298	0.6100	-1.1658	1.2253
12	0.6942	2.8043	-4.8021	6.1904
13	1.1355	2.3284	-3.4281	5.6991
14	1.3625	3.4131	-5.3270	8.0519
15	-2.0699	1.7160	-5.4331	1.2933
16	2.2743	3.6182	-4.8172	9.3658
17	0.1072	4.5338	-8.7789	8.9934

Signif. codes: '*' confidence band does not cover 0


Control Group: Never Treated, Anticipation Periods: 0

Estimation Method: Outcome Regression

```

cs.output <- cbind.data.frame(Estimate = cs.att.1$att.egt,
                              SE = cs.att.1$se.egt,
                              time = cs.att.1$egt + 1)
p.cs.1 <- esplot(cs.output, Period = 'time', Estimate = 'Estimate',
                 SE = 'SE', xlim = c(-12,10))
p.cs.1

```



21.2.1. Imai, Kim, and Wang

PanelMatch is an R package implementing a set of methodological tools proposed by Imai, Kim, and Wang (2021) that enables researchers to apply matching methods for causal inference on time-series cross-sectional data with binary treatments. The package includes implementations of matching methods based on propensity scores and Mahalanobis distance, as well as weighting methods. PanelMatch enables users to easily calculate a variety of possible quantities of interest, along with standard errors. The software is flexible, allowing users to tune the matching, refinement, and estimation procedures with a large number of parameters. The package also offers a variety of visualization and diagnostic tools for researchers to better understand their data and assess their results.

```
library(PanelMatch)
df.pm <- as.data.frame(df.twfe)
# we need to convert the unit and time indicator to integer
df.pm[, "bfs"] <- as.integer(as.factor(df.pm[, "bfs"]))
df.pm[, "year"] <- as.integer(as.factor(df.pm[, "year"]))
df.pm <- df.pm[, c("bfs", "year", "nat_rate_ord", "indirect")]

# Create PanelData object (new API)
pd <- PanelData(df.pm, unit.id = "bfs", time.id = "year",
               treatment = "indirect", outcome = "nat_rate_ord")

PM.results <- PanelMatch(lag=3,
                        panel.data = pd,
                        refinement.method = "none",
                        qoi = "att",
```

```

        lead = c(0:3),
        match.missing = TRUE)

## For pre-treatment dynamic effects
PM.results.placebo <- PanelMatch(lag=3,
                                panel.data = pd,
                                refinement.method = "none",
                                qoi = "att",
                                lead = c(0:3),
                                match.missing = TRUE,
                                placebo.test = TRUE)
PE.results.pool <- PanelEstimate(PM.results, panel.data = pd, pooled = TRUE)
summary(PE.results.pool)

```

```


      estimate std.error      2.5%      97.5%
[1,] 1.177674 0.3492237 0.4959534 1.872191

```

```

PE.results <- PanelEstimate(PM.results, panel.data = pd)
PE.results.placebo <- placebo_test(PM.results.placebo, panel.data = pd, plot = F)
est_lead <- as.vector(estimates(PE.results))
est_lag <- tryCatch(as.vector(estimates(PE.results.placebo)),
                   error = function(e) as.vector(PE.results.placebo$estimates))
sd_lead <- tryCatch(apply(PE.results$bootstrapped.estimates,2,sd),
                   error = function(e) rep(NA, length(est_lead)))
sd_lag <- tryCatch(apply(PE.results.placebo$bootstrapped.estimates,2,sd),
                  error = function(e) rep(NA, length(est_lag)))
coef <- c(est_lag, 0, est_lead)
sd <- c(sd_lag, 0, sd_lead)
pm.output <- cbind.data.frame(ATT=coef, se=sd, t=c(-2:4))
p.pm <- esplot(data = pm.output, Period = 't',
               Estimate = 'ATT', SE = 'se')
p.pm

```



21.2.2. Liu, Wang and Xu

Liu, Wang, and Xu (2022) proposes a method for estimating the fixed effect counterfactual model by utilizing the untreated observations. It is essentially an imputation method that uses the untreated observations to estimate the counterfactual outcome for the treated.

```
library(fect)
out.fect <- fect(nat_rate_ord~indirect, data = df,
                 index = c("bfs","year"),
                 method = 'fe', se = TRUE)
print(out.fect$est.avg)
```

```
      ATT.avg      S.E. CI.lower CI.upper      p.value
[1,] 1.506416 0.2033546 1.107849 1.904984 1.283418e-13
```


```
fect.output <- as.matrix(out.fect$est.att)
print(fect.output)
```

	ATT	S.E.	CI.lower	CI.upper	p.value	count
-17	-0.01985591	0.5249937	-1.04882465	1.00911284	9.698302e-01	89
-16	0.07269638	0.2600444	-0.43698129	0.58237405	7.798199e-01	157
-15	-0.21365489	0.2524156	-0.70838039	0.28107061	3.973068e-01	162
-14	-0.13153052	0.1533090	-0.43201065	0.16894962	3.909235e-01	350
-13	0.46902947	0.2244936	0.02903007	0.90902887	3.668274e-02	371
-12	0.24357664	0.1960828	-0.14073858	0.62789186	2.141579e-01	398
-11	-0.24399485	0.1332452	-0.50515055	0.01716086	6.707482e-02	417

-10	-0.02765866	0.1742036	-0.36909137	0.31377405	8.738485e-01	434
-9	-0.09347047	0.1417648	-0.37132431	0.18438337	5.096807e-01	451
-8	-0.10638473	0.1648035	-0.42939365	0.21662420	5.185872e-01	464
-7	0.17476094	0.2031886	-0.22348141	0.57300329	3.897382e-01	468
-6	-0.40297604	0.1439751	-0.68516202	-0.12079006	5.127241e-03	503
-5	-0.24326436	0.1767746	-0.58973630	0.10320759	1.687823e-01	503
-4	-0.23687833	0.1716243	-0.57325570	0.09949904	1.675205e-01	503
-3	0.44244127	0.2155461	0.01997871	0.86490383	4.010627e-02	503
-2	-0.11927690	0.1817101	-0.47542211	0.23686831	5.115583e-01	508
-1	0.17612194	0.1932931	-0.20272562	0.55496950	3.622084e-01	512
0	0.22491529	0.1918671	-0.15113738	0.60096795	2.410987e-01	514
1	0.83813813	0.2884721	0.27274312	1.40353314	3.667387e-03	514
2	1.81274486	0.3130701	1.19913872	2.42635100	7.029422e-09	425
3	1.91941052	0.4049619	1.12569975	2.71312129	2.140022e-06	357
4	1.19567464	0.2979622	0.61167949	1.77966979	5.999252e-05	352
5	0.93520961	0.4429749	0.06699484	1.80342438	3.475491e-02	164
6	2.52517848	0.6411037	1.26863840	3.78171855	8.189048e-05	143
7	2.84129736	0.7735045	1.32525638	4.35733835	2.394584e-04	116
8	2.32288547	0.7838386	0.78658995	3.85918100	3.041877e-03	97
9	2.19105962	0.7439582	0.73292831	3.64919092	3.228106e-03	80
10	1.71382906	0.7761588	0.19258578	3.23507233	2.723795e-02	63
11	1.07651548	1.0454215	-0.97247296	3.12550392	3.031306e-01	50
12	-0.30828629	0.6163128	-1.51623712	0.89966453	6.169267e-01	46
13	0.35079211	2.0582844	-3.68337111	4.38495534	8.646725e-01	11
14	0.81696815	2.5111791	-4.10485243	5.73878873	7.449294e-01	11
15	0.97427612	3.8070109	-6.48732812	8.43588037	7.980155e-01	11
16	-2.49173503	1.7363824	-5.89498197	0.91151190	1.512828e-01	11
17	1.42334105	5.3589459	-9.07999999	11.92668210	7.905466e-01	6
18	-0.01384226	4.7462662	-9.31635305	9.28866853	9.976730e-01	2

```
fect.output <- as.data.frame(fect.output)
fect.output$Time <- c(-17:18)
p.fect <- esplot(fect.output, Period = 'Time', Estimate = 'ATT',
                 SE = 'S.E.', CI.lower = "CI.lower",
                 CI.upper = 'CI.upper', xlim = c(-12, 10))
p.fect
```

21. Extended TWFE and other DiD estimators



Borusyak, Jaravel, and Spiess (2021) with package `didimputation` almost gives the same result as Liu, Wang, and Xu (2022).

21.2.3. Wooldridge

Wooldridge 2021, also called extended TWFE (ETWFE), is to allow more flexibility to make TWFE work properly. The problem with TWFE is that we do not specify it flexibly enough. If we allow heterogeneous treatment effect, then we cannot expect the canonical TWFE to work properly.

```
df.wool <- df.twfe
df.wool[which(is.na(df.wool$FirstTreat)), "FirstTreat"] <- 0 # replace NA with 0

library(etwfe)

model.wool = etwfe(
  fml = nat_rate_ord ~ 1, # outcome ~ controls
  tvar = year,            # time variable
  gvar = FirstTreat,      # group variable
  data = df.wool,         # dataset
  vcov = ~bfs,            # vcov adjustment (here: clustered)
  cgroup = "never"
)
model.wool
```

```
OLS estimation, Dep. Var.: nat_rate_ord
Observations: 17,594
Fixed-effects: FirstTreat: 16, year: 19
```

Standard-errors: Clustered (bfs)

```

                                Estimate Std. Error  t value Pr(>|t|)
.Dtreat:FirstTreat::1992:year::1992  2.18863      1.36923   1.59844 0.110287
.Dtreat:FirstTreat::1992:year::1993 -4.18925      3.10938  -1.34730 0.178214
.Dtreat:FirstTreat::1992:year::1994 -2.50997      1.63352  -1.53654 0.124747
.Dtreat:FirstTreat::1992:year::1995 -5.23221      3.11282  -1.68086 0.093129 .
.Dtreat:FirstTreat::1992:year::1996 -5.42014      3.11207  -1.74165 0.081902 .
.Dtreat:FirstTreat::1992:year::1997 -5.15575      3.11248  -1.65648 0.097965 .
.Dtreat:FirstTreat::1992:year::1998 -5.08210      3.11345  -1.63231 0.102955
.Dtreat:FirstTreat::1992:year::1999 -5.13290      3.11447  -1.64808 0.099676 .
... 262 coefficients remaining (display them with summary() or use
argument n)
... 15 variables were removed because of collinearity
(.Dtreat:FirstTreat::1992:year::1991,
.Dtreat:FirstTreat::1993:year::1992 and 13 others [full set in
$collin.var])
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
RMSE: 4.6642      Adj. R2: 0.088474
                Within R2: 0.040642

```

```
emfx(model.wool)
```

```

.Dtreat Estimate Std. Error    z Pr(>|z|)    S 2.5 % 97.5 %
TRUE      1.33      0.288 4.62  <0.001 18.0 0.766    1.9

Term: .Dtreat
Type: response
Comparison: TRUE - FALSE

```

```
mod_es = emfx(model.wool, type = "event", post_only=FALSE)
mod_es
```

```

event Estimate Std. Error    z Pr(>|z|)    S 2.5 % 97.5 %
-18   -0.266     0.789 -0.3371   0.736 0.4 -1.812  1.280
-17    0.178     0.503  0.3538   0.723 0.5 -0.808  1.164
-16    0.114     0.519  0.2198   0.826 0.3 -0.903  1.131
-15   -0.104     0.408 -0.2550   0.799 0.3 -0.903  0.695
-14    0.410     0.413  0.9927   0.321 1.6 -0.400  1.220
--- 26 rows omitted. See ?print.marginaleffects ---
13    1.135     1.998  0.5683   0.570 0.8 -2.781  5.052
14    1.362     3.129  0.4355   0.663 0.6 -4.770  7.495
15   -2.070     1.439 -1.4387   0.150 2.7 -4.890  0.750
16    2.274     3.080  0.7385   0.460 1.1 -3.762  8.311


```

21. Extended TWFE and other DiD estimators

```
17      0.107      4.576  0.0234      0.981 0.0 -8.861  9.076
Term: .Dtreat
Type: response
Comparison: TRUE - FALSE
```

```
wool.output <- mod_es |>
  filter(event != -1) |>
  mutate(Time = event + 1) |>
  as.data.frame()

p.wool <- esplot(wool.output, Period = 'Time', Estimate = 'estimate',
  SE = 'std.error', xlim = c(-12, 10))
p.wool
```



Note that in order to get pre-event estimates, we need to set “cgroup=never” in the `etwfe` function. Otherwise, ETWFE will only give us post-event estimates. Under the default setting, the control group comprises the “not-yet” treated units, therefore all pre-treatment effects are mechanically set to zero. By setting “cgroup=never”, we only use never treated units as the control group. This way we can get pre-treatment estimates, but we are not using all available information for the control group, therefore the estimates may not be as efficient as the default setting. That cost is visible in the two runs below: the never-treated version gives an ATT of 1.33 with a standard error of 0.288, while the default not-yet-treated version gives 1.51 with a standard error of 0.188 — a little over a third narrower.

Also we don’t have control variable here. If we have, it is very easy to incorporate into ETWFE. Wooldridge used the Mundlak device to incorporate control variables, which is a very simple and intuitive way to do it.

If we do the default setting and plot only post-event estimates:

```

model.wool2 = etwfe(
  fml = nat_rate_ord ~ 1, # outcome ~ controls
  tvar = year,           # time variable
  gvar = FirstTreat, # group variable
  data = df.wool,       # dataset
  vcov = ~bfs # vcov adjustment (here: clustered)
)
model.wool2

```

OLS estimation, Dep. Var.: nat_rate_ord

Observations: 17,594

Fixed-effects: FirstTreat: 16, year: 19

Standard-errors: Clustered (bfs)

	Estimate	Std. Error	t value	Pr(> t)
.Dtreat:FirstTreat::1992:year::1992	1.99670	1.35883	1.46943	0.142057
.Dtreat:FirstTreat::1992:year::1993	-4.52291	3.09436	-1.46166	0.144174
.Dtreat:FirstTreat::1992:year::1994	-2.62350	1.61688	-1.62257	0.105022
.Dtreat:FirstTreat::1992:year::1995	-5.17241	3.09454	-1.67146	0.094969
.Dtreat:FirstTreat::1992:year::1996	-5.08905	3.09404	-1.64479	0.100353
.Dtreat:FirstTreat::1992:year::1997	-5.17232	3.09482	-1.67128	0.095004
.Dtreat:FirstTreat::1992:year::1998	-5.39134	3.09606	-1.74135	0.081954
.Dtreat:FirstTreat::1992:year::1999	-5.09733	3.09495	-1.64698	0.099901

... 121 coefficients remaining (display them with summary() or use argument n)

... 141 variables were removed because of collinearity
 (.Dtreat:FirstTreat::1993:year::1992,
 .Dtreat:FirstTreat::1994:year::1992 and 139 others [full set in \$collin.var])

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

RMSE: 4.6999 Adj. R2: 0.081954

 Within R2: 0.0259

```
emfx(model.wool2)
```

	Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
.Dtreat	1.51	0.188	8.02	<0.001	49.7	1.14	1.87

Term: .Dtreat

Type: response

Comparison: TRUE - FALSE

```

mod_es2 = emfx(model.wool2, type = "event")
mod_es2

```

21. Extended TWFE and other DiD estimators

event	Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
0	0.8381	0.288	2.90643	0.00366	8.1	0.273	1.403
1	1.8127	0.316	5.74107	< 0.001	26.7	1.194	2.432
2	1.9194	0.399	4.81257	< 0.001	19.4	1.138	2.701
3	1.1957	0.320	3.73394	< 0.001	12.4	0.568	1.823
4	0.9352	0.423	2.20947	0.02714	5.2	0.106	1.765
5	2.5252	0.641	3.93965	< 0.001	13.6	1.269	3.781
6	2.8413	0.738	3.84931	< 0.001	13.0	1.395	4.288
7	2.3229	0.710	3.27278	0.00106	9.9	0.932	3.714
8	2.1911	0.751	2.91737	0.00353	8.1	0.719	3.663
9	1.7138	0.795	2.15674	0.03103	5.0	0.156	3.271
10	1.0765	0.941	1.14366	0.25277	2.0	-0.768	2.921
11	-0.3083	0.498	-0.61885	0.53602	0.9	-1.285	0.668
12	0.3508	1.892	0.18538	0.85293	0.2	-3.358	4.060
13	0.8170	2.012	0.40611	0.68466	0.5	-3.126	4.760
14	0.9743	3.351	0.29076	0.77123	0.4	-5.593	7.542
15	-2.4917	1.413	-1.76355	0.07781	3.7	-5.261	0.278
16	1.4233	4.430	0.32126	0.74801	0.4	-7.260	10.107
17	-0.0138	4.556	-0.00304	0.99758	0.0	-8.944	8.916

Term: .Dtreat

Type: response


Comparison: TRUE - FALSE

```

wool.output2 <- mod_es2 |>
  filter(event != -1) |>
  mutate(Time = event + 1) |>
  as.data.frame()

p.wool2 <- esplot(wool.output2, Period = 'Time', Estimate = 'estimate',
                  SE = 'std.error', xlim = c(0,10))
p.wool2

```



21.3. Staggered DiD with treatment reversal

With treatment reversal (ie., the treated units go back to control), Wooldridge’s method still works. However, I don’t think package “etwfe” has it implemented.

So far, packages “fect”, “PanelMatch” can handle treatment reversal.

We can probably manually do ETWFE with treatment reversal. See Wooldridge’s dropbox for the code in Stata. The basic idea is to treat cohorts as different treatment beginning time and exit time combinations. Then do the same regression with interactions of cohort dummies with the time dummies. It is very flexible. The tricky part is to get the marginal effect and its standard error right.

Systematic treatment: [R](#) · [Julia](#).

22. Causal Panel (why DDDiD)

22.1. Why DDDiD

Guido Imbens has been saying DDDiD (Don't Do Diff in Diff). But why?

To see that, DiD objective function (in the form of TWFE) is:

$$(\hat{\tau}^{did}, \hat{\mu}, \hat{\alpha}, \hat{\beta}) = \underset{\tau, \mu, \alpha, \beta}{argmin} \sum_{i=1}^N \sum_{t=1}^T (Y_{it} - \mu - \alpha_i - \beta_t - W_{it}\tau)^2 \quad (22.1)$$

We can see that in this objective function, we are allowing only α_i and β_t to capture the difference between units and time periods. The obvious drawback is that control units that are less similar to a treated unit is given the same weight as those that are more similar. A time period that is far from the beginning of the treatment period is given the same weight as those that are closer. Say the treatment period is 2000, we have observations from 1988 to 2005, then 1988 observations are the same as 1999 observations. This is not ideal. The same logic applies to the units.

22.2. Synthetic Control

SC objective function:

$$(\hat{\tau}^{sc}, \hat{\mu}, \hat{\beta}) = \underset{\tau, \mu, \beta}{argmin} \sum_{i=1}^N \sum_{t=1}^T (Y_{it} - \mu - \beta_t - W_{it}\tau)^2 \hat{\omega}_i^{sc} \quad (22.2)$$

It's a weighted version of DiD, but note carefully that it is **not** a nesting of it: setting every $\hat{\omega}_i^{sc}$ to 1 here does not recover the DiD objective above, because SC has no α_i term — it would give a pooled regression with time effects only. Dropping the unit fixed effects is a substantive restriction, not a re-weighting. (DiD *is* nested in synthetic DiD below, which keeps α_i and sets both weights to be constant; that is the right way to see DiD as a special case.) The weights are set to optimally match donor units to the treated unit so that they are as close as possible at each time point. There is no time weight.

22.3. Synthetic DiD

Arkhangelsky et al (2021) tries to combine the idea of DiD and SC. SC assigns different weights to different control units. The standard DiD is a TWFE, assigning equal weights to all time periods and units.

22. Causal Panel (why DDDiD)

$$(\hat{\tau}^{sdid}, \hat{\mu}, \hat{\alpha}, \hat{\beta}) = \underset{\tau, \mu, \alpha, \beta}{\operatorname{argmin}} \sum_{i=1}^N \sum_{t=1}^T \hat{\omega}_i^{sdid} \lambda_t^{sdid} (Y_{it} - \mu - \alpha_i - \beta_t - W_{it}\tau)^2 \quad (22.3)$$

SDiD sets another weight in addition to SC weights, which changes over time. The SC weights are trying to construct a control unit that is close to the treated unit; the SDiD weights are trying to put more weights on pre-treatment periods that are more similar to post-treatment periods.

22.4. Example from synthdid

```
library(synthdid)
data('california_prop99')
setup = panel.matrices(california_prop99)
tau.hat = synthdid_estimate(setup$Y, setup$N0, setup$T0)
summary(tau.hat)
```

```
$estimate
[1] -15.60383
```

```
$se
      [,1]
[1,]    NA
```

```
$controls
      estimate 1
Nevada      0.124
New Hampshire 0.105
Connecticut 0.078
Delaware     0.070
Colorado     0.058
Illinois     0.053
Nebraska     0.048
Montana      0.045
Utah         0.042
New Mexico   0.041
Minnesota    0.039
Wisconsin    0.037
West Virginia 0.034
North Carolina 0.033
Idaho        0.031
Ohio         0.031
Maine        0.028
Iowa         0.026
```

```
$periods
```


```
estimate 1
1988      0.427
1986      0.366
1987      0.206
```

```
$dimensions
      N1      NO NO.effective      T1      T0 T0.effective
      1.000    38.000    16.388    12.000    19.000     2.783
```

```
tau.sc = sc_estimate(setup$Y, setup$N0, setup$T0)
tau.did = did_estimate(setup$Y, setup$N0, setup$T0)
estimates = list(tau.did, tau.sc, tau.hat)
names(estimates) = c('Diff-in-Diff', 'Synthetic Control', 'Synthetic Diff-in-Diff')
print(unlist(estimates))
```

```
Diff-in-Diff      Synthetic Control Synthetic Diff-in-Diff
      -27.34911           -19.61966           -15.60383
```

```
plot <- synthdid_plot(estimates, facet.vertical=FALSE,
  control.name='control', treated.name='california',
  lambda.comparable=TRUE, se.method = 'none',
  trajectory.linetype = 1, line.width=.75, effect.curvature=-.4,
  trajectory.alpha=.7, effect.alpha=.7,
  diagram.alpha=1, onset.alpha=.7) +
  theme(legend.position=c(.26,.07), legend.direction='horizontal',
    legend.key=element_blank(), legend.background=element_blank(),
    strip.background=element_blank(), strip.text.x = element_blank())
plot
```



22. Causal Panel (why DDDiD)

Note: λ_t is plotted at the bottom.

22.5. Summary

To me, “DDDiD” is saying DiD is too simple. Even if all the identification assumptions are correct for DiD, namely parallel trend and no anticipation, there are still better ways to do it, in the sense of putting weights on units and time periods, such as synthetic DiD.

Systematic treatment: *R* · *Julia*.

Part V.

Count Data & Specialized Models

23. Which count data model to use

23.1. A comparison of various count data models with extra zeros

In empirical studies, data sets with a lot of zeros are often hard to model. There are various models to deal with it: zero-inflated Poisson model, Negative Binomial (NB) model, hurdle model, etc.

Here we are following a zero-inflated model's thinking: model the data with two processes. One is a Bernoulli process, the other one is a count data process (Poisson or NB).

We'd like to see, in this simulation exercise, how different models perform with changes of sample size and percentage of zeros (we expect the less zero, the better a plain Poisson model would perform). Therefore we vary sample size n and an indicator of how much percentage of zeros in the data θ .

For the count data process (y_c):

$$\log(y_c) = 2x + u \quad (23.1)$$

For the Bernoulli process (y_b):

$$z_1 = 4z + \theta \quad (23.2)$$

$$\text{logit}(y_b) = z_1 \quad (23.3)$$

$$p_y = \frac{e^{z_1}}{1 + e^{z_1}} \quad (23.4)$$

where $y_b \sim \text{Bernoulli}(p_y)$ is the realized 0/1 draw. Combining these two processes:

$$y = y_c \cdot y_b \quad (23.5)$$

so $y = y_c$ when $y_b = 1$ and $y = 0$ when $y_b = 0$.

23.1.1. Zero-inflated Poisson models

A zero-inflated Poisson needs specifying both the binary process and the count process correctly. Often than not, we don't have a model for the binary process. Many people simply use the same explanatory variables for both processes. We simulate both situations. Case 1: suppose we observe z , and case 2: suppose we don't observe z . In the graph below, they are labeled zip1 and zip2.

23. Which count data model to use

23.1.2. Poisson model

A plain Poisson model returns a consistent estimator for the coefficients, with or without Poisson-distributed data. We expect Poisson model's performance improve with sample size. Note that the standard errors from a Poisson model needs adjustment, which we do not discuss in this post.

23.1.3. NB model

NB model is used widely to handle “overdispersion” problem. That is, the variance far exceeds the mean, therefore the Poisson model is considered inappropriate. NB model addresses that by allowing an extra parameter. However, many people also use it to model “extra zero” situation, we'll see in our simulation it may not be better than a plain Poisson model.

23.1.4. Log-linear model

What about an OLS model with $\log(y + 1)$?

23.1.5. hurdle model

A hurdle model models the zero's and other values separately; that is, the zero's are from a binomial process only, the other positive values are from a truncated count data process. We assume here, in the simulation, we don't observe z . Therefore, x is determining both binary and count processes. In the graph below, it's labeled hurdle.

```
library(MASS)
library(psc1)
library(parallel)
set.seed(666)

gen.sim <- function(df) {
  nobs <- df["nobs"]; th <- df["th"]
  z    <- rnorm(nobs, 0, 1)
  x    <- rnorm(nobs, 0, 1)
  u    <- rnorm(nobs, 0, 1)
  y.count <- floor(exp(2 * x + u))
  z1     <- 4 * z + th
  prob   <- plogis(z1)
  y.logit <- rbinom(nobs, size = 1, prob = prob)
  y      <- ifelse(y.logit == 1, y.count, y.logit)

  m1 <- zeroinfl(y ~ x | z)
  m4 <- zeroinfl(y ~ x | x)
  m2 <- glm(y ~ x, family = "poisson")
  m3 <- lm(log(y + 1) ~ x)
  m5 <- glm.nb(y ~ x)
```



```

m6 <- hurdle(y ~ x)

c(zip1      = coef(m1, "count")["x"] - 2,
  poisson    = coef(m2, "x") - 2,
  log.linear = coef(m3, "x") - 2,
  zip2       = coef(m4, "count")["x"] - 2,
  nb         = coef(m5, "x") - 2,
  hurdle     = coef(m6, "count")["x"] - 2)
}

data.grid <- expand.grid(
  nobs = ceiling(exp(seq(4, 9, 1)) / 100) * 100,
  nsim = seq(1, 100, 1),
  th   = seq(-4, 4, 2)
)

cl      <- makeCluster(detectCores() - 1)
clusterEvalQ(cl, { library(MASS); library(psc1) })
clusterExport(cl, "gen.sim")
results <- parApply(cl, data.grid, 1, gen.sim)
stopCluster(cl)

forshiny <- cbind(data.grid, t(results))
write.csv(forshiny, "results.csv", row.names = FALSE)

```

Count data models can be used even if data is not “counts”; for example, some non-negative non-integer numbers. In fact, Poisson model is consistent even if data is not Poisson-distributed, if the model specification is correct on modeling the log of expected counts. We simulate both scenarios: Case 1, data is generated from a Poisson process. Case 2, data is generated from a Normal distribution, but we use count data models to model it. The above code is for case 2.

We simulate 100 times with θ ranging from -4 to 4, lower number means higher percentage of zeros; number of observations from e^4 to e^9 (roughly from 50 to 8000).

Since there are many simulations, we use base R’s `parallel` library to speed things up.

For raw code, please visit [case1: poisson](#) and [case2: normal](#). The code above updates the original `snowfall`-based version to use base R `parallel`.

Summary of simulation findings: With large samples, plain Poisson outperforms NB and log-linear even with extra zeros. Zero-inflated Poisson with correct specification of the zero process is best, but that relies on knowing what drives the zeros. Hurdle and ZIP without correct zero-model specification are similar and not much worse than Poisson. Log-linear ($\log(y+1)$) is consistently the worst.

(Note: the `log.linear` bias line above previously read `exp(coef(m3)["x"]) - 2` instead of `coef(m3)["x"] - 2`, which — since the log-linear coefficient is already on the log scale — manufactured a spurious 5.4 apparent bias regardless of the model’s actual performance. That has been corrected here to match the other rows. Log-linear regression on $\log(y+1)$ is still expected to

23. Which count data model to use

be biased relative to Poisson-family models in a zero-inflated count setting, for real reasons (the log transform's nonlinearity near zero, and the fact that $E[\log(y + 1)] \neq \log(E[y] + 1)$), but the *magnitude* of “consistently the worst” above reflects the original (buggy) run and has not been re-verified against a corrected simulation.)

23.2. Quick worked example

All five models on a single small dataset ($n = 1000$, ~40% zeros):

```
library(MASS)
library(pscl)
set.seed(42)

n <- 1000
x <- rnorm(n)
z <- rnorm(n)                # zero-inflation driver
mu <- exp(1.5 + 0.8 * x)      # count mean
y_c <- rpois(n, lambda = mu)
y_b <- rbinom(n, 1, plogis(1 + 2 * z)) # 1 = non-zero
y   <- y_c * y_b              # zero-inflated count

cat("Zero fraction:", round(mean(y == 0), 2), "\n")
```

Zero fraction: 0.39

```
cat("Mean (non-zero):", round(mean(y[y > 0]), 2), "\n\n")
```

Mean (non-zero): 6.79

```
# 1. Poisson
m_pois <- glm(y ~ x, family = poisson)

# 2. Negative binomial
m_nb <- glm.nb(y ~ x)

# 3. Zero-inflated Poisson (ZIP) - correct zero model
m_zip <- zeroinfl(y ~ x | z)

# 4. ZIP - wrong zero model (only x)
m_zip2 <- zeroinfl(y ~ x | x)

# 5. Hurdle
m_hurdle <- hurdle(y ~ x)
```

```
# 6. Log-linear OLS on log(y + 1)
m_loglin <- lm(log(y + 1) ~ x)

# Compare x coefficient (true = 0.8)
coefs <- c(
  Poisson      = coef(m_pois)["x"],
  NB           = coef(m_nb)["x"],
  `ZIP (correct z)` = coef(m_zip, "count")["x"],
  `ZIP (wrong)` = coef(m_zip2, "count")["x"],
  Hurdle       = coef(m_hurdle, "count")["x"],
  `Log-linear log(y+1)` = coef(m_loglin)["x"]
)

knitr::kable(
  data.frame(Model = names(coefs), `x coefficient` = round(coefs, 3),
    Bias = round(coefs - 0.8, 3)),
  caption = "True x coefficient = 0.8"
)
```

Table 23.1.: True x coefficient = 0.8

	Model	x.coefficient	Bias
Poisson.x	Poisson.x	0.850	0.050
NB.x	NB.x	0.824	0.024
ZIP (correct z).x	ZIP (correct z).x	0.803	0.003
ZIP (wrong).x	ZIP (wrong).x	0.802	0.002
Hurdle.x	Hurdle.x	0.803	0.003
Log-linear log(y+1).x	Log-linear log(y+1).x	0.458	-0.342

The first five models recover the true x slope of 0.8 reasonably well; the log-linear row is included to make the earlier claim checkable, and needs reading with care — see the note after the table. This is the key point of *this* DGP: the zero-inflation process depends on z , which is generated **independently of x** , so omitting the zero model distorts the zero mass and the intercept, not the slope on x . The hurdle model and even the ZIP with a *misspecified* (x -only) zero model still recover the slope closely; the plain Poisson shows a small upward bias *in this particular sample* ($n = 1000$), but that is finite-sample noise, not a systematic property of the model. Since the zero-inflation driver z is independent of x , the conditional mean here is correctly specified for Poisson QMLE — increasing n from 1,000 to 1,000,000 shrinks the Poisson x -coefficient bias from about +0.04 to effectively 0, confirming consistency. The unmodeled excess zeros create overdispersion, but overdispersion under a correctly-specified conditional mean affects only the standard errors (which should be corrected with a robust/sandwich estimator), not the consistency of the slope.

About the log-linear row: its coefficient is **not** estimating the same thing as the others, which is the real reason to avoid it rather than any particular bias number. The count models target $\log E[y \mid x]$, whereas OLS on $\log(y + 1)$ targets $E[\log(y + 1) \mid x]$, and $E[\log(y + 1)] \neq \log(E[y] + 1)$ by Jensen's inequality. The +1 offset makes the discrepancy worst exactly where the mass of zeros is. So

23. Which count data model to use

comparing its coefficient to 0.8 is an apples-to-oranges comparison, and the honest statement is that log-linear answers a different question — not that it answers this one badly by some measured amount coefficient. A misspecified zero process biases the count *slope* only when the zero process is correlated with the count covariates — which is not the case here. (Change *z* to depend on *x* and the omitted-zero-process slope bias appears.)

23.3. Mixed count models with glmmTMB

When data has both overdispersion *and* random effects (clustered counts, panel data, repeated measures), `glmmTMB` handles the full range in one package:

```
library(glmmTMB)

# Zero-inflated negative binomial with random intercept per subject
m_glmmtmb <- glmmTMB(
  count ~ treatment + time + (1 | subject),    # count model
  ziformula = ~ treatment,                    # zero-inflation model
  family = nbinom2,
  data = your_data
)

summary(m_glmmtmb)

# Conditional vs marginal predictions
predict(m_glmmtmb, type = "conditional") # excluding zero-inflation
predict(m_glmmtmb, type = "response")    # including zero-inflation
```

`glmmTMB` supports: Poisson, NB1, NB2, zero-inflated variants of all of the above, hurdle models, beta-binomial, tweedie, and more. The `ziformula` argument takes a one-sided formula for the zero-inflation component — use `~ 1` for a constant zero-inflation rate, or covariates to model what drives the structural zeros.

24. What model to use for rare events

24.1. Introduction

In empirical studies, people are worried about rare event situation. That is, when you have, for example, lots of 0's and only a few 1's, or vice versa. Do you run a logit model, or do you use a "rare event logit"? When should you use either approach? Or there is a third approach?

Paul Allison said in his blog (<https://statisticalhorizons.com/logistic-regression-for-rare-events>):

"Prompted by a 2001 article by King and Zeng, many researchers worry about whether they can legitimately use conventional logistic regression for data in which events are rare. Although King and Zeng accurately described the problem and proposed an appropriate solution, there are still a lot of misconceptions about this issue.

The problem is not specifically the rarity of events, but rather the possibility of a small number of cases on the rarer of the two outcomes. If you have a sample size of 1000 but only 20 events, you have a problem. If you have a sample size of 10,000 with 200 events, you may be OK. If your sample has 100,000 cases with 2000 events, you're golden."

In general I agree with him. However, when exactly should we use King and Zeng's "relogit"?

Allison also mentioned two other methods. One is called the Firth method, a penalized likelihood approach. The other one is the exact logistic regression, which is for small samples.

In this simulation exercise, I compare logit against bias-reduced logistic regression (`brglm2` with `method="brglmFit", type="AS_mean"`) and the Firth model (`logistf`).

One thing to be clear about before reading any comparison of the latter two: for a logistic regression **they are the same estimator**. Firth's penalized likelihood (penalizing by the Jeffreys prior) coincides with mean bias reduction for GLMs with a canonical link, and logit is canonical for the binomial. So `brglm2`'s `AS_mean` and `logistf` solve the same estimating equations and differ only in numerical tolerance and in how each package builds standard errors and intervals (`logistf` uses penalized profile likelihood). A demonstration:

```
library(brglm2); library(logistf)
set.seed(11)
n <- 500; x <- rnorm(n); y <- rbinom(n, 1, plogis(-3.5 + 0.8 * x))
d <- data.frame(y, x)

rbind(
  `plain logit`      = coef(glm(y ~ x, data = d, family = binomial)),
  `brglm2 AS_mean`  = coef(glm(y ~ x, data = d, family = binomial,
                                method = "brglmFit", type = "AS_mean")),
```

24. What model to use for rare events

```
`logistf (Firth)` = coef(logistf(y ~ x, data = d))
)
```

	(Intercept)	x
plain logit	-3.176125	0.7378324
brglm2 AS_mean	-3.134453	0.7277674
logistf (Firth)	-3.134453	0.7277675

They agree to about 10^{-7} . This matters for reading the findings below: what the simulation can tell us is how **logit** compares with **Firth-type bias reduction**, not which of two implementations of the same estimator is better. If a genuinely distinct third method is wanted, **brglm2**'s **type = "AS_median"** (median bias reduction) is one, and King and Zeng's **relogit** correction is another.

24.2. Simulation

Here I have some code for using multiple cores to run these three models. The bias-reduced logistic regression is implemented in R via the **brglm2** package, the Firth method is implemented in **logistf**.

```
library(brglm2)
library(logistf)
require(snowfall)
set.seed(666)

# initialize parallel cores.
sfInit(parallel=TRUE, cpus=16)

gen.sim <- function(df){
  x <- rnorm(df['nobs'], 0, 1)

  # generate binary data
  p <- df['p']
  alpha <- -log((1-p)/p)
  z = alpha + 2*x
  pr = 1/(1+exp(-z))
  y = rbinom(df['nobs'], 1, pr)
  df = data.frame(y=y, x=x)

  # With small nobs and small p (e.g. nobs=10, p=.01), y can easily come out
  # all-zero (>90% chance in that cell), which makes glm/logistf unable to
  # estimate a coefficient on x. Guard against that rather than letting a
  # single degenerate draw crash the whole parallel simulation loop.
  if (var(y) == 0) {
    return(c(logit=NA, brglm2=NA, logistf=NA))
  }
}
```

```

# logit
m1 <- tryCatch(glm(y ~ x, family='binomial'), error = function(e) NULL)
m1.x <- if (!is.null(m1)) summary(m1)$coefficients['x','Estimate'] - 2 else NA

# bias-reduced logistic regression (brglm2)
m2 <- tryCatch(
  glm(y ~ x, family = binomial(link = "logit"), data = df,
    method = "brglmFit", type = "AS_mean"),
  error = function(e) NULL
)
m2.x <- if (!is.null(m2)) coef(m2)['x'] - 2 else NA

# logistf (Firth)
m3 <- tryCatch(logistf(y ~ x, data=df), error = function(e) NULL)
m3.x <- if (!is.null(m3)) coef(m3)['x'] - 2 else NA

return(c(logit=m1.x, brglm2=m2.x, logistf=m3.x))
}

# set parameter space
sim.grid = seq(1, 100, 1)
p.grid = c(.01, .05, .1)
nobs.grid = c(10, 30, 50, 100, 200, 500, 1000, 10000)

data.grid <- expand.grid(nobs.grid, sim.grid, p.grid)
names(data.grid) <- c('nobs', 'nsim', 'p')

# export functions and libraries to parallel workers
sfExport(list=list("gen.sim"))
sfLibrary(brglm2)
sfLibrary(logistf)

results <- data.frame(t(sfApply(data.grid, 1, gen.sim)))

# stop the cluster
sfStop()

forshiny <- cbind(data.grid, results)
write.csv(forshiny, 'results.csv')

```

We simulate 100 times with sample size from 10 to 10000, event probability .01, .05, and .1.

Since there are many simulations, we used the `snowfall` library to speed things up.

(The original post plotted bias and MSE by sample size and event probability from the `results.csv` output above; those figures are not reproduced here — the findings are summarized in prose below.)

24. What model to use for rare events

In the case of small sample and rare event (for example, any situation that the product of sample size and probability is less than 5), none of these three models perform well. This is understandable, after all, the rarer of the two groups has only less than 5 observations. When the product is more than 50, there is not much difference between these three models. For the situations in between, that is, the product of sample size and probability is greater than 5 but less than 50, we found that Firth-type bias reduction — whether computed by `brglm2` or by `logistf` — performs better than plain logit. Any apparent ranking *between* those two should be discounted, since as shown above they are the same estimator; differences at that level reflect numerical tolerance or differing standard-error conventions, not a substantive advantage.

In the small sample situation, maybe it's better to use the exact logistic regression.

25. Poisson model with exposure

25.1. Poisson with exposure or Poisson rate

Poisson models for count data often arise when units are observed for different lengths of time. A hospital admission count for a patient followed for 30 days is not directly comparable to a count for a patient followed for 365 days. The natural quantity of interest is the *rate* — admissions per unit time — not the raw count.

There are three related but distinct specifications, each making a different assumption:

Model 1 — Poisson with offset (rate model):

$$\log(E[Y]/t) = \beta_0 + \beta_1 X \quad (25.1)$$

Equivalently, $\log E[Y] = \beta_0 + \beta_1 X + \log t$, where $\log t$ is included as a *fixed* offset (coefficient constrained to 1). This is the standard exposure model: the expected count is proportional to exposure time.

Model 2 — Poisson on the ratio Y/t :

$$\log(E[Y/t]) = \beta_0 + \beta_1 X \quad (25.2)$$

This looks similar, but it is not equivalent to the offset model, and strictly speaking Y/t is **not a Poisson outcome at all**: a Poisson likelihood is defined for non-negative integer counts, and the continuous ratio Y/t is not one. Because t is a fixed constant for each unit, $E[Y/t] = E[Y]/t$ exactly (linearity of expectation), so the *conditional mean* is still log-linear in X ; that is why the point estimates are often similar. But the Poisson variance assumption fails: $\text{Var}(Y/t) = \lambda/t^2$, whereas the mean is $E[Y/t] = \lambda/t$, so the variance is $1/t$ times the mean rather than equal to it. (Writing it as λ/t would make variance equal mean and quietly restore equidispersion, which is precisely what fails here.) So fitting `glm(rate ~ x, family = poisson)` is a quasi-likelihood *mean model* rather than a correct Poisson likelihood. If you use it, treat it as a workaround: weight by exposure (`weights = t`) so the implied variance is right, and/or use robust standard errors. The offset model (Model 1) avoids the issue entirely by keeping the integer count as the outcome.

Model 3 — Unconstrained log-exposure:

$$\log E[Y] = \beta_0 + \beta_1 X + \gamma \log t \quad (25.3)$$

Here γ is estimated from the data. If $\hat{\gamma} \approx 1$, the proportionality assumption is supported. If $\hat{\gamma} < 1$, counts grow sub-linearly with exposure — a testable, substantive claim.

25.2. Simulation

```
library(tidyverse)
set.seed(42)

n <- 800

# Simulate a health study: count of adverse events
# Covariate: age group (0 = younger, 1 = older)
age_old <- rbinom(n, 1, 0.45)

# Exposure: days under observation (varies 7-365)
exposure <- sample(seq(7, 365, by = 7), n, replace = TRUE)

# True model: rate = exp(-3 + 1.2 * age_old)
# Count ~ Poisson(rate * exposure)
true_rate <- exp(-3 + 1.2 * age_old)
count <- rpois(n, lambda = true_rate * exposure)

df <- data.frame(count, age_old, exposure, log_exposure = log(exposure))
summary(df[, c("count", "exposure")])
```

count	exposure
Min. : 0.00	Min. : 7.0
1st Qu.: 7.00	1st Qu.: 112.0
Median : 13.00	Median : 196.0
Mean : 18.95	Mean : 192.9
3rd Qu.: 27.00	3rd Qu.: 281.8
Max. : 71.00	Max. : 364.0

The data has a wide range of exposure times, so ignoring them would badly confound the comparison between age groups.

25.3. Fitting the three models

```
# Model 1: offset - coefficient on log(exposure) constrained to 1
m1 <- glm(count ~ age_old + offset(log(exposure)),
          family = poisson, data = df)

# Model 2: Poisson on the rate Y/t (different response)
df$rate <- df$count / df$exposure
m2 <- glm(rate ~ age_old,
          family = poisson, data = df)
```

```
# Model 3: log(exposure) as a free covariate
m3 <- glm(count ~ age_old + log_exposure,
          family = poisson, data = df)

# Collect coefficients
results <- data.frame(
  Model = c("1. Offset (constrained =1)",
            "2. Rate Y/t as response",
            "3. Unconstrained log(t)"),
  Intercept = c(coef(m1)[1], coef(m2)[1], coef(m3)[1]),
  age_old = c(coef(m1)[2], coef(m2)[2], coef(m3)[2]),
  gamma_log_t = c(1, NA, coef(m3)[3])
)

knitr::kable(results, digits = 3,
  caption = "Coefficient estimates: true intercept = -3, true age effect = 1.2, true  $\gamma$  = 1")
```

Table 25.1.: Coefficient estimates: true intercept = -3 , true age effect = 1.2 , true $\gamma = 1$

Model	Intercept	age_old	gamma_log_t
1. Offset (constrained =1)	-3.009	1.212	1.000
2. Rate Y/t as response	-2.980	1.176	NA
3. Unconstrained log(t)	-2.801	1.212	0.962

Reading the table: Model 1 and Model 3 give nearly identical age-group coefficients because the true $\gamma = 1$ (counts are indeed proportional to exposure time). Model 2 differs slightly in the intercept because the response variable changes. The estimated $\hat{\gamma}$ from Model 3 should be close to 1.

25.4. When $\gamma \neq 1$

To see why the unconstrained model matters, repeat the simulation with a sub-linear count process ($\gamma = 0.6$: an initial burst of events that doesn't grow proportionally with long follow-up):

```
count2 <- rpois(n, lambda = true_rate * exposure^0.6)
df2 <- data.frame(count2, age_old, exposure, log_exposure = log(exposure))

m1b <- glm(count2 ~ age_old + offset(log(exposure)),
          family = poisson, data = df2)
m3b <- glm(count2 ~ age_old + log_exposure,
          family = poisson, data = df2)

cat("Offset model age coefficient:      ", round(coef(m1b)[2], 3), "\n")
```

25. Poisson model with exposure

Offset model age coefficient: 1.171

```
cat("Unconstrained model age coefficient:", round(coef(m3b)[2], 3), "\n")
```

Unconstrained model age coefficient: 1.168

```
cat("Estimated gamma (log-exposure):    ", round(coef(m3b)[3], 3),  
    " (true = 0.6)\n")
```

Estimated gamma (log-exposure): 0.566 (true = 0.6)

Two things to read off this, and the first is not what one might expect.

The offset model's age coefficient (1.171) and the unconstrained model's (1.168) are essentially the same, and both sit just below the true 1.2. Forcing $\gamma = 1$ when the truth is 0.6 does **not** bias the treatment coefficient here — and the reason is structural: `exposure` is drawn by `sample(...)` independently of `age_old`, so the misspecified exposure elasticity is absorbed by the intercept and cannot contaminate a coefficient on a covariate orthogonal to it. The constraint would bite if exposure were *correlated* with the covariate of interest — for instance if older patients were systematically followed longer, which is common in real cohort data and is the case worth worrying about.

What the unconstrained model does buy, unambiguously, is γ itself: it recovers $\hat{\gamma} = 0.566$ against a truth of 0.6, whereas the offset model cannot report it at all because it fixes it at 1. So Model 3's role is diagnostic — it tells you whether the proportionality assumption holds — rather than protective of the treatment estimate.

25.5. Stata equivalent

```
* Model 1: offset  
poisson count age_old, exposure(exposure)  
  
* Model 2: rate as response  
gen rate = count / exposure  
* Stata will warn that the dependent variable is non-integer (poisson expects  
* integer counts by default) -- this is expected and can be ignored here,  
* since Poisson QMLE is valid for continuous non-negative outcomes too.  
poisson rate age_old  
  
* Model 3: unconstrained log(t)  
gen log_exposure = log(exposure)  
poisson count age_old log_exposure
```

25.6. Which model to use?

Situation	Recommended model
Counts grow proportionally with time	Model 1 (offset)
Uncertain about proportionality	Model 3 (unconstrained), test $\hat{\gamma} = 1$
Rate is the natural estimand	Model 1 (offset) — it estimates a rate directly; Model 2 only as a weighted/robust quasi-model
Long follow-up with declining event rate	Model 3 or negative binomial with offset

In most health and social-science applications, Model 1 (Poisson with offset) is the default. Model 3 is the diagnostic check: fit it, look at $\hat{\gamma}$, and return to Model 1 if the constraint is supported by the data.

Systematic treatment: [R](#) · [Julia](#).

26. IV in Fixed effect Poisson with many fixed effects

26.1. Poisson regression with many fixed effects

I have a few posts about Poisson regression for non-negative outcomes, including continuous variable, count variable, or corner solution (such as non-negative outcome with many zeros). The benefits include that we can use QMLE Poisson, which means we only need to have the mean of outcome specified correctly, as an exponentiated function of covariates. There is no assumption that the distribution needs to be a Poisson distribution. When we want to model non-negative outcome with fixed effects, fixed effect Poisson should be our first choice also, because fixed effect Poisson is shown to be consistent without the assumption of the distribution of the error term, and it does not have the “incidental parameter” problem.

Stata has “xtpoisson” to do fixed effect Poisson. However, we would recommend “ppmlhdfe”. “ppmlhdfe is a Stata package that implements Poisson pseudo-maximum likelihood regressions (PPML) with multi-way fixed effects”. It is faster than xtpoisson. More importantly, it does report the predicted value correctly. I mentioned in past post that we need to be cautious when we need to use predicted values after “xtpoisson, fe”, or “xtreg, fe”. The reason is that these are models estimated after “absorbing” the fixed effects, meaning fixed effects are not actually estimated. Therefore the predicted values are without fixed effects. But in “ppmlhdfe” we get the correct predicted values with fixed effect.

On a side note, the same author has “reghdfe” as a replacement for “areg” in Stata. “reghdfe” is faster than “areg”, or “xtreg, fe”. It also allows for many levels of fixed effects. As “ppmlhdfe”, it also provides correct predicted values with fixed effects.

26.2. IV poisson with many fixed effects

Stata has a “ivpoisson” function with handles IV in a Poisson regression setting. However, it does not have the option to include fixed effects.

This estimator has a published treatment. Lin and Wooldridge (2019), “Testing and correcting for endogeneity in nonlinear unobserved effects models,” treats exactly this setup: their **Procedure 3** is a reduced-form residual added to an FE Poisson, with a robust Wald test on the residual’s coefficient serving as a test of endogeneity. Several points below come from it.

One warning about that test, since it is easy to misuse. It is tempting to run it and then decide which model to report: if the residual’s coefficient is insignificant, fall back on the plain fixed-effects estimate. Do not do this. Guggenberger (2010) studies exactly this two-stage procedure and finds that its asymptotic size **equals one** — the two-stage test rejects with probability approaching

one — because the pretest lacks power against correlations local to zero, so you frequently fail to reject when the regressor is in fact endogenous and then report the non-robust estimate. He examines both the linear IV version, where the pretest is a test of exogeneity of a regressor, and the random-versus-fixed-effects panel version, and recommends against the two-stage procedure in both. Both of his results are for linear models, so the transfer to a nonlinear control function is an analogy rather than a theorem, but the mechanism does not depend on linearity. Report the coefficient on the residual as an estimate of the degree of endogeneity, and report both models side by side; do not let the test choose for you.

In the case that the endogenous variable is continuous and we can use a linear model for the first stage, Wooldridge 2010 textbook (chapter 18) suggested using control function approach, which is very simple:

1. Run the first stage regression (the reduced form for the endogenous variable), and get the residual. This is the first-stage **control-function** residual. Note in the many fixed effects case, we need to include all fixed effects in the first stage regression. In the case of stata, we would use “`reghdfe`” here.

A terminology note, because two different objects are easy to conflate. The **Mundlak device** means adding group *means* of covariates as regressors, as in the [Mundlak chapter](#). The **control-function residual** here is the residual from a reduced form. They are different constructions. But a third object borrows the name, and it matters below: Lin and Wooldridge (2019) call the residual from a reduced form *that itself includes the Mundlak means* — from regressing y_{it2} on z_{it} and $\bar{z}_i$ — the **Mundlak residual**, as against the **FE residual** obtained by within-transforming. So in that paper “Mundlak residual” denotes a control-function residual, not the device itself.

2. There is a useful result about which of those two residuals to use. Lin and Wooldridge (2019, appendices A.1 and A.2) show that in FE Poisson they give **numerically identical** estimates of both the structural coefficient and the coefficient on the control function. The reason is that the two residuals differ only by a term constant within unit, and the FE Poisson score is unaffected by adding anything unit-constant.

This has a practical consequence. It means the first stage does not need the thousands of unit dummies: the unit means will do, at no cost in the estimate. That is what makes a **nonlinear** first stage practical — you cannot put 2,000 dummies into a logit without an incidental-parameters problem, but you can put in 2,000 unit means. Note the result as stated is for a *linear* reduced form; see the caveat on discrete endogenous variables below.

3. Include both the endogenous variable and the control-function residual as additional regressors in the second stage, the original Poisson regression. Use “`ppmlhdfe`” here.

Note that the standard errors would be incorrect if we use the standard errors from the second stage regression. We need to use the bootstrap method to get the correct standard errors. This is not merely a precaution: Lin and Wooldridge (2019, p. 30) recommend exactly this, namely a panel bootstrap in which both estimating steps are done with each bootstrap sample. The reason is that the control-function residual is generated from the first stage, so we need to bootstrap the whole process to get the correct standard errors. In the case of clustered standard errors, we need to bootstrap the clusters.

26.3. Example

```
* control function approach
webuse website
reghdfe time phone frfam, absorb(ad female) resid
predict double u2h_fe, resid
ppmlhdfe visits time u2h_fe frfam, absorb(ad female)
```

(Visits to website)

(dropped 1 singleton observations)

(MWFE estimator converged in 3 iterations)

HDFE Linear regression	Number of obs	=	499
Absorbing 2 HDFE groups	F(2, 484)	=	29.02
	Prob > F	=	0.0000
	R-squared	=	0.3137
	Adj R-squared	=	0.2938
	Within R-sq.	=	0.1071
	Root MSE	=	2.2239

time	Coefficient	Std. err.	t	P> t	[95% conf. interval]	
phone	.2705429	.0626834	4.32	0.000	.1473778	.393708
frfam	.2539231	.0619373	4.10	0.000	.1322239	.3756222
_cons	1.406026	.2549245	5.52	0.000	.9051305	1.906921

Absorbed degrees of freedom:

Absorbed FE	Categories	- Redundant	= Num. Coefs
ad	12	0	12
female	2	1	1

(1 missing value generated)

```
Iteration 1:  deviance = 3.7303e+02  eps = .  iters = 3  tol = 1.0e-0
> 4
> min(eta) = -0.57  P
Iteration 2:  deviance = 3.5849e+02  eps = 4.06e-02  iters = 3  tol = 1.0e-0
> 4
> min(eta) = -0.59
Iteration 3:  deviance = 3.5840e+02  eps = 2.46e-04  iters = 3  tol = 1.0e-0
```

26. IV in Fixed effect Poisson with many fixed effects

```

> 4
> min(eta) = -0.60
Iteration 4: deviance = 3.5840e+02 eps = 3.17e-08  iters = 2    tol = 1.0e-0
> 4
> min(eta) = -0.60
Iteration 5: deviance = 3.5840e+02 eps = 9.72e-16  iters = 2    tol = 1.0e-0
> 5
> min(eta) = -0.60 S
Iteration 6: deviance = 3.5840e+02 eps = 1.62e-16  iters = 1    tol = 1.0e-0
> 9
> min(eta) = -0.60 S 0
-----
> -----
(legend: p: exact partial-out    s: exact solver    h: step-halving    o: epsilon
> below tolerance)
Converged in 6 iterations and 14 HDFE sub-iterations (tol = 1.0e-08)

HDFE PPML regression                               No. of obs      =      499
Absorbing 2 HDFE groups                           Residual df    =      483
                                                    Wald chi2(3)   =     169.93
Deviance = 358.4027719                             Prob > chi2    =      0.0000
Log pseudolikelihood = -993.8547331                 Pseudo R2     =      0.2976
-----

      |               Robust
visits | Coefficient std. err.      z    P>|z|    [95% conf. interval]
-----+-----
      time |   .0448103   .0448665     1.00   0.318   -.0431265   .1327471
      u2h_fe |   .0765144   .0459404     1.67   0.096   -.0135272   .166556
      frfam |   .0004108   .0205161     0.02   0.984    -.0398     .0406215
      _cons |   1.510566   .1033122    14.62   0.000    1.308078   1.713054
-----

Absorbed degrees of freedom:
-----+-----
Absorbed FE | Categories - Redundant = Num. Coefs |
-----+-----
      ad |      12      0      12 |
      female |      2      1      1 |
-----+-----

```

In this example, “time” is the endogenous variable, and “phone” is the instrument. We first run the first stage regression with “reghdfe”, and get the control-function residual “u2h_fe”. Then we include “time” and “u2h_fe” in the second stage regression with “ppmlhdfe”. “ad” and “female” are two fixed effects. The standard errors are incorrect, we need to use bootstrap to get the correct standard errors.

To verify our control function approach, we do the same with a linear model. A word of caution on “it just works”: for the *linear* model, inserting the first-stage residual is the standard 2SLS-equivalent

control function. For the *Poisson* model it is not automatic — residual inclusion is valid here under the specific structure Wooldridge (2010, ch. 18) assumes: a **continuous** endogenous variable, a **linear** first stage, and an exponential conditional mean with the control-function residual entering the index (an additive-error / Mundlak-type assumption). Outside that structure (e.g. a discrete endogenous regressor, or a different error structure) plugging in a linear residual is not generally consistent, and even when it is, the second-stage standard errors must be bootstrapped because the residual is generated (as noted above). There is a package “ivreghdfe” we use to verify our two step process.

```
clear all
webuse website
ivreghdfe visits frfam (time=phone), absorb(ad female)
* now the control function approach
reghdfe time phone frfam , absorb(ad female) resid
predict u2h_fe, resid
reghdfe visits time u2h_fe frfam , absorb(ad female)
```

(Visits to website)

(dropped 1 singleton observations)

(MWFE estimator converged in 3 iterations)

IV (2SLS) estimation

Estimates efficient for homoskedasticity only
Statistics consistent for homoskedasticity only

		Number of obs =	499	
		F(2, 484) =	8.31	
		Prob > F =	0.0003	
Total (centered) SS	=	3335.918099	Centered R2 =	0.3971
Total (uncentered) SS	=	3335.918099	Uncentered R2 =	0.3971
Residual SS	=	2011.119506	Root MSE =	2.038

visits	Coefficient	Std. err.	t	P> t	[95% conf. interval]
time	.5642802	.2123745	2.66	0.008	.1469904 .98157
frfam	-.04045	.0922873	-0.44	0.661	-.2217832 .1408833

Underidentification test (Anderson canon. corr. LM statistic): 18.494
Chi-sq(1) P-val = 0.0000

Weak identification test (Cragg-Donald Wald F statistic): 18.628
Stock-Yogo weak ID test critical values: 10% maximal IV size 16.38
15% maximal IV size 8.96

26. IV in Fixed effect Poisson with many fixed effects

20% maximal IV size 6.66
25% maximal IV size 5.53

Source: Stock-Yogo (2005). Reproduced by permission.

Sargan statistic (overidentification test of all instruments): 0.000
(equation exactly identified)

Instrumented: time
Included instruments: frfam
Excluded instruments: phone
Partialled-out: _cons
nb: total SS, model F and R2s are after partialling-out;
any small-sample adjustments include partialled-out
variables in regressor count K

Absorbed degrees of freedom:

Absorbed FE	Categories	- Redundant	= Num. Coefs
ad	12	0	12
female	2	1	1

(dropped 1 singleton observations)

(MWFE estimator converged in 3 iterations)

HDFE Linear regression	Number of obs	=	499
Absorbing 2 HDFE groups	F(2, 484)	=	29.02
	Prob > F	=	0.0000
	R-squared	=	0.3137
	Adj R-squared	=	0.2938
	Within R-sq.	=	0.1071
	Root MSE	=	2.2239

time	Coefficient	Std. err.	t	P> t	[95% conf. interval]
phone	.2705429	.0626834	4.32	0.000	.1473778 .393708
frfam	.2539231	.0619373	4.10	0.000	.1322239 .3756222
_cons	1.406026	.2549245	5.52	0.000	.9051305 1.906921

Absorbed degrees of freedom:

Absorbed FE	Categories	- Redundant	= Num. Coefs
ad	12	0	12

```

      female |           2           1           1           |
-----+-----+
(1 missing value generated)

(MWFE estimator converged in 3 iterations)

HDFE Linear regression      Number of obs   =          499
Absorbing 2 HDFE groups    F(   3,   483) =       117.11
                           Prob > F         =        0.0000
                           R-squared         =        0.7677
                           Adj R-squared     =        0.7605
                           Within R-sq.     =        0.4211
                           Root MSE       =        1.9996

-----+-----+
      visits | Coefficient  Std. err.      t    P>|t|     [95% conf. interval]
-----+-----+
      time   |   .5642802   .2083276     2.71  0.007     .15494   .9736205
      u2h_fe |   .1827169   .2122987     0.86  0.390    -.2344262 .5998601
      frfam   |   -.04045    .0905287    -0.45  0.655    -.2183288 .1374288
      _cons   |   3.357879   .4427337     7.58  0.000     2.487957 4.227801
-----+-----+

Absorbed degrees of freedom:
-----+-----+
Absorbed FE | Categories - Redundant = Num. Coefs |
-----+-----+
      ad     |         12         0         12      |
      female |          2          1          1      |
-----+-----+

```

We can see that “ivreghdfe” gives the same result as our two step process. The standard errors are wrong. Now we do bootstrapping.

26.4. Bootstrapping standard errors

Suppose we need clustered standard errors on “ad”.

```

* bootstrap, suppose we need to cluster by "ad"
* Note: bootstrap's idcluster(newid) below gives each resampled COPY of a
* cluster its own unique id -- necessary because a cluster can be drawn more
* than once within a replicate, and those duplicates must be treated as
* distinct groups. The regressions inside this program must absorb "newid",
* not the original "ad": absorbing "ad" would pool resampled duplicates of
* the same cluster back into a single fixed effect, which violates the

```

26. IV in Fixed effect Poisson with many fixed effects

```
* cluster-bootstrap logic (verified empirically: within a replicate, the
* number of distinct "ad" values is consistently below the true cluster
* count, while "newid" always has exactly that many distinct values).
clear all
capture program drop ppmlhdfe_cf
program ppmlhdfe_cf, rclass
    reghdfe time phone frfam , absorb(newid female) resid
    predict double u2h_fe, resid
    ppmlhdfe visits time u2h_fe frfam , absorb(newid female)
    return scalar b_time = _b[time]
    drop u2h_fe
    xtset, clear
end

webuse website

xtset, clear
bootstrap r(b_time) , reps(100) seed(123) cluster(ad) idcluster(newid) : ppmlhdfe_cf
```

(Visits to website)

(running ppmlhdfe_cf on estimation sample)

```
Bootstrap replications (100): .....10.....20.....30.....40.....
> ....50.....60.....70.....80.....90.....100 done
```

```
Bootstrap results                                     Number of obs = 499
                                                    Replications = 100
```

```
Command: ppmlhdfe_cf
       _bs_1: r(b_time)
```

(Replications based on 12 clusters in ad)

	Observed	Bootstrap			Normal-based	
	coefficient	std. err.	z	P> z	[95% conf. interval]	
_bs_1	.0448103	.0476634	0.94	0.347	-.0486083	.1382288

A troubleshooting note: we need to run `xtset, clear` before each regression inside the bootstrap. The reason is that Stata's cluster bootstrap generates a new id variable for the resampled clusters and leaves the panel `xtset` to that id; if it is not cleared, a later regression can pick up the stale panel setting even when it does not use the id. Clearing it each iteration avoids this.

26.5. How to do it with R

R has similar packages https://cran.r-project.org/web/packages/fixest/vignettes/fixest_walkthrough.html.

We can use “fixest” and “fepois” to do the same thing as “reghdfe” and “ppmlhdfe”. Again, at the moment, we need to manually bootstrap the two stages for correct standard errors.

26.6. What if the endogenous variable is binary?

Wooldridge treats this case only in a section on extensions and future directions, and declines to vouch for the approximation, so it is the least settled material in this chapter. Both procedures above assume otherwise: Procedure 3’s first stage is *linear*, and the companion probit procedure in the same paper states outright that the endogenous variable is assumed continuous. The binary case appears only in that paper’s final section, on extensions and future directions.

What is offered there is the **generalized residual** from a reduced-form probit in $w_{it} = (1, z_{it}, \bar{z}_i)$ — the instrument, the time-varying exogenous controls, and their unit means:

$$\widehat{gr}_{it2} = y_{it2} \lambda(w_{it} \hat{\theta}_2) - (1 - y_{it2}) \lambda(-w_{it} \hat{\theta}_2) \quad (26.1)$$

where λ is the inverse Mills ratio. For a **logit** first stage the generalized residual collapses algebraically to the plain response residual $y_{it2} - \hat{p}_{it2}$, so the familiar “predicted probability, then subtract” recipe is the generalized residual, not an approximation to it.

The reason for that collapse is worth knowing, because it tells you when the shortcut stops working. The generalized residual is the score of the log-likelihood with respect to the index. For a logit that score is $y_{it2} - \hat{p}_{it2}$. For a probit it is the inverse Mills expression above, which is a genuinely different variable: in a simulated sample of 5,000 the two differ by as much as 1.59, while correlating at 0.999. So the subtraction recipe is exact for a logit and simply wrong for a probit, and the high correlation means the mistake will not announce itself.

One practical warning follows. What you want is the **response** residual — predict the probability, then subtract — and **predict** on its own returns a probability, not a residual. The other residuals on offer after a binary model are not the generalized residual either: Pearson residuals divide through by $\sqrt{\hat{p}(1 - \hat{p})}$, and deviance residuals are a different transformation again. In the same simulation those correlate with the response residual at 0.991 and 0.995 — close enough to look correct if you glance at them, and the wrong variable to put in the second stage.

A note on the name, because it misleads. This construction is usually filed under **correlated random effects** (CRE), and that is how to find it in the literature — see the **CRE chapter**. The name describes the *assumption*: heterogeneity is modelled as a function of the group means and allowed to correlate with the regressors, instead of being treated as a free parameter. It does not describe the estimator. The first stage below is a **pooled logit** with the group means added as regressors. No random effect is estimated, no variance component appears, and nothing is integrated out. Below it is labelled “CRE (Mundlak device)” so that both names stay visible: the first tells you where to look it up, the second tells you what is actually computed.

Two caveats should be stated plainly. The construction is introduced as “a simple example”, and is immediately followed by: “The Mundlak device, or Chamberlain’s version of it, might work reasonably well, but neither might be flexible enough. We leave investigations into the quality of CF approximations in discrete cases to future research.” So this is the construction to use, but it is not a settled result. If you are worried the specification is too rigid, the same paper suggests adding nonlinear functions of the conditioning set — in particular interactions of the covariates with the unit means, and with the control function itself — or letting the variance depend on the unit means via a heteroskedastic probit.

26.6.1. Several fixed effects in the first stage

The expression $w_{it} = (1, z_{it}, \bar{z}_i)$ above is written for a single unobserved effect. If the second stage absorbs more than one set of fixed effects — units, years, and something else — the first stage needs the treatment set out in the [Mundlak chapter](#), where the CRE construction is the Mundlak device: take means over the *one* dimension you are eliminating, of everything else in the model, retained dummies included. In practice that means unit means of the instrument and the time-varying controls, plus unit means of the year and other dummies, with those dummies themselves kept in levels. You do not want one set of means per dimension; that construction is exact only for orthogonal designs, which observational panels do not have.

26.6.2. Why not just put the dummies in the first stage?

Because a logit with many thin groups is badly biased, and this is the whole reason the Mundlak route is needed. The size of the problem depends on observations per group, not on the number of groups. In a simulation with the total sample held near 12,000 and a true slope of 1, a logit with group dummies returns about 2.0 at two observations per group — matching the known result that the fixed-effects logit slope is inflated by a factor of two in that case — about 1.37 at four observations per group, 1.15 at six, 1.04 at twenty-five, and essentially the truth by a hundred. A panel with a few thousand units and four or five observations each therefore sits in the worst part of that range, while year effects with hundreds of observations apiece are harmless and can simply be estimated.

The CRE (Mundlak device) first stage is better but not unbiased: in the same simulation it ran roughly ten per cent low at small group sizes, because integrating a logit over normal heterogeneity does not give a logit. That is the same approximation error the previous section’s caveat refers to, arriving from a different direction.

26.6.3. A worked example

The `website` data used earlier does not have the shape that causes trouble — its fixed effects have 13 and 2 levels, with roughly 90 observations each, so a logit with dummies would be perfectly well behaved there. This simulates something more awkward instead: a count outcome, a binary endogenous regressor, an instrument, and two thousand units with between two and six observations each.


```

preserve
clear
set seed 20260731
set obs 2000
gen i = _n
gen double ci = rnormal()*0.5
gen T = 2 + int(runiform()*5)      /* unbalanced: 2-6 observations per unit */
expand T
bysort i: gen t = _n

gen double z = rnormal()          /* the instrument */
gen double e = rnormal()
gen double u = 0.8*e + rnormal()*0.6 /* correlated with e, hence the endogeneity */
gen byte d = (-0.2 + 0.9*z + ci + e > 0)
gen double mu = exp(0.3*d + ci + 0.5*u)
gen y = rpoisson(mu)              /* true coefficient on d is 0.3 */

* 1. naive FE Poisson, ignoring the endogeneity
quietly ppmlhdfe y d, absorb(i)
scalar b_naive = _b[d]

* 2. control function with a CRE (Mundlak device) logit first stage: a POOLED
*   logit with the group mean of the instrument added. No random effect is
*   estimated. For a logit the generalized residual is just d - phat.
bysort i: egen double z_i = mean(z)
quietly logit d z z_i
predict double phat, pr
gen double gr = d - phat
quietly ppmlhdfe y d gr, absorb(i)
scalar b_cre = _b[d]

* 3. control function with a linear first stage, unit FE absorbed exactly
quietly reghdfe d z, absorb(i) resid
predict double v_lin, resid
quietly ppmlhdfe y d v_lin, absorb(i)
scalar b_lin = _b[d]

di "true coefficient on d          = 0.300000"
di "naive FE Poisson              = " %8.6f b_naive
di "CF, CRE (Mundlak) logit       = " %8.6f b_cre
di "CF, linear FE first stage     = " %8.6f b_lin
restore

```

Number of observations (_N) was 0, now 2,000.

26. IV in Fixed effect Poisson with many fixed effects

(5,969 observations created)

true coefficient on d	= 0.300000
naive FE Poisson	= 0.747206
CF, CRE (Mundlak) logit	= 0.293829
CF, linear FE first stage	= 0.277644

The naive estimate is inflated by a factor of about two and a half; both control-function versions recover something close to the truth, and they agree with each other. That agreement is the point. When the endogenous variable is binary I would report both: the linear first stage rests on the exact result in point 2 above and drops no observations, while the logit version is better specified for a binary variable but rests on the unsettled argument in this section. If they diverge materially, that is the most important thing you have learned, and it belongs in the write-up.

The standard errors in the chunk above are wrong, for the reason given earlier — the residual is generated. Bootstrap both stages together, as in the earlier section.

Systematic treatment: [R](#) · [Julia](#).

Part VI.

Causal Inference Methods

27. Using machine learning for causal effect in observational study

27.1. A simulation for an OLS model

In an observational study, we need to assume we have the functional form to get causal effect estimated correctly, in addition to the assumption of treatment being exogenous.

```
library(MASS)
library(ggplot2)
library(dplyr)
library(tmle)
library(glmnet)
set.seed(366)

nobs <- 2000
xw <- .8
xz <- .5
zw <- .6
nrow <- 3
ncol <- 3
covarMat = matrix( c(1^2, xz^2, xw^2, xz^2, 1^2, zw^2, xw^2, zw^2, 1^2) , nrow=ncol , ncol=nobs)

mu <- rep(0,3)
rawvars <- mvrnorm(n=nobs, mu=mu, Sigma=covarMat)
df <- as_tibble(rawvars, .name_repair = "minimal")
names(df) <- c('x','z','w')
df <- df %>%
  # A small constant inside each log() keeps log(v^2) from diverging to
  # large negative values when v is close to 0 (e.g. log(w^2) alone ranges
  # down to about -15 here), which otherwise pushes the propensity score
  # P(A=1|W) toward 0/1 for those units and violates positivity.
  mutate(log.x=log(x^2 + 0.01), log.z=log(z^2 + 0.01), log.w=log(w^2 + 0.01), z.sqr=z^2, w.sqr=w^2)
  mutate(g.var= log.w + rnorm(nobs)) %>%
  mutate(A = rbinom(nobs, 1, 1/(1+exp((g.var)))) %>%
  mutate(y0=rnorm(nobs) + log.x) %>%
  mutate(tau.true = 2 + rnorm(nobs), y1=y0+tau.true, treat=A, y = treat*y1 + (1-treat)*y0)
lm1 <- lm(y ~ A + log.w + log.x , data=df)
summary(lm1)
```

27. Using machine learning for causal effect in observational study

Call:

```
lm(formula = y ~ A + log.w + log.x, data = df)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-4.5207	-0.8338	-0.0089	0.8386	4.1534

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.01311	0.04774	0.275	0.784
A	1.95487	0.06703	29.166	<2e-16 ***
log.w	0.02835	0.01945	1.458	0.145
log.x	1.00090	0.01748	57.264	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.253 on 1996 degrees of freedom

Multiple R-squared: 0.6746, Adjusted R-squared: 0.6741

F-statistic: 1380 on 3 and 1996 DF, p-value: < 2.2e-16

```
lm2 <- lm(y ~ A , data=df)
summary(lm2)
```

Call:

```
lm(formula = y ~ A, data = df)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-6.1667	-1.4416	0.1372	1.4760	6.1625

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.75733	0.07539	-10.04	<2e-16 ***
A	1.47007	0.09571	15.36	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.077 on 1998 degrees of freedom

Multiple R-squared: 0.1056, Adjusted R-squared: 0.1052

F-statistic: 235.9 on 1 and 1998 DF, p-value: < 2.2e-16

```
lm3 <- lm(y ~ A + w, data=df)
summary(lm3)
```

```
Call:
lm(formula = y ~ A + w, data = df)

Residuals:
    Min       1Q   Median       3Q      Max
-6.1657 -1.4531  0.1338  1.4780  6.0957

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.75631    0.07541 -10.030  <2e-16 ***
A             1.46876    0.09573  15.343  <2e-16 ***
w            -0.03874    0.04464  -0.868    0.386
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.077 on 1997 degrees of freedom
Multiple R-squared:  0.1059,    Adjusted R-squared:  0.105
F-statistic: 118.3 on 2 and 1997 DF,  p-value: < 2.2e-16
```

```
lm4 <- lm(y ~ A + w + x, data=df)
summary(lm4)
```

```
Call:
lm(formula = y ~ A + w + x, data = df)

Residuals:
    Min       1Q   Median       3Q      Max
-6.167 -1.445  0.147  1.473  6.205

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.75360    0.07539  -9.996  <2e-16 ***
A             1.46387    0.09573  15.292  <2e-16 ***
w            -0.09981    0.05764  -1.732   0.0835 .
x             0.10281    0.06141   1.674   0.0943 .
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.076 on 1996 degrees of freedom
Multiple R-squared:  0.1072,    Adjusted R-squared:  0.1059
F-statistic: 79.88 on 3 and 1996 DF,  p-value: < 2.2e-16
```

In this example, treatment assignment process is determined by logged w , and outcome is determined by logged x and treatment. However, what we observe is w and x . In observational studies, this happens all the time. In fact, this is an ideal situation, that we observe variables that are determinants of outcome, although we are not sure about the functional form that determines the

27. Using machine learning for causal effect in observational study

outcome. However, this example shows that unless we have observed exactly the factors themselves (in this case logged x , w , which determines the DGP), we have biased estimates of the true treatment effect.

Model 1 is the only model with reasonable estimate of treatment effect (which is 2 in this case). Model 2 is a model with endogeneity: A is correlated with the missing variable logged x . Model 3 and 4 we have x and w , but not logged, therefore still biased.

The lesson here is the functional form does matter. However, we have no way of knowing the functional form. What can we do here?

```
# Algorithm set trimmed to 7 fast learners for render speed.
# A production analysis would also include SL.randomForest, SL.gbm, SL.gam.
Q.SL.library <- c("SL.glmnet","SL.glm","SL.glm.interaction", "SL.rpart","SL.bayesglm","SL.step
g.SL.library <- c("SL.glmnet","SL.glm","SL.glm.interaction", "SL.rpart","SL.bayesglm","SL.step

# tmle1 uses x and w; tmle2 additionally includes z. Note z is not in the true
# treatment or outcome DGP -- it is only correlated with x and w -- so adding it
# should not change the estimate much.
tmle1 <- tmle(Y = df$y, A = df$treat, W = df[,c('x','w')], g.SL.library = g.SL.library , Q.SL.
tmle1
```

Marginal mean under treatment (EY1)

Parameter Estimate: 0.92554
Estimated Variance: 0.0033412
p-value: <2e-16
95% Conf Interval: (0.81225, 1.0388)

Marginal mean under comparator (EY0)

Parameter Estimate: -1.0196
Estimated Variance: 0.0035588
p-value: <2e-16
95% Conf Interval: (-1.1365, -0.90269)

Additive Effect

Parameter Estimate: 1.9451
Estimated Variance: 0.004266
p-value: <2e-16
95% Conf Interval: (1.8171, 2.0732)

Additive Effect among the Treated

Parameter Estimate: 1.8969
Estimated Variance: 0.0052492
p-value: <2e-16
95% Conf Interval: (1.7549, 2.0389)

Additive Effect among the Controls

Parameter Estimate: 2.0351


```

Estimated Variance: 0.0051986
      p-value: <2e-16
95% Conf Interval: (1.8938, 2.1765)

```

```

tmle2 <- tmle(Y = df$y, A = df$treat, W = df[,c('x', 'w', 'z')], g.SL.library = g.SL.library , (
tmle2

```

Marginal mean under treatment (EY1)

```

Parameter Estimate: 0.9072
Estimated Variance: 0.0033185
      p-value: <2e-16
95% Conf Interval: (0.79429, 1.0201)

```

Marginal mean under comparator (EY0)

```

Parameter Estimate: -1.0068
Estimated Variance: 0.003566
      p-value: <2e-16
95% Conf Interval: (-1.1238, -0.88975)

```

Additive Effect

```

Parameter Estimate: 1.914
Estimated Variance: 0.004266
      p-value: <2e-16
95% Conf Interval: (1.786, 2.042)

```

Additive Effect among the Treated

```

Parameter Estimate: 1.8705
Estimated Variance: 0.0052652
      p-value: <2e-16
95% Conf Interval: (1.7283, 2.0127)

```

Additive Effect among the Controls

```

Parameter Estimate: 1.9817
Estimated Variance: 0.0051984
      p-value: <2e-16
95% Conf Interval: (1.8404, 2.123)

```

We use [Mark van der Laan's TMLE method](#). It uses [SuperLearner](#) as the initial estimator. It's an ensemble of multiple machine learning algorithms. Therefore it does not need to assume the functional form of the DGP. Even if we don't have the variables that determines the DGP of outcome, if we observe some functions (even nonlinear functions) of these variables, we can still get reasonable estimates of the treatment effect.

In this example, we used multiple popular machine learning algorithms in modeling both treatment assignment process and the outcome process. The first TMLE model is with x and w (note not the logged x and w which are in the true DGP), the second one with an additional variable z.

27. Using machine learning for causal effect in observational study

It seems that TMLE results are less biased than the linear models with x and w . It may not be better than the linear model with logged x and w , but in empirical studies, we often cannot assume we have the variables in the DGP, but only some proxy of the variables in the DGP. I'll do more simulations to see whether TMLE does perform better in the situation that we are not sure about the functional form. We should expect that is the case.

The simple `tmle` package example used here is for a binary treatment. TMLE itself is not limited to binary treatments: there are TMLE estimators for continuous, multivalued, longitudinal, stochastic, and survival settings (e.g. the `ltmle` and `lmtp` packages), though which estimand and package you use differs by setting.

It's about time we embrace machine learning techniques into studies of causal effect in observational studies.

Systematic treatment: [R](#) · [Julia](#).

28. Mediation analysis in R and Stata

28.1. Mediation analysis

Traditionally mediation model can be represented in the following equations:

$$Y = aX + bM + \epsilon_1 \quad (28.1)$$

$$M = cX + \epsilon_2 \quad (28.2)$$

That is, we'd like to study the effect of X on Y , and we see the effect can be a direct effect, and an indirect effect, through M .

Baron and Kenny's (<http://davidakenny.net/cm/mediate.htm>) method is done in four steps. Modern approach tends to use SEM (structural equation modeling) to model these two equations directly.

R's lavaan and Stata's sem commands are powerful tools.

A simple example:

```
library(lavaan)
set.seed(1234)
n <- 10000
X <- rnorm(n)
M <- 0.5*X + rnorm(n)
Y <- 0.7*M + 0.3*X + rnorm(n)
Data <- data.frame(X = X, Y = Y, M = M)
model <- '
    Y ~ a*X + b*M
    M ~ c*X
    # direct effect (a)
    # indirect effect (b*c)
    bc := b*c
    # total effect
    total := a + (b*c)
'
fit <- sem(model, data = Data)
summary(fit)
```

lavaan 0.7-2 ended normally after 1 iteration

Estimator

ML

Optimization method

NLMINB

28. Mediation analysis in R and Stata

Number of model parameters 5

Number of observations 10000

Model Test User Model:

Test statistic 0.000

Degrees of freedom 0

Parameter Estimates:

Standard errors	Standard
Information	Expected
Information saturated (h1) model	Structured

Regressions:

		Estimate	Std.Err	z-value	P(> z)
Y ~					
X	(a)	0.286	0.011	25.233	0.000
M	(b)	0.699	0.010	70.384	0.000
M ~					
X	(c)	0.511	0.010	50.161	0.000

Variances:

	Estimate	Std.Err	z-value	P(> z)
.Y	0.998	0.014	70.711	0.000
.M	1.012	0.014	70.711	0.000

Defined Parameters:

	Estimate	Std.Err	z-value	P(> z)
bc	0.357	0.009	40.849	0.000
total	0.643	0.012	51.956	0.000

Now let's do it in stata.

```
set obs 10000
gen x= rnormal()
gen m= .5*x + rnormal()
gen y= .7*m + .3*x +rnormal()
sem (y <- x m) (m <- x)
estat teffects
```

Number of observations (_N) was 0, now 10,000.

Endogenous variables

Observed: y m

Exogenous variables

Observed: x

Fitting target model:

Iteration 0: Log likelihood = -42567.918

Iteration 1: Log likelihood = -42567.918

Structural equation model

Number of obs = 10,000

Estimation method: ml

Log likelihood = -42567.918

		OIM		z	P> z	[95% conf. interval]	
		Coefficient	std. err.				
Structural y							
	m	.7136545	.0101594	70.25	0.000	.6937424	.7335666
	x	.2971221	.0113774	26.12	0.000	.2748228	.3194215
	_cons	-.0000391	.0101082	-0.00	0.997	-.0198508	.0197725
m							
	x	.5038605	.0100014	50.38	0.000	.4842581	.5234628
	_cons	.0187204	.0099478	1.88	0.060	-.0007769	.0382177
var(e.y)		1.021391	.0144446			.9934684	1.050098
var(e.m)		.9895863	.0139949			.9625336	1.017399

LR test of model vs. saturated: chi2(0) = 0.00

Prob > chi2 = .

Direct effects

		OIM		z	P> z	[95% conf. interval]	
		Coefficient	std. err.				
Structural y							
	m	.7136545	.0101594	70.25	0.000	.6937424	.7335666
	x	.2971221	.0113774	26.12	0.000	.2748228	.3194215
m							

28. Mediation analysis in R and Stata

x		.5038605	.0100014	50.38	0.000	.4842581	.5234628
---	--	----------	----------	-------	-------	----------	----------

Indirect effects

			OIM				
		Coefficient	std. err.	z	P> z	[95% conf. interval]	
Structural							
y							
	m	0	(no path)				
	x	.3595823	.0087834	40.94	0.000	.3423672	.3767974
m							
	x	0	(no path)				

Total effects

			OIM				
		Coefficient	std. err.	z	P> z	[95% conf. interval]	
Structural							
y							
	m	.7136545	.0101594	70.25	0.000	.6937424	.7335666
	x	.6567044	.0124172	52.89	0.000	.6323672	.6810417
m							
	x	.5038605	.0100014	50.38	0.000	.4842581	.5234628

The above examples should have direct effect of .3 and indirect effect of .35, and total effect of .65.

28.2. Causal Mediation analysis

The traditional mediation analysis has been criticized for the lack of causal interpretation. Without manipulation of the mediator, it is hard to interpret the effects causally, because even if the treatment is from random experiments, the mediator is often not. Therefore there could be an unmeasured confounder that is causing both M and Y .

R's "mediation" package is for causal mediation analysis. It uses simulation to estimate the causal effects of treatment, under assumptions of sequential ignorability.

i Identifying assumptions for natural (in)direct effects

Before reading the `mediate()` output causally, the following must hold (this is “sequential ignorability” plus the standard support and link conditions):

1. **Treatment ignorability** — no unmeasured treatment–outcome confounding given covariates W .
2. **Mediator ignorability** — no unmeasured mediator–outcome confounding given W and treatment.
3. **No treatment-induced confounding** — no variable affected by the treatment confounds the mediator–outcome relationship (this is the strong, often-violated condition).
4. **Positivity** — both treatment and mediator have positive probability across W .
5. **Consistency** — observed outcomes equal the corresponding potential outcomes.
6. **Cross-world independence** — the condition underlying natural effects, which no single experiment can test.

Mediator ignorability and the no-treatment-induced-confounding condition are the ones that usually break, because the mediator is rarely randomized.

It estimates the following quantities:

$$\tau_i = Y_i(1, M_i(1)) - Y_i(0, M_i(0)) \quad (28.3)$$

This is the total treatment effect, which is to say, what’s the change in Y if we change each unit from control to treated, hypothetically?

Then this can be decomposed into the causal mediation effects:

$$\delta_i(t) = Y_i(t, M_i(1)) - Y_i(t, M_i(0)) \quad (28.4)$$

for treatment status $t = 0, 1$. This is to say, given treatment status, what’s the mediation effect?

$$\eta_i(t) = Y_i(1, M_i(t)) - Y_i(0, M_i(t)) \quad (28.5)$$

for treatment status $t = 0, 1$. This is to say, given mediator status for each treatment status, what’s the direct effect?

Therefore there are four quantities estimated, direct and mediation effect for treated and control.

R’s “mediation” needs users to feed two models, outcome model and mediation model.

If we study the same data, we would expect it returns the same estimates as the traditional methods. However, the causal mediation models can be much more flexible in outcome and mediation models.

```
library(mediation)
med.fit <- lm(M ~ X, data=Data)
summary(med.fit)
```

28. Mediation analysis in R and Stata

Call:

```
lm(formula = M ~ X, data = Data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-4.1134	-0.6972	0.0086	0.6837	3.7309

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.007814	0.010060	-0.777	0.437
X	0.510954	0.010187	50.156	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.006 on 9998 degrees of freedom

Multiple R-squared: 0.201, Adjusted R-squared: 0.2009

F-statistic: 2516 on 1 and 9998 DF, p-value: < 2.2e-16

```
out.fit <- lm(Y ~ X + M, data=Data)
summary(out.fit)
```

Call:

```
lm(formula = Y ~ X + M, data = Data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-4.0680	-0.6771	-0.0081	0.6643	3.8304

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.004170	0.009992	0.417	0.676
X	0.285582	0.011319	25.229	<2e-16 ***
M	0.699006	0.009933	70.373	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9991 on 9997 degrees of freedom

Multiple R-squared: 0.4734, Adjusted R-squared: 0.4733

F-statistic: 4494 on 2 and 9997 DF, p-value: < 2.2e-16

```
set.seed(123)
med.out <- mediate(med.fit, out.fit, treat = "X", mediator = "M", sims = 100)
summary(med.out)
```


Causal Mediation Analysis

Quasi-Bayesian Confidence Intervals

	Estimate	95% CI Lower	95% CI Upper	p-value
ACME	0.35684	0.34208	0.37741	< 2.2e-16 ***
ADE	0.28540	0.26522	0.30604	< 2.2e-16 ***
Total Effect	0.64224	0.62091	0.66752	< 2.2e-16 ***
Prop. Mediated	0.55530	0.53571	0.57897	< 2.2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Sample Size Used: 10000

Simulations: 100

28.3. Binary outcome

For example, in the case of binary outcome, the traditional approach will have difficulties. We can estimate the outcome model and mediator model jointly, but the total effects are not easy to decompose into direct and indirect effect (see Imai et al, page 320 <https://imai.fas.harvard.edu/research/files/BaronKenny.pdf>).

The causal mediation analysis framework is much more general.

```
library(mediation)
data(framing)
med.fit <- lm(emo ~ treat + age + educ + gender + income, data = framing)
out.fit <- glm(cong_mesg ~ emo + treat + age + educ + gender + income,
               data = framing, family = binomial("probit"))
set.seed(123)
med.out <- mediate(med.fit, out.fit, treat = "treat", mediator = "emo",
                  robustSE = TRUE, sims = 100)
summary(med.out)
```

Causal Mediation Analysis

Quasi-Bayesian Confidence Intervals

	Estimate	95% CI Lower	95% CI Upper	p-value
ACME (control)	0.081341	0.037680	0.122101	<2e-16 ***
ACME (treated)	0.082417	0.036606	0.126460	<2e-16 ***
ADE (control)	0.018299	-0.105898	0.158469	0.72
ADE (treated)	0.019375	-0.115315	0.172002	0.72

28. Mediation analysis in R and Stata

```
Total Effect          0.100716    -0.028036    0.259054    0.24
Prop. Mediated (control) 0.632609   -13.705651   4.305615    0.24
Prop. Mediated (treated) 0.661761   -12.422330   4.036601    0.24
ACME (average)         0.081879    0.037143    0.124080   <2e-16 ***
ADE (average)          0.018837   -0.110606    0.165198    0.72
Prop. Mediated (average) 0.647185   -13.063991   4.171108    0.24
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Sample Size Used: 265

Simulations: 100

“mediation” package has more functionalities, such as multilevel, interaction of treatment and mediator, etc.

Stata’s `sem` and `gsem` commands can model different situations, but the direct effect and indirect effects are not easy to compute, especially when you have binary outcome, or other non-continuous outcome situations. They are not designed for causal mediation analysis.

28.4. Going further: formal causal mediation

The `mediation` package above implements sequential ignorability — a specific set of identifying assumptions. The [causal mediation chapter](#) covers the assumptions more formally, introduces the controlled direct effect (CDE) and natural direct/indirect effects (NDE/NIE) using potential-outcome notation, and discusses identification with unmeasured mediator–outcome confounding. If your setting involves a non-binary outcome, treatment–mediator interaction, or you need a doubly-robust estimator, start there.

Systematic treatment: [R](#) · [Julia](#).

29. Endogeneity with censored variable

29.1. Bunching for the outcome variable

Recently I read this paper, Caetano et al (2020) https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3699644. It's about bunching.

Bunching is a phenomenon where we see a mass point in the distribution of some variable. For example, Saez (2010) studies a mass point in the distribution of earned income at the kink point of the tax schedule. This is because people bunch at the kink point to take advantage of the tax break.

Bunching is similar to regression discontinuity. In this blog, <https://blogs.worldbank.org/en/impactevaluations/ready-set-bunch>, “As noted by Kleven (2016), regression discontinuity (RD) is a close cousin of bunching estimators. In regression discontinuity, we maintain the assumption that there is no such “manipulation” as described above. Bunching relaxes this assumption — instead, we estimate the fraction of manipulators by estimating what densities of individuals would have been without manipulation, that is the “manipulation-free counterfactual”. With both the observed and manipulation-free counterfactual distributions of individuals estimated, it may be possible to compare the two distributions to recover the fraction of individuals who manipulated.”

The main idea is to compare the bunching with the counterfactual distribution without bunching to estimate elasticities, for example, how people react to tax rate.

29.2. Bunching for the treatment variable

29.2.1. A test of endogeneity

Caetano (2015) showed that, if the distribution of T_i has bunching at $T_i = \bar{t}$, it is possible to test the exogeneity of T_i . When one compares the outcome of observations at the bunching point and those around it, the treatment itself is very similar. Therefore, there cannot be more than a marginal difference in the outcome that is due to treatment variation, since the treatment hardly varies.

Caetano (2015)'s estimation can be done in a two step process:

- (1) estimate $E[Y_i|T_i = \bar{T}, X_i]$ non-parametrically
- (2) do a local linear regression $E[Y_i|T_i = \bar{T}, X_i] - Y_i$ onto T_i at $\bar{T}$, using only observations such that $T_i > \bar{T}$.

The approach is known as the Discontinuity Test.

29.2.2. Dummy test

Caetano and Maheshri (2018) proposed a simpler test, which is simply regressing Y_i on T_i , X_i and a dummy variable for $T_i > \bar{T}$. The coefficient of the dummy variable is the test statistic. If the coefficient is significant, which is to say there is a discontinuity on outcome, then T_i is endogenous.

This is simpler, but with linear functional form assumption.

The logic of the endogeneity test is that around the bunching point, the treatment is very similar, so the outcome should be similar. If there is a discontinuity in the outcome, then it has to be one of two reasons. One is that treatment has discontinuous effect on outcome around the bunching point. Usually we can reasonably rule out that possibility. Second is that there is an unobserved variable that is causing both the treatment and the outcome. That unobserved variable has a discontinuous effect on the outcome. If that's the case, then the treatment is endogenous.

29.2.3. Treatment effect estimation

Caetano et al (2020) proposed a method to estimate the treatment effect when there is bunching. It seems to me there is no reason that we cannot use it for the censoring situation. There are a lot of situations that we have censored treatment variable, when we study continuous treatment variable.

Suppose we have

$$Y = \beta T + Z\gamma + \delta\eta + \epsilon \quad (29.1)$$

where T is the treatment variable, η is the unobserved variable that is causing both the treatment and the outcome, Z is the control variables, and ϵ is the error term. We only observe Y, T, Z .

Suppose T^* is the latent variable that is censored at $T^* = 0$. We observe $T = \max(T^*, 0)$. We assume that T^* is continuous over its support.

$$T^* = Z\pi + \eta \quad (29.2)$$

From the two equations, we can derive:

$$E[Y|T, Z] = (\beta + \delta)T + Z(\gamma - \pi\delta) + \delta E[T^*|T^* \leq 0, Z]1(T = 0) \quad (29.3)$$

or

$$E[Y|T, Z] = \beta T + Z(\gamma - \pi\delta) + \delta(T + E[T^*|T^* \leq 0, Z]1(T = 0)) \quad (29.4)$$

Therefore we can “correct” for endogeneity by simply including an additional term $\delta(T + E[T^*|T^* \leq 0, Z]1(T = 0))$ in the regression.

However, the prediction of $\delta(T + E[T^*|T^* \leq 0, Z]1(T = 0))$ is out of sample. We have to make some additional assumptions for the distribution of T^* . For example, we can assume that T^* is normally distributed. Or, we can assume it is symmetric. Then we can use the right tail to estimate the left tail.

29.3. Simulation

Let's do a simulation to see how it works.

```
library(MASS)
library(ggplot2)
library(dplyr)
# set seed
set.seed(123)
# number of observations
n=1000
# generate Z and make it dataframe
Z=as.data.frame(mvrnorm(n=n, mu=c(0,0), Sigma=matrix(c(1,0.5,0.5,1),2,2)))
data <- Z
names(data) <- c("Z1", "Z2")
# generate eta
# eta is normally distributed
data$eta=rnorm(n, 1,1)
# generate T*
# T* is normally distributed
# Y depends on the OBSERVED (censored) treatment T, not the latent T_star --
# matching the structural equation Y = beta*T + Z*gamma + delta*eta + epsilon
# above. If Y depended on T_star directly, T=0 would not fully "block" the
# latent variable's effect on Y, and the treatment effect of T would not even
# be well defined.
data <- data %>%
  mutate(T_star=Z1 + 2*Z2 + 1.5*eta) %>%
  mutate(T=ifelse(T_star>0, T_star, 0)) %>%
  mutate(Y=2*T + 0.5*Z1 + 0.5*Z2 + 2*eta + rnorm(n, 0, 1))

# regress Y on T, Z1, Z2
lm1 <- lm(Y ~ T + Z1 + Z2, data=data)
summary(lm1)
```

Call:

```
lm(formula = Y ~ T + Z1 + Z2, data = data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-8.2706	-0.8700	0.1009	1.0955	4.0476

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.42394	0.09588	-4.422	1.09e-05 ***
T	3.14972	0.03906	80.631	< 2e-16 ***
Z1	-0.29574	0.06667	-4.436	1.02e-05 ***

29. Endogeneity with censored variable

```
Z2          -1.14069      0.08060 -14.153  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 1.639 on 996 degrees of freedom
Multiple R-squared:  0.9343,    Adjusted R-squared:  0.9341
F-statistic:  4719 on 3 and 996 DF,  p-value: < 2.2e-16
```

Without including *eta*, the effect estimate is biased. Let's do a dummy test.

```
lm2 <- lm(Y ~ T + Z1 + Z2 + I(T>0), data=data)
summary(lm2)
```

Call:

```
lm(formula = Y ~ T + Z1 + Z2 + I(T > 0), data = data)
```

Residuals:

Min	1Q	Median	3Q	Max
-8.0768	-0.8363	0.0406	0.9126	4.8213

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.68077	0.11961	-14.052	<2e-16 ***
T	3.06030	0.03570	85.728	<2e-16 ***
Z1	-0.52092	0.06189	-8.416	<2e-16 ***
Z2	-1.48038	0.07601	-19.476	<2e-16 ***
I(T > 0)TRUE	2.10995	0.13884	15.197	<2e-16 ***

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 1.477 on 995 degrees of freedom
Multiple R-squared:  0.9466,    Adjusted R-squared:  0.9464
F-statistic:  4414 on 4 and 995 DF,  p-value: < 2.2e-16
```

It shows that the dummy variable is significant. It means that T is endogenous.

29.3.1. Correcting for endogeneity

The correction term is $E[T^* | T^* \leq 0, Z]$, on the scale of the *treatment* T^* — so we need to model T^* given Z , not Y given T and Z . Let's assume T^* is normally distributed given Z : $T^* | Z \sim N(Z\pi, \sigma^2)$. Since we observe $T = T^*$ whenever $T^* > 0$ and $T = 0$ otherwise, this is exactly a (left-)censored (Tobit) regression of T on Z , using **all** observations (not just the right tail): the censored observations still inform π and σ through the likelihood, they just don't reveal T^* 's value directly. Given $\hat{\pi}$ and $\hat{\sigma}$, the standard truncated-normal mean formula gives us $E[T^* | T^* \leq 0, Z]$ directly, with no separate mirroring step needed:

$$E[T^* \mid T^* \leq 0, Z] = Z\hat{\pi} - \hat{\sigma} \frac{\phi(-Z\hat{\pi}/\hat{\sigma})}{\Phi(-Z\hat{\pi}/\hat{\sigma})} \quad (29.5)$$

```
library(AER)

tob <- tobit(T ~ Z1 + Z2, left = 0, data = data)
lin_pred <- coef(tob)["(Intercept)"] + coef(tob)["Z1"] * data$Z1 + coef(tob)["Z2"] * data$Z2
sigma_hat <- tob$scale
alpha <- -lin_pred / sigma_hat
E_Tstar_below0 <- lin_pred - sigma_hat * dnorm(alpha) / pnorm(alpha)

# here I create an extra term for the correction:
# when T is 0, it is T (=0); otherwise it's the truncated-normal conditional
# mean of T* given T* <= 0 and Z, computed above.
data <- data %>%
  mutate(extra = if_else(T > 0, T, E_Tstar_below0))

lm4 <- lm(Y ~ T + Z1 + Z2 + extra, data=data)
summary(lm4)
```

Call:

```
lm(formula = Y ~ T + Z1 + Z2 + extra, data = data)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.8902	-0.7686	0.0480	0.7935	5.4430

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.03693	0.07979	0.463	0.644
T	1.98735	0.05917	33.585	<2e-16 ***
Z1	-0.82513	0.05838	-14.133	<2e-16 ***
Z2	-2.31062	0.08224	-28.097	<2e-16 ***
extra	1.36256	0.05873	23.200	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.321 on 995 degrees of freedom

Multiple R-squared: 0.9573, Adjusted R-squared: 0.9572

F-statistic: 5582 on 4 and 995 DF, p-value: < 2.2e-16

Under the normality assumption, this recovers β (the true coefficient on T is 2) very closely — much better than a naive symmetric-mirroring heuristic that estimates the right-tail conditional mean by an unweighted linear regression of T on Z restricted to a somewhat arbitrary right-tail cutoff: that heuristic's accuracy turns out to be quite sensitive to the cutoff and sample used, because the true

29. *Endogeneity with censored variable*

conditional mean of T^* given T^* beyond a cutoff is a nonlinear (inverse-Mills-ratio) function of Z , not linear — so a plain linear regression on the selected right-tail subsample is itself misspecified. The parametric (Tobit) route sidesteps that by modeling the untruncated latent distribution directly and using the exact conditional-mean formula.

29.4. Conclusion

We are making assumptions for the distribution of the continuous treatment variable. Symmetry is the minimum assumption. We can make other parametric assumptions on the distribution of T^* . This is a trade-off between making distribution assumptions and ignoring endogeneity. The good news is that we don't need an instrument. We are using the censoring part of the treatment variable to deal with endogeneity.

30. G-computation (g-formula) with time varying covariates

30.1. The problem


The g-formula (g-computation algorithm) is a method to estimate the causal effect of a time varying treatment or covariates, developed by Robins (1986). Note this is distinct from **g-estimation** of structural nested models (Robins and Rotnitzky, 1992), a different g-methods estimator that uses instrumental-variable-like estimating equations and does not require modeling the joint distribution of all time-varying confounders. This chapter — and the `gfoRmula/gmethods` packages used below — implements the g-formula, not g-estimation.

We can imagine such a scenario. Suppose A_0 and A_1 are time 0 and time 1 treatment, can be binary or continuous. L_1 is a time 1 covariate. It can be some intermediate outcome. Say $HbA1c$. $HbA1c$ can be a result of treatment A_0 . Then it can be a cause for the next period treatment A_1 . Y is the outcome. U is the unmeasured confounder. The causal diagram is shown below. How do we estimate the effect of treatment?

```
library(dagitty)
library(ggdag)
library(ggplot2)
library(gfoRmula)
library(gmethods)
library(tidyverse)

dag <- dagify(
  Y ~ A1 + A0 + L1 ,
  A1 ~ A0 + L1,
  L1 ~ A0 ,
  coords = time_ordered_coords()
)

dag %>%
  ggplot(aes(x = x, y = y, xend = xend, yend = yend)) +
  geom_dag_point() +
  geom_dag_edges_arc() +
  geom_dag_text() +
  theme_dag()
```



30.2. Assumptions

There are four potential outcomes, with A_0 and A_1 can be 0 or 1. The causal estimand is

$$\tau_{a_0, a_1, a'_0, a'_1} = E[Y^{a_0, a_1} - Y^{a'_0, a'_1}] \quad (30.1)$$

To identify the effects, we need assumptions.

1. Sequential ignorability

$$Y^{a_0, a_1} \perp A_0 \quad (30.2)$$

$$Y^{a_0, a_1} \perp A_1^{a_0} | A_0, L_1^{a_0} \quad (30.3)$$

In general, treatment at time t is randomized with probabilities depending on the observed past (covariates and intermediate outcomes): at any time t , for every regime $\bar{a}$,

$$Y^{\bar{a}} \perp A_t | H_t \quad (30.4)$$

where H_t is the **observed** history of treatment and covariates up to time t ,

$$H_t = (A_0, A_1, \dots, A_{t-1}; L_1, \dots, L_t). \quad (30.5)$$

Read this in terms of the *observed* treatment A_t and the *observed* history H_t : conditional on what has actually happened up to time t , the next treatment is independent of the future potential outcome. The superscripted objects in the two-period statement above (e.g. $A_1^{a_0}$, $L_1^{a_0}$) denote the values that would arise under setting $A_0 = a_0$; by consistency these equal the observed A_1 , L_1 on the subset of units with $A_0 = a_0$, which is what lets the conditional-independence statement be checked on observed data.

2. Positivity

$$0 < P(A_0 = 1) < 1 \quad (30.6)$$

$$0 < P(A_1 = 1 | A_0 = 1, L_1) < 1 \quad (30.7)$$

In general,

$$0 < P(A_t = 1 | H_t) < 1 \quad (30.8)$$

3. Consistency

$$Y^{a_0, a_1} = Y | (A_0 = a_0, A_1 = a_1) \quad (30.9)$$

30.3. Identifiability

$$\begin{aligned} E[Y^{a_0, a_1}] &= E[Y^{a_0, a_1} | A_0 = a_0] \\ &= E[Y^{A_0, a_1} | A_0 = a_0] \\ &= \sum_{l=0,1} E[Y^{A_0, a_1} | A_0 = a_0, L = l] P(L = l | A_0 = a_0) \\ &= \sum_{l=0,1} E[Y^{A_0, a_1} | A_0 = a_0, L = l, A_1 = a_1] P(L = l | A_0 = a_0) \\ &= \sum_{l=0,1} E[Y^{A_0, A_1} | A_0 = a_0, L = l, A_1 = a_1] P(L = l | A_0 = a_0) \end{aligned} \quad (30.10)$$

This can be expanded to the general case. Basically start from the first period, then the second, based on the first, and so on.

30.4. Estimation

There are two types of components in the g-estimation. The first is the outcome model, then there are models for time-varying covariates, or intermediate outcomes (confounders).

30.5. Parametric g-formula

In this simple case, if Y is continuous, we can use a linear model. If Y is binary, we can use a logistic model.

$$E[Y | A_0, A_1, L_1] = \beta_0 + \beta_1 A_0 + \beta_2 L_1 + \beta_3 A_1 \quad (30.11)$$

Suppose L_1 is binary,

$$P(L_1 = 1 | A_0) = \text{logit}^{-1}(\gamma_0 + \gamma_1 A_0) \quad (30.12)$$

$$L_1|A_0 \sim \text{Bernoulli}(\text{logit}^{-1}(\gamma_0 + \gamma_1 A_0)) \quad (30.13)$$

We may need to simulate the distribution of the time-varying covariates for the expected value of the potential outcomes.

30.6. an example

I am following Christopher Boyer's example in his teaching page: <https://christopherb-boyer.com/about>.

The first example is from his lab3 teaching page.

First we set up a simulated data set.

```
# setting up the data -----

# The code below creates the frequency table, and then turns it into long format
# data. Read through the code carefully and run it.

# re-create the frequency table for homework 3
hw3_freq <- tribble(
  ~A_0, ~L_1, ~A_1, ~N, ~Y_1,
  0, 0, 0, 6000, 60,
  0, 0, 1, 2000, 60,
  0, 1, 0, 2000, 210,
  0, 1, 1, 6000, 210,
  1, 0, 0, 3000, 240,
  1, 0, 1, 1000, 240,
  1, 1, 0, 3000, 120,
  1, 1, 1, 9000, 120
)

# expand the frequency table to 1 row per person
# (first just assign everyone the average Y value)
dat <- uncount(hw3_freq, N) %>%
  # give those rows an id number
  rowid_to_column(var = "id") %>%
  # for each of the variables except for id, split it into two
  # rows, one for each time point
  pivot_longer(-id,
    names_to = c(".value", "time"),
    names_sep = "_"
  ) %>%
  # make sure that time is read as a number
  mutate(time = parse_number(time),
    # add some random error to Y
    # (nrow(.) means a different value for each row of the dataset in use
```

```

      Y = Y + rnorm(nrow(.), 0, 10))

# create lagged variables
# group_by(id) makes this robust to row order -- lag() would otherwise
# silently lag across a different person's rows if dat were ever re-sorted
# (e.g. by time first). Currently harmless here since dat is in id-major
# order with exactly one row per id per time, but safer to be explicit.
t1 <- dat %>%
  group_by(id) %>%
  mutate(lag_A = lag(A)) %>%
  ungroup() %>%
  filter(time == 1)

```

Then calculate the g-formula by hand.

First for $E[Y^{1,1}]$.

```

# fit models for covariate and outcome given past on t1
L_mod <- glm(L ~ lag_A, data = t1, family = binomial())
Y_mod <- glm(Y ~ A * lag_A * L, data = t1)

# create new data set with space for 10,000 simulated entries
nsim <- 10000
baseline_pop <- tibble(id = 1:nsim)

# fix intervention values in the new data set to be consistent with the regime
new_dat_11 <- baseline_pop %>%
  mutate(lag_A = 1, A = 1)

# predict covariate means at time 1
new_dat_11$pL <- predict(L_mod, newdata = new_dat_11, type = "response")

# simulate covariate values at time 1
new_dat_11$L <- rbinom(nsim, 1, new_dat_11$pL)

# predict outcome at time 1 based on simulated covariates and fixed treatments
new_dat_11$Ey <- predict(Y_mod, newdata = new_dat_11, type = "response")

# take mean across all simulations to get g-formula estimate!
mean(new_dat_11$Ey)

```

```
[1] 150.3459
```

Then for $E[Y^{0,0}]$.

30. G-computation (g-formula) with time varying covariates

```
# for 0, 0
new_dat_00 <- baseline_pop %>%
  mutate(lag_A = 0, A = 0)

# predict covariate means at time 1
new_dat_00$pL <- predict(L_mod, newdata = new_dat_00, type = "response")

# simulate covariate values at time 1
new_dat_00$L <- rbinom(nsim, 1, new_dat_00$pL)

# predict outcome at time 1 based on simulated covariates and fixed treatments
new_dat_00$Ey <- predict(Y_mod, newdata = new_dat_00, type = "response")

# take mean across all simulations to get g-formula estimate!
mean(new_dat_00$Ey)
```

```
[1] 135.1991
```

Now with the gfoRmula package.

```
# running the g-formula -----

# add your comments to the lines below
id <- "id"
time_name <- "time"
covnames <- c("L", "A")
outcome_name <- "Y"
#
outcome_type <- "continuous_eof"
#
histories <- c(lagged)
#
histvars <- list(c("A"))
#
covtypes <- c("binary", "binary")
covparams <- list(covmodels = c(
  L ~ lag1_A,
  A ~ lag1_A * L
))
ymodel <- Y ~ A * L * lag1_A
#
intvars <- list("A", "A", "A", "A")
#
interventions <- list(
  list(c(static, c(0, 0))),
  list(c(static, c(0, 1))),
```

```

    list(c(static, c(1, 0))),
    list(c(static, c(1, 1)))
  )
int_descript <- c(
  "0, 0",
  "0, 1",
  "1, 0",
  "1, 1"
)

# put it all together
# (the warning "obs_data was coerced to a data table" is fine!)
gform_res <- gfoRmula::gformula(
  obs_data = dat,
  id = id,
  time_name = time_name,
  covnames = covnames,
  outcome_name = outcome_name,
  outcome_type = outcome_type,
  covtypes = covtypes,
  covparams = covparams,
  ymodel = ymodel,
  intvars = intvars,
  interventions = interventions,
  int_descript = int_descript,
  ref_int = 1,
  histories = histories,
  histvars = histvars,
  sim_data_b = TRUE,
  seed = 345636
)

# Use the following code to explore the results. (You can ignore Intervention 0,
# the natural course intervention, for now.) How many models were fit? How can
# you tell which of the data in the `sim_data` object was simulated or predicted
# from a model? Does these results match your answers to the earlier questions,
gform_res

```

PREDICTED RISK UNDER MULTIPLE INTERVENTIONS

Intervention	Description
0	Natural course
1	0, 0
2	0, 1
3	1, 0
4	1, 1

Sample size = 32000, Monte Carlo sample size = 32000

30. G-computation (g-formula) with time varying covariates

Number of bootstrap samples = 0

Reference intervention = 1

k	Interv.	NP mean	g-form mean	Mean ratio	Mean difference	% Intervened	On
<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>
1	0	142.5619	141.6934	1.051544	6.9454932	0.00000	
1	1	NA	134.7479	1.000000	0.0000000	74.87812	
1	2	NA	134.6241	0.999081	-0.1238386	75.12188	
1	3	NA	150.3413	1.115723	15.5934188	80.85625	
1	4	NA	150.0710	1.113716	15.3230335	69.14375	

Aver % Intervened On

<num>
0.00000
49.74531
50.25469
55.99375
44.00625

`gform_res$coeffs`

\$L

(Intercept)	lag1_A
8.985178e-14	1.098612e+00

\$A

(Intercept)	lag1_A	L	lag1_A:L
-1.098612e+00	-2.273084e-13	2.197225e+00	2.983558e-14

\$Y

(Intercept)	A	L	lag1_A	A:L
60.07560004	-0.06212277	150.06685483	180.12328736	-0.12402849
A:lag1_A	L:lag1_A	A:L:lag1_A		
-0.52497353	-270.03686749	0.54687356		

`gform_res$sim_data`

\$`Natural course`

	id	time	L	A	eligible_pt	intervened	lag1_A	Ey
	<int>	<num>	<num>	<num>	<lgcl>	<num>	<num>	<num>
1:	1	0	NA	0	TRUE	0	0	NA
2:	1	1	1	1	TRUE	0	0	209.9563
3:	2	0	NA	0	TRUE	0	0	NA
4:	2	1	1	1	TRUE	0	0	209.9563
5:	3	0	NA	0	TRUE	0	0	NA


```

---
63996: 31998      1      1      1      TRUE      0      1 120.0646
63997: 31999      0     NA      1      TRUE      0      0      NA
63998: 31999      1      1      1      TRUE      0      1 120.0646
63999: 32000      0     NA      1      TRUE      0      0      NA
64000: 32000      1      1      1      TRUE      0      1 120.0646

```

\$`0, 0`

```

      id  time      L      A eligible_pt intervened lag1_A      Ey
      <int> <num> <num> <num>      <lgcl>      <lgcl> <num>      <num>
1:      1      0     NA      0      TRUE      FALSE      0      NA
2:      1      1      1      0      TRUE      TRUE      0 210.1425
3:      2      0     NA      0      TRUE      FALSE      0      NA
4:      2      1      1      0      TRUE      TRUE      0 210.1425
5:      3      0     NA      0      TRUE      FALSE      0      NA

```

```

---
63996: 31998      1      0      0      TRUE      FALSE      0 60.0756
63997: 31999      0     NA      0      TRUE      TRUE      0      NA
63998: 31999      1      0      0      TRUE      FALSE      0 60.0756
63999: 32000      0     NA      0      TRUE      TRUE      0      NA
64000: 32000      1      1      0      TRUE      TRUE      0 210.1425

```

\$`0, 1`

```

      id  time      L      A eligible_pt intervened lag1_A      Ey
      <int> <num> <num> <num>      <lgcl>      <lgcl> <num>      <num>
1:      1      0     NA      0      TRUE      FALSE      0      NA
2:      1      1      1      1      TRUE      FALSE      0 209.95630
3:      2      0     NA      0      TRUE      FALSE      0      NA
4:      2      1      1      1      TRUE      FALSE      0 209.95630
5:      3      0     NA      0      TRUE      FALSE      0      NA

```

```

---
63996: 31998      1      0      1      TRUE      TRUE      0 60.01348
63997: 31999      0     NA      0      TRUE      TRUE      0      NA
63998: 31999      1      0      1      TRUE      TRUE      0 60.01348
63999: 32000      0     NA      0      TRUE      TRUE      0      NA
64000: 32000      1      1      1      TRUE      FALSE      0 209.95630

```

\$`1, 0`

```

      id  time      L      A eligible_pt intervened lag1_A      Ey
      <int> <num> <num> <num>      <lgcl>      <lgcl> <num>      <num>
1:      1      0     NA      1      TRUE      TRUE      0      NA
2:      1      1      1      0      TRUE      TRUE      1 120.2289
3:      2      0     NA      1      TRUE      TRUE      0      NA
4:      2      1      1      0      TRUE      TRUE      1 120.2289
5:      3      0     NA      1      TRUE      TRUE      0      NA

```

```

---
63996: 31998      1      1      0      TRUE      TRUE      1 120.2289
63997: 31999      0     NA      1      TRUE      FALSE      0      NA

```

30. G-computation (g-formula) with time varying covariates

```
63998: 31999      1      1      0      TRUE      TRUE      1 120.2289
63999: 32000      0     NA      1      TRUE     FALSE      0      NA
64000: 32000      1      1      0      TRUE      TRUE      1 120.2289
```

```
$`1, 1`
```

	id	time	L	A	eligible_pt	intervened	lag1_A	Ey
	<int>	<num>	<num>	<num>	<lgcl>	<lgcl>	<num>	<num>
1:	1	0	NA	1	TRUE	TRUE	0	NA
2:	1	1	1	1	TRUE	FALSE	1	120.0646
3:	2	0	NA	1	TRUE	TRUE	0	NA
4:	2	1	1	1	TRUE	FALSE	1	120.0646
5:	3	0	NA	1	TRUE	TRUE	0	NA

63996:	31998	1	1	1	TRUE	FALSE	1	120.0646
63997:	31999	0	NA	1	TRUE	FALSE	0	NA
63998:	31999	1	1	1	TRUE	FALSE	1	120.0646
63999:	32000	0	NA	1	TRUE	FALSE	0	NA
64000:	32000	1	1	1	TRUE	FALSE	1	120.0646

30.7. A second example

Christopher Boyer has a package “gmethods” that can be used to estimate the g-formula. To me it’s a bit more intuitive than the gfoRmula package. Here is an example in the “gmethods” package, with the same model in the “gfoRmula” package.

The example dataset continuous_cofdata again consists of 7,500 observations on 2,500 individuals with a maximum of 7 follow-up times, where the outcome corresponds to a characteristic only in the last interval (e.g., systolic blood pressure in interval 7). The variables in the dataset are: t0: The time index id: A unique identifier for each individual L1: A time-varying covariate; categorical L2: A time-varying covariate; continuous L3: A baseline covariate; continuous A: The treatment variable; binary Y: The outcome of interest; continuous


The goal is to estimate the mean outcome at time $K + 1 = 7$ under “never treat” versus “always treat.”

The DAG would be something like:

```
dag <- dagify(
  Y ~ A + L1 + L2 + L3 + lag1_A + lag1_L1 + lag1_L2 ,
  L1 ~ lag1_A + lag1_L1 + lag1_L2 + L3 ,
  L2 ~ lag1_A + L1 + lag1_L1 + lag1_L2 + L3,
  A ~ lag1_A + L1 + L2 + lag1_L1 + lag1_L2 + L3,
  coords = time_ordered_coords()
)

dag %>%
  ggplot(aes(x = x, y = y, xend = xend, yend = yend)) +
  geom_dag_point() +
```

```
geom_dag_edges_arc() +
geom_dag_text() +
theme_dag()
```



30.7.1. gFormula results

First use the gFormula package.

```
id <- 'id'
time_name <- 't0'
covnames <- c('L1', 'L2', 'A')
outcome_name <- 'Y'
covtypes <- c('categorical', 'normal', 'binary')
histories <- c(lagged)
histvars <- list(c('A', 'L1', 'L2'))
covparams <- list(covmodels = c(L1 ~ lag1_A + lag1_L1 + L3 + t0 +
                                lag1_L2,
                                L2 ~ lag1_A + L1 + lag1_L1 + lag1_L2 + L3 + t0,
                                A ~ lag1_A + L1 + L2 + lag1_L1 + lag1_L2 + L3 + t0))
ymodel <- Y ~ A + L1 + L2 + lag1_A + lag1_L1 + lag1_L2 + L3
intvars <- list('A', 'A')
interventions <- list(list(c(static, rep(0, 7))),
                      list(c(static, rep(1, 7))))
int_descript <- c('Never treat', 'Always treat')
nsimul <- 50000
```

30. G-computation (g-formula) with time varying covariates

```
gform_cont_eof <- gformula_continuous_eof(obs_data = continuous_eofdata,
                                         id = id,
                                         time_name = time_name,
                                         covnames = covnames,
                                         outcome_name = outcome_name,
                                         covtypes = covtypes,
                                         covparams = covparams, ymodel = ymodel,
                                         intvars = intvars,
                                         interventions = interventions,
                                         int_descript = int_descript,
                                         histories = histories,
                                         histvars = histvars,
                                         model_fits = TRUE,
                                         basecovs = c("L3"),
                                         nsimul = nsimul, seed = 1234)

gform_cont_eof
```

PREDICTED RISK UNDER MULTIPLE INTERVENTIONS

Intervention	Description
0	Natural course
1	Never treat
2	Always treat

Sample size = 2500, Monte Carlo sample size = 50000
Number of bootstrap samples = 0
Reference intervention = natural course (0)

k	Interv.	NP mean	g-form mean	Mean ratio	Mean difference	% Intervened	On
<num>	<num>	<num>	<num>	<num>	<num>	<num>	<num>
6	0	-4.414543	-4.363027	1.0000000	0.0000000	0.000	
6	1	NA	-3.133081	0.7180981	1.2299454	100.000	
6	2	NA	-4.602952	1.0549904	-0.2399248	74.038	
Aver % Intervened On							
		<num>					
		0.00000					
		82.35200					
		17.49571					

30.7.2. gmethods results

Now use the gmethods package.

```

gf <- gmethods::gformula(
  outcome_model = list(
    formula = Y ~ A + L1 + L2 + lag1_A + lag1_L1 + lag1_L2 + L3,
    link = "identity",
    family = "normal"
  ),
  covariate_model = list(
    "L1" = list(
      formula = L1 ~ lag1_A + lag1_L1 + L3 + t0 + lag1_L2,
      family = "categorical"
    ),
    "L2" = list(
      formula = L2 ~ lag1_A + L1 + lag1_L1 + lag1_L2 + L3 + t0,
      link = "identity",
      family = "normal"
    ),
    "A" = list(
      formula = A ~ lag1_A + L1 + L2 + lag1_L1 + lag1_L2 + L3 + t0,
      link = "logit",
      family = "binomial"
    )
  ),
  data = continuous_eofdata,
  survival = FALSE,
  id = 'id',
  time = 't0'
)

s <- simulate(
  gf,
  interventions = list(
    "Never treat" = function(data, time) set(data, j = "A", value = 0),
    "Always treat" = function(data, time) set(data, j = "A", value = 1)
  ),
  n_samples = 50000
)

s$results

```

```

$means
  intervention estimate conf.low conf.high
1             0 -4.360358      NA      NA
2             1 -3.139260      NA      NA
3             2 -4.602994      NA      NA

```

```

$ratios
  intervention estimate conf.low conf.high

```

30. G-computation (g-formula) with time varying covariates

1	0	NA	NA	NA
2	1	0.7199547	NA	NA
3	2	1.0556459	NA	NA

```
$diffs
  intervention estimate conf.low conf.high
1           0         NA         NA         NA
2           1  1.221098         NA         NA
3           2 -0.242636         NA         NA
```

Note the results are very close. Also note that in the covariates and treatment models, there is a time variable t0. This is because these models are pooled regression with all time points. Therefore a time fixed effect is included. In the outcome model, there is only the last period data, so the time variable is not included.

Standard errors can be calculated with bootstrapping.

Systematic treatment: [R](#) · [Julia](#).

31. Policy learning by policytree

31.1. Optimal policy

I have a few posts about estimating ATE or CATE. The target has been getting the treatment effect or conditional treatment effect accurately. However, sometimes after estimating the treatment effects, we may be interested in the optimal treatment assignment. For example, a doctor may be interested in how to assign treatment based on patient characteristics. OSHA might be interested in which firms to send inspector to, based on firm characteristics, or past history. We may think we can just assign the treatment to those who benefited most (largest CATE), or any patient who has a positive CATE. However, the objective function for policy is different from the CATE target.

Suppose we have $\{Y(0), Y(1), X\}$ and we have estimated the CATE as $\tau(x) = E[Y(1) - Y(0) | X = x]$. The policy learning problem is to find a policy function that assigns treatment to maximize some utility function.

Suppose that policy is

$$\pi : X \rightarrow \{0, 1\} \quad (31.1)$$

This policy is to assign treatment or control based on X .

The natural objective function is a value function $V(\pi) = E[Y_i(\pi(X_i))]$, which is the expected value of outcome with this policy. We'd like to maximize this function.

$$V(\pi) = E[Y_i(\pi(X_i))] = E[Y_i(0)] + E[\tau(X)\pi(X)] \quad (31.2)$$

Since π only enters through $E[\tau(X)\pi(X)]$, the maximizer is exactly $\pi(X) = 1\{\tau(X) > 0\}$ — treat everyone with a positive effect. A threshold $c \neq 0$ is not a free choice here: it appears once treatment is costly, where maximizing $E[(\tau(X) - c)\pi(X)]$ gives the rule $\tau(X) > c$ with c the per-unit cost. However, estimating $\tau(X)$ is different from maximizing V .

First, we don't have the true τ . The loss function in estimating $\tau(X)$ is different from the loss function in maximizing V . The loss function in estimating $\tau(X)$ is the squared error loss. We'll see maximizing V is a different problem.

Second, the X we use to estimate τ is different from the ones we can use for policy. For example, we have some variables used in CATE estimation that we cannot use for policy, like gender, or age, for discrimination reasons. Or there are variables we have after the experiment that we cannot use for CATE estimation, but we'd like to use for policy learning.

31.2. IPW loss

Kitagawa and Tetenov (2018) proposed a loss function for policy learning. The loss function is based on the inverse propensity weighting.

$$\hat{\pi} = \operatorname{argmax}\{\hat{V}(\pi) : \pi \in \Pi\} \quad (31.3)$$

$$\hat{V}(\pi) = \frac{1}{n} \sum_{i=1}^n \frac{1(W_i = \pi(X_i))}{P[W_i = \pi(X_i)|X_i]} Y_i \quad (31.4)$$

This is to say when we have $W_i = \pi(X_i)$, we can use the outcome Y_i to estimate the value function. That is, when we have a policy, we then select the observations that matches that policy then weight the outcome by the inverse propensity score. This is to adjust the fact the there is selection bias for that policy. Therefore we re-weight it by IPW.

Obviously we don't have the true propensity score of $W_i = \pi(X_i)$, so we need to estimate it. We can estimate it by machine learning.

31.3. AIPW loss

The IPW loss can be not robust to model misspecification. Athey and Wager (2021) proposed a loss function that is robust to model misspecification. The loss function is based on the doubly robust estimator, AIPW.

$$\hat{\pi} = \operatorname{argmax}\{\hat{V}(\pi) : \pi \in \Pi\} \quad (31.5)$$

$$\hat{V}(\pi) = \frac{1}{n} \sum_{i=1}^n (2\pi(X_i) - 1)\hat{\Gamma}_i \quad (31.6)$$

$$\Gamma_i = \hat{\mu}(1)(X_i) - \hat{\mu}(0)(X_i) + \frac{W_i}{\hat{e}(X_i)}(Y_i - \hat{\mu}(1)(X_i)) - \frac{1 - W_i}{1 - \hat{e}(X_i)}(Y_i - \hat{\mu}(0)(X_i)) \quad (31.7)$$

Note the Γ_i is the score from the AIPW estimator. It's basically the efficient influence function we learned from the nonparametric estimation of ATE. The value function $\hat{V}(\pi)$ is the score if treated, minus the score if control for each observation. We'd like to maximize this value function.

The basic idea is for each policy, calculate the score from AIPW estimator. Then calculate the value function. Then search over the policy space for the policy that maximizes the value function.

This is very computationally intensive task. Right now seems it would work for "shallow" trees. That is, we allow 2 or 3 levels of trees. Otherwise, the search space could be too large to search.

31.4. Policytree


```


library(grf)
library(policytree)
library(DiagrammeR)
n <- 1000
p <- 10

X <- matrix(rnorm(n * p), n, p)
ee <- 1 / (1 + exp(X[, 3]))
tt <- 1 / (1 + exp((X[, 1] + X[, 2]) / 2)) - 0.5
W <- rbinom(n, 1, ee)
Y <- X[, 3] + W * tt + rnorm(n)

cf <- causal_forest(X, Y, W)

plot(tt, predict(cf)$predictions)

```



```

dr <- double_robust_scores(cf)
tree <- policy_tree(X, dr, 2)
tree

```

```


policy_tree object
Tree depth: 2
Actions: 1: control 2: treated
Variable splits:
(1) split_variable: X3 split_value: -0.563092
(2) split_variable: X1 split_value: -0.599028
(4) * action: 2

```


31. Policy learning by *policytree*

```
(5) * action: 1  
(3) split_variable: X10  split_value: 1.09723  
(6) * action: 1  
(7) * action: 2
```

```
pp <- predict(tree, X)  
boxplot(tt ~ pp)
```




```
plot(tree)
```



31. Policy learning by `policytree`

```
plot(X[, 1], X[, 2], col = pp)
abline(0, -1, lwd = 4, col = 4)
```



In the example, first “`causal_forest`” is used to estimate the CATE. “`double_robust_scores`” is used to calculate the AIPW score. Then “`policy_tree`” is used to estimate the policy tree. The level of tree is set to 2.

Systematic treatment: [R](#) · [Julia](#).

32. Recent causal inference tools

In this post, I'll use simulated data to see how a few recently developed causal inference packages work.

In other posts, I have discussed some methods involving machine learning, such as TMLE, or causal forest. This time we'll talk about nonparametric methods and double machine learning.

32.1. Identification

First suppose we have a “target” parameter in mind, which is some function of the some unknown distribution P^* , called a functional, say $\psi^*(P^*)$. For example ATE, $E(Y^1 - Y^0)$.

The goal is to have $\psi^*(P^*)$ estimated by $\psi(\hat{P})$, where P is some observational distribution, and ψ is some known function. This is called identification.

For example, suppose we are interested in the ATE, $E[Y^1 - Y^0]$. The identification is done by deriving from $E(Y^a|X) = E(Y|X, A = a)$; basically from the target parameter, which is counterfactual, to something estimable from the observed variables in the sample. This identification can only be done with assumptions such as ignorability, positivity and consistency.

In this case, $\psi^* = E(Y^a)$, and $\psi = E\{E(Y|X)\}$.

32.2. Nonparametric estimation in a nutshell

After identification, we then need to find a way to construct a good estimator $\hat{\psi}$ for $\psi(P)$.

Suppose we have a sample Z , which can be A, W, Y (treatment, covariates, and outcome) in causal inference situation.

We can use a parametric model, assuming we know what kind of distribution P is. We assume $P = P_\theta$. Then we can use MLE to get $\psi(P_\theta)$. This is what we do in parametric modeling. But this can be too restrictive, we usually have no way of knowing how good P_θ is; in other words, the model can be misspecified thus the target parameter is estimated with bias.

If we don't make any assumptions about the distribution, then we are in nonparametric world. The natural way to approximate P is to use the empirical distribution $\hat{P}$. Then $\psi^*(\hat{P})$ is the “plug-in” estimator. For example we can use sample mean to estimate the population mean, etc. The plug-in is often not the best choice *when the nuisance functions are estimated flexibly* (e.g. with machine learning): the plug-in then inherits the slow convergence rate of those nuisance estimates and carries a first-order regularization (“plug-in”) bias, so it is generally not $\sqrt{n}$ consistent in that regime. (This is not a universal indictment of plug-in estimators: with a correctly specified, smooth parametric

model a plug-in / g-computation estimator can itself be $\sqrt{n}$ efficient. The problem is specifically the flexible / high-dimensional nuisance case.)

Another choice is nonparametric estimator based on influence function. The nice thing about this type of estimators based on efficient influence function is that they are often $\sqrt{n}$ consistent, meaning they converge to the truth as fast as $\sqrt{n}$ rate. This means $\hat{\psi} - \psi$ not only goes to 0 with n goes to infinity, but also as fast as $1/\sqrt{n}$ goes to 0. Why does this matter? Because we have limited sample size, when we are based on asymptotics, the faster the bias goes to 0 the smaller sample size we need.

If we have an estimator such that

$$\sqrt{n}(\hat{\psi} - \psi) = \sqrt{n}P_n(\phi) + o_P(1) \rightarrow N(0, \text{var}(\phi)) \quad (32.1)$$

Then this estimator is $\sqrt{n}$ consistent and asymptotically normal; and ϕ is the influence function. The efficient influence function is the one with variance at the lower bound (similar to parametric case for Cramer-Rao lower bound).

Consider the mean of the treated potential outcome, $\psi = E\{E(Y|X, A = 1)\} = E[Y^1]$. Its efficient influence function is

$$\phi_1(Z; P) = \frac{A}{\pi(X)}\{Y - \mu_1(X)\} + \mu_1(X) - \psi \quad (32.2)$$

where $\pi(X) = P(A = 1|X)$ the propensity score model, $\mu_1(X) = E(Y|X, A = 1)$, the outcome model. The ATE $E[Y^1 - Y^0]$ has efficient influence function $\phi_1 - \phi_0$, where ϕ_0 is the analogous control-arm term $\frac{1-A}{1-\pi(X)}\{Y - \mu_0(X)\} + \mu_0(X) - E[Y^0]$ with $\mu_0(X) = E(Y|X, A = 0)$.

The idea of IF-based estimator is to use so called “one-step correction” to estimate the “plug-in” bias. The bias is approximated by first deriving the influence function; then based on the influence function, the bias can be estimated by using the influence function as a slope to the path from “plug-in” estimate” to true parameter. This path can be non-linear; therefore this one-step correction may not be perfect, but it should be closer to the truth than the original estimate.

For example, an IF-based bias-corrected estimator

$$\hat{\psi} = P_n\left[\frac{A}{\hat{\pi}(X)}\{Y - \hat{\mu}(X)\} + \hat{\mu}(X)\right] \quad (32.3)$$

where P_n means sample mean. This is the familiar AIPW estimator. It is known to be doubly robust. Since we know the influence function, we can calculate its variance. This variance is the same as variance of $\hat{\psi}$ since they differ by a constant ψ . We can use any estimators to estimate the nuisance parameters, then get $\hat{\psi}$. Then the sample variance is the variance of the influence function; and the sample mean is the ATE.

In this estimator, both $\hat{\mu}$ and $\hat{\pi}$ can be done using any estimator, such as machine learning.

The IF based estimator only works if either $\phi(P)$ is not too complex (empirical processes theory involved, a lot of difficult situations won't apply), or we separate $\hat{P}$ and P_n to prevent overfitting. The latter is by sample splitting. Sample splitting is usually preferred because it needs less assumption than empirical processes. The idea is to split the sample, do estimation of nuisance parameters (such as μ and π) in one sample, and construct estimator in the other. This way it is shown the EIF based estimator is $\sqrt{n}$ consistent and asymptotically normal; the variance is the variance of the influence function.

Unfortunately influence function is hard to derive for different situations. Ed Kennedy has done a lot of work on using IF on causal inference <http://www.ehkennedy.com/>. Here we use his “npcausal” package.

32.3. Simulation with a known treatment effect

The simulated data is from <https://migariane.github.io/TMLE.nb.html>

```
library(npcausal)
library(boot)
library(MASS)
library(SuperLearner)
library(tmle)
library(tidyverse)
library(ranger)
```

```
set.seed(123)
n=1000
w1 <- rbinom(n, size=1, prob=0.5)
w2 <- rbinom(n, size=1, prob=0.65)
w3 <- round(runif(n, min=0, max=4), digits=3)
w4 <- round(runif(n, min=0, max=5), digits=3)
A <- rbinom(n, size=1, prob= plogis(-0.4 + 0.2*w2 + 0.15*w3 + 0.2*w4 + 0.15*w2*w4))
Y.1 <- rbinom(n, size=1, prob= plogis(-1 + 1 -0.1*w1 + 0.3*w2 + 0.25*w3 + 0.2*w4 + 0.15*w2*w4))
Y.0 <- rbinom(n, size=1, prob= plogis(-1 + 0 -0.1*w1 + 0.3*w2 + 0.25*w3 + 0.2*w4 + 0.15*w2*w4))
Y <- Y.1*A + Y.0*(1 - A)
W<-data.frame(cbind(w1,w2,w3,w4))
data <- data.frame(w1, w2, w3, w4, A, Y, Y.1, Y.0)
True_Psi <- mean(data$Y.1-data$Y.0);
cat(" True_Psi:", True_Psi)
```

```
True_Psi: 0.203
```

We first include a few learners. “npcausal” has a set of default learners.

```
SL.library <- c("SL.earth","SL.gam","SL.glm","SL.glmnet","SL.glm.interaction", "SL.mean","SL.ranger")
```

32.3.1. plug-in estimator from outcome regression

```
SL.outcome<- SuperLearner(Y=data$Y, X=data %>% select(w1,w2,w3,w4,A),
                          SL.library=SL.library, family="binomial")
SL.outcome.obs<- predict(SL.outcome, X=data %>% select(w1,w2,w3,w4,A))$pred
# predict the PO  $Y^1$ 
```

32. Recent causal inference tools

```
SL.outcome.1<- predict(SL.outcome, newdata=data %>% select(w1,w2,w3,w4,A) %>% mutate(A=1))$pred  
# predict the PO  $Y^0$   
SL.outcome.0<- predict(SL.outcome, newdata=data %>% select(w1,w2,w3,w4,A) %>% mutate(A=0))$pred
```

This is to do machine learning on the outcome equation; then get prediction on treated vs control, to get $\hat{Y}^1$ and $\hat{Y}^0$. The sample average of the difference is our plug-in estimator.

```
SL.plugin<-mean(SL.outcome.1-SL.outcome.0)  
SL.plugin
```

```
[1] 0.2050486
```

It's not too far from the truth in this case.

```
Q=data.frame(QAW=SL.outcome.obs, Q0W=SL.outcome.0, Q1W=SL.outcome.1)
```

32.3.2. AIPW with machine learning

We can use “npcausal” to do AIPW, or we can manually calculate it.

First we get propensity score. Then, based on both $\hat{\mu}$ and $\hat{\pi}$, we can get the AIPW estimator.

```
set.seed(123)  
SL.g<- SuperLearner(Y=data$A, X=data %>% select(w1,w2,w3,w4),  
                   SL.library=SL.library, family="binomial")  
g1W <- SL.g$SL.predict  
g0W<- 1- g1W  
IF.1<-((data$A/g1W)*(data$Y-Q$Q1W)+Q$Q1W)  
IF.0<-(((1-data$A)/g0W)*(data$Y-Q$Q0W)+Q$Q0W)  
IF<-IF.1-IF.0  
aipw.1<-mean(IF.1);aipw.0<-mean(IF.0)  
aipw.manual<-aipw.1-aipw.0  
ci.lb<-mean(IF)-qnorm(.975)*sd(IF)/sqrt(length(IF))  
ci.ub<-mean(IF)+qnorm(.975)*sd(IF)/sqrt(length(IF))  
res.manual.aipw<-c(aipw.manual,ci.lb, ci.ub)  
res.manual.aipw
```

```
[1] 0.2123436 0.1511466 0.2735405
```

Now we use “npcausal”.

```
library(npcausal)  
aipw<- ate(y=Y, a=A, x=W, nsplits=1, sl.lib=SL.library)
```



```

|
|
|
|=====
|
|=====
|
|=====
|
|=====
| 100%

```

	parameter	est	se	ci.ll	ci.ul	pval
1	E{Y(0)}	0.5719531	0.02594164	0.5211075	0.6227987	0
2	E{Y(1)}	0.7860664	0.01580281	0.7550929	0.8170399	0
3	E{Y(1)-Y(0)}	0.2141133	0.03004191	0.1552311	0.2729955	0

```
aipw$res
```

	parameter	est	se	ci.ll	ci.ul	pval
1	E{Y(0)}	0.5719531	0.02594164	0.5211075	0.6227987	0
2	E{Y(1)}	0.7860664	0.01580281	0.7550929	0.8170399	0
3	E{Y(1)-Y(0)}	0.2141133	0.03004191	0.1552311	0.2729955	0

The results from “npcausal” are almost the same as the manual results. Two clarifications. First, the manual estimator above *is* bias-corrected in the relevant sense: AIPW is exactly the one-step correction of the plug-in estimator using the efficient influence function. What we did *not* do is **cross-fitting** (sample splitting), which is a separate device. Second, that matters for inference: because the manual example fits the SuperLearner nuisances on the *full* sample and then uses $\text{sd}(\text{IF})/\sqrt{n}$, the resulting confidence interval is only valid under empirical-process (Donsker) conditions that flexible ML learners typically violate. With cross-fitting, the influence-function variance is valid without those conditions. “npcausal” does the cross-fitting for you when you set `nsplits > 1` (here `nsplits = 1` turns it off, for comparison with the manual full-sample version).


```
aipw2<- ate(y=Y, a=A, x=W, nsplits=10)
```

```

|
|
|
|==
|
|====
|
|=====
|
|=====
| 10%

```

32. Recent causal inference tools



```

|
|=====| 72%
|
|=====| 75%
|
|=====| 78%
|
|=====| 80%
|
|=====| 82%
|
|=====| 85%
|
|=====| 88%
|
|=====| 90%
|
|=====| 92%
|
|=====| 95%
|
|=====| 98%
|
|=====| 100%

```

	parameter	est	se	ci.ll	ci.ul	pval
1	E{Y(0)}	0.5701831	0.02904995	0.5132452	0.6271210	0
2	E{Y(1)}	0.7864167	0.01628375	0.7545005	0.8183328	0
3	E{Y(1)-Y(0)}	0.2162336	0.03300309	0.1515475	0.2809197	0

32.3.3. TMLE

We can use TMLE for ATE.

```

library(tmle)
TMLE<- tmle(Y=data$Y,A=data$A,W=W, family="binomial", Q.SL.library=SL.library, g.SL.library=SL
TMLE$estimates$ATE

```

```

$psi
[1] 0.2117905

```

```

$var.psi
[1] 0.001019485

```

```

$CI
[1] 0.1492101 0.2743709

```

```
$pvalue
[1] 3.287231e-11

$bs.var
[1] NA

$bs.CI.twosided
[1] NA NA

$bs.CI.onesided.lower
[1] -Inf    NA

$bs.CI.onesided.upper
[1]    NA Inf
```

32.3.4. DoubleML

The R package DoubleML implements the double/debiased machine learning framework of Chernozhukov et al. (2018). It provides functionalities to estimate parameters in causal models based on machine learning methods. The double machine learning framework consist of three key ingredients: Neyman orthogonality, High-quality machine learning estimation and Sample splitting.

They consider a partially linear model:

$$y_i = \theta d_i + g_0(x_i) + \eta_i \quad (32.4)$$

$$d_i = m_0(x_i) + v_i \quad (32.5)$$

This model is quite general, except it does not allow interaction of d and x ; therefore no heterogeneous treatment effect across x . But “DoubleML” implements more than partially linear model, it actually allows for heterogeneous treatment effects, in models such as interactive regression model.

The basic idea of doubleML is to use any machine learning algorithm to estimate outcome equation ($l_0(X) = E(Y|X)$) and treatment equation ($m_0(X) = E(D|X)$). Then get the residuals, namely $\hat{W} = Y - \hat{l}_0(X)$ and $\hat{V} = D - \hat{m}_0(X)$.

Then regress $\hat{W}$ on $\hat{V}$. Based on FWL theorem, you get $\hat{\theta}$.

An important component here is to specify a Neyman-orthogonal score function. In the case of PLR, it can be the product of the two residuals:

$$\psi(W; \theta, \eta) = (Y - D\theta - g(X))(D - m(X)) \quad (32.6)$$

The estimators $\hat{\theta}$ is to solve the equation that the sample mean of this score function being 0.

And the variance of this score function is used as the variance of the doubleML estimator’s variance.

```
library(DoubleML)
dml_data = DoubleMLData$new(data,
                             y_col = "Y",
                             d_cols = "A",
                             x_cols = c("w1", "w2", "w3", "w4"))
print(dml_data)
```

===== DoubleMLData Object =====

```
----- Data summary -----
Outcome variable: Y
Treatment variable(s): A
Covariates: w1, w2, w3, w4
Instrument(s):
Selection variable:
No. Observations: 1000
```

```
library(mlr3)
library(mlr3learners)
# suppress messages from mlr3 package during fitting
lgr::get_logger("mlr3")$set_threshold("warn")

learner = lrn("regr.ranger", num.trees=500, max.depth=5, min.node.size=2)
ml_g = learner$clone()
ml_m = learner$clone()

learner = lrn("regr.glmnet")
ml_g_sim = learner$clone()
ml_m_sim = learner$clone()

set.seed(3141)
obj_dml_plr = DoubleMLPLR$new(dml_data, ml_g=ml_g, ml_m=ml_m)
obj_dml_plr$fit()
print(obj_dml_plr)
```

===== DoubleMLPLR Object =====

```
----- Data summary -----
Outcome variable: Y
Treatment variable(s): A
Covariates: w1, w2, w3, w4
Instrument(s):
Selection variable:
No. Observations: 1000
```

```

----- Score & algorithm -----
Score function: partialling out
DML algorithm: dml2

----- Machine learner -----
ml_l: regr.ranger
ml_m: regr.ranger

----- Resampling -----
No. folds: 5
No. repeated sample splits: 1
Apply cross-fitting: TRUE

----- Fit summary -----
Estimates and significance testing of the effect of target variables
  Estimate. Std. Error t value Pr(>|t|)
A    0.23271    0.03224    7.219 5.25e-13 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Note: this example (and the EAGeR example below) both have a **binary** treatment **A** and outcome **Y**, but **DoubleMLPLR** assumes a partially linear model with a single constant effect **theta** and continuous-style regression learners (**ml_g**, **ml_m**) for both nuisance functions. The more appropriate DoubleML estimator for a binary treatment/outcome is the Interactive Regression Model, **DoubleMLIRM**, which uses classification learners for the propensity score and allows heterogeneous treatment effects. **DoubleMLPLR** will still run and often gives a reasonable ATE-like summary, but **DoubleMLIRM** is the estimator actually designed for this setting.

32.4. EAGeR trial simulation (AIPW package)

This simulation is based on a package “AIPW”, which has a simulation data “eager_sim_rct”. A denotes the binary treatment assignment, Y is the binary outcome, and W_g represents the covariate that affects the treatment assignment, which in our case was deemed to be the eligibility stratum indicator, sampled with replacement from the EAGeR Trial. Similarly, W_Q is a set of baseline prognostic covariates, which were also sampled with replacement from the EAGeR Trial, and includes the number of prior pregnancy losses, age, number of months of trying to conceive prior to randomization, body mass index (weight (kg)/height (m)²), and mean arterial blood pressure (denoted $W_{1...5}$, respectively).

Data and the true risk difference:

```

library(AIPW)
data("eager_sim_rct")
glimpse(eager_sim_rct)

```

```

Rows: 1,228
Columns: 8
$ sim_Y      <int> 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0~
$ sim_Tx     <int> 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0~
$ eligibility <dbl> 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0~
$ loss_num   <dbl> 2, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 1, 1, 2, 2, 2, 1~
$ age        <dbl> 21.93, 30.65, 30.96, 32.24, 31.49, 27.99, 31.60, 25.~
$ time_try_pregnant <dbl> 2, 1, 0, 9, 1, 10, 8, 6, 1, 0, 1, 1, 3, 3, 0, 10, 10~
$ BMI        <dbl> 19.51963, 25.56474, 20.75446, 21.65728, 24.27480, 19~
$ meanAP     <dbl> 73.33333, 79.00000, 79.33333, 63.83333, 97.44444, 85~

```

```
attributes(eager_sim_rct)$true_rd
```

```
[1] 0.1324569
```

```

data <- eager_sim_rct %>%
  rename(Y=sim_Y, A=sim_Tx) %>%
  tibble()

```

```

W <- data %>%
  select(-Y,-A)

```

32.4.1. plug-in estimator from outcome regression

```

SL.outcome<- SuperLearner(Y=data$Y, X=data %>% select(-Y),
                        SL.library=SL.library, family="binomial")
SL.outcome.obs<- predict(SL.outcome, X=data %>% select(-Y))$pred
# predict the PO  $Y^1$ 
SL.outcome.1<- predict(SL.outcome, newdata=data %>% select(-Y) %>% mutate(A=1))$pred
# predict the PO  $Y^0$ 
SL.outcome.0<- predict(SL.outcome, newdata=data %>% select(-Y) %>% mutate(A=0))$pred

```

This is to do machine learning on the outcome equation; then get prediction on treated vs control, to get $\hat{Y}^1$ and $\hat{Y}^0$. The sample average of the difference is our plug-in estimator.

```

SL.plugin<-mean(SL.outcome.1-SL.outcome.0)
SL.plugin

```

```
[1] 0.1185897
```

```
Q=data.frame(QAW=SL.outcome.obs, Q0W=SL.outcome.0, Q1W=SL.outcome.1)
```

32.4.2. AIPW with machine learning

```

set.seed(123)
SL.g<- SuperLearner(Y=data$A, X=data %>% select(-Y, -A),
                    SL.library=SL.library, family="binomial")
g1W <- SL.g$SL.predict
g0W<- 1- g1W
IF.1<-((data$A/g1W)*(data$Y-Q$Q1W)+Q$Q1W)
IF.0<-(((1-data$A)/g0W)*(data$Y-Q$Q0W)+Q$Q0W)
IF<-IF.1-IF.0
aipw.1<-mean(IF.1);aipw.0<-mean(IF.0)
aipw.manual<-aipw.1-aipw.0
ci.lb<-mean(IF)-qnorm(.975)*sd(IF)/sqrt(length(IF))
ci.ub<-mean(IF)+qnorm(.975)*sd(IF)/sqrt(length(IF))
res.manual.aipw<-c(aipw.manual,ci.lb, ci.ub)
res.manual.aipw

```

```
[1] 0.13103847 0.06938821 0.19268873
```

Now we use “npcausal”.

```

library(npcausal)
aipw<- ate(y=data$Y, a=data$A, x=W, nsplits=1, sl.lib=SL.library)

```

```

|
|
|
|=====| 25%
|
|=====| 50%
|
|=====| 75%
|
|=====| 100%

```

	parameter	est	se	ci.ll	ci.ul	pval
1	E{Y(0)}	0.4455985	0.02240994	0.40167501	0.4895220	0
2	E{Y(1)}	0.5753483	0.02105948	0.53407176	0.6166249	0
3	E{Y(1)-Y(0)}	0.1297499	0.03059579	0.06978212	0.1897176	0

The results from “npcausal” are almost the same as the manual results. Two clarifications. First, the manual estimator above *is* bias-corrected in the relevant sense: AIPW is exactly the one-step correction of the plug-in estimator using the efficient influence function. What we did *not* do is **cross-fitting** (sample splitting), which is a separate device. Second, that matters for inference: because the manual example fits the SuperLearner nuisances on the *full* sample and then uses

$\text{sd}(\text{IF})/\sqrt{n}$, the resulting confidence interval is only valid under empirical-process (Donsker) conditions that flexible ML learners typically violate. With cross-fitting, the influence-function variance is valid without those conditions. “npcausal” does the cross-fitting for you when you set `nsplits > 1` (here `nsplits = 1` turns it off, for comparison with the manual full-sample version).

```
aipw2<- ate(y=data$Y, a=data$A, x=W, nsplits=10)
```

	0%
==	2%
====	5%
=====	8%
=====	10%
=====	12%
=====	15%
=====	18%
=====	20%
=====	22%
=====	25%
=====	28%
=====	30%
=====	32%
=====	35%
=====	38%
=====	40%
=====	42%
=====	45%

32. Recent causal inference tools

							48%
							50%
							52%
							55%
							58%
							60%
							62%
							65%
							68%
							70%
							72%
							75%
							78%
							80%
							82%
							85%
							88%
							90%
							92%
							95%
							98%
							100%
	parameter	est	se	ci.ll	ci.ul	pval	
1	E{Y(0)}	0.4448128	0.02529634	0.39523200	0.4943936	0	
2	E{Y(1)}	0.5775070	0.02255959	0.53329021	0.6217238	0	
3	E{Y(1)-Y(0)}	0.1326942	0.03377055	0.06650392	0.1988845	0	

32.4.3. DoubleML

```
library(DoubleML)
dml_data = DoubleMLData$new(as.data.frame(data),
                             y_col = "Y",
                             d_cols = "A",
                             x_cols = c("eligibility" , "loss_num","age", "time_try_pregnant",
print(dml_data)
```

```
===== DoubleMLData Object =====
```

```
----- Data summary -----
Outcome variable: Y
Treatment variable(s): A
Covariates: eligibility, loss_num, age, time_try_pregnant, BMI, meanAP
Instrument(s):
Selection variable:
No. Observations: 1228
```

```
library(mlr3)
library(mlr3learners)
# suppress messages from mlr3 package during fitting
lgr::get_logger("mlr3")$set_threshold("warn")

learner = lrn("regr.ranger", num.trees=500, max.depth=5, min.node.size=2)
ml_g = learner$clone()
ml_m = learner$clone()

learner = lrn("regr.glmnet")
ml_g_sim = learner$clone()
ml_m_sim = learner$clone()

set.seed(3141)
obj_dml_plr = DoubleMLPLR$new(dml_data, ml_g=ml_g, ml_m=ml_m)
obj_dml_plr$fit()
print(obj_dml_plr)
```

```
===== DoubleMLPLR Object =====
```

```
----- Data summary -----
Outcome variable: Y
Treatment variable(s): A
Covariates: eligibility, loss_num, age, time_try_pregnant, BMI, meanAP
```

32. Recent causal inference tools

```
Instrument(s):
Selection variable:
No. Observations: 1228

----- Score & algorithm -----
Score function: partialling out
DML algorithm: dml2

----- Machine learner -----
ml_l: regr.ranger
ml_m: regr.ranger

----- Resampling -----
No. folds: 5
No. repeated sample splits: 1
Apply cross-fitting: TRUE

----- Fit summary -----
Estimates and significance testing of the effect of target variables
Estimate. Std. Error t value Pr(>|t|)
A 0.12505 0.03292 3.798 0.000146 ***
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Note: as with the earlier PLR example, **A** and **Y** here are **binary** (this is the EAGeR trial data), so `DoubleMLIRM` (which uses classification learners for the propensity score and allows heterogeneous effects) would be the more appropriate DoubleML estimator than `DoubleMLPLR`'s partially linear, constant-effect model — see the note above.


Systematic treatment: [R](#) · [Julia](#).

33. Intro to proximal causal inference

This is based on papers of Tchetgen Tchetgen and co-authors, mostly based on Xu Shi's talk on proximal causal inference.

33.1. Unobserved confounder

Unobserved confounder is the problem for observational studies. Assuming controlling for all relevant confounders (X), we have an unobserved confounder U that affects both treatment A and outcome Y . The causal diagram is as follows:




In this case, we have biased estimates of the treatment effect, regressing Y on A controlling for X .

33.2. Proximal causal inference

Proximal causal inference (also called negative controls) is a framework that allows us to estimate the causal effect of treatment A on outcome Y while accounting for unobserved confounders.

We'll need a DAG like this:




Basically we need two proxies for U to control for confounding. Z is called Negative Control for Exposure (NCE) and W is called Negative Control for Outcome (NCO). There is no link from Z to Y , and no link from W to A or Z .


Z is an NCE if $Y(a, z) = Y(a)$ and $Z \perp Y(a)|U$, meaning that Z satisfies the exclusion restriction that there is no direct link from Z to Y , and that Z is independent of Y given U .

W is an NCO if $W(a, z) = W$ and $W \perp (A, Z)|U$, meaning that A and Z have no direct link to W .

33.2.1. Why this works




If we regress Y on A , then we get $\beta_{ay} + \beta_{uy}\beta_{ua}$, which is biased. How do we recover β_{ay} ?



Suppose we have W as proxy of U and $\beta_{uw} = \beta_{uy}$, then we can get the true effect. We can do regression of Y on A , and then W on A . The first regression gives us $\beta_{ay} + \beta_{uy}\beta_{ua}$, and the second regression gives us $\beta_{uw}\beta_{ua}$. We can then get β_{ay} from the difference of the two coefficients. This is similar to a DiD setting. Say W is Y_{t-1} , and assume the effect of U is the same on Y and Y_{t-1} , then effect of A on Y by the difference of regressing Y on A and Y_{t-1} on A .

33. Intro to proximal causal inference

However, most likely $\beta_{uw} \neq \beta_{uy}$. Then we need another proxy Z of U to recover the true effect. This is also called “double proxy”, or “double negative control”.



33.2.2. A linear model

Suppose we have the following linear model:

$$E[Y|A, Z, U] = \beta_{y0} + \beta_{ay}A + \beta_{uy}U \quad (33.1)$$

$$E[W|A, Z, U] = \beta_{w0} + \beta_{uw}U \quad (33.2)$$

$$E[U|A, Z] = \beta_{u0} + \beta_{zu}Z + \beta_{au}A \quad (33.3)$$

β_{ay} is the true effect of A on Y .

$$E[Y|A, Z] = \beta_{y0} + \beta_{ay}A + \beta_{uy}(\beta_{u0} + \beta_{zu}Z + \beta_{au}A) \quad (33.4)$$

$$E[W|A, Z] = \beta_{w0} + \beta_{uw}(\beta_{u0} + \beta_{zu}Z + \beta_{au}A) \quad (33.5)$$

From the last two regressions we can recover β_{ay} :

If we define $\delta_a^y = \beta_{ay} + \beta_{uy}\beta_{au}$ (regression coefficient of A on the Y regression) and $\delta_a^w = \beta_{uw}\beta_{au}$, (regression coefficient of A on the W regression), similarly for δ_z^y and δ_z^w , then we have

$$\beta_{ay} = \delta_a^y - \delta_a^w \frac{\delta_z^y}{\delta_z^w} \quad (33.6)$$

In this **linear-Gaussian special case only**, we can look at this as a two-stage-least-squares-style regression and replace U with the predicted value of W :

$$E[Y|A, Z] = \left(\beta_{y0} - \frac{\beta_{uy}}{\beta_{uw}} \beta_{w0} \right) + \beta_{ay}A + \frac{\beta_{uy}}{\beta_{uw}} E[W|A, Z] \quad (33.7)$$

since $E[U | A, Z] = (E[W | A, Z] - \beta_{w0})/\beta_{uw}$ from the W equation above. The coefficient on the predicted value of W is β_{uy}/β_{uw} , not β_{uy} alone — it only reduces to β_{uy} when $\beta_{uw} = 1$. So we first regress W on A and Z , then regress Y on A and the predicted value of W .

This “predict W from A and Z , then plug it in” shortcut is **not** the general proximal identification argument, and it is not the same as a standard IV/2SLS specification (we are not instrumenting an endogenous regressor with `ivreg`). In general proximal causal inference, the outcome bridge function $h(A, W)$ is identified by solving the conditional moment restriction $E[Y | A, Z] = E[h(A, W) | A, Z]$, which requires completeness (relevance) conditions linking the treatment-inducing proxy Z and the outcome-inducing proxy W to the hidden confounder U . The generated regressor $S = \hat{E}[W | A, Z]$ is only a valid substitute under the linear-Gaussian structure above; even then the two-stage standard errors are wrong as written, because S is generated in the first stage (bootstrap the whole procedure).

33.2.3. A nonparametric model

We don’t have to have a linear model, this can be done with a nonparametric model as well.

$$E[Y(a)] = E[h(a, W)] \quad (33.8)$$

where $E[Y|A, Z] = E[h(A, W)|A, Z]$.


That is, we can estimate h as a function of A and W , identified using Z , nonparametrically. This h function is called outcome bridge function. In the linear case, it is $h(A, W) = \beta_0 + \beta_a A + \beta_w W$.

33.2.4. How to find proximal variables

Tchetgen Tchetgen and co-authors come up with a Data-driven Automated Negative Control Estimation (DANCE) algorithm to find negative controls. The idea is to use a machine learning algorithm to find the variables that are correlated with the treatment and the outcome, but not with the treatment and the outcome together. This is similar to the idea of finding instrumental variables.

33.3. example

Let’s generated data based on this DAG:



```

set.seed(123)
n <- 5000
# Set coefficients
b_uw <- 0.5 # effect of U on W
b_uy <- 0.5 # effect of U on Y
b_ay <- 0.5 # effect of A on Y
b_uz <- 0.5 # effect of U on Z
b_ua <- 0.5 # effect of U on A
b_z_a <- 0.5 # effect of Z on A

U <- rnorm(n)
W <- b_uw * U + rnorm(n) # W is a function of U
Z <- b_uz * U + rnorm(n) # Z is a function of U
A <- b_ua * U + b_z_a * Z + rnorm(n) # A is a function of U and Z
Y <- b_uy * U + b_ay * A + rnorm(n) # Y is a function of U and A
# A <- rbinom(n, 1, plogis(-1 + U / 2)) # A is a binary treatment

# Create a data frame with the generated variables

data <- data.frame(W = W, Z = Z, A = A, Y = Y)

```

If we run regression of Y on A , Z , W ,

```

m1 <- lm(Y ~ A + Z + W, data = data)
summary(m1)

```

```
Call:
lm(formula = Y ~ A + Z + W, data = data)

Residuals:
    Min       1Q   Median       3Q      Max
-4.5836 -0.7167  0.0005  0.7305  4.4866

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.006347   0.015279  -0.415   0.678
A             0.660021   0.014123  46.734 < 2e-16 ***
Z             0.085004   0.016650   5.105 3.43e-07 ***
W             0.132798   0.014165   9.375 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.08 on 4996 degrees of freedom
Multiple R-squared:  0.4556,    Adjusted R-squared:  0.4553
F-statistic: 1394 on 3 and 4996 DF,  p-value: < 2.2e-16
```

We don't get the true effect of A on Y .

If we do the 2sls using W and Z , then we get the true effect.

```
# Linear-Gaussian (Y, W) toy case ONLY. This two-stage plug-in identifies
# the effect under the linear structure above; it is not the general
# proximal estimator, and the reported SEs below are NOT valid because S is
# a generated regressor (bootstrap the whole two-step procedure for inference).
# Perform the first-stage linear regression W ~ A + Z
m1 <- lm(W ~ A + Z, data = data)
# Compute the estimated E[W|A,Z]
S <- m1$fitted.values
# Perform the second-stage linear regression Y ~ A + E[W|A,Z]
m2 <- lm(Y ~ A + S, data = data)
summary(m2) # point estimate OK under the linear model; SEs need bootstrap
```

```
Call:
lm(formula = Y ~ A + S, data = data)

Residuals:
    Min       1Q   Median       3Q      Max
-4.5477 -0.7309  0.0059  0.7395  4.3715

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.005049   0.015408  -0.328   0.743
```

33. Intro to proximal causal inference

```
A          0.489342    0.043210  11.325  < 2e-16 ***
S          1.143480    0.199171   5.741  9.96e-09 ***
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.089 on 4997 degrees of freedom

Multiple R-squared: 0.446, Adjusted R-squared: 0.4458

F-statistic: 2012 on 2 and 4997 DF, p-value: < 2.2e-16


Systematic treatment: *R* · *Julia*.

34. Simulation on proximal causal inference with panel data


This is to continue the discussion on proximal causal inference. The idea is to use negative controls to estimate the causal effect of treatment on outcome while accounting for unobserved confounders. This is a follow-up of the previous chapter on [proximal causal inference](#).

In this post, I am trying to see whether we can use lagged D and lagged Y as NCE and NCO, in a panel data setting. The idea is to use the lagged treatment as a negative control exposure (NCE) and the lagged outcome as a negative control outcome (NCO).

Let's see a two period DAG:



Suppose that's the case for the causal structure, then I can reasonably combine $U1$ and $U2$ into a single unobserved confounder U :



That would fit into the double negative control framework. The problem is that we have to assume there is no link from A1 to Y1, which is odd. Can we assume in period 1 there is no effect, but in period 2, there is an effect?

34.1. A simulation with panel data

I generate two period data based on the DAG above.

```

library(dplyr)
library(rootSolve)
library(tidyverse)

gen_data <- function(seed, n){
  set.seed(seed)
  u <- rnorm(n, mean = 0, sd = 1) # unobserved confounder
  x1 <- rnorm(n, mean = 0, sd = 1) # observed confounder
  x2 <- rnorm(n, mean = 0, sd = 1) # observed confounder
  a1 <- rbinom(n, 1, plogis(2*u + 1.5*x1)) # treatment at time 1
  a2 <- rbinom(n, 1, plogis(u + a1 + x1)) # treatment at time 2
  y1 <- 1.5*u + .5*x1 + rnorm(n, mean = 0, sd = 1) # outcome at time 1
  y2 <- y1 + 2.5*u + 2.5*x1 + 3*a2 + rnorm(n, mean = 0, sd = 1) # outcome at time 2
  data <- data.frame(u=u, x1=x1, x2=x2, a1=a1, a2=a2, y1=y1, y2=y2)
  return(data)
}

data <- gen_data(333, 100000)
  
```

```
data_sim <- data |>
  select(a1, a2, y2, x1, x2, y1) |>
  rename(a = a2, y = y2, x = x1, w = y1, z=a1)
```

We can use the simple 2sls:

```
lm1 <- lm(w ~ a + x + z, data = data_sim)
data_sim$what <- lm1$fitted.values
lm2 <- lm(y ~ a + x + what, data = data_sim)
summary(lm2)
```

Call:

```
lm(formula = y ~ a + x + what, data = data_sim)
```

Residuals:

Min	1Q	Median	3Q	Max
-16.4284	-2.3083	0.0071	2.3020	20.7590

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.01293	0.02315	-0.558	0.577
a	3.02113	0.03528	85.625	<2e-16 ***
x	1.67203	0.01245	134.327	<2e-16 ***
what	2.66578	0.01747	152.631	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.456 on 99996 degrees of freedom

Multiple R-squared: 0.6822, Adjusted R-squared: 0.6822

F-statistic: 7.155e+04 on 3 and 99996 DF, p-value: < 2.2e-16

Or we can use more sophisticated methods. Here I use the m-estimation code by Paul Zivich: https://github.com/pzivich/publications-code/blob/master/ProximalCI/proximal_sims_r.R. This code is to construct estimating equations, then use root finding to find the parameter estimates.

```
pcl_m2 <- function(theta, mydata, out = "ef"){
  # read in variables
  a <- mydata$a
  x <- mydata$x
  w <- mydata$w
  z <- mydata$z
  y <- mydata$y
  # model for what
  what <- theta[1] + theta[2] * a + theta[3] * x + theta[4] * z
```

34. Simulation on proximal causal inference with panel data

```

# matrix of covars in model for what
wmat <- as.matrix(cbind(rep(1, nrow(mydata)), a, x, z), ncol = 4,
                  nrow = nrow(mydata))

# estimating function for what
ef1 <- (t(wmat) %*% (w - what))          # gives px1 vector (summed over units)
ef1_b <- wmat * (w - what)               # gives nxp matrix (rows are units)

# model for yhat
yhat <- theta[5] + theta[6]*a + theta[7]*x + theta[8] * what

# matrix of covars in model for yhat
ymat <- as.matrix(cbind(rep(1, nrow(mydata)), a, x, what), ncol = 4,
                  nrow = nrow(mydata))

#estimating function for yhat
ef2 <- t(ymat) %*% (y - yhat)            # px1 vector, summed over units
ef2_b <- (ymat) * (y - yhat)            # nxp matrix (rows are units)
f <- c(ef1, ef2)                        # p-dimensional vector
if(out == "full") return(cbind(ef1_b, ef2_b)) else return(f) # full used in var calculation
}

# call m-estimator for PCL
mpcl_mest <- function(dat){
  theta_start <- rep(.05, 8)
  results <- mestimator(pcl_m2, init = theta_start, dat, mydata = dat)
  return(results)
}

##' M-estimation engine
mestimator <- function(ef, init, dat, ...){
  # get estimated coefficients
  fit <- multiroot(ef, start = init, out = "ef",
                  rtol = 1e-9, atol = 1e-12, ctol = 1e-12, ... )
  betahat <- fit$root
  n <- nrow(dat)

  # bread
  # derivative of estimating function at betahat
  pd <- gradient(ef, betahat, out = "ef", ...)
  pd <- as.matrix(pd)
  bread <- -pd/n
  se_mod <- sqrt(abs(diag(-solve(pd))))

  # meat1
  efhat <- ef(betahat, out = "full", ...)
  meat <- list()

```



```

meat1 <- (t(efhat) %*% efhat)/n

# sandwich
sandwich1 <- (solve(bread) %*% meat1 %*% t(solve(bread)))/n
se_sandwich1 <- sqrt(diag(sandwich1))

results <- data.frame(betahat, se_sandwich1, se_mod)
return(results)
}

rmpcl <- unname(mpcl_mest(data_sim)[6, c(1:2)])
rmpcl

```

```
6 3.021126 0.0198433
```

```
#mpcl_mest(data_sim)
```

Both the 2sls and the m-estimation give us the same estimate of the causal effect of treatment on outcome, which is 3. The m estimator can be more general. 2sls is good for the linear case.

However, the main issue is the assumption that there is no link from A_1 to Y_1 .

To think about it from a different angle: this is similar to DiD setting. Basically we need A_1 to be 0. There is no effect in period 1. If we only have A_1 as 0, then it makes sense to have no link from A_1 to Y_1 (no anticipation assumption).

This is a restrictive identifying assumption, not a technical footnote. Using A_{t-1} as the negative control exposure and Y_{t-1} as the negative control outcome only identifies the effect if two separate exclusions hold: A_1 must have no direct effect on Y_1 , so that Y_1 is a clean proxy for U rather than a partly post-treatment quantity; and A_1 must have no direct effect on Y_2 except through U , which is the exclusion restriction that makes A_1 admissible as a negative control exposure in the first place. If either link is actually present — for example, if treatment has an immediate effect in the same period it starts, which is the empirically common case — this proximal construction is misspecified and the estimate above should not be trusted. Before applying this approach, researchers should have a substantive reason to believe treatment effects only appear with a lag.

Systematic treatment: [R](#) · [Julia](#).

35. longitudinal modified treatment policy (LMTP)

I read a few papers with longitudinal modified treatment policy (LMTP), and found it interesting. It has been used in epidemiology and biostatistics, but I have not seen it in applied econometrics yet.

Here I am mostly following Nicholas Williams: <https://beyondtheate.com/>

We are usually interested in ATE, the average treatment effect. However, there could be more complicated situations that the treatment is continuous, or the treatment is multivalued, or the treatment is time-varying. The static interventions have problems. For example, the hypothetical interventions that treatment applies to everyone might be inconceivable. Or such intervention could make positivity assumption fail.


Suppose we have such a DAG:

```
library(ggplot2)
library(ggdag)

g <- dagify(
  A1 ~ L1,
  L2 ~ L1,
  L2 ~ A1,
  A2 ~ A1,
  A2 ~ L2,
  Y ~ A2,
  Y ~ L2,
  Y ~ A1,
  exposure = "A2",
  outcome = "Y",
  coords = list(x = c(L1 = 1, L2 = 2, Y = 3, A1 = 1.5, A2 = 2.5),
                y = c(L1 = 1, L2 = 1, Y = 1, A1 = 1.5, A2 = 1.5))
)

ggdag(g) +
  theme_dag()
```

35. longitudinal modified treatment policy (LMTP)



We have multiple time points, and the treatment A is time-varying. We are interested in not only the ATE of A_2 on Y , but also some other hypothetical interventions.

35.1. Definitions and assumptions

35.1.1. Assumptions

1. Positivity. Basically if there is a unit with a_t and h_t , then there is a unit with $d(a_t, h_t)$ and h_t .
2. Sequential unconfoundedness. There is no unmeasured confounders.

35.1.2. dynamic treatment regime

Some notations commonly used in this literature: We observe $Z = (L_0, A_0, L_1, A_1, Y)$, where L is the observed confounder, A is the treatment, and Y is the outcome. The treatment A can be time-varying. We can also have multiple treatments at different time points. For example, we can have A_2, A_3 , etc. In this graph, baseline L_0 affects A_0 , which in turn affects L_1 , which then affects A_1 , and finally A_1 affects the outcome Y .

History H_t is the history of data up to time t , right before A_t . for example, $H_1 = (L_0, A_0, L_1)$, and $H_2 = (L_0, A_0, L_1, A_1)$. d is the hypothetical intervention function, or shift function, which is a function of the history H_t and the treatment A_t . For example, $d_0(a_0, h_0, \epsilon_0)$ is a user-given function to map a_0, h_0 , and ϵ_0 to a potential treatment value. The function d can be deterministic, or it can be stochastic. Then we can replace A_0 with $A_0^d = d_0(A_0, H_0, \epsilon_0)$. Then after that $A_1(A_0^d)$ is called the natural value of treatment.

This is very general, comparing to the static treatment regime.

Suppose treatment A is a function of the history of treatment and confounders, for example, $A = d(A_1, L_1, A_2, L_2)$. A can be set to a fixed value, say 1 or 0, or some value A^d . This function d can be anything, it can be taking a deterministic value, or it can be a function that takes the natural treatment value A as input. In the package “lmtp”, this is called a shift function, or hypothetical intervention. For example, d can be set to 1 if $age < 30$, or d can be set to double the natural value of A . Many possibilities.

In comparison, for ATE, we only need to set A to 0 or 1.

Under this LMTP, the causal parameter is

$$\theta = E[Y^{\bar{A}^d}] \quad (35.1)$$

$Y^{\bar{A}^d}$ is the potential outcome under the hypothetical intervention $\bar{A}^d$. At time 1, $A_1^d = d(A_1, H_1)$.

35.1.3. modified treatment policy

Let's look at a simulated data set to see how exactly we can estimate it.

This simulation is from Susmann et al. (2024) “Longitudinal Generalizations of the Average Treatment Effect on the Treated for Multi-valued and Continuous Treatments”. I modified slightly to fit the DAG above.

```
library(tidyverse)
library(tidyr)

mtp <- function(data, trt) {
  a <- data[[trt]]
  a * 0 + 1
}

simulate_data <- function(seed, N, tau, sigma = 0.5) {
  set.seed(seed)

  data <- tibble(id = 1:N)

  for(t in 1:tau) {
    Lt <- paste0("L_", t)
    Ltd <- paste0("L_", t, "d")

    At <- paste0("A_", t)
    Atd <- paste0("A_", t, "d")

    Lt1 <- paste0("L_", t - 1)
    Lt1d <- paste0("L_", t - 1, "d")

    At1 <- paste0("A_", t - 1)
    At1d <- paste0("A_", t - 1, "d")
  }
}
```

35. longitudinal modified treatment policy (LMTP)

```

if(t == 1) {
  data[[Lt]] <- runif(N, 0, 1)
  data[[Ltd]] <- data[[Lt]]

  data[[At]] <- rbinom(N, size = 1, prob = 0.5)
}
else {
  # L_t depends on the previous treatment A_{t-1} (the DAG edge A1 -> L2),
  # which is the time-varying confounding this chapter is about.
  data[[Lt]] <- rnorm(N, mean = 0.25 * data[[Lt1]] + 0.4 * data[[At1]], 0.5)
  data[[Ltd]] <- rnorm(N, mean = 0.25 * data[[Ltd1]] + 0.4 * data[[At1d]], 0.5)

  data[[At]] <- rbinom(N, size = 1, prob = plogis(0.5 - 0.2 * data[[At1]] + 0.1 * data[[Lt1]])
}

data[[Atd]] <- mtp(data, At)
}
# Y depends on the final treatment A_2 and confounder L_2, and also directly
# on the first-period treatment A_1 (the DAG edge A1 -> Y).
data$Y <- rnorm(N, data[[At]] + data[[Lt]] + 0.5 * data[["A_1"]], sigma)
data$Yd <- rnorm(N, data[[Atd]] + data[[Ltd]] + 0.5 * data[["A_1d"]], sigma)
data
}

simulated_data1 <- simulate_data(seed = 123, N = 10000, tau = 2)

mean(simulated_data1$Yd)

```

```
[1] 2.028593
```

```

mtp <- function(data, trt) {
  a <- data[[trt]]
  a * 0 + 0
}

simulated_data2 <- simulate_data(seed = 123, N = 10000, tau = 2)

mean(simulated_data2$Yd)

```

```
[1] 0.1285934
```

```
mean(simulated_data1$Yd) - mean(simulated_data2$Yd)
```

```
[1] 1.9
```

Note in this simulation the variables ending with “d” are the variables under hypothetical intervention, or modified treatment policy. L_1 is from $\text{Uniform}(0, 1)$, and L_2 is from $N(0.25L_1 + 0.4A_1, 0.5)$ – so the time-varying confounder L_2 depends on the earlier treatment A_1 , which is the time-varying confounding this chapter is about. The treatment A_1 is from a Bernoulli distribution with probability 0.5, and the treatment A_2 is from a Bernoulli distribution with probability $\text{plogis}(0.5 - 0.2A_1 + 0.1L_1)$. The outcome Y is from a normal distribution with mean $A_2 + L_2 + 0.5A_1$, so Y depends on both the current treatment A_2 and the lagged treatment A_1 . In the code, the first modified treatment policy sets treatment to 1 in every period (`simulated_data1`) and the second sets it to 0 in every period (`simulated_data2`); the reported difference $E[Y_{\text{always } 1}^d] - E[Y_{\text{always } 0}^d]$ is the effect of the always-treat versus never-treat static regimes.

In this case, we can just do a linear regression to get the effect of A_2 on Y , knowing the exact DAG.

```
m1 <- glm(Y ~ L_2 + A_1 + A_2, data = simulated_data1)
summary(m1)
```

Call:

```
glm(formula = Y ~ L_2 + A_1 + A_2, data = simulated_data1)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-3.453e-05	9.615e-03	-0.004	0.997
L_2	9.998e-01	9.889e-03	101.103	<2e-16 ***
A_1	4.940e-01	1.079e-02	45.779	<2e-16 ***
A_2	1.000e+00	1.027e-02	97.380	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.2500364)

Null deviance: 9257.2 on 9999 degrees of freedom
 Residual deviance: 2499.4 on 9996 degrees of freedom
 AIC: 14523

Number of Fisher Scoring iterations: 2

Let’s try a different MTP: set half of the time to 0, the other half remain unchanged.

$$d(a_t, \epsilon_t) = \begin{cases} 0, & \text{if } \epsilon_t < .5 \text{ and } a_t = 1 \\ a_t, & \text{otherwise} \end{cases} \quad (35.2)$$

This is, say, to set half of smokers to non-smokers, and the other half remain smokers.

35. longitudinal modified treatment policy (LMTP)

```
# mtp <- function(data, trt) {  
#   a <- data[[trt]]  
#   epsilon <- rbinom(nrow(data), size = 1, prob = 0.5)  
#   ifelse(epsilon < 0.5 & a == 1, 0, a)  
# }  
d <- function(a) {  
  epsilon <- runif(length(a))  
  ifelse(epsilon < 0.5 & a == 1, 0, a)  
}  
simulated_data1$m3_d <- simulated_data1$Y  
  
m2 <- glm(m3_d ~ L_1 + A_1 + L_2 + A_2, data = simulated_data1)  
summary(m2)
```

Call:

```
glm(formula = m3_d ~ L_1 + A_1 + L_2 + A_2, data = simulated_data1)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.005815	0.012842	-0.453	0.651
L_1	0.011964	0.017621	0.679	0.497
A_1	0.494512	0.010814	45.729	<2e-16 ***
L_2	0.998831	0.009987	100.015	<2e-16 ***
A_2	0.999873	0.010273	97.334	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.2500499)

Null deviance: 9257.2 on 9999 degrees of freedom
Residual deviance: 2499.2 on 9995 degrees of freedom
AIC: 14525

Number of Fisher Scoring iterations: 2

```
simulated_data1$m2_d <- predict(m2, mutate(simulated_data1, A_2 = d(A_2)))  
  
m1 <- glm(m2_d ~ L_1 + A_1, data = simulated_data1)  
summary(m1)
```

Call:

```
glm(formula = m2_d ~ L_1 + A_1, data = simulated_data1)
```

Coefficients:


```

              Estimate Std. Error t value Pr(>|t|)
(Intercept)  0.30716    0.01531   20.07  <2e-16 ***
L_1          0.26079    0.02384   10.94  <2e-16 ***
A_1          0.87604    0.01367   64.08  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.4669449)

Null deviance: 6633.7  on 9999  degrees of freedom
Residual deviance: 4668.0  on 9997  degrees of freedom
AIC: 20768

Number of Fisher Scoring iterations: 2

```

```

m1_d <- predict(m1, mutate(simulated_data1, A_1 = d(A_1)))
mean(m1_d)

```

```
[1] 0.6518997
```

Note this recursive process is based on Diaz, et al. (2023) “Nonparametric Causal Effects Based on Longitudinal Modified Treatment Policies”.

Here is how exactly we estimate it: we start with the last time point. Regress it on previous treatment and confounders, and then get the predicted value with A changed based on the MTP. Regress that predicted value on the previous treatment and confounders, get predicted values with A changed based on MTP. Repeat until the first time point. The average of the predicted value at time 1 is the expected value under this MTP.

The shift function `d()` above is a *stochastic* intervention: at each call it draws a fresh $\epsilon \sim \text{Uniform}(0, 1)$ and, with probability 0.5, switches a treated unit ($a = 1$) to 0, otherwise leaves a unchanged. The regime is “at each time point, draw independently conditional on history and apply this rule”; because the draw is independent across time points and units, calling `d()` separately for A_2 and A_1 is what implements that regime. (By contrast, the `d1/d2` functions used with `lmtm_tmle` below are *deterministic* static regimes – always 1 and always 0.)

This is basically g-formula extended to longitudinal data.

In a general case, this is the generalized g-formula to estimate θ :

Set $m_{\tau+1} = Y$, let $A_t^d = d(A_t, H_t)$. For $t = \tau, \dots, 1$, recursively define:

$$m_t(a_t, h_t) = E[m_{t+1}(A_{t+1}^d, H_{t+1}) | A_t = a_t, H_t = h_t] \quad (35.3)$$

Then $\theta = E[m_1(A_1^d, L_1)]$.

We start from the last period. Regress $m_{\tau+1}$, which is Y on A_τ and H_τ . Then get the predicted value with A_τ changed based on the MTP. Then regress that predicted value on $A_{\tau-1}$ and $H_{\tau-1}$. Repeat until the first time point. The average of the predicted value at time 1 is the expected value under this MTP.

35. longitudinal modified treatment policy (LMTP)

35.1.4. Estimators

The authors advocate two estimators, TMLE and SDR (sequentially doubly robust estimator). The procedures are the same, starting from the last time point, then apply TMLE or SDR, iterate to the first time point.

35.2. Using lmtip package

```
library(tidyverse)
library(lmtip)
d1 <- function(data, trt) {
  rep(1, nrow(data))
}
A <- "A_2"
Y <- "Y"
W <- c("L_1", "L_2", "A_1")

set.seed(34465)

treat <- lmtip_tmle(
  data = simulated_data1,
  trt = "A_2",
  outcome = "Y",
  baseline = W,
  outcome_type = "continuous",
  shift = d1,
  folds = 1,
  learners_trt = "SL.glm",
  learners_outcome = "SL.glm"
)

print(treat)
```

```
::: {.cell-output .cell-output-stdout}
```

```
      Estimate: 1.558
Std. error: 0.009
```

```
:::
```

```
d2 <- function(data, trt) {
  rep(0, nrow(data))
}
A <- "A_2"
```

```

Y <- "Y"
W <- c("L_1", "L_2", "A_1")

set.seed(34465)

control <- lmtmle(
  data = simulated_data1,
  trt = "A_2",
  outcome = "Y",
  baseline = W,
  outcome_type = "continuous",
  shift = d2,
  folds = 1,
  learners_trt = "SL.glm",
  learners_outcome = "SL.glm"
)

print(control)

```

```
:: { .cell-output .cell-output-stdout }
```

```

      Estimate: 0.559
Std. error: 0.01

```

```
lmtm_contrast(treat, ref = control)
```

```
:: { .cell-output .cell-output-stdout }
```

```

      shift   ref estimate std.error conf.low conf.high p.value
1  1.56 0.559          1    0.0102     0.98     1.02 <0.001

```

```
:: ::
```

Note that `folds = 1` turns off cross-fitting (sample splitting). That is fine for a small, fast demonstration with `SL.glm`, but in applied work with flexible learners you should use several folds so the nuisance fits are cross-fitted; otherwise the standard errors can be anti-conservative.

Note on the estimand: this `lmtmle` example only shifts A_2 (`trt = "A_2"`), treating A_1 as an observed baseline covariate (via `baseline = W`) rather than also intervening on it. That is a different, single-period policy from the always-1-vs-always-0 *both-period* regime computed manually via `simulated_data1$Yd/simulated_data2$Yd` above (lines 139-152) — the two numbers are not expected to match, and should not be used as a sanity check against each other. To replicate the both-period always-treat/never-treat contrast with `lmt`, shift both time points by passing `trt = c("A_1", "A_2")` and supplying `time_vary` (the time-varying covariates measured before each treatment time, e.g. `list(character(0), "L_2")`), with `shift` functions that map each treatment in the vector to 1 (or 0).

35.3. another example

The data set `bmi`, from the `DynTxRegime` package, are simulated to reflect a two-stage RCT {A1, A2} that studied the effect of meal replacement (MR) shakes versus a calorie deficit (CD) diet on adolescent

35.3.1. shift function 1

Consider a shift function that assigns meal replacement to all observations at time 1, but only meal replacement at time 2 to those observations whose 4-month BMI is greater than 30.

```
library(DynTxRegime)
data(bmiData)
bmi <- bmiData
glimpse(bmi)
```

```
Rows: 210
Columns: 8
$ gender      <int> 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1~
$ race        <int> 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0~
$ parentBMI   <dbl> 31.59683, 30.17564, 30.27918, 27.49256, 26.42350, 29.30970~
$ baselineBMI <dbl> 35.84005, 37.30396, 36.83889, 36.70679, 34.84207, 36.68640~
$ month4BMI   <dbl> 34.22717, 36.38014, 34.42168, 32.52011, 33.72922, 32.06622~
$ month12BMI  <dbl> 34.27263, 36.38401, 34.41447, 32.52397, 33.73546, 32.15977~
$ A1          <chr> "CD", "CD", "MR", "CD", "CD", "MR", "CD", "CD", "CD", "CD"~
$ A2          <chr> "MR", "MR", "CD", "CD", "CD", "MR", "MR", "CD", "MR", "MR"~
```

```
d_dtr <- function(data, trt) {
  if (trt == "A1") return(rep("MR", nrow(data)))

  ifelse(data$month4BMI > 30, "MR", "CD")
}

fit_dtr <- lmtp_sdr(
  data = bmi,
  trt = c("A1", "A2"),
  outcome = "month12BMI",
  baseline = c("gender", "race", "parentBMI"),
  time_vary = list("baselineBMI", "month4BMI"),
  shift = d_dtr,
  outcome_type = "continuous",
  folds = 1,
  learners_trt = "SL.glm",
  learners_outcome = c("SL.mean", "SL.glm", "SL.gam")
)
```

```
fit_dtr
```

```
::: {.cell-output .cell-output-stdout}
```

```
      Estimate: 35.854
Std. error: 0.362
```

```
:::
```

35.3.2. shift function 2

Suppose we are interested in comparing the dynamic treatment regime to a static treatment regime where all patients receive meal replacement at both time points. Using the SDR estimator, estimate the effect of this static intervention.

```
fit_MR <- lmtp_sdr(
  data = bmi,
  trt = c("A1", "A2"),
  outcome = "month12BMI",
  baseline = c("gender", "race", "parentBMI"),
  time_vary = list("baselineBMI", "month4BMI"),
  shift = \(data, trt) rep("MR", nrow(data)),
  outcome_type = "continuous",
  folds = 1,
  learners_trt = "SL.glm",
  learners_outcome = c("SL.mean", "SL.glm", "SL.gam")
)

fit_MR
```

```
::: {.cell-output .cell-output-stdout}
```

```
      Estimate: 35.831
Std. error: 0.361
```

```
:::
```

Let's also estimate the effect of an intervention where all patients receive a calorie deficit diet at both time points.

35. longitudinal modified treatment policy (LMTP)

```
fit_CD <- lmtp_sdr(  
  data = bmi,  
  trt = c("A1", "A2"),  
  outcome = "month12BMI",  
  baseline = c("gender", "race", "parentBMI"),  
  time_vary = list("baselineBMI", "month4BMI"),  
  shift = \(data, trt) rep("CD", nrow(data)),  
  outcome_type = "continuous",  
  folds = 1,  
  learners_trt = "SL.glm",  
  learners_outcome = c("SL.mean", "SL.glm", "SL.gam")  
)  
  
fit_CD
```

```
::: {.cell-output .cell-output-stdout}
```

```
      Estimate: 35.053  
Std. error: 0.301
```

```
:::
```

Finally, we can compare the three treatment regimes using the `lmtp_contrast` function.

```
lmtp_contrast(fit_dtr, fit_MR, ref = fit_CD)
```

```
::: {.cell-output .cell-output-stdout}
```

	shift	ref	estimate	std.error	conf.low	conf.high	p.value
1	35.9	35.1	0.801	0.335	0.144	1.46	0.0169
2	35.8	35.1	0.778	0.335	0.122	1.43	0.0201

```
:::
```

Systematic treatment: *R* · *Julia*.

Part VII.

Advanced Topics

36. Doing the same multi-level model with R and stata

This post is a demo for doing the same multi-level models with R and stata. R's multilevel models are mostly done with `lme4` and `nlme`. Stata's multilevel models are mostly done with `mixed` and `menl`. The two packages are not exactly the same. `mixed` does in fact support several residual covariance structures via its `residuals()` option (e.g. `residuals(exponential, t(year))`), as well as `ar`, `ma`, `toeplitz`, etc.) — so it's not that `mixed` cannot handle a structured residual at all, but `menl`'s nonlinear-mixed-model framework gives more flexibility to combine that with custom nonlinear mean/random-effect specifications (like the CEM model below), which is why we reach for it here. Likewise, in R `lme4`'s “`lmer`” does not allow this specification. We have to use `nlme`. But it's good to know how to do the same thing in both packages.

For example, we want to do a consensus emergence model (CEM) as described in [Lang, Bliese and de Voogt \(2018\)](#), *Personnel Psychology*. The CEM model is basically a growth model with the residual being an exponential function of the time variable. We want to do the same thing in both R and stata.

We start with simpler models.

36.1. Stata and R syntax comparison for the same three level model

We run a three level model specified in stata with `mixed` and `menl`. In R with `nlme` and `lme4`. This is a random intercept model. Only intercept is varying for each unit at each level. In this data set supposedly state is nested within region. There are situations that levels are not nested, then we need to specify cross effect models. We do not discuss that here. One point to watch: `nlme`'s `lme` function expects the grouping factors in the `random` specification ordered from the *highest* level down. Here `state` is nested within `region`, so `region` must be listed before `state`; reversing the order leads `lme` to treat regions as nested within states, which is not the intended structure.

There are different ways to specify this model in `nlme`. I show two ways here (including `lme2` which is commented out).

36.1.1. Stata

```
* let's see how to use menl
webuse productivity
* to replicate mixed with menl
```

36. Doing the same multi-level model with R and stata

```
mixed gsp || region: || state:
menl gsp = {b1:}, define(b1: U1[region] UU1[region>state])
```

(Public capital productivity)

Performing EM optimization ...

Performing gradient-based optimization:

Iteration 0: Log likelihood = 200.86162

Iteration 1: Log likelihood = 200.86162

Computing standard errors ...

Mixed-effects ML regression

Number of obs = 816

Grouping information

Group variable		No. of groups	Observations per group		
			Minimum	Average	Maximum
region		9	51	90.7	136
state		48	17	17.0	17

Log likelihood = 200.86162

Wald chi2(0) = .
Prob > chi2 = .

gsp	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
_cons	10.65961	.2503806	42.57	0.000	10.16887	11.15035

Random-effects parameters		Estimate	Std. err.	[95% conf. interval]	
region: Identity					
	var(_cons)	.4376123	.2697616	.13073	1.464887
state: Identity					
	var(_cons)	.6080626	.1382733	.3893904	.9495357
	var(Residual)	.0246633	.0012586	.0223159	.0272577

LR test vs. linear model: chi2(2) = 2750.56

Prob > chi2 = 0.0000

36.1. Stata and R syntax comparison for the same three level model

Note: LR test is conservative and provided only for reference.

Obtaining starting values by EM:

Alternating PNLS/LME algorithm:

Iteration 1: Linearization log likelihood = 200.86162

Iteration 2: Linearization log likelihood = 200.86162

Computing standard errors:

Mixed-effects ML nonlinear regression Number of obs = 816

Grouping information

Path	No. of groups	Observations per group		
		Minimum	Average	Maximum
region	9	51	90.7	136
region>state	48	17	17.0	17

Linearization log likelihood = 200.86162

b1: U1[region] UU1[region>state]

	gsp	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
b1							
	_cons	10.65961	.2505341	42.55	0.000	10.16857	11.15065

Random-effects parameters		Estimate	Std. err.	[95% conf. interval]	
region: Identity					
	var(U1)	.4376119	.2686479	.1313834	1.457598
region>state: Identity					
	var(UU1)	.6080625	.1381527	.3895418	.9491666
	var(Residual)	.0246633	.0012586	.0223159	.0272577

36.1.2. R

These two R packages return same results. Note that lmer reports variance, while lme reports standard deviation.

```
library(tidyverse)
library(haven)
library(dplyr)
library(readxl)
library(nlme)
library(lme4)
data <- read_dta('https://www.stata-press.com/data/r18/productivity.dta')
# lme4 for three level model
# (1|region) + (1|state) specifies CROSSED random effects; it numerically
# coincides with the intended nested structure here only because each
# state's label happens to be unique across the whole dataset (never reused
# across regions) -- verified: identical logLik to the explicitly nested
# (1|region/state) below. The nested form is the conceptually correct one
# and is robust even if state labels were ever reused across regions.
lmer1 <- lmer(gsp ~ 1 + (1|region/state), data=data)
summary(lmer1)
```

Linear mixed model fit by REML ['lmerMod']

Formula: gsp ~ 1 + (1 | region/state)

Data: data

REML criterion at convergence: -400.9

Scaled residuals:

Min	1Q	Median	3Q	Max
-2.8328	-0.6757	0.0474	0.6382	3.3928

Random effects:

Groups	Name	Variance	Std.Dev.
state:region	(Intercept)	0.60775	0.7796
region	(Intercept)	0.50977	0.7140
Residual		0.02466	0.1570

Number of obs: 816, groups: state:region, 48; region, 9

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	10.665	0.266	40.1

```
lme1 <- lme(gsp ~ 1, random= list( region=~1, state=~1), data=data)
#or
#lme1 <- lme(gsp ~ 1, random= list(region = pdSymm(~1), state=pdSymm(~1)), data=data)
#but
```

```
#lme1 <- lme(gsp ~ 1, random= list( state=~1, region=~1), data=data)
#would be a different result!
#The other way to specify:
#lme2 <- lme(gsp ~ 1, random= ~ 1 | region/state, data=data)
summary(lme1)
```

Linear mixed-effects model fit by REML

```
Data: data
      AIC      BIC   logLik
-392.8509 -374.0381 200.4254
```

Random effects:

```
Formula: ~1 | region
(Intercept)
StdDev:   0.7139792
```

```
Formula: ~1 | state %in% region
(Intercept) Residual
StdDev:     0.7795814 0.1570457
```

Fixed effects: gsp ~ 1

```
      Value Std.Error DF t-value p-value
(Intercept) 10.66483 0.2659842 768 40.09574      0
```

Standardized Within-Group Residuals:

```
      Min      Q1      Med      Q3      Max
-2.83282326 -0.67572390 0.04744387 0.63817953 3.39281358
```

Number of Observations: 816

Number of Groups:

```
      region state %in% region
          9          48
```

36.2. A growth model in R and stata

We implement a growth model in R and stata. The results are somewhat different. The results may not be that reliable given that the year slope random component is tiny.

36.2.1. Stata

```
* same growth model with mixed and menl
* Stata's "mixed ... || region: year" defaults to an INDEPENDENT covariance
* structure for region's random intercept and slope (var(year), var(_cons),
```

36. Doing the same multi-level model with R and stata

```
* cov=0 by construction) -- verified via `estat recovariance`. R's lmer
* "(year|region)" and lme's "region=~year" both default to an UNSTRUCTURED
* (correlated) covariance. Add covariance(unstructured) to match R.
webuse productivity
mixed gsp year || region: year, covariance(unstructured) || state:
menl gsp = {b2:} + {U1[region]}*year, define(b2: year U0[region] U11[region>state])
```

(Public capital productivity)

Performing EM optimization ...

Performing gradient-based optimization:

```
Iteration 0: Log likelihood = 782.11459 (not concave)
Iteration 1: Log likelihood = 803.58071 (not concave)
Iteration 2: Log likelihood = 813.1017 (not concave)
Iteration 3: Log likelihood = 816.56892
Iteration 4: Log likelihood = 885.27225 (not concave)
Iteration 5: Log likelihood = 888.29208
Iteration 6: Log likelihood = 888.79472
Iteration 7: Log likelihood = 888.82395
Iteration 8: Log likelihood = 888.82404
Iteration 9: Log likelihood = 888.82404
```

Computing standard errors ...

Mixed-effects ML regression

Number of obs = 816

Grouping information

Group variable		No. of groups	Observations per group		
			Minimum	Average	Maximum
region		9	51	90.7	136
state		48	17	17.0	17

Log likelihood = 888.82404

Wald chi2(1) = 107.46

Prob > chi2 = 0.0000

gsp	Coefficient	Std. err.	z	P> z	[95% conf. interval]	
year	.0263275	.0025397	10.37	0.000	.0213497	.0313053
_cons	-41.43216	5.20712	-7.96	0.000	-51.63793	-31.22639

```

-----
Random-effects parameters | Estimate Std. err. [95% conf. interval]
-----+-----
region: Unstructured    |
      var(year) |      .000056   .0000274   .0000215   .000146
      var(_cons) |     235.7658   115.06     90.58819   613.6066
      cov(year,_cons) |    -.1148543   .056104    -.2248162   -.0048923
-----+-----
state: Identity         |
      var(_cons) |      .6121886   .139357     .39185     .9564242
-----+-----
      var(Residual) |      .0039862   .0002046     .0036046   .0044081
-----
LR test vs. linear model: chi2(4) = 4112.15          Prob > chi2 = 0.0000

```

Note: LR test is conservative and provided only for reference.

Obtaining starting values by EM:

Alternating PNLS/LME algorithm:

```

Iteration 1: Linearization log likelihood = 784.65134
Iteration 2: Linearization log likelihood = 784.65343
Iteration 3: Linearization log likelihood = 784.65343

```

Computing standard errors:

```

Mixed-effects ML nonlinear regression          Number of obs      =          816

```

Grouping information

```

-----
              |      No. of      Observations per group
              |      groups      Minimum    Average    Maximum
-----+-----
      region |           9         51       90.7       136
region>state |          48         17       17.0        17
-----

```

```

                                          Wald chi2(1)      =      2737.78
Linearization log likelihood = 784.65343    Prob > chi2      =      0.0000

```

b2: year U0[region] U11[region>state]

```

-----
gsp | Coefficient Std. err.      z    P>|z|    [95% conf. interval]
-----+-----
b2      |

```

36. Doing the same multi-level model with R and stata

year	.0274903	.0005254	52.32	0.000	.0264605	.02852
_cons	-43.71615	1.069039	-40.89	0.000	-45.81143	-41.62087

Random-effects parameters	Estimate	Std. err.	[95% conf. interval]			
-----+-----						
region: Independent						
var(U0)	.4378219	.2687857		.1314413	1.458355	
var(U1)	5.40e-13	2.67e-10		0	.	
-----+-----						
region>state: Identity						
var(U11)	.6091976	.1381402		.3906086	.9501115	
-----+-----						
var(Residual)	.0053926	.0002752		.0048793	.0059598	

36.2.2. R

```
# lme4 for three level model, growth model
lmer_growth1 <- lmer(gsp ~ 1 + year + (year|region) + (1|state), data=data, REML=FALSE)
summary(lmer_growth1)
```

```
Linear mixed model fit by maximum likelihood ['lmerMod']
Formula: gsp ~ 1 + year + (year | region) + (1 | state)
Data: data
```

AIC	BIC	logLik	-2*log(L)	df.resid
-1455.8	-1422.9	734.9	-1469.8	809

Scaled residuals:

Min	1Q	Median	3Q	Max
-2.9207	-0.5232	0.0180	0.5541	3.8399

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
state	(Intercept)	1.327e-01	3.643e-01	
region	(Intercept)	2.068e-01	4.547e-01	
	year	4.495e-09	6.705e-05	1.00
Residual		6.664e-03	8.163e-02	

Number of obs: 816, groups: state, 48; region, 9

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	-4.373e+01	1.165e+00	-37.53


```
year          2.751e-02  5.838e-04  47.13
```

Correlation of Fixed Effects:

```
(Intr)
year -0.985
optimizer (nloptwrap) convergence code: 0 (OK)
boundary (singular) fit: see help('isSingular')
```

```
# nlme for three level model, growth model
#lme_growth1<-lme(gsp ~ year, random = list(region=pdSymm(~year), state=pdIdent(~1)),data=data)
lme_growth1<-lme(gsp ~ year, random = list(region=~year, state=~1),data=data, method="ML")
summary(lme_growth1)
```

Linear mixed-effects model fit by maximum likelihood

```
Data: data
      AIC      BIC   logLik
-1555.307 -1522.376 784.6534
```

Random effects:

```
Formula: ~year | region
Structure: General positive-definite, Log-Cholesky parametrization
      StdDev      Corr
(Intercept) 6.61523e-01 (Intr)
year        7.96178e-09 0.001
```

```
Formula: ~1 | state %in% region
      (Intercept) Residual
StdDev:  0.7805102 0.0734343
```

```
Fixed effects:  gsp ~ year
              Value Std.Error DF   t-value p-value
(Intercept) -43.71617 1.0690287 767 -40.89336      0
year         0.02749 0.0005254 767  52.32366      0
Correlation:
      (Intr)
year -0.972
```

Standardized Within-Group Residuals:

```
      Min      Q1      Med      Q3      Max
-3.24606517 -0.59087173  0.03087456  0.62062441  4.27790886
```

Number of Observations: 816

```
Number of Groups:
      region state %in% region
      9          48
```



```

-----+-----
region>state: Identity      |
              var(U11) |   .6046714   .1369848   .3878686   .942658
-----+-----
Residual variance:        |
  Exponential year        |
              sigma2 |   9.96e-99   .           .           .
              gamma  |   .0556293   .000013   .0556039   .0556548
-----+-----

```

36.3.2. R

As in stata, it has some convergence problem. So the results can differ somewhat between R and stata.

```

lme_growth1<-lme(gsp ~ year, random = list(region=pdSymm(~year), state=pdIdent(~1)),data=data,
lme_growth2<-update(lme_growth1,weights=varExp( form = ~ year), method="ML", control = lmeCont
summary(lme_growth2)

```

Linear mixed-effects model fit by maximum likelihood

Data: data

	AIC	BIC	logLik
	-1548.244	-1510.608	782.1218

Random effects:

Formula: ~year | region

Structure: General positive-definite

	StdDev	Corr
(Intercept)	0.2937615870	(Intr)
year	0.0001436049	0.995

Formula: ~1 | state %in% region

	(Intercept)	Residual
StdDev:	0.8046618	0.0736431

Variance function:

Structure: Exponential of variance covariate

Formula: ~year

Parameter estimates:

	expon
	-2.165027e-16

Fixed effects: gsp ~ year

	Value	Std.Error	DF	t-value	p-value
(Intercept)	-43.80206	1.0533833	767	-41.58226	0
year	0.02752	0.0005291	767	52.01290	0

Correlation:

36. *Doing the same multi-level model with R and stata*

```
(Intr)
year -0.976
```

Standardized Within-Group Residuals:

Min	Q1	Med	Q3	Max
-3.24679735	-0.58919213	0.03104216	0.61843888	4.27977261

Number of Observations: 816

Number of Groups:

region	state	%in%	region
9			48

```
delta1<-lme_growth2$modelStruct$varStruct
delta1
```

Variance function structure of class varExp representing

```
expon
-2.165027e-16
```

37. Modeling variation, not just the mean, in Likert-scale data

37.1. Modeling the spread, not the mean

Usually we regress a Likert item on covariates and model only its mean. A 1–7 satisfaction item, a 1–5 agreement item, a peer rating — we ask what moves the average.

But many questions are about the spread. Do evaluators agree more about some employees than others? Does one group of raters disagree more than another? Does a training program tighten the distribution of responses, or only shift its center?

The natural move is to compute a spread for each unit — each subject’s ratings across raters, each respondent’s answers across items — and regress *that* on covariates, the way we regress a mean. On a bounded scale, that move breaks. This chapter is about why it breaks, and what to do instead.

37.2. The ceiling is exact, not empirical

On a scale with a hard floor and ceiling, the largest variance a set of responses can have depends on where their mean sits. This is not a tendency in typical data. It is a mathematical fact, true for any distribution on a bounded interval.

Take any X confined to $[a, b]$. Both $(X - a)$ and $(b - X)$ are non-negative for every draw, so their product is too:

$$(X - a)(b - X) \geq 0 \quad \text{for every } X \in [a, b]. \quad (37.1)$$

Take expectations, expand, and substitute into $\text{Var}(X) = E[X^2] - \mu^2$ (with $\mu = E[X]$):

$$\text{Var}(X) \leq (a + b)\mu - ab - \mu^2 = (\mu - a)(b - \mu). \quad (37.2)$$

This is the **Bhatia–Davis inequality** (Bhatia and Davis 2000). In words: fix the mean, and the variance cannot exceed a value set entirely by that mean and the scale’s endpoints. Equality is reached, not just approached — put all the mass at the two endpoints, $P(X = a) = (b - \mu)/(b - a)$ and $P(X = b) = (\mu - a)/(b - a)$, and the variance hits $(\mu - a)(b - \mu)$ exactly. That is the most polarized a set of responses can be, given the mean.

On a 1–7 item that is

$$\text{SD}_{\max}(\mu) = \sqrt{(\mu - 1)(7 - \mu)}, \quad (37.3)$$

37. Modeling variation, not just the mean, in Likert-scale data

a downward parabola. It is widest at the midpoint ($\mu = 4$, where $SD_{\max} = 3$) and pinched to exactly zero at both ends. A respondent whose raters all pick 7 has zero disagreement, by definition. There is nowhere else for a rating to go.

```
mu_grid <- seq(1, 7, length.out = 200)
sd_max <- sqrt(pmax((mu_grid - 1) * (7 - mu_grid), 0))

ggplot(data.frame(mu = mu_grid, sd_max = sd_max), aes(mu, sd_max)) +
  geom_line(linewidth = 1, color = "#2d5fa3") +
  labs(x = "mean rating (.)", y = expression(SD[max](mu)),
       title = "Maximum possible SD, as a function of the mean") +
  theme_minimal(base_size = 12)
```




Figure 37.1.: The mechanical ceiling on SD, given the mean, for a 1–7 scale.

The consequence is direct. Two groups whose *means* differ have different *ceilings* on their SD, whether or not their true dispersion differs at all. And skew makes this the common case, not a corner case. Satisfaction items, performance ratings, and health scales usually sit near one end of the scale — exactly where the curve is steepest, and where raw SD comparisons mislead most.

i A finite-sample refinement

The curve is the *population* bound. For a finite sample of n real numbers with mean μ , the true maximum is tighter still. Ellis (2025) gives it exactly as $\text{Var}_{\max}(n, \mu) = (\mu - 1)(7 - \mu) - \frac{36}{n}a(1 - a)$, where a is the fractional part of $n(\mu - 1)/6$; this sits between $(\mu - 1)(7 - \mu) - 9/n$ and $(\mu - 1)(7 - \mu)$. At $n = 5$ raters that is up to 1.8 points² of extra slack — worth knowing when only a handful of people rate each unit, since the population curve then understates how tight the real ceiling

is.

37.3. Dividing by the mean does not fix it

The first instinct is to divide the SD by the mean, the way you would for a coefficient of variation. It does not work here. That correction is built for a *different* mechanism — variance scaling proportionally with the mean, as in count data — not a ceiling that pinches to zero at *both* ends. $\text{SD}_{\max}(\mu)/\mu$ is still a curve: zero at $\mu = 1$, rising to a peak of about 1.13 at $\mu = 7/4$, then falling back to zero at $\mu = 7$ (it happens to pass through 0.75 at $\mu = 4$). Note that it is not even shaped like the $\text{SD}_{\max}$ curve it was meant to flatten — dividing by μ pulls the maximum well below mid-scale, so the correction trades a symmetric mean-dependence for a skewed one. Dividing by the mean does not remove the mean-dependence. It relabels it.

37.4. The right idea: spread after removing the ceiling effect

Every fix worth having does the same thing — measure dispersion *after* dividing out the ceiling. Beta regression names this directly. Rescale the outcome to $(0, 1)$ (e.g. $(y - 1)/6$ for a 1–7 item) and parameterize by a mean μ and a precision φ (Ferrari and Cribari-Neto 2004; `betareg` in Cribari-Neto and Zeileis 2010). On the 1–7 scale,

$$\text{Var}(Y) = \frac{(\mu - 1)(7 - \mu)}{1 + \varphi} \quad \Rightarrow \quad \varphi = \frac{(\mu - 1)(7 - \mu)}{\text{Var}(Y)} - 1. \quad (37.4)$$

Set $\varphi = 0$ and this *is* the Bhatia–Davis bound. In words: φ is dispersion measured after the boundary is divided out, so $\varphi \sim$ group compares spread without inheriting the ceiling the way raw SD does.

φ itself is awkward to report — it lives on $(0, \infty)$ with no natural reference point. Its bounded twin is easier:

$$R := \frac{\text{Var}(Y)}{\text{Var}_{\max}(\mu)} = \frac{1}{1 + \varphi}. \quad (37.5)$$

By Bhatia–Davis, $\text{Var}(Y) \leq \text{Var}_{\max}(\mu)$ for *any* distribution on $[1, 7]$, not just Beta ones, so $R \in [0, 1]$ always, with no distributional assumption. $R = 1$ means a subject’s raters are as polarized as the ceiling allows, given their mean; $R = 0$ means perfect consensus. It reads like a percentage because it is one.

Two costs come with it. First, it is a continuous model grafted onto discrete, ordinal data — a 1–7 response is not a continuous proportion. Second, both forms put a rater-estimated variance in a denominator, and that variance comes from only a handful of raters. For φ the denominator is $\text{Var}(Y)$: when a few raters agree by chance — more likely near the ceiling, where most of the sample lives — $\text{Var}(Y) \rightarrow 0$ and $\varphi \rightarrow \infty$. For R the denominator is $\text{Var}_{\max}(\mu)$, which is zero whenever the mean sits at a scale endpoint; if all raters pick 7, then $\text{Var}(Y)$ and $\text{Var}_{\max}$ vanish together and $R = 0/0$ is undefined outright. We return to this in the worked example, because it bites hard.

37.5. Two models: one with the boundary, one without

37.5.1. Gaussian location-scale model (MELSM)

The location-scale model is the fix people reach for first. Model the mean as usual, and *also* model the residual SD as a function of covariates.

```
library(brms)

fit_gaussian <- brm(
  bf(rating ~ group + (1 | subject),
    sigma ~ group),          # model the SD, not just the mean
  data = df, family = gaussian()
)
```

brms puts a log link on **sigma**, so the **group** coefficient on the **sigma** part is a difference in *log*-SD. Exponentiate it for an SD ratio between groups.

The catch is fatal for our question. A Gaussian model has no boundary in it. It does not know the scale stops at 1 and 7. So it cannot tell “this group’s raters genuinely agree more” from “this group’s mean sits closer to the ceiling.” Both produce the identical signature in the fitted **sigma**.

37.5.2. Ordinal cumulative-probit MELSM: put the boundary inside the model

The clean fix treats the 1–7 response as what it is: a discretized slice of a latent continuous trait, cut by six ordered thresholds.

Posit an unobserved continuous Y_i^* behind each rating,

$$Y_i^* = \eta_i + \sigma_i \epsilon_i, \quad \epsilon_i \sim N(0, 1), \quad (37.6)$$

where η_i is the linear predictor for *location* (**group** plus a subject random intercept, say) and σ_i is the latent residual SD, which we also let depend on covariates. The observed rating is Y_i^* coarsened at six fixed thresholds $\theta_1 < \dots < \theta_6$ (with $\theta_0 = -\infty$, $\theta_7 = \infty$):

$$Y_i = k \iff \theta_{k-1} < Y_i^* \leq \theta_k. \quad (37.7)$$

So the cumulative probability is

$$P(Y_i \leq k) = P\left(\epsilon_i \leq \frac{\theta_k - \eta_i}{\sigma_i}\right) = \Phi\left(\frac{\theta_k - \eta_i}{\sigma_i}\right), \quad (37.8)$$

with Φ the standard normal CDF, for the probit link. Write $\text{disc}_i := 1/\sigma_i$ and this is **brms**’s parameterization exactly:

$$P(Y_i \leq k) = \Phi(\text{disc}_i (\theta_k - \eta_i)). \quad (37.9)$$

In words: discrimination is the reciprocal of the latent residual SD. Not an analogy — an algebraic identity, reached by rescaling the argument of Φ rather than dividing by σ_i directly. Letting `disc` depend on covariates is the ordinal analogue of `sigma ~ group`, written in reciprocal form: a positivity-friendly way to let dispersion vary (`disc > 0` everywhere, with no constraint on σ_i itself). `brms` puts a log link on `disc`, so `disc ~ 0 + group` means $\log(\text{disc}_i) = \gamma_{\text{group}[i]}$ — a group difference in *log*-discrimination, parallel to the Gaussian log-SD coefficient (Bürkner and Vuorre 2019).

```
fit_ordinal <- brm(
  bf(rating ~ group + (1 | subject),
     disc ~ 0 + group), # log(disc) = 0 + group; disc = 1 / latent SD
  data = df, family = cumulative("probit"), init = 0,
  prior = c(
    prior(constant(0), class = "b", coef = "groupA", dpar = "disc"),
    prior(normal(0, 1), class = "b", coef = "groupB", dpar = "disc")
  )
)
```

The `constant(0)` prior pins the reference group's *log*-discrimination — the linear predictor γ_A , not `disc` itself — at zero, so $\text{disc}_A = \exp(0) = 1$ and its latent SD is $1/\text{disc}_A = 1$. (Fixing $\gamma_A = 0$ is what makes `disc = 1` and `SD = 1` coincide for the reference group; it is not a claim that discrimination is zero, which would mean an infinite latent SD.) Only γ_B is estimated: $\text{disc}_B = \exp(\gamma_B)$, so $\gamma_B = 0$ means identical dispersion, $\gamma_B > 0$ means Group B is *more* discriminating (a *smaller* latent SD), and $\gamma_B < 0$ the reverse. The boundary is now explicit — six estimated cutpoints on an *unbounded* latent scale — so a group whose mean sits near 7 no longer shows a tighter latent SD by mechanics alone.

37.5.3. What belongs in the dispersion submodel?

A fair question: why does the dispersion side carry only `group`, when the mean model might carry more? Should the two not mirror each other? Not by default. They answer different questions. The mean model asks what predicts the average rating. The dispersion model asks what predicts the spread. A covariate earns its place on each side on its own merits.

In this chapter's simulation `group` is the only variable in the data-generating process, so there is nothing else *to* include, on either side. In real data the choice is substantive, and getting it wrong has a familiar cost. Suppose some covariate Z (rater experience, item wording, a subject-level trait) genuinely shifts dispersion and is correlated with `group`, but is left out of the dispersion submodel. Its effect is absorbed into the `group` coefficient on `sigma` (or `disc`, or φ) — omitted-variable bias for the *variance*, exactly analogous to the ordinary one for the mean. The remedy is the same one you already use on the mean side: include the covariates you have a substantive reason to think shift dispersion, and give that choice the same scrutiny — neither copying the mean formula reflexively nor stripping the dispersion submodel to the one comparison you happen to care about.

37.5.4. Four candidates, side by side

Two questions separate the approaches: does it remove the mechanical ceiling, and does it stay stable when only a handful of raters rate each unit? Only one passes both.

Approach	Removes the mechanical ceiling?	Stable at $n \approx 5$ raters?
SD / mean	No — still a curve in μ , just a different one	No — inherits raw SD's sampling noise
Hand-built $R = \text{Var}/\text{Var}_{\max}$	Yes, by construction	No — the denominator collapses near the ceiling, and R is undefined at a scale endpoint
Gaussian MELSM's <code>sigma</code>	No — no boundary in the model, so it inherits the confound	Yes — the joint likelihood pools across units
Ordinal MELSM's <code>disc</code>	Yes — the boundary lives in six estimated thresholds	Yes — the joint likelihood pools across units

The useful way to remember it: R is the right idea computed the wrong way, and the Gaussian MELSM is the wrong idea computed the right way — stable and precise, but blind to the boundary, so it reports the ceiling artifact as a confident finding. `disc` is the same idea as R — spread with the ceiling effect removed — estimated coherently inside one model. The worked example shows all of this.

37.6. Worked example: same dispersion, different means

Simulate two groups of 100 subjects each, 6 raters per subject. Both groups' raters draw from a **latent** continuous opinion with the *same* true SD ($\sigma = 1$) — there is no real disagreement difference to find. Group A's latent mean sits high (near the ceiling once discretized); Group B's sits in the middle of the range.

```
set.seed(42)

n_per_group <- 100
n_raters    <- 6
sigma_true  <- 1.0 # SAME latent SD in both groups -- no real difference
mu_A <- 2.4      # group A's latent mean sits near the ceiling
mu_B <- 0.6      # group B's sits mid-range

subjects <- data.frame(
  subject = 1:(2 * n_per_group),
  group   = rep(c("A", "B"), each = n_per_group),
  mu      = rep(c(mu_A, mu_B), each = n_per_group)
)

df <- subjects[rep(seq_len(nrow(subjects)), each = n_raters), ]
df$latent <- rnorm(nrow(df), mean = df$mu, sd = sigma_true)

# six fixed cutpoints -> seven observed categories, 1 through 7
```

```
cuts <- c(-2, -1.2, -0.5, 0.2, 1.0, 1.8)
df$rating <- as.integer(cut(df$latent, breaks = c(-Inf, cuts, Inf), labels = 1:7))
df$rating_ord <- factor(df$rating, ordered = TRUE, levels = 1:7)
```

The group difference enters in exactly one place, and it is easy to miss. `rnorm(..., sd = sigma_true)` uses the same `sigma_true` for every row, in both groups — `group` never touches the standard deviation, only the mean, via `df$mu`. The two groups differ *only* in where their latent mean sits relative to the six fixed cutpoints, which are the same `cuts` object for both. That alone produces very different observed spreads:

```
edges <- c(-Inf, cuts, Inf)
tibble(
  category = 1:7,
  prob_A = round(diff(pnorm(edges, mean = mu_A, sd = sigma_true)), 3),
  prob_B = round(diff(pnorm(edges, mean = mu_B, sd = sigma_true)), 3)
) |>
  knitr::kable(caption = "Category probabilities implied by each group's latent mean, with the
```

Table 37.2.: Category probabilities implied by each group’s latent mean, with the same cutpoints and the same sigma.

category	prob_A	prob_B
1	0.000	0.005
2	0.000	0.031
3	0.002	0.100
4	0.012	0.209
5	0.067	0.311
6	0.193	0.230
7	0.726	0.115

Group A’s latent mean (2.4) sits 0.6 SDs above the highest cutpoint (1.8), so `pnorm(0.6)  $\approx$  0.73`: nearly three-quarters of its ratings collapse into the top category, and under 2% fall below category 5. Group B’s mean (0.6) sits mid-range, so its ratings spread across all seven categories, none above about a third.

That imbalance — not a difference in σ — shrinks Group A’s raw SD. When three-quarters of a subject’s six ratings are forced into one bin by `cut()`, the *discrete* variance across them is mechanically small, even though the *continuous* latent draws behind them are as dispersed as Group B’s ($\sigma = 1$ for both). It is the same Bhatia–Davis mechanism from the introduction, now visible subject by subject through `cut()`’s fixed boundaries rather than through the population $SD_{\max}(\mu)$ curve.

The raw per-subject summary already shows the pattern that tempts a reader into seeing a real disagreement gap:

37. Modeling variation, not just the mean, in Likert-scale data

```
subj_summary <- df |>
  group_by(subject, group) |>
  summarise(mean_rating = mean(rating), sd_rating = sd(rating), .groups = "drop")

knitr::kable(
  subj_summary |> group_by(group) |>
    summarise(mean_of_means = round(mean(mean_rating), 2),
              mean_of_sds = round(mean(sd_rating), 2)),
  caption = "Raw per-subject summary by group -- true latent SD is IDENTICAL in both."
)
```

Table 37.3.: Raw per-subject summary by group – true latent SD is IDENTICAL in both.

group	mean_of_means	mean_of_sds
A	6.64	0.58
B	4.89	1.22

```
ggplot(subj_summary, aes(mean_rating, sd_rating, color = group)) +
  geom_jitter(width = 0.05, height = 0, alpha = 0.5) +
  geom_line(data = data.frame(mean_rating = mu_grid, sd_rating = sd_max),
            aes(mean_rating, sd_rating), color = "grey40", inherit.aes = FALSE) +
  labs(x = "subject mean rating", y = "subject SD",
       title = "Group A's raw SD looks smaller -- purely because its mean sits\nnear the grey ceiling",
       theme_minimal(base_size = 12))
```




Figure 37.2.: Subject-level mean vs. SD, with the theoretical ceiling overlaid.

Now fit both a Gaussian MELSM and the ordinal cumulative-probit MELSM to the same data.

```
fit_gaussian <- brm(
  bf(rating ~ group + (1 | subject), sigma ~ group),
  data = df, family = gaussian(),
  chains = 2, iter = 2000, warmup = 1000, cores = 2, seed = 1,
  refresh = 0, silent = 2
)

fit_ordinal <- brm(
  bf(rating_ord ~ group + (1 | subject), disc ~ 0 + group),
  data = df, family = cumulative("probit"), init = 0,
  prior = c(
    prior(constant(0), class = "b", coef = "groupA", dpar = "disc"),
    prior(normal(0, 1), class = "b", coef = "groupB", dpar = "disc")
  ),
  chains = 2, iter = 2000, warmup = 1000, cores = 2, seed = 1,
  refresh = 0, silent = 2
)
```

```
knitr::kable(round(fixef(fit_gaussian), 2), caption = "Gaussian MELSM fixed effects (sigma part is on the log-SD scale)")
```

Table 37.4.: Gaussian MELSM fixed effects (sigma part is on the log-SD scale)

	Estimate	Est.Error	Q2.5	Q97.5
Intercept	6.64	0.03	6.58	6.69
sigma_Intercept	-0.41	0.03	-0.46	-0.35
groupB	-1.75	0.06	-1.86	-1.63
sigma_groupB	0.68	0.04	0.60	0.76

```
knitr::kable(round(fixef(fit_ordinal), 2), caption = "Ordinal MELSM fixed effects (disc part is on the log scale: 0 = same latent SD as the reference group)")
```

Table 37.5.: Ordinal MELSM fixed effects (disc part is on the log scale: 0 = same latent SD as the reference group)

	Estimate	Est.Error	Q2.5	Q97.5
Intercept[1]	-4.39	0.32	-5.06	-3.81
Intercept[2]	-3.65	0.23	-4.12	-3.22
Intercept[3]	-2.98	0.17	-3.34	-2.66
Intercept[4]	-2.27	0.12	-2.52	-2.04
Intercept[5]	-1.44	0.07	-1.59	-1.30
Intercept[6]	-0.63	0.06	-0.74	-0.52
groupB	-1.89	0.10	-2.11	-1.69
disc_groupA	0.00	0.00	0.00	0.00

	Estimate	Est.Error	Q2.5	Q97.5
disc_groupB	-0.01	0.08	-0.17	0.15

On the raw scale Group A averages 6.64 with an SD of 0.58; Group B averages 4.89 with an SD of 1.22. The pooled correlation between subject mean and SD is -0.74 . That looks like a real, sizable disagreement gap, and the Gaussian MELSM agrees: `sigma_groupB` of 0.68 on the log-SD scale puts Group B's SD at **1.97 times** Group A's (95% CI: 1.82–2.14), nowhere near zero.

But there is no true difference. Both groups' raters were drawn with the same $\sigma = 1$ by construction; only the means differ. The ordinal model reads it correctly: `disc_groupB` is $\gamma_B = -0.015$ on the log scale (95% CI: -0.171 to 0.146), so $\text{disc}_B = e^{-0.015} \approx 0.985$ — a latent SD of ≈ 1.015 against the reference group's 1, recovering the fact that the two groups' dispersion is identical. The Gaussian model's "Group B disagrees twice as much" is not a finding about the raters. It is the ceiling effect from the previous section, wearing a regression coefficient.

In fairness, this is the ordinal model's best case: the simulation *is* its functional form, with one set of cutpoints shared by both groups and a constant latent σ . That is the right way to show the mechanism in isolation, but it does not test the assumption the model actually needs, which is that the thresholds are common across groups. If the two groups use the scale differently — differential item functioning, in the measurement literature — then threshold shifts and `disc` trade off against each other, and a dispersion difference can be read as a location or cutpoint difference instead. On real data that assumption deserves the same scrutiny the Gaussian model's missing boundary gets here.

37.6.1. What about R here?

R is mean-independent by construction, so it should land where the ordinal model does: no group difference. Compute it per subject, using the exact finite-sample bound from the earlier callout ($n = 6$ raters):

```
var_max_finite <- function(n, mu) {
  c <- (mu - 1) / 6
  a <- (n * c) %% 1
  36 * (c * (1 - c) - (1 / n) * a * (1 - a))
}

subj_R <- df |>
  group_by(subject, group) |>
  summarise(mean_rating = mean(rating),
            var_pop = mean((rating - mean(rating))^2), # population (1/n) convention
            .groups = "drop") |>
  mutate(var_max = var_max_finite(n_raters, mean_rating),
         R = ifelse(var_max > 1e-8, var_pop / var_max, NA_real_))

knitr::kable(
  subj_R |> group_by(group) |>
```

```
summarise(n_undefined = sum(is.na(R)), mean_R = round(mean(R, na.rm = TRUE), 3)),
caption = "Per-subject R, using each subject's own 6 ratings for both mean and variance."
)
```

Table 37.6.: Per-subject R , using each subject's own 6 ratings for both mean and variance.

group	n_undefined	mean_R
A	14	0.626
B	0	0.200

This is worse than the raw SD comparison, not better. Group A's per-subject R averages 0.63, Group B's 0.20 — a larger apparent gap, and pointing the *opposite* way from the ceiling artifact. Fourteen Group A subjects even have an *undefined* R : all six raters picked 7, so observed variance and mechanical ceiling are both zero at once. Has R failed?

No. The guarantee — $R \in [0, 1]$ at every μ , no forced link to the mean — is about the *population* quantity. Plug in a mean from only 6 raters and it breaks, because the finite-sample bound in the denominator is itself estimated from that same noisy mean, and it is far more sensitive near the ceiling than mid-scale:

```
knitr::kable(
  subj_R |> group_by(group) |>
    summarise(mean_var_pop = round(mean(var_pop), 3),
              mean_var_max = round(mean(var_max, na.rm = TRUE), 3)),
  caption = "The denominator (mean_var_max) is nearly 7x smaller for Group A -- not because Gr
)
```

Table 37.7.: The denominator (mean_var_max) is nearly 7x smaller for Group A — not because Group A's true ceiling is that much tighter, but because a 6-rater sample mean near 7 makes the *estimated* ceiling collapse.

group	mean_var_pop	mean_var_max
A	0.374	0.984
B	1.366	6.903

A Group A subject's six raters average near 7 because that is where their mean sits — but exactly *how* near is noisy, and the finite-sample bound shrinks fast as the sample mean approaches the boundary. A modest observed variance over an unluckily tiny denominator gives a large R , for no reason connected to real disagreement. Swapping in the plain continuous bound shrinks the gap a great deal — 0.181 vs. 0.173 — but each value still averages 100 noisy single-subject estimates. Only pooling each group's full 600 ratings into one mean and one variance removes the noise entirely:

37. Modeling variation, not just the mean, in Likert-scale data

```
pooled <- df |> group_by(group) |>
  summarise(mean_rating = mean(rating), var_pop = mean((rating - mean(rating))^2)) |>
  mutate(var_max_cont = (mean_rating - 1) * (7 - mean_rating),
         R_pooled = round(var_pop / var_max_cont, 3))

knitr::kable(pooled |> select(group, mean_rating, R_pooled),
             caption = "R computed once per group, from all 600 ratings pooled.")
```

Table 37.8.: R computed once per group, from all 600 ratings pooled.

group	mean_rating	R_pooled
A	6.638333	0.219
B	4.893333	0.211

Pooled, R is 0.22 for Group A and 0.21 for Group B — essentially identical, agreeing with the ordinal model’s null finding. The lesson is not “avoid R .” It is that R , like φ , is only as good as the sample behind the mean and variance you feed it. A hand-computed per-subject R at 6 raters can manufacture a *larger* spurious gap than raw SD, in the wrong direction. The ordinal model never faces this, because it never forms a per-subject summary at all. `disc_groupB` comes straight from the full, pooled 600-observation likelihood per group — the “pooled” calculation above done coherently, with a proper posterior interval instead of an average of point estimates.

37.7. Takeaways

i Summary

- On any bounded scale, the *maximum possible* variance is a function of the mean — the Bhatia–Davis inequality, $\text{Var}(X) \leq (\mu - a)(b - \mu)$. Not a tendency to check for, a fact to design around.
- Dividing SD (or variance) by the mean does not remove the confound. The mechanism is not proportional scaling; it is a ceiling that pinches at both ends.
- A Gaussian location-scale model can estimate `sigma ~ x`, but inherits the confound, because it has no boundary in it.
- Beta regression’s precision φ (and its bounded twin R) remove the confound algebraically, at the cost of a continuous approximation to discrete data. And R is only as good as the sample behind it: with few raters per subject, a hand-computed per-subject R can manufacture a *larger* spurious gap than raw SD, in the wrong direction — and is undefined outright when every rater picks the same endpoint of the scale, since then its numerator and denominator vanish together.
- An ordinal cumulative-link mixed model puts the boundary where it belongs — estimated thresholds on a latent scale — and is the cleanest way to ask “is this dispersion difference real?” on Likert data, because it never needs a per-unit summary statistic at all.

37.8. References

- Bhatia, R., and Davis, C. (2000). “A Better Bound on the Variance.” *American Mathematical Monthly* 107(4): 353–357.
- Bürkner, P.-C., and Vuorre, M. (2019). “Ordinal Regression Models in Psychology: A Tutorial.” *Advances in Methods and Practices in Psychological Science* 2(1): 77–101.
- Cribari-Neto, F., and Zeileis, A. (2010). “Beta Regression in R.” *Journal of Statistical Software* 34(2): 1–24.
- Ellis, J. L. (2025). “The maximum variance of a finite dataset, given its mean, minimum, and maximum.” [arXiv:2508.17525](https://arxiv.org/abs/2508.17525).
- Ferrari, S. L. P., and Cribari-Neto, F. (2004). “Beta Regression for Modelling Rates and Proportions.” *Journal of Applied Statistics* 31(7): 799–815.

38. Conjoint design and analysis using R

This post is to illustrate how to use R to do basic conjoint analysis.

38.1. Conjoint Analysis

Marketing people used to use Sawtooth to do conjoint analysis.

What is conjoint analysis? Here is the intro from Sawtooth:

“Conjoint analysis methods use statistical analysis to compute mathematical representations of survey respondents’ preferences for product features, simulate how attribute changes affect demand, and even predict market acceptance of new products before they launch.

This powerful research methodology has become very popular with market researchers and economists for the valuable insight that few other research methods can provide.”

“Consider the task of purchasing a new vehicle.

When searching for the “right” vehicle, you typically consider a multitude of variables: type, make (brand), model, year, mileage, price, color, accessories packages, etc. You can see how this is a complex choice to make.

How does anyone possibly make such a complicated decision?

You think about what features matter most to you and then make tradeoffs.

For example, one vehicle could be an acceptable make and model, but at a high price point, whereas a second option is not the ideal make and model but has a more palatable price.

The tradeoff in this simple example is whether the make and model is more preferable to a reasonable price, or vice versa.

A conjoint survey captures the relationship between different attributes of a product and how preferable they are both comparatively and when combined.

Because of this, conjoint analysis is the gold standard solution for testing multiple features simultaneously, when limited A/B testing doesn’t go far enough.”

“In simple terms, the conjoint survey design algorithm generates a balanced survey design with product profiles, also known as concepts, formulated to have ideal statistical properties (such as level balance and independence of attributes).

Each profile is made up of multiple conjoined product features (hence, conjoint analysis), such as brand, type, engine, and color, each with systematically varied levels.

These product concepts are then included in a series of questions, usually 8-20, called choice tasks, that make up the conjoint analysis portion of the survey.”

38.2. Example

I am following the example here: https://bookdown.org/lien_lamey/R_tutorial/8-1-data-5.html

```
library(tidyverse)
library(openxlsx)

url <- "https://feb.kuleuven.be/lien.lamey/public/rmodule/icecream.xlsx"
icecream <- read.xlsx(url)
icecream <- icecream %>%
  pivot_longer(cols = starts_with("Individual"), names_to = "respondent", values_to = "rating") %>%
  rename("profile" = "Observations") %>% # rename Observations to profile
  mutate(profile = factor(profile), respondent = factor(respondent), # factorize identifiers
         Flavor = factor(Flavor), Packaging = factor(Packaging), Light = factor(Light), Organic = factor(Organic))

icecream
```

```
# A tibble: 150 × 7
  profile Flavor Packaging Light Organic respondent rating
  <fct>   <fct>   <fct>   <fct>   <fct>   <fct>   <dbl>
1 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 1
2 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 6
3 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 5
4 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 1
5 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 2
6 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 7
7 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 7
8 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 5
9 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 1
10 Profile 1 Raspberry Homemade waffle No low fat Not organic Individual... 10
# 140 more rows
```

In this example, we have 4 attributes: Flavor, Packaging, Light, and Organic. There are 5, 3, 2, 2 levels correspondingly. Usually we are interested in how much does each level of attributes predict the outcome, say the rating of the icecream. To do this, we would think we'll need all $5 \times 3 \times 2 \times 2 = 60$ combinations of attribute levels to estimate the effect. In that case, we would need data for each combination. This is called full factorial design. The problem of full factorial design is that it might be too costly or not possible if the combination is huge. Instead, conjoint analysis allows us to estimate the effect of each level of each attribute by using a fraction of the combinations. This is done by creating a set of profiles, each of which contains a subset of the levels of the attributes. The profiles are then presented to the respondents, who rate the profiles. The ratings are then used to estimate the effect of each level of each attribute. This is done by regressing the ratings on the levels of the attributes. The coefficients of the regression are the estimates of the effects of the levels of the attributes.

The disadvantage of a fractional factorial design is that some effects are confounded; that is, the effect of one level of an attribute is confounded with the effect of another level of another attribute.

```
# attribute1, attribute2, etc. are vectors with one element in which we first provide the name
attribute1 <- "Flavor; Raspberry; Chocolate; Strawberry; Mango; Vanilla"
attribute2 <- "Package; Homemade waffle; Cone; Pint"
attribute3 <- "Light; Low fat; No low fat"
attribute4 <- "Organic; Organic; Not organic"

# now combine these different attributes into one vector with c()
attributes <- c(attribute1, attribute2, attribute3, attribute4)
attributes
```

```
[1] "Flavor; Raspberry; Chocolate; Strawberry; Mango; Vanilla"
[2] "Package; Homemade waffle; Cone; Pint"
[3] "Light; Low fat; No low fat"
[4] "Organic; Organic; Not organic"
```

38.3. Conjoint design of experiment

The first step in conjoint analysis is to create a design of experiment. This is to use a fraction of all possible combinations of attribute levels to estimate the effect of each level of each attribute. For example, instead of using all 60 icecream profiles, we use 10 profiles.

We use “radiant” package’s “doe” function to create the design of experiment. The “doe” function takes the attributes as input and returns a data frame with the profiles. The “seed” argument is used to fix the random number generator. This is useful if you want to reproduce the same results.

```
library(radiant)

ice_doe <- doe(attributes, seed = 123)
ice_doe$eff
```

	Trials	D-efficiency	Balanced
1	9	0.105	FALSE
2	10	0.389	FALSE
3	11	0.411	FALSE
4	12	0.614	FALSE
5	13	0.542	FALSE
6	14	0.479	FALSE
7	15	0.762	FALSE
8	16	0.738	FALSE
9	17	0.748	FALSE
10	18	0.756	FALSE
11	19	0.644	FALSE
12	20	0.895	FALSE
13	21	0.848	FALSE
14	22	0.833	FALSE

38. Conjoint design and analysis using R

15	23	0.790	FALSE
16	24	0.827	FALSE
17	25	0.787	FALSE
18	26	0.768	FALSE
19	27	0.759	FALSE
20	28	0.736	FALSE
21	29	0.702	FALSE
22	30	0.984	TRUE
23	31	0.952	FALSE
24	32	0.933	FALSE
25	33	0.928	FALSE
26	34	0.900	FALSE
27	35	0.871	FALSE
28	36	0.893	FALSE
29	37	0.866	FALSE
30	38	0.843	FALSE
31	39	0.836	FALSE
32	40	0.922	FALSE
33	41	0.899	FALSE
34	42	0.904	FALSE
35	43	0.882	FALSE
36	44	0.861	FALSE
37	45	0.949	FALSE
38	46	0.919	FALSE
39	47	0.912	FALSE
40	48	0.911	FALSE
41	49	0.891	FALSE
42	50	0.959	FALSE
43	51	0.939	FALSE
44	52	0.944	FALSE
45	53	0.925	FALSE
46	54	0.924	FALSE
47	55	0.906	FALSE
48	56	0.902	FALSE
49	57	0.884	FALSE
50	58	0.872	FALSE
51	59	0.855	FALSE
52	60	1.000	TRUE

```
#summary(doe(attributes, seed = 123, trials = 10))
```

This is a list of efficiency of different number of “trials”. Trials mean combinations of attribute levels. For example, with 10 trials, the efficiency is .389 and it’s unbalanced.

We can see that the only two designs with balanced design are 30 and 60. That’s because 30 and 60 are the only two numbers that can be divided by 5, 3 and 2, all unique number of 5, 3, 3, 2. If the factors are, say, 3, 2, 3. Then the integers that can be divided by 3 and 2 are 6, 12, and 18.

Efficiency is a measure of the information content a design can capture. Efficiency are typically stated in relative terms, as in “design A is 80% as efficient as design B.” In practical terms this means you will need 25% more (the reciprocal of 80%) design A observations (respondents, choice sets per respondent or a combination of both) to get the same standard errors and significance as with the more efficient design B. We used a measure of efficiency called D-Efficiency (Kuhfeld et al. 1994). In this case, the only orthogonal design is the full factorial design of 60 trials.

For example, a trial of 30 has a efficiency of .984, that means it’s 98.4% as efficient as the full factorial design of 60 trials.

```
library(radiant)
ice_doe2 <- doe(attributes, seed = 123, trials = 30)
ice_doe2$cor_mat
```

	Flavor	Package	Light	Organic
Flavor	1.000000e+00	-1.942890e-16	-1.665335e-16	-1.665335e-16
Package	-1.942890e-16	1.000000e+00	-1.665335e-16	-1.665335e-16
Light	-1.665335e-16	-1.665335e-16	1.000000e+00	-1.045299e-01
Organic	-1.665335e-16	-1.665335e-16	-1.045299e-01	1.000000e+00

From the covariance matrix, we can see that in this 30 profile design, there is a correlation of -.105 between “Light” and “Organic”. That means these two effects are confounded; the design is not orthogonal. That’s why the efficiency is not 1. The D-efficiency is calculated as a function of the determinant of the correlation matrix. The D-efficiency of the full factorial design is 1. The D-efficiency of the 30 profile design is .984. That means the 30 profile design is 98.4% as efficient as the full factorial design.

The partial factorial design of 30 profile is

```
ice_doe2$part
```

	trial	Flavor	Package	Light	Organic
1	1	Raspberry	Homemade_waffle	Low_fat	Organic
2	4	Raspberry	Homemade_waffle	No_low_fat	Not_organic
3	6	Raspberry	Cone	Low_fat	Not_organic
4	7	Raspberry	Cone	No_low_fat	Organic
5	10	Raspberry	Pint	Low_fat	Not_organic
6	11	Raspberry	Pint	No_low_fat	Organic
7	13	Chocolate	Homemade_waffle	Low_fat	Organic
8	14	Chocolate	Homemade_waffle	Low_fat	Not_organic
9	19	Chocolate	Cone	No_low_fat	Organic
10	20	Chocolate	Cone	No_low_fat	Not_organic
11	22	Chocolate	Pint	Low_fat	Not_organic
12	23	Chocolate	Pint	No_low_fat	Organic
13	26	Strawberry	Homemade_waffle	Low_fat	Not_organic
14	28	Strawberry	Homemade_waffle	No_low_fat	Not_organic
15	29	Strawberry	Cone	Low_fat	Organic

38. Conjoint design and analysis using R

16	32	Strawberry	Cone	No_low_fat	Not_organic
17	33	Strawberry	Pint	Low_fat	Organic
18	35	Strawberry	Pint	No_low_fat	Organic
19	39	Mango	Homemade_waffle	No_low_fat	Organic
20	40	Mango	Homemade_waffle	No_low_fat	Not_organic
21	41	Mango	Cone	Low_fat	Organic
22	42	Mango	Cone	Low_fat	Not_organic
23	46	Mango	Pint	Low_fat	Not_organic
24	47	Mango	Pint	No_low_fat	Organic
25	49	Vanilla	Homemade_waffle	Low_fat	Organic
26	51	Vanilla	Homemade_waffle	No_low_fat	Organic
27	53	Vanilla	Cone	Low_fat	Organic
28	56	Vanilla	Cone	No_low_fat	Not_organic
29	58	Vanilla	Pint	Low_fat	Not_organic
30	60	Vanilla	Pint	No_low_fat	Not_organic

38.4. Conjoint analysis

In the icecream data set we read in, it's a 10 profile design. There are 15 respondents. They rate 10 profiles of icecreams.

```
conjoint <- conjoint(icecream, rvar = "rating", evar = c("Flavor","Packaging","Light","Organic"),
summary(conjoint)
```

Conjoint analysis

Data : icecream

Response variable : rating

Explanatory variables: Flavor, Packaging, Light, Organic

Conjoint part-worths:

Attributes	Levels	PW
Flavor	Chocolate	0.000
Flavor	Mango	1.522
Flavor	Raspberry	0.522
Flavor	Strawberry	0.767
Flavor	Vanilla	1.389
Packaging	Cone	0.000
Packaging	Homemade waffle	-0.244
Packaging	Pint	-0.100
Light	Low fat	0.000
Light	No low fat	0.478
Organic	Not organic	0.000
Organic	Organic	0.307
Base utility	~	4.358


Conjoint importance weights:

Attributes	IW
Flavor	0.597
Packaging	0.096
Light	0.187
Organic	0.120

Conjoint regression results:

	coefficient
(Intercept)	4.358
Flavor Mango	1.522
Flavor Raspberry	0.522
Flavor Strawberry	0.767
Flavor Vanilla	1.389
Packaging Homemade waffle	-0.244
Packaging Pint	-0.100
Light No low fat	0.478
Organic Organic	0.307

```
plot(conjoint)
```



The partworths are just coefficient from a linear regression:

```
summary(lm(rating ~ Flavor + Packaging + Light + Organic, data = icecream))
```

Call:

```
lm(formula = rating ~ Flavor + Packaging + Light + Organic, data = icecream)
```

Residuals:

Min	1Q	Median	3Q	Max
-5.2867	-2.2244	-0.1911	2.3128	5.4089

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.3578	0.8172	5.332	3.76e-07 ***
FlavorMango	1.5222	0.9453	1.610	0.110
FlavorRaspberry	0.5222	0.9453	0.552	0.582
FlavorStrawberry	0.7667	0.8399	0.913	0.363
FlavorVanilla	1.3889	0.9453	1.469	0.144
PackagingHomemade waffle	-0.2444	0.8675	-0.282	0.779

PackagingPint	-0.1000	0.8399	-0.119	0.905
LightNo low fat	0.4778	0.5738	0.833	0.406
OrganicOrganic	0.3067	0.4751	0.645	0.520

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.91 on 141 degrees of freedom

Multiple R-squared: 0.03541, Adjusted R-squared: -0.01932

F-statistic: 0.6469 on 8 and 141 DF, p-value: 0.7371

The conjoint importance weights are the relative importance of each attribute. First the range of partworths for each attribute is calculated. Then the range of partworths for each attribute is divided by the sum of the ranges of all attributes. In this example, Flavor is the most important attribute, with an importance weight of .597.

There is another package called “conjoint” which I have not played with.

In summary, we can use R packages to set up a conjoint experiment, in the sense that we can have partial (fractional) factorial design. After collecting the data, we can use either the conjoint packages, or just run linear or logit models to analyze it. The frequentist approach used in conjoint analysis seems nothing special. Bayesian approach can also be used, but that’s another topic.

38.5. Causal effect in conjoint analysis

Hainmueller, Hopkins, and Yamamoto (2014) defined ACME (average marginal component effect). The basic idea is that once all the assumptions are met (usual assumptions in causal inference), we can have the effect of changing one attribute level to another with a causal interpretation. Fortunately it is usually just coefficients from a linear model, or linear combinations of the coefficients. There is a package “cregg” for each calculation and plotting of ACME, but we can just do it manually.

39. Stata 19 CATE command

39.1. Stata 19 new CATE command

Fernando Rios-Avila has a new blog about the new CATE command in Stata 19, which implements Conditional Average Treatment Effect (CATE) estimation.

https://friosavila.github.io/app_metrics/app_metrics12.html

⚠ If a Stata chunk here renders with no output

The Stata chunks below are cached (`cache: true`) because the `cate` command does lasso cross-fitting plus a random forest, and a full re-execution of this chapter takes over five minutes. The cache is safe for prose edits, which leave every chunk served uniformly from cache. It is **not** safe if you edit the *code* of one Stata chunk: the edited chunk re-executes while the earlier `collectcode: true` chunks are still served from cache, their code never enters Statamarkdown’s accumulated buffer, every command runs against missing data, and the chunk renders silently blank – and that empty result is then cached in turn.

If that happens, delete `stata-cate_cache/` **entirely** (not just the offending chunk’s entry) along with `_freeze/stata-cate/`, then re-render so all chunks run in order. Clearing `_freeze/` alone is not enough. See the note in this book’s `CLAUDE.md`.

```
clear all
webuse assets3
global catecovars age educ i.(incomecat pension married twoearn ira ownhome)
cate po (assets $catecovars) (e401k), group(incomecat) nolog
```

(Excerpt from Chernozhukov and Hansen (2004))

Conditional average treatment effects	Number of observations	=	9,913
Estimator: Partialing out	Number of folds in cross-fit	=	10
Outcome model: Linear lasso	Number of outcome controls	=	17
Treatment model: Logit lasso	Number of treatment controls	=	17
CATE model: Random forest	Number of CATE variables	=	17

		Robust			
	assets	Coefficient	std. err.	z	P> z
					[95% conf. interval]

39. Stata 19 CATE command

GATE							
incomecat							
0		3999.888	989.8742	4.04	0.000	2059.77	5940.006
1		1424.931	1668.081	0.85	0.393	-1844.447	4694.31
2		5092.801	1344.769	3.79	0.000	2457.102	7728.499
3		8700.512	2272.126	3.83	0.000	4247.226	13153.8
4		20350.97	4717.093	4.31	0.000	11105.64	29596.3

ATE							
e401k							
(Eligible							
vs							
Not elig..)		7914.24	1150.79	6.88	0.000	5658.734	10169.75

P0mean							
e401k							
Not eligi..		14012.54	834.0963	16.80	0.000	12377.74	15647.34

He said “We can achieve comparable results using standard regression with full interactions:”

```
clear all
webuse assets3
global catecovars age educ i.(incomecat pension married twoearn ira ownhome)

* you do not want to see all interactions.
qui:reg assets i.incomecat##i.e401k##c($catecovars), robust

* Average Treatment Effect (ATE)
margins, at(e401k=(0 1)) noestimcheck contrast(atcontrast(r))

* Group Average Treatment Effects (GATEs)
margins, at(e401k=(0 1)) noestimcheck contrast(atcontrast(r)) over(incomecat)
```

(Excerpt from Chernozhukov and Hansen (2004))

Contrasts of predictive margins
Model VCE: Robust

Number of obs = 9,913

Expression: Linear prediction, predict()
1._at: e401k = 0
2._at: e401k = 1

	df	F	P>F
_at	1	44.72	0.0000
Denominator	9833		

	Delta-method			
	Contrast	std. err.	[95% conf. interval]	
_at				
(2 vs 1)	7642.407	1142.797	5402.291	9882.524

Contrasts of predictive margins
Model VCE: Robust

Number of obs = 9,913

Expression: Linear prediction, predict()

Over: incomecat

1._at: 0.incomecat
e401k = 0

1.incomecat
e401k = 0

2.incomecat
e401k = 0

3.incomecat
e401k = 0

4.incomecat
e401k = 0

2._at: 0.incomecat
e401k = 1

1.incomecat
e401k = 1

2.incomecat
e401k = 1

3.incomecat
e401k = 1

4.incomecat
e401k = 1

	df	F	P>F
_at@incomecat			
(2 vs 1) 0	1	16.08	0.0001
(2 vs 1) 1	1	0.67	0.4123

39. Stata 19 CATE command

(2 vs 1) 2		1	16.03	0.0001
(2 vs 1) 3		1	13.77	0.0002
(2 vs 1) 4		1	17.26	0.0000
Joint		5	12.76	0.0000
Denominator		9833		

		Delta-method		
		Contrast	std. err.	[95% conf. interval]

_at@incomecat				
(2 vs 1) 0		3594.182	896.3298	1837.191 5351.172
(2 vs 1) 1		1283.3	1565.151	-1784.718 4351.317
(2 vs 1) 2		5056.144	1262.728	2580.938 7531.35
(2 vs 1) 3		8610.622	2320.156	4062.641 13158.6
(2 vs 1) 4		19665.61	4733.551	10386.88 28944.34

Numerically these two sets of estimations may be close, but I don't think they are from the same model, even the same idea. The CATE command with "PO" option is to estimate the CATE with "partialling out" approach.

What Fernando did is a "regression adjustment" type of model, which is modeling outcome model and allow interaction of treatment and covariates. The "partialling out" approach is sometimes called double ML.

Let's see how to replicate "teffects ra" with "reg":

```
clear all
webuse assets3
global catecovars age educ i.(incomecat pension married twoearn ira ownhome)

teffects ra (assets age educ i.(incomecat pension married twoearn ira ownhome))(e401k)

qui: reg assets i.e401k##c.(age educ) i.e401k##i.(incomecat pension married twoearn ira ownhome)
margins, at(e401k=(0 1)) noestimcheck contrast(atcontrast(r))
```

(Excerpt from Chernozhukov and Hansen (2004))

```
Iteration 0: EE criterion = 1.291e-21
Iteration 1: EE criterion = 1.046e-23
```

```
Treatment-effects estimation      Number of obs      =      9,913
Estimator      : regression adjustment
```


Outcome model : linear
 Treatment model: none

			Robust				
	assets	Coefficient	std. err.	z	P> z	[95% conf. interval]	
ATE							
	e401k						
	(Eligible						
	vs						
	Not elig..)	7929.242	1216.273	6.52	0.000	5545.39	10313.09
POmean							
	e401k						
	Not eligi..	14094.12	863.7641	16.32	0.000	12401.18	15787.07

Contrasts of predictive margins
 Model VCE: OLS

Number of obs = 9,913

Expression: Linear prediction, predict()

1._at: e401k = 0

2._at: e401k = 1

		df	F	P>F
	_at	1	34.39	0.0000
Denominator		9889		

			Delta-method		
		Contrast	std. err.	[95% conf. interval]	
	_at				
	(2 vs 1)	7929.242	1352.153	5278.746	10579.74

39.1.1. DoubleML

The R package DoubleML implements the double/debiased machine learning framework of Chernozhukov et al. (2018). It provides functionalities to estimate parameters in causal models based on

39. Stata 19 CATE command

machine learning methods. The double machine learning framework consist of three key ingredients: Neyman orthogonality, High-quality machine learning estimation and Sample splitting.

They consider a partially linear model:

$$y_i = \theta d_i + g_0(x_i) + \eta_i \quad (39.1)$$

$$d_i = m_0(x_i) + v_i \quad (39.2)$$

This model is quite general, except it does not allow interaction of d and x ; therefore no heterogeneous treatment effect across x . But “DoubleML” implements more than partially linear model, it actually allows for heterogeneous treatment effects, in models such as interactive regression model.

The basic idea of doubleML is to use any machine learning algorithm to estimate outcome equation ($l_0(X) = E(Y|X)$) and treatment equation ($m_0(X) = E(D|X)$). Then get the residuals, namely $\hat{W} = Y - \hat{l}_0(X)$ and $\hat{V} = D - \hat{m}_0(X)$.

Then regress $\hat{W}$ on $\hat{V}$. Based on FWL theorem, you get $\hat{\theta}$.

An important component here is to specify a Neyman-orthogonal score function. In the case of PLR, it can be the product of the two residuals:

$$\psi(W; \theta, \eta) = (Y - D\theta - g(X))(D - m(X)) \quad (39.3)$$

The estimators $\hat{\theta}$ is to solve the equation that the sample mean of this score function being 0.

And the variance of this score function is used as the variance of the doubleML estimator’s variance.

Now we try to replicate the CATE estimation in Stata using the R package DoubleML. I have not found a way to include factor variables in the lasso regression in R, so I included them as numeric variables.

The example in stata’s cate command is to use lasso on both outcome and treatment regression, then random forest on the individual treatment effect to get ATE. Here I try to replicate it in R.

```
library(haven)
library(DoubleML)
library(mlr3)
library(mlr3learners)
library(data.table)
library(dplyr)
# get rid of labels because doubleML does not like them
data <- read_dta("https://www.stata-press.com/data/r19/assets3.dta") |>
  zap_labels()

# pension, married, twoearn, ira, ownhome are already binary 0/1, so
# passing them as numeric is fine. incomecat, however, has 5 unordered
# categories (0-4) -- passing it as a single numeric column would force
```

```

# the ML learners (especially the linear regr.glmnet below) to treat
# category 2 as twice category 1, which is not a meaningful assumption
# for a nominal variable. One-hot encode it into dummy columns instead.
incomecat_dummies <- model.matrix(~ as.factor(incomecat) - 1, data = data)
colnames(incomecat_dummies) <- paste0("incomecat_", 0:(ncol(incomecat_dummies) - 1))
data <- cbind(data, incomecat_dummies)

assets3 <- data.table::setDT(data)
dml_data = DoubleMLData$new(data = assets3,
                             y_col = "assets",
                             d_cols = "e401k",
                             x_cols = c("age", "educ", colnames(incomecat_dummies), "pension", "r"))

print(dml_data)

```

```
===== DoubleMLData Object =====
```

```

----- Data summary -----
Outcome variable: assets
Treatment variable(s): e401k
Covariates: age, educ, incomecat_0, incomecat_1, incomecat_2, incomecat_3, incomecat_4, pension
Instrument(s):
Selection variable:
No. Observations: 9913

```

```

lgr::get_logger("mlr3")$set_threshold("warn")
learner = lrn("regr.glmnet", lambda=1)
ml_g_sim = learner$clone()
ml_m_sim = learner$clone()
set.seed(123)
obj_dml_plr = DoubleMLPLR$new(dml_data, ml_l=ml_g_sim, ml_m=ml_m_sim, n_folds=10)
obj_dml_plr$fit()
obj_dml_plr$summary()

```

Estimates and significance testing of the effect of target variables

	Estimate.	Std. Error	t value	Pr(> t)
e401k	7553	1281	5.897	3.69e-09 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

If we use random forest on the two models:

39. Stata 19 CATE command

```
# suppress messages from mlr3 package during fitting
lgr::get_logger("mlr3")$set_threshold("warn")
learner = lrn("regr.ranger", num.trees=2000, max.depth=5, min.node.size=2)
ml_l = learner$clone()
ml_m = learner$clone()
# learner = lrn("regr.glmnet")
# ml_g_sim = learner$clone()
# ml_m_sim = learner$clone()
set.seed(123)
obj_dml_plr = DoubleMLPLR$new(dml_data, ml_l=ml_l, ml_m=ml_m, n_folds=5)
obj_dml_plr$fit()
obj_dml_plr$summary()
```

Estimates and significance testing of the effect of target variables

```
Estimate. Std. Error t value Pr(>|t|)
e401k      9114      1375   6.627 3.43e-11 ***
---
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

We can do the same in Stata:

```
clear all
webuse assets3
global catecovars age educ incomecat pension married twoearn ira ownhome
cate po (assets age educ incomecat pension married twoearn ira ownhome) (e401k), omethod(rfor
```

(Excerpt from Chernozhukov and Hansen (2004))

Conditional average treatment effects	Number of observations	= 9,913
Estimator: Partialing out	Number of folds in cross-fit	= 5
Outcome model: Random forest	Number of outcome controls	= 8
Treatment model: Random forest	Number of treatment controls	= 8
CATE model: Random forest	Number of CATE variables	= 8

	assets	Coefficient	Robust std. err.	z	P> z	[95% conf. interval]	
ATE							
	e401k						
	(Eligible						
	vs						
	Not elig..)	8367.508	1180.69	7.09	0.000	6053.398	10681.62

P0mean							
e401k							
Not eligi..		13999.41	842.9698	16.61	0.000	12347.22	15651.6

The results are somewhat similar, but not exactly the same. There are some subtle differences in the way these two packages running random forest on the two models that I am not aware of.

39.2. AIPW

Stata's cate also has "aipw" option.

```
clear all
webuse assets3
global catecovars age educ incomecat pension married twoearn ira ownhome
cate aipw (assets $catecovars) (e401k), nolog
```

(Excerpt from Chernozhukov and Hansen (2004))

Conditional average treatment effects	Number of observations	= 9,913
Estimator: Augmented IPW	Number of folds in cross-fit	= 10
Outcome model: Linear lasso	Number of outcome controls	= 8
Treatment model: Logit lasso	Number of treatment controls	= 8
CATE model: Random forest	Number of CATE variables	= 8

		Robust				
	assets	Coefficient	std. err.	z	P> z	[95% conf. interval]

ATE						
e401k						
(Eligible						
vs						
Not eligi..)		7126.343	1182.229	6.03	0.000	4809.217 9443.469

P0mean						
e401k						
Not eligi..		14136.14	857.5555	16.48	0.000	12455.36 15816.92

I'll try to do this with "npcausal":

```

library(npcausal)
#SL.library <- c("SL.earth","SL.gam","SL.glmnet","SL.glm.interaction", "SL.mean","SL.ranger",
#SL.library <- c("SL.glmnet")
SL.lasso = function(...) {
  SL.glmnet(..., alpha=1)
}
SL.library <- c("SL.lasso")

Y <- data$assets
A <- data$e401k
# incomecat is nominal (5 categories); use the one-hot dummies created above
# rather than the raw numeric column, for the same reason as in the
# DoubleML example.
# as.data.frame() is required: `data` is a tibble, so `data[, cols]` returns a
# tibble, and npcausal::ate() fails on one with
#   Error in rep(): invalid 'times' argument
# It needs a plain data.frame.
# Two coercions are needed here, both for npcausal/SuperLearner's benefit:
#   * as.data.frame(): `data` is a tibble, and ate() fails on one with
#     Error in rep(): invalid 'times' argument
#   * as.numeric on every column: read_dta() returns haven_labelled columns, and
#     SuperLearner then errors with
#     The variable names and order in newX must be identical to ... X
W <- as.data.frame(data[, c("age", "educ", grep("^incomecat_", names(data), value = TRUE), "per
W <- as.data.frame(lapply(W, as.numeric))
# npcausal::ate() currently fails on this design with
#   Error in SuperLearner(): The variable names and order in newX must be
#   identical to the variable names and order in X
# raised from inside npcausal's own cross-fitting loop, not from anything in this
# chunk (the equivalent call works on a synthetic data.frame of the same shape).
# Wrapped so the chapter degrades gracefully rather than publishing a raw error;
# the DoubleML comparison above is the working cross-fitted estimate.
aipw <- tryCatch(
  ate(y = Y, a = A, x = W, nsplits = 10, sl.lib = SL.library),
  error = function(e) {
    message("npcausal::ate() failed: ", conditionMessage(e))
    NULL
  }
)

```

|
|

| 0%

```
if (!is.null(aipw)) aipw$res
```

Note here to match the Stata results, I included only “glmnet” in the SuperLearner library. The results are not exactly the same, but close enough.

Or we can use a similar package with more functionalities: “AIPW”.

```
library(AIPW)
library(SuperLearner)

#SuperLearner libraries for outcome (Q) and exposure models (g)
#sl.lib <- c("SL.glmnet")

AIPW_sl <- AIPW$new(Y= Y,
                   A= A,
                   W= W,
                   Q.SL.library = SL.library,
                   g.SL.library = SL.library,
                   k_split = 10,
                   verbose=FALSE)
```

```
Error in `value[[3L]]()` :
! Covariates dimension error: nrow(W) != length(A)
```

```
AIPW_sl$fit()
```

```
Error:
! object 'AIPW_sl' not found
```

```
AIPW_sl$summary()
```

```
Error:
! object 'AIPW_sl' not found
```

```
AIPW_sl$result
```

```
Error:
! object 'AIPW_sl' not found
```

```
# library(ggplot2)
# AIPW_sl$plot.p_score()
# AIPW_sl$plot.ip_weights()
```

Systematic treatment: [R](#) · [Julia](#).

40. Spatial econometrics introduction

40.1. Why do we need spatial econometrics?

‘Everything is related to everything else, but near things are more related than distant things’ (Tobler 1970). This is the first law of geography. This is the reason why we need spatial econometrics. The spatial econometrics is to account for the spatial dependence in the data. The spatial dependence is the correlation between the observations that are close to each other. This is a common feature in many datasets. For example, in the housing market, the price of a house is related to the prices of the houses nearby. In the crime data, the crime rate in one neighborhood is related to the crime rate in the neighboring neighborhoods. In the health data, the health outcomes of people in one area are related to the health outcomes of people in the neighboring areas. In the environmental data, the pollution level in one area is related to the pollution level in the neighboring areas. In the economic data, the economic performance of one region is related to the economic performance of the neighboring regions. In the social network data, the behavior of one person is related to the behavior of the people in the neighboring areas.

Econometricians love OLS. It is worth being precise about *what* spatial dependence breaks. Spatially correlated **errors** alone (the term $u = \rho W u + \epsilon$ below, with exogenous regressors) do not make the OLS coefficients inconsistent — they make the conventional standard errors wrong, so inference is invalid even though the point estimates are fine. OLS coefficients become **biased/inconsistent** in the genuinely simultaneous or omitted-variable cases: a spatially lagged *outcome* $W y$ on the right-hand side (endogenous by construction), omitted spatially correlated confounders, or endogenous spillovers. The sections below treat these cases separately: S2SLS for the spatial-lag-of-outcome model, and spatial-error models for the correlated-error case.

40.2. A general linear spatial model

$$y = \lambda W y + X\beta + W X\gamma + u \quad (40.1)$$

$$u = \rho W u + \epsilon \quad (40.2)$$

where Y is the dependent variable, X is the matrix of independent variables, β is the vector of coefficients, and u is the error term. The error term u is spatially correlated.

We can see the major difference is W . The term $W y$ is called “spatially lagged y ”, which in my opinion is misleading. This term is likely from time-series counter-part of lagged y , but there is nothing about lags in the spatial data. W is the spatial weights matrix. It is a matrix that describes the spatial structure of the data. The simplest example of W is the case that we define an element of W , $w_{i,j}$ to be 1 if unit i and j are neighbors, 0 if not.

In this simplest case, our model is saying we have a pattern we need to take account of: neighboring units are related, they seem to display similar pattern in terms of our study. This commonality of neighbors can be causal, for example, if we study the effect of minimum wage effect on employment, then if the neighboring state's wage increase might negatively impact the focal state's employment, because people go for higher wages. This is the spatial spillover effect. This can also be non-causal, for example, if we study the effect of temperature on crop yield, then the neighboring state's temperature might be a good predictor of the focal state's temperature. This is the spatial autocorrelation. In either case, we need to take account of the spatial structure in the data.

The reason behind the spatial pattern can be all kinds. It can be x_i causing y_j , or y_i causing y_j , or some unobserved factors causing both y_i and y_j , or some other reasons. Ignoring this spatial structure could make OLS biased. Betz et al (2020) show that ignoring spatial interdependence in the outcome results in asymptotically biased estimates even when instruments are randomly assigned.

40.3. Estimation with S2SLS for a less general model

Because we have the term Wy , we introduced endogeneity. Wy is necessarily correlated with u , even if u is not spatially correlated. Suppose u is not spatially correlated.

The spatial two stage least squares (S2SLS) estimator is one solution. Usually we use X and spatial lags of X as instruments. Typically instruments H is (X, WX, W^2X) .

S2SLS is similar to 2SLS. The first stage is

$$\hat{Z} = PZ = [X, WX, \hat{W}y] \quad (40.3)$$

where

$$\hat{W}y = PWy \quad (40.4)$$

and

$$P = H(H'H)^{-1}H'$$

{#eq-spatial-econometrics-5}.

Then

$$\hat{\theta}_{S2SLS} = (\hat{Z}'Z)^{-1}\hat{Z}'y \quad (40.5)$$

where $\hat{\theta}_{S2SLS}$ is the full coefficient vector (β, γ, λ) .

The first stage is to regress WY on H . The second stage is to regress Y on the predicted WY from the first stage and X and WX . The coefficient of X is the S2SLS estimator.

40.4. Estimation with a more general model

Suppose instead the error term is spatially correlated

$$u = \rho Wu + \epsilon \quad (40.6)$$

Then the estimation is more complicated. The generalized spatial two stage least squares (GS2SLS) estimator is needed. See Kelejian and Prucha (1998) for details.

40.5. an example

This example is from R package “sphet”.

The Boston data contain information on property values and characteristics in the area of Boston, Massachusetts and has been widely used for illustrating spatial models. Specifically, there is a total of 506 units of observation for each of which a variety of attributes are available, such as: (corrected) median values of owner-occupied homes (CMEDV); per-capita crime rate (CRIM); nitric oxides concentration (NOX); average number of rooms (RM); proportions of residential land zoned for lots over 25,000 sq. ft (ZN); proportions of non-retail business acres per town (INDUS); Charles River dummy variable (CHAS); proportions of units built prior to 1940 (AGE); weighted distances to five Boston employment centers (DIS); index of accessibility to highways (RAD); property-tax rate (TAX); pupil-teacher ratios (PTRATIO); proportion of blacks (B); and % of the lower status of the population (LSTAT).

The dataset with Pace’s tract coordinates is available to the R community as part of “spdep”.

The spatial weights matrix is a sphere of influence neighbors list also available from spdep once the Boston data are loaded. Here the weight matrix comes with the data. Usually it’s the most heavy lifting part of the study.

```
library(sphet)
library(spdep)
data("boston", package = "spData")
listw <- nb2listw(boston.soi)
```

The function read.gwt2dist reads a ‘GWT’ file (e.g., generated using the function distance). In this example we are using the matrix of tract point coordinates projected to UTM zone 19 boston.utm available from spdep to generate a ‘GWT’ file of the 10 nearest neighbors.

```
id1 <- seq(1, nrow(boston.utm))
tmp <- distance(boston.utm, region.id = id1, output = TRUE,
               type = "NN", nn = 10, shape.name = "shapefile",
               region.id.name = "id1", firstline = TRUE,
               file.name = "boston_nn_10.GWT")
coldist <- read.gwt2dist(file = "boston_nn_10.GWT", region.id = id1, skip = 1)
class(coldist)
```

```
[1] "sphet"      "distance" "nb"        "GWT"
```

```
summary(coldist)
```

```
Number of observations:
n: 506
```

```
Distance summary:
```

40. Spatial econometrics introduction

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.5441	0.9588	1.5843	2.0848	2.6389	11.6388

Neighbors summary:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
10	10	10	10	10	10

Spatial models in `sphet` are fitted by the functions `stslshac` and `gstslshet`. “`stslshac`” is Spatial two stages least square with HAC standard errors. “`gstslshet`” is GM estimation of a Cliff-Ord type model with Heteroskedastic Innovations, which is the general model we discussed before, with spatially correlated errors.

The theoretical model above is the Spatial Durbin Model (SDM), which includes WX as regressors (with their own coefficient γ), not just as instruments for Wy . If we only pass X (no WX) in the formula below, `stslshac`’s `W2X` argument still uses spatial lags of X internally, but only as *instruments* for the endogenous Wy – γ is never estimated, so the fitted model is actually the simpler Spatial Autoregressive (SAR) model ($\gamma = 0$), not SDM. To match the SDM specification, include WX explicitly as regressors:

```
boston_x_vars <- c("CRIM","ZN","INDUS","CHAS","NOX","RM","AGE","DIS","RAD","TAX","PTRATIO","B"
boston.c2 <- boston.c
for (v in boston_x_vars) {
  xv <- boston.c[[v]]
  if (is.factor(xv)) xv <- as.numeric(as.character(xv))
  boston.c2[[paste0("W_", v)]] <- lag.listw(listw, xv)
}

sdm_formula <- as.formula(paste(
  "log(CMEDV) ~ CRIM + ZN + INDUS + CHAS + I(NOX^2) + I(RM^2) + AGE + log(DIS) + log(RAD) + TA",
  paste0("W_", boston_x_vars, collapse = " + ")
))

res <- stslshac(sdm_formula,
  data = boston.c2, listw,
  distance = coldist, type = "Triangular", HAC = TRUE)
summary(res)
```

```
=====
=====
                        Spatial Lag Model
                        HAC standard errors
=====
=====
```

Call:

```
stslshac(formula = sdm_formula, data = boston.c2, listw = listw,
  HAC = TRUE, distance = coldist, type = "Triangular")
```

Residuals:

Min.	1st Qu.	Median	3rd Qu.	Max.
-0.71074	-0.06258	-0.00291	0.06279	0.66519

Coefficients:

	Estimate	SHAC	St.Er.	t-value	Pr(> t)	
Wy	0.71491435	0.09797746	7.2967	2.949e-13	***	
(Intercept)	1.74597474	0.46867690	3.7253	0.0001951	***	
CRIM	-0.00548837	0.00153211	-3.5822	0.0003407	***	
ZN	0.00081016	0.00045400	1.7845	0.0743454	.	
INDUS	-0.00070720	0.00285254	-0.2479	0.8041961		
CHAS1	-0.05854830	0.03984589	-1.4694	0.1417328		
I(NOX^2)	-0.03240666	0.21228222	-0.1527	0.8786677		
I(RM^2)	0.00846160	0.00289595	2.9219	0.0034793	**	
AGE	-0.00150525	0.00047265	-3.1847	0.0014489	**	
log(DIS)	-0.16117615	0.07522729	-2.1425	0.0321515	*	
log(RAD)	0.04219650	0.01648649	2.5595	0.0104835	*	
TAX	-0.00045089	0.00012647	-3.5652	0.0003636	***	
PTRATIO	-0.01336079	0.00581140	-2.2991	0.0215013	*	
B	0.00058124	0.00011965	4.8578	1.187e-06	***	
log(LSTAT)	-0.21713640	0.03831649	-5.6669	1.454e-08	***	
W_CRIM	-0.00451534	0.00300618	-1.5020	0.1330927		
W_ZN	-0.00133281	0.00058617	-2.2738	0.0229807	*	
W_INDUS	0.00079183	0.00351396	0.2253	0.8217152		
W_CHAS	0.09604488	0.06127241	1.5675	0.1169964		
W_NOX	-0.34866680	0.32382701	-1.0767	0.2816112		
W_RM	-0.06511363	0.05723288	-1.1377	0.2552474		
W_AGE	0.00175310	0.00062322	2.8130	0.0049086	**	
W_DIS	0.01948831	0.01383400	1.4087	0.1589164		
W_RAD	0.00203459	0.00345353	0.5891	0.5557722		
W_TAX	0.00028960	0.00019481	1.4866	0.1371284		
W_PTRATIO	0.00732717	0.00807107	0.9078	0.3639669		
W_B	-0.00054543	0.00015750	-3.4630	0.0005342	***	
W_LSTAT	0.00895071	0.00427009	2.0961	0.0360698	*	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The function `gstslshet` is used to estimate the GS2SLS model. Unlike `stslshac`, `gstslshet`'s GMM machinery hits an exactly-singular matrix on this dataset once the WX regressors are added (likely due to near-collinearity introduced by the additional lagged terms combined with its own internal instrument construction), so we keep the original (SAR, no WX) specification here — this estimates the spatial-error model with $\gamma = 0$, not the full Durbin extension shown above for `stslshac`.

```
res <- gstslshet(log(CMEDV) ~ CRIM + ZN + INDUS + CHAS + I(NOX^2) + I(RM^2)
                  + AGE + log(DIS) + log(RAD) + TAX + PTRATIO + B + log(LSTAT),
                  data = boston.c, listw = listw, initial.value = 0.2)
summary(res)
```

Call:

```
gstslshet(formula = log(CMEDV) ~ CRIM + ZN + INDUS + CHAS + I(NOX^2) +
  I(RM^2) + AGE + log(DIS) + log(RAD) + TAX + PTRATIO + B +
  log(LSTAT), data = boston.c, listw = listw, initial.value = 0.2)
```

Residuals:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-0.569385	-0.073165	-0.001675	0.000531	0.071500	0.740309

Coefficients:

	Estimate	Std. Error	t-value	Pr(> t)	
(Intercept)	2.51316605	0.26749367	9.3952	< 2.2e-16	***
CRIM	-0.00662744	0.00144522	-4.5858	4.523e-06	***
ZN	0.00038299	0.00036563	1.0475	0.2948733	
INDUS	0.00159352	0.00179772	0.8864	0.3753967	
CHAS1	-0.00447974	0.03689065	-0.1214	0.9033480	
I(NOX^2)	-0.27295899	0.11561412	-2.3609	0.0182283	*
I(RM^2)	0.00744059	0.00199637	3.7270	0.0001937	***
AGE	-0.00045400	0.00045572	-0.9962	0.3191333	
log(DIS)	-0.16517174	0.03484858	-4.7397	2.140e-06	***
log(RAD)	0.07453521	0.01752830	4.2523	2.116e-05	***
TAX	-0.00041956	0.00010763	-3.8981	9.697e-05	***
PTRATIO	-0.01412661	0.00410143	-3.4443	0.0005725	***
B	0.00035970	0.00011182	3.2168	0.0012964	**
log(LSTAT)	-0.24593826	0.03213364	-7.6536	1.954e-14	***
lambda	0.42407826	0.04463747	9.5005	< 2.2e-16	***
rho	0.29587455	0.08614291	3.4347	0.0005932	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

In stata, you would use “spivregress” which replaced user-written “spivreg”. Stata has a set of commands for spatial econometrics. Check out stata’s “sp” manual.

41. Causal simulation

Evans, R. J., & Didelez, V. (2024). “Parameterizing and simulating from causal models” is about simulating from a causal model. I found that interesting.

The usual causal inference is to have a observational data, then assume causal structure, then estimate the parameters. This paper is about the opposite: Suppose we have a causal model, what kind of distribution would generate it?

Suppose we are interested in joint distribution $p(x, a, y)$, x is covariates, a is treatment, y is outcome.

$$\begin{aligned} p(x, a, y) &= p(x, a)p(y|x, a) \\ &= p(x, a)p_a(y|x) \\ &= p(x, a)\frac{p_a(x, y)}{p_a(x)} \\ &= p(x, a)\frac{p_a(x)p_a(y)c(x, y|a)}{p_a(x)} \\ &= p(x, a)p_a(y)c(x, y|a) \end{aligned} \tag{41.1}$$

Basically the joint distribution of (x, a, y) can be factorized into the joint distribution of (x, a) , the *interventional marginal* $p_a(y)$ — that is, $p(y | do(a))$, marginal in y rather than conditional on x — and the copula $c(x, y|a)$. It is worth being careful on that middle term: the derivation starts from $p(y|x, a)$ but ends at $p_a(y)$, and those are different objects; the x -dependence that was in $p(y|x, a)$ has been moved into the copula. The copula is a function that captures the dependence between x and y given a . (x, a) is the “past”, which can be specified. $p_a(y)$ is the “marginal structure model” which can be specified. The copula model can also be specified.

About copula: copula is a function can link joint distribution to marginal distributions.

$$p(x, y) = f(x)g(y)c(F(x), G(y)) \tag{41.2}$$

where $F(x)$ and $G(y)$ are the marginal distributions of x and y , respectively, and c is the copula function that captures the dependence between x and y . The copula function is a multivariate distribution with uniform marginals. It can be used to generate joint distributions from marginal distributions.


Here x , a and y can all be vectors.

41.1. examples

What do we use this for? We can use this to simulate data from a causal model, and then fit different models to the simulated data. This allows to compare the performance of different models.

41.1.1. example 1

Suppose we have a causal graph like this:



We need to specify $P(y|do(a)) = \sum_x P(x)P(y|x, a)$, where $P(x)$ is the marginal distribution of x , and $P(y|x, a)$ is the conditional distribution of y given x and a . This is the “marginal structural model” (Robins, 2000). Also the “g-formula”.

We also need to specify $P(x, a)$, the “past”.

Finally we need to specify the copula $c(x, y|a)$, which captures the dependence between x and y given a . Depending on different situations, different copula can be used.

In the “causl” package example, we specify:

$$X \sim N(0, 1) \tag{41.3}$$

$$A|X = x \sim N(x/2, 1) \tag{41.4}$$

$$Y|do(A = a) \sim N((a - 1)/2, 1) \tag{41.5}$$

and Gaussian copula with correlation $\rho = 2\expit(1) - 1$.


```

library(causl)
# formulae corresponding to covariates, treatments, outcomes and the dependence
forms <- list(X ~ 1, A ~ X, Y ~ A, ~ 1)
# vector of model families (3=gamma/exponential, 1=normal/Gaussian)
fam <- c(1, 1, 1, 1)
# list of parameters, including 'beta' (regression params) and 'phi' dispersion
pars <- list(X = list(beta=0, phi=1),
             A = list(beta=c(0,0.5), phi=1),
             Y = list(beta=c(-0.5,0.5), phi=1),
             cop = list(beta=1))

## now create a `causl_model` object
cm <- causl_model(formulas=forms, family=fam, pars=pars,method="inversion")

# now simulate 1000 observations
set.seed(123456)
data <- rfrugal(n=1000, causl_model=cm)
head(data)

```

	X	A	Y
1	0.83373317	1.10709111	1.1416713
2	-0.27604777	-0.27218853	-1.0495996
3	-0.35500184	0.77575209	1.3179678
4	0.08748742	-0.05959823	-0.6106023
5	2.25225573	0.61909216	2.0789188
6	0.83446013	0.52019822	0.3229069

In this example, first we specify the structure of the model using a list of formulas. The first formula is for the covariates, the second for the treatments, the third for the outcomes, and the fourth for the dependence structure (copula). Note that outcome Y only depends on A in this interventional distribution.

Then we specify the families of the random variables. Here we use normal distribution for X , normal distribution for A and Y , and Gaussian copula for the dependence structure.

Finally we specify the parameters of the model, including regression coefficients and dispersion parameters. For example, $Y \sim A$ has two coefficients, intercept -0.5 and slope $.5$.

The `causl_model` object is then created using these components.

In this simple example, if we run a regression with A , X in the model, on the simulated data, we would get the correct effect back:

```

lm1 <- lm(Y ~ A*X, data=data)
summary(lm1)

```

Call:

41. Causal simulation

```
lm(formula = Y ~ A * X, data = data)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.6326	-0.6212	-0.0105	0.6051	3.3132

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.491922	0.030413	-16.175	<2e-16 ***
A	0.493471	0.027820	17.738	<2e-16 ***
X	0.468038	0.031878	14.682	<2e-16 ***
A:X	-0.009247	0.022360	-0.414	0.679

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.8953 on 996 degrees of freedom

Multiple R-squared: 0.4875, Adjusted R-squared: 0.4859

F-statistic: 315.8 on 3 and 996 DF, p-value: < 2.2e-16

We can also use MLE:

```
out <- fit_causl(data, formulas = list(X ~ 1, Y ~ A, ~ 1), family = c(1, 1, 1))
out
```

log-likelihood: -2716.696

X ~ 1

	est.	s.e.	sandwich
(intercept)	0.0108	0.0314	0.0314
residual s.e.:	0.983	0.044	0.0447

Y ~ A

	est.	s.e.	sandwich
(intercept)	-0.491	0.0319	0.0319
A	0.494	0.0278	0.0295
residual s.e.:	1.01	0.0472	0.0475

copula parameters:

cop ~ 1

	est.	s.e.	sandwich
(intercept)	0.998	0.069	0.0734

41.1.2. Plasmode simulation

In reality, we have a data set and we have a causal model. We want to simulate data from the causal model, but we want the simulated data to have the same distribution as the original data. This is called “plasmode simulation”. For example, we don’t want to use the real outcome variable.

In stead, we simulate X's from the original data, Then simulate Y and A from the causal model. We can then test which method works better in estimating the causal effect of A on Y.

```
# A tibble: 4,802 x 58
  x_1 x_2 x_3 x_4 x_5 x_6 x_7 x_8 x_9 x_10 x_11 x_12 x_13
  <int> <fct> <dbl> <dbl> <int> <int> <int> <int> <int> <int> <int> <int> <int>
1    29 C      1    7   60   85    0    0    1    0    0    1    0
2    27 C      0    0   64  178    0    0    0    0    0    2    1
3    27 C      0    0   60  102    0    0    0    0    0    1    0
4    37 C      0    0   65  174    0    0    0    0    0    1    0
5    24 C     20   14   63  129    0    0    0    0    0    1    0
6    27 C     40   15   63  135    0    0    0    0    0    1    1
7    26 C     20    8   69  140    0    0    0    1    0    0    0
8    33 C      0    0   60  110    0    0    0    1    0    0    1
9    28 C      3    5   61  160    0    0    0    0    0    3    1
10   31 C      0    0   63  114    0    1    0    0    0    0    0
# i 4,792 more rows
# i 45 more variables: x_14 <int>, x_15 <int>, x_16 <int>, x_17 <int>,
# x_18 <int>, x_19 <int>, x_20 <int>, x_21 <fct>, x_22 <int>, x_23 <int>,
# x_24 <fct>, x_25 <int>, x_26 <int>, x_27 <int>, x_28 <int>, x_29 <int>,
# x_30 <int>, x_31 <int>, x_32 <int>, x_33 <int>, x_34 <int>, x_35 <int>,
# x_36 <int>, x_37 <int>, x_38 <int>, x_39 <int>, x_40 <int>, x_41 <int>,
# x_42 <int>, x_43 <int>, x_44 <int>, x_45 <int>, x_46 <int>, x_47 <int>, ...
```

```
# Model for the causal effect of smoking on birthweight
forms <- list(list(),
              list(A ~ x_1 + x_3 + x_4),
              list(Y ~ A),
              list(~ 1))
# fams <- list(integer(0), 5, 1, 1)
fams <- list(integer(0), "binomial", "gaussian", 1)
pars <- list(A = list(beta=c(-1.5,0.03,0.02,0.05)),
             Y = list(beta=c(3200, -500), phi=400^2),
             cop = list(beta=-1))

cm2 <- causl_model(formulas=forms, family=fams, pars=pars, dat = dat)

set.seed(123456)
data2 <- rfrugal(causl_model=cm2)
head(data2)
```

```
  x_1 x_2 x_3 x_4 x_5 x_6 x_7 x_8 x_9 x_10 x_11 x_12 x_13 x_14 x_15 x_16 x_17
1  29  C   1   7  60  85   0   0   1   0   0   1   0   0   2   0   0
2  27  C   0   0  64 178   0   0   0   0   0   2   1   0   2   0   0
3  27  C   0   0  60 102   0   0   0   0   0   1   0   0   0   0   0
4  37  C   0   0  65 174   0   0   0   0   0   1   0   0   1   0   0
5  24  C  20  14  63 129   0   0   0   0   0   1   0   0   0   0   0
```

41. Causal simulation

```

6  27  C  40  15  63 135  0  0  0  0  0  1  1  2  2  0  0
   x_18 x_19 x_20 x_21 x_22 x_23 x_24 x_25 x_26 x_27 x_28 x_29 x_30 x_31 x_32
1   12  35  10   J   1  43   B  17  64 105  30  11   0   0   0
2   12  75  12   J   1  50   E  12  90 183  40   8   0   0   0
3   12  35  10   J   1  57   E  14  70 121  27  12   0   0   0
4   11  35  12   J   1  43   E  10  80 185  39  10   0   0   0
5    9  35  15   J   1  33   E  14  70 160  28   7   0   0   0
6    9  45  10   J   1  40   E  15  64 160  31   9   0   0   0
   x_33 x_34 x_35 x_36 x_37 x_38 x_39 x_40 x_41 x_42 x_43 x_44 x_45 x_46 x_47
1    0  79  28 340  48   1   9   9   3   0   3  54  16   0   0
2    0  69  30 430  73   0   9   9   3   0   5  51  16   0   0
3    0  84  29 450  76   0   3   6   2   0   4  41  13   0   0
4    0  82  30 440  78   1   8   9   2   0   8  70  16   0   0
5    1  85  35 360  71   1   5   8   0   2   3  65  18   0   0
6    1  81  28 515  56   1   7   9   2   0   1  63  20   0   0
   x_48 x_49 x_50 x_51 x_52 x_53 x_54 x_55 x_56 x_57 x_58 A      Y
1    0   0   0   0   0   0   0   0   0   0  45  39 1 2125.677
2    0   0   0   0   0   0   0   0   0   0  46  42 1 2689.114
3    0   1   0   0   0   0   0   0   0   0  45  40 0 3442.800
4    0   0   0   0   0   0   0   0   0   0  47  40 0 3212.584
5    0   2   0   0   0   0   0   0   0   0  47  43 0 2992.775
6    0   0   0   0   0   0   0   0   0   0  45  44 0 2738.893

```

```

# we can also use rfrugalParam(), which is older version of rfrugal().
#datAY <- rfrugalParam(formulas=forms, family=fams, pars=pars, dat=dat)

```

In this specification, we have a binary treatment A and a continuous outcome Y . The covariates are x_1 , x_3 and x_4 . As in the first example above, the copula's `beta` parameter maps to the actual Gaussian correlation via $\rho = 2 \text{expit}(\beta) - 1$; with `beta=-1` this gives $\rho = 2 \text{expit}(-1) - 1 \approx -0.46$, not -1 (a correlation of exactly -1 would require $\beta \rightarrow -\infty$ and would make the joint distribution singular). We can then simulate the data.

41.2. a more complicated model

Let's simulate a more complicated model:

```

library(ggplot2)
library(ggdag)

g <- dagify(
  A1 ~ L1,
  L2 ~ L1,
  L2 ~ A1,
  A2 ~ A1,
  A2 ~ L2,


```

```

Y ~ A2,
Y ~ L2,
Y ~ A1,
exposure = "A2",
outcome = "Y",
coords = list(x = c(L1 = 1, L2 = 2, Y = 3, A1 = 1.5, A2 = 2.5),
              y = c(L1 = 1, L2 = 1, Y = 1, A1 = 1.5, A2 = 1.5))
)

ggdag(g) +
  theme_dag()

```



Let's see how to write a causal model:

```

# formulae corresponding to covariates, treatments, outcomes and the dependence
forms <- list(list(L1 ~ 1, L2 ~ L1*A1), # covariates
              list(A1 ~ L1, A2 ~ L2*A1), # treatments
              Y ~ A1*A2*L2, # outcome
              ~ 1)

# vector of model families (3=gamma/exponential, 1=normal/Gaussian)

fams <- list(c(1, 1), c(5,5), 1, 1)
pars <- list(L1 = list(beta=0, phi=1),
             L2 = list(beta=c(0.3,0.5,-0.2,-0.1), phi=1),
             A1 = list(beta=c(-0.3,0.4), phi=1),
             A2 = list(beta=c(0.5,0.3,0.1,0), phi=1),
             Y = list(beta=c(0,1,2,0.5,1,0.2,0,0.8), phi=1),

```

41. Causal simulation

```
cop = list(beta = 0.5))

# we can use rfrugalParam() to simulate data from the model, or rfrugal().
#dat <- rfrugalParam(n=1e4, formulas=forms, family = fams, pars=pars)
#head(dat)

cm3 <- causl_model(formulas=forms, family=fams, pars=pars)

set.seed(123456)
data3 <- rfrugal(n=1e4, causl_model=cm3)
head(data3)
```

	L1	L2	A1	A2	Y
1	0.83373317	1.91691936	1	1	8.1680612
2	-0.27604777	0.17008294	1	1	5.2276976
3	-0.35500184	0.16907708	1	1	2.8194835
4	0.08748742	0.19688985	0	1	2.6789661
5	2.25225573	-0.29257431	0	0	0.3180188
6	0.83446013	0.02661196	0	1	1.5694426

If we have the correct specification for the outcome model, we'll get it right by linear model:

```
m1 <- glm(Y ~ A1*A2*L2, data = data3)
summary(m1)
```

Call:

```
glm(formula = Y ~ A1 * A2 * L2, data = data3)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.11524	0.02110	-5.461	4.85e-08 ***
A1	1.11589	0.03245	34.391	< 2e-16 ***
A2	2.02332	0.02682	75.454	< 2e-16 ***
L2	0.80424	0.01929	41.682	< 2e-16 ***
A1:A2	0.98345	0.04090	24.044	< 2e-16 ***
A1:L2	0.16028	0.02976	5.386	7.37e-08 ***
A2:L2	-0.03446	0.02401	-1.435	0.151
A1:A2:L2	0.83541	0.03693	22.620	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for gaussian family taken to be 0.9100255)

Null deviance: 53902 on 9999 degrees of freedom
Residual deviance: 9093 on 9992 degrees of freedom

AIC: 27446

Number of Fisher Scoring iterations: 2

```
library(marginaleffects)
avg_comparisons(
  model = m1,
  variables = "A2")
```

Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
2.51	0.0202	124	<0.001	Inf	2.47	2.55

Term: A2
 Type: response
 Comparison: 1 - 0

So the ATE is 2.5.

Let's use "lmt" package to estimate the treatment effect of A_2 on Y .

```
library(lmt)

d1 <- function(data, trt) {
  rep(1, nrow(data))
}

A <- "A2"
Y <- "Y"
W <- c("L1", "L2", "A1")

set.seed(34465)

treat <- lmt_tmle(
  data = data3,
  trt = "A2",
  outcome = "Y",
  baseline = W,
  outcome_type = "continuous",
  shift = d1,
  folds = 1,
  learners_trt = "SL.glm",
  learners_outcome = "SL.glm"
)

print(treat)
```

::: {.cell-output .cell-output-stdout}

41. Causal simulation

```
Estimate: 3.046
Std. error: 0.022
```

```
d2 <- function(data, trt) {
  rep(0, nrow(data))
}
control <- lmtp_tmle(
  data = data3,
  trt = "A2",
  outcome = "Y",
  baseline = W,
  outcome_type = "continuous",
  shift = d2,
  folds = 1,
  learners_trt = "SL.glm",
  learners_outcome = "SL.glm"
)

print(control)
```

```
::: {.cell-output .cell-output-stdout}
```

```
Estimate: 0.542
Std. error: 0.021
```

```
:::
```

```
lmtp_contrast(treat, ref = control)
```

```
::: {.cell-output .cell-output-stdout}
```

	shift	ref	estimate	std.error	conf.low	conf.high	p.value
1	3.05	0.542	2.5	0.0227	2.46	2.55	<0.001

```
::: :::
```


42. Frengression and Engression

In the [causal simulation chapter](#), I used Evans and Didelez (2024)’s decomposition of the joint distribution $p(x, a, y)$ into three pieces:

$$p(x, a, y) = p(x, a) \cdot p_a(y) \cdot c(x, y|a) \quad (42.1)$$

where $p(x, a)$ is the “past” distribution of covariates and treatment, $p_a(y)$ is the causal margin, and $c(x, y|a)$ is the copula linking covariates and outcome given treatment. The `causl` package implements this with parametric families. We specify Gaussian, binomial, and so on for each piece.

But real data can be hard to specify parametrically. We may not know the right copula family. The marginal distributions may have mass points, skewness, or other features that are annoying to write down.

Frengression (Yang, Evans and Shen, 2025) keeps the same decomposition but learns the pieces with neural networks.

There are two uses: causal simulation with a learned generative model, and interventional prediction using learned confounding structure. The second part uses engression, so I start there.

42.1. Engression

Engression (Shen and Meinshausen, 2024) stands for energy regression. Standard regression estimates $E[Y|X]$. Engression tries to learn the whole conditional distribution $P(Y|X)$ using neural networks trained with the energy score.

42.1.1. The Energy Score

The energy score is a proper scoring rule for distributions. Given observations Y and samples $\hat{Y}, \hat{Y}'$ from the model:

$$S(Y, \hat{Y}, \hat{Y}') = E\|Y - \hat{Y}\| - \frac{1}{2}E\|\hat{Y} - \hat{Y}'\| \quad (42.2)$$

The first term rewards samples close to the truth. The second term keeps the simulated distribution from collapsing to one point. This is why the method targets the full distribution rather than only the mean.

42.1.2. The Stochastic Network (StoNet)

The building block is a **StoNet**, a neural network that adds random noise $\epsilon \sim N(0, I)$ to the input at each layer:

$$\text{StoNet}(x) = f(x, \epsilon), \quad \epsilon \sim N(0, I) \quad (42.3)$$

Different noise draws produce different outputs. Each output is a sample from the learned conditional distribution. So we can:

- **Predict the mean:** average over many noise draws $\rightarrow E[Y|X = x]$
- **Predict quantiles:** take quantiles of the noise draws $\rightarrow$ conditional quantile regression
- **Sample:** each forward pass with fresh noise is a sample from $P(Y|X = x)$

42.1.3. Why Engression?

Compared with quantile regression or mixture density networks, engression is useful because:

1. it handles multivariate responses directly;
2. it does not require us to specify a Gaussian mixture or a list of quantiles;
3. it gives samples from the conditional distribution, which is convenient for simulation.

42.1.4. Quick Example

The **engression** R package, separate from **Rfrengression**, provides the standalone estimator. The code below shows the API. The **Rfrengression** examples later are the ones I actually run here.

```
library(engression)

# Nonlinear DGP with heteroskedastic noise
n <- 1000
X <- matrix(rnorm(n * 2), ncol = 2)
Y <- (X[,1])^2 + X[,2] + rnorm(n) * (0.5 + abs(X[,1]))

# Fit engression; it learns the full conditional distribution
engr <- engression(X, Y)

# Predict conditional mean (like lm, but nonlinear)
Yhat_mean <- predict(engr, X, type = "mean")

# Predict conditional quantiles (like quantile regression, but joint)
Yhat_quant <- predict(engr, X, type = "quantiles")

# Sample from the learned distribution (generative)
Ysample <- predict(engr, X, type = "sample", nsample = 1)
```

The point is that engression gives a distributional model of $Y|X$, not just a fitted mean.

42.2. From Engression to Frengression

Frengression is frugal simulation plus engression. It uses three engression models to represent the decomposition of $P(X, A, Y)$. Each sub-network is trained with the energy score.

42.3. The Architecture

Frengression has three sub-networks, matching the three pieces of the decomposition:

Decomposition	causl (parametric)	Frengression (neural net)
$P(X, A)$ — the past	Specified by family + params	<code>model_xz</code> : noise-only generator
$P_a(Y)$ — causal margin	Specified by family + params	<code>model_y</code> : takes (x, eta) , outputs y
$c(X, Y A)$ — copula	Gaussian/Frank/etc. copula	<code>model_eta</code> : takes (x, z) , outputs eta

The important part is `model_eta`. It takes (x, z) and outputs a latent variable `eta` that carries the dependence between confounders and outcome. During training, a Z-permutation trick forces `eta` to be marginally $N(0, I)$. This means:

- When `eta` comes from `model_eta(x, z)`, it carries the observational confounding information.
- When `eta` is sampled independently from $N(0, 1)$, it gives the interventional prediction.

Same network, different source of `eta`. That is the trick.

42.4. Using Rfrengression

The `Rfrengression` package is a native R implementation using `torch`.

```
library(Rfrengression)
```

42.4.1. Example: Binary Treatment DGP

Use a DGP with binary treatment, binary outcome, and four confounders:

```
set.seed(42)
n <- 2000

w1 <- rbinom(n, 1, 0.5)
w2 <- rbinom(n, 1, 0.5)
w3 <- round(runif(n, 0, 4), 3)
w4 <- round(runif(n, 0, 5), 3)
```

```

A <- rbinom(n, 1, plogis(-0.4 + 0.2*w2 + 0.15*w3 + 0.2*w4 + 0.15*w2*w4))
Y.1 <- rbinom(n, 1, plogis(-1 + 1 - 0.1*w1 + 0.3*w2 + 0.25*w3 + 0.2*w4 + 0.15*w2*w4))
Y.0 <- rbinom(n, 1, plogis(-1 + 0 - 0.1*w1 + 0.3*w2 + 0.25*w3 + 0.2*w4 + 0.15*w2*w4))
Y <- Y.1 * A + Y.0 * (1 - A)

true_ate <- mean(Y.1 - Y.0)
cat("True ATE:", round(true_ate, 4), "\n")

```

True ATE: 0.1935

42.4.2. Train Frengression

Training has two stages: first the outcome model (`model_y + model_eta`), then the marginal model (`model_xz`).

```

x_mat <- matrix(A, ncol = 1)
z_mat <- as.matrix(data.frame(w1, w2, w3, w4))
y_mat <- matrix(Y, ncol = 1)

model <- frengression(x_dim = 1, y_dim = 1, z_dim = 4,
                      x_binary = TRUE, y_binary = TRUE,
                      z_binary_dims = 2, noise_dim = 10)

model <- train_y(model, x_mat, z_mat, y_mat,
                 num_iters = 1000, lr = 1e-3, print_every = 500)

```

Epoch 1: loss 1.2123, loss_y 0.4705 (0.4861, 0.0312), loss_eta 0.7418 (0.8311, 0.1786)
 Stopping at iter 81

```

model <- train_xz(model, x_mat, z_mat,
                  num_iters = 1000, lr = 1e-4, print_every = 500)

```

Epoch 1: loss 3.3615, loss1 3.4682, loss2 0.2134
 Epoch 500: loss 1.3539, loss1 2.7097, loss2 2.7115
 Epoch 1000: loss 1.3635, loss1 2.7116, loss2 2.6963

42.4.3. Estimate ATE

For interventional prediction, `eta` is sampled from $N(0,1)$ instead of being generated from the observed confounding structure:

```

y1 <- predict(model, matrix(1, ncol = 1), type = "mean", nsample = 2000)
y0 <- predict(model, matrix(0, ncol = 1), type = "mean", nsample = 2000)

cat("Frengression ATE:", round(y1 - y0, 4), "\n")

```

Frengression ATE: 0.1911

```
cat("True ATE:", round(true_ate, 4), "\n")
```

True ATE: 0.1935

42.5. Frengression as DGP

Frengression can estimate an ATE, but that is not its main advantage. If I only want an ATE, I would usually use `npcausal`, `TMLE`, or `DoubleML`. The more interesting use is to train a data-generating process for benchmarking.

The idea is:

1. Train frengression on observational data, so it learns the confounding structure.
2. Replace the causal margin with a known function via `specify_causal()`.
3. Generate synthetic data with realistic confounding and known ground truth.

```
# Specify a known causal effect: Y = sigmoid(0.2*X + eta)
model_bench <- specify_causal(model, function(x, eta) {
  torch::torch_sigmoid(0.2 * x + eta)
})

# True ATE computed empirically from the model
y1_true <- predict(model_bench, matrix(1, ncol = 1), type = "mean", nsample = 5000)
y0_true <- predict(model_bench, matrix(0, ncol = 1), type = "mean", nsample = 5000)
cat("True ATE (specified):", round(as.numeric(y1_true - y0_true), 4), "\n")
```

True ATE (specified): 0.0913

Now we can generate synthetic data and benchmark any method:

```
syn <- sample_joint(model_bench, sample_size = 1000)
cat("Synthetic data: n =", nrow(syn$x), "\n")
```

Synthetic data: n = 1000

```
cat("Treatment prevalence:", mean(syn$x), "\n")
```

Treatment prevalence: 0.925

```
cat("Outcome prevalence:", mean(syn$y), "\n")
```

Outcome prevalence: 0.621

42.6. Real Data: LaLonde

This is most useful with real data. The LaLonde job-training data have many of the problems that make simulations hard to write by hand: zeros in prior earnings, uneven covariate distributions, and poor overlap.

```
library(MatchIt)
data("lalonge", package = "MatchIt")
lalonge$black <- as.integer(lalonge$race == "black")
lalonge$hispan <- as.integer(lalonge$race == "hispan")

x_lal <- matrix(lalonge$treat, ncol = 1)
y_lal <- matrix(lalonge$re78, ncol = 1)
# frengression's z_binary_dims assumes the FIRST z_binary_dims columns of z
# are the binary ones -- so the binary covariates (black, hispan, married,
# nodegree) must come first, followed by the continuous ones (age, educ,
# re74, re75). The original column order (age, educ first) put two
# continuous variables where binary ones were expected, and vice versa.
z_lal <- as.matrix(lalonge[, c("black", "hispan", "married", "nodegree",
                              "age", "educ", "re74", "re75")])

# Standardize for training
z_sc <- scale(z_lal)
y_sc <- (y_lal - mean(y_lal)) / sd(y_lal)

model_lal <- frengression(x_dim = 1, y_dim = 1, z_dim = 8,
                          x_binary = TRUE, y_binary = FALSE,
                          z_binary_dims = 4, noise_dim = 10)

model_lal <- train_y(model_lal, x_lal, z_sc, y_sc,
                     num_iters = 1000, lr = 1e-3, print_every = 500)
```

Epoch 1: loss 1.4842, loss_y 0.7154 (0.8012, 0.1717), loss_eta 0.7689 (0.8510, 0.1643)
Stopping at iter 38

```
model_lal <- train_xz(model_lal, x_lal, z_sc,
                     num_iters = 1000, lr = 1e-4, print_every = 500)
```

Epoch 1: loss 2.7885, loss1 2.9670, loss2 0.3571
Epoch 500: loss 1.9783, loss1 3.8119, loss2 3.6673
Epoch 1000: loss 1.9913, loss1 3.8261, loss2 3.6696

Now specify a known effect and generate data:

```

model_lal_bench <- specify_causal(model_lal, function(x, eta) {
  0.3 * x + eta
})

# Generate synthetic LaLonde-like data with known ATE = 0.3 (standardized)
syn_lal <- sample_joint(model_lal_bench, sample_size = 2000)
cat("Synthetic LaLonde: n =", nrow(syn_lal$x), "\n")

```

Synthetic LaLonde: n = 2000

```
cat("Treatment prevalence:", round(mean(syn_lal$x), 3), "\n")
```

Treatment prevalence: 0.023

```
cat("True ATE:", 0.3, "(standardized) =", round(0.3 * sd(y_lal), 0), "dollars\n")
```

True ATE: 0.3 (standardized) = 2241 dollars

The synthetic data inherits LaLonde's confounding structure, but now the true effect is known. This is the part that is hard to do with a hand-written DGP.

42.7. Comparison: *causal* vs *frengression*

Both packages use the same decomposition:

Feature	<i>causal</i>	<i>Frengression</i>
Marginal $P(X, A)$	Parametric families	Learned by neural net
Causal margin $P_a(Y)$	Parametric families	Learned (or specified)
Copula	Gaussian, Frank, etc.	Learned by <code>model_eta</code>
Requires specification	Yes — families + params	No — learns from data
Can learn from real data	Limited (plasmode)	Yes — full joint
Distributional flexibility	Limited by family choice	Arbitrary

The *causal* approach is more transparent because we know exactly what we specified. *Frengression* is more flexible because it can learn structure we did not write down. I would treat them as complements.

42.8. Benchmarking: Frengression vs Semiparametric Methods

Now compare frengression against standard semiparametric causal inference tools on the same data. The three comparison methods are:

- `npcausal`: AIPW estimator with SuperLearner nuisance models
- `tmle`: targeted minimum loss-based estimation with SuperLearner
- `DoubleML`: partially linear regression with `ranger`

These methods are built for ATE estimation. They should usually be better for this specific estimand. The comparison is still useful because frengression is trying to learn much more than the ATE.

```
library(npcausal)
library(tmle)
library(SuperLearner)
library(DoubleML)
library(mlr3)
library(mlr3learners)
library(tidyverse)
library(knitr)
```

42.8.1. Single-Dataset Comparison

First, run all four methods on the binary DGP data created above.

```
W <- data.frame(w1, w2, w3, w4)
SL.library <- c("SL.earth", "SL.glm.interaction", "SL.mean",
               "SL.ranger", "SL.glmnet")

set.seed(123)
aipw_fit <- ate(y = Y, a = A, x = W, nsplits = 2, sl.lib = SL.library)
```

		0%
=====		12%
=====		25%
=====		38%
=====		50%
=====		62%


```

|=====| 75%
|
|=====| 88%
|
|=====| 100%

```

	parameter	est	se	ci.ll	ci.ul	pval
1	E{Y(0)}	0.5697312	0.02038159	0.5297832	0.6096791	0
2	E{Y(1)}	0.7608322	0.01207030	0.7371744	0.7844900	0
3	E{Y(1)-Y(0)}	0.1911010	0.02335102	0.1453330	0.2368690	0

```

ate_npcausal <- aipw_fit$res$est[3]
se_npcausal  <- aipw_fit$res$se[3]
ci_npcausal  <- c(aipw_fit$res$ci.ll[3], aipw_fit$res$ci.ul[3])

cat("npcausal ATE:", round(ate_npcausal, 4), "\n")

```

npcausal ATE: 0.1911

```
cat("SE:", round(se_npcausal, 4), "\n")
```

SE: 0.0234

```

set.seed(123)
tmle_fit <- tmle(Y = Y, A = A, W = W, family = "binomial",
                 Q.SL.library = SL.library, g.SL.library = SL.library)

ate_tmle <- tmle_fit$estimates$ATE$psi
se_tmle  <- sqrt(tmle_fit$estimates$ATE$var.psi)
ci_tmle  <- tmle_fit$estimates$ATE$CI

cat("TMLE ATE:", round(ate_tmle, 4), "\n")

```

TMLE ATE: 0.1916

```
cat("SE:", round(se_tmle, 4), "\n")
```

SE: 0.0209

```

lgr::get_logger("mlr3")$set_threshold("warn")
data_obs <- data.frame(w1, w2, w3, w4, A, Y)

dml_data <- DoubleMLData$new(data_obs,
                             y_col = "Y", d_cols = "A",
                             x_cols = c("w1", "w2", "w3", "w4"))

```

42. Frengression and Engression

```
learner_l <- lrn("regr.ranger", num.trees = 500, max.depth = 5, min.node.size = 2)
learner_m <- learner_l$clone()

set.seed(123)
dml_plr <- DoubleMLPLR$new(dml_data, ml_l = learner_l, ml_m = learner_m)
dml_plr$fit()

ate_dml <- dml_plr$coef
se_dml <- dml_plr$se
ci_dml <- c(dml_plr$confint()[1], dml_plr$confint()[2])

cat("DoubleML ATE:", round(ate_dml, 4), "\n")
```

DoubleML ATE: 0.1872

```
cat("SE:", round(se_dml, 4), "\n")
```

SE: 0.0229

```
# Frengression ATE already computed above (y1 - y0)
ate_freng <- as.numeric(y1 - y0)

# Bootstrap SE via repeated prediction (Monte Carlo variability)
set.seed(123)
ate_boot <- replicate(200, {
  predict(model, matrix(1, ncol = 1), type = "mean", nsample = 500) -
  predict(model, matrix(0, ncol = 1), type = "mean", nsample = 500)
})
se_freng <- sd(ate_boot)
```

```
results <- data.frame(
  Method = c("True ATE", "Frengression", "npcausal (AIPW)", "TMLE", "DoubleML (PLR)"),
  ATE = c(true_ate, ate_freng, ate_npcausal, ate_tmle, ate_dml),
  SE = c(NA, se_freng, se_npcausal, se_tmle, se_dml),
  CI_lower = c(NA,
    ate_freng - 1.96 * se_freng,
    ci_npcausal[1], ci_tmle[1], ci_dml[1]),
  CI_upper = c(NA,
    ate_freng + 1.96 * se_freng,
    ci_npcausal[2], ci_tmle[2], ci_dml[2])
)

results$Bias <- results$ATE - true_ate
results$Covers <- ifelse(is.na(results$CI_lower), NA,
  results$CI_lower <= true_ate & true_ate <= results$CI_upper)
```

```
kable(results, digits = 4,
       caption = "ATE Estimation: Frengression vs Semiparametric Methods")
```

Table 42.3.: ATE Estimation: Frengression vs Semiparametric Methods

Method	ATE	SE	CI_lower	CI_upper	Bias	Covers
True ATE	0.1935	NA	NA	NA	0.0000	NA
Frengression	0.1656	0.0316	0.1036	0.2275	-0.0279	TRUE
npcausal (AIPW)	0.1911	0.0234	0.1453	0.2369	-0.0024	TRUE
TMLE	0.1916	0.0209	0.1506	0.2326	-0.0019	TRUE
DoubleML (PLR)	0.1872	0.0229	0.1424	0.2320	-0.0063	TRUE

The semiparametric estimators target the ATE directly. Frengression learns the joint distribution and then computes the ATE as one functional of that distribution. So I would not expect it to be the most precise ATE estimator in this table.

42.8.2. Monte Carlo Benchmarking (Simulated DGP)

This is where frengression is more useful. Use the model trained above, plug in a known causal effect with `specify_causal()`, generate 10 synthetic datasets, and benchmark all four estimators.

```
B <- 10
n_syn <- 2000

methods <- c("Frengression", "npcausal", "TMLE", "DoubleML")
ate_mat <- matrix(NA, nrow = B, ncol = length(methods),
                  dimnames = list(NULL, methods))
se_mat <- ate_mat
cover_mat <- ate_mat

SL.lib_fast <- c("SL.glm.interaction", "SL.ranger", "SL.mean")

# True ATE from the specified model (already computed above as y1_true - y0_true)
true_ate_bench <- as.numeric(y1_true - y0_true)

set.seed(2024)
for (b in seq_len(B)) {
  syn <- sample_joint(model_bench, sample_size = n_syn)

  x_syn <- as.numeric(syn$x)
  y_syn <- as.numeric(syn$y)
  z_syn <- as.data.frame(syn$z)
  colnames(z_syn) <- paste0("z", 1:ncol(z_syn))
  df_syn <- cbind(data.frame(A = x_syn, Y = y_syn), z_syn)
```

```

# --- Frengression ---
tryCatch({
  mod_f <- frengression(x_dim = 1, y_dim = 1, z_dim = ncol(z_syn),
                        x_binary = TRUE, y_binary = TRUE,
                        noise_dim = 10, hidden_dim = 100, num_layer = 3)
  mod_f <- train_y(mod_f, matrix(x_syn, ncol=1), as.matrix(z_syn),
                  matrix(y_syn, ncol=1),
                  num_iters = 500, lr = 1e-3, silent = TRUE)
  mod_f <- train_xz(mod_f, matrix(x_syn, ncol=1), as.matrix(z_syn),
                  num_iters = 500, lr = 1e-4, silent = TRUE)

  y1_f <- predict(mod_f, matrix(1, ncol=1), type="mean", nsample=1000)
  y0_f <- predict(mod_f, matrix(0, ncol=1), type="mean", nsample=1000)
  ate_mat[b, "Frengression"] <- y1_f - y0_f

  ate_mc <- replicate(100, {
    predict(mod_f, matrix(1, ncol=1), type="mean", nsample=200) -
    predict(mod_f, matrix(0, ncol=1), type="mean", nsample=200)
  })
  se_mat[b, "Frengression"] <- sd(ate_mc)
  ci_f <- ate_mat[b, "Frengression"] + c(-1, 1) * 1.96 * se_mat[b, "Frengression"]
  cover_mat[b, "Frengression"] <- (ci_f[1] <= true_ate_bench & true_ate_bench <= ci_f[2])
}, error = function(e) message("Frengression failed on rep ", b, ": ", e$message))

# --- npcausal ---
tryCatch({
  fit_np <- ate(y = y_syn, a = x_syn, x = z_syn,
               nsplits = 2, sl.lib = SL.lib_fast)
  ate_mat[b, "npcausal"] <- fit_np$res$est[3]
  se_mat[b, "npcausal"] <- fit_np$res$se[3]
  ci <- c(fit_np$res$ci.ll[3], fit_np$res$ci.ul[3])
  cover_mat[b, "npcausal"] <- (ci[1] <= true_ate_bench & true_ate_bench <= ci[2])
}, error = function(e) message("npcausal failed on rep ", b, ": ", e$message))

# --- TMLE ---
tryCatch({
  fit_tmle <- tmle(Y = y_syn, A = x_syn, W = z_syn, family = "binomial",
                  Q.SL.library = SL.lib_fast, g.SL.library = SL.lib_fast)
  ate_mat[b, "TMLE"] <- fit_tmle$estimates$ATE$psi
  se_mat[b, "TMLE"] <- sqrt(fit_tmle$estimates$ATE$var.psi)
  ci_t <- fit_tmle$estimates$ATE$CI
  cover_mat[b, "TMLE"] <- (ci_t[1] <= true_ate_bench & true_ate_bench <= ci_t[2])
}, error = function(e) message("TMLE failed on rep ", b, ": ", e$message))

# --- DoubleML ---
tryCatch({
  dml_d <- DoubleMLData$new(df_syn, y_col = "Y", d_cols = "A",

```

```

                                x_cols = colnames(z_syn))
lr_l <- lrn("regr.ranger", num.trees = 300, max.depth = 5, min.node.size = 2)
lr_m <- lr_l$clone()
dml_obj <- DoubleMLPLR$new(dml_d, ml_l = lr_l, ml_m = lr_m)
dml_obj$fit()
ate_mat[b, "DoubleML"] <- dml_obj$coef
se_mat[b, "DoubleML"] <- dml_obj$se
ci_d <- c(dml_obj$confint()[1], dml_obj$confint()[2])
cover_mat[b, "DoubleML"] <- (ci_d[1] <= true_ate_bench & true_ate_bench <= ci_d[2])
}, error = function(e) message("DoubleML failed on rep ", b, ": ", e$message))

if (b %% 5 == 0) cat("Completed replication", b, "of", B, "\n")
}

```

Completed replication 5 of 10

Completed replication 10 of 10

```

bench_summary <- data.frame(
  Method = colnames(ate_mat),
  Mean_ATE = colMeans(ate_mat, na.rm = TRUE),
  Bias = colMeans(ate_mat, na.rm = TRUE) - true_ate_bench,
  SD = apply(ate_mat, 2, sd, na.rm = TRUE),
  RMSE = sqrt(colMeans((ate_mat - true_ate_bench)^2, na.rm = TRUE)),
  Mean_SE = colMeans(se_mat, na.rm = TRUE),
  Coverage_95 = colMeans(cover_mat, na.rm = TRUE),
  N_valid = colSums(!is.na(ate_mat))
)

kable(bench_summary, digits = 4,
      caption = paste0("Monte Carlo Benchmarking (B=", B,
                        ", n=", n_syn, ", True ATE=",
                        round(true_ate_bench, 4), ")"))

```

Table 42.4.: Monte Carlo Benchmarking (B=10, n=2000, True ATE=0.0871)

	Method	Mean_ATE	Bias	SD	RMSE	Mean_SE	Cover- age_95	N_valid
Frengres- sion	Frengres- sion	0.0490	-0.0381	0.0304	0.0478	0.0562	1.0	10
npcausal	npcausal	0.2006	0.1134	0.3818	0.3796	0.0960	0.6	10
TMLE	TMLE	0.0136	-0.0735	0.4469	0.4303	0.0127	0.0	10
DoubleML	DoubleML	0.0450	-0.0422	0.0691	0.0780	0.0821	1.0	10

```

ate_long <- as.data.frame(ate_mat) |>
  pivot_longer(everything(), names_to = "Method", values_to = "ATE") |>
  filter(!is.na(ATE))

# Trim extreme outliers for readable plot
q_bounds <- quantile(ate_long$ATE, c(0.02, 0.98), na.rm = TRUE)
ate_long_trim <- ate_long |> filter(ATE >= q_bounds[1] & ATE <= q_bounds[2])

ggplot(ate_long_trim, aes(x = Method, y = ATE, fill = Method)) +
  geom_boxplot(alpha = 0.7) +
  geom_hline(yintercept = true_ate_bench, linetype = "dashed", color = "red",
    linewidth = 1) +
  annotate("text", x = 0.5, y = true_ate_bench, label = "True ATE",
    hjust = 0, vjust = -0.5, color = "red") +
  labs(title = "ATE Estimates Across Frengression-Generated Datasets",
    subtitle = paste0("B=", B, " replications, n=", n_syn, " per dataset"),
    y = "Estimated ATE") +
  theme_minimal() +
  theme(legend.position = "none")

```




Figure 42.1.: ATE Estimates Across Synthetic Datasets

42.8.3. Monte Carlo Benchmarking (LaLonde)

Now do the same exercise with the LaLonde-trained model. Here the confounding structure comes from the real data rather than from my hand-written simulation.

```
# Compute true ATE empirically from the specified model
y1_lal_true <- predict(model_lal_bench, matrix(1, ncol=1), type="mean", nsample=5000)
y0_lal_true <- predict(model_lal_bench, matrix(0, ncol=1), type="mean", nsample=5000)
true_ate_lal <- as.numeric(y1_lal_true - y0_lal_true)
cat("True ATE (standardized):", round(true_ate_lal, 4), "\n")
```

True ATE (standardized): 0.3101

```
cat("True ATE (dollars):", round(true_ate_lal * sd(y_lal), 0), "\n")
```

True ATE (dollars): 2317

```
B_lal <- 10
n_syn_lal <- 2000

methods <- c("Frengression", "npcausal", "TMLE", "DoubleML")
ate_mat_lal <- matrix(NA, nrow = B_lal, ncol = length(methods),
                     dimnames = list(NULL, methods))
se_mat_lal <- ate_mat_lal
cover_mat_lal <- ate_mat_lal

SL.lib_fast <- c("SL.glm.interaction", "SL.ranger", "SL.mean")

set.seed(2024)
for (b in seq_len(B_lal)) {
  syn <- sample_joint(model_lal_bench, sample_size = n_syn_lal)

  x_syn <- as.numeric(syn$x)
  y_syn <- as.numeric(syn$y)
  z_syn <- as.data.frame(syn$z)
  colnames(z_syn) <- paste0("z", 1:ncol(z_syn))
  df_syn <- cbind(data.frame(A = x_syn, Y = y_syn), z_syn)

  # --- Frengression ---
  tryCatch({
    mod_f <- frengression(x_dim = 1, y_dim = 1, z_dim = ncol(z_syn),
                        x_binary = TRUE, y_binary = FALSE,
                        noise_dim = 10, hidden_dim = 100, num_layer = 3)
    mod_f <- train_y(mod_f, matrix(x_syn, ncol=1), as.matrix(z_syn),
                    matrix(y_syn, ncol=1),
                    num_iters = 500, lr = 1e-3, silent = TRUE)
    mod_f <- train_xz(mod_f, matrix(x_syn, ncol=1), as.matrix(z_syn),
                    num_iters = 500, lr = 1e-4, silent = TRUE)

    y1_f <- predict(mod_f, matrix(1, ncol=1), type="mean", nsample=1000)
    y0_f <- predict(mod_f, matrix(0, ncol=1), type="mean", nsample=1000)
```

```

ate_mat_lal[b, "Frengression"] <- y1_f - y0_f

ate_mc <- replicate(100, {
  predict(mod_f, matrix(1,ncol=1), type="mean", nsample=200) -
  predict(mod_f, matrix(0,ncol=1), type="mean", nsample=200)
})
se_mat_lal[b, "Frengression"] <- sd(ate_mc)
ci_f <- ate_mat_lal[b, "Frengression"] + c(-1,1) * 1.96 * se_mat_lal[b, "Frengression"]
cover_mat_lal[b, "Frengression"] <- (ci_f[1] <= true_ate_lal & true_ate_lal <= ci_f[2])
}, error = function(e) message("Frengression failed on rep ", b, ": ", e$message))

# --- npcausal ---
tryCatch({
  fit_np <- ate(y = y_syn, a = x_syn, x = z_syn,
               nsplits = 2, sl.lib = SL.lib_fast)
  ate_mat_lal[b, "npcausal"] <- fit_np$res$est[3]
  se_mat_lal[b, "npcausal"] <- fit_np$res$se[3]
  ci <- c(fit_np$res$ci.ll[3], fit_np$res$ci.ul[3])
  cover_mat_lal[b, "npcausal"] <- (ci[1] <= true_ate_lal & true_ate_lal <= ci[2])
}, error = function(e) message("npcausal failed on rep ", b, ": ", e$message))

# --- TMLE ---
tryCatch({
  fit_tmle <- tmle(Y = y_syn, A = x_syn, W = z_syn, family = "gaussian",
                  Q.SL.library = SL.lib_fast, g.SL.library = SL.lib_fast)
  ate_mat_lal[b, "TMLE"] <- fit_tmle$estimates$ATE$psi
  se_mat_lal[b, "TMLE"] <- sqrt(fit_tmle$estimates$ATE$var.psi)
  ci_t <- fit_tmle$estimates$ATE$CI
  cover_mat_lal[b, "TMLE"] <- (ci_t[1] <= true_ate_lal & true_ate_lal <= ci_t[2])
}, error = function(e) message("TMLE failed on rep ", b, ": ", e$message))

# --- DoubleML ---
tryCatch({
  dml_d <- DoubleMLData$new(df_syn, y_col = "Y", d_cols = "A",
                           x_cols = colnames(z_syn))
  lr_l <- lrn("regr.ranger", num.trees = 300, max.depth = 5, min.node.size = 2)
  lr_m <- lr_l$clone()
  dml_obj <- DoubleMLPLR$new(dml_d, ml_l = lr_l, ml_m = lr_m)
  dml_obj$fit()
  ate_mat_lal[b, "DoubleML"] <- dml_obj$coef
  se_mat_lal[b, "DoubleML"] <- dml_obj$se
  ci_d <- c(dml_obj$confint()[1], dml_obj$confint()[2])
  cover_mat_lal[b, "DoubleML"] <- (ci_d[1] <= true_ate_lal & true_ate_lal <= ci_d[2])
}, error = function(e) message("DoubleML failed on rep ", b, ": ", e$message))

if (b %% 5 == 0) cat("Completed replication", b, "of", B_lal, "\n")
}

```


Completed replication 5 of 10

Completed replication 10 of 10

```
bench_lal <- data.frame(
  Method = colnames(ate_mat_lal),
  Mean_ATE = colMeans(ate_mat_lal, na.rm = TRUE),
  Bias = colMeans(ate_mat_lal, na.rm = TRUE) - true_ate_lal,
  SD = apply(ate_mat_lal, 2, sd, na.rm = TRUE),
  RMSE = sqrt(colMeans((ate_mat_lal - true_ate_lal)^2, na.rm = TRUE)),
  Mean_SE = colMeans(se_mat_lal, na.rm = TRUE),
  Coverage_95 = colMeans(cover_mat_lal, na.rm = TRUE),
  N_valid = colSums(!is.na(ate_mat_lal))
)

kable(bench_lal, digits = 4,
      caption = paste0("LaLonde-Based Benchmarking (B=", B_lal,
                        ", n=", n_syn_lal,
                        ", True ATE=", round(true_ate_lal, 4),
                        " [", round(true_ate_lal * sd(y_lal), 0), " dollars]"))
```

Table 42.5.: LaLonde-Based Benchmarking (B=10, n=2000, True ATE=0.318 [2376 dollars])

	Method	Mean_ATE	Bias	SD	RMSE	Mean_SE	Cover- age_95	N_valid
Frengres- sion	Frengres- sion	0.1006	-0.2174	0.1379	0.2537	0.1070	0.5	10
npcausal	npcausal	0.0667	-0.2513	0.2768	0.3634	0.0557	0.1	10
TMLE	TMLE	0.2599	-0.0581	0.2449	0.2395	0.0312	0.1	10
DoubleML	DoubleML	-0.0884	-0.4064	0.0480	0.4089	0.0821	0.0	10

```
ate_long_lal <- as.data.frame(ate_mat_lal) |>
  pivot_longer(everything(), names_to = "Method", values_to = "ATE") |>
  filter(!is.na(ATE))

q_bounds_lal <- quantile(ate_long_lal$ATE, c(0.02, 0.98), na.rm = TRUE)
ate_long_lal_trim <- ate_long_lal |>
  filter(ATE >= q_bounds_lal[1] & ATE <= q_bounds_lal[2])

ggplot(ate_long_lal_trim, aes(x = Method, y = ATE, fill = Method)) +
  geom_boxplot(alpha = 0.7) +
  geom_hline(yintercept = true_ate_lal, linetype = "dashed", color = "red",
            linewidth = 1) +
  annotate("text", x = 0.5, y = true_ate_lal, label = "True ATE",
          hjust = 0, vjust = -0.5, color = "red") +
```

```
labs(title = "ATE Estimates: LaLonde-Based Frengression DGP",
     subtitle = paste0("B=", B_lal, " replications, n=", n_syn_lal,
                        " per dataset (confounding learned from real data)"),
     y = "Estimated ATE (standardized)") +
theme_minimal() +
theme(legend.position = "none")
```




Figure 42.2.: ATE Estimates from LaLonde-Based Synthetic Data

The semiparametric methods show more bias on the LaLonde-generated data than on the hand-written simulation. That is the point of this exercise. The synthetic data has a more realistic propensity surface and prior-earnings distribution, so it is harder than the usual clean simulation.

42.9. Summary

Frengression uses the Evans and Didelez simulation decomposition, but learns the pieces nonparametrically. The practical use is not that it beats TMLE or AIPW for a simple ATE. The practical use is that it can learn a realistic observational distribution, let us impose a known causal effect, and then generate synthetic data where the truth is known.

That makes it useful for benchmarking. The simulation is still artificial, but it is less artificial than writing down a clean logistic model and pretending that it looks like real observational data.

Systematic treatment: [Julia](#).

43. Partial Interference

Most causal inference methods assume SUTVA: one unit’s treatment does not affect another unit’s outcome. This is often not true. Vaccinating your neighbor can reduce your risk of infection. A job training program for one worker may affect another worker’s job prospect. A student’s test score may depend on whether classmates received tutoring.

Partial interference (Sobel, M.E. (2006), “What Do Randomized Studies of Housing Mobility Demonstrate?”, *Journal of the American Statistical Association*, 101(476); Hudgens, M.G. & Halloran, M.E. (2008), “Toward Causal Inference With Interference”, *Journal of the American Statistical Association*, 103(482)) is one way to make this manageable. Units are organized into clusters. Interference can happen within a cluster, but not between clusters. This is still a strong assumption, but it is weaker than no interference.

Qu et al. (2021) develop clustered IPW and AIPW estimators for this setting – see the `CausalModel` R package’s documentation/vignette for the exact reference. Here I just simulate data and see how the functions work.

43.1. The Setup

43.1.1. Clusters and Groups

Each unit i belongs to a cluster c , such as a household, classroom, or village. Within each cluster, units are also organized into groups indexed by $j = 0, 1, \dots, J - 1$. The group structure $\mathbf{s} = (s_0, s_1, \dots, s_{J-1})$ gives the number of units per group.

For example, with `group_struct = c(2, 3)`:

- Group 0 has 2 units
- Group 1 has 3 units
- Each cluster has 5 units total

The groups can be real roles, such as parents and children, or they can just be a modeling device.

43.1.2. The Neighborhood Configuration $\mathbf{G}$

The key object is the neighborhood configuration $G_i = (g_0, g_1, \dots, g_{J-1})$. Here g_j counts how many other units in group j within unit i ’s cluster are treated. “Other” means excluding unit i itself.

With `group_struct = c(2, 3)`:

- If unit i is in group 0, then $g_0 \in \{0, 1\}$ (the other unit in group 0 is treated or not) and $g_1 \in \{0, 1, 2, 3\}$ (0 to 3 units in group 1 are treated)

43. Partial Interference

- G records the treatment environment that unit i experiences

43.1.3. Heterogeneous Treatment Effects (G)

Under partial interference, the treatment effect for unit i depends on its neighborhood G_i :

$$\beta(G) = E[Y_i(1, G) - Y_i(0, G)] \quad (43.1)$$

where $Y_i(z, G)$ is unit i 's potential outcome when its own treatment is z and its neighborhood configuration is G .

A simple model is to make the treatment effect linear in the neighborhood:

$$\beta(G) = \tau + \sum_{j=0}^{J-1} \gamma_j \cdot g_j \quad (43.2)$$

where:

- τ is the direct effect when no neighbors are treated
- γ_j is the spillover from group j , per treated neighbor

43.1.3.1. Example

With `group_struct = c(2, 3)`, $\tau = 1$, $\gamma_0 = 0.5$, $\gamma_1 = 0.1$:

Neighborhood G	Formula	$\beta(G)$
(0, 0)	$1 + 0 \times 0.5 + 0 \times 0.1$	1.0
(1, 0)	$1 + 1 \times 0.5 + 0 \times 0.1$	1.5
(0, 1)	$1 + 0 \times 0.5 + 1 \times 0.1$	1.1
(1, 1)	$1 + 1 \times 0.5 + 1 \times 0.1$	1.6

Having one treated neighbor in group 0 adds 0.5 to the treatment effect; having one in group 1 adds only 0.1.

43.2. Estimation

43.2.1. Two Propensity Models

Standard IPW uses one propensity score, $P(Z_i = 1|X_i)$. Under interference, we need two:

1. **Individual propensity:** $P(Z_i = 1|X_i)$ — probability of being treated given covariates (logistic regression)
2. **Neighborhood propensity:** $P(G_i = g|Z_i, X_i)$ — probability of neighborhood configuration g given **own treatment** and covariates (multinomial logistic regression)

The second piece is the new part. It models how likely a particular treatment environment is, conditional on covariates.

Note the conditioning on Z_i in the second piece. Own treatment Z_i and the neighborhood exposure G_i are usually **dependent by design** — for a network or cluster exposure mapping, G_i is built from the treatments of i 's neighbours, which are correlated with i 's own assignment mechanism. The joint exposure probability must therefore be modeled as a joint distribution $P(Z_i = z, G_i = g | X_i)$, or via the valid factorization $P(Z_i = z | X_i) P(G_i = g | Z_i = z, X_i)$. A product of the two *marginals* $P(Z_i | X_i)P(G_i | X_i)$ is only correct when own and neighborhood treatment are conditionally independent given X_i , which is not the typical case.

43.2.2. Clustered IPW

The clustered IPW estimator for $\beta(g)$ reweights outcomes by the joint exposure probability, written here as the conditional factorization:

$$\hat{\beta}_{\text{IPW}}(g) = \frac{1}{N} \sum_i \left[\frac{\mathbf{1}(Z_i = 1, G_i = g)}{P(Z_i = 1 | X_i) \cdot P(G_i = g | Z_i = 1, X_i)} - \frac{\mathbf{1}(Z_i = 0, G_i = g)}{P(Z_i = 0 | X_i) \cdot P(G_i = g | Z_i = 0, X_i)} \right] Y_i \quad (43.3)$$

This gives a separate treatment-effect estimate at each neighborhood configuration g .

43.2.3. Clustered AIPW (Doubly Robust)

The AIPW version adds outcome models:

- Fit separate outcome models $E[Y | X, Z, G = g]$ for each configuration g
- Combine with propensity weighting so the estimator is consistent if either the outcome model or the (joint) exposure-propensity model is correct. Here the propensity part must correctly model the *joint* exposure mechanism $P(Z_i, G_i | X_i)$, not two independent marginals.

This is the estimator I would usually try first, for the same reason standard AIPW is attractive.

43.2.4. Variance Estimation

Standard errors are computed by a matching-based variance estimator. Within each group and neighborhood configuration, units are matched on standardized covariates, and the variance is estimated from the matched pairs. This avoids a cluster bootstrap, which can be expensive here.

43.3. Using CausalModel

```
library(CausalModel)
library(knitr)
set.seed(42)
```

43.3.1. Generate Clustered Data

The `generate_fixed_cluster()` function creates synthetic data where the interference structure is known.

```
cdata <- generate_fixed_cluster(
  clusters = 200,
  group_struct = c(2, 3), # 2 units in group 0, 3 in group 1
  tau = 1, # direct effect
  gamma = c(0.5, 0.1) # spillover: 0.5 per treated in group 0, 0.1 in group 1
)

cat("Total units:", length(cdata$Y), "\n")
```

Total units: 1000

```
cat("Units per cluster:", sum(c(2, 3)), "\n")
```

Units per cluster: 5

```
cat("Treatment prevalence:", round(mean(cdata$Z), 3), "\n")
```

Treatment prevalence: 0.477

Each unit has:

- `Y` — observed outcome
- `Z` — treatment indicator (0/1)
- `X` — individual covariates
- `cluster_labels` — which cluster the unit belongs to
- `group_labels` — which group within the cluster (0-indexed)
- `ingroup_labels` — position within the group (0-indexed)

43.3.2. Construct the Clustered Object

```
cl_obj <- clustered(
  Y = cdata$Y,
  Z = cdata$Z,
  X = cdata$X,
  cluster_labels = cdata$cluster_labels,
  group_labels = cdata$group_labels,
  ingroup_labels = cdata$ingroup_labels,
  n_matches = 50
)
```

The `clustered()` constructor does several things:

1. Augments individual covariates with cluster-level aggregates (mean of neighbors' covariates)
2. Fits the individual propensity model (logistic regression)
3. Fits the neighborhood propensity model (multinomial logistic regression)
4. Computes G for each unit — the number of treated neighbors per group

43.3.3. Estimate Treatment Effects

43.3.3.1. Clustered AIPW

```
result_aipw <- est_via_aipw(cl_obj)
```

The result is a list with one element per group. Each element contains `beta_g`, the treatment effects at each neighborhood configuration, and `se`, the standard errors:

```
for (j in seq_along(result_aipw)) {
  cat(sprintf("Group %d:\n", j - 1))
  valid <- !is.nan(result_aipw[[j]]$beta_g) & result_aipw[[j]]$se > 0
  df <- data.frame(
    g_index = which(valid) - 1,
    beta_g = round(result_aipw[[j]]$beta_g[valid], 3),
    se = round(result_aipw[[j]]$se[valid], 3)
  )
  print(df, row.names = FALSE)
  cat("\n")
}
```

Group 0:

g_index	beta_g	se
0	0.677	0.364
1	1.572	0.347
3	1.447	0.315
4	1.870	0.257
6	1.141	0.276
7	2.003	0.276
9	1.015	0.428
10	2.674	0.384

Group 1:

g_index	beta_g	se
0	1.528	0.311
1	1.634	0.228
2	2.126	0.331
3	1.280	0.266

43. Partial Interference

```

4  1.799 0.188
5  2.180 0.269
6  1.272 0.419
7  2.137 0.247
8  2.850 0.389

```

43.3.3.2. Clustered IPW

```
result_ipw <- est_via_ipw(cl_obj)
```

```
cat("Clustered IPW beta_g (group 0):\n")
```

Clustered IPW beta_g (group 0):

```
valid <- !is.nan(result_ipw[[1]]$beta_g) & result_ipw[[1]]$se > 0
round(result_ipw[[1]]$beta_g[valid], 3)
```

```
[1] 0.491 1.591 1.404 1.865 0.984 2.002 1.182 2.327
```

43.3.4. Interpreting the Results

The neighborhood configurations are encoded as integers. For `group_struct = c(2, 3)`, the encoding for group 0 is:

$$G_{\text{encoded}} = g_0 + g_1 \times (s_0 + 1) \quad (43.4)$$

where $s_0 = 2$ is the size of group 0. The stride for the second dimension is $s_0 + 1 = 3$ because the package allocates slots for $g_0 \in \{0, 1, 2\}$, even though $g_0 = 2$ is impossible for units in group 0.

g_0	g_1	Encoded	True β
0	0	0	1.0
1	0	1	1.5
0	1	3	1.1
1	1	4	1.6
0	2	6	1.2
1	2	7	1.7

Some configurations do not have enough observations. The `se > 0` filter removes degenerate cases.

43.4. Monte Carlo Validation

Now run a small Monte Carlo. Generate data 100 times, estimate $\beta(G)$ each time, and check bias, SE calibration, and coverage.

```
n_clusters <- 200
group_struct <- c(2, 3)
tau_int <- 1
gamma_int <- c(0.5, 0.1)
n_reps <- 100

# True beta(g) function
true_beta <- function(g) tau_int + sum(g * gamma_int)

# Track results for group 0 at three neighborhood configurations
g_indices <- list(c(0, 0), c(1, 0), c(0, 1))
g_encoded <- sapply(g_indices, function(g) {
  g[1] + g[2] * (group_struct[1] + 1) + 1 # 1-indexed; stride = max_g0 + 1
})
true_values <- sapply(g_indices, true_beta)

int_results <- array(NA, dim = c(n_reps, length(g_indices), 2),
  dimnames = list(NULL, NULL, c("beta", "se")))
```

```
set.seed(2024)
for (i in seq_len(n_reps)) {
  cdata <- generate_fixed_cluster(
    clusters = n_clusters, group_struct = group_struct,
    tau = tau_int, gamma = gamma_int
  )
  cl_obj <- clustered(
    Y = cdata$Y, Z = cdata$Z, X = cdata$X,
    cluster_labels = cdata$cluster_labels,
    group_labels = cdata$group_labels,
    ingroup_labels = cdata$ingroup_labels,
    n_matches = 50
  )
  res <- est_via_aipw(cl_obj)

  for (j in seq_along(g_indices)) {
    idx <- g_encoded[j]
    int_results[i, j, "beta"] <- res[[1]]$beta_g[idx]
    int_results[i, j, "se"] <- res[[1]]$se[idx]
  }
}
```

```

int_rows <- list()
for (j in seq_along(g_indices)) {
  betas <- int_results[, j, "beta"]
  ses <- int_results[, j, "se"]

  valid <- !is.nan(betas) & !is.nan(ses) & ses > 0
  betas <- betas[valid]
  ses <- ses[valid]
  tv <- true_values[j]

  ci_lo <- betas - 1.96 * ses
  ci_hi <- betas + 1.96 * ses

  int_rows[[j]] <- data.frame(
    g = paste0("(", paste(g_indices[[j]], collapse = ","), ")"),
    True_beta = tv,
    Mean_est = mean(betas),
    Bias = mean(betas) - tv,
    Emp_SE = sd(betas),
    Mean_SE = mean(ses),
    Coverage = mean(ci_lo <= tv & tv <= ci_hi),
    N_valid = length(betas)
  )
}
tbl_int <- do.call(rbind, int_rows)
kable(tbl_int, digits = 3, row.names = FALSE,
      caption = sprintf("Clustered AIPW (group 0, %d clusters, %d reps)",
                        n_clusters, n_reps))

```

Table 43.3.: Clustered AIPW (group 0, 200 clusters, 100 reps)

g	True_beta	Mean_est	Bias	Emp_SE	Mean_SE	Coverage	N_valid
(0,0)	1.0	2.107	1.107	8.760	0.448	0.820	100
(1,0)	1.5	1.569	0.069	0.549	0.415	0.899	99
(0,1)	1.1	1.083	-0.017	0.248	0.246	0.930	100

```

par(mfrow = c(1, length(g_indices)), mar = c(3, 3, 2, 1))
for (j in seq_along(g_indices)) {
  betas <- int_results[, j, "beta"]
  ses <- int_results[, j, "se"]
  valid <- !is.nan(betas) & !is.nan(ses) & ses > 0
  t_stats <- (betas[valid] - true_values[j]) / ses[valid]
  g_label <- paste0("(", paste(g_indices[[j]], collapse = ","), ")")
  hist(t_stats, breaks = 20, freq = FALSE, col = "lightblue",
       main = sprintf("g = %s", g_label),
       xlab = "", xlim = c(-4, 4))
}

```

```

curve(dnorm(x), add = TRUE, col = "red", lwd = 2)
abline(v = 0, lty = 2)
}

```




Figure 43.1.: Studentized statistics under clustered AIPW

If both the estimator and SE work well, the studentized statistics should look roughly normal.

43.5. What If You Ignore Interference?

What happens if we ignore interference and use a standard AIPW estimator? We get one number that averages over neighborhood configurations. That number is hard to interpret.

```

# Generate one dataset with interference
set.seed(42)
cdata <- generate_fixed_cluster(
  clusters = 200,
  group_struct = c(2, 3),
  tau = 1,
  gamma = c(0.5, 0.1)
)

# Naive: ignore clusters, treat as standard observational data

```

43. Partial Interference

```
obs_naive <- observational(cdata$Y, cdata$Z, cdata$X)
naive_ate <- est_via_aipw(obs_naive)

cat("Naive AIPW (ignoring interference):", round(naive_ate$ate, 3), "\n")
```

Naive AIPW (ignoring interference): 1.611

```
cat("SE:", round(naive_ate$se, 3), "\n\n")
```

SE: 0.068

```
# Correct: use clustered estimator
cl_obj <- clustered(
  Y = cdata$Y, Z = cdata$Z, X = cdata$X,
  cluster_labels = cdata$cluster_labels,
  group_labels = cdata$group_labels,
  ingroup_labels = cdata$ingroup_labels,
  n_matches = 50
)
result <- est_via_aipw(cl_obj)

cat("Clustered AIPW beta_g (group 0):\n")
```

Clustered AIPW beta_g (group 0):

```
valid <- !is.nan(result[[1]]$beta_g) & result[[1]]$se > 0
df <- data.frame(
  g_index = which(valid) - 1,
  beta_g = round(result[[1]]$beta_g[valid], 3),
  se = round(result[[1]]$se[valid], 3)
)
print(df, row.names = FALSE)
```

g_index	beta_g	se
0	0.677	0.364
1	1.572	0.347
3	1.447	0.315
4	1.870	0.257
6	1.141	0.276
7	2.003	0.276
9	1.015	0.428
10	2.674	0.384

The naive estimate gives some weighted average of the $\beta(g)$ values, but it does not tell us how much is direct effect and how much is spillover. The clustered estimator separates those pieces. If spillovers are large, treating clusters may be more effective than treating isolated individuals.

43.6. Summary

Partial interference is a compromise. We do not assume away spillovers, but we do assume they stay inside clusters. Once we make that assumption, the neighborhood configuration G becomes part of the estimand. The effect is no longer just treatment versus control; it is treatment versus control under a particular neighborhood treatment environment.

The `CausalModel` package estimates these effects with clustered IPW and AIPW. The important practical lesson is that a standard AIPW estimate is not a good substitute. It collapses direct effects and spillovers into one number.

44. Power analysis with list experiment

This post is based on Ulrich, et al 2012.

44.1. List experiment

List experiment is also called item count technique (ICT), which was developed by Miller 1984, also called unmatched count technique, the unmatched block design, or the block total response. This technique requires two groups of respondents. First, a control group with n_c respondents receives a list of k neutral questions. The probability of answering the neutral question N_i with “yes” is $\pi_i (i = 1, \dots, k)$. Second, an experimental group with n_e respondents receives the same set of questions as the control group plus the sensitive question of interest. In each group, the respondent is asked to indicate the total count of “yes” responses. Because each item is a binary variable, the expected number of total “yes” responses is $\mu_c = \sum_{i=1}^k \pi_i$ for the control group and $\mu_e = \sum_{i=1}^k \pi_i + \pi_s$ for the experimental group. Therefore, the difference $\hat{\mu}_e - \hat{\mu}_c$ between the observed means of the two groups provides an estimate of π_s , that is,

$$\hat{\pi}_s = \hat{\mu}_e - \hat{\mu}_c. \quad (44.1)$$

44.2. Power analysis

44.2.1. Define power

Suppose π_s is the proportion of the population that has the sensitive attribute, or the proportion of the population that would answer the sensitive question with “yes”, such as the proportion of people who have committed a crime. The null hypothesis H_0 is that $\pi_s = 0$, and the alternative hypothesis H_1 is that $\pi_s > 0$.

The expected mean and the sampling variance of $\hat{\pi}_s$ under the null hypothesis H_0 are denoted μ_0 and σ_0^2 , respectively. Likewise, let μ_1 and σ_1^2 be the mean and the variance under the alternative hypothesis H_1 that π_s has a specific value that is greater than zero, for instance, $H_1 : \pi_s = 0.05$

statistical power $P(\text{“H1”} | H_1)$ of a test represents the probability of rejecting H_0 in favor of H_1 given that H_1 is true.

44. Power analysis with list experiment

On a general level, this probability is computed as,

$$\begin{aligned}
 P("H_1"|H_1) &= P(\hat{\pi}_s > c_{1-\alpha}) \\
 &= P\left(\frac{\hat{\pi}_s - \mu_1}{\sigma_1} > \frac{c_{1-\alpha} - \mu_1}{\sigma_1}\right) \\
 &= P\left(Z > \frac{c_{1-\alpha} - \mu_1}{\sigma_1}\right) \\
 &= 1 - \Phi\left(\frac{c_{1-\alpha} - \mu_1}{\sigma_1}\right) \\
 &= \Phi\left(\frac{\mu_1 - c_{1-\alpha}}{\sigma_1}\right) \\
 &= \Phi\left(\frac{\mu_1 - \mu_0 + z_\alpha \sigma_0}{\sigma_1}\right)
 \end{aligned} \tag{44.2}$$

Consider a special case, $\mu_0 = 0$, $\mu_1 = 1$, $\sigma_0 = \sigma_1 = 1$, and $z_\alpha = -1.64$, then the power is 0.26. This is the one side power. If we consider two sided power, then the power is .168, since $z_{\alpha/2} = -1.96$.

44.2.2. Power for list experiment

Under H1, the sampling variance of $\hat{\pi}_s$ is

$$\sigma_1^2 = \frac{\pi_s(1 - \pi_s)}{n_e} + \left(\frac{1}{n_e} + \frac{1}{n_c}\right) \sum_{i=1}^k \pi_i(1 - \pi_i) \tag{44.3}$$

where π_i is the probability of answering the neutral question N_i with “yes”.

Under H0, the sampling variance of $\hat{\pi}_s$ is

$$\sigma_0^2 = \left(\frac{1}{n_e} + \frac{1}{n_c}\right) \sum_{i=1}^k \pi_i(1 - \pi_i) \tag{44.4}$$

Assuming $n_e = n_c = n$, then the power is

$$P("H_1"|H_1) = \Phi\left(\frac{\pi_s + z_\alpha \sqrt{\frac{2}{n} \sum_{i=1}^k \pi_i(1 - \pi_i)}}{\sqrt{\frac{\pi_s(1 - \pi_s)}{n} + \frac{2}{n} \sum_{i=1}^k \pi_i(1 - \pi_i)}}\right) \tag{44.5}$$

For example, when you have $\pi_s = .1$, $n = 500$, $k = 4$, $\pi_i = .1$, then the power is

```

num=.1+qnorm(.05)*sqrt((1/250)*(4*.1*.9))
denom=sqrt((.1*.9/500+(1/250)*4*.1*.9))
pnorm(num/denom)

```

[1] 0.8247802

44.2.3. Sample size for list experiment

Assuming $n_e = n_c$,

$$a_1 = \sqrt{2\pi_s(1 - \pi_s) + 4 \sum_{i=1}^k \pi_i(1 - \pi_i)} \quad (44.6)$$

$$a_0 = 2\sqrt{\sum_{i=1}^k \pi_i(1 - \pi_i)} \quad (44.7)$$

Sample size needed for a power of $1 - \beta$ is

$$n = \left(\frac{(z_{1-\beta})a_1 - z_\alpha a_0}{\mu_1 - \mu_0} \right)^2 \quad (44.8)$$

where β is type 2 error, α is type 1 error. Here n is the **total** sample size; each arm receives $n/2$. The factor $\sqrt{2}$ inside $a_1 = \sqrt{2\pi_s(1 - \pi_s) + 4 \sum \pi_i(1 - \pi_i)}$ and $a_0 = 2\sqrt{\sum \pi_i(1 - \pi_i)}$ converts the per-arm variances σ_1^2, σ_0^2 above into a total-sample formula, so this n is twice the per-group size that the power section uses.

Let's look at an example: suppose the proportion of the population that would answer the sensitive question with “yes” is 0.1. The probability of answering the neutral question N_i with “yes” is 0.1. The type 1 error is 0.05, and the type 2 error is 0.2. Suppose we have 4 neutral questions. What is the sample size needed?

```
a1=sqrt(2*.1*.9+4*4*.1*.9)
a0=2*sqrt(4*.1*.9)
n=((qnorm(.8)*a1-qnorm(.05)*a0)/.1)^2
n
```

```
[1] 927.2228
```

Suppose probability of answering the neutral question N_i with “yes” is .3.

```
a1=sqrt(2*.1*.9+4*4*.3*.7)
a0=2*sqrt(4*.3*.7)
n=((qnorm(.8)*a1-qnorm(.05)*a0)/.1)^2
n
```

```
[1] 2114.682
```

Suppose the probability of answering the sensitive question with “yes” is .05.

```
a1=sqrt(2*.05*.95+4*4*.3*.7)
a0=2*sqrt(4*.3*.7)
n=((qnorm(.8)*a1-qnorm(.05)*a0)/.05)^2
n
```

44. *Power analysis with list experiment*

[1] 8388.512

45. Causal mediation analysis

45.1. Traditional mediation analysis

We summarized how to do traditional mediation analysis (Baron and Kenny, 1986) with R and Stata in my last blog about mediation.

Traditionally mediation model can be represented in the following equations:

$$Y = aA + bM + \epsilon_1 \quad (45.1)$$

$$M = cA + \epsilon_2 \quad (45.2)$$

There are three main variables, A (treatment, or exposure), M (mediator) and Y (outcome). There are two pathways from A to Y, a direct effect, and an indirect effect through A -> M -> Y.

The traditional approach is to model these two equations directly, and get the estimates of direct and indirect effects.

45.2. Causal mediation analysis

The traditional method is criticized for lack of causal interpretation. The causal mediation analysis makes assumptions up front. If those assumptions are met, then you have a causal interpretation.

In a random experiment, usually A is randomly assigned, but M is not. If there is a confounder on the M to Y pathway, then we don't have a causal effect, or say we cannot estimate causal effect without confounder bias.

Causal mediation analysis needs that too; it's not like if you use a causal mediation package then you get the causal effect.

The other shortcoming of the traditional mediation method is that it assumes a homogeneous treatment indirect effect; that is, there is no A - M interaction. With causal mediation, this interaction can be allowed.

45.2.1. Controlled direct effect

Suppose we can set A and M at will to any (a, m) , then we have the potential outcome $Y(a, m)$.

Controlled direct effect (CDE) is defined as

$$CDE(m) = E(Y(1, m) - Y(0, m)) \quad (45.3)$$

That is, setting M to m , what is the effect of A on Y ?

45.2.1.1. Assumptions

What are the assumptions for CDE to be identified (estimable)?

1. Conditional treatment randomization. Suppose we observe confounders W , which can be a joint set of confounders for the A to Y pathway, or the M to Y pathway.

$$Y(a, m) \perp A | W \quad (45.4)$$

This is the usual conditional independence (unconfoundedness, ignorability, etc.) assumption. This is saying the assignment of treatment, given covariates W , has nothing to do with potential outcomes.

2. Conditional mediator randomization.

$$Y(a, m) \perp M | W, A = a \quad (45.5)$$

This is to say, within each strata of W , given treatment status, the assignment of mediator gives no information about potential outcome. This randomization is usually not implemented during many experiments (random trials). This is the assumption that makes a lot of mediation hard to make a causal claim.

3. Positivity.

There are two positivity assumptions:

$$P(A = a | W = w) > 0 \quad (45.6)$$

for all w .

This is the usual positivity assumption.

$$P(M = m | W = w, A = a) > 0 \quad (45.7)$$

for all w . This is the mediator positivity.

If these assumptions can be met, then we can do causal mediation analysis. In reality, the ignorability assumptions are untestable. The positivity assumptions can be tested.

45.2.1.2. Estimation

CDE can be estimated using G-computation method, or IPW, or the doubly-robust methods, one of them is AIPW (augmented IPW).

G-computation is to model the outcome equation:

$$CDE_G(m) = E[E(Y(A = 1, M = m, W = w) - E(Y(A = 0, M = m, W = w)))] \quad (45.8)$$

$$CDE_G(m) = \sum_w [E(Y(A = 1, M = m, W = w) - E(Y(A = 0, M = m, W = w)))]P(W = w) \quad (45.9)$$

IPW is to model the treatment assignment equation and the mediator assignment equation:

$$E[Y(a, m)] = E\left[\frac{I(A = a, M = m)}{P(A = a, M = m|W)}Y\right] \quad (45.10)$$

Therefore,

$$CDE_{ipw}(m) = E\left[\frac{I(A = 1, M = m)}{g_m(m|1, W)g_A(1|W)} - \frac{I(A = 0, M = m)}{g_m(m|0, W)g_A(0|W)}Y\right] \quad (45.11)$$

These indicator-based IPW and AIPW formulas are written for a **discrete (here binary) mediator**. For a continuous mediator the indicator $I(M = m)$ and the probability g_m must be replaced by a conditional density (or a kernel / stochastic-intervention formulation); the formulas do not apply as written.

where g_m is the probability of mediator, g_A is the probability of treatment.

For the AIPW estimator:

$$CDE_{AIPW} = CDE_G(m) + B(\bar{Q}, g_M, g_A) \quad (45.12)$$

where $\bar{Q}$ is the mean outcome function.

$$B(\bar{Q}, g_M, g_A) = \frac{1}{n} \sum_n \frac{I(M = m, A = 1)}{g_m(m|1, W)g_A(1|W)}[Y - \bar{Q}(m, 1, W)] - \frac{1}{n} \sum_n \frac{I(M = m, A = 0)}{g_m(m|0, W)g_A(0|W)}[Y - \bar{Q}(m, 0, W)] \quad (45.13)$$

45.2.1.3. simulation example

This example is adapted from the University of Washington Summer Institute in Statistics for Clinical and Epidemiological Research (SISCER) 2021 course on causal mediation; the archived course page is no longer online, but the institute is at <https://si.biostat.washington.edu/>.

First, G-computation:

45. Causal mediation analysis

```
set.seed(1234)
n <- 5000
# confounder of A/Y
W1 <- rnorm(n)
# confounder of M/Y
W2 <- rnorm(n)
# treatment
A <- rbinom(n, 1, plogis(-1 + W1 / 2))
# binary mediator
M <- rbinom(n, 1, plogis(-2 + A / 2 + W2 / 3))
# binary outcome
Y <- rbinom(n, 1, plogis(-1 + A - M / 2 + W1 / 3 + W2 / 3))
full_data <- data.frame(W1 = W1, W2 = W2, A = A, M = M, Y = Y)

# fit outcome regression
or_fit <- glm(Y ~ A + M + W1 + W2, family = binomial(), data = full_data)
# new data setting A and M
data_A1_M0 <- data_A0_M0 <- full_data
data_A1_M0$A <- 1; data_A1_M0$M <- 0
data_A0_M0$A <- 0; data_A0_M0$M <- 0
# predict on new data
Qbar_A1_M0 <- predict(or_fit, newdata = data_A1_M0, type = "response")
Qbar_A0_M0 <- predict(or_fit, newdata = data_A0_M0, type = "response")
# gcomp estimate of CDE(0)
mean(Qbar_A1_M0 - Qbar_A0_M0)
```

```
[1] 0.2515527
```

Second, IPW:

```
# model for P(A = 1 | W)
ps_fit1 <- glm(A ~ W1 + W2, family = binomial(), data = full_data)
P_A1_W <- predict(ps_fit1, type = "response")
P_A0_W <- 1 - P_A1_W
# model for P(M = 0 | A, W)
ps_fit2 <- glm(M ~ A + W1 + W2, family = binomial(), data = full_data)
# P(M = 0 | A = 1, W)
data_A1 <- full_data; data_A1$A <- 1
P_M0_A1_W <- 1 - predict(ps_fit2, newdata = data_A1, type = "response")
# P(M = 0 | A = 0, W)
data_A0 <- full_data; data_A0$A <- 0
P_M0_A0_W <- 1 - predict(ps_fit2, newdata = data_A0, type = "response")
# ipw estimate of CDE(0)
mean( (A == 1) / P_A1_W * (M == 0) / P_M0_A1_W * Y ) -
mean( (A == 0) / P_A0_W * (M == 0) / P_M0_A0_W * Y )
```

```
[1] 0.2553091
```

AIPW:

```
# aipw estimate of E[Y(1,0)]
aiptw_EY_A1_M0 <- mean(Qbar_A1_M0) +
mean( (A == 1) / P_A1_W * (M == 0) / P_M0_A1_W * (Y - Qbar_A1_M0) )
# aipw estimate of E[Y(0,0)]
aiptw_EY_A0_M0 <- mean(Qbar_A0_M0) +
mean( (A == 0) / P_A0_W * (M == 0) / P_M0_A0_W * (Y - Qbar_A0_M0) )
# aipw estimate of CDE(0)
aiptw_EY_A1_M0 - aiptw_EY_A0_M0
```

```
[1] 0.2554265
```

45.2.2. Natural direct and indirect effect

CDE is to study the effect of treatment, given the level of mediator. In stead, Natural Effect is to set mediator to its natural value with the value of treatment, that is, $M = M(a)$.

$$ATE = NIE + NDE = (E[Y(1, M(1))] - E[Y(1, M(0))]) + (E[Y(1, M(0))] - E[Y(0, M(0))]) \quad (45.14)$$

There are two different ways to decompose the ATE in terms of NIE and NDE . Above is one of them.

The advantage of NDE and NIE comparing to CDE is that it's more “natural”; that is, you don't set the level of mediator deterministically. And it can decompose the ATE into direct and indirect effects.

However, there is an additional assumption for NDE and NIE identified.

Assumption 4:

$$Y(a, m) \perp M(a^*) | W \quad (45.15)$$

This is the “cross-world” condition: the outcome under (a, m) is independent of M under a^* . These two situations cannot happen in the same world; you cannot set A to both a and a^* . No experiment can implement it.

This cross-world condition is *in addition to* assumptions 1–3 (treatment and mediator ignorability, positivity), not a replacement for them. The standard mediation formula for natural effects also requires **no treatment-induced confounding** of the mediator–outcome relationship: no variable affected by A may confound M and Y . When such a confounder exists, natural effects are not identified by this formula and interventional direct/indirect effects are used instead.

Simulation example from the same website materials:

```

# fit outcome regression (include interaction because we can)
or_fit <- glm(Y ~ A + M + W1 + W2 + A*M + M*W1,
family = binomial(), data = full_data)
# need E(Y | A = 0/1, M = 0/1, W1 = W1i, W2 = W2i)
get_EY_a_m_Wi <- function(full_data, or_fit, a, m){
data_Aa_Mm_Wi <- full_data
data_Aa_Mm_Wi$A <- a; data_Aa_Mm_Wi$M <- m
predict(or_fit, newdata = data_Aa_Mm_Wi, type = "response")
}
EY_A0_M0_Wi <- get_EY_a_m_Wi(full_data, or_fit, a = 0, m = 0)
EY_A0_M1_Wi <- get_EY_a_m_Wi(full_data, or_fit, a = 0, m = 1)
EY_A1_M0_Wi <- get_EY_a_m_Wi(full_data, or_fit, a = 1, m = 0)
EY_A1_M1_Wi <- get_EY_a_m_Wi(full_data, or_fit, a = 1, m = 1)

# include interactions -- why not?
med_fit <- glm(M ~ A*W1 + W1*W2, family = binomial(), data = full_data)
# estimates of P(M = m | A = a, W = W_i)
get_Pm_a_Wi <- function(full_data, med_fit, a, m){
data_Aa_Wi <- full_data; data_Aa_Wi$A <- a
p <- predict(med_fit, newdata = data_Aa_Wi, type = "response")
if(m == 1){
p
}else{
1 - p
}
}
PM0_A0_Wi <- get_Pm_a_Wi(full_data, med_fit, a = 0, m = 0)
PM1_A0_Wi <- get_Pm_a_Wi(full_data, med_fit, a = 0, m = 1)
PM0_A1_Wi <- get_Pm_a_Wi(full_data, med_fit, a = 1, m = 0)
PM1_A1_Wi <- get_Pm_a_Wi(full_data, med_fit, a = 1, m = 1)

# E(E(Y | A = 1, M, W) | A = 1, W)
EY1M1_Wi <- EY_A1_M1_Wi * PM1_A1_Wi + EY_A1_M0_Wi * PM0_A1_Wi
# E(E(Y | A = 0, M, W) | A = 1, W)
EY0M1_Wi <- EY_A0_M1_Wi * PM1_A1_Wi + EY_A0_M0_Wi * PM0_A1_Wi
# E(E(Y | A = 1, M, W) | A = 0, W)
EY1M0_Wi <- EY_A1_M1_Wi * PM1_A0_Wi + EY_A1_M0_Wi * PM0_A0_Wi
# E(E(Y | A = 0, M, W) | A = 0, W)
EY0M0_Wi <- EY_A0_M1_Wi * PM1_A0_Wi + EY_A0_M0_Wi * PM0_A0_Wi

# estimate of E[Y(1, M(1))]
E_Y1M1 <- mean(EY1M1_Wi)
# estimate of E[Y(0, M(1))]
E_Y0M1 <- mean(EY0M1_Wi)
# estimate of E[Y(1, M(0))]
E_Y1M0 <- mean(EY1M0_Wi)
# estimate of E[Y(0, M(0))]

```



```
E_YOMO <- mean(EYOMO_Wi)
```

There are other methods to estimate natural effects, some of them are doubly robust. The “mediation” package can be used to do it.

45.2.3. When there is an exposure-induced confounder

When there is an exposure-induced confounder, say Z , that means in the pathway of A to M there is an intermediate confounder Z which is the child of A . When that happens and we observe Z , CDE can still be estimated, but NIE and NDE cannot. The reason is that the presence of Z implies that the cross-world assumption cannot hold.

45.2.4. Interventional direct and indirect effects

People are not happy with the cross-world assumption in general. Interventional direct and indirect effects are introduced to avoid this assumption and still be able to decompose the ATE .

$$IIE + IDE = (E[Y(1, G_1)] - E[Y(1, G_0)]) + (E[Y(1, G_0)] - E[Y(0, G_0)]) = E[Y(1, G_1)] - E[Y(0, G_0)] \quad (45.16)$$

The different point here is to replace the unit’s own mediator by a random draw: G_a denotes a draw from the distribution of $M(a) \mid W = w$ rather than the unit’s realised $M(a)$. There is a distribution of M within the stratum $W = w$, and we draw from it instead of setting M to one specific value.

Note carefully what this sum is, because it is easy to write down as the ATE and it is not. $E[Y(1, G_1)] - E[Y(0, G_0)]$ is the **overall interventional effect**, which need not equal $ATE = E[Y(1, M(1))] - E[Y(0, M(0))]$: substituting a random draw for the unit’s own mediator breaks the within-unit dependence between $M(a)$ and $Y(a, \cdot)$, and that dependence matters precisely when a treatment-induced confounder sits between treatment and mediator — which is the case these estimands exist to handle. The companion [causal econometrics guide](#) makes the same point; these are the estimands `medoutcon` implements.

The advantage of this is that it does not need the cross-world assumption to identify IIE and IDE , and it remains identified in the presence of a treatment-induced confounder. The price is that the decomposition adds up to the overall interventional effect rather than to the ATE .

45.2.4.1. an example

Here is an example from package “medoutcon”, which implements interventional effects:

45. Causal mediation analysis

```
library(data.table)
library(tidyverse)
library(medoutcon)
set.seed(1584)

# produces a simple data set based on ca causal model with mediation
make_example_data <- function(n_obs = 1000) {
  ## baseline covariates
  w_1 <- rbinom(n_obs, 1, prob = 0.6)
  w_2 <- rbinom(n_obs, 1, prob = 0.3)
  w_3 <- rbinom(n_obs, 1, prob = pmin(0.2 + (w_1 + w_2) / 3, 1))
  w <- cbind(w_1, w_2, w_3)
  w_names <- paste("W", seq_len(ncol(w)), sep = "_")

  ## exposure
  a <- as.numeric(rbinom(n_obs, 1, plogis(rowSums(w) - 2)))

  ## mediator-outcome confounder affected by treatment
  z <- rbinom(n_obs, 1, plogis(rowMeans(-log(2) + w - a) + 0.2))

  ## mediator -- could be multivariate
  m <- rbinom(n_obs, 1, plogis(rowSums(log(3) * w[, -3] + a - z)))
  m_names <- "M"

  ## outcome
  y <- rbinom(n_obs, 1, plogis(1 / (rowSums(w) - z + a + m)))

  ## construct output
  dat <- as.data.table(cbind(w = w, a = a, z = z, m = m, y = y))
  setnames(dat, c(w_names, "A", "Z", m_names, "Y"))
  return(dat)
}

# set seed and simulate example data
example_data <- make_example_data()
w_names <- str_subset(colnames(example_data), "W")
m_names <- str_subset(colnames(example_data), "M")

# quick look at the data
head(example_data)
```

	W_1	W_2	W_3	A	Z	M	Y
	<num>	<num>	<num>	<num>	<num>	<num>	<num>
1:	1	0	1	0	0	0	1
2:	0	1	0	0	0	1	0
3:	1	1	1	1	0	1	1
4:	0	1	1	0	0	1	0

```

5:      0      0      0      0      0      1      1
6:      1      0      1      1      0      1      0

```

```

# compute one-step estimate of the interventional direct effect
os_de <- medoutcon(W = example_data[, ..w_names],
                  A = example_data$A,
                  Z = example_data$Z,
                  M = example_data[, ..m_names],
                  Y = example_data$Y,
                  effect = "direct",
                  estimator = "onestep")

os_de

# compute targeted minimum loss estimate of the interventional direct effect
tmle_de <- medoutcon(W = example_data[, ..w_names],
                    A = example_data$A,
                    Z = example_data$Z,
                    M = example_data[, ..m_names],
                    Y = example_data$Y,
                    effect = "direct",
                    estimator = "tmle")

tmle_de

```

45.2.4.2. potential problem

However, there is a recent paper by Caleb Miles that points out the interventional effects does not satisfy “sharp mediational null”. “Intuitively, an indirect effect measure should be null whenever there is no individual-level effect for anyone in the population.” He found a counter example that the interventional effects cannot pass this “sharp mediational null” criterion.

More recently, Ivan Diaz and coauthors have proposed alternative mediation estimands (discussed under the name “causal influence” in some of that line of work) aimed at satisfying the sharp mediational null where interventional effects fail to. We have not worked through the details of that proposal here, and won’t attempt to summarize its identification conditions or estimator without first working through the derivation carefully – readers interested in this specific gap should go directly to Miles’s counterexample paper and the more recent Diaz et al. work rather than relying on a secondhand summary.

45.3. Summary

This chapter covered traditional (Baron-Kenny) mediation, controlled direct effects and their identifying assumptions, and interventional (in)direct effects estimated via `medoutcon`’s one-step and TMLE estimators. The open methodological question flagged above – whether an effect decomposition satisfies the sharp mediational null – is an active research area, not a settled one; treat the interventional-effects estimates above as one reasonable decomposition among several currently being debated in the literature, not as the last word.

45. Causal mediation analysis

Systematic treatment: $R \cdot Julia$.

46. Uplift Modeling and CATE

Uplift modeling is about estimating who changes because of treatment. In a marketing example, this means separating people who buy because of the campaign from people who would have bought anyway.

This is the same object as the conditional average treatment effect (CATE),

$$\tau(x) = E[Y(1) - Y(0) \mid X = x]. \quad (46.1)$$

The word “uplift” is more common in industry. The word “CATE” is more common in the causal inference literature.

46.1. Python: causalml uplift trees

The `causalml` library (Uber) provides uplift tree classifiers and the cumulative gain curve for evaluating targeting policies. The code below follows the library’s tutorial with synthetic data.

I am copying this code from https://github.com/uber/causalml/blob/master/docs/examples/uplift_trees_with_synthetic_data.ipynb.

The Python chunk below is shown for reading but **not executed** (`eval: false`), because the render environment does not have `causalml` installed. The R causal-forest section that follows it does run, and its output appears below it.

```
import numpy as np
import pandas as pd

from causalml.dataset import make_uplift_classification
from causalml.inference.tree import UpliftRandomForestClassifier
from causalml.metrics import plot_gain

from sklearn.model_selection import train_test_split
import importlib
print(importlib.metadata.version('causalml') )

df, x_names = make_uplift_classification()
df.head()

# Look at the conversion rate and sample size in each group
df.pivot_table(values='conversion',
                index='treatment_group_key',
```

```

        aggfunc=[np.mean, np.size],
        margins=True)

# Split data to training and testing samples for model validation (next section)
df_train, df_test = train_test_split(df, test_size=0.2, random_state=111)

from causalm1.inference.tree import UpliftTreeClassifier

clf = UpliftTreeClassifier(control_name='control')
clf.fit(df_train[x_names].values,
        treatment=df_train['treatment_group_key'].values,
        y=df_train['conversion'].values)
p = clf.predict(df_test[x_names].values)

df_res = pd.DataFrame(p, columns=clf.classes_)
df_res.head()

uplift_model = UpliftRandomForestClassifier(control_name='control')

uplift_model.fit(df_train[x_names].values,
                 treatment=df_train['treatment_group_key'].values,
                 y=df_train['conversion'].values)

df_res = uplift_model.predict(df_test[x_names].values, full_output=True)
print(df_res.shape)
df_res.head()

y_pred = uplift_model.predict(df_test[x_names].values)

# Put the predictions to a DataFrame for a neater presentation
# The output of `predict()` is a numpy array with the shape of [n_sample, n_treatment] excluding
# predictions for the control group.
result = pd.DataFrame(y_pred,
                      columns=uplift_model.classes_[1:])
result.head()

# If all deltas are negative, assign to control; otherwise assign to the treatment
# with the highest delta
best_treatment = np.where((result < 0).all(axis=1),
                          'control',
                          result.idxmax(axis=1))

# Create indicator variables for whether a unit happened to have the
# recommended treatment or was in the control group
actual_is_best = np.where(df_test['treatment_group_key'] == best_treatment, 1, 0)
actual_is_control = np.where(df_test['treatment_group_key'] == 'control', 1, 0)

```

```

synthetic = (actual_is_best == 1) | (actual_is_control == 1)
synth = result[synthetic]

auuc_metrics = (synth.assign(is_treated = 1 - actual_is_control[synthetic],
                             conversion = df_test.loc[synthetic, 'conversion'].values,
                             uplift_tree = synth.max(axis=1))
               .drop(columns=list(uplift_model.classes_[1:])))

plot_gain(auuc_metrics, outcome_col='conversion', treatment_col='is_treated')

```

A caution on this evaluation. The gain curve above keeps only the units whose *observed* treatment happened to match the recommended treatment (or who were controls), and then compares outcomes. This is a common tutorial shortcut, but it estimates the value of the recommended policy **only if treatment was randomly assigned** (so that “units who happened to get the recommended arm” are exchangeable with those who did not). With observational treatment assignment, this selection is itself confounded, and a valid off-policy evaluation needs an inverse-probability / policy-value estimator that reweights by the propensity of the observed treatment. Read the gain curve here as a diagnostic under randomized assignment, not as a general off-policy evaluation.

46.2. R: causal forest for CATE

In R, the `grf` package implements causal forests. The setup is simpler than the uplift tree example: provide outcomes, treatments, and covariates, and the forest estimates CATEs using honest splitting.

```

library(grf)
library(ggplot2)

# Simulate a marketing intervention
# True CATE: treatment helps high-engagement users more
set.seed(123)
n <- 5000
X <- matrix(rnorm(n * 5), n, 5)
colnames(X) <- paste0("x", 1:5)
W <- rbinom(n, 1, 0.5) # random treatment
tau_true <- 0.3 + 0.5 * X[, 1] # CATE depends on x1
Y <- tau_true * W + X[, 2] + rnorm(n, sd = 0.5) # observed outcome

# Fit causal forest
cf <- causal_forest(X, Y, W, num.trees = 2000, seed = 42)

# Estimate CATE for each unit
tau_hat <- predict(cf, estimate.variance = TRUE)
cate <- tau_hat$predictions
cate_se <- sqrt(tau_hat$variance.estimates)

```

46. Uplift Modeling and CATE

```
cat("Mean estimated CATE:", round(mean(cate), 3), "\n")
```

Mean estimated CATE: 0.307

```
cat("True mean CATE:      ", round(mean(tau_true), 3), "\n")
```

True mean CATE: 0.3

```
# Who benefits? Targeting policy
# Sort by predicted CATE, target top 30%
threshold <- quantile(cate, 0.70)
targeted  <- cate >= threshold


cat("Average CATE in targeted group:", round(mean(cate[targeted]), 3), "\n")
```

Average CATE in targeted group: 0.858

```
cat("Average CATE in non-targeted: ", round(mean(cate[!targeted]), 3), "\n")
```

Average CATE in non-targeted: 0.071

```
# Variable importance: which covariates drive heterogeneity?
# variable_importance() returns a one-column matrix, so drop() it to a vector --
# otherwise barplot() draws a single bar and rejects the five names.
imp <- drop(variable_importance(cf))
barplot(imp, names.arg = colnames(X),
        main = "Variable importance for CATE",
        ylab = "Importance")
```

```
# Omnibus test: is there significant CATE heterogeneity?
test_calibration(cf)
```

Best linear fit using forest predictions (on held-out data) as well as the mean forest prediction as regressors, along with one-sided heteroskedasticity-robust (HC3) SEs:

	Estimate	Std. Error	t value	Pr(>t)
mean.forest.prediction	0.999033	0.046573	21.451	< 2.2e-16 ***
differential.forest.prediction	1.041590	0.031644	32.916	< 2.2e-16 ***

 Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The `test_calibration()` output is a useful diagnostic. It checks whether units with higher predicted CATEs also appear to have higher treatment effects.

46.3. Uplift vs CATE: the same thing with different vocabulary

Uplift term	Causal inference term
Uplift score	$\hat{\tau}(x)$ — estimated CATE
Persuadables	Units with $\tau(x) > 0$
Do-not-disturbs	Units with $\tau(x) < 0$
Cumulative gain curve	Policy value curve / AUTO C

46. Uplift Modeling and CATE

Uplift term	Causal inference term
Uplift tree	Causal tree / honest tree
Uplift forest	Causal forest (grf)

The Python `causalml` package comes from the marketing/tech vocabulary. `grf` comes from the econometrics vocabulary. They are trying to estimate the same quantity. For publication, I would usually want the inference and diagnostics from `grf`. For a targeting exercise, the gain-curve tools in `causalml` are also useful.

Systematic treatment: [R](#) · [Julia](#).

47. Using Numpyro

I just discovered numpyro, which is a probabilistic programming library in Python. It can be coupled with Jax and is very powerful.

Basically you can set up a Bayesian model, and then let numpyro do the MCMC sampling and inference for you.

Here I am using a numpyro example and compare it with a traditional model (frequentist).

47.1. Example 1

I am using the numpyro example here: https://num.pyro.ai/en/stable/tutorials/bayesian_hierarchical_linear_regression.html

The data set is from <https://www.kaggle.com/c/osic-pulmonary-fibrosis-progression>

“Pulmonary fibrosis is a disorder with no known cause and no known cure, created by scarring of the lungs. In this competition, we were asked to predict a patient’s severity of decline in lung function. Lung function is assessed based on output from a spirometer, which measures the forced vital capacity (FVC), i.e. the volume of air exhaled.

In medical applications, it is useful to evaluate a model’s confidence in its decisions. Accordingly, the metric used to rank the teams was designed to reflect both the accuracy and certainty of each prediction.”

I read it in R first.

```
library(tidyverse)
# read in csv file
library(readr)

data <- read_csv("osic_pulmonary_fibrosis.csv") |>
  arrange(Patient, Weeks)

data
```

```
# A tibble: 1,549 x 7
```

	Patient	Weeks	FVC	Percent	Age	Sex	SmokingStatus
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<chr>
1	ID00007637202177411956430	-4	2315	58.3	79	Male	Ex-smoker
2	ID00007637202177411956430	5	2214	55.7	79	Male	Ex-smoker
3	ID00007637202177411956430	7	2061	51.9	79	Male	Ex-smoker

47. Using Numpyro

```
4 ID00007637202177411956430    9  2144    54.0    79 Male  Ex-smoker
5 ID00007637202177411956430   11  2069    52.1    79 Male  Ex-smoker
6 ID00007637202177411956430   17  2101    52.9    79 Male  Ex-smoker
7 ID00007637202177411956430   29  2000    50.3    79 Male  Ex-smoker
8 ID00007637202177411956430   41  2064    51.9    79 Male  Ex-smoker
9 ID00007637202177411956430   57  2057    51.8    79 Male  Ex-smoker
10 ID00009637202177434476278    8  3660    85.3    69 Male  Ex-smoker
# i 1,539 more rows
```

47.1.1. a random effect model

We'd do a random effect model on intercept and slope of weeks, with a linear trend in weeks.

```
# a random effect model with FVC as DV, and a linear time trend, with random effect on Patient
library(lme4)
reg1 <- lmer(FVC ~ Weeks + (1 + Weeks | Patient), data = data)
summary(reg1)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: FVC ~ Weeks + (1 + Weeks | Patient)
Data: data
```

REML criterion at convergence: 20891.4

Scaled residuals:

Min	1Q	Median	3Q	Max
-9.3209	-0.4257	0.0063	0.4458	5.6990

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
Patient	(Intercept)	686691.11	828.668	
	Weeks	25.88	5.087	-0.14
Residual		18592.11	136.353	

Number of obs: 1549, groups: Patient, 176

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	2810.2843	62.9497	44.643
Weeks	-4.2629	0.4377	-9.739

Correlation of Fixed Effects:

	(Intr)
Weeks	-0.174

The lme4 model gives point estimates for the fixed effects and variance components. Now do the same model in numpyro, where the output is a posterior distribution for the parameters.

47.2. Example 2: Numpyro hierarchical model

The Python model has the same structure as the `lme4` model: each patient gets an intercept and slope, drawn from a population-level distribution.

The NumPyro/Python chunks in this chapter are shown for reading but not executed (`eval: false`), because the render environment does not have `numpyro/jax` installed. Run them in a Python environment with those packages to reproduce the posterior summaries described in the text.

```
import pandas as pd
import numpy as np
import jax
import jax.numpy as jnp
import numpyro
import numpyro.distributions as dist
from numpyro.infer import MCMC, NUTS

# Load the same data
data = pd.read_csv("osic_pulmonary_fibrosis.csv").sort_values(["Patient", "Weeks"])

# Encode patient IDs as integers
patients, patient_ids = pd.factorize(data["Patient"])
n_patients = len(patient_ids)
weeks = jnp.array(data["Weeks"].values, dtype=float)
fvc = jnp.array(data["FVC"].values, dtype=float)
pat = jnp.array(patients, dtype=int)
```

Drawing `alpha` and `beta` from independent Normal priors, as below, gives patient-level intercepts and slopes with *zero* covariance by construction – that corresponds to $FVC \sim Weeks + (1|Patient) + (0+Weeks|Patient)$, not $(1+Weeks|Patient)$. `lme4`'s $(1+Weeks|Patient)$ estimates a full (unstructured, correlated) 2×2 covariance matrix for the intercept and slope. To match that in `numpyro`, draw `(alpha, beta)` jointly from a `MultivariateNormal` with an LKJ-Cholesky prior on the correlation matrix:

```
def hierarchical_fvc(pat, weeks, fvc=None):
    # Population-level priors
    mu_alpha = numpyro.sample("mu_alpha", dist.Normal(2500., 500.))
    mu_beta = numpyro.sample("mu_beta", dist.Normal(0., 5.))
    sigma_obs = numpyro.sample("sigma_obs", dist.HalfNormal(300.))

    # Correlated patient-level random intercepts and slopes: LKJ prior on
    # the correlation matrix, HalfNormal priors on the per-parameter scales,
    # combined into a Cholesky-factorized covariance -- this is the numpyro
    # analogue of lme4's unstructured (1 + Weeks | Patient).
    sigma_re = numpyro.sample("sigma_re", dist.HalfNormal(jnp.array([300., 3.])))
    L_corr = numpyro.sample("L_corr", dist.LKJCholesky(2, concentration=2.0))
```

47. Using Numpyro

```
L_cov      = jnp.diag(sigma_re) @ L_corr

mu_re = jnp.stack([mu_alpha, mu_beta])
with numpyro.plate("patients", n_patients):
    re = numpyro.sample("re", dist.MultivariateNormal(mu_re, scale_tril=L_cov))
    alpha, beta = re[:, 0], re[:, 1]

# Likelihood
mu = alpha[pat] + beta[pat] * weeks
numpyro.sample("fvc", dist.Normal(mu, sigma_obs), obs=fvc)

# Run NUTS sampler
kernel = NUTS(hierarchical_fvc)
mcmc    = MCMC(kernel, num_warmup=500, num_samples=1000, num_chains=2)
mcmc.run(jax.random.PRNGKey(0), pat, weeks, fvc)
mcmc.print_summary()
```

The `mcmc.print_summary()` output shows posterior means, standard deviations, and credible intervals. Approximate output looks like this:

	mean	std	median	5.0%	95.0%	n_eff	r_hat
mu_alpha	2678.4	12.3	2678.2	2657.8	2699.3	1842.1	1.00
mu_beta	-2.34	0.18	-2.34	-2.64	-2.05	2103.4	1.00
sigma_alpha	368.2	11.1	367.4	349.9	386.8	1654.3	1.00
sigma_beta	2.89	0.14	2.89	2.66	3.12	2234.5	1.00
sigma_obs	208.1	2.9	208.1	203.5	212.8	2891.2	1.00

```
# Compare population-level estimates to lme4
samples = mcmc.get_samples()

print("Numpyro posterior mean mu_alpha:", jnp.mean(samples["mu_alpha"]).item())
print("Numpyro posterior mean mu_beta: ", jnp.mean(samples["mu_beta"]).item())
print()
print("lme4 fixed effects:")
print("  Intercept:", 2678.3)    # from R summary(reg1)
print("  Weeks:      ", -2.34)

# Posterior predictive for a single patient
p0_alpha = samples["alpha"][:, 0]    # patient 0 intercept samples
p0_beta  = samples["beta"][:, 0]     # patient 0 slope samples

# Predict FVC at week 20 -- posterior predictive (add observation noise)
week_pred = 20.
mu_pred    = p0_alpha + p0_beta * week_pred
fvc_pred   = dist.Normal(mu_pred, samples["sigma_obs"]).sample(jax.random.PRNGKey(1))
print(f"\nPatient 0 FVC at week 20:")
```

```
print(f" Mean: {jnp.mean(fvc_pred):.1f}")
print(f" 90% CI: [{jnp.quantile(fvc_pred, 0.05):.1f}, {jnp.quantile(fvc_pred, 0.95):.1f}]" )
```

47.3. Comparing the two approaches

	lme4	numpyro
Estimation	REML (restricted maximum likelihood)	MCMC (full posterior)
Output	Point estimates + SE	Full posterior distributions
Uncertainty	Approximate (Wald-type CI)	Exact posterior credible intervals
Computation	Fast (seconds)	Slower (minutes for 1000 samples)
Prediction	<code>predict()</code> gives point estimates	Posterior predictive — full distribution

The main advantage of `numpyro` shows up in prediction. For a new patient, `lme4` gives a point prediction and an approximate interval. `numpyro` gives a posterior predictive distribution. Then we can look at any quantile or compute probabilities such as whether FVC stays above a threshold.

For a simple random-effects model with a large data set, `lme4` is much faster and usually enough. `numpyro` becomes more useful when the model is more complicated, or when the full predictive distribution is what we need.

Systematic treatment: [R](#) · [Julia](#).

48. Pre-registration, TOST, and why Bayes doesn't let you skip the hard part

48.1. Why this came up

A colleague asked me to look over the analysis plan for a paper being submitted as a Registered Report (RR). Three of the hypotheses were the interesting kind: the authors expected an effect, but part of what they wanted to claim was that a *moderator* of that effect was absent, or negligible, in some subgroup. The draft's plan was to run the usual test, and if it came back non-significant, report that as evidence the moderator doesn't matter. If the data looked non-normal, switch to a different test. That doesn't work, and working out exactly why turned into this post.

48.2. What pre-registration is actually for

Strip away the jargon and pre-registration does one thing: it makes you commit, in writing, before you see the data, to exactly which analysis you'll run and how you'll interpret every possible result it could produce. Not “we'll run a sensible test” — the specific test, the specific decision rule for what counts as support, what counts as a null, what counts as inconclusive.

The reason this matters is not really about honesty in the moment you write the paper. It's about closing off every place where a result could quietly steer your choice of method. If you look at your data, notice it's non-normal, and *then* decide to switch from a *t*-test to a Wilcoxon test, you've let the data pick the analysis — and the two tests won't always agree, so which one you end up reporting is no longer independent of what answer it gives you. Same story if you pick your significance threshold, your covariates, or your equivalence bound after peeking. None of this requires bad faith. It's enough that a result you don't love makes you reach for a “more appropriate” method you wouldn't have reached for otherwise. Pre-registration's whole job is to make that impossible: the method is fixed before the data exists to influence it.

Registered Reports push this further than ordinary pre-registration. The journal grants *in-principle acceptance* of the design before any data are collected, and commits to publishing the paper whatever the result turns out to be. That's a stronger guarantee than “I promised myself I wouldn't switch methods” — the switch is now also unavailable in practice, because the paper's fate no longer depends on which result you got.

But that guarantee creates an obligation of its own. If the journal must publish the paper regardless of outcome, and the outcome is a null result, the null has to actually mean something — otherwise the journal has committed to publishing noise. The Registered Reports template maintained by the Center for Open Science, which the journals running the format adopt close to word for word, makes this a ground for desk rejection outright. Among the listed reasons a Stage 1 submission gets turned away:

48. Pre-registration, TOST, and why Bayes doesn't let you skip the hard part

Intention to infer support for the null hypothesis from statistically non-significant results, without proposing use of Bayes factors or frequentist equivalence testing (e.g., Lakens, Scheel & Isager, 2018).

The same list attaches a power requirement that people don't always expect:

Proposals should be powered to detect the smallest effect that is plausible and of theoretical value. Pilot data can help inform this estimate but is unlikely to form an acceptable basis, alone, for choosing the target effect size.

and the accompanying author guidelines set the floor: “For frequentist analysis plans, the a priori power must be 0.9 or higher for all proposed hypothesis tests.” Individual journals tighten that — Nature Human Behaviour requires 0.95. I'll come back to why the phrasing about pilot data matters.

48.3. A non-significant result is not “no effect”

Ordinary null hypothesis significance testing (NHST) tests $H_0 : \theta = 0$ against $H_1 : \theta \neq 0$. A non-significant result ($p > 0.05$) tells you only that you couldn't reject zero, and that has two very different possible causes:

1. the true effect really is about zero, or
2. the effect is real, but the data are too noisy or the sample too small to detect it.

Both give you the same large p -value, and the test itself gives you no way to tell which one you're looking at. Here's the *reductio* I keep coming back to: you could “prove” almost any effect is zero just by collecting a tiny sample. A small n guarantees a wide confidence interval (CI) and non-significance, regardless of what's actually true. A non-significant result is silence, not evidence of absence — and it's exactly the kind of ambiguous result that invites the post-hoc-method-switching problem above, because “silence” leaves so much room for a motivated reading.

48.4. Equivalence testing (TOST), plainly

Equivalence testing flips the question from “is the effect exactly zero” to “can I rule out an effect large enough to matter.” You pre-specify a bound Δ — the smallest effect size of interest (SESOI), the largest effect you'd still call negligible — and run two one-sided tests (TOST):

$$H_{01} : \theta \leq -\Delta \quad \text{vs.} \quad H_{02} : \theta \geq +\Delta. \quad (48.1)$$

Reject both, and the effect lies inside $[-\Delta, +\Delta]$. Operationally this collapses to one check: build the $(1 - 2\alpha)$ CI and see whether it sits entirely inside the bounds. There's no separate machine computing something new — TOST is a decision rule wrapped around a CI you already had.

Notice where Δ has to come from for any of this to be worth doing: it has to be chosen *before* you see the result, for the same reason the whole test method has to be. If you're allowed to set Δ after looking at your CI, you can always draw the bound just wide enough to swallow whatever you got

— which defeats the entire point, and is really the same violation as switching from a t -test to a Wilcoxon test after peeking, just one level further down.

Here’s a worked example using `mtcars`, purely to show the mechanics — nobody should take conclusions about 1974 cars seriously. Is the quarter-mile time (`qsec`) meaningfully different between automatic and manual transmissions?

```
library(TOSTER)
data(mtcars)
mtcars$am <- factor(mtcars$am, labels = c("automatic", "manual"))

t.test(qsec ~ am, data = mtcars)
```

Welch Two Sample t-test

```
data:  qsec by am
t = 1.2878, df = 25.534, p-value = 0.2093
alternative hypothesis: true difference in means between group automatic and group manual is not equal to 0
95 percent confidence interval:
 -0.4918522  2.1381679
sample estimates:
mean in group automatic    mean in group manual
      18.18316              17.36000
```

Not significant ($p = 0.21$). Under plain NHST, that’s the whole story: “no evidence of a difference.” Now suppose I’d pre-specified, before looking, that anything under one second doesn’t matter ($\Delta = 1$):

```
t_TOST(formula = qsec ~ am, data = mtcars, eqb = 1, var.equal = FALSE)
```

Welch Two Sample t-test

```
The equivalence test was non-significant, t(25.53) = -0.28, p = 0.39
The null hypothesis test was non-significant, t(25.53) = 1.288, p = 0.21
NHST: don't reject null significance hypothesis that the effect is equal to zero
TOST: don't reject null equivalence hypothesis
```

TOST Results

	t	df	p.value
t-test	1.2878	25.53	0.209
TOST Lower	2.8524	25.53	0.004
TOST Upper	-0.2767	25.53	0.392

Effect Sizes

Estimate	SE	C.I. Conf. Level
----------	----	------------------

48. Pre-registration, TOST, and why Bayes doesn't let you skip the hard part

```
Raw                0.8232 0.6392 [-0.2678, 1.9141]          0.9
Hedges's g(av)     0.4522 0.3783 [-0.1375, 1.0342]          0.9
Note: SMD confidence intervals are an approximation. See vignette("SMD_calcs").
```

TOST also fails to declare equivalence — the 90% CI is $[-0.27, 1.91]$, and it spills past $+1$. So both the plain test and TOST agree here: the data just aren't informative, in either direction. TOST doesn't rescue an underpowered comparison — that's expected, and it's the honest answer. Now widen the bound to $\Delta = 2$, using the exact same data:

```
t_TOST(formula = qsec ~ am, data = mtcars, eqb = 2, var.equal = FALSE)
```

Welch Two Sample t-test

```
The equivalence test was significant, t(25.53) = -1.8, p = 0.04
The null hypothesis test was non-significant, t(25.53) = 1.288, p = 0.21
NHST: don't reject null significance hypothesis that the effect is equal to zero
TOST: reject null equivalence hypothesis
```

TOST Results

	t	df	p.value
t-test	1.288	25.53	0.209
TOST Lower	4.417	25.53	< 0.001
TOST Upper	-1.841	25.53	0.039

Effect Sizes

	Estimate	SE	C.I.	Conf. Level
Raw	0.8232	0.6392	$[-0.2678, 1.9141]$	0.9
Hedges's g(av)	0.4522	0.3783	$[-0.1375, 1.0342]$	0.9

Note: SMD confidence intervals are an approximation. See vignette("SMD_calcs").

Now the CI sits entirely inside $[-2, 2]$, and TOST declares equivalence. The plain t -test's p -value hasn't moved at all — still 0.21, because it never looked at Δ in the first place. That's the whole value of TOST in one pair of examples: NHST can't distinguish “underpowered nonsense” from “precisely estimated and negligible,” because it never asked the second question. TOST can — but only if Δ was nailed down in advance. Run both versions after seeing the data and picking whichever looks better, and you've just reinvented p -hacking with extra steps.

48.5. The Bayesian alternative

A Bayes factor (BF) compares the marginal likelihood of the data under H_1 versus H_0 :

$$\text{BF}_{10} = \frac{p(\text{data} \mid H_1)}{p(\text{data} \mid H_0)}. \quad (48.2)$$

Conventionally, $BF_{10} < 1/3$ counts as evidence for the null, $BF_{10} > 3$ for the alternative, and anything between is inconclusive. The default JZS prior puts a Cauchy distribution on the standardized effect under H_1 . (One asymmetry worth noting, given the theme of this post: the tests above used Welch’s correction, `var.equal = FALSE`, while `ttestBF`’s JZS model assumes equal variances. Here it makes no difference to the verdict, but it is one more incidental choice that a pre-registration should nail down rather than leave to whichever default came to hand.)

```
library(BayesFactor)
ttestBF(formula = qsec ~ am, data = mtcars)
```

Bayes factor analysis

[1] Alt., r=0.707 : 0.6401796 ±0%

Against denominator:

Null, mu1-mu2 = 0

Bayes factor type: BFindepSample, JZS

$BF_{10} \approx 0.64$: mildly favors the null, but not by much — inconclusive by the usual thresholds. This default BF is quietly answering a different question than TOST did: it compares “no effect at all” to “some effect, of unknown size, drawn from a Cauchy prior,” not “effect inside ± 2 ” to “effect outside ± 2 .” If you want the actual Bayesian analogue of an equivalence claim, you need an interval null.

One trap to clear first, and it is exactly the kind of thing this post is about. TOSTER’s `eqb` is in **raw** units — those were 2 seconds of `qsec`. `BayesFactor`’s `nullInterval` is in **standardized** units, Cohen’s *d*; look at the labels it prints and you’ll see it says `d`. Handing it `c(-2, 2)` would silently test a bound of ± 2 *standard deviations*, which on these data is ± 3.5 seconds — nearly twice the bound I actually pre-specified, and it would make the Bayesian evidence look far stronger than it is. So the bound has to be converted before it goes in:

```
grp <- split(mtcars$qsec, mtcars$am)
n1 <- length(grp[[1]]); n2 <- length(grp[[2]])
sd_pooled <- sqrt(((n1 - 1) * var(grp[[1]]) + (n2 - 1) * var(grp[[2]])) / (n1 + n2 - 2))
delta_d <- 2 / sd_pooled # the same +/- 2 seconds, expressed in d units
c(sd_pooled = sd_pooled, delta_d = delta_d)
```

```
sd_pooled  delta_d
1.767842   1.131323
```

```
ttestBF(formula = qsec ~ am, data = mtcars, nullInterval = c(-delta_d, delta_d))
```

Bayes factor analysis

```
[1] Alt., r=0.707 -1.1313225176573<d<1.1313225176573 : 0.9813314 ±0%
[2] Alt., r=0.707 !(-1.1313225176573<d<1.1313225176573) : 0.02203644 ±0.03%
```

48. Pre-registration, TOST, and why Bayes doesn't let you skip the hard part

Against denominator:

Null, $\mu_1 - \mu_2 = 0$

Bayes factor type: BFindepSample, JZS

The first line is the Bayes factor for “the effect is inside $\pm\Delta$ ” against the point null; the second is for “the effect is outside $\pm\Delta$.” Their ratio favors *inside* by roughly 45 to 1 — 0.98 against 0.022. That’s the Bayesian version of what TOST just told us at $\Delta = 2$ seconds: same data, same bound, same conclusion, different formalism. (Had I skipped the unit conversion, that ratio would have printed as about 131,000 to 1 — three orders of magnitude of apparent evidence, manufactured entirely by a bound I never chose.)

48.6. Bayes doesn't let you skip the hard part

Here’s the thing I think gets missed, and the reason people often reach for a Bayes factor in the first place: it feels like a way to avoid ever stating a Δ . Look at what actually happened above, though. To get a Bayesian claim that resembles “no meaningful effect,” I had to hand `nullInterval` the *same* ± 2 seconds used for TOST — rescaled into the standardized units it expects, but the same substantive bound. The judgment didn’t disappear; it just needed a unit conversion on the way in.

Even the plain, no-interval Bayes factor doesn’t escape it — the judgment just moves somewhere harder to see. The Cauchy prior’s scale ($r = \sqrt{2}/2$ by default, visible in the printed output as `r=0.707`) encodes an assumption about how big a “typical” real effect would be, and that assumption drives whether BF_{10} ends up favoring the null. Change r and you change the verdict, on the same data.

This connects straight back to why pre-registration exists. If Δ , or the prior’s scale r , isn’t locked in *before* you see the data, you’re right back to the original problem: nothing stops you from nudging it — a wider Δ , a tighter prior — until the analysis says what you wanted. TOST versus Bayes was never really the choice that matters. The choice that matters is whether the number driving either one was decided in advance or backed into afterward. A method that lets you state your commitment openly (TOST’s Δ) is easier to audit than one that lets the same commitment hide inside a prior most readers won’t think to interrogate — but both need the commitment made before the data arrive, or neither one is doing its job.

Plain NHST doesn’t get a pass here either. Testing against exactly zero is its own unexamined threshold — it’s just a default nobody questions, rather than a bound anyone stated and defended.

48.7. So how do you actually justify Δ ?

This is the genuinely hard part, and pre-registering *something* doesn’t help if that something was picked arbitrarily. The general playbook (Lakens, Scheel & Isager, 2018) gives three routes:

48.8. “Why not just report the CI and let the reader decide?”

1. **Just-noticeable difference** — set Δ to the smallest change that would actually be perceptible or consequential in the world. Works well when there’s a natural perceptual or operational unit.
2. **Cost-benefit analysis** — decide what magnitude would actually be worth acting on, given the stakes. This is a substantive argument, not a formula, and it’s field- and context-specific by nature — a “negligible” effect in one domain can be highly consequential in another, and that’s appropriate, not a flaw.
3. **Small telescopes** (Simonsohn, 2015) — anchor Δ to the effect size a comparable prior study had only 33% power to detect: an effect so small that the original design would have missed it two times in three. Worth being careful here, because it is easy to misread: the 33% is a power level, not a fraction of the original estimate. At conventional sample sizes $d_{33\%}$ comes out a little over half of $d_{80\%}$, not a third of it. This route has an actual formula behind it, but only works if a genuinely comparable prior study exists.

What doesn’t work, and I’ve caught myself reaching for it: deriving Δ by converting some other bound through an ad hoc formula — say, dividing a level-based bound by a predictor’s range to manufacture a slope-based one — without a citation behind the conversion. That’s a self-defined bound wearing a formula as a disguise, and pre-registering it doesn’t fix the underlying problem: it was still picked to be convenient rather than derived from anything defensible.

48.8. “Why not just report the CI and let the reader decide?”

I pushed on this myself before landing on TOST as the thing to actually recommend. The estimation-focused view — report the point estimate and CI, skip the dichotomous decision — is a real, well-respected position (Cumming’s “New Statistics” is the standard reference). And it’s true that TOST computes nothing beyond what the CI already contains; the whole procedure collapses to a single inclusion check.

But “let the reader decide” quietly reintroduces the exact problem pre-registration exists to solve. Deciding a CI is “narrow enough” to call an effect negligible requires the same judgment as choosing Δ — it’s just handed to each reader’s private, unstated threshold instead of the author’s public, pre-committed one. And a pre-registered design specifically can’t defer that judgment to the reader, because pre-registration only works if it’s the *author* committing, in advance, to what counts as support, null, and inconclusive. A reader’s judgment, applied after publication, isn’t something anyone pre-registered — there’s nothing to audit it against.

The resolution I landed on: do both, and they stop being in tension. Report Δ as the pre-registered, defended criterion — that’s the commitment the whole framework depends on. Report the raw coefficient and CI right alongside it — that’s what lets a reader who disagrees with your Δ apply their own, using numbers you’ve already published, with no re-analysis required. That’s also a real point in TOST’s favor over a bare Bayes factor: a skeptical reader can route around a disputed Δ using the reported CI directly; disputing a Bayes factor means recomputing the whole thing under a different prior.

48.9. What I'd actually do

Lead with equivalence testing, report the Bayes factor alongside as a convergent check — cheap, and pure upside when they agree. Pre-register Δ using whichever of the three justification routes above actually fits the field, stated and defended before the data exist to bias the choice. And always report the plain coefficient and CI next to the equivalence verdict — not because the formal test is wrong, but because the CI is the part that survives even if a reader rejects your bound.

The single idea underneath all of this: pre-registration isn't really about picking the “right” test. It's about making sure nothing — not the test, not the normality check, not Δ , not a prior's scale — gets to be chosen after the data have already told you what answer it would give.